



Oracle Database 11g SQL

Master SQL and PL/SQL in the Oracle Database



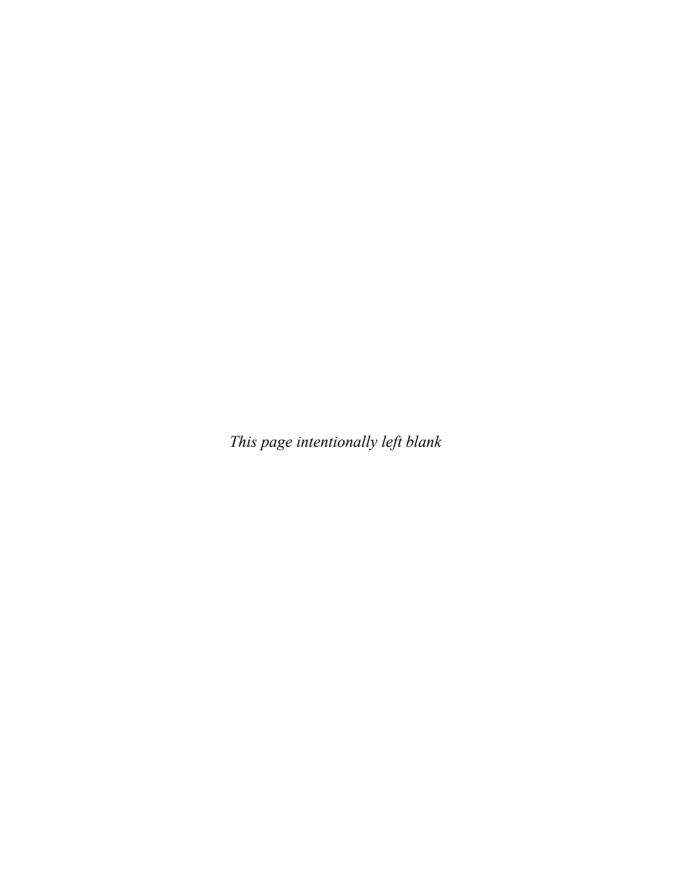
Jason Price

Oracle Certified Professional Database Administrator and Application Developer





Oracle Database 11*g* SQL







Oracle Database 11*g* SQL

Jason Price



New York Chicago San Francisco Lisbon London Madrid Mexico City Milan New Delhi San Juan Seoul Singapore Sydney Toronto Copyright © 2008 by The McGraw-Hill Companies, Inc. All rights reserved. Manufactured in the United States of America. Except as permitted under the United States Copyright Act of 1976, no part of this publication may be reproduced or distributed in any form or by any means, or stored in a database or retrieval system, without the prior written permission of the publisher.

0-07-159613-5

The material in this eBook also appears in the print version of this title: 0-07-149850-8.

All trademarks are trademarks of their respective owners. Rather than put a trademark symbol after every occurrence of a trademarked name, we use names in an editorial fashion only, and to the benefit of the trademark owner, with no intention of infringement of the trademark. Where such designations appear in this book, they have been printed with initial caps.

McGraw-Hill eBooks are available at special quantity discounts to use as premiums and sales promotions, or for use in corporate training programs. For more information, please contact George Hoare, Special Sales, at george_hoare@mcgraw-hill.com or (212) 904-4069.

TERMS OF USE

This is a copyrighted work and The McGraw-Hill Companies, Inc. ("McGraw-Hill") and its licensors reserve all rights in and to the work. Use of this work is subject to these terms. Except as permitted under the Copyright Act of 1976 and the right to store and retrieve one copy of the work, you may not decompile, disassemble, reverse engineer, reproduce, modify, create derivative works based upon, transmit, distribute, disseminate, sell, publish or sublicense the work or any part of it without McGraw-Hill's prior consent. You may use the work for your own noncommercial and personal use; any other use of the work is strictly prohibited. Your right to use the work may be terminated if you fail to comply with these terms.

THE WORK IS PROVIDED "AS IS." McGRAW-HILL AND ITS LICENSORS MAKE NO GUARANTEES OR WARRANTIES AS TO THE ACCURACY, ADEQUACY OR COMPLETENESS OF OR RESULTS TO BE OBTAINED FROM USING THE WORK, INCLUDING ANY INFORMATION THAT CAN BE ACCESSED THROUGH THE WORK VIA HYPERLINK OR OTHERWISE, AND EXPRESSLY DISCLAIM ANY WARRANTY, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. McGraw-Hill and its licensors do not warrant or guarantee that the functions contained in the work will meet your requirements or that its operation will be uninterrupted or error free. Neither McGraw-Hill nor its licensors shall be liable to you or anyone else for any inaccuracy, error or omission, regardless of cause, in the work or for any damages resulting therefrom. McGraw-Hill has no responsibility for the content of any information accessed through the work. Under no circumstances shall McGraw-Hill and/or its licensors be liable for any indirect, incidental, special, punitive, consequential or similar damages that result from the use of or inability to use the work, even if any of them has been advised of the possibility of such damages. This limitation of liability shall apply to any claim or cause whatsoever whether such claim or cause arises in contract, tort or otherwise.





Want to learn more?

We hope you enjoy this McGraw-Hill eBook! If

you'd like more information about this book, its author, or related books and websites, please click here.

This book is dedicated to my family. Even though you're far away, you are still in my heart.

About the Author

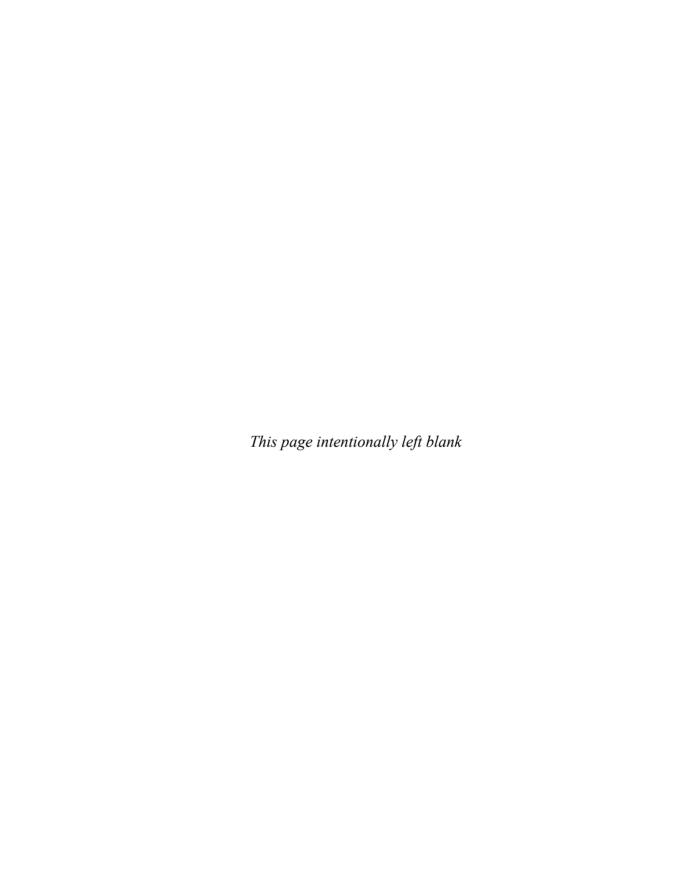
Jason Price is a freelance consultant and former product manager of Oracle Corporation. He has contributed to many of Oracle's products, including the database, the application server, and several of the CRM applications. Jason is an Oracle Certified Database Administrator and Application Developer, and has more than 15 years of experience in the software industry. Jason has written many books on Oracle, Java, and .NET. Jason holds a Bachelor of Science degree (with honors) in physics from the University of Bristol, England.

About the Technical Editor

Scott Mikolaitis is an applications architect at Oracle Corporation and has worked at Oracle for over ten years. He performs prototyping and standards development for the SOA technology in Oracle Fusion. Scott also enjoys working with web services in Java as well as Jabber for human and system interaction patterns. He spends his spare time on DIY home improvement and gas-fueled RC cars.

Contents at a Glance

1	Introduction
2	Retrieving Information from Database Tables
3	Using SQL*Plus
4	Using Simple Functions
5	Storing and Processing Dates and Times
6	Subqueries
7	Advanced Queries
8	Changing Table Contents
9	Users, Privileges, and Roles
10	Creating Tables, Sequences, Indexes, and Views
11	Introducing PL/SQL Programming
12	Database Objects
13	Collections
14	Large Objects
15	Running SQL Using Java
16	SQL Tuning
17	XML and the Oracle Database
Α	Oracle Data Types
	Index 639





	Introduction	XXV
1	Introduction	1
•	What Is a Relational Database?	
	Introducing the Structured Query Language (SQL)	
	Using SQL*Plus	
	Starting SQL*Plus	4
	Starting SQL*Plus from the Command Line	
	Performing a SELECT Statement Using SQL*Plus	6
	SQL Developer	7
	Creating the Store Schema	
	Running the SQL*Plus Script to Create the Store Schema	
	Data Definition Language (DDL) Statements Used to Create the Store Schema .	11
	Adding, Modifying, and Removing Rows	20
	Adding a Row to a Table	20
	Modifying an Existing Row in a Table	22
	Removing a Row from a Table	22
	The BINARY_FLOAT and BINARY_DOUBLE Types	23
	Benefits of BINARY_FLOAT and BINARY_DOUBLE	23
	Using BINARY_FLOAT and BINARY_DOUBLE in a Table	24
	Special Values	24
	Quitting SQL*Plus	25
	Introducing Oracle PL/SQL	25
	Summary	
	Summary	20
2	Retrieving Information from Database Tables	27
	Performing Single Table SELECT Statements	28
	Retrieving All Columns from a Table	29
	Specifying Rows to Retrieve Using the WHERE Clause	
	Row Identifiers	30
	Row Numbers	30

\mathbf{X} Oracle Database 11g SQL

	Performing Arithmetic	31
	Performing Date Arithmetic	
	Using Columns in Arithmetic	
	Using Column Aliases	34
	Combining Column Output Using Concatenation	35
	Null Values	35
	Displaying Distinct Rows	37
	Comparing Values	37
	Using the SQL Operators	39
	Using the LIKE Operator	40
	Using the IN Operator	41
	Using the BETWEEN Operator	42
	Using the Logical Operators	43
	Operator Precedence	44
	Sorting Rows Using the ORDER BY Clause	44
	Performing SELECT Statements That Use Two Tables	45
	Using Table Aliases	47
	Cartesian Products	48
	Performing SELECT Statements That Use More than Two Tables	49
	Join Conditions and Join Types	50
	Non-equijoins	50
	Outer Joins	51
	Self Joins	55
	Performing Joins Using the SQL/92 Syntax	56
	Performing Inner Joins on Two Tables Using SQL/92	56
	Simplifying Joins with the USING Keyword	57
	Performing Inner Joins on More than Two Tables Using SQL/92	58
	Performing Inner Joins on Multiple Columns Using SQL/92	58
	Performing Outer Joins Using SQL/92	59
	Performing Self Joins Using SQL/92	60
	Performing Cross Joins Using SQL/92	61
	Summary	61
3	Using SQL*Plus	63
•	Viewing the Structure of a Table	
	Editing SQL Statements	65
	Saving, Retrieving, and Running Files	66
	Formatting Columns	
	Setting the Page Size	
	Setting the Line Size	
	Clearing Column Formatting	
	Using Variables	74
	Temporary Variables	74
	Defined Variables	77
	Creating Simple Reports	79
	Using Temporary Variables in a Script	80
	Using Defined Variables in a Script	80
	Passing a Value to a Variable in a Script	81
	Adding a Header and Footer	82
	Computing Subtotals	83

	Getting Help from SQL*Plus	85
	Automatically Generating SQL Statements	
	Disconnecting from the Database and Exiting SQL*Plus	86
	Summary	
4	Using Simple Functions	89
•	Using Single-Row Functions	
	Character Functions	
	Numeric Functions	
	Conversion Functions	
	Regular Expression Functions	
	Using Aggregate Functions	
	0 00 0	
	AVG()	
	COUNT()	
	MAX() and MIN()	
	STDDEV()	
	SUM()	
	VARIANCE()	
	Grouping Rows	
	Using the GROUP BY Clause to Group Rows	
	Incorrect Usage of Aggregate Function Calls	
	Using the HAVING Clause to Filter Groups of Rows	
	Using the WHERE and GROUP BY Clauses Together	
	Using the WHERE, GROUP BY, and HAVING Clauses Together	
	Summary	127
_	Staving and Discossing Dates and Times	120
5	Storing and Processing Dates and Times	129
	Converting Datetimes Using TO_CHAR() and TO_DATE()	
	Using TO_CHAR() to Convert a Datetime to a String	
	Using TO_DATE() to Convert a String to a Datetime	
	Setting the Default Date Format	
	How Oracle Interprets Two-Digit Years	
	Using the YY Format	
	Using the RR Format	
	Using Datetime Functions	
	ADD_MONTHS()	
	LAST_DAY()	
	MONTHS_BETWEEN()	
	NEXT_DAY()	
	ROUND()	
	SYSDATE	145
	TRUNC()	146
	Using Time Zones	146
	Time Zone Functions	147
	The Database Time Zone and Session Time Zone	147
	The Battabase time Botte and bession time Botte 11111111111111111111111111111111111	,
	Obtaining Time Zone Offsets	149
		149

Xii Oracle Database 11*g* SQL

	Using Timestamps	
	Using the Timestamp Types	150
	Timestamp Functions	154
	Using Time Intervals	
	Using the INTERVAL YEAR TO MONTH Type	160
	Using the INTERVAL DAY TO SECOND Type	162
	Time Interval Functions	
	Summary	166
,	Cultura de la companya della companya della companya de la companya de la companya della company	167
6	Subqueries	
	Types of Subqueries	
	Writing Single-Row Subqueries	
	Subqueries in a WHERE Clause	
	Using Other Single-Row Operators	
	Subqueries in a HAVING Clause	
	Subqueries in a FROM Clause (Inline Views)	
	Errors You Might Encounter	
	Writing Multiple-Row Subqueries	
	Using IN with a Multiple-Row Subquery	
	Using ANY with a Multiple-Row Subquery	
	Using ALL with a Multiple-Row Subquery	
	Writing Multiple-Column Subqueries	
	Writing Correlated Subqueries	
	A Correlated Subquery Example	
	Using EXISTS and NOT EXISTS with a Correlated Subquery	
	Writing Nested Subqueries	
	Writing UPDATE and DELETE Statements Containing Subqueries	
	Writing an UPDATE Statement Containing a Subquery	
	Writing a DELETE Statement Containing a Subquery	
	Summary	181
7	Advanced Queries	183
	Using the Set Operators	
	The Example Tables	
	Using the UNION ALL Operator	
	Using the UNION Operator	
	Using the INTERSECT Operator	188
	Using the MINUS Operator	
	Combining Set Operators	
	Using the TRANSLATE() Function	
	Using the DECODE() Function	
		193
	Using Simple CASE Expressions	193
	Using Searched CASE Expressions	194
	Hierarchical Queries	196
	The Example Data	196
	Using the CONNECT BY and START WITH Clauses	198
	Using the LEVEL Pseudo Column	198
	Formatting the Results from a Hierarchical Query	
	Starting at a Node Other than the Root	
	Using a Subquery in a START WITH Clause	

Traversing Upward Through the Tree	
Eliminating Nodes and Branches from a Hierarchical Query	
Including Other Conditions in a Hierarchical Query	
Using the Extended GROUP BY Clauses	
The Example Tables	203
Using the ROLLUP Clause	
Using the CUBE Clause	207
Using the GROUPING() Function	
Using the GROUPING SETS Clause	
Using the GROUPING_ID() Function	
Using a Column Multiple Times in a GROUP BY Clause	
Using the GROUP_ID() Function	215
Using the Analytic Functions	
The Example Table	
Using the Ranking Functions	
Using the Inverse Percentile Functions	
Using the Window Functions	225
Using the Reporting Functions	
Using the LAG() and LEAD() Functions	
Using the FIRST and LAST Functions	
Using the Linear Regression Functions	233
Using the Hypothetical Rank and Distribution Functions	
Using the MODEL Clause	
An Example of the MODEL Clause	
Using Positional and Symbolic Notation to Access Cells	
Accessing a Range of Cells Using BETWEEN and AND	
Accessing All Cells Using ANY and IS ANY	
Getting the Current Value of a Dimension Using CURRENTV()	
Accessing Cells Using a FOR Loop	240
Handling Null and Missing Values	
Updating Existing Cells	
Using the PIVOT and UNPIVOT Clauses	
A Simple Example of the PIVOT Clause	244
Pivoting on Multiple Columns	246
Using Multiple Aggregate Functions in a Pivot	247
Using the UNPIVOT Clause	
Summary	249
Changing Table Contents	251
Adding Rows Using the INSERT Statement	
Omitting the Column List	
Specifying a Null Value for a Column	
Including Single and Double Quotes in a Column Value	
Copying Rows from One Table to Another	
Modifying Rows Using the UPDATE Statement	
The RETURNING Clause	
Removing Rows Using the DELETE Statement	
Database Integrity	
Enforcement of Primary Key Constraints	
Enforcement of Foreign Key Constraints	257

8

xiv Oracle Database 11*g* SQL

	Using Default Values	258
	Merging Rows Using MERGE	259
	Database Transactions	
	Committing and Rolling Back a Transaction	262
	Starting and Ending a Transaction	
	Savepoints	
	ACID Transaction Properties	
	Concurrent Transactions	
	Transaction Locking	
	Transaction Isolation Levels	
	A SERIALIZABLE Transaction Example	
	Query Flashbacks	
	Granting the Privilege for Using Flashbacks	
	Time Query Flashbacks	
	System Change Number Query Flashbacks	
	Summary	
	·	
9	Users, Privileges, and Roles	
	Users	
	Creating a User	
	Changing a User's Password	
	Deleting a User	
	System Privileges	
	Granting System Privileges to a User	
	Checking System Privileges Granted to a User	
	Making Use of System Privileges	
	Revoking System Privileges from a User	
	Object Privileges	
	Granting Object Privileges to a User	
	Checking Object Privileges Made	
	Checking Object Privileges Received	
	Making Use of Object Privileges	
	Synonyms	
	Public Synonyms	
	Revoking Object Privileges	
	Roles	
	Creating Roles	
	Granting Privileges to Roles	
	Granting Roles to a User	
	Checking Roles Granted to a User	
	Checking System Privileges Granted to a Role	
	Checking Object Privileges Granted to a Role	
	Making Use of Privileges Granted to a Role	
	Default Roles	
	Revoking a Role	294
	Revoking Privileges from a Role	294
	Dropping a Role	
	Auditing	295
	Privileges Required to Perform Auditing	295
	Auditing Examples	295

	Audit Trail Views	297
	Summary	297
10	Creating Tables, Sequences, Indexes, and Views	299
	Tables	
	Creating a Table	300
	Getting Information on Tables	
	Getting Information on Columns in Tables	303
	Altering a Table	
	Renaming a Table	
	Adding a Comment to a Table	
	Truncating a Table	
	Dropping a Table	
	Sequences	
	Creating a Sequence	
	Retrieving Information on Sequences	
	Using a Sequence	
	Populating a Primary Key Using a Sequence	
	Modifying a Sequence	
	Dropping a Sequence	
	Indexes	
	Creating a B-tree Index	
	Creating a Function-Based Index	
	Retrieving Information on Indexes	
	Retrieving Information on the Indexes on a Column	
	Modifying an Index	
	Dropping an Index	
	Creating a Bitmap Index	
	Views	
	Creating and Using a View	
	Modifying a View	
	Dropping a View	
	Flashback Data Archives	
	Summary	337
	,	
11	Introducing PL/SQL Programming	
	Block Structure	
	Variables and Types	
	Conditional Logic	
	Loops	
	Simple Loops	
	WHILE Loops	
	FOR Loops	
	Cursors	345
	Step 1: Declare the Variables to Store the Column Values	346
	Step 2: Declare the Cursor	346
	Step 3: Open the Cursor	346
	Step 4: Fetch the Rows from the Cursor	347
	Step 5: Close the Cursor	347 348
	COMPLETE EXAMPLE, DECORICE CHESOL 201	348

XVi Oracle Database 11*g* SQL

	Cursors and FOR Loops	
	OPEN-FOR Statement	350
	Unconstrained Cursors	352
	Exceptions	353
	ZERO_DIVIDE Exception	
	DUP_VAL_ON_INDEX Exception	
	INVALID_NUMBER Exception	
	OTHERS Exception	
	Creating a Procedure	
	Calling a Procedure	
	Getting Information on Procedures	
	Dropping a Procedure	362
	Viewing Errors in a Procedure	362
	Functions	
	Creating a Function	
	Calling a Function	
	Getting Information on Functions	
	Dropping a Function	
	Packages	
	Creating a Package Specification	
	Creating a Package Body	
	Calling Functions and Procedures in a Package	
	Getting Information on Functions and Procedures in a Package	368
	Dropping a Package	368
	Triggers	
	When a Trigger Fires	
	Set Up for the Example Trigger	
	Creating a Trigger	
	Firing a Trigger	
	Getting Information on Triggers	
	Disabling and Enabling a Trigger	
	Dropping a Trigger	
	New Oracle Database 11 <i>g</i> PL/SQL Features	
	SIMPLE_INTEGER Type	375
	Sequences in PL/SQL	
	PL/SQL Native Machine Code Generation	
	Summary	
	,	
12	Database Objects	379
	Introducing Objects	380
	Creating Object Types	381
	Using DESCRIBE to Get Information on Object Types	
	Using Object Types in Database Tables	383
	Column Objects	383
	Object Tables	386
	Object Identifiers and Object References	390
	Comparing Object Values	392
	Using Objects in PL/SQL	394
	The get_products() Function	395
	The display_product() Procedure	396

		~ ~ -
	The insert_product() Procedure	
	The update_product_price() Procedure	
	The get_product() Function	
	The update_product() Procedure	
	The get_product_ref() Function	
	The delete_product() Procedure	
	The product_lifecycle() Procedure	
	The product_lifecycle2() Procedure	402
	Type Inheritance	
	Using a Subtype Object in Place of a Supertype Object	405
	SQL Examples	405
	PL/SQL Examples	406
	NOT SUBSTITUTABLE Objects	407
	Other Useful Object Functions	
	IS OF()	
	TREAT()	412
	SYS_TYPEID()	
	NOT INSTANTIABLE Object Types	
	User-Defined Constructors	
	Overriding Methods	
	Generalized Invocation	
	Summary	
	·	
13	Collections	427
	Introducing Collections	428
	Creating Collection Types	429
	Creating a Varray Type	429
	Creating a Nested Table Type	429
	Using a Collection Type to Define a Column in a Table	430
	Using a Varray Type to Define a Column in a Table	430
	Using a Nested Table Type to Define a Column in a Table	430
	Getting Information on Collections	431
	Getting Information on a Varray	
	Getting Information on a Nested Table	
	Populating a Collection with Elements	434
	Populating a Varray with Elements	
	Populating a Nested Table with Elements	
	Retrieving Elements from Collections	
	Retrieving Elements from a Varray	
	Retrieving Elements from a Nested Table	
	Using TABLE() to Treat a Collection as a Series of Rows	
	Using TABLE() with a Varray	437
	Using TABLE() with a Nested Table	
	Modifying Elements of Collections	
		430
	Madificing Flamoute of a Variance	
	Modifying Elements of a Varray	438
	Modifying Elements of a Nested Table	438 439
	Modifying Elements of a Nested Table	438 439 440
	Modifying Elements of a Nested Table Using a Map Method to Compare the Contents of Nested Tables Using CAST() to Convert Collections from One Type to Another	438 439 440 443
	Modifying Elements of a Nested Table	438 439 440 443 443

XVIII Oracle Database 11*g* SQL

	Using Collections in PL/SQL	444
	Manipulating a Varray	444
	Manipulating a Nested Table	
	PL/SQL Collection Methods	
	Multilevel Collections	458
	Oracle Database 10g Enhancements to Collections	
	Associative Arrays	
	Changing the Size of an Element Type	
	Increasing the Number of Elements in a Varray	
	Using Varrays in Temporary Tables	
	Using a Different Tablespace for a Nested Table's Storage Table	
	ANSI Support for Nested Tables	
	Summary	
14	Large Objects	
	Introducing Large Objects (LOBs)	
	The Example Files	
	Large Object Types	
	Creating Tables Containing Large Objects	
	Using Large Objects in SQL	
	Using CLOBs and BLOBs	
	Using BFILEs	
	Using Large Objects in PL/SQL	
	APPEND()	
	CLOSE()	
	COMPARE()	
	COPY()	
	CREATETEMPORARY()	
	ERASE()	
	FILECLOSE()	
	FILECLOSEALL()	
	FILEEXISTS()	
	FILEGETNAME()	
	FILEISOPEN()	
	FILEOPEN()	
	FREETEMPORARY()	
	GETCHUNKSIZE()	
	GET_STORAGE_LIMIT()	
	GETLENGTH()	
	INSTR()	
	ISOPEN()	
	ISTEMPORARY()	495
	LOADFROMFILE()	
	LOADBLOBFROMFILE()	
	LOADCLOBFROMFILE()	497
	OPEN()	498
	READ()	499
	SUBSTR()	
	TRIM()	501

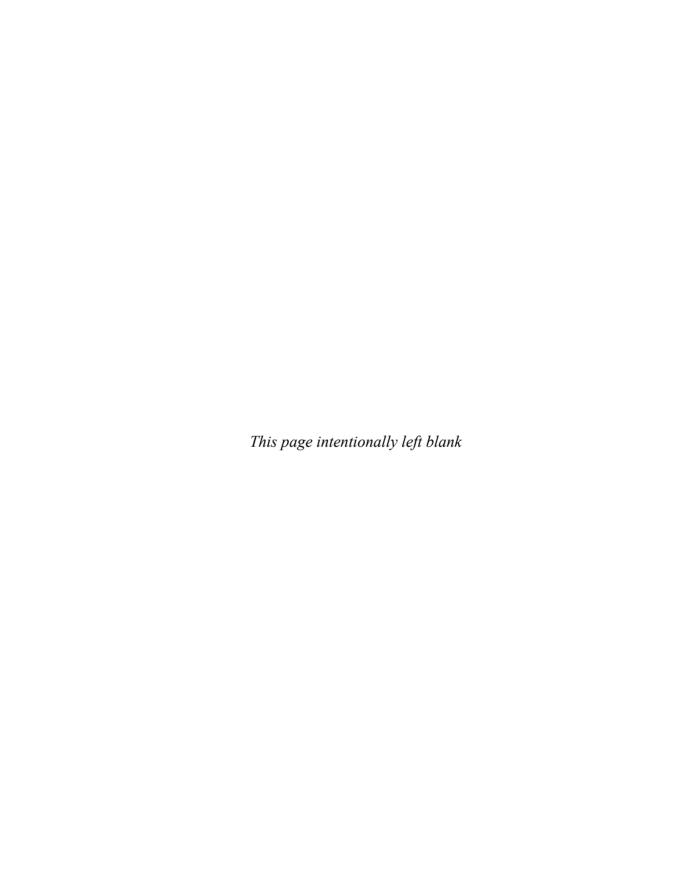
	WRITE()	502
	WRITEAPPEND()	503
	Example PL/SQL Procedures	503
	LONG and LONG RAW Types	521
	The Example Tables	
	Adding Data to LONG and LONG RAW Columns	
	Converting LONG and LONG RAW Columns to LOBs	
	Oracle Database 10g Enhancements to Large Objects	
	Implicit Conversion Between CLOB and NCLOB Objects	
	Use of the :new Attribute When Using LOBs in a Trigger	
	Oracle Database 11g Enhancements to Large Objects	
	Encrypting LOB Data	
	Compressing LOB Data	
	Removing Duplicate LOB Data	
	Summary	
	•	
15	Running SQL Using Java	
	Getting Started	
	Configuring Your Computer	
	Setting the ORACLE_HOME Environment Variable	
	Setting the JAVA_HOME Environment Variable	
	Setting the PATH Environment Variable	
	Setting the CLASSPATH Environment Variable	
	Setting the LD_LIBRARY_PATH Environment Variable	
	The Oracle JDBC Drivers	
	The Thin Driver	535
	The OCI Driver	536
	The Server-Side Internal Driver	536
	The Server-Side Thin Driver	536
	Importing the JDBC Packages	
	Registering the Oracle JDBC Drivers	537
	Opening a Database Connection	537
	Connecting to the Database Using getConnection()	537
	The Database URL	538
	Connecting to the Database Using an Oracle Data Source	539
	Creating a JDBC Statement Object	
	Retrieving Rows from the Database	543
	Step 1: Create and Populate a ResultSet Object	
	Step 2: Read the Column Values from the ResultSet Object	
	Step 3: Close the ResultSet Object	546
	Adding Rows to the Database	547
	Modifying Rows in the Database	548
	Deleting Rows from the Database	548
	Handling Numbers	549
	Handling Database Null Values	550
	Controlling Database Transactions	
	Performing Data Definition Language Statements	
	Handling Exceptions	
	Closing Your JDBC Objects	
	· · · · · · · · · · · · · · · · · · ·	

XX Oracle Database 11*g* SQL

	Example Program: BasicExample1.java	
	Compile BasicExample1	. 560
	Run BasicExample1	. 561
	Prepared SQL Statements	. 562
	Example Program: BasicExample2.java	. 565
	The Oracle JDBC Extensions	. 567
	The oracle.sql Package	. 568
	The oracle.jdbc Package	. 571
	Example Program: BasicExample3.java	. 575
	Summary	
	, , , , , , , , , , , , , , , , , , ,	
16	SQL Tuning	
	Introducing SQL Tuning	
	Use a WHERE Clause to Filter Rows	
	Use Table Joins Rather than Multiple Queries	
	Use Fully Qualified Column References When Performing Joins	
	Use CASE Expressions Rather than Multiple Queries	
	Add Indexes to Tables	
	Use WHERE Rather than HAVING	
	Use UNION ALL Rather than UNION	
	Use EXISTS Rather than IN	
	Use EXISTS Rather than DISTINCT	
	Use GROUPING SETS Rather than CUBE	. 588
	Use Bind Variables	
	Non-Identical SQL Statements	. 588
	Identical SQL Statements That Use Bind Variables	. 588
	Listing and Printing Bind Variables	
	Using a Bind Variable to Store a Value Returned by a PL/SQL Function	. 590
	Using a Bind Variable to Store Rows from a REFCURSOR	. 590
	Comparing the Cost of Performing Queries	. 591
	Examining Execution Plans	
	Comparing Execution Plans	
	Passing Hints to the Optimizer	. 598
	Additional Tuning Tools	
	Oracle Enterprise Manager Diagnostics Pack	
	Automatic Database Diagnostic Monitor	. 600
	Summary	
	YALLIN DELDEL	
17	XML and the Oracle Database	
	Introducing XML	
	Generating XML from Relational Data	
	XMLELEMENT()	
	XMLATTRIBUTES()	
	XMLFOREST()	
	XMLAGG()	
	XMLCOLATTVAL()	
	XMLCONCAT()	
	XMLPARSE()	
	XMLPI()	
	XMLCOMMENT()	. 614

	XMLSEQUENCE() XMLSERIALIZE() A PL/SQL Example That Writes XML Data to a File	616 616
	XMLQUERY()	
	Saving XML in the Database	
	The Example XML File	
	Creating the Example XML Schema	623
	Retrieving Information from the Example XML Schema	
	Updating Information in the Example XML Schema	630
	Summary	633
A	Oracle Data Types	635
	Oracle SQL Types	636
	Oracle PL/SQL Types	638

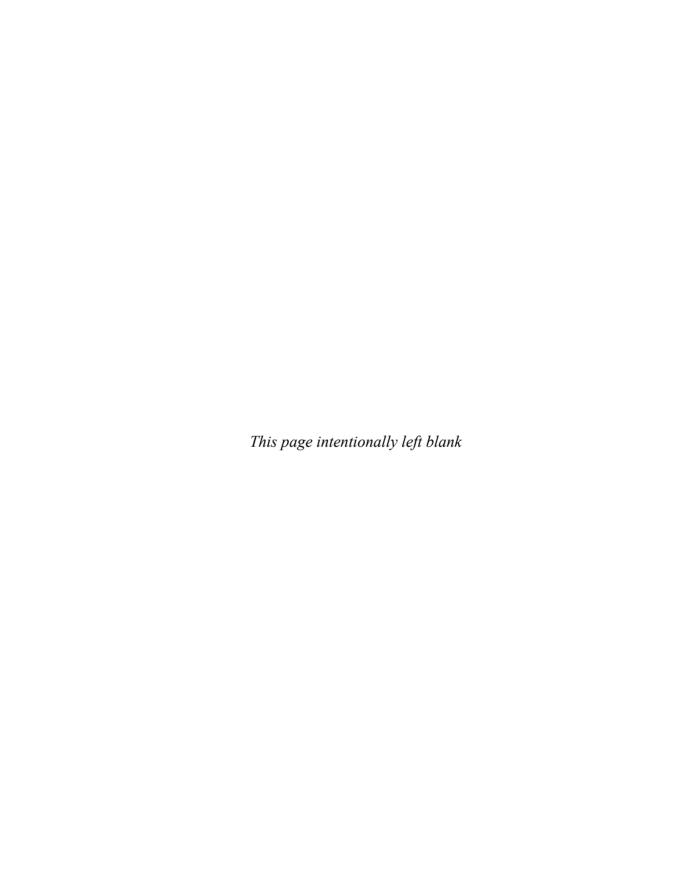
Contents **xxi**



Acknowledgments



hanks to the wonderful people at McGraw-Hill, including Lisa McClain, Mandy Canales, Carl Wikander, and Laura Stone. Thanks also to Scott Mikolaitis for his thorough technical review.



Introduction

oday's database management systems are accessed using a standard language known as *Structured Query Language*, or SQL. Among other things, SQL allows you to retrieve, add, update, and delete information in a database. In this book, you'll learn how to master SQL, and you'll find a wealth of practical examples. You can also get all the scripts and programs featured in this book

online (see the last section, "Retrieving the Examples," for details).

With this book, you will

- Master standard SQL, as well as the extensions developed by Oracle Corporation for use with the specific features of the Oracle database.
- Explore PL/SQL (Procedural Language/SQL), which is built on top of SQL and enables you to write programs that contain SQL statements.
- Use SQL*Plus to execute SQL statements, scripts, and reports; SQL*Plus is a tool that allows you to interact with the database.
- Execute queries, inserts, updates, and deletes against a database.
- Create database tables, sequences, indexes, views, and users.
- Perform transactions containing multiple SQL statements.
- Define database object types and create object tables to handle advanced data.
- Use large objects to handle multimedia files containing images, music, and movies.
- Perform complex calculations using analytic functions.
- Use all the very latest Oracle Database 11*g* features such as PIVOT and UNPIVOT, flashback archives, and much more.
- Implement high-performance tuning techniques to make your SQL statements really fly.

- Write Java programs to access an Oracle database using JDBC.
- Explore the XML capabilities of the Oracle database.

This book contains 17 chapters and one appendix.

Chapter 1: Introduction

In this chapter, you'll learn about relational databases, be introduced to SQL, see a few simple queries, use SQL*Plus and SQL Developer to execute queries, and briefly see PL/SQL.

Chapter 2: Retrieving Information from Database Tables

You'll explore how to retrieve information from one or more database tables using SELECT statements, use arithmetic expressions to perform calculations, filter rows using a WHERE clause, and sort the rows retrieved from a table.

Chapter 3: Using SQL*Plus

In this chapter, you'll use SQL*Plus to view a table's structure, edit a SQL statement, save and run scripts, format column output, define and use variables, and create reports.

Chapter 4: Using Simple Functions

In this chapter, you'll learn about some of the Oracle database's built-in functions. A function can accept input parameters and returns an output parameter. Functions allow you to perform tasks such as computing averages and square roots of numbers.

Chapter 5: Storing and Processing Dates and Times

You'll learn how the Oracle database processes and stores dates and times, collectively known as datetimes. You'll also learn about timestamps that allow you to store a specific date and time, and time intervals that allow you to store a length of time.

Chapter 6: Subqueries

You'll learn how to place a SELECT statement within an outer SQL statement. The inner SELECT statement is known as a subquery. You'll learn about the different types of subqueries and see how subqueries allow you to build up very complex statements from simple components.

Chapter 7: Advanced Queries

In this chapter, you'll learn how to perform queries containing advanced operators and functions such as: set operators that combine rows returned by multiple queries, the TRANSLATE() function to convert characters in one string to characters in another string, the DECODE() function to search a set of values for a certain value, the CASE expression to perform if-then-else logic, and the ROLLUP and CUBE clauses to return rows containing subtotals. You'll learn about the analytic functions that enable you to perform complex calculations such as finding the top-selling product type for each month, the top salespersons, and so on. You'll see how to perform queries against data that is organized into a hierarchy. You'll also explore the MODEL clause, which performs inter-row calculations. Finally, you'll see the new Oracle Database 11g PIVOT and UNPIVOT clauses, which are useful for seeing overall trends in large amounts of data.

Chapter 8: Changing Table Contents

You'll learn how to add, modify, and remove rows using the INSERT, UPDATE, and DELETE statements, and how to make the results of your transactions permanent using the COMMIT statement or undo their results entirely using the ROLLBACK statement. You'll also learn how an Oracle database can process multiple transactions at the same time.

Chapter 9: Users, Privileges, and Roles

In this chapter, you'll learn about database users and see how privileges and roles are used to enable users to perform specific tasks in the database.

Chapter 10: Creating Tables, Sequences, Indexes, and Views

You'll learn about tables and sequences, which generate a series of numbers, and indexes, which act like an index in a book and allow you quick access to rows. You'll also learn about views, which are predefined queries on one or more tables; among other benefits, views allow you to hide complexity from a user, and implement another layer of security by only allowing a view to access a limited set of data in the tables. You'll also examine flashback data archives, which are new for Oracle Database 11g. A flashback data archive stores changes made to a table over a period of time.

Chapter 11: Introducing PL/SQL Programming

In this chapter, you'll explore PL/SQL, which is built on top of SQL and enables you to write stored programs in the database that contain SQL statements. PL/SQL contains standard programming constructs.

Chapter 12: Database Objects

You'll learn how to create database object types, which may contain attributes and methods. You'll use object types to define column objects and object tables, and see how to manipulate objects using SQL and PL/SQL.

Chapter 13: Collections

In this chapter, you'll learn how to create collection types, which may contain multiple elements. You'll use collection types to define columns in tables. You'll see how to manipulate collections using SQL and PL/SQL.

Chapter 14: Large Objects

You'll learn about large objects, which can be used to store up to 128 terabytes of character and binary data or point to an external file. You'll also learn about the older LONG types, which are still supported in Oracle Database 11g for backward compatibility.

Chapter 15: Running SQL Using Java

In this chapter, you'll learn the basics of running SQL using Java through the Java Database Connectivity (JDBC) applications programming interface, which is the glue that allows a Java program to access a database.

Chapter 16: SQL Tuning

You'll see SQL tuning tips that you can use to shorten the length of time your queries take to execute. You'll also learn about the Oracle optimizer and examine how to pass hints to the optimizer.

Chapter 17: XML and the Oracle Database

The Extensible Markup Language (XML) is a general-purpose markup language. XML enables you to share structured data across the Internet, and can be used to encode data and other documents. In this chapter, you'll see how to generate XML from relational data and how to save XML in the database.

Appendix: Oracle Data Types

This appendix shows the data types available in Oracle SQL and PL/SQL.

Intended Audience

This book is suitable for the following readers:

- Developers who need to write SQL and PL/SQL.
- Database administrators who need in-depth knowledge of SQL.
- Business users who need to write SQL queries to get information from their organization's database.
- Technical managers or consultants who need an introduction to SQL and PL/SQL.

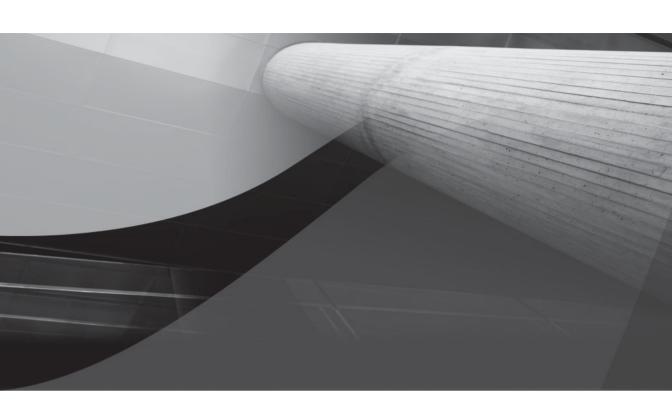
No prior knowledge of the Oracle database, SQL, or PL/SQL is assumed; you can find everything you need to know to become a master in this book.

Retrieving the Examples

All the SQL scripts, programs, and other files used in this book can be downloaded from the Oracle Press website at www.OraclePressBooks.com. The files are contained in a Zip file. Once you've downloaded the Zip file, you need to extract its contents. This will create a directory named sql book that contains the following subdirectories:

- Java Contains the Java programs used in Chapter 15
- **sample_files** Contains the sample files used in Chapter 14
- **SQL** Contains the SQL scripts used throughout the book, including scripts to create and populate the example database tables
- xml_files Contains the XML used in Chapter 17

I hope you enjoy this book!



CHAPTER 1

Introduction



n this chapter, you will learn about the following:

- Relational databases.
- The Structured Query Language (SQL), which is used to access a database.
- SQL*Plus, Oracle's interactive text-based tool for running SQL statements.
- SQL Developer, which is a graphical tool for database development.
- PL/SQL, Oracle's procedural programming language. PL/SQL allows you to develop programs that are stored in the database.

Let's plunge in and consider what a relational database is.

What Is a Relational Database?

The concept of a relational database was originally developed back in 1970 by Dr. E.F. Codd. He laid down the theory of relational databases in his seminal paper entitled "A Relational Model of Data for Large Shared Data Banks," published in *Communications of the ACM* (Association for Computing Machinery), Vol. 13, No. 6, June 1970.

The basic concepts of a relational database are fairly easy to understand. A *relational database* is a collection of related information that has been organized into *tables*. Each table stores data in *rows*; the data is arranged into *columns*. The tables are stored in database *schemas*, which are areas where users may store their own tables. A user may grant *permissions* to other users so they can access their tables.

Most of us are familiar with data being stored in tables—stock prices and train timetables are sometimes organized into tables. One example table used in this book records customer information for an imaginary store; the table stores the customer first names, last names, dates of birth (dobs), and phone numbers:

first_name	last_name	dob	phone
John	Brown	01-JAN-1965	800-555-1211
Cynthia	Green	05-FEB-1968	800-555-1212
Steve	White	16-MAR-1971	800-555-1213
Gail	Black		800-555-1214
Doreen	Blue	20-MAY-1970	

This table could be stored in a variety of forms:

- A card in a box
- An HTML file on a web page
- A table in a database

An important point to remember is that the information that makes up a database is different from the system used to access that information. The software used to access a database is known as a *database management system*. The Oracle database is one such piece of software; other examples include SQL Server, DB2, and MySQL.

Of course, every database must have some way to get data in and out of it, preferably using a common language understood by all databases. Database management systems implement a standard language known as Structured Query Language, or SQL. Among other things, SQL allows you to retrieve, add, modify, and delete information in a database.

Introducing the Structured Query Language (SQL)

Structured Query Language (SQL) is the standard language designed to access relational databases. SQL should be pronounced as the letters "S-Q-L."



NOTE

"S-Q-L" is the correct way to pronounce SQL according to the American National Standards Institute. However, the single word "sequel" is frequently used instead.

SQL is based on the groundbreaking work of Dr. E.F. Codd, with the first implementation of SQL being developed by IBM in the mid-1970s. IBM was conducting a research project known as System R, and SQL was born from that project. Later, in 1979, a company then known as Relational Software Inc. (known today as Oracle Corporation) released the first commercial version of SQL. SQL is now fully standardized and recognized by the American National Standards Institute.

SQL uses a simple syntax that is easy to learn and use. You'll see some simple examples of its use in this chapter. There are five types of SQL statements, outlined in the following list:

- Query statements retrieve rows stored in database tables. You write a query using the SQL SELECT statement.
- Data Manipulation Language (DML) statements modify the contents of tables. There are three DML statements:
 - **INSERT** adds rows to a table.
 - **UPDATE** changes rows.
 - **DELETE** removes rows.
- **Data Definition Language (DDL) statements** define the data structures, such as tables, that make up a database. There are five basic types of DDL statements:
 - **CREATE** creates a database structure. For example, CREATE TABLE is used to create a table; another example is CREATE USER, which is used to create a database user.
 - **ALTER** modifies a database structure. For example, ALTER TABLE is used to modify a table.
 - **DROP** removes a database structure. For example, DROP TABLE is used to remove a
 - **RENAME** changes the name of a table.
 - **TRUNCATE** deletes all the rows from a table.

- Transaction Control (TC) statements either permanently record any changes made to rows, or undo those changes. There are three TC statements:
 - **COMMIT** permanently records changes made to rows.
 - **ROLLBACK** undoes changes made to rows.
 - **SAVEPOINT** sets a "save point" to which you can roll back changes.
- Data Control Language (DCL) statements change the permissions on database structures. There are two DCL statements:
 - **GRANT** gives another user access to your database structures.
 - **REVOKE** prevents another user from accessing your database structures.

There are many ways to run SQL statements and get results back from the database, some of which include programs written using Oracle Forms and Reports. SQL statements may also be embedded within programs written in other languages, such as Oracle's Pro*C++, which allows you to add SQL statements to a C++ program. You can also add SQL statements to a Java program using JDBC; for more details, see my book *Oracle9i JDBC Programming* (Oracle Press, 2002).

Oracle also has a tool called SQL*Plus that allows you to enter SQL statements using the keyboard or to run a script containing SQL statements. SQL*Plus enables you to conduct a "conversation" with the database; you enter SQL statements and view the results returned by the database. You'll be introduced to SQL*Plus next.

Using SQL*Plus

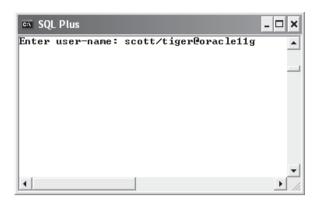
If you're at all familiar with the Oracle database, chances are that you're already familiar with SQL*Plus. If you're not, don't worry: you'll learn how to use SQL*Plus in this book.

In the following sections, you'll learn how to start SQL*Plus and run a guery.

Starting SQL*Plus

If you're using Windows XP Professional Edition and Oracle Database 11*g*, you can start SQL*Plus by clicking start and selecting All Programs | Oracle | Application Development | SQL Plus.

Figure 1-1 shows SQL*Plus running on Windows XP. SQL*Plus asks you for a username. Figure 1-1 shows the scott user connecting to the database (scott is an example user that is contained in many Oracle databases; scott has a default password of tiger). The host string after the @ character tells SQL*Plus where the database is running. If you are running the database on your own computer, you'll typically omit the host string (that is, you enter scott/tiger)—doing this causes SQL*Plus to attempt to connect to a database on the same machine on which SQL*Plus is running. If the database isn't running on your machine, you should speak with your database administrator (DBA) to get the host string. If the scott user doesn't exist or is locked, ask your DBA for an alternative user and password (for the examples in the first part of this chapter, you can use any user; you don't absolutely have to use the scott user).



Oracle Database 11g SQL*Plus Running on Windows XP FIGURE 1-1

If you're using Windows XP and Oracle Database 10g or below, you can run a special Windows-only version of SQL*Plus. You start this version of SQL*Plus by clicking Start and selecting All Programs | Oracle | Application Development | SQL Plus. The Windows-only version of SQL*Plus is deprecated in Oracle Database 11g (that is, it doesn't ship with 11g), but it will still connect to an 11g database. Figure 1-2 shows the Windows-only version of Oracle Database 10g SQL*Plus running on Windows XP.



NOTE

The Oracle Database 11g version of SQL*Plus is slightly nicer than the Windows-only version. In the 11g version, you can scroll through previous commands you've run by pressing the UP and DOWN ARROW keys on the keyboard.

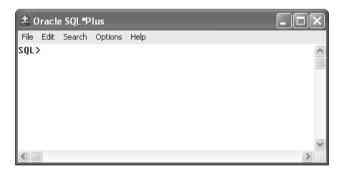


FIGURE 1-2 Oracle Database 10g SQL*Plus Running on Windows XP

Starting SQL*Plus from the Command Line

You can also start SQL*Plus from the command line. To do this, you use the sqlplus command. The full syntax for the sqlplus command is

sqlplus [user name[/password[@host string]]]

where

- user name is the name of the database user.
- **password** is the password for the database user.
- *host string* is the database you want to connect to.

The following examples show sqlplus commands:

```
sqlplus scott/tiger
sqlplus scott/tiger@orcl
```

If you're using SQL*Plus with a Windows operating system, the Oracle installer automatically adds the directory for SQL*Plus to your path. If you're using a non-Windows operating system (for example, Unix or Linux), either you must be in the same directory as the SQL*Plus program to run it or, better still, you should add the directory to your path. If you need help with that, talk to your system administrator.

For security, you can hide the password when connecting to the database. For example, you can enter

```
sqlplus scott@orcl
```

SQL*Plus then prompts you to enter the password. As you type in the password, it is hidden from prying eyes. This also works when starting SQL*Plus in Windows.

You can also just enter

sqlplus

SQL*Plus then prompts you for the user name and password. You can specify the host string by adding it to the user name (for example, scott@orcl).

Performing a SELECT Statement Using SQL*Plus

Once you're logged onto the database using SQL*Plus, go ahead and run the following SELECT statement (it returns the current date):

SELECT SYSDATE FROM dual;

SYSDATE is a built-in database function that returns the current date, and the dual table is a table that contains a single row. The dual table is useful when you need the database to evaluate an expression (e.g., 2 * 15 / 5), or when you want to get the current date.

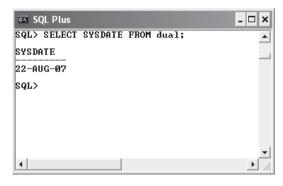


NOTE

SQL statements directly entered into SQL*Plus are terminated using a semicolon character (;).

This illustration shows the results of this SELECT statement in SOL*Plus running on Windows. As you can see, the query displays the current date from the database.

You can edit your last SQL statement in SQL*Plus by entering EDIT. Doing this is useful when you make a mistake or you want to make a change to your SQL statement. On Windows, when you enter EDIT you are taken to the Notepad application; you then use Notepad to edit your SQL statement. When you exit Notepad and save your statement, the new statement is



passed back to SQL*Plus, where you can re-execute it by entering a forward slash (/). On Linux or Unix, the default editor is typically set to vi or emacs.

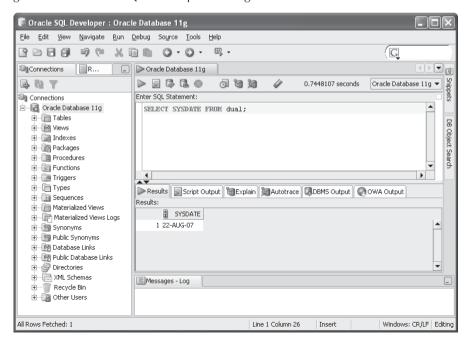


NOTE

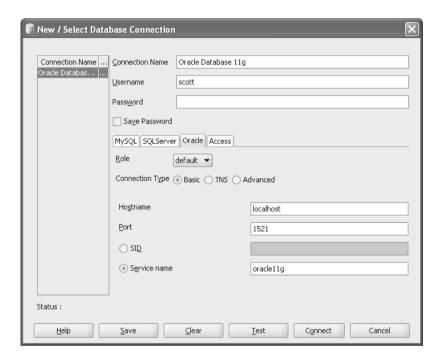
You'll learn more about editing SQL statements using SQL*Plus in Chapter 3.

SQL Developer

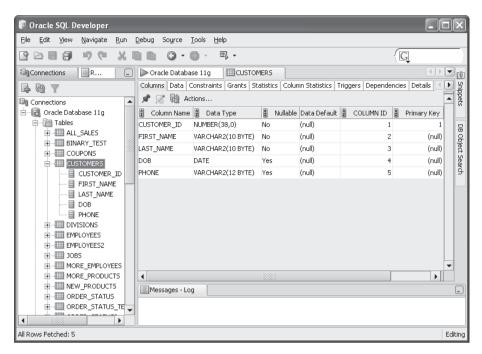
You can also enter SQL statements using SQL Developer. SQL Developer uses a very nice graphical user interface through which you can enter SQL statements, examine database tables, run scripts, edit and debug PL/SQL code, and much more. SQL Developer can connect to any Oracle Database, version 9.2.0.1 and higher, and runs on Windows, Linux, and Mac OSX. The following illustration shows SQL Developer running.



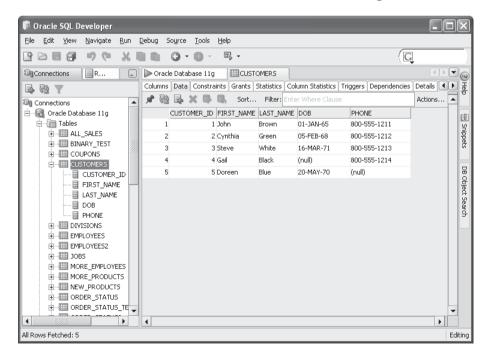
You need to have Java installed on your computer before you can run SQL Developer. If you're using Windows XP Professional Edition and Oracle Database 11g, you start SQL Developer by clicking Start and selecting All Programs | Oracle | Application Development | SQL Developer. SQL Developer will prompt you to select the Java executable. You then browse to the location where you have installed it and select the executable. Next, you need to create a connection by right-clicking Connections and selecting New Connection, as shown in the following illustration.



Once you've created a connection and tested it, you can use it to connect to the database and run queries, examine database tables, and so on. The following illustration shows the details for a database table named customers.



You can also view the data stored in a table, as shown in the following illustration.



You can see full details on using SQL Developer by selecting Help | Table of Contents from the menu bar in SQL Developer.

In the next section, you'll learn how to create the imaginary store schema used throughout this book.

Creating the Store Schema

The imaginary store sells items such as books, videos, DVDs, and CDs. The database for the store will hold information about the customers, employees, products, and sales. The SQL*Plus script to create the database is named store_schema.sql, which is located in the SQL directory where you extracted the Zip file for this book. The store_schema.sql script contains the DDL and DML statements used to create the store schema. You'll now learn how to run the store_schema.sql script.

Running the SQL*Plus Script to Create the Store Schema

You perform the following steps to create the store schema:

- 1. Start SQL*Plus.
- 2. Log into the database as a user with privileges to create new users, tables, and PL/SQL packages. I run scripts in my database using the system user; this user has all the required privileges. You may need to speak with your database administrator about setting up a user for you with the required privileges (they might also run the store_schema.sql script for you).
- 3. Run the store schema.sql script from within SQL*Plus using the @ command.

The @ command has the following syntax:

@ directory\store_schema.sql

where *directory* is the directory where your store_schema.sql script is located. For example, if the script is stored in E:\sql book\SQL, then you enter

@ E:\sql book\SQL\store schema.sql

If you have placed the store_schema.sql script in a directory that contains spaces, then you must place the directory and script in quotes after the @ command. For example:

@ "E:\Oracle SQL book\sql book\SQL\store schema.sql"

If you're using Unix or Linux and you saved the script in a directory named SQL in the tmp file system, then you enter

@ /tmp/SQL/store schema.sql



NOTE

Windows uses backslash characters (\) in directory paths, whereas Unix and Linux use forward slash characters (/).

The first executable line in the store schema.sql script attempts to drop the store user, generating an error because the user doesn't exist yet. Don't worry about the error: the line is there so you don't have to manually drop the store user when recreating the schema later in the book.

When the store schema.sql script has finished running, you'll be connected as the store user. If you want to, open the store schema.sql script using a text editor like Windows Notepad and examine the statements contained in it. Don't worry about the details of the statements contained in the script—you'll learn the details as you progress through this book.



NOTE

To end SQL*Plus, you enter EXIT. To reconnect to the store schema in SQL*Plus, you enter store as the user name with a password of store password. While you're connected to the database, SQL*Plus maintains a database session for you. When you disconnect from the database, your session is ended. You can disconnect from the database and keep SQL*Plus running by entering DISCONNECT. You can then reconnect to a database by entering CONNECT.

Data Definition Language (DDL) Statements Used to Create the Store Schema

As mentioned earlier, Data Definition Language (DDL) statements are used to create users and tables, plus many other types of structures in the database. In this section, you'll see the DDL statements used to create the store user and some of the tables.



NOTE

The SQL statements you'll see in the rest of this chapter are the same as those contained in the store schema.sql script. You don't have to type the statements in yourself: you just run the store schema .sql script.

The next sections describe the following:

- How to create a database user
- The commonly used data types used in an Oracle database
- Some of the tables in the imaginary store

Creating a Database User

To create a user in the database, you use the CREATE USER statement. The simplified syntax for the CREATE USER statement is as follows:



CREATE USER user name IDENTIFIED BY password;

where

- user name is the user name
- password is the password for the user

For example, the following CREATE USER statement creates the store user with a password of store password:

CREATE USER store IDENTIFIED BY store password;

If you want the user to be able to work in the database, the user must be granted the necessary *permissions* to do that work. In the case of store, this user must be able to log onto the database (which requires the connect permission) and create items like database tables (which requires the resource permission). Permissions are granted by a privileged user (for example, the system user) using the GRANT statement.

The following example grants the connect and resource permissions to store:

GRANT connect, resource TO store;

Once a user has been created, the database tables and other database objects can be created in the associated schema for that user. Many of the examples in this book use the store schema. Before I get into the details of the store tables, you need to know about the commonly used Oracle database types.

The Common Oracle Database Types

There are many types that may be used to handle data in an Oracle database. Some of the commonly used types are shown in Table 1-1.

You can see all the data types in the appendix. The following table illustrates a few examples of how numbers of type NUMBER are stored in the database.

Format	Number Supplied	Number Stored
NUMBER	1234.567	1234.567
NUMBER(6, 2)	123.4567	123.46
NUMBER(6, 2)	12345.67	Number exceeds the specified precision and is therefore rejected by the database.

Examining the Store Tables

In this section, you'll learn how the tables for the store schema are created. Some of the information held in the store schema includes

- Customer details
- Types of products sold
- Product details
- A history of the products purchased by the customers
- Employees of the store
- Salary grades

The following tables are used to hold the information:

- **customers** holds the customer details.
- **product types** holds the types of products sold by the store.

- **products** holds the product details.
- purchases holds which products were purchased by which customers.
- **employees** holds the employee details.
- **salary_grades** holds the salary grade details.

Oracle Type	Meaning
CHAR(length)	Stores strings of a fixed length. The <code>length</code> parameter specifies the length of the string. If a string of a smaller length is stored, it is padded with spaces at the end. For example, <code>CHAR(2)</code> may be used to store a fixed-length string of two characters; if <code>'C'</code> is stored in a <code>CHAR(2)</code> , then a single space is added at the end; <code>'CA'</code> is stored as is, with no padding.
VARCHAR2(length)	Stores strings of a variable length. The <code>length</code> parameter specifies the maximum length of the string. For example, <code>VARCHAR2(20)</code> may be used to store a string of up to 20 characters in length. No padding is used at the end of a smaller string.
DATE	Stores dates and times. The DATE type stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second. The DATE type may be used to store dates and times between January 1, 4712 B.C. and December 31, 4712 A.D.
INTEGER	Stores integers. An integer doesn't contain a floating point: it is a whole number, such as 1, 10, and 115.
<pre>NUMBER(precision, scale)</pre>	Stores floating point numbers, but may also be used to store integers. The <code>precision</code> is the maximum number of digits (left and right of a decimal point, if used) that may be used for the number. The maximum precision supported by the Oracle database is 38. The <code>scale</code> is the maximum number of digits to the right of a decimal point (if used). If neither <code>precision</code> nor <code>scale</code> is specified, any number may be stored up to a precision of 38 digits. Any attempt to store a number that exceeds the <code>precision</code> is rejected by the database.
BINARY_FLOAT	Introduced in Oracle Database 10g, stores a single precision 32-bit floating point number. You'll learn more about BINARY_FLOAT later in the section "The BINARY_FLOAT and BINARY_DOUBLE Types."
BINARY_DOUBLE	Introduced in Oracle Database 10 <i>g</i> , stores a double precision 64-bit floating point number. You'll learn more about BINARY_DOUBLE later in the section "The BINARY_FLOAT and BINARY_DOUBLE Types."

 TABLE 1-1
 Commonly Used Oracle Data Types



NOTE

The store_schema.sql script creates other tables and database items not mentioned in the previous list. You'll learn about these items in later chapters.

In the following sections, you'll see the details of some of the tables, and you'll see the CREATE TABLE statements included in the store schema.sql script that create the tables.

The customers Table The customers table holds the details of the customers. The following items are held in this table:

- First name
- Last name
- Date of birth (dob)
- Phone number

Each of these items requires a column in the customers table. The customers table is created by the store_schema.sql script using the following CREATE TABLE statement:

```
CREATE TABLE customers (
    customer_id INTEGER CONSTRAINT customers_pk PRIMARY KEY,
    first_name VARCHAR2(10) NOT NULL,
    last_name VARCHAR2(10) NOT NULL,
    dob DATE,
    phone VARCHAR2(12)
);
```

As you can see, the customers table contains five columns, one for each item in the previous list, and an extra column named customer id. The columns are

- customer_id Contains a unique integer for each row in the table. Each table should have one or more columns that uniquely identifies each row; the column(s) are known as the primary key. The CONSTRAINT clause indicates that the customer_id column is the primary key. A CONSTRAINT clause restricts the values stored in a column, and, for the customer_id column, the PRIMARY KEY keywords indicate that the customer_id column must contain a unique value for each row. You can also attach an optional name to a constraint, which must immediately follow the CONSTRAINT keyword—for example, customers_pk. You should always name your primary key constraints, so that when a constraint error occurs it is easy to spot where it happened.
- first_name Contains the first name of the customer. You'll notice the use of the NOT NULL constraint for this column—this means that a value must be supplied for first_name when adding or modifying a row. If a NOT NULL constraint is omitted, a user doesn't need to supply a value and the column can remain empty.

- last name Contains the last name of the customer. This column is NOT NULL, and therefore a value must be supplied when adding or modifying a row.
- **dob** Contains the date of birth for the customer. Notice that no NOT NULL constraint is specified for this column; therefore, the default NULL is assumed, and a value is optional when adding or modifying a row.
- **phone** Contains the phone number of the customer. This is an optional value.

The store schema.sql script populates the customers table with the following rows:

```
customer id first name last name dob phone
      ______ ____
                   1 John Brown 01-JAN-65 800-555-1211
2 Cynthia Green 05-FEB-68 800-555-1212
3 Steve White 16-MAR-71 800-555-1213
4 Gail Black 800-555-1214
5 Doreen Blue 20-MAY-70
```

Notice that customer #4's date of birth is null, as is customer #5's phone number. You can see the rows in the customers table for yourself by executing the following SELECT statement using SQL*Plus:

SELECT * FROM customers;

The asterisk (*) indicates that you want to retrieve all the columns from the customers table.



In this book, SQL statements shown in **bold** are statements you should type in and run if you want to follow along with the examples. Nonbold statements are statements you don't need to type in.

The product_types Table The product types table holds the names of the product types sold by the store. This table is created by the store schema.sql script using the following CREATE TABLE statement:

```
CREATE TABLE product types (
     product type id INTEGER CONSTRAINT product types pk PRIMARY KEY,
     name VARCHAR2(10) NOT NULL
   );
```

The product types table contains the following two columns:

- product type id uniquely identifies each row in the table; the product type id column is the primary key for this table. Each row in the product types table must have a unique integer value for the product type id column.
- name contains the product type name. It is a NOT NULL column, and therefore a value must be supplied when adding or modifying a row.

The store_schema.sql script populates the product_types table with the following rows:

```
product_type_id name

1 Book
2 Video
3 DVD
4 CD
5 Magazine
```

The product_types table contains the product types for the store. Each product sold by the store must be one of these types.

You can see the rows in the product_types table for yourself by executing the following SELECT statement using SQL*Plus:

```
SELECT * FROM product types;
```

The products Table The products table holds the products sold by the store. The following pieces of information are held for each product:

- Product type
- Name
- Description
- Price

The store_schema.sql script creates the products table using the following CREATE TABLE statement:

```
CREATE TABLE products (

product_id INTEGER CONSTRAINT products_pk PRIMARY KEY,

product_type_id INTEGER

CONSTRAINT products_fk_product_types

REFERENCES product_types(product_type_id),

name VARCHAR2(30) NOT NULL,

description VARCHAR2(50),

price NUMBER(5, 2)
);
```

The columns in this table are as follows:

- product_id uniquely identifies each row in the table. This column is the primary key of the table.
- product_type_id associates each product with a product type. This column is a reference to the product_type_id column in the product_types table; it is known as a foreign key because it references a column in another table. The table containing the foreign key (the products table) is known as the detail or child table, and the table that is referenced (the product_types table) is known as the master or parent table.

This type of relationship is known as a master-detail or parent-child relationship. When you add a new product, you associate that product with a type by supplying a matching product types.product type id value in the products.product type id column (you'll see an example later).

- name contains the product name, which must be specified, as the name column is NOT NULL.
- **description** contains an optional description of the product.
- price contains an optional price for a product. This column is defined as NUMBER (5, 2)—the precision is 5, and therefore a maximum of 5 digits may be supplied for this number. The scale is 2; therefore 2 of those maximum 5 digits may be to the right of the decimal point.

The following is a subset of the rows stored in the products table:

product_id	product_type_id	name	description	price
1	1	Modern Science	A description of modern science	19.95
2	1	Chemistry	Introduction to Chemistry	30
3	2	Supernova	A star explodes	25.99
4	2	Tank War	Action movie about a future war	13.95

The first row in the products table has a product type id of 1, which means the product is a book (this product type id matches the "book" product type in the product types table). The second product is also a book, but the third and fourth products are videos (their product type id is 2, which matches the "video" product type in the product types table).

You can see all the rows in the products table for yourself by executing the following SELECT statement using SQL*Plus:

SELECT * FROM products;

The purchases Table The purchases table holds the purchases made by a customer. For each purchase made by a customer, the following information is held:

- Product ID
- Customer ID
- Number of units of the product that were purchased by the customer

The store_schema.sql script uses the following CREATE TABLE statement to create the purchases table:

```
CREATE TABLE purchases (
    product_id INTEGER
        CONSTRAINT purchases_fk_products
        REFERENCES products(product_id),
        customer_id INTEGER
        CONSTRAINT purchases_fk_customers
        REFERENCES customers(customer_id),
        quantity INTEGER NOT NULL,
        CONSTRAINT purchases_pk PRIMARY KEY (product_id, customer_id)
);
```

The columns in this table are as follows:

- **product_id** contains the ID of the product that was purchased. This must match a product_id column value in the products table.
- **customer_id** contains the ID of a customer who made the purchase. This must match a customer_id column value in the customers table.
- quantity contains the number of units of the product that were purchased by the customer.

The purchases table has a primary key constraint named purchases_pk that spans two columns: product_id and customer_id. The combination of the two column values must be unique for each row. When a primary key consists of multiple columns, it is known as a *composite* primary key.

The following is a subset of the rows that are stored in the purchases table:

product_id	customer_id	quantity
1	1	1
2	1	3
1	4	1
2	2	1
1	3	1

As you can see, the combination of the values in the product_id and customer_id columns is unique for each row.

You can see all the rows in the purchases table for yourself by executing the following SELECT statement using SQL*Plus:

```
SELECT * FROM purchases;
```

The employees Table The employees table holds the details of the employees. The following information is held in the table:

■ Employee ID

- The ID of the employee's manager (if applicable)
- First name
- Last name
- Title
- Salary

The store schema.sql script uses the following CREATE TABLE statement to create the employees table:

```
CREATE TABLE employees (
       employee id INTEGER CONSTRAINT employees pk PRIMARY KEY,
      manager id INTEGER,
      first name VARCHAR2(10) NOT NULL,
      last name VARCHAR2(10) NOT NULL,
      title VARCHAR2(20),
      salary NUMBER(6, 0)
     );
```

The store schema.sql script populates the employees table with the following rows:

employee_id	manager_id	first_name	last_name	title	salary
1		James	Smith	CEO	800000
2	1	Ron	Johnson	Sales Manager	600000
3	2	Fred	Hobbs	Salesperson	150000
4	2	Susan	Jones	Salesperson	500000

As you can see, James Smith doesn't have a manager. That's because he is the CEO of the store.

The salary_grades Table The salary grades table holds the different salary grades available to employees. The following information is held:

- Salary grade ID
- Low salary boundary for the grade
- High salary boundary for the grade

The store schema.sql script uses the following CREATE TABLE statement to create the salary grades table:

```
CREATE TABLE salary grades (
     salary grade id INTEGER CONSTRAINT salary grade pk PRIMARY KEY,
     low salary NUMBER(6, 0),
     high salary NUMBER(6, 0)
   );
```

The store schema.sql script populates the salary grades table with the following rows:

salary_c	grade_id	low_salary	high_salary
	1	1	250000
	2	250001	500000
	3	500001	750000
	4	750001	999999

Adding, Modifying, and Removing Rows

In this section, you'll learn how to add, modify, and remove rows in database tables by using the SQL INSERT, UPDATE, and DELETE statements. You can make your row changes permanent in the database using the COMMIT statement, or you can undo them using the ROLLBACK statement. This section doesn't exhaustively cover all the details of using these statements; you'll learn more about them in Chapter 8.

Adding a Row to a Table

You use the INSERT statement to add new rows to a table. You can specify the following information in an INSERT statement:

- The table into which the row is to be inserted
- A list of columns for which you want to specify column values
- A list of values to store in the specified columns

When inserting a row, you need to supply a value for the primary key and all other columns that are defined as NOT NULL. You don't have to specify values for the other columns if you don't want to; those columns will be automatically set to null if you omit values for them.

You can tell which columns are defined as NOT NULL using the SQL*Plus DESCRIBE command. The following example DESCRIBEs the customers table:

SOL> DESCRIBE customers

Name	Null?		Туре
CUSTOMER_ID	NOT	NULL	NUMBER (38)
FIRST_NAME	NOT	NULL	VARCHAR2(10)
LAST_NAME	NOT	NULL	VARCHAR2(10)
DOB			DATE
PHONE			VARCHAR2 (12)

As you can see, the customer_id, first_name, and last_name columns are NOT NULL, meaning that you must supply a value for these columns. The dob and phone columns don't require a value; you could omit the values if you wanted, and they would be automatically set to null.

Go ahead and run the following INSERT statement, which adds a row to the customers table; notice that the order of values in the VALUES list matches the order in which the columns are specified in the column list:

```
SQL> INSERT INTO customers (
2 customer_id, first_name, last_name, dob, phone
```

```
3 ) VALUES (
 4 6, 'Fred', 'Brown', '01-JAN-1970', '800-555-1215'
 5):
1 row created.
```



SQL*Plus automatically numbers lines after you hit ENTER at the end of

In the previous example, SQL*Plus responds that one row has been created after the INSERT statement is executed. You can verify this by running the following SELECT statement:

SELECT *

FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	
6	Fred	Brown	01-JAN-70	800-555-1215

Notice the new row that has been added to the end of the table.

By default, the Oracle database displays dates in the format DD-MON-YY, where DD is the day number, MON is the first three characters of the month (in uppercase), and YY is the last two digits of the year. The database actually stores all four digits for the year, but by default it only displays the last two digits.

When a row is added to the customers table, a unique value for the customer id column must be given. The Oracle database will prevent you from adding a row with a primary key value that already exists in the table; for example, the following INSERT statement causes an error because a row with a customer id of 1 already exists:

```
SQL> INSERT INTO customers (
          customer id, first name, last name, dob, phone
      3 ) VALUES (
         1, 'Lisa', 'Jones', '02-JAN-1971', '800-555-1225'
      5);
     INSERT INTO customers (
    ERROR at line 1:
    ORA-00001: unique constraint (STORE.CUSTOMERS_PK) violated
```

Notice that the name of the constraint is shown in the error (CUSTOMERS PK). That's why you should always name your primary key constraints; otherwise, the Oracle database assigns an unfriendly system-generated name to a constraint (for example, SYS C0011277).

Modifying an Existing Row in a Table

You use the UPDATE statement to change rows in a table. Normally, when you use the UPDATE statement, you specify the following information:

- The table containing the rows that are to be changed
- A WHERE clause that specifies the rows that are to be changed
- A list of column names, along with their new values, specified using the SET clause

You can change one or more rows using the same UPDATE statement. If more than one row is specified, the same change will be made for all the rows. The following example updates customer #2's last name to Orange:

```
UPDATE customers

SET last_name = 'Orange'

WHERE customer_id = 2;

1 row updated.
```

SQL*Plus confirms that one row was updated.



CAUTION

If you forget to add a WHERE clause, then all the rows will be updated.

The following query confirms the update worked:

```
SELECT *

FROM customers
WHERE customer_id = 2;

CUSTOMER_ID FIRST_NAME LAST_NAME DOB PHONE
```

Removing a Row from a Table

You use the DELETE statement to remove rows from a table. You typically use a WHERE clause to limit the rows you wish to delete; if you don't, *all* the rows will be deleted from the table.

2 Cynthia Orange 05-FEB-68 800-555-1212

The following DELETE statement removes customer #2:

```
DELETE FROM customers
WHERE customer_id = 2;

1 row deleted.
```

To undo the changes you've made to the rows, you use ROLLBACK:

ROLLBACK;

Rollback complete.

Go ahead and run the ROLLBACK to undo any changes you've made so far. That way, your results will match those shown in subsequent chapters.



NOTE

You can make changes to rows permanent using COMMIT. You'll see how to do that in Chapter 8.

The BINARY_FLOAT and BINARY_DOUBLE Types

Oracle Database 10g introduced two new data types: BINARY FLOAT and BINARY DOUBLE. BINARY FLOAT stores a single precision 32-bit floating point number; BINARY DOUBLE stores a double precision 64-bit floating point number. These new data types are based on the IEEE (Institute of Electrical and Electronics Engineers) standard for binary floating-point arithmetic.

Benefits of BINARY FLOAT and BINARY DOUBLE

BINARY FLOAT and BINARY DOUBLE are intended to complement the existing NUMBER type. BINARY FLOAT and BINARY DOUBLE offer the following benefits over NUMBER:

- **Smaller storage required** BINARY FLOAT and BINARY DOUBLE require 5 and 9 bytes of storage space, whereas NUMBER might use up to 22 bytes.
- Greater range of numbers represented BINARY FLOAT and BINARY DOUBLE support numbers much larger and smaller than can be stored in a NUMBER.
- Faster performance of operations Operations involving BINARY FLOAT and BINARY DOUBLE are typically performed faster than NUMBER operations. This is because BINARY FLOAT and BINARY DOUBLE operations are typically performed in the hardware, whereas NUMBERS must first be converted using software before operations can be performed.
- **Closed operations** Arithmetic operations involving BINARY FLOAT and BINARY DOUBLE are closed, which means that either a number or a special value is returned. For example, if you divide a BINARY FLOAT by another BINARY FLOAT, a BINARY FLOAT is returned.
- Transparent rounding BINARY FLOAT and BINARY DOUBLE use binary (base 2) to represent a number, whereas NUMBER uses decimal (base 10). The base used to represent a number affects how rounding occurs for that number. For example, a decimal floatingpoint number is rounded to the nearest decimal place, but a binary floating-point number is rounded to the nearest binary place.



TIP

If you are developing a system that involves a lot of numerical computations, you should use BINARY FLOAT and BINARY DOUBLE to represent numbers. Of course, you must be using Oracle Database 10g or higher.

Using BINARY_FLOAT and BINARY_DOUBLE in a Table

The following statement creates a table named binary_test that contains a BINARY_FLOAT and a BINARY DOUBLE column:

```
CREATE TABLE binary_test (
    bin_float BINARY_FLOAT,
    bin_double BINARY_DOUBLE
);
```



NOTE

You'll find a script named oracle_10g_examples.sql in the SQL directory that creates the binary_test table in the store schema. The script also performs the INSERT statements you'll see in this section. You can run this script if you are using Oracle Database 10g or higher.

The following example adds a row to the binary test table:

```
INSERT INTO binary_test (
   bin_float, bin_double
) VALUES (
   39.5f, 15.7d
);
```

Notice that f indicates a number is a BINARY_FLOAT, and d indicates a number is a BINARY_DOUBLE.

Special Values

You can also use the special values shown in Table 1-2 with a BINARY_FLOAT or BINARY_DOUBLE.

The following example inserts BINARY_FLOAT_INFINITY and BINARY_DOUBLE_
INFINITY into the binary test table:

```
INSERT INTO binary_test (
    bin_float, bin_double
) VALUES (
    BINARY_FLOAT_INFINITY, BINARY_DOUBLE_INFINITY
);
```

Special Value

BINARY_FLOAT_NAN
BINARY_FLOAT_INFINITY
BINARY_DOUBLE_NAN
BINARY_DOUBLE_INFINITY

Description

Not a number (NaN) for the BINARY_FLOAT type
Infinity (INF) for the BINARY_FLOAT type
Not a number (NaN) for the BINARY_DOUBLE type
Infinity (INF) for the BINARY_DOUBLE type

The following query retrieves the rows from binary test:

SELECT *

FROM binary_test;

```
BIN FLOAT BIN DOUBLE
3.95E+001 1.57E+001
    Inf Inf
```

Quitting SQL*Plus

You use the EXIT command to quit from SQL*Plus. The following example quits SQL*Plus using the EXIT command:

EXIT



NOTE

When you exit SQL*Plus in this way, it automatically performs a COMMIT for you. If SQL*Plus terminates abnormally—for example, if the computer on which SQL*Plus is running crashes—a ROLLBACK is automatically performed. You'll learn more about this in Chapter 8.

Introducing Oracle PL/SQL

PL/SQL is Oracle's procedural language that allows you to add programming constructs around SQL statements. PL/SQL is primarily used for creating procedures and functions in a database that contain business logic. PL/SQL contains standard programming constructs such as

- Variable declarations
- Conditional logic (if-then-else, and so on)
- Loops
- Procedures and functions

The following CREATE PROCEDURE statement creates a procedure named update product price(). The procedure multiplies the price of a product by a factor—the product ID and the factor are passed as parameters to the procedure. If the specified product doesn't exist, the procedure takes no action; otherwise, it updates the product price.

NOTE

Don't worry about the details of the PL/SQL shown in the following listing—you'll learn all about PL/SQL in Chapter 11. I just want you to get a feel for PL/SQL at this stage.

```
CREATE PROCEDURE update product price (
    p product id IN products.product id%TYPE,
    p factor IN NUMBER
    product count INTEGER;
```

```
BEGIN
 -- count the number of products with the
 -- supplied product id (will be 1 if the product exists)
 SELECT COUNT(*)
  INTO product count
  FROM products
 WHERE product id = p product id;
  -- if the product exists (i.e. product count = 1) then
  -- update that product's price
  IF product count = 1 THEN
   UPDATE products
    SET price = price * p factor
   WHERE product id = p product id;
    COMMIT;
 END IF;
EXCEPTION
 WHEN OTHERS THEN
   ROLLBACK;
END update product price;
```

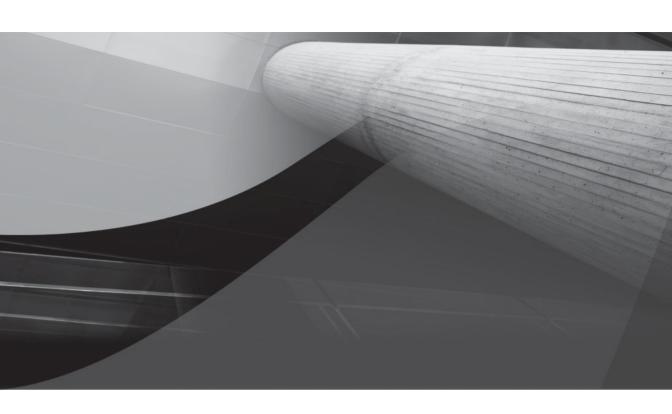
Exceptions are used to handle errors that occur in PL/SQL code. The EXCEPTION block in the previous example performs a ROLLBACK if an exception is thrown in the code.

Summary

In this chapter, you have learned the following:

- A relational database is a collection of related information that has been organized into structures known as tables. Each table contains rows that are further organized into columns. These tables are stored in the database in structures known as schemas, which are areas where database users may store their objects (such as tables and PL/SQL procedures).
- Structured Query Language (SQL) is the standard language designed to access relational databases.
- SQL*Plus allows you to run SQL statements and SQL*Plus commands.
- SQL Developer is a graphical tool for database development.
- How to run SELECT, INSERT, UPDATE, and DELETE statements.
- PL/SQL is Oracle's procedural language that contains programming statements.

In the next chapter, you'll learn more about retrieving information from database tables.



CHAPTER

2

Retrieving Information from Database Tables



n this chapter, you will see how to

- Retrieve information from one or more database tables using SELECT statements
- Use arithmetic expressions to perform calculations
- Limit the retrieval of rows to just those you are interested in using a WHERE clause
- Sort the rows retrieved from a table

The examples in this section use the store schema. If you want to follow along with the examples, you should start SQL*Plus and log in as the store user.

Performing Single Table SELECT Statements

You use the SELECT statement to retrieve information from database tables. In the statement's simplest form, you specify the table and columns from which you want to retrieve data. The following SELECT statement retrieves the customer_id, first_name, last_name, dob, and phone columns from the customers table:

```
SELECT customer_id, first_name, last_name, dob, phone FROM customers;
```

Immediately after the SELECT keyword, you supply the column names that you want to retrieve; after the FROM keyword, you supply the table name. The SQL statement is ended using a semicolon (;). SELECT statements are also known as *queries*.

You don't tell the database management system software exactly how to access the information you want. You just tell it what you want and let the software worry about how to actually get it. The items that immediately follow the SELECT keyword needn't always be columns from a table: They can be any valid expression. You'll see examples of expressions later in this chapter.

After you press ENTER at the end of the SQL statement, the statement is executed and the results are returned to SQL*Plus for display on the screen:

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

The rows returned by the database are known as a *result set*. As you can see from the example, the Oracle database converts the column names into their uppercase equivalents. Character and date columns are left-justified; number columns are right-justified. By default, the Oracle database displays dates in the format DD-MON-YY, where DD is the day number, MON is the first three characters of the month (in uppercase), and YY is the last two digits of the year. The database actually stores all four digits for the year, but by default it displays only the last two digits.



NOTE

A database administrator can change the default display format for dates by setting an Oracle database parameter called NLS DATE FORMAT. You'll learn more about dates in Chapter 5.

Although you can specify column names and table names using either lowercase or uppercase text, it is better to stick with one style. The examples in this book use uppercase for SQL and Oracle keywords, and lowercase for everything else.

Retrieving All Columns from a Table

If you want to retrieve all columns in a table, you can use the asterisk character (*) in place of a list of columns. In the following query, the asterisk is used to retrieve all columns from the customers table:

SELECT * FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
.5	Doreen	Blue	20-MAY-70	

As you can see, all the columns in the customers table are retrieved.

Specifying Rows to Retrieve Using the WHERE Clause

You use the WHERE clause in a query to specify the rows you want to retrieve. This is very important, as Oracle has the capacity to store large numbers of rows in a table, and you may be interested in only a very small subset of those rows. You place the WHERE clause after the FROM clause:

```
SELECT list of items
    FROM list of tables
    WHERE list of conditions;
```

In the following query, the WHERE clause is used to retrieve the row from the customers table where the customer id column is equal to 2:

SELECT * FROM customers WHERE customer id = 2;

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB PHONE
______ _____
      2 Cynthia Green 05-FEB-68 800-555-1212
```

Row Identifiers

Each row in an Oracle database has a unique row identifier, or *rowid*, which is used internally by the Oracle database to store the physical location of the row. A rowid is an 18-digit number that is represented as a base-64 number. You can view the rowid for rows in a table by retrieving the ROWID column in a query. For example, the following query retrieves the ROWID and customer id columns from the customers table; notice the base-64 number in the output:

SELECT ROWID, customer_id FROM customers;

ROWID	CUSTOMER_ID
АААҒ4 УААВАААН ЕКААА	1
AAAF4yAABAAAHeKAAB	2
AAAF4yAABAAAHeKAAC	3
AAAF4yAABAAAHeKAAD	4
AAAF4yAABAAAHeKAAE	5

When you describe a table using the SQL*Plus DESCRIBE command, ROWID doesn't appear in the output from the command because it is only used internally by the database. ROWID is known as a *pseudo* column. The following example describes the customers table; notice ROWID doesn't appear in the output:

DESCRIBE customers

Name	Null?		Type
CUSTOMER_ID	NOT	NULL	NUMBER (38)
FIRST_NAME	NOT	NULL	VARCHAR2(10)
LAST_NAME	NOT	NULL	VARCHAR2(10)
DOB			DATE
PHONE			VARCHAR2(12)

Row Numbers

Another pseudo column is ROWNUM, which returns the row number in a result set. The first row returned by a query has a row number of 1, the second has a row number of 2, and so on. For example, the following query includes ROWNUM when retrieving the rows from the customers table:

SELECT ROWNUM, customer_id, first_name, last_name FROM customers;

ROWNUM	CUSTOMER_ID	FIRST_NAME	LAST_NAME
1	1	John	Brown
2	2	Cynthia	Green
3	3	Steve	White
4	4	Gail	Black
5	5	Doreen	Blue

Here's another example:

```
SELECT ROWNUM, customer id, first name, last name
   FROM customers
   WHERE customer id = 3;
      ROWNUM CUSTOMER ID FIRST NAME LAST NAME
    ----- -----
         1
             3 Steve White
```

Performing Arithmetic

Oracle allows you to perform arithmetic in SQL statements using arithmetic expressions, consisting of addition, subtraction, multiplication, and division. Arithmetic expressions consist of two operands—numbers or dates—and an arithmetic operator. The four arithmetic operators are shown in the following table:

Operator	Description
+	Addition
_	Subtraction
*	Multiplication
/	Division

The following query shows how to use the multiplication operator (*) to calculate 2 multiplied by 6 (the numbers 2 and 6 are the operands):

```
SELECT 2*6
     FROM dual:
           2*6
            12
```

As you can see from this guery, the correct result of 12 is displayed. The use of 2*6 in the query is an example of an expression. An expression may contain a combination of columns, literal values, and operators.

Performing Date Arithmetic

You can use the addition and subtraction operators with dates. You can add a number—representing a number of days—to a date. The following example adds two days to July 25, 2007, and displays the resulting date:

```
SELECT TO_DATE('25-JUL-2007') + 2
    FROM dual:
    TO DATE (
    27-JUL-07
```

The dual Table

You'll notice the use of the dual table in the previous example. I mentioned the dual table in the previous chapter—dual is a table that contains a single row. The following output from the DESCRIBE command shows the structure of the dual table, along with a query that retrieves the row from the dual table:

Null? Type

VARCHAR2(1)

Name

SELECT *
FROM dual;

DUMMY

D -X

Notice the dual table has one VARCHAR2 column named dummy and contains a single row with the value X.



NOTE

TO_DATE() is a function that converts a string to a date. You'll learn more about TO DATE() in Chapter 5.

The next example subtracts three days from August 2, 2007:

SELECT TO_DATE('02-AUG-2007') - 3 FROM dual;

```
TO_DATE('
-----
30-JUL-07
```

You can also subtract one date from another, yielding the number of days between the two dates. The following example subtracts July 25, 2007, from August 2, 2007:

SELECT TO_DATE('02-AUG-2007') - TO_DATE('25-JUL-2007') FROM dual;

Using Columns in Arithmetic

Operands do not have to be literal numbers or dates; they may also be columns from a table. In the following query, the name and price columns are retrieved from the products table; notice that 2 is added to the value in the price column using the addition operator (+) to form the expression price + 2:

SELECT name, price + 2 FROM products;

NAME	PRICE+2
Modern Science	21.95
Chemistry	32
Supernova	27.99
Tank War	15.95
Z Files	51.99
2412: The Return	16.95
Space Force 9	15.49
From Another Planet	14.99
Classical Music	12.99
Pop 3	17.99
Creative Yell	16.99
My Front Line	15.49

You can also combine more than one operator in an expression. In the following query, the price column is multiplied by 3, and then 1 is added to the resulting value:

SELECT name, price * 3 + 1 FROM products;

NAME	PRICE*3+1
Modern Science	60.85
Chemistry	91
Supernova	78.97
Tank War	42.85
Z Files	150.97
2412: The Return	45.85
Space Force 9	41.47
From Another Planet	39.97
Classical Music	33.97
Pop 3	48.97
Creative Yell	45.97
My Front Line	41.47

The normal rules of arithmetic operator precedence apply in SQL: multiplication and division are performed first, followed by addition and subtraction. If operators of the same precedence are used, they are performed from left to right. For example, in the expression 10*12/3-1, the first calculation would be 10 multiplied by 12, yielding a result of 120; then 120 would be divided by 3, yielding 40; finally, 1 would be subtracted from 40, yielding 39:

SELECT 10 * 12 / 3 - 1 FROM dual;

10*12/3-1 39 You can also use parentheses () to specify the order of execution for the operators, as in the following:

In this example, the parentheses are used to force calculation of 12/3-1 first, the result of which is then multiplied by 10, yielding 30 as the final answer.

Using Column Aliases

As you've seen, when you select a column from a table, Oracle uses the uppercase version of the column name as the header for the column in the output. For example, when you select the price column, the header in the resulting output is PRICE. When you use an expression, Oracle strips out the spaces and uses the expression as the header. You aren't limited to using the header generated by Oracle; you can provide your own using an *alias*. In the following query, the expression price * 2 is given the alias DOUBLE PRICE:

```
SELECT price * 2 DOUBLE_PRICE
FROM products;
```

If you want to use spaces and preserve the case of your alias text, you must place the text within double quotation marks (""):

You can also use the optional AS keyword before the alias, as shown in the following query:

```
SELECT 10 * (12 / 3 - 1) AS "Computation"
   FROM dual;
   Computation
           30
```

Combining Column Output Using Concatenation

You can combine the column values retrieved by a query using concatenation, which allows you to create more friendly and meaningful output. For example, in the customers table, the first name and last name columns contain the customer name, and in the previous queries the column values were displayed independently. But wouldn't it be nice to combine the first name and last name columns? You can do this using the concatenation operator (||), as shown in the following query; notice that a space character is added after the first name column, and then the last name column is added:

```
SELECT first name || ' ' || last name AS "Customer Name"
   FROM customers;
```

_____ John Brown Cynthia Green Steve White Gail Black Doreen Blue

Customer Name

The first name and last name column values are combined together in the output under the "Customer Name" alias.

Null Values

How does a database represent a value that is unknown? It uses a special value called a *null* value. A null value is not a blank string—it is a distinct value. A null value means the value for the column is unknown.

When you retrieve a column that contains a null value, you see nothing in the output for that column. You saw this (or rather, didn't see it!) in the earlier examples that retrieved rows from the customers table: customer #4 has a null value in the dob column, and customer #5 has a null value in the phone column. In case you missed it, here's the query again:

SELECT * FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

You can also check for null values using IS NULL in a query. In the following example, customer #4 is retrieved because its dob value is null:

SELECT customer_id, first_name, last_name, dob FROM customers WHERE dob IS NULL;

In the next example, customer #5 is retrieved because its phone value is null:

SELECT customer_id, first_name, last_name, phone FROM customers
WHERE phone IS NULL;

```
CUSTOMER_ID FIRST_NAME LAST_NAME PHONE

5 Doreen Blue
```

Since null values don't display anything, how do you tell the difference between a null value and a blank string? The answer is to use the Oracle NVL() built-in function. NVL() returns another value in place of a null. NVL() accepts two parameters: a column (or, more generally, any expression that results in a value) and the value to be returned if the first parameter is null. In the following query, NVL() returns string 'Unknown phone number' when the phone column contains a null value:

SELECT customer_id, first_name, last_name,

NVL(phone, 'Unknown phone number') AS PHONE_NUMBER
FROM customers:

```
CUSTOMER_ID FIRST_NAME LAST_NAME PHONE_NUMBER

1 John Brown 800-555-1211
2 Cynthia Green 800-555-1212
3 Steve White 800-555-1213
4 Gail Black 800-555-1214
5 Doreen Blue Unknown phone number
```

You can also use NVL() to convert null numbers and dates. In the following query, NVL() returns the date 01-JAN-2000 when the dob column contains a null value:

SELECT customer_id, first_name, last_name, NVL(dob, '01-JAN-2000') AS DOB FROM customers;

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB

1 John Brown 01-JAN-65
2 Cynthia Green 05-FEB-68
```

3 Steve	White	16-MAR-71
4 Gail	Black	01-JAN-00
5 Doreen	Blue	20-MAY-70

Notice that customer #4's dob is now displayed as 01-JAN-00.

Displaying Distinct Rows

Suppose you wanted to get the list of customers who purchased products from our imaginary store. You can get that list using the following query, which retrieves the customer id column from the purchases table:

SELECT customer id FROM purchases;

CUSTOMER_ID
1
2
3
4
1
2
3
4
3

The customer id column contains the IDs of customers who purchased a product. As you can see from the output returned by the query, some customers made more than one purchase and therefore appear twice. Wouldn't it be great if you could throw out the duplicate rows that contain the same customer ID? You do this using the DISTINCT keyword. In the following query, DISTINCT is used to suppress the duplicate rows:

SELECT DISTINCT customer id FROM purchases;

```
CUSTOMER ID
```

From this list, it's easy to see that customers #1, #2, #3, and #4 made purchases; the duplicate rows are suppressed.

Comparing Values

The following table lists the operators you can use to compare values:

Operator	Description
=	Equal
<> or !=	Not equal
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
ANY	Compares one value with any value in a list
SOME	Identical to the ANY operator; you should use ANY rather than SOME because ANY is more widely used and readable.
ALL	Compares one value with all values in a list

The following query uses the not equal (<>) operator in the WHERE clause to retrieve the rows from the customers table whose customer id is not equal to 2:

SELECT * FROM customers WHERE customer id <> 2;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

The next query uses the > operator to retrieve the product_id and name columns from the products table where the product id column is greater than 8:

```
SELECT product_id, name
FROM products
WHERE product_id > 8;

PRODUCT_ID NAME

9 Classical Music
10 Pop 3
11 Creative Yell
12 My Front Line
```

The following query uses the ROWNUM pseudo column and the <= operator to retrieve the first 3 rows from the products table:

```
SELECT ROWNUM, product_id, name
FROM products
WHERE ROWNUM <= 3;
```

ROWNUM	PRODUCT_ID	NAME
1	1	Modern Science
2	2	Chemistry
3	3	Supernova

You use the ANY operator in a WHERE clause to compare a value with any of the values in a list. You must place an =, <>, <, >, <=, or >= operator before ANY. The following query uses ANY to retrieve rows from the customers table where the value in the customer id column is greater than any of the values 2, 3, or 4:

SELECT *

FROM customers WHERE customer id > ANY (2, 3, 4);

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

You use the ALL operator in a WHERE clause to compare a value with all of the values in a list. You must place an =, <>, <, >, <=, or >= operator before ALL. The following query uses ALL to retrieve rows from the customers table where the value in the customer id column is greater than all of the values 2, 3, and 4:

SELECT *

FROM customers WHERE customer id > ALL (2, 3, 4);

```
CUSTOMER ID FIRST NAME LAST NAME DOB PHONE
______ ____
     5 Doreen Blue 20-MAY-70
```

Only customer #5 is returned because 5 is greater than 2, 3, and 4.

Using the SQL Operators

The SQL operators allow you to limit rows based on pattern matching of strings, lists of values, ranges of values, and null values. The SQL operators are listed in the following table:

Operator	Description
LIKE	Matches patterns in strings
IN	Matches lists of values
BETWEEN	Matches a range of values
IS NULL	Matches null values
IS NAN	Matches the NAN special value, which means "not a number" (introduced in Oracle Database 10 <i>g</i>)
IS INFINITE	Matches infinite BINARY_FLOAT and BINARY_DOUBLE values (introduced in Oracle Database 10 <i>g</i>)

You can also use NOT to reverse the meaning of an operator:

- NOT LIKE
- NOT IN
- NOT BETWEEN
- IS NOT NULL
- IS NOT NAN
- IS NOT INFINITE

You'll learn about the LIKE, IN, and BETWEEN operators in the following sections.

Using the LIKE Operator

You use the LIKE operator in a WHERE clause to search a string for a pattern. You specify patterns using a combination of normal characters and the following two wildcard characters:

- Underscore (_) Matches one character in a specified position
- Percent (%) Matches any number of characters beginning at the specified position

For example, consider the following pattern:

'_०%'

The underscore (_) matches any one character in the first position, the o matches an o character in

the second position, and the percent (%) matches any characters following the o character. The following query uses the LIKE operator to search the first_name column of the customers table for the pattern '_o%':

SELECT *

FROM customers
WHERE first name LIKE ' o%';

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB PHONE

1 John Brown 01-JAN-65 800-555-1211
5 Doreen Blue 20-MAY-70
```

As you can see from the results, two rows are returned, because the strings John and Doreen both have o as the second character.

The next query uses NOT LIKE to get the rows not retrieved by the previous query:

SELECT *

```
FROM customers
WHERE first_name NOT LIKE '_o%';
```

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214

If you need to search for actual underscore or percent characters in a string, you can use the ESCAPE option to identify those characters. For example, consider the following pattern:

```
'%\%%' ESCAPE '\'
```

10% off Chemistry 20% off Creative Yell 15% off My Front Line

The character after the ESCAPE tells the database how to differentiate between characters to search for and wildcards, and in the example the backslash character (\) is used. The first % is treated as a wildcard and matches any number of characters; the second % is treated as an actual character to search for; the third % is treated as a wildcard and matches any number of characters. The following query uses the promotions table, which contains the details for products being discounted by the store (you'll learn more about this table later in this book). The guery uses the LIKE operator to search the name column of the promotions table for the pattern '%\%%' ESCAPE '\':

```
SELECT name
    FROM promotions
    WHERE name LIKE '%\%%' ESCAPE '\';
    NAME
    _____
    10% off Z Files
    20% off Pop 3
    30% off Modern Science
    20% off Tank War
```

As you can see from the results, the query returns the rows whose names contain a percentage character.

Using the IN Operator

You may use the IN operator in a WHERE clause to retrieve the rows whose column value is in a list. The following query uses IN to retrieve rows from the customers table where the customer id is 2, 3, or 5:

```
SELECT *
     FROM customers
    WHERE customer_id IN (2, 3, 5);
```

```
CUSTOMER ID FIRST NAME LAST NAME DOB
                                                               PHONE
             2 Cynthia Green 05-FEB-68 800-555-1212
3 Steve White 16-MAR-71 800-555-1213
5 Doreen Blue 20-MAY-70
```

NOT IN retrieves the rows not retrieved by IN:

SELECT *

FROM customers
WHERE customer id NOT IN (2, 3, 5);

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB PHONE

1 John Brown 01-JAN-65 800-555-1211
4 Gail Black 800-555-1214
```

One important point to be aware of is that NOT IN returns false if a value in the list is null. This is illustrated by the following query, which doesn't return any rows because null is included in the list:

SELECT *

FROM customers
WHERE customer id NOT IN (2, 3, 5, NULL);

no rows selected



CAUTION

NOT IN returns false if a value in the list is null. This is important because, since you can use any expression in the list and not just literal values, it can be difficult to spot when a null value occurs. Consider using NVL() with expressions that might return a null value.

Using the BETWEEN Operator

You may use the BETWEEN operator in a WHERE clause to retrieve the rows whose column value is in a specified range. The range is *inclusive*, which means the values at both ends of the range are included. The following query uses BETWEEN to retrieve the rows from the customers table where the customer id is between 1 and 3:

SELECT *

FROM customers

WHERE customer_id BETWEEN 1 AND 3;

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB PHONE

1 John Brown 01-JAN-65 800-555-1211
2 Cynthia Green 05-FEB-68 800-555-1212
3 Steve White 16-MAR-71 800-555-1213
```

NOT BETWEEN retrieves the rows not retrieved by BETWEEN:

SELECT *

FROM customers

WHERE customer id NOT BETWEEN 1 AND 3;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

Using the Logical Operators

The logical operators allow you to limit rows based on logical conditions. The logical operators are listed in the following table:

Operator	Description
x AND y	Returns true when both x and y are true
x OR y	Returns true when either x or y is true
NOT x	Returns true if x is false, and returns false if x is true

The following query illustrates the use of the AND operator to retrieve the rows from the customers table where both of the following conditions are true:

- The dob column is greater than January 1, 1970.
- The customer id column is greater than 3.

SELECT *

FROM customers WHERE dob > '01-JAN-1970' AND customer id > 3;

```
CUSTOMER ID FIRST NAME LAST NAME DOB PHONE
______ ____
     5 Doreen Blue 20-MAY-70
```

The next query illustrates the use of the OR operator to retrieve rows from the customers table where either of the following conditions is true:

- The dob column is greater than January 1, 1970.
- The customer id column is greater than 3.

SELECT *

FROM customers WHERE dob > '01-JAN-1970' OR customer_id > 3;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

You can also use AND and OR to combine expressions in a WHERE clause, as you'll see in the following section.

Operator Precedence

If you combine AND and OR in the same expression, the AND operator takes precedence over the OR operator ("takes precedence over" means it's executed first). The comparison operators take precedence over AND. Of course, you can override the default precedence by using parentheses to indicate the order in which you want to execute the expressions.

The following example retrieves the rows from the customers table where *either* of the following two conditions is true:

- The dob column is greater than January 1, 1970.
- The customer id column is less than 2 and the phone column has 1211 at the end.

SELECT *

```
FROM customers
WHERE dob > '01-JAN-1970'
OR customer_id < 2
AND phone LIKE '%1211';
```

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
3	Steve	White	16-MAR-71	800-555-1213
5	Doreen	Blue	20-MAY-70	

As mentioned earlier, AND takes precedence over OR, so you can think of the WHERE clause in the previous query as follows:

```
dob > '01-JAN-1970' OR (customer id < 2 AND phone LIKE '%1211')
```

Therefore, customers #1, #3, and #5 are returned by the guery.

Sorting Rows Using the ORDER BY Clause

You use the ORDER BY clause to sort the rows retrieved by a query. The ORDER BY clause may specify one or more columns on which to sort the data; also, the ORDER BY clause must follow the FROM clause or the WHERE clause (if a WHERE clause is supplied).

The following query uses ORDER BY to sort the rows retrieved from the customers table by the last name:

SELECT *

```
FROM customers
ORDER BY last name;
```

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213

By default, ORDER BY sorts the columns in ascending order (lower values appear first). You can use the DESC keyword to sort the columns in descending order (higher values appear first). You can also use the ASC keyword to explicitly specify an ascending sort—as I mentioned, ascending order is the default, but you can still specify it if you want to make it clear what the order is for the sort.

The next query uses ORDER BY to sort the rows retrieved from the customers table by ascending first name and descending last name:

SELECT * FROM customers ORDER BY first name ASC, last name DESC;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
2	Cynthia	Green	05-FEB-68	800-555-1212
5	Doreen	Blue	20-MAY-70	
4	Gail	Black		800-555-1214
1	John	Brown	01-JAN-65	800-555-1211
3	Steve	White	16-MAR-71	800-555-1213

You can also use a column position number in the ORDER BY clause to indicate which column to sort: Use 1 to sort by the first column selected, 2 to sort by the second column selected, and so on. In the following query, column 1 (the customer id column) is used to sort the rows:

SELECT customer id, first name, last name FROM customers ORDER BY 1:

CUSTOMER_ID	FIRST_NAME	LAST_NAME
1	John	Brown
2	Cynthia	Green
3	Steve	White
4	Gail	Black
5	Doreen	Blue

Because the customer id column is in position 1 after the SELECT keyword, it is the column used in the sort.

Performing SELECT Statements That Use Two Tables

Database schemas have more than one table, with those tables storing different data. For example, the store schema has tables that store information on customers, products, employees, and so on. Up to now, all the gueries shown in this book retrieve rows from only one table. In the real world, you will often need to get information from multiple tables; for example, you might need to get the name of a product and the name of its product type. In this section, you'll learn how to perform queries that use two tables; later, you'll see queries that use more than two tables.

Let's return to the example where you want to get the name of product #3 and its product type. The name of the product is stored in the name column of the products table, and the name of the product type is stored in the name column of the product types table. The

products and product_types tables are related to each other via the foreign key column product_type_id; the product_type_id column (the foreign key) of the products table points to the product type id column (the primary key) of the product types table.

The following query retrieves the name and product_type_id columns from the products table for product #3:

The next query retrieves the name column from the product_types table for the product type id of 2:

```
SELECT name

FROM product_types

WHERE product_type_id = 2;

NAME

-----
Video
```

From this, you know that product #3 is a video. Nothing new so far, but what you really want is to retrieve the product name and its product type name using one query. You do this using a *table join* in the query. To join two tables in a query, you include both tables in the query's FROM clause and include the related columns from each table in the WHERE clause.

For our example query, the FROM clause becomes

```
FROM products, product types
```

And the WHERE clause is

```
WHERE products.product_type_id = product_types.product_type_id
AND products.product id = 3;
```

The join is the first condition in the WHERE clause (products.product_type_id = product_types.product_type_id); typically, the columns used in the join are a primary key from one table and a foreign key from the other table. Because the equality operator (=) is used in the join condition, the join is known as an *equijoin*. The second condition in the WHERE clause (products.product id = 3) gets product #3, the product we are interested in viewing.

You'll notice the tables as well as their columns are included in the WHERE clause. This is because there is a product_type_id column in both the products and product_types tables, and you need to tell the database which table the column you want to use is in. If the columns had different names you could omit the table names, but you should always include them to make it clear where the columns come from.

The SELECT clause for the query is

```
SELECT products.name, product types.name
```

Notice the tables and their columns are specified again.

Putting everything together, the completed query is

```
SELECT products.name, product_types.name
    FROM products, product types
    WHERE products.product type id = product types.product type id
    AND products.product id = 3;
   NAME
    ______
                             Video
```

Perfect! This is exactly what we wanted: the name of the product and the name of the product type. The next query gets all the products and orders them by the products.name column:

```
SELECT products.name, product types.name
     FROM products, product types
    WHERE products.product type id = product types.product type id
    ORDER BY products.name;
```

NAME	NAME
2412: The Return	Video
Chemistry	Book
Classical Music	CD
Creative Yell	CD
From Another Planet	DVD
Modern Science	Book
Pop 3	CD
Space Force 9	DVD
Supernova	Video
Tank War	Video
Z Files	Video

Notice the product with the name "My Front Line" is missing from the results. The product type id for this product row is null, and the join condition does not return the row. You'll see how to include this row later in the section "Outer Joins."

The join syntax you've seen in this section uses Oracle's syntax for joins, which is based on the American National Standards Institute (ANSI) SQL/86 standard. With the introduction of Oracle Database 9i, the database also implements the ANSI SQL/92 standard syntax for joins; you'll see this new syntax later in the section "Performing Joins Using the SQL/92 Syntax." You should use the SQL/92 standard in your queries when working with Oracle Database 9i and above, and you should use SQL/86 queries only when you're using Oracle Database 8i and below.

Using Table Aliases

In the previous section you saw the following query:

```
SELECT products.name, product types.name
     FROM products, product types
     WHERE products.product type id = product types.product type id
     ORDER BY products.name;
```

Notice that the products and product_types table names are used in the SELECT and WHERE clauses. Repeating the table names is redundant typing. A better way is to define table aliases in the FROM clause and then use the aliases when referencing the tables elsewhere in the query. For example, the following query uses the alias p for the products table and pt for the product_types table; notice the table aliases are specified in the FROM clause, and the aliases are placed before the columns in the rest of the query:

```
SELECT p.name, pt.name

FROM products p, product_types pt

WHERE p.product_type_id = pt.product_type_id

ORDER BY p.name;
```



TIP

Table aliases also make your queries more readable, especially when you start writing longer queries that use many tables.

Cartesian Products

If a join condition is missing, you will end up joining all rows from one table with all the rows from the other table; this result set is known as a *Cartesian product*. When this occurs, you will end up with a lot of rows being returned by the query. For example, assume you had one table containing 50 rows and a second table containing 100 rows. If you select columns from those two tables without a join, you would get 5,000 rows returned. This result happens because each row from table 1 would be joined to each row in table 2, which would yield a total of 50 rows multiplied by 100 rows, or 5,000 rows.

The following example shows a subset of the rows for a Cartesian product between the product types and products tables:

SELECT pt.product_type_id, p.product_id FROM product types pt, products p;

PRODUCT_TYPE_ID	PRODUCT_ID
1	1
1	2
1	3
1	4
1	5
5	8
5	9
5	10
5	11
5	12

60 rows selected.

A total of 60 rows are selected because the product types and products tables contain 5 and 12 rows, respectively, and 5*12 = 60.

Performing SELECT Statements That Use More than Two Tables

Joins can be used to connect any number of tables. You use the following formula to calculate the number of joins you will need in your WHERE clause:

Number of joins = the number of tables used in the guery -1

For example, the following query uses two tables and one join:

```
SELECT p.name, pt.name
   FROM products p, product types pt
   WHERE p.product type id = pt.product type id
   ORDER BY p.name;
```

Let's consider a more complicated example that uses four tables—and therefore requires three joins. Let's say you want to see the following information:

- The purchases each customer has made (comes from the purchases table)
- The customer's first and last name (comes from the customers table)
- The name of the product they purchased (comes from the products table)
- The name of the product type (comes from the product types table)

In order to view this information, you need to query the customers, purchases, products, and product types tables, and the joins use the foreign key relationships between these tables. The following list shows the required joins:

- 1. To get the customer who made the purchase, join the customers and purchases tables using the customer id columns (customers.customer id = purchases .customer id).
- 2. To get the product purchased, join the products and purchases tables using the product id columns (products.product id = purchases.product id).
- To get the product type name for the product, join the products and product types tables using the product type id columns (products.product type id = product types.product type id).

The following query uses these joins; notice I've used table aliases and renamed the headings for the product name to PRODUCT and the product type name to TYPE:

```
SELECT c.first name, c.last name, p.name AS PRODUCT, pt.name AS TYPE
   FROM customers c, purchases pr, products p, product types pt
   WHERE c.customer id = pr.customer id
   AND p.product id = pr.product id
   AND p.product type id = pt.product type id
   ORDER BY p.name;
```

FIRST_NAME	LAST_NAME	PRODUCT	TYPE
John	Brown	Chemistry	Book
Cynthia	Green	Chemistry	Book
Steve	White	Chemistry	Book
Gail	Black	Chemistry	Book
John	Brown	Modern Science	Book
Cynthia	Green	Modern Science	Book
Steve	White	Modern Science	Book
Gail	Black	Modern Science	Book
Steve	White	Supernova	Video

The multi-table queries you've seen so far use the equality operator (=) in the join conditions; these joins are known as *equijoins*. As you'll see in the next section, the equijoin is not the only join you can use.

Join Conditions and Join Types

In this section, you'll explore join conditions and join types that allow you to create more advanced queries.

There are two types of join conditions, which are based on the operator you use in your join:

- **Equijoins** use the equality operator (=). You've already seen examples of equijoins.
- **Non-equijoins** use an operator other than the equality operator, such as <, >, BETWEEN, and so on. You'll see an example of a non-equijoin shortly.

There are also three different types of joins:

- Inner joins return a row *only* when the columns in the join contain values that satisfy the join condition. This means that if a row has a null value in one of the columns in the join condition, that row isn't returned. The examples you've seen so far have been inner joins.
- Outer joins return a row even when one of the columns in the join condition contains a null value.
- **Self joins** return rows joined on the same table.

You'll learn about non-equijoins, outer joins, and self joins next.

Non-equijoins

A non-equijoin uses an operator other than the equality operator (=) in the join. These operators are not-equal (<>), less than (<), greater than (>), less than or equal to (<=), greater than or equal to (>=), LIKE, IN, and BETWEEN. It's pretty rare to encounter situations where you need to use a non-equijoin, but I have come across a few occasions when I've needed to use one; on those occasions I've had to use the BETWEEN operator.

For example, let's say you want to get the salary grades for the employees. First, the following query retrieves the salary grades from the salary grades table:

SELECT * FROM salary grades;

SALARY_GRADE_ID	LOW_SALARY	HIGH_SALARY
1	1	250000
2	250001	500000
3	500001	750000
4	750001	999999

The next query uses a non-equijoin to retrieve the salary and salary grades for the employees; the salary grade is determined using the BETWEEN operator:

SELECT e.first name, e.last name, e.title, e.salary, sq.salary grade id FROM employees e, salary grades sg WHERE e.salary BETWEEN sg.low salary AND sg.high salary ORDER BY salary grade id;

FIRST_NAME	LAST_NAME	TITLE	SALARY	SALARY_GRADE_ID
Fred	Hobbs	Salesperson	150000	1
Susan	Jones	Salesperson	500000	2
Ron	Johnson	Sales Manager	600000	3
James	Smith	CEO	800000	4

In this query, the BETWEEN operator returns true if the employee's salary is between the low salary and high salary for the salary grade; when true is returned, the salary grade found is the salary grade for the employee. For example, Fred Hobbs' salary is \$150,000; this is between the low salary of \$1 and the high salary of \$250,000 in the salary grades table for the salary grade id of 1; therefore, Fred Hobbs' salary grade is 1. Similarly, Susan Jones' salary is \$500,000; this is between the low salary of \$250,001 and the high salary of \$500,000 for the salary grade ID of 2; therefore, Susan Jones' salary grade is 2. Ron Johnson and James Smith have salary grades of 3 and 4 respectively.

Outer Joins

An outer join retrieves a row even when one of the columns in the join contains a null value. You perform an outer join by supplying the outer join operator in the join condition; the Oracle proprietary outer join operator is a plus character in parentheses (+).

Let's take a look at an example. Remember the guery earlier that didn't show the "My Front Line" product because its product type id is null? You can use an outer join to get that row. In the following query, notice that the Oracle outer join operator (+) is on the opposite side of the product type id column in the product table (this is the column that contains the null value):

```
SELECT p.name, pt.name
   FROM products p, product types pt
   WHERE p.product type id = pt.product type id (+)
   ORDER BY p.name;
```

NAME	NAME
2412: The Return	Video
Chemistry	Book
Classical Music	CD
Creative Yell	CD
From Another Planet	DVD
Modern Science	Book
My Front Line	
Pop 3	CD
Space Force 9	DVD
Supernova	Video
Tank War	Video
Z Files	Video

Notice that "My Front Line"—the product with the null product_type_id—is now retrieved, even though its product type id is null.



NOTE

You can place the outer join operator on either side of the join operator, but you always place it on the opposite side of the column that contains the null value.

The following query returns the same results as the previous one, but notice that the column with the null value (pt.product_type_id) and the Oracle outer join operator are on the left of the equality operator:

```
SELECT p.name, pt.name
FROM products p, product_types pt
WHERE pt.product_type_id (+) = p.product_type_id
ORDER BY p.name;
```

Left and Right Outer Joins

Outer joins can be split into two types:

- Left outer joins
- Right outer joins

To understand the difference between left and right outer joins, consider the following syntax:

```
SELECT ...
FROM table1, table2
```

Assume the tables are to be joined on table1.column1 and table2.column2. Also, assume table1 contains a row with a null value in column1. To perform a left outer join, the WHERE clause is

```
WHERE table1.column1 = table2.column2 (+);
```



In a left outer join, the outer join operator is actually on the right of the equality operator.

Next, assume table2 contains a row with a null value in column2. To perform a right outer join, you switch the position of the Oracle outer join operator to the *left* of the equality operator, and the WHERE clause becomes

WHERE table1.column1 (+) = table2.column2;



NOTE

As you'll see, if table1 and table2 both contain rows with null values, you get different results depending on whether you use a left or right outer join.

Let's take a look at some concrete examples to make left and right outer joins clearer.

An Example of a Left Outer Join The following query uses a left outer join; notice that the Oracle outer join operator appears on the *right* of the equality operator:

```
SELECT p.name, pt.name
    FROM products p, product types pt
    WHERE p.product type id = pt.product type id (+)
    ORDER BY p.name;
```

NAME	NAME
2412: The Return	Video
Chemistry	Book
Classical Music	CD
Creative Yell	CD
From Another Planet	DVD
Modern Science	Book
My Front Line	
Pop 3	CD
Space Force 9	DVD
Supernova	Video
Tank War	Video
Z Files	Video

Notice all the rows from the products table are retrieved, including the "My Front Line" row, which has a null value in the product type id column.

An Example of a Right Outer Join The product types table contains a type of product not referenced in the products table (there are no magazines in the products table); the magazine product type appears at the end of the following example:

SELECT * FROM product_types;

37736

```
PRODUCT_TYPE_ID NAME

1 Book
2 Video
3 DVD
4 CD
5 Magazine
```

You can retrieve the magazine in a join on the products and product_types tables by using a right outer join, as shown in the following query; notice that the Oracle outer join operator appears on the *left* of the equality operator:

37336

```
SELECT p.name, pt.name

FROM products p, product_types pt

WHERE p.product_type_id (+) = pt.product_type_id

ORDER BY p.name;
```

NAME	NAME
2412: The Return	Video
Chemistry	Book
Classical Music	CD
Creative Yell	CD
From Another Planet	DVD
Modern Science	Book
Pop 3	CD
Space Force 9	DVD
Supernova	Video
Tank War	Video
Z Files	Video
	Magazine

Limitations on Outer Joins

There are limitations when using outer joins, and you will learn about some of them in this section. You may only place the outer join operator on one side of the join (not both). If you try to place the Oracle outer join operator on both sides, you get an error, as shown in the following example:

You cannot use an outer join condition with the IN operator:

```
SQL> SELECT p.name, pt.name

2 FROM products p, product_types pt

3 WHERE p.product_type_id (+) IN (1, 2, 3, 4);

WHERE p.product_type_id (+) IN (1, 2, 3, 4)

*

ERROR at line 3:

ORA-01719: outer join operator (+) not allowed in operand of OR or IN
```

You cannot use an outer join condition with another join using the OR operator:

```
SQL> SELECT p.name, pt.name
       2 FROM products p, product types pt
       3 WHERE p.product type id (+) = pt.product type id
       4 OR p.product type id = 1;
     WHERE p.product type id (+) = pt.product type id
     ERROR at line 3:
     ORA-01719: outer join operator (+) not allowed in operand of OR or IN
```



These are the commonly encountered limitations when using the outer join operator. For all the limitations, read the Oracle Database SQL Reference manual from Oracle Corporation.

Self Joins

A self join is a join made on the same table. To perform a self join, you must use a different table alias to identify each reference to the table in the query. Let's consider an example. The employees table has a manager id column that contains the employee id of the manager for each employee; if the employee has no manager, then the manager id is null. The employees table contains the following rows:

EMPLOYEE_ID	MANAGER_ID	FIRST_NAME	LAST_NAME	TITLE	SALARY
1		James	Smith	CEO	800000
2	1	Ron	Johnson	Sales Manager	600000
3	2	Fred	Hobbs	Salesperson	150000
4	2	Susan	Jones	Salesperson	500000

James Smith—the CEO—has a null value for the manager id, meaning that he doesn't have a manager. Susan Jones and Fred Hobbs are managed by Ron Johnson, and Ron Johnson is managed by James Smith.

You can use a self join to display the names of each employee and their manager. In the following query, the employees table is referenced twice, using two aliases w and m. The w alias is used to get the worker name, and the m alias is used to get the manager name. The self join is made between w.manager id and m.employee id:

```
SELECT w.first_name || ' ' || w.last_name || ' works for '||
      m.first name || ' ' || m.last name
    FROM employees w, employees m
    WHERE w.manager id = m.employee id
    ORDER BY w.first_name;
    W.FIRST NAME | | ' ' | | W.LAST NAME | | 'WORKSFOR' | | M.FIRST NA
    _____
    Fred Hobbs works for Ron Johnson
    Ron Johnson works for James Smith
    Susan Jones works for Ron Johnson
```

Because James Smith's manager id is null, the join condition does not return the row.

You can perform outer joins in combination with self joins. In the following query, an outer join is used with the self join shown in the previous example to retrieve the row for James Smith. You'll notice the use of the NVL() function to provide a string indicating that James Smith works for the shareholders (he's the CEO, so he reports to the store's shareholders):

```
SELECT w.last_name || 'works for '||

NVL(m.last_name, 'the shareholders')

FROM employees w, employees m

WHERE w.manager_id = m.employee_id (+)

ORDER BY w.last_name;

W.LAST_NAME||'WORKSFOR'||NVL(M.LAST_N

Hobbs works for Johnson

Johnson works for Smith

Jones works for Johnson

Smith works for the shareholders
```

Performing Joins Using the SQL/92 Syntax

The joins you've seen so far use Oracle's syntax for joins, which is based on the ANSI SQL/86 standard. With the introduction of Oracle Database 9*i*, the database implements the ANSI SQL/92 standard syntax for joins, and you should use SQL/92 in your queries. You'll see how to use SQL/92 in this section, including its use in avoiding unwanted Cartesian products.



NOTE

You can visit the ANSI website at www.ansi.org.

Performing Inner Joins on Two Tables Using SQL/92

Earlier, you saw the following query, which uses the SQL/86 standard for performing an inner join:

```
SELECT p.name, pt.name

FROM products p, product_types pt

WHERE p.product_type_id = pt.product_type_id

ORDER BY p.name;
```

SQL/92 introduces the INNER JOIN and ON clauses for performing an inner join. The following example rewrites the previous query using the INNER JOIN and ON clauses:

```
SELECT p.name, pt.name

FROM products p INNER JOIN product_types pt

ON p.product_type_id = pt.product_type_id

ORDER BY p.name;
```

You can also use non-equijoin operators with the ON clause. Earlier, you saw the following query, which uses the SQL/86 standard for performing a non-equijoin:

```
SELECT e.first_name, e.last_name, e.title, e.salary, sg.salary_grade_id FROM employees e, salary_grades sg
WHERE e.salary BETWEEN sg.low_salary AND sg.high_salary
ORDER BY salary_grade_id;
```

The following example rewrites this guery to use the SQL/92 standard:

```
SELECT e.first name, e.last name, e.title, e.salary, sq.salary grade id
    FROM employees e INNER JOIN salary grades sq
    ON e.salary BETWEEN sg.low salary AND sg.high salary
    ORDER BY salary grade id;
```

Simplifying Joins with the USING Keyword

SQL/92 allows you to further simplify the join condition with the USING clause, but with the following limitations:

- The guery must use an equijoin.
- The columns in the equijoin must have the same name.

Most of the joins you'll perform will be equijoins, and if you always use the same name as the primary key for your foreign keys, you'll satisfy these limitations.

The following query uses the USING clause instead of ON:

```
SELECT p.name, pt.name
    FROM products p INNER JOIN product types pt
    USING (product type id);
```

If you wanted to retrieve the product type id, you must provide only this column name on its own without a table name or alias in the SELECT clause, as for example:

```
SELECT p.name, pt.name, product type id
    FROM products p INNER JOIN product types pt
    USING (product type id);
```

If you try to provide a table alias with the column, such as p.product type id for example, you'll get an error:

```
SQL> SELECT p.name, pt.name, p.product type id
      2 FROM products p INNER JOIN product types pt
      3 USING (product_type_id);
    SELECT p.name, pt.name, p.product type id
    ERROR at line 1:
    ORA-25154: column part of USING clause cannot have qualifier
```

Also, you only use the column name on its own within the USING clause. For example, if you specify USING (p.product type id) in the previous query instead of USING (product type id), you'll get an error:

```
SQL> SELECT p.name, pt.name, p.product_type_id
      2 FROM products p INNER JOIN product types pt
      3 USING (p.product type id);
    USING (p.product type id)
    ERROR at line 3:
    ORA-01748: only simple column names allowed here
```



CAUTION

Don't use a table name or alias when referencing columns used in a USING clause. You'll get an error if you do.

Performing Inner Joins on More than Two Tables Using SQL/92

Earlier you saw the following SQL/86 query, which retrieves rows from the customers, purchases, products, and product types tables:

```
SELECT c.first_name, c.last_name, p.name AS PRODUCT, pt.name AS TYPE FROM customers c, purchases pr, products p, product_types pt WHERE c.customer_id = pr.customer_id

AND p.product_id = pr.product_id

AND p.product_type_id = pt.product_type_id

ORDER BY p.name;
```

The following example rewrites this query using SQL/92; notice how the foreign key relationships are navigated using multiple INNER JOIN and USING clauses:

```
SELECT c.first_name, c.last_name, p.name AS PRODUCT, pt.name AS TYPE
FROM customers c INNER JOIN purchases pr
USING (customer_id)
INNER JOIN products p
USING (product_id)
INNER JOIN product_types pt
USING (product_type_id)
ORDER BY p.name;
```

Performing Inner Joins on Multiple Columns Using SQL/92

If your join uses more than one column from the two tables, you provide those columns in your ON clause and use the AND operator. For example, let's say you have two tables named table1 and table2 and you want to join these tables using columns named column1 and column2 in both tables. Your query would be

```
FROM table1 INNER JOIN table2
ON table1.column1 = table2.column1
AND table1.column2 = table2.column2;
```

You can further simplify your query with the USING clause, but only if you're performing an equijoin and the column names are identical. For example, the following query rewrites the previous example with the USING clause:

```
FROM table1 INNER JOIN table2
USING (column1, column2);
```

Performing Outer Joins Using SQL/92

Earlier you saw how to perform outer joins using the outer join operator (+), which is Oracle proprietary syntax. SQL/92 uses a different syntax for performing outer joins. Instead of using (+), you specify the type of join in the FROM clause using the following syntax:

```
FROM table1 { LEFT | RIGHT | FULL } OUTER JOIN table2
```

where

- table1 and table2 are the tables you want to join.
- LEFT means you want to perform a left outer join.
- RIGHT means you want to perform a right outer join.
- FULL means you want to perform a full outer join. A full outer join uses all rows in table1 and table2, including those that have null values in the columns used in the join. You cannot directly perform a full outer join using the (+) operator.

You'll see how to perform left, right, and full outer joins using the SQL/92 syntax in the following sections.

Performing Left Outer Joins Using SQL/92

Earlier you saw the following query, which performed a left outer join using the Oracle proprietary (+) operator:

```
SELECT p.name, pt.name
    FROM products p, product types pt
    WHERE p.product type id = pt.product type id (+)
    ORDER BY p.name;
```

The next example rewrites this query using the SQL/92 LEFT OUTER JOIN keywords:

```
SELECT p.name, pt.name
    FROM products p LEFT OUTER JOIN product types pt
    USING (product type id)
    ORDER BY p.name;
```

Performing Right Outer Joins Using SQL/92

Earlier you saw the following query, which performed a right outer join using the Oracle proprietary (+) operator:

```
SELECT p.name, pt.name
    FROM products p, product types pt
    WHERE p.product type id (+) = pt.product type id
    ORDER BY p.name;
```

The next example rewrites this query using the SQL/92 RIGHT OUTER JOIN keywords:

```
SELECT p.name, pt.name

FROM products p RIGHT OUTER JOIN product_types pt
USING (product_type_id)

ORDER BY p.name;
```

Performing Full Outer Joins Using SQL/92

A full outer join uses all rows in the joined tables, including those that have null values in either of the columns used in the join. The following example shows a query that uses the SQL/92 FULL OUTER JOIN keywords:

```
SELECT p.name, pt.name

FROM products p FULL OUTER JOIN product_types pt

USING (product_type_id)

ORDER BY p.name;
```

NAME	NAME
2412: The Return	Video
Chemistry	Book
Classical Music	CD
Creative Yell	CD
From Another Planet	DVD
Modern Science	Book
My Front Line	
Pop 3	CD
Space Force 9	DVD
Supernova	Video
Tank War	Video
Z Files	Video
	Magazine

Notice that both "My Front Line" from the products table and "Magazine" from the product_types table are returned.

Performing Self Joins Using SQL/92

The following example uses SOL/86 to perform a self join on the employees table:

```
SELECT w.last_name || ' works for ' || m.last_name
FROM employees w, employees m
WHERE w.manager_id = m.employee_id;
```

The next example rewrites this query to use the SQL/92 INNER JOIN and ON keywords:

```
SELECT w.last_name || ' works for ' || m.last_name
FROM employees w INNER JOIN employees m
ON w.manager id = m.employee id;
```

Performing Cross Joins Using SQL/92

Earlier you saw how omitting a join condition between two tables leads to a Cartesian product. By using the SQL/92 join syntax, you avoid inadvertently producing a Cartesian product because you must always provide an ON or USING clause to join the tables—this is a good thing because you usually don't want a Cartesian product.

If you really want a Cartesian product, the SQL/92 standard requires that you explicitly state this in your query using the CROSS JOIN keywords. In the following query, a Cartesian product between the product types and products tables is generated using the CROSS JOIN keywords:

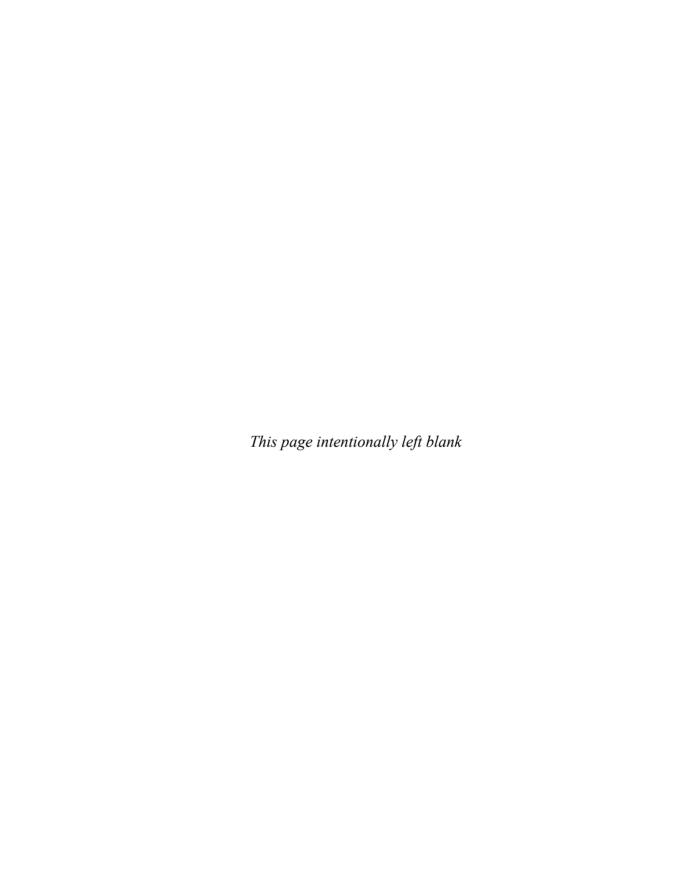
```
SELECT *
FROM product types CROSS JOIN products;
```

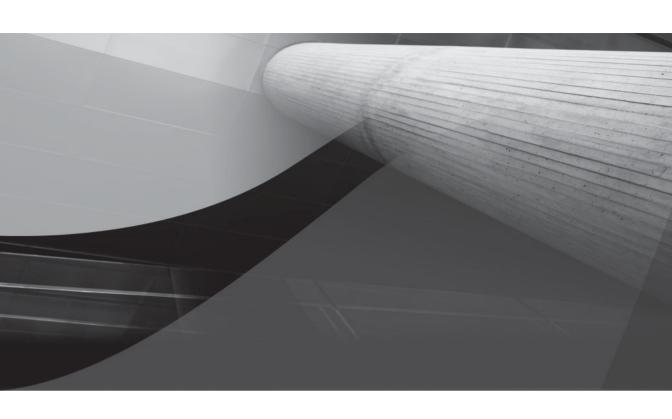
Summary

In this chapter, you have learned the following:

- How to perform single and multiple table queries
- How to select all columns from a table using an asterisk (*) in a query
- How a row identifier (rowid) is used internally by the Oracle database to store the location of a row
- How to perform arithmetic in SQL
- How to use addition and subtraction operators with dates
- How to reference tables and columns using aliases
- How to merge column output using the concatenation operator (||)
- How nulls are used to represent unknown values
- How to display distinct rows using the DISTINCT operator
- How to limit the retrieval of rows using the WHERE clause
- How to sort rows using the ORDER BY clause
- How to perform inner, outer, and self joins using the SQL/86 and SQL/92 syntax

In the next chapter, you'll learn about SQL*Plus.





CHAPTER 3

Using SQL*Plus



n this chapter, you will see how to do the following:

- View the structure of a table.
- Edit an SQL statement.
- Save and run scripts containing SQL statements and SQL*Plus commands.
- Format the results returned by SQL*Plus.
- Use variables in SQL*Plus.
- Create simple reports.
- Get help from SQL*Plus.
- Automatically generate SQL statements.
- Disconnect from a database and exit SQL*Plus.

Let's plunge in and examine how you view the structure of a table.

Viewing the Structure of a Table

Knowing the structure of a table is useful because you can use the information to formulate an SQL statement. For example, you can figure out the columns you want to retrieve in a query. You use the DESCRIBE command to view the structure of a table.

The following example uses DESCRIBE to view the structure of the customers table; notice that the semicolon character (;) is omitted from the end of the command:

SQL> DESCRIBE customers

Name	Null	L?	Type
CUSTOMER_ID	NOT	NULL	NUMBER (38)
FIRST_NAME	NOT	NULL	VARCHAR2(10)
LAST_NAME	NOT	NULL	VARCHAR2(10)
DOB			DATE
PHONE			VARCHAR2(12)

The output from DESCRIBE has three columns that show the structure of the table. These columns are as follows:

- Name lists the names of the columns contained in the table. In the example, you can see the customers table has five columns: customer_id, first_name, last_name, dob, and phone.
- Null? indicates whether the column can store null values. If set to NOT NULL, the column cannot store a null value; if blank, the column can store a null value. In the preceding example, you can see that the customer_id, first_name, and last_name columns cannot store null values, but the dob and phone columns can.

Type indicates the type of the column. In the preceding example, you can see that the type of the customer id column is NUMBER (38) and that the type of the first name is VARCHAR2 (10).

You can save some typing by shortening the DESCRIBE command to DESC (DESC [RIBE]). The following command uses DESC to view the structure of the products table:

SQL> DESC products

Name	Null	L?	Type
PRODUCT_ID	NOT	NULL	NUMBER (38)
PRODUCT_TYPE_ID			NUMBER (38)
NAME	NOT	NULL	VARCHAR2(30)
DESCRIPTION			VARCHAR2(50)
PRICE			NUMBER (5,2)

Editing SQL Statements

As you may have noticed, it becomes tedious to have to repeatedly type similar SQL statements into SQL*Plus. You will be pleased to know SQL*Plus stores your previous SQL statement in a buffer. You can then edit the lines that make up your SQL statement stored in the buffer.

Some of the editing commands are listed in the following table; notice the optional part of each command in square brackets (for example, you can abbreviate the APPEND command to A).

Command	Description
A[PPEND] text	Appends text to the current line.
C[HANGE] /old/new	Changes the text specified by old to new in the current line.
CL[EAR] BUFF[ER]	Clears all lines from the buffer.
DEL	Deletes the current line.
DEL x	Deletes the line specified by the line number \boldsymbol{x} (line numbers start with 1).
L[IST]	Lists all the lines in the buffer.
L[IST] x	Lists line number x.
R[UN]	Runs the statement stored in the buffer. You can also use / to run
or	the statement.
/	
X	Makes the line specified by the line number x the current line.

Let's take a look at some examples of using the SQL*Plus editing commands. The following example shows a query in SQL*Plus:

```
SQL> SELECT customer id, first name, last name
      2 FROM customers
      3 WHERE customer id = 1;
```

SQL*Plus automatically increments the line number when you press ENTER. You can make line 1 the current line by entering 1 at the prompt:

```
SQL> 1

1* SELECT customer id, first name, last name
```

Notice that SQL*Plus displays the current line and the line number. The following example uses APPEND to add ", dob" to the end of the line:

```
SQL> APPEND , dob

1* SELECT customer id, first name, last name, dob
```

The next example uses LIST to show all the lines in the buffer:

```
SQL> LIST

1 SELECT customer_id, first_name, last_name, dob
2 FROM customers
3* WHERE customer id = 1
```

Notice that the current line has been changed to the last line, as indicated by the asterisk character (*). The following example uses CHANGE to replace "customer_id = 1" with "customer id = 2" in this last line:

```
SQL> CHANGE /customer_id = 1/customer_id = 2

3* WHERE customer_id = 2
```

The next example uses RUN to execute the query:

You can also use a forward slash character (/) to run the SQL statement; for example:

SQL> /

```
CUSTOMER_ID FIRST_NAME LAST_NAME DOB

2 Cynthia Green 05-FEB-68
```

Saving, Retrieving, and Running Files

SQL*Plus allows you to save, retrieve, and run scripts containing SQL*Plus commands and SQL statements. You've already seen one example of running an SQL*Plus script: the store_schema.sql script file that was run in Chapter 1.

Some of the file commands are listed in the following table.

Command	Description
SAV[E] filename [{REPLACE APPEND }]	Saves the contents of the SQL*Plus buffer to a file specified by filename. You append the content of the buffer to an existing file using the APPEND option. You overwrite an existing file using the REPLACE option.
GET filename	Retrieves the contents of the file specified by filename into the SQL*Plus buffer.
STA[RT] filename	Retrieves the contents of the file specified by filename into the SQL*Plus buffer and then attempts to run the contents of the buffer.
@ filename	Same as the START command.
ED[IT]	Copies the contents of the SQL*Plus buffer to a file named afiedt.buf and then starts the default editor for the operating system. When you exit the editor, the contents of the edited file are copied to the SQL*Plus buffer.
ED[IT] filename	Same as the EDIT command, but you can specify a file to start editing. You specify the file to edit using the filename parameter.
SPO[OL] filename	Copies the output from SQL*Plus to the file specified by filename.
SPO[OL] OFF	Stops the copying of output from SQL*Plus to the file, then closes that file.

Let's take a look at some examples of using these SQL*Plus commands. If you want to follow along with the examples, go ahead and enter the following query into SQL*Plus:

```
SQL> SELECT customer id, first name, last name
```

2 FROM customers

3 WHERE customer id = 1;

The following example uses SAVE to save the contents of the SQL*Plus buffer to a file named cust query.sql:

SQL> SAVE cust_query.sql Created file cust query.sql



On my computer, the <code>cust_query.sql</code> file is saved in the <code>E:\</code> oracle 11g\product\11.1.0\db 1\BIN directory.

The next example uses GET to retrieve the contents of the cust guery.sql file:

SQL> GET cust_query.sql

- 1 SELECT customer id, first_name, last_name
- 2 FROM customers
- 3* WHERE customer id = 1

The following example runs the query using /:

SQL> /

```
CUSTOMER_ID FIRST_NAME LAST_NAME

1 John Brown
```

The next example uses START to load and run the contents of the <code>cust_query.sql</code> file in one step:

SQL> START cust_query.sql

You can edit the contents of the SQL*Plus buffer using the EDIT command:

SQL> EDIT

The EDIT command starts the default editor for your operating system. On Windows the default editor is Notepad. On Unix and Linux the default editors are vi or emacs, respectively.

Figure 3-1 shows the contents of the SQL*Plus buffer in Notepad. Notice the SQL statement is terminated using a slash character (/) rather than a semicolon.

In your editor, change the WHERE clause to WHERE customer_id = 2, then save and exit the editor (in Notepad, you select File | Exit, then click Yes to save your query). SQL*Plus displays the following output containing your modified query; notice that the WHERE clause has been changed:

```
1 SELECT customer_id, first_name, last_name
2 FROM customers
3* WHERE customer id = 2
```

```
File Edit Format Help

SELECT customer_id, first_name, last_name
FROM customers
WHERE customer_id = 1
```

FIGURE 3-1 Editing the SQL*Plus buffer contents using Notepad

Changing the Default Editor

You can change the default editor using the SQL*Plus DEFINE command:

```
DEFINE EDITOR = 'editor'
```

where editor is the name of your preferred editor.

For example, the following command sets the default editor to vi:

```
DEFINE EDITOR = 'vi'
```

You can also change the default editor SQL*Plus uses by adding the line DEFINE EDITOR = 'editor' to a new file named login.sql, where editor is the name of your preferred editor. You can add any SQL*Plus commands you want to this file. SQL*Plus will check the current directory for a login.sql file and execute it when SQL*Plus starts. If there is no login.sql file in the current directory, then SQL*Plus will check all directories (and their subdirectories) in the SOLPATH environment variable for a login.sql file. On Windows, SQLPATH is defined as a registry entry in HKEY LOCAL MACHINE\SOFTWARE\ ORACLE\oracle home key (where oracle home key is the key for the associated installation of the Oracle database). On my Windows XP computer running Oracle Database 11g, SQLPATH is set to E:\oracle 11g\product\11.1.0\db 1\dbs. On Unix or Linux, there is no default SQLPATH defined, and you will need to add it as an environment variable. For further details on setting up a login.sql file, you can read the SQL*Plus User's Guide and Reference, published by Oracle Corporation.

You run your modified guery using the slash character (/):

SQL> /

```
CUSTOMER ID FIRST NAME LAST NAME
_____
      2 Cynthia Green
```



TIP

In the Oracle Database 11g version of SQL*Plus, you can also scroll through your previously executed statements using the UP and DOWN ARROW keys on the keyboard. Once you've selected a statement, you can use the LEFT and RIGHT ARROW keys to move the cursor to a specific point in the statement.

You use the SPOOL command to copy the output from SQL*Plus to a file. The following example spools the output to a file named cust results.txt, runs the query again, and then turns spooling off by executing SPOOL OFF:

```
SQL> SPOOL cust results.txt
   SOL> /
   CUSTOMER ID FIRST NAME LAST NAME
   _____
          2 Cynthia Green
```

SOL> SPOOL OFF

Feel free to examine the <code>cust_results.txt</code> file; it will contain the previous output between the slash (/) and <code>SPOOL OFF.</code> On my Windows XP computer, the file is stored in <code>E:\oracle_11g\product\11.1.0\db_1\BIN</code>; the directory used is the current directory you are in when you start SQL*Plus.

You can also specify the full directory path where you want the spool file to be written; for example:

SPOOL C:\my files\spools\cust results.txt

Formatting Columns

You use the COLUMN command to format the display of column headings and column data. The simplified syntax for the COLUMN command is as follows:

COL[UMN] {column | alias} [options]

where

- column is the column name.
- alias is the column alias to be formatted. In Chapter 2, you saw that you can "rename" a column using a column alias; you can reference an alias in the COLUMN command.
- options are one or more options to be used to format the column or alias.

There are a number of options you can use with the COLUMN command. The following table shows some of these options.

Option	Description
FOR[MAT] format	Sets the format for the display of the column or alias to the format string.
HEA[DING] heading	Sets the heading of the column or alias to the heading string.
JUS[TIFY] [{ LEFT CENTER RIGHT }]	Places the column output to the left, center, or right.
WRA[PPED]	Wraps the end of a string onto the next line of output. This option may cause individual words to be split across multiple lines.
WOR[D_WRAPPED]	Similar to the WRAPPED option, except individual words are not split across two lines.
CLE[AR]	Clears any formatting of columns (that is, sets the formatting back to the default).

The format string in the previous table may take a number of formatting parameters. The parameters you specify depend on the data stored in your column:

If your column contains characters, you use Ax to format the characters, where x specifies the width for the characters. For example, A12 sets the width to 12 characters.

- If your column contains numbers, you can use any of a variety of number formats, which are shown later in Table 4-4 of Chapter 4. For example, \$99.99 sets the format to a dollar sign, followed by two digits, the decimal point, plus another two digits.
- If your column contains a date, you can use any of the date formats shown later in Table 5-2 of Chapter 5. For example, MM-DD-YYYY sets the format to a two-digit month (MM), a two-digit day (DD), and a four-digit year (YYYY).

Let's take a look at an example. You'll see how to format the output of a guery that retrieves the product id, name, description, and price columns from the products table. The display requirements, format strings, and COLUMN commands are shown in the following table:

Column	Display As	Format	COLUMN Command
product_id	Two digits	99	COLUMN product_id FORMAT 99
name	Thirteen-character word- wrapped strings and set the column heading to PRODUCT_NAME	A13	COLUMN name HEADING PRODUCT_NAME FORMAT A13 WORD_WRAPPED
description	Thirteen-character word- wrapped strings	A13	COLUMN description FORMAT A13 WORD_WRAPPED
price	Dollar symbol, with two digits before and after the decimal point	\$99.99	COLUMN price FORMAT \$99.99

The following example shows the COLUMN commands in SQL*Plus:

- SQL> COLUMN product id FORMAT 99
 - SQL> COLUMN name HEADING PRODUCT NAME FORMAT A13 WORD WRAPPED
 - SQL> COLUMN description FORMAT A13 WORD WRAPPED
 - SQL> COLUMN price FORMAT \$99.99

The next example runs a query to retrieve some rows from the products table; notice the formatting of the columns in the output:

- SQL> SELECT product id, name, description, price
 - 2 FROM products
 - 3 WHERE product id < 6;</pre>

PRODUCT_ID PRODUC	CT_NAME DI	ESCRIPTION	PRICE
1 Modern Science	ce o:	description f modern cience	\$19.95
2 Chemis	4	ntroduction	\$30.00

3	Supernova	A star explodes	\$25.99
4	Tank War	Action movie	\$13.95
PRODUCT_ID	PRODUCT_NAME	DESCRIPTION	PRICE
		about a future war	
5	Z Files	Series on mysterious activities	\$49.99

This output is readable, but wouldn't it be nice if you could display the headings just once at the top? You do that by setting the page size, as you'll see next.

Setting the Page Size

You set the number of lines in a page using the SET PAGESIZE command. This command sets the number of lines that SQL*Plus considers one "page" of output, after which SQL*Plus will display the headings again.

The following example sets the page size to 100 lines using the SET PAGESIZE command and runs the query again using /:

SQL> SET PAGESIZE 100 SQL> /

PRODUCT_ID	PRODUCT_NAME	DESCRIPTION	PRICE
1	Modern Science	A description of modern science	\$19.95
2	Chemistry	Introduction to Chemistry	\$30.00
3	Supernova	A star explodes	\$25.99
4	Tank War	Action movie about a future war	\$13.95
5	Z Files	Series on mysterious activities	\$49.99

Notice that the headings are shown only once, at the top, and the resulting output looks better.



NOTE

The maximum number for the page size is 50,000.

Setting the Line Size

You set the number of characters in a line using the SET LINESIZE command. The following example sets the line size to 50 lines and runs another query:

SQL> SET LINESIZE 50 SQL> SELECT * FROM customers;

CUSTOMER_	ID FIRST_NAME	LAST_NAME	DOB
PHONE			
800-555-1	 1 John 211	Brown	01-JAN-65
800-555-1	- 2	Green	05-FEB-68
800-555-1	3 Steve 213	White	16-MAR-71
800-555-1	4 Gail 214	Black	
	5 Doreen	Blue	20-MAY-70

The lines don't span more than 50 characters.



NOTE

The maximum number for the line size is 32,767.

Clearing Column Formatting

You clear the formatting for a column using the CLEAR option of the COLUMN command. For example, the following COLUMN command clears the formatting for the product id column:

SQL> COLUMN product_id CLEAR

You can clear the formatting for all columns using CLEAR COLUMNS. For example:

SOL> CLEAR COLUMNS

Once you've cleared the columns, the output from the queries will use the default format.

Using Variables

In this section, you'll see how to create variables that may be used in place of actual values in SQL statements. These variables are known as *substitution variables* because they are used as substitutes for values. When you run an SQL statement, you enter values for the variables; the values are then "substituted" into the SQL statement.

There are two types of substitution variables:

- **Temporary variables** A temporary variable is valid only for the SQL statement in which it is used—it doesn't persist.
- **Defined variables** A defined variable persists until you explicitly remove it, redefine it, or exit SQL*Plus.

You'll learn how to use these types of variables in this section.

Temporary Variables

You define a temporary variable using the ampersand character (&) in an SQL statement, followed by the name you want to call your variable. For example, &v_product_id defines a variable named v product id.

When you run the following query, SQL*Plus prompts you to enter a value for v_product_id and then uses that value in the WHERE clause. If you enter the value 2 for v_product_id, the details for product #2 will be displayed.

```
SQL> SELECT product_id, name, price

2 FROM products

3 WHERE product_id = &v_product_id;

Enter value for v_product_id: 2

old 3: WHERE product_id = &v_product_id

new 3: WHERE product_id = 2

PRODUCT_ID NAME PRICE

2 Chemistry 30
```

Notice that SQL*Plus does the following:

- Prompts you to enter a value for v_product_id.
- Substitutes the value you entered for v_product_id in the WHERE clause.

SQL*Plus shows you the substitution in the old and new lines in the output, along with the line number in the query where the substitution was performed. In the previous example, you can see that the old and new lines indicate that v product id is set to 2 in the WHERE clause.

Why Are Variables Useful?

Variables are useful because they allow you to create scripts that a user who doesn't know SQL can run. Your script would prompt the user to enter the value for a variable and use that value in an SQL statement. Let's take a look at an example.

Suppose you wanted to create a script for a user who doesn't know SQL, but who wants to see the details of a single specified product in the store. To do this, you could hard code the product id value in the WHERE clause of a guery and place it in an SQL*Plus script. For example, the following query retrieves product #1:

```
SELECT product id, name, price
FROM products
WHERE product id = 1;
```

This query works, but it only retrieves product #1. What if you wanted to change the product id value to retrieve a different row? You could modify the script, but this would be tedious. Wouldn't it be great if you could supply a variable for the product id? You can do that using a substitution variable.

If you rerun the query using the slash character (/), SQL*Plus will ask you to enter a new value for v product id. For example:

```
SOL> /
    Enter value for v product id: 3
    old 3: WHERE product id = &v product id
    new 3: WHERE product id = 3
    PRODUCT ID NAME
    ______
                                      25.99
          3 Supernova
```

Once again, SQL*Plus echoes the old line of the SQL statement (old 3: WHERE product id = &v product id), followed by the new line containing the variable value you entered (new 3: WHERE product id = 3).

Controlling Output Lines

You may control the output of the old and new lines using the SET VERIFY command. If you enter SET VERIFY OFF, the old and new lines are suppressed. For example:

```
SOL> SET VERIFY OFF
   SQL> /
   Enter value for v product id: 4
   PRODUCT ID NAME
                                  PRICE
   ______
         4 Tank War
                                   13.95
```

To turn the echoing of the lines back on, you enter SET VERIFY ON. For example:

SQL> SET VERIFY ON

Changing the Variable Definition Character

You can use the SET DEFINE command to specify a character other than an ampersand (&) for defining a variable. The following example shows how to set the variable character to the pound character (#) and shows a new query:

```
SQL> SET DEFINE '#'

SQL> SELECT product_id, name, price

2 FROM products

3 WHERE product_id = #v_product_id;

Enter value for v_product_id: 5

old 3: WHERE product_id = #v_product_id

new 3: WHERE product_id = 5

PRODUCT_ID NAME PRICE

5 Z Files 49.99
```

The next example uses SET DEFINE to change the character back to an ampersand:

SQL> SET DEFINE '&'

Substituting Table and Column Names Using Variables

You can also use variables to substitute the names of tables and columns. For example, the following query defines variables for a column name (v_col), a table name (v_table), and a column value (v_val):

```
SQL> SELECT name, &v col
       2 FROM &v table
       3 WHERE &v col = &v val;
      Enter value for v col: product type id
      old 1: SELECT name, &v col
      new 1: SELECT name, product_type_id
      Enter value for v table: products
      old 2: FROM &v table
      new 2: FROM products
      Enter value for v col: product_type_id
      Enter value for v val: 1
      old 3: WHERE &v col = &v val
      new 3: WHERE product type id = 1
     NAME
                                  PRODUCT TYPE ID
     Modern Science
     Chemistry
                                                 1
```

You can avoid having to repeatedly enter a variable by using &&. For example:

Chapter 3: Using SQL*Plus **77**

```
SQL> SELECT name, &&v col
       2 FROM &v table
       3 WHERE &&v col = &v_val;
     Enter value for v col: product type id
     old 1: SELECT name, &&v col
     new 1: SELECT name, product type id
     Enter value for v table: products
     old 2: FROM &v table
     new 2: FROM products
     Enter value for v val: 1
     old 3: WHERE &&v col = &v val
     new 3: WHERE product type id = 1
     NAME.
                                  PRODUCT TYPE ID
     Modern Science
     Chemistry
                                                 1
```

Variables give you a lot of flexibility in writing queries that another user may run. You can give the user a script and have them enter the variable values.

Defined Variables

You can define a variable prior to using it in an SQL statement. You may use these variables multiple times within an SQL statement. A defined variable persists until you explicitly remove it, redefine it, or exit SOL*Plus.

You define a variable using the DEFINE command. When you're finished with your variable, you remove it using UNDEFINE. You'll learn about these commands in this section. You'll also learn about the ACCEPT command, which allows you to define a variable and set its data type.

You can also define variables in an SQL*Plus script and pass values to those variables when the script is run. This feature enables you to write generic reports that any user can run—even if they're unfamiliar with SQL. You'll learn how to create simple reports later in the section "Creating Simple Reports."

Defining and Listing Variables Using the DEFINE Command

You use the DEFINE command to both define a new variable and list the currently defined variables. The following example defines a variable named v product id and sets its value to 7:

```
SQL> DEFINE v product id = 7
```

You can view the definition of a variable using the DEFINE command followed by the name of the variable. The following example displays the definition of v product id:

```
SQL> DEFINE v product id
    DEFINE V_PRODUCT ID = "7" (CHAR)
```

Notice that v product id is defined as a CHAR variable.

You can see all your session variables by entering DEFINE on its own. For example:

```
SQL> DEFINE

DEFINE _DATE = "12-AUG-07" (CHAR)

DEFINE _CONNECT_IDENTIFIER = "Oracle11g" (CHAR)

DEFINE _USER = "STORE" (CHAR)

DEFINE _PRIVILEGE = "" (CHAR)

DEFINE _SQLPLUS_RELEASE = "1101000400" (CHAR)

DEFINE _EDITOR = "Notepad" (CHAR)

DEFINE _O_VERSION = "Oracle Database 11g ..." (CHAR)

DEFINE _O_RELEASE = "1101000500" (CHAR)

DEFINE _RC = "0" (CHAR)

DEFINE _V_PRODUCT_ID = "7" (CHAR)
```

You can use a defined variable to specify an element such as a column value in an SQL statement. For example, the following query uses references v_product_id in the WHERE clause:

```
SQL> SELECT product_id, name, price

2 FROM products

3 WHERE product_id = &v_product_id;

old 3: WHERE product_id = &v_product_id

new 3: WHERE product_id = 7

PRODUCT_ID NAME PRICE

7 Space Force 9 13.49
```

You're not prompted for the value of v_product_id; that's because v_product_id was set to 7 when the variable was defined earlier.

Defining and Setting Variables Using the ACCEPT Command

The ACCEPT command waits for a user to enter a value for a variable. You can use the ACCEPT command to set an existing variable to a new value or to define a new variable and initialize it with a value. The ACCEPT command also allows you to specify the data type for the variable. The simplified syntax for the ACCEPT command is as follows:

```
ACCEPT variable_name [type] [FORMAT format] [PROMPT prompt] [HIDE]
```

where

- variable name is the variable name.
- type is the data type for the variable. You can use the CHAR, NUMBER, and DATE types. By default, variables are defined using the CHAR type. DATE variables are actually stored as CHAR variables.
- format is the format used for the variable. Some examples include A15 (15 characters), 9999 (a four-digit number), and DD-MON-YYYY (a date). You can view the number formats in Table 4-4 of Chapter 4; you can view the date formats in Table 5-2 of Chapter 5.
- prompt is the text displayed by SQL*Plus as a prompt to the user to enter the variable's value.

HIDE means hide the value as it is entered. For example, you might want to hide passwords or other sensitive information.

Let's take a look at some examples of the ACCEPT command. The following example defines a variable named v customer id as a two-digit number:

```
SQL> ACCEPT v customer id NUMBER FORMAT 99 PROMPT 'Customer id: '
    Customer id: 5
```

The next example defines a DATE variable named v date; the format is DD-MON-YYYY:

```
SQL> ACCEPT v date DATE FORMAT 'DD-MON-YYYY' PROMPT 'Date: '
   Date: 12-DEC-2006
```

The next example defines a CHAR variable named v password; the value entered is hidden using HIDE:

```
SQL> ACCEPT v password CHAR PROMPT 'Password: ' HIDE
     Password:
```

In Oracle Database 9i and below, the value appears as a string of asterisk characters (*) to hide the value as you enter it. In Oracle Database 10g and above, nothing is displayed as you

You can view your variables using the DEFINE command. For example:

```
SQL> DEFINE
```

```
DEFINE V CUSTOMER ID = 5 (NUMBER)
DEFINE V_DATE = "12-DEC-2006" (CHAR)
DEFINE V_PASSWORD = "1234567" (CHAR)
DEFINE V PRODUCT ID = "7" (CHAR)
```

Notice that v date is stored as a CHAR.

Removing Variables Using the UNDEFINE Command

You remove variables using the UNDEFINE command. The following example uses UNDEFINE to remove v customer id, v date, v password, and v product id:

```
SQL> UNDEFINE v customer id
     SQL> UNDEFINE v date
     SQL> UNDEFINE v password
     SQL> UNDEFINE v product id
```



NOTE

All your variables are removed when you exit SQL*Plus, even if you don't explicitly remove them using the UNDEFINE command.

Creating Simple Reports

You can use variables in an SQL*Plus script to create reports that a user can run. You'll find the SQL*Plus scripts referenced in this section in the SQL directory.



TIP

SQL*Plus was not specifically designed as a full-fledged reporting tool. If you have complex reporting requirements, you should use software like Oracle Reports.

Using Temporary Variables in a Script

The following report1.sql script uses a temporary variable named v_product_id in the WHERE clause of a query:

```
-- suppress display of the statements and verification messages
SET ECHO OFF
SET VERIFY OFF

SELECT product_id, name, price
FROM products
WHERE product_id = &v_product_id;
```

The SET ECHO OFF command stops SQL*Plus from displaying the SQL statements and commands in the script. SET VERIFY OFF suppresses display of the verification messages. I put these two commands in to minimize the number of extra lines displayed by SQL*Plus when you run the script.

You can run report1.sql in SQL*Plus using the @ command. For example:

```
Enter value for v_product_id: 2

PRODUCT_ID NAME PRICE

2 Chemistry 30
```

You'll need to replace the directory in the example with the directory where you saved the files for this book. Also, if you have spaces in the directory, you'll need to put everything after the @ command in quotes; for example:

@ "C:\my directory\sql book\SQL\report1.sql"

SQL> @ C:\sql book\SQL\report1.sql

Using Defined Variables in a Script

The following report2.sql script uses the ACCEPT command to define a variable named v_product_id:

```
SET ECHO OFF
SET VERIFY OFF

ACCEPT v_product_id NUMBER FORMAT 99 PROMPT 'Enter product id: '

SELECT product_id, name, price
FROM products
WHERE product_id = &v_product_id;
```

```
-- clean up
UNDEFINE v product id
```

Notice that a user-friendly prompt is specified for the entry of v product id and that v product id is removed at the end of the script—doing this makes the script cleaner. You can run the report2.sql script using SQL*Plus:

```
SQL> @ C:\sql book\SQL\report2.sql
```

```
Enter product id: 4
```

```
PRODUCT ID NAME
_____
   4 Tank War
                   13.95
```

Passing a Value to a Variable in a Script

You can pass a value to a variable when you run your script. When you do this, you reference the variable in the script using a number. The following script report3.sql shows an example of this; notice that the variable is identified using &1:

```
SET ECHO OFF
     SET VERIFY OFF
     SELECT product id, name, price
     FROM products
     WHERE product id = &1;
```

When you run report3.sql, you supply the variable's value after the script name. The following example passes the value 3 to report3.sql:

```
SQL> @ C:\sql book\SQL\report3.sql 3
     PRODUCT ID NAME
                                               PRICE
             3 Supernova
```

If you have spaces in the directory where you saved the scripts, you'll need to put the directory and script name in quotes; for example:

```
@ "C:\my directory\sql book\SQL\report3.sql" 3
```

You can pass any number of parameters to a script, with each value corresponding to the matching number in the script. The first parameter corresponds to &1, the second to &2, and so on. The following report4.sql script shows an example with two parameters:

```
SET ECHO OFF
     SET VERIFY OFF
     SELECT product id, product type id, name, price
     FROM products
     WHERE product type id = &1
     AND price > &2;
```

The following example run of report4.sql shows the addition of two values for &1 and &2, which are set to 1 and 9.99, respectively:

SQL> @ C:\sql book\SQL\report4.sql 1 9.99

PRODUCT_ID	PRODUCT_TYPE_ID	NAME	PRICE
1	1	Modern Science	19.95
2	1	Chemistry	30

Because &1 is set to 1, the product_type_id column in the WHERE clause is set to 1. Also, because &2 is set to 9.99, the price column in the WHERE clause is set to 9.99. Therefore, rows with a product type id of 1 and a price greater than 9.99 are displayed.

Adding a Header and Footer

You add a header and footer to your report using the TTITLE and BTITLE commands. The following is an example TTITLE command:

```
TTITLE LEFT 'Run date: '_DATE CENTER 'Run by the ' SQL.USER ' user' RIGHT 'Page: ' FORMAT 999 SOL.PNO SKIP 2
```

The following list explains the contents of this command:

- _DATE displays the current date.
- SQL.USER displays the current user.
- SQL. PNO displays the current page (FORMAT is used to format the number).
- LEFT, CENTER, and RIGHT justify the text.
- SKIP 2 skips two lines.

If the example is run on August 12, 2007 by the store user, it displays

Run date: 12-AUG-07 Run by the STORE user Page: 1

The next example shows a BTITLE command:

BTITLE CENTER 'Thanks for running the report' RIGHT 'Page: ' FORMAT 999 SQL.PNO

This command displays

```
Thanks for running the report Page: 1
```

The following report5.sql script contains the TTITLE and BTITLE commands:

```
TTITLE LEFT 'Run date: '_DATE CENTER 'Run by the ' SQL.USER ' user'
RIGHT 'Page: ' FORMAT 999 SQL.PNO SKIP 2

BTITLE CENTER 'Thanks for running the report' RIGHT 'Page: '
FORMAT 999 SQL.PNO

SET ECHO OFF
```

```
SET VERIFY OFF
SET PAGESIZE 30
SET LINESIZE 70
CLEAR COLUMNS
COLUMN product id HEADING ID FORMAT 99
COLUMN name HEADING 'Product Name' FORMAT A20 WORD WRAPPED
COLUMN description HEADING Description FORMAT A30 WORD WRAPPED
COLUMN price HEADING Price FORMAT $99.99
SELECT product id, name, description, price
FROM products;
CLEAR COLUMNS
TTITLE OFF
BTITLE OFF
```

The last two lines switch off the header and footer set by the TTITLE and BTITLE commands. The following example shows a run of report5.sql:

SQL> @ C:\sql_book\SQL\report5.sql

Run	date: 12-AUG-07	Run by the STORE user	Page:	1	
	Product Name	Description	Price		
		A description of modern science			
3	Supernova	_	\$25.99		
5	Z Files	Series on mysterious activities	\$49.99		
7		Aliens return Adventures of heroes Alien from another planet lands on Earth			
10 11	Pop 3 Creative Yell	The best classical music The best popular music Debut album Their greatest hits	\$15.99 \$14.99		
	Thanks for running the report Page: 1				

Computing Subtotals

You can add a subtotal for a column using a combination of the BREAK ON and COMPUTE commands. BREAK ON causes SQL*Plus to break up output based on a change in a column value, and COMPUTE causes SQL*Plus to compute a value for a column.

The following report6.sql script shows how to compute a subtotal for products of the same type:

```
BREAK ON product_type_id

COMPUTE SUM OF price ON product_type_id

SET ECHO OFF
SET VERIFY OFF
SET PAGESIZE 50
SET LINESIZE 70

CLEAR COLUMNS
COLUMN price HEADING Price FORMAT $999.99

SELECT product_type_id, name, price
FROM products
ORDER BY product_type_id;

CLEAR COLUMNS
```

The following example shows a run of report 6.sql:

SQL> @ C:\sql_book\SQL\report6.sql

PRODUCT_TYPE_ID	NAME	Price
1	Modern Science Chemistry	\$19.95 \$30.00

sum		\$49.95
2	Supernova	\$25.99
	Tank War	\$13.95
	Z Files	\$49.99
	2412: The Return	\$14.95

sum		\$104.88
3	Space Force 9	\$13.49
	From Another Planet	\$12.99

sum		\$26.48
4	Classical Music	\$10.99
	Pop 3	\$15.99
	Creative Yell	\$14.99

sum		\$41.97
	My Front Line	\$13.49

sum		\$13.49

Notice that whenever a new value for product_type_id is encountered, SQL*Plus breaks up the output and computes a sum for the price columns for the rows with the same

product type id. The product type id value is shown only once for rows with the same product type id. For example, "Modern Science" and "Chemistry" are both books and have a product type id of 1, and 1 is shown once for "Modern Science." The sum of the prices for these two books is \$49.95. The other sections of the report contain the sum of the prices for products with different product type id values.

Getting Help from SQL*Plus

You can get help from \$QL*Plus using the HELP command. The following example runs HELP:

SQL> HELP

HELP ____

Accesses this command line help system. Enter HELP INDEX or ? INDEX for a list of topics. In iSQL*Plus, click the Help button to display iSQL*Plus online help.

You can view SQL*Plus resources at http://otn.oracle.com/tech/sql plus/ and the Oracle Database Library at http://otn.oracle.com/documentation/

HELP|? [topic]

The next example runs HELP INDEX:

SOL> HELP INDEX

Enter Help [topic] for help.

@	COPY	PAUSE	SHUTDOWN
@ @	DEFINE	PRINT	SPOOL
/	DEL	PROMPT	SQLPLUS
ACCEPT	DESCRIBE	QUIT	START
APPEND	DISCONNECT	RECOVER	STARTUP
ARCHIVE LOG	EDIT	REMARK	STORE
ATTRIBUTE	EXECUTE	REPFOOTER	TIMING
BREAK	EXIT	REPHEADER	TTITLE
BTITLE	GET	RESERVED WORDS (SQL)	UNDEFINE
CHANGE	HELP	RESERVED WORDS (PL/SQL)	VARIABLE
CLEAR	HOST	RUN	WHENEVER OSERROR
COLUMN	INPUT	SAVE	WHENEVER SQLERROR
COMPUTE	LIST	SET	
CONNECT	PASSWORD	SHOW	

The following example runs HELP EDIT:

SQL> HELP EDIT

EDIT

Invokes an operating system text editor on the contents of the specified file or on the contents of the SQL buffer. The buffer has no command history list and does not record SQL*Plus commands.

```
ED[IT] [file_name[.ext]]
Not available in iSQL*Plus
```

Automatically Generating SQL Statements

In this section, I'll briefly show you a technique of writing SQL statements that produce other SQL statements. This capability is very useful and can save you a lot of typing when writing SQL statements that are similar. One simple example is an SQL statement that produces DROP TABLE statements, which remove tables from a database. The following query produces a series of DROP TABLE statements that drop the tables from the store schema:

```
SELECT 'DROP TABLE ' || table_name || ';'

FROM user_tables;

'DROPTABLE' || TABLE_NAME || ';'

DROP TABLE COUPONS;
DROP TABLE CUSTOMERS;
DROP TABLE EMPLOYEES;
DROP TABLE PRODUCTS;
DROP TABLE PRODUCT_TYPES;
DROP TABLE PROMOTIONS;
DROP TABLE PURCHASES;
DROP TABLE PURCHASES;
DROP TABLE PURCHASES_TIMESTAMP_WITH_TZ;
DROP TABLE PURCHASES_WITH_LOCAL_TZ;
DROP TABLE PURCHASES_WITH_TIMESTAMP;
DROP TABLE SALARY GRADES;
```



NOTE

user_tables contains the details of the tables in the user's schema. The table name column contains names of the tables.

You can spool the generated SQL statements to a file and run them later.

Disconnecting from the Database and Exiting SQL*Plus

You can disconnect from the database and keep SQL*Plus running by entering DISCONNECT (SQL*Plus also automatically performs a COMMIT for you). While you're connected to the database, SQL*Plus maintains a database session for you. When you disconnect from the database, your session is ended. You can reconnect to a database by entering CONNECT.

To end SQL*Plus, you enter EXIT (SQL*Plus also automatically performs a COMMIT for you).

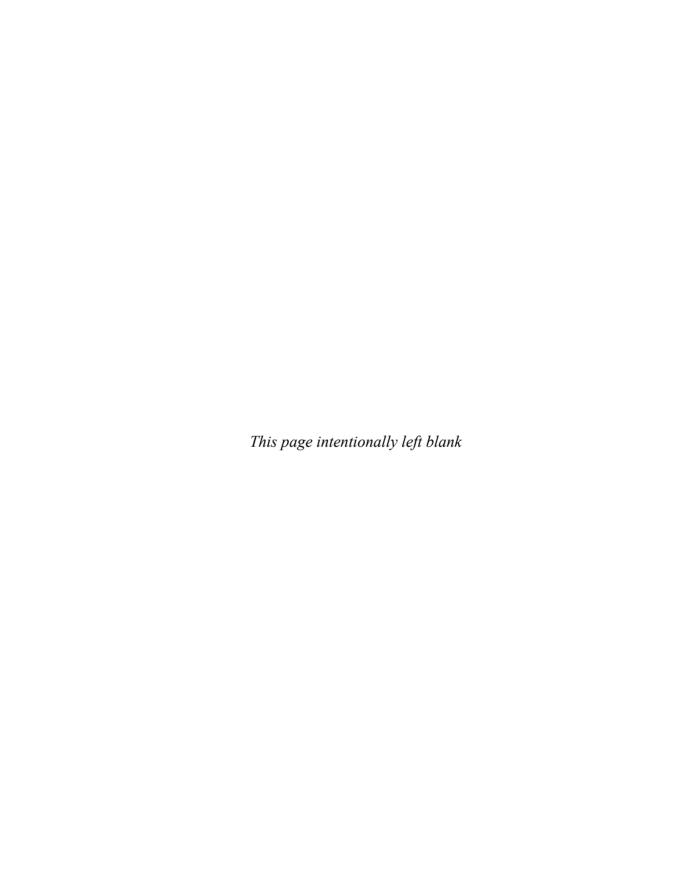
Summary

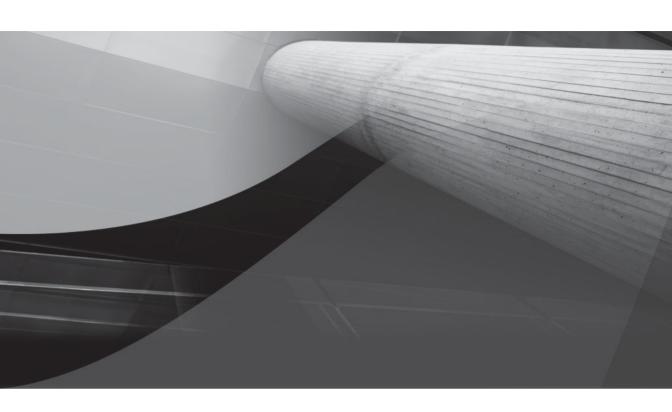
In this chapter, you learned how to do the following:

- View the structure of a table
- Edit an SQL statement
- Save, retrieve, and run files containing SQL and SQL*Plus commands
- Format the results returned by SQL*Plus
- Set the page and line size for SQL*Plus output
- Use variables in SQL*Plus
- Create simple reports
- Get help from SQL*Plus
- Write SQL statements that generate other SQL statements
- Disconnect from the database and exit SQL*Plus

For further details on SQL*Plus, you can read the SQL*Plus User's Guide and Reference, published by Oracle Corporation.

In the next chapter, you'll learn how to use functions.





CHAPTER 1

Using Simple Functions



n this chapter, you'll learn about some of the Oracle database's built-in functions. A function accepts zero or more input parameters and returns an output parameter. There are two main types of functions you can use in an Oracle database:

- **Single-row functions** operate on one row at a time and return one row of output for each input row. An example single-row function is CONCAT (x, y), which appends y to x and returns the resulting string.
- **Aggregate functions** operate on multiple rows at the same time and return one row of output. An example aggregate function is AVG(x), which returns the average of x where x may be a column or, more generally, any expression.

You'll learn about single-row functions first, followed by aggregate functions. You'll see more advanced functions as you progress through this book.

Using Single-Row Functions

A single-row function operates on one row at a time and returns one row of output for each row. There are five main types of single-row functions:

- Character functions manipulate strings of characters.
- Numeric functions perform calculations.
- **Conversion functions** convert a value from one database type to another.
- **Date functions** process dates and times.
- **Regular expression functions** use regular expressions to search data. These functions were introduced in Oracle Database 10*g* and are expanded in 11*g*.

You'll learn about character functions first, followed by numeric functions, conversion functions, and regular expression functions. You'll learn about date functions in the next chapter.

Character Functions

Character functions accept character input, which may come from a column in a table or, more generally, from any expression. This input is processed and a result returned. An example character function is $\mathtt{UPPER}()$, which converts the letters in an input string to uppercase and returns the new string. Another example is $\mathtt{NVL}()$, which converts a null value to another value. In Table 4-1, which shows some of the character functions, and in all the syntax definitions that follow, x and y may represent columns from a table or, more generally, any valid expressions.

You'll learn more about some of the functions shown in Table 4-1 in the following sections.

Function	Description
ASCII(x)	Returns the ASCII value for the character <i>x</i> .
CHR(x)	Returns the character with the ASCII value of <i>x</i> .
CONCAT(x, y)	Appends <i>y</i> to <i>x</i> and then returns the new string.
INITCAP(x)	Converts the initial letter of each word in <i>x</i> to uppercase and returns the new string.
<pre>INSTR(x, find_string [, start] [, occurrence])</pre>	Searches for <code>find_string</code> in <code>x</code> and returns the position at which <code>find_string</code> occurs. You can supply an optional <code>start</code> position to begin the search. You can also supply an optional <code>occurrence</code> that indicates which occurrence of <code>find_string</code> should be returned.
LENGTH(x)	Returns the number of characters in x.
LOWER(x)	Converts the letters in <i>x</i> to lowercase and returns the new string.
LPAD(x, width [, pad_string])	Pads <i>x</i> with spaces to left to bring the total length of the string up to <i>width</i> characters. You can supply an optional <i>pad_string</i> , which specifies a string to be repeated to the left of <i>x</i> to fill up the padded space. The resulting padded string is returned.
<pre>LTRIM(x [, trim_string])</pre>	Trims characters from the left of x. You can supply an optional trim_string, which specifies the characters to trim; if no trim_string is supplied, then spaces are trimmed by default.
NANVL(x, value)	Returns $value$ if x matches the NAN (not a number) special value; otherwise x is returned. (This function was introduced in Oracle Database 10 g .)
NVL(x, value)	Returns <i>value</i> if <i>x</i> is null; otherwise <i>x</i> is returned.
NVL2(x, value1, value2)	Returns value1 if x is not null; otherwise value2 is returned.
REPLACE(x, search_string, replace_string)	Searches x for search_string and replaces it with replace_string.
<pre>RPAD(x, width [, pad_string])</pre>	Same as $LPAD()$, but x is padded to the right.
<pre>RTRIM(x [, trim_string])</pre>	Same as $\mathtt{LTRIM}()$, but x is trimmed from the right.
SOUNDEX(x)	Returns a string containing the phonetic representation of <i>x</i> . This lets you compare words that sound alike in English but are spelled differently.
SUBSTR(x, start	Returns a substring of x that begins at the position specified by
[, length])	start. You can supply an optional length for the substring.
<pre>TRIM([trim_char FROM) x)</pre>	Trims characters from the left and right of x. You can supply an optional trim_char, which specifies the characters to trim; if no trim_char is supplied, spaces are trimmed by default.
UPPER(x)	Converts the letters in <i>x</i> to uppercase and returns the new string.

 TABLE 4-1
 Character Functions

ASCII() and CHR()

You use ASCII(x) to get the ASCII value for the character x. You use CHR(x) to get the character with the ASCII value of x.

The following query gets the ASCII value of a, A, z, Z, 0, and 9 using ASCII ():

SELECT ASCII('a'), ASCII('A'), ASCII('z'), ASCII('Z'), ASCII(0), ASCII(9)
FROM dual;



NOTE

The dual table is used in this query. As you saw in Chapter 2, the dual table contains a single row through which you may perform queries that don't go against a particular table.

The following query gets the characters with the ASCII values of 97, 65, 122, 90, 48, and 57 using CHR ():

SELECT CHR(97), CHR(65), CHR(122), CHR(90), CHR(48), CHR(57) FROM dual;

Notice the characters returned from CHR() in this query are the same as those passed to ASCII() in the previous query. This shows that CHR() and ASCII() have the opposite effect.

CONCAT()

You use CONCAT (x, y) to append y to x and then return the new string. The following query appends last name to first name using CONCAT():

SELECT CONCAT(first_name, last_name)
FROM customers:

```
CONCAT (FIRST_NAME, LA
------
JohnBrown
CynthiaGreen
SteveWhite
GailBlack
DoreenBlue
```



NOTE

CONCAT () is the same as the | | operator you saw in Chapter 2.

INITCAP()

You use INITCAP (x) to convert the initial letter of each word in x to uppercase.

The following query retrieves the product id and description columns from the products table, then uses INITCAP() to convert the first letter of each word in description to uppercase:

```
SELECT product id, INITCAP(description)
    FROM products
    WHERE product id < 4;
    PRODUCT ID INITCAP (DESCRIPTION)
             1 A Description Of Modern Science
             2 Introduction To Chemistry
             3 A Star Explodes
```

INSTR()

You use INSTR(x, find string[, start][, occurrence]) to search for find string in x. INSTR() returns the position at which find string occurs. You can supply an optional start position to begin the search. You can also supply an optional occurrence that indicates which occurrence of find string should be returned.

The following query gets the position where the string Science occurs in the name column for product #1:

```
SELECT name, INSTR(name, 'Science')
    FROM products
    WHERE product id = 1;
                                  INSTR(NAME, 'SCIENCE')
    Modern Science
```

The next guery displays the position where the second occurrence of the e character occurs, starting from the beginning of the product name:

```
SELECT name, INSTR(name, 'e', 1, 2)
   FROM products
   WHERE product id = 1;
   NAME
                         INSTR (NAME, 'E', 1, 2)
    -----
   Modern Science
```

Notice the second e in Modern Science is the eleventh character.

You can also use dates with character functions. The following query gets the position where the string JAN occurs in the dob column for customer #1:

```
SELECT customer id, dob, INSTR(dob, 'JAN')
    FROM customers
    WHERE customer id = 1;
    CUSTOMER ID DOB INSTR(DOB, 'JAN')
    -----
            1 01-JAN-65
```

LENGTH()

You use LENGTH (x) to get the number of characters in x. The following query gets the length of the strings in the name column of the products table using LENGTH ():

SELECT name, LENGTH(name) FROM products;

NAME	LENGTH (NAME)
Modern Science	14
Chemistry	9
Supernova	9
Tank War	8
Z Files	7
2412: The Return	16
Space Force 9	13
From Another Planet	19
Classical Music	15
Pop 3	5
Creative Yell	13
My Front Line	13

The next query gets the total number of characters that make up the product price; notice that the decimal point (.) is counted in the number of price characters:

SELECT price, LENGTH(price) FROM products WHERE product_id < 3;

PRICE	LENGTH (PRICE)
19.95	5
30	2

LOWER() and **UPPER()**

You use LOWER (x) to convert the letters in x to lowercase. Similarly, you use UPPER (x) to convert the letters in x to uppercase.

The following query converts the strings in the first_name column to uppercase using the UPPER() function and the strings in the last_name column to lowercase using the LOWER() function:

SELECT UPPER(first_name), LOWER(last_name) FROM customers;

```
UPPER(FIRS LOWER(LAST
-----
JOHN brown
CYNTHIA green
```

STEVE	white
GAIL	black
DOREEN	blue

LPAD() and RPAD()

You use LPAD (x, width [, pad string]) to pad x with spaces to the left to bring the total length of the string up to width characters. You can supply an optional pad string, which specifies a string to be repeated to the left of x to fill up the padded space. The resulting padded string is then returned. Similarly, you use RPAD(x, width [, pad string]) to pad x with strings to the right.

The following query retrieves the name and price columns from the products table. The name column is right-padded using RPAD() to a length of 30 characters, with periods filling up the padded space. The price column is left-padded using LPAD() to a length of 8, with the string *+ filling up the padded space.

```
SELECT RPAD (name, 30, '.'), LPAD (price, 8, '*+')
  FROM products
  WHERE product id < 4;
  RPAD (NAME, 30, '.')
  _____
  Modern Science..... *+*19.95
  Chemistry..... *+*+*+30
  Supernova..... *+*25.99
```



NOTE

This example shows that character functions can use numbers. Specifically, the price column in the example contains a number that was left-padded by LPAD().

LTRIM(), RTRIM(), and TRIM()

LTRIM(' Hello Gail Seymour!'),

You use LTRIM (x [, trim string]) to trim characters from the left of x. You can supply an optional trim string, which specifies the characters to trim; if no trim string is supplied; spaces are trimmed by default. Similarly, you use RTRIM() to trim characters from the right of x. You use TRIM() to trim characters from the left and right of x. The following query uses these three functions:

SELECT

```
RTRIM('Hi Doreen Oakley!abcabc', 'abc'),
 TRIM('0' FROM '000Hey Steve Button!00000')
FROM dual;
LTRIM('HELLOGAILSEY RTRIM('HIDOREENOA TRIM('0'FROM'000H
____________
Hello Gail Seymour! Hi Doreen Oakley! Hey Steve Button!
```

NVL()

You use NVL() to convert a null value to another value. NVL(x, value) returns value if x is null; otherwise x is returned.

The following query retrieves the customer_id and phone columns from the customers table. Null values in the phone column are converted to the string Unknown Phone Number by NVL():

SELECT customer_id, NVL(phone, 'Unknown Phone Number') FROM customers;

The phone column for customer #5 is converted to Unknown Phone Number because the phone column is null for that row.

NVL2()

NVL2 (x, value1, value2) returns value1 if x is not null; otherwise value2 is returned.

The following query retrieves the customer ind and phone columns from the customer.

The following query retrieves the customer_id and phone columns from the customers table. Non-null values in the phone column are converted to the string Known, and null values are converted to Unknown:

SELECT customer_id, NVL2(phone, 'Known', 'Unknown') FROM customers;

Notice that the phone column values are converted to Known for customers #1 through #4 because the phone column values for those rows are not null. For customer #5 the phone column value is converted to Unknown because the phone column is null for that row.

REPLACE()

You use REPLACE (x, search_string, replace_string) to search x for search_string and replace it with replace string.

The following example retrieves the name column from the products table for product #1 (whose name is Modern Science) and replaces the string Science with Physics using REPLACE():

```
SELECT REPLACE(name, 'Science', 'Physics')
   FROM products
   WHERE product id = 1;
   REPLACE (NAME, 'SCIENCE', 'PHYSICS')
   Modern Physics
```



NOTE

REPLACE() doesn't modify the actual row in the database; only the row returned by the function is modified.

SOUNDEX()

You use SOUNDEX(x) to get a string containing the phonetic representation of x. This lets you compare words that sound alike in English but are spelled differently.

The following query retrieves the last name column from the customers table where last name sounds like "whyte":

```
SELECT last name
     FROM customers
     WHERE SOUNDEX(last name) = SOUNDEX('whyte');
     LAST NAME
     _____
     White
```

The next query gets last names that sound like "bloo":

```
SELECT last name
    FROM customers
    WHERE SOUNDEX(last name) = SOUNDEX('bloo');
    LAST NAME
    _____
    Blue
```

SUBSTR()

upernov

You use SUBSTR(x, start [, length]) to return a substring of x that begins at the position specified by start. You can also provide an optional length for the substring.

The following query uses SUBSTR() to get the 7-character substring starting at position 2 from the name column of the products table:

```
SELECT SUBSTR(name, 2, 7)
    FROM products
    WHERE product id < 4;
    SUBSTR (
    odern S
    hemistr
```

Using Expressions with Functions

You're not limited to using columns in functions: you can supply any valid expression that evaluates to a string. The following query uses the SUBSTR() function to get the substring little from the string Mary had a little lamb:

SELECT SUBSTR('Mary had a little lamb', 12, 6) FROM dual;

SUBSTR ----little

Combining Functions

You can use any valid combination of functions in an SQL statement. The following query combines the <code>UPPER()</code> and <code>SUBSTR()</code> functions; notice that the output from <code>SUBSTR()</code> is passed to <code>UPPER()</code>:

SELECT name, UPPER(SUBSTR(name, 2, 8))
FROM products
WHERE product id < 4;

NAME	UPPER(SU
Modern Science	ODERN SC
Chemistry	HEMISTRY
Supernova	UPERNOVA



NOTE

This ability to combine functions is not limited to character functions. Any valid combination of functions will work.

Numeric Functions

You use the numeric functions to perform calculations. These functions accept an input number, which may come from a numeric column or any expression that evaluates to a number. A calculation is then performed and a number returned. An example of a numeric function is SQRT(x), which returns the square root of x.

Table 4-2 shows some of the numeric functions.

You'll learn more about some of the functions shown in Table 4-2 in the following sections.

Function	Description	Examples
ABS(x)	Returns the absolute value of <i>x</i> .	ABS $(10) = 10$ ABS $(-10) = 10$
ACOS(x)	Returns the arccosine of <i>x</i> .	ACOS(1) = 0 ACOS(-1) = 3.14159265
ASIN(x)	Returns the arcsine of <i>x</i> .	ASIN(1) = 1.57079633 ASIN(-1) = -1.5707963
ATAN(x)	Returns the arctangent of <i>x</i> .	ATAN(1) = $.785398163$ ATAN(-1) = 78539816

Function	Description	Examples
ATAN2 (x, y)	Returns the arctangent of x and y .	ATAN2 $(1, -1) = 2.35619449$
BITAND(x, y)	Returns the result of performing a bitwise AND on x and y .	BITAND(0, 0) = 0 BITAND(0, 1) = 0 BITAND(1, 0) = 0 BITAND(1, 1) = 1 BITAND(1010, 1100) = 64
COS(x)	Returns the cosine of x , where x is an angle in radians.	COS(90 * 3.1415926) = 1 COS(45 * 3.1415926) = -1
COSH(x)	Returns the hyperbolic cosine of <i>x</i> .	COSH(3.1415926) = 11.5919527
CEIL(x)	Returns the smallest integer greater than or equal to x .	CEIL(5.8) = 6 CEIL(-5.2) = -5
EXP(x)	Returns the result of the number e raised to the power x , where e is approximately 2.71828183.	EXP(1) = 2.71828183 EXP(2) = 7.3890561
FLOOR(x)	Returns the largest integer less than or equal to x .	FLOOR(5.8) = 5 FLOOR(-5.2) = -6
LOG(x, y)	Returns the logarithm, base x , of y .	LOG(2, 4) = 2 LOG(2, 5) = 2.32192809
LN(x)	Returns the natural logarithm of x.	LN(2.71828183) = 1
MOD(x, y)	Returns the remainder when x is divided by y .	MOD(8, 3) = 2 MOD(8, 4) = 0
POWER(x, y)	Returns the result of x raised to the power y .	POWER $(2, 1) = 2$ POWER $(2, 3) = 8$
ROUND(x [, y])	Returns the result of rounding x to an optional y decimal places. If y is omitted, x is rounded to zero decimal places. If y is negative, x is rounded to the left of the decimal point.	ROUND(5.75) = 6 ROUND(5.75, 1) = 5.8 ROUND(5.75, -1) = 10
SIGN(x)	Returns -1 if x is negative, 1 if x is positive, or 0 if x is zero.	SIGN(-5) = -1 SIGN(5) = 1 SIGN(0) = 0
SIN(x)	Returns the sine of x.	SIN(0) = 0
SINH(x)	Returns the hyperbolic sine of x.	SINH(1) = 1.17520119
SQRT(x)	Returns the square root of x.	SQRT(25) = 5 SQRT(5) = 2.23606798
TAN(x)	Returns the tangent of x.	TAN(0) = 0
TANH(x)	Returns the hyperbolic tangent of x.	TANH(1) = .761594156
TRUNC(x [, y])	Returns the result of truncating x to an optional y decimal places. If y is omitted, x is truncated to zero decimal places. If y is negative, x is truncated to the left of the decimal point.	TRUNC $(5.75) = 5$ TRUNC $(5.75, 1) = 5.7$ TRUNC $(5.75, -1) = 0$

 TABLE 4-2
 Numeric Functions (continued)

ABS()

You use ABS (x) to get the absolute value of x. The absolute value of a number is that number without any positive or negative sign. The following query gets the absolute value of 10 and -10:

SELECT ABS(10), ABS(-10) FROM dual;

The absolute value of 10 is 10. The absolute value of -10 is 10.

Of course, the parameters that are input to any of the number functions don't have to be literal numbers. The input may also be a numeric column from a table or, more generally, any valid expression. The following query gets the absolute value of subtracting 30 from the price column from the products table for the first three products:

PRODUCT_ID	PRICE	PRICE-30	ABS(PRICE-30)
1	19.95	-10.05	10.05
2	30	0	0
3	25.99	-4.01	4.01

CEIL()

You use CEIL(x) to get the smallest integer greater than or equal to x. The following query uses CEIL() to get the absolute values of 5.8 and -5.2:

The ceiling for 5.8 is 6, because 6 is the smallest integer greater than 5.8. The ceiling for –5.2 is –5, because –5.2 is negative, and the smallest integer greater than this is –5.

FLOOR()

You use FLOOR(x) to get the largest integer less than or equal to x. The following query uses FLOOR() to get the absolute value of 5.8 and -5.2:

The floor for 5.8 is 5; because 5 is the largest integer less than 5.8. The floor for –5.2 is –6, because –5.2 is negative, and the largest integer less than this is –6.

MOD()

You use MOD(x, y) to get the remainder when x is divided by y. The following query uses MOD () to get the remainder when 8 is divided by 3 and 4:

The remainder when 8 is divided by 3 is 2: 3 goes into 8 twice, leaving 2 left over—the remainder. The remainder when 8 is divided by 4 is 0: 4 goes into 8 twice, leaving nothing left over.

POWER()

You use POWER (x, y) to get the result of x raised to the power y. The following query uses POWER () to get 2 raised to the power 1 and 3:

```
POWER(2,1) POWER(2,3)
_____
```

When 2 is raised to the power 1, which is equivalent to 2*1, the result is 2; 2 raised to the power 3 is equivalent to 2*2*2, the result of which is 8.

ROUND()

You use ROUND (x, [y]) to get the result of rounding x to an optional y decimal places. If y is omitted, x is rounded to zero decimal places. If y is negative, x is rounded to the left of the decimal point.

The following query uses ROUND() to get the result of rounding 5.75 to zero, 1, and -1 decimal places:

```
ROUND(5.75) ROUND(5.75,1) ROUND(5.75,-1)
_____
            5.8
```

5.75 rounded to zero decimal places is 6; 5.75 rounded to one decimal place (to the right of the decimal point) is 5.8; and 5.75 rounded to one decimal place to the left of the decimal point (as indicated using a negative sign) is 10.

SIGN()

You use SIGN (x) to get the sign of x. SIGN () returns -1 if x is negative, 1 if x is positive, or 0 if x is zero. The following query gets the sign of -5, 5, and 0:

SELECT SIGN(-5), SIGN(5), SIGN(0) FROM dual;

The sign of -5 is -1; the sign of 5 is 1; the sign of 0 is 0.

SQRT()

You use SQRT (x) to get the square root of x. The following query gets the square root of 25 and 5:

The square root of 25 is 5; the square root of 5 is approximately 2.236.

TRUNC()

You use TRUNC (x, [y]) to get the result of truncating the number x to an optional y decimal places. If y is omitted, x is truncated to zero decimal places. If y is negative, x is truncated to the left of the decimal point. The following query truncates 5.75 to zero, 1, and –1 decimal places:

```
TRUNC (5.75) TRUNC (5.75,1) TRUNC (5.75,-1)

5 5.7 0
```

In the above, 5.75 truncated to zero decimal places is 5; 5.75 truncated to one decimal place (to the right of the decimal point) is 5.7; and 5.75 truncated to one decimal place to the left of the decimal point (as indicated using a negative sign) is 0.

Conversion Functions

Sometimes you need to convert a value from one data type to another. For example, you might want to reformat the price of a product that is stored as a number (e.g., 1346.95) to a string containing dollar signs and commas (e.g., \$1,346.95). For this purpose, you use a conversion function to convert a value from one data type to another.

Table 4-3 shows some of the conversion functions.

You'll learn more about the TO_CHAR() and TO_NUMBER() functions in the following sections. You'll learn about some of the other functions in Table 4-3 as you progress through this book. You can find out more about national language character sets and Unicode in the *Oracle Database Globalization Support Guide* from Oracle Corporation.

Function	Description
ASCIISTR(x)	Converts x to an ASCII string, where x may be a string in any character set.
BIN_TO_NUM(x)	Converts a binary number x to a NUMBER.
CAST(x AS type)	Converts x to a compatible database type specified in $type$.
CHARTOROWID(x)	Converts x to a ROWID.
COMPOSE (x)	Converts x to a Unicode string in its fully normalized form in the same character set as x . Unicode uses a 2-byte character set and can represent over 65,000 characters; it may also be used to represent non-English characters.
<pre>CONVERT(x, source_char_set, dest_char_set)</pre>	Converts x from source_char_set to dest_char_set.
DECODE(x, search, result, default)	Compares x with the value in search; if equal, DECODE() returns the value in result; otherwise the value in default is returned.
DECOMPOSE (x)	Converts x to a Unicode string after decomposition of the string into the same character set as x .
HEXTORAW(x)	Converts the character x containing hexadecimal digits (base-16) to a binary number (RAW). This function then returns the RAW number.
NUMTODSINTERVAL(x)	Converts the number x to an INTERVAL DAY TO SECOND. You'll learn about date and time interval—related functions in the next chapter.
NUMTOYMINTERVAL(x)	Converts the number x to an INTERVAL YEAR TO MONTH.
RAWTOHEX(x)	Converts the binary number (RAW) x to a VARCHAR2 string containing the equivalent hexadecimal number.
RAWTONHEX(x)	Converts the binary number (RAW) x to an NVARCHAR2 string containing the equivalent hexadecimal number. (NVARCHAR2 stores a string using the national character set.)
ROWIDTOCHAR(x)	Converts the ROWID x to a VARCHAR2 string.
ROWIDTONCHAR(x)	Converts the ROWID x to an NVARCHAR2 string.
TO_BINARY_DOUBLE(x)	Converts x to a BINARY_DOUBLE. (This function was introduced in Oracle Database 10 g .)
TO_BINARY_FLOAT(x)	Converts x to a BINARY_FLOAT. (This function was introduced in Oracle Database 10 g .)
TO_BLOB(x)	Converts x to a binary large object (BLOB). A BLOB is used to store large amounts of binary data. You'll learn about large objects in Chapter 14.

 TABLE 4-3
 Conversion Functions

Function	Description
TO_CHAR(x [, format])	Converts x to a VARCHAR2 string. You can supply an optional $format$ that indicates the format of x .
TO_CLOB(x)	Converts x to a character large object (CLOB). A CLOB is used to store large amounts of character data.
$TO_DATE(x [, format])$	Converts x to a DATE.
TO_DSINTERVAL(x)	Converts the string x to an INTERVAL DAY TO SECOND.
TO_MULTI_BYTE(x)	Converts the single-byte characters in x to their corresponding multi-byte characters. The return type is the same as the type for x .
TO_NCHAR(x)	Converts x in the database character set to an NVARCHAR2 string.
TO_NCLOB(x)	Converts x to a large object NCLOB. An NCLOB is used to store large amounts of national language character data.
TO_NUMBER(x [, format])	Converts x to a NUMBER.
TO_SINGLE_BYTE(x)	Converts the multi-byte characters in x to their corresponding single-byte characters. The return type is the same as the type for x .
TO_TIMESTAMP(x)	Converts the string x to a TIMESTAMP.
TO_TIMESTAMP_TZ(x)	Converts the string x to a TIMESTAMP WITH TIME ZONE.
TO_YMINTERVAL(x)	Converts the string x to an INTERVAL YEAR TO MONTH.
<pre>TRANSLATE(x, from_string, to_string)</pre>	Converts all occurrences of from_string in x to to_string.
UNISTR(x)	Converts the characters in x to an NCHAR character. (NCHAR stores a character using the national language character set.)

 TABLE 4-3
 Conversion Functions (continued)

TO_CHAR()

You use ${\tt TO_CHAR}(x \ [, format])$ to convert x to a string. You can also provide an optional format that indicates the format of x. The structure format depends on whether x is a number or date. You'll learn how to use ${\tt TO_CHAR}()$ to convert a number to a string in this section, and you'll see how to convert a date to a string in the next chapter.

Let's take a look at a couple of simple queries that use TO_CHAR() to convert a number to a string. The following query converts 12345.67 to a string:

```
SELECT TO_CHAR(12345.67)
FROM dual;
```

```
TO_CHAR(1 ----- 12345.67
```

The next query uses TO_CHAR() to convert 12345678.90 to a string and specifies this number is to be converted using the format 99, 999. 99. This results in the string returned by TO CHAR () having a comma to delimit the thousands:

SELECT TO_CHAR(12345.67, '99,999.99') FROM dual;

TO CHAR (12 12,345.67

The optional format string you may pass to TO CHAR() has a number of parameters that affect the string returned by TO_CHAR(). Some of these parameters are listed in Table 4-4.

Parameter	Format Examples	Description
9	999	Returns digits in specified positions, with a leading negative sign if the number is negative.
0	0999 9990	0999: Returns a number with leading zeros. 9990: Returns a number with trailing zeros.
	999.99	Returns a decimal point in the specified position.
,	9,999	Returns a comma in the specified position.
\$	\$999	Returns a leading dollar sign.
В	В9.99	If the integer part of a fixed point number is zero, returns spaces for the zeros.
С	C999	Returns the ISO currency symbol in the specified position. The symbol comes from the NLS_ISO_CURRENCY database parameter set by a DBA.
D	9D99	Returns the decimal point symbol in the specified position. The symbol comes from the NLS_NUMERIC_CHARACTER database parameter (the default is a period character).
EEEE	9.99EEEE	Returns the number using the scientific notation.
FM	FM90.9	Removes leading and trailing spaces from the number.
G	9G999	Returns the group separator symbol in the specified position. The symbol comes from the NLS_NUMERIC_CHARACTER database parameter.
L	L999	Returns the local currency symbol in the specified position. The symbol comes from the NLS_CURRENCY database parameter.

 TABLE 4-4
 Numeric Formatting Parameters

Parameter	Format Examples	Description
MI	999MI	Returns a negative number with a trailing minus sign. Returns a positive number with a trailing space.
PR	999PR	Returns a negative number in angle brackets (< >). Returns a positive number with leading and trailing spaces.
RN	RN	Returns the number as Roman numerals. RN returns
rn	rn	uppercase numerals; rn returns lowercase numerals. The number must be an integer between 1 and 3999.
S	S999	S999: Returns a negative number with a leading negative
	999S	sign; returns a positive number with a leading positive sign. 9998: Returns a negative number with a trailing negative sign; returns a positive number with a trailing positive sign.
TM	TM	Returns the number using the minimum number of characters. The default is TM9, which returns the number using fixed notation unless the number of characters is greater than 64. If greater than 64, the number is returned using scientific notation.
U	U999	Returns the dual currency symbol (Euro, for example) in the specified position. The symbol comes from the NLS_DUAL_CURRENCY database parameter.
V	99V99	Returns the number multiplied by 10^x where x is the number of 9 characters after the V . If necessary, the number is rounded.
X	XXXX	Returns the number in hexadecimal. If the number is not an integer, the number is rounded to an integer.

 TABLE 4-4
 Numeric Formatting Parameters (continued)

Let's look at some more examples that convert numbers to strings using ${\tt TO_CHAR}$ (). The following table shows examples of calling ${\tt TO_CHAR}$ (), along with the output returned.

TO_CHAR() Function	Call	Output
TO_CHAR(12345.67	, '99999.99')	12345.67
TO_CHAR(12345.67	, '99,999.99')	12,345.67
TO_CHAR(-12345.6	7, '99,999.99')	-12,345.67
TO_CHAR(12345.67	, '099,999.99')	012,345.67
TO_CHAR(12345.67	, '99,999.9900')	12,345.6700
TO_CHAR(12345.67	, '\$99,999.99')	\$12,345.67
TO_CHAR(0.67, 'B	9.99')	.67

TO CHAR() Function Call Output TO CHAR (12345.67, 'C99,999.99') USD12,345.67 TO CHAR (12345.67, '99999D99') 12345.67 TO CHAR(12345.67, '99999.99EEEE') 1.23E+04 TO CHAR(0012345.6700, 'FM99999.99') 12345.67 TO CHAR (12345.67, '99999G99') 123,46 TO CHAR(12345.67, 'L99,999.99') \$12,345.67 TO CHAR (-12345.67, '99,999.99MI') 12,345.67 TO CHAR(-12345.67, '99,999.99PR') 12,345.67 TO CHAR (2007, 'RN') MMVII TO CHAR (12345.67, 'TM') 12345.67 TO CHAR(12345.67, 'U99,999.99') \$12,345.67 TO CHAR (12345.67, '99999V99') 1234567

TO CHAR() will return a string of pound characters (#) if you try to format a number that contains too many digits for the format. For example:

```
SELECT TO CHAR (12345678.90, '99,999.99')
   FROM dual:
   TO CHAR (12
   _____
   ##########
```

Pound characters are returned by ${\tt TO_CHAR}$ () because the number 12345678.90 has more digits than those allowed in the format 99, 999.99.

You can also use TO CHAR () to convert columns containing numbers to strings. For example, the following query uses TO CHAR() to convert the price column of the products table to a string:

```
SELECT product id, 'The price of the product is' || TO CHAR(price, '$99.99')
    FROM products
    WHERE product id < 5;
```

```
PRODUCT ID 'THEPRICEOFTHEPRODUCTIS' | | TO CHAR (
_____
       1 The price of the product is $19.95
       2 The price of the product is $30.00
       3 The price of the product is $25.99
       4 The price of the product is $13.95
```

TO NUMBER()

You use TO NUMBER (x [, format]) to convert x to a number. You can provide an optional format string to indicate the format of x. Your format string may use the same parameters as those listed earlier in Table 4-4.

The following query converts the string 970.13 to a number using TO NUMBER():

```
SELECT TO_NUMBER('970.13')
FROM dual;
```

The next query converts the string 970.13 to a number using TO_NUMBER() and then adds 25.5 to that number:

SELECT TO_NUMBER('970.13') + 25.5 FROM dual;

The next query converts the string -\$12,345.67 to a number, passing the format string \$99,999.99 to TO NUMBER():

SELECT TO_NUMBER('-\$12,345.67', '\$99,999.99')
FROM dual;

CAST()

You use CAST (x AS type) to convert x to a compatible database type specified by type. The following table shows the valid type conversions (valid conversions are marked with an X):

			From				
То	BINARY_ FLOAT BINARY_ DOUBLE	CHAR VARCHAR2	NUMBER	DATE TIMESTAMP INTERVAL	RAW	ROWID UROWID	NCHAR NVARCHAR2
BINARY_FLOAT BINARY_DOUBLE	X	X	X				X
CHAR VARCHAR2	X	Х	Х	X	Х	Х	
NUMBER	X	X	X				X
DATE TIMESTAMP INTERVAL		X		X			
RAW		X			X		
ROWID UROWID		X				X	
NCHAR NVARCHAR2	X		X	X	X	X	X

The following query shows the use of CAST() to convert literal values to specific types:

```
SELECT
     CAST (12345.67 AS VARCHAR2 (10)),
     CAST('9A4F' AS RAW(2)),
     CAST('05-JUL-07' AS DATE),
     CAST (12345.678 AS NUMBER (10,2))
    FROM dual:
    CAST (12345 CAST CAST ('05- CAST (12345.678ASNUMBER (10,2))
    ______
    12345.67 9A4F 05-JUL-07
```

You can also convert column values from one type to another, as shown in the following query:

```
SELECT
    CAST (price AS VARCHAR2(10)),
    CAST(price + 2 AS NUMBER(7,2)),
    CAST (price AS BINARY DOUBLE)
    FROM products
    WHERE product id = 1;
    CAST (PRICE CAST (PRICE+2ASNUMBER (7,2)) CAST (PRICEASBINARY DOUBLE)
    ______
                              21.95
                                                1.995E+001
```

You'll see additional examples in Chapter 5 that show how to use CAST() to convert dates, times, and intervals. Also, Chapter 13 shows how you use CAST() to convert collections.

Regular Expression Functions

In this section, you'll learn about regular expressions and their associated Oracle database functions. These functions allow you to search for a pattern of characters in a string. For example, let's say you have the following list of years,

1965 1968 1971 1970

> and want to get the years 1965 through to 1968. You can do this using the following regular expression:

```
^196[5-8]$
```

The regular expression contains a number of *metacharacters*. In this example, ^, [5–8], and \$ are the metacharacters; ^ matches the beginning position of a string; [5-8] matches characters between 5 and 8; \$ matches the end position of a string. Therefore ^196 matches a string that begins with 196, and [5-8]\$ matches a string that ends with 5, 6, 7, or 8. So ^196[5-8]\$ matches 1965, 1966, 1967, and 1968, which are the years you wanted to get from the list.

110 Oracle Database 11*g* SQL

The next example uses the following string, which contains a quote from Shakespeare's *Romeo and Juliet*:

But, soft! What light through yonder window breaks?

Let's say you want to get the substring light. You do this using the following regular expression:

l[[:alpha:]]{4}

In this regular expression, [[:alpha:]] and {4} are the metacharacters. [[:alpha:]] matches an alphanumeric character A-Z and a-z; {4} repeats the previous match four times. When 1, [[:alpha:]], and {4} are combined, they match a sequence of five letters starting with 1. Therefore, the regular expression 1[[:alpha:]]{4} matches light in the string.

Table 4-5 lists some the metacharacters you can use in a regular expression, along with their meaning and an example of their use.

Metacharacters	Meaning	Examples
\	Matches a special character or a literal or performs a backreference. (A backreference repeats the previous match.)	\n matches the newline character \\ matches \ \(matches (\) matches)
^	Matches the position at the start of the string.	^A matches A if A is the first character in the string.
\$	Matches the position at the end of the string.	\$B matches B if B is the last character in the string.
*	Matches the preceding character zero or more times.	ba*rk matches brk, bark, baark, and so on.
+	Matches the preceding character one or more times.	ba+rk matches bark, baark, and so on, but not brk.
?	Matches the preceding character zero or one time.	ba?rk matches brk and bark only.
{n}	Matches a character exactly n times, where n is an integer.	hob{2}it matches hobbit.
{n,m}	Matches a character at least n times and at most m times, where n and m are both integers.	hob{2,3}it matches hobbit and hobbbit only.
•	Matches any single character except null.	hob.it matches hobait, hobbit, and so on.

 TABLE 4-5
 Regular Expression Metacharacters

Metacharacters	Meaning	Examples
(pattern)	A subexpression that matches the specified pattern. You use subexpressions to build up complex regular expressions. You can access the individual matches, known as captures, from this type of subexpression.	anatom(y ies) matches anatomy and anatomies.
$x \mid y$	Matches x or y , where x and y are one or more characters.	war peace matches war or peace.
[abc]	Matches any of the enclosed characters.	[ab] be matches abe and bbe.
[a-z]	Matches any character in the specified range.	[a-c]bc matches abc, bbc, and cbc.
[::]	Specifies a character class and matches any character in that class.	[:alphanum:] matches alphanumeric characters 0–9, A–Z, and a–z. [:alpha:] matches alphabetic characters A–Z and a–z. [:blank:] matches space or tab. [:digit:] matches digits 0–9. [:graph:] matches non-blank characters. [:lower:] matches lowercase alphabetic characters a–z. [:print:] is similar to [:graph:] except [:print:] includes the space character. [:punct:] matches punctuation characters.,", and so on. [:space:] matches all whitespace characters. [:upper:] matches all uppercase alphabetic characters A–Z. [:xdigit:] matches characters permissible in a hexadecimal number 0–9, A–F, and a–f.
[]	Matches one collation element, like a multicharacter element.	No example.
[==]	Specifies equivalence classes.	No example.

 TABLE 4-5
 Regular Expression Metacharacters (continued)

Metacharacters	Meaning	Examples
\n	This is a backreference to an earlier capture, where n is a positive integer.	(.)\1 matches two consecutive identical characters. The (.) captures any single character except null, and the \1 repeats the capture, matching the same character again, therefore matching two consecutive identical characters.

 TABLE 4-5
 Regular Expression Metacharacters (continued)

Oracle Database 10g Release 2 introduced a number of Perl-influenced metacharacters, which are shown in Table 4-6.

Table 4-7 shows the regular expression functions. Regular expression functions were introduced in Oracle Database 10*g*, and additional items have been added to 11*g*, as shown in the table.

Metacharacters	Meaning
\d	Digit character
\ D	Non-digit character
\w	Word character
/W	Non-word character
\s	Whitespace character
\S	Non-whitespace character
\A	Matches only at the beginning of a string or before a newline character at the end of a string
\Z	Matches only at the end of a string
*?	Matches the preceding pattern element 0 or more times
+?	Matches the preceding pattern element 1 or more times
??	Matches the preceding pattern element 0 or 1 time
{ n }	Matches the preceding pattern element exactly n times
{n,}	Matches the preceding pattern element at least n times
{ n, m}	Matches the preceding pattern element at least n but not more than m times

 TABLE 4-6
 Perl-Influenced Metacharacters

Function

REGEXP LIKE(x, pattern [, match option])

REGEXP INSTR(x, pattern [, start [, occurrence [, return option [, match option [, subexp option]]])

Description

Searches x for the regular expression defined in the pattern parameter. You can also provide an optional match option, which may be set to one of the following characters:

- 'c', which specifies case-sensitive matching (this is the default)
- 'I', which specifies case-insensitive matching
- 'n', which allows you to use the match-any-character operator
- \blacksquare 'm', which treats x as a multiple line

Searches *x* for the *pattern* and returns the position at which the pattern occurs. You can supply an optional:

- **start** position to begin the search. The default is 1, which is the first character in x.
- occurrence that indicates which occurrence of pattern should be returned. The default is 1, which means the function returns the position of the first occurrence of pattern in x.
- return option that indicates what integer to return. 0 specifies the integer to return is the position of the first character in x; 1 specifies the integer to return is the position of the character in *x* after the occurrence.
- match option to change the default matching. Works in the same way as specified in REGEXP LIKE().
- subexp option (new for Oracle Database 11g) works as follows: for a pattern with subexpressions, subexp option is a nonnegative integer from 0 to 9 indicating which subexpression in pattern is the target of the function. For example, consider the following expression: 0123(((abc)(de)f)ghi)45(678) This expression has five subexpressions in the following order: "abcdefghi", "abcdef", "abc", "de", and "678".

If subexp option is 0, the position of pattern is returned. If pattern does not have the correct number of subexpressions, then the function returns 0. A null subexp option value returns null. The default value for subexp option is 0.

Function

```
REGEXP_REPLACE(x, pattern
[, replace_string
[, start
[, occurrence
[, match_option]]]])
REGEXP_SUBSTR(x, pattern
[, start
[, occurrence
[, match_option
[, subexp_option]]]])
REGEXP_COUNT(x, pattern
[, start
[, match_option]])
```

Description

Searches x for the pattern and replaces it with replace_string. The other options have the same meaning as those shown earlier.

Returns a substring of x that matches pattern; the search begins at the position specified by start. The other options have the same meaning as those shown earlier. The $subexp_option$ (new for Oracle Database 11g) works in the same way as shown for REGEXP INSTR().

New for Oracle Database 11g. Searches in x for the pattern and returns the number of times the pattern is found in x. You can supply an optional:

- *start* position to begin the search. The default is 1, which is the first character in *x*.
- match_option to change the default matching. Works in the same way as shown for REGEXP_LIKE().

 TABLE 4-7
 Regular Expression Functions (continued)

You'll learn more about the regular expression functions in the following sections.

REGEXP_LIKE()

You use REGEXP_LIKE(x, pattern[, $match_option$]) to search x for the regular expression defined in the pattern parameter. You can also provide an optional $match_option$, which may be set to one of the following characters:

- 'c', which specifies case-sensitive matching (this is the default)
- 'I', which specifies case-insensitive matching
- 'n', which allows you to use the match-any-character operator
- \blacksquare 'm', which treats x as a multiple line

The following query retrieves customers whose date of birth is between 1965 and 1968 using REGEXP LIKE():

The next query retrieves customers whose first name starts with J or j. Notice the regular expression passed to REGEXP LIKE() is ^j and the match option is i (i indicates caseinsensitive matching and so in this example ^ i matches J or i).

```
SELECT customer id, first name, last name, dob
    FROM customers
    WHERE REGEXP LIKE (first name, '^j', 'i');
    CUSTOMER ID FIRST NAME LAST NAME DOB
    _____
           1 John Brown 01-JAN-65
```

REGEXP_INSTR()

You use REGEXP INSTR(x, pattern[, start[, occurrence[, return option[, match option]]]]) to search x for the pattern;. This function returns the position at which pattern occurs (positions start at number 1).

The following query returns the position that matches the regular expression 1 [: alpha:]]{4} using REGEXP INSTR():

SELECT

```
REGEXP INSTR('But, soft! What light through yonder window breaks?',
  'l[[:alpha:]]{4}') AS result
FROM dual;
   RESULT
_____
       17
```

Notice that 17 is returned, which is the position of the 1 in light.

The next query returns the position of the second occurrence that matches the regular expression s [[:alpha:]] {3} starting at position 1:

SELECT

```
REGEXP INSTR('But, soft! What light through yonder window softly breaks?',
  's[[:alpha:]]{3}', 1, 2) AS result
FROM dual;
   RESULT
-----
```

The next query returns the position of the second occurrence that matches the letter o starting the search at position 10:

SELECT

```
REGEXP INSTR('But, soft! What light through yonder window breaks?',
  'o', 10, 2) AS result
FROM dual:
   RESULT
_____
       32
```

REGEXP_REPLACE()

You use REGEXP_REPLACE(x, pattern[, replace_string[, start[, occurrence[, match_option]]]]) to search x for the pattern and replace it with replace_string. The following query replaces the substring that matches the regular expression 1[[: alpha:]]{4} with the string 'sound' using REGEXP_REPLACE():

SELECT

```
REGEXP_REPLACE('But, soft! What light through yonder window breaks?',
  'l[[:alpha:]]{4}', 'sound') AS result
FROM dual;
```

```
RESULT
```

But, soft! What sound through yonder window breaks?

Notice that light has been replaced by sound.

REGEXP_SUBSTR()

You use REGEXP_SUBSTR(x, pattern [, start [, occurrence [, match_option]]]) to get a substring of x that matches pattern; the search begins at the position specified by start.

The following query returns the substring that matches the regular expression 1[[:alpha:]]{4} using REGEXP SUBSTR():

SELECT

```
REGEXP_SUBSTR('But, soft! What light through yonder window breaks?',
  'l[[:alpha:]]{4}') AS result
FROM dual;
```

RESUL

light

REGEXP COUNT()

REGEXP_COUNT() is new for Oracle Database 11g. You use REGEXP_COUNT(x, pattern [, start [, match_option]]) to search in x for the pattern and get the number of times pattern is found in x. You can provide an optional start number to indicate the character in x to begin searching for pattern and an optional match_option string to indicate the match option.

The following query returns the number of times the regular expression s[[:alpha:]]{3} occurs in a string using REGEXP COUNT():

SELECT

```
REGEXP_COUNT('But, soft! What light through yonder window softly breaks?',
   's[[:alpha:]]{3}') AS result
FROM dual;
```

```
RESULT
```

Notice that 2 is returned, which means the regular expression has two matches in the supplied string.

Using Aggregate Functions

The functions you've seen so far operate on a single row at a time and return one row of output for each input row. In this section, you'll learn about aggregate functions, which operate on a group of rows and return one row of output.



NOTE

Aggregate functions are also known as group functions because they operate on groups of rows.

Table 4-8 lists some of the aggregate functions, all of which return a NUMBER. Here are some points to remember when using aggregate functions:

- You can use the aggregate functions with any valid expression. For example, you can use the COUNT(), MAX(), and MIN() functions with numbers, strings, and datetimes.
- Null values are ignored by aggregate functions, because a null value indicates the value is unknown and therefore cannot be used in the aggregate function's calculation.
- You can use the DISTINCT keyword with an aggregate function to exclude duplicate entries from the aggregate function's calculation.

You'll learn more about some of the aggregate functions shown in Table 4-8 in the following sections. In Chapters 7 and 8, you'll see how to use these functions in conjunction with the SELECT statement's ROLLUP and RETURNING clauses. As you'll see, ROLLUP allows you to get a subtotal for a group of rows, where the subtotal is calculated using one of the aggregate functions; RETURNING allows you to store the value returned by an aggregate function in a variable.

Function	Description
AVG(x)	Returns the average value of x
COUNT (x)	Returns the number of rows returned by a query involving x
MAX (x)	Returns the maximum value of x
MEDIAN(x)	Returns the median value of x
MIN(x)	Returns the minimum value of x
STDDEV(x)	Returns the standard deviation of x
SUM(x)	Returns the sum of x
VARIANCE(x)	Returns the variance of x

 TABLE 4-8
 Aggregate Functions

AVG()

You use AVG(x) to get the average value of x. The following query gets the average price of the products; notice that the price column from the products table is passed to the AVG() function:

SELECT AVG(price) FROM products;

```
AVG(PRICE)
-----
19.7308333
```

You can use the aggregate functions with any valid expression. For example, the following query passes the expression price + 2 to AVG(); this adds 2 to each row's price and then returns the average of those values.

SELECT AVG(price + 2) FROM products;

```
AVG(PRICE)
------
21.7308333
```

You can use the DISTINCT keyword to exclude identical values from a computation. For example, the following query uses the DISTINCT keyword to exclude identical values in the price column when computing the average using AVG():

SELECT AVG(DISTINCT price) FROM products;

```
AVG (DISTINCTPRICE)
-----
20.2981818
```

Notice that the average in this example is slightly higher than the average returned by the first query in this section. This is because the value for product #12 (13.49) in the price column is the same as the value for product #7; it is considered a duplicate and excluded from the computation performed by AVG (). Therefore, the average is slightly higher in this example.

COUNT()

You use COUNT(x) to get the number of rows returned by a query. The following query gets the number of rows in the products table using COUNT():

SELECT COUNT(product_id) FROM products;

```
COUNT (PRODUCT_ID)
```



You should avoid using the asterisk (*) with the COUNT () function, as it may take longer for COUNT () to return the result. Instead, you should use a column in the table or use the ROWID pseudo column. (As you saw in Chapter 2, the ROWID column contains the internal location of the row in the Oracle database.)

The following example passes ROWID to COUNT () and gets the number of rows in the products table:

SELECT COUNT (ROWID) FROM products;

```
COUNT (ROWID)
_____
```

MAX() and MIN()

You use MAX (x) and MIN (x) to get the maximum and minimum values for x. The following query gets the maximum and minimum values of the price column from the products table using MAX() and MIN():

SELECT MAX(price), MIN(price) FROM products;

```
MAX(PRICE) MIN(PRICE)
    49.99 10.99
```

You may use MAX() and MIN() with any type, including strings and dates. When you use MAX() with strings, the strings are ordered alphabetically with the "maximum" string being at the bottom of a list and the "minimum" string being at the top of the list. For example, the string Albert would appear before Zeb in such a list. The following example gets the maximum and minimum name strings from the products table using MAX() and MIN():

SELECT MAX (name) , MIN (name) FROM products;

```
MAX (NAME)
Z Files
                                 2412: The Return
```

In the case of dates, the "maximum" date occurs at the latest point in time, and the "minimum" date at the earliest point in time. The following query gets the maximum and minimum dob from the customers table using MAX() and MIN():

SELECT MAX (dob), MIN (dob) FROM customers;

```
MAX (DOB) MIN (DOB)
_____
16-MAR-71 01-JAN-65
```

STDDEV()

You use STDDEV (x) to get the standard deviation of x. Standard deviation is a statistical function and is defined as the square root of the variance (you'll learn about variance shortly).

The following query gets the standard deviation of the price column values from the products table using STDDEV():

SELECT STDDEV(price) FROM products;

```
STDDEV(PRICE)
-----
11.0896303
```

SUM()

SUM(x) adds all the values in x and returns the total. The following query gets the sum of the price column from the products table using SUM():

SELECT SUM(price) FROM products;

```
SUM(PRICE)
-------
236.77
```

VARIANCE()

You use VARIANCE (x) to get the variance of x. Variance is a statistical function and is defined as the spread or variation of a group of numbers in a sample. Variance is equal to the square of the standard deviation.

The following example gets the variance of the price column values from the products table using VARIANCE():

SELECT VARIANCE(price) FROM products;

```
VARIANCE (PRICE)
```

122.979899

Grouping Rows

Sometimes you need to group blocks of rows in a table and get some information on those groups of rows. For example, you might want to get the average price for the different types of products in the products table. You'll see how to do this the hard way first, then you'll see the easy way, which involves using the GROUP BY clause to group similar rows together.

To do it the hard way, you limit the rows passed to the AVG() function using a WHERE clause. For example, the following query gets the average price for books from the products table (books have a product_type_id of 1):

```
SELECT AVG(price)
   FROM products
   WHERE product type id = 1;
   AVG (PRICE)
   _____
       24.975
```

To get the average price for the other types of products, you would need to perform additional queries with different values for the product type id in the WHERE clause. As you can imagine, this is very labor intensive. You'll be glad to know there's an easier way to do this through the use of the GROUP BY clause.

Using the GROUP BY Clause to Group Rows

You use the GROUP BY clause to group rows into blocks with a common column value. For example, the following query groups the rows from the products table into blocks with the same product type id:

```
SELECT product type id
    FROM products
    GROUP BY product type id;
```

```
PRODUCT_TYPE_ID
               3
```

Notice that there's one row in the result set for each block of rows with the same product type id and that there's a gap between 1 and 2 (you'll see why this gap occurs shortly). In the result set, there's one row for products with a product type id of 1, another for products with a product type id of 2, and so on. There are actually two rows in the products table with a product type id of 1, four rows with a product type id of 2, and so on for the other rows in the table. These rows are grouped together into separate blocks by the GROUP BY clause, one block for each product type id. The first block contains two rows, the second contains four rows, and so on.

The gap between 1 and 2 is caused by a row whose product type id is null. This row is shown in the following example:

```
SELECT product id, name, price
     FROM products
     WHERE product type id IS NULL;
```

```
PRODUCT ID NAME
12 My Front Line
                   13.49
```

Because this row's product_type_id is null, the GROUP BY clause in the earlier query groups this row into a single block. The row in the result set is blank because the product_type_id is null for the block, so there's a gap between 1 and 2.

Using Multiple Columns in a Group

You can specify multiple columns in a GROUP BY clause. For example, the following query includes the product_id and customer_id columns from the purchases table in a GROUP BY clause:

SELECT product_id, customer_id FROM purchases GROUP BY product_id, customer_id;

PRODUCT_ID	CUSTOMER_ID
1	1
1	2
1	3
1	4
2	1
2	2
2	3
2	4
3	3

Using Groups of Rows with Aggregate Functions

You can pass blocks of rows to an aggregate function. The aggregate function performs its computation on the group of rows in each block and returns one value per block. For example, to get the number of rows with the same product_type_id from the products table, you do the following:

- Use the GROUP BY clause to group rows into blocks with the same product type_id.
- Use COUNT (ROWID) to get the number of rows in each block.

The following query shows this:

Notice that there are five rows in the result set, with each row corresponding to one or more rows in the products table grouped together with the same product type id. From the result set, you can see there are two rows with a product type id of 1, four rows with a product type id of 2, and so on. The last line in the result set shows there is one row with a null product type id (this is caused by the "My Front Line" product mentioned earlier).

Let's take a look at another example. To get the average price for the different types of products in the products table, you do the following:

- Use the GROUP BY clause to group rows into blocks with the same product type id.
- Use AVG (price) to get the average price for each block of rows.

The following query shows this:

```
SELECT product type id, AVG(price)
     FROM products
     GROUP BY product type id
     ORDER BY product type id;
     PRODUCT TYPE ID AVG(PRICE)
                 1 24.975
                      26.22
                  3 13.24
4 13.99
                        13.49
```

Each group of rows with the same product type id is passed to the AVG() function. AVG () then computes the average price for each group. As you can see from the result set, the average price for the group of products with a product type id of 1 is 24.975. Similarly, the average price of the products with a product type id of 2 is 26.22. Notice that the last row in the result set shows an average price of 13.49; this is simply the price of the "My Front Line" product, the only row with a null product type id.

You can use any of the aggregate functions with the GROUP BY clause. For example, the next query gets the variance of product prices for each product type id:

```
SELECT product_type_id, VARIANCE(price)
      FROM products
      GROUP BY product type id
      ORDER BY product type id;
      PRODUCT_TYPE_ID VARIANCE(PRICE)
                   50.50125
2 280.8772
3 4
```

One point to remember is that you don't have to include the columns used in the GROUP BY in the list of columns immediately after the SELECT. For example, the following query is the same as the previous one except product type id is omitted from the SELECT clause:

You can also include an aggregate function call in the ORDER BY, as shown in the following query:

```
SELECT VARIANCE (price)
FROM products
GROUP BY product_type_id
ORDER BY VARIANCE (price);

VARIANCE (PRICE)

0
.125
7
50.50125
280.8772
```

Incorrect Usage of Aggregate Function Calls

When your query contains an aggregate function—and retrieves columns not placed within an aggregate function—those columns must be placed in a GROUP BY clause. If you forget to do this, you'll get the following error: ORA-00937: not a single-group group function. For example, the following query attempts to retrieve the product_type_id column and AVG (price) but omits a GROUP BY clause for product type id:

```
SQL> SELECT product_type_id, AVG(price)

2 FROM products;

SELECT product_type_id, AVG(price)

*

ERROR at line 1:

ORA-00937: not a single-group group function
```

The error occurs because the database doesn't know what to do with the product_type_id column. Think about it: the query attempts to use the AVG() aggregate function, which operates

on multiple rows, but also attempts to get the product type id column values for each individual row. You can't do both at the same time. You must provide a GROUP BY clause to tell the database to group multiple rows with the same product type id together; then the database passes those groups of rows to the AVG () function.



CAUTION

When a query contains an aggregate function—and retrieves columns not placed within an aggregate function—then those columns must be placed in a GROUP BY clause.

Also, you cannot use an aggregate function to limit rows in a WHERE clause. If you try to do so, you will get the following error: ORA-00934: group function is not allowed here. For example:

```
SQL> SELECT product type id, AVG(price)
     2 FROM products
     3 WHERE AVG(price) > 20
     4 GROUP BY product_type_id;
   WHERE AVG(price) > 20
   ERROR at line 3:
   ORA-00934: group function is not allowed here
```

The error occurs because you may only use the WHERE clause to filter individual rows, not groups of rows. To filter groups of rows, you use the HAVING clause, which you'll learn about next.

Using the HAVING Clause to Filter Groups of Rows

You use the HAVING clause to filter groups of rows. You place the HAVING clause after the GROUP BY clause:

```
SELECT ...
   FROM ...
   WHERE
   GROUP BY ...
   HAVING ...
   ORDER BY ...;
```



GROUP BY can be used without HAVING, but HAVING must be used in conjunction with GROUP BY.

Let's take a look at an example. Say you want to view the types of products that have an average price greater than \$20. To do this, you do the following:

- Use the GROUP BY clause to group rows into blocks with the same product type id.
- Use the HAVING clause to limit the returned results to those groups that have an average price greater than \$20.

The following query shows this:

As you can see, only the groups of rows having an average price greater than \$20 are displayed.

Using the WHERE and GROUP BY Clauses Together

You can use the WHERE and GROUP BY clauses together in the same query. When you do this, first the WHERE clause filters the rows returned, then the GROUP BY clause groups the remaining rows into blocks. For example, the following query uses

- A WHERE clause to filter the rows from the products table to select those whose price is less than \$15.
- A GROUP BY clause to group the remaining rows by the product type id column.

Using the WHERE, GROUP BY, and HAVING Clauses Together

You can use the WHERE, GROUP BY, and HAVING clauses together in the same query. When you do this, the WHERE clause first filters the rows, the GROUP BY clause then groups the remaining rows into blocks, and finally the HAVING clause filters the row groups. For example, the following query uses

- A WHERE clause to filter the rows from the products table to select those whose price is less than \$15.
- A GROUP BY clause to group the remaining rows by the product type id column.
- A HAVING clause to filter the row groups to select those whose average price is greater than \$13.

```
SELECT product type id, AVG(price)
   FROM products
   WHERE price < 15
   GROUP BY product type id
   HAVING AVG(price) > 13
   ORDER BY product type id;
   PRODUCT TYPE ID AVG(PRICE)
                2 14.45
3 13.24
                       13.49
```

Compare these results with the previous example. Notice that the group of rows with the product type id of 4 is filtered out. That's because the group of rows has an average price less than \$13.

The final query uses ORDER BY AVG (price) to re-order the results by the average price:

```
SELECT product type id, AVG(price)
    FROM products
    WHERE price < 15
    GROUP BY product type id
    HAVING AVG(price) > 13
    ORDER BY AVG(price);
    PRODUCT TYPE ID AVG(PRICE)
```

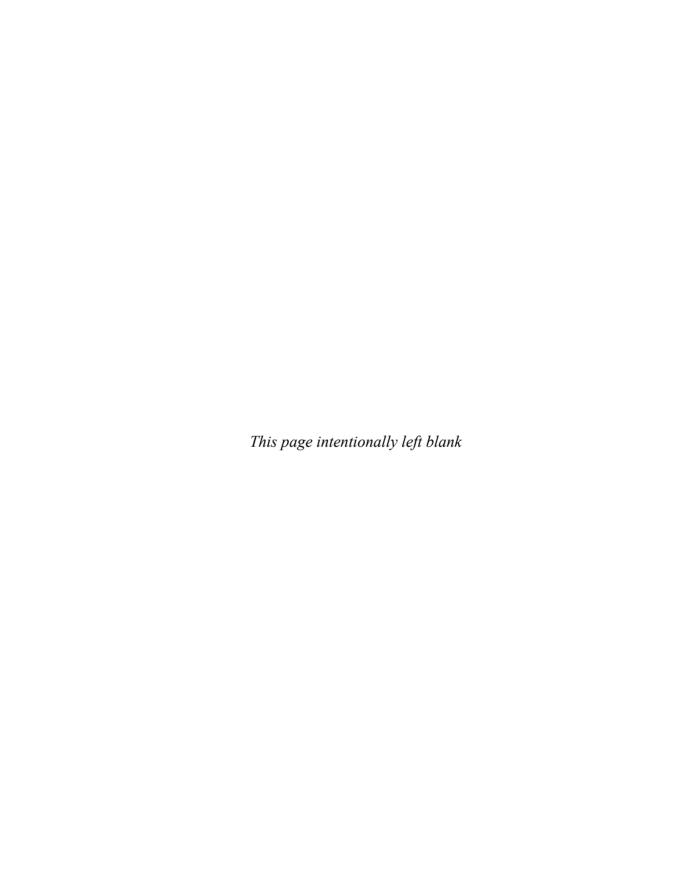
_	-		
	3	13.	24
		13.	49
	2	14.	45

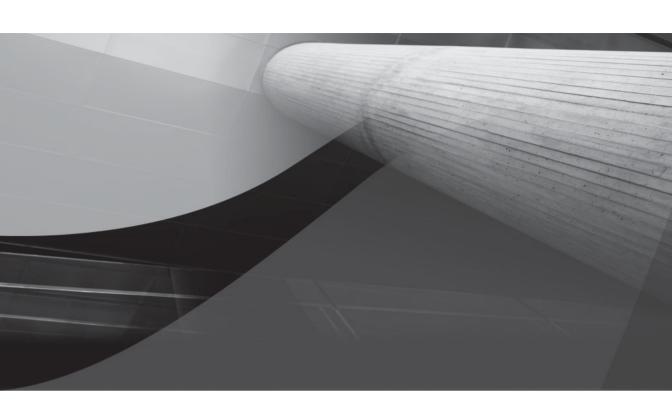
Summary

In this chapter, you have learned the following:

- The Oracle database has two main groups of functions: single-row functions and aggregate functions.
- Single-row functions operate on one row at a time and return one row of output for each input row. There are five main types of single-row functions: character functions, numeric functions, conversion functions, date functions, and regular expression functions.
- Aggregate functions operate on multiple rows and return one row of output.
- Blocks of rows may be grouped together using the GROUP BY clause.
- Groups of rows may be filtered using the HAVING clause.

In the next chapter, you'll learn about dates and times.





CHAPTER 5

Storing and Processing Dates and Times



n this chapter, you will see how to

- Process and store a specific date and time, collectively known as a datetime. An example of a datetime is 7:15:30 p.m. on October 10, 2007. You store a datetime using the DATE type. The DATE type stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second.
- Use *timestamps* to store a specific date and time. A timestamp stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second. The advantages of a timestamp over a DATE are that a timestamp can store a fractional second and a time zone.
- Use time *intervals* to store a length of time. An example of a time interval is 1 year 3 months.

Let's plunge in and see some simple examples of storing and retrieving dates.

Simple Examples of Storing and Retrieving Dates

By default the database uses the format DD-MON-YYYY to represent a date, where

- DD is a two-digit day, e.g., 05
- MON is the first three letters of the month, e.g., FEB
- YYYY is a four-digit year, e.g., 1968

Let's take a look at an example of adding a row to the customers table, which contains a DATE column named dob. The following INSERT adds a row to the customers table, setting the dob column to 05-FEB-1968:

```
INSERT INTO customers (
    customer_id, first_name, last_name, dob, phone
) VALUES (
    6, 'Fred', 'Brown', '05-FEB-1968', '800-555-1215'
);
```

You can also use the DATE keyword to supply a date literal to the database. The date must use the ANSI standard date format YYYY-MM-DD, where

- YYYY is a four-digit year.
- MM is a two-digit month from 1 to 12.
- DD is a two-digit day.



TIP

Using ANSI standard dates in SQL statements has the advantage that those statements could potentially run against non-Oracle databases.

For example, to specify a date of October 25, 1972, you use DATE '1972-10-25'. The following INSERT adds a row to the customers table, specifying DATE '1972-10-25' for the dob column:

```
INSERT INTO customers (
  customer id, first name, last name, dob, phone
) VALUES (
  7, 'Steve', 'Purple', DATE '1972-10-25', '800-555-1215'
);
```

By default, the database returns dates in the format DD-MON-YY, where YY are the last two digits of the year. For example, the following example retrieves rows from the customers table and then performs a ROLLBACK to undo the results of the previous two INSERT statements; notice the two-digit years in the dob column returned by the query:

SELECT * FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	
6	Fred	Brown	05-FEB-68	800-555-1215
7	Steve	Purple	25-OCT-72	800-555-1215

ROLLBACK;

Customer #4's dob is null and is therefore blank in the previous result set.



NOTE

If you actually ran the two INSERT statements, make sure you undo the changes by executing the ROLLBACK. That way, you'll keep the database in its initial state, and the results from your queries will match those in this chapter.

In this section, you saw some simple examples of using dates that use default formats. You'll learn how to provide your own date formats in the following section and see how to convert a datetime to another database type.

Converting Datetimes Using TO_CHAR() and TO_DATE()

The Oracle database has functions that enable you to convert a value in one data type to another. You saw some of these functions in the previous chapter. In this section, you'll see how to use the TO CHAR() and TO DATE() functions to convert a datetime to a string and vice versa. Table 5-1 summarizes the TO CHAR() and TO DATE() functions.

Function	Description
TO_CHAR(x [, format])	Converts x to a string. You can also supply an optional format for x. You saw how to use TO_CHAR() to convert a number to a string in the previous chapter. In this chapter, you'll see how to convert a datetime to a string.
TO_DATE(x [, format])	Converts the string x to a DATE.

 TABLE 5-1
 TO CHAR () and TO DATE ()
 Conversion Functions

Let's start off by examining how you use TO_CHAR() to convert a datetime to a string. Later, you'll see how to use TO_DATE() to convert a string to a DATE.

Using TO_CHAR() to Convert a Datetime to a String

You can use $TO_CHAR(x [, format])$ to convert the datetime x to a string. You can also provide an optional format for x. An example format is MONTH DD, YYYY, where

- MONTH is the full name of the month in uppercase, e.g., JANUARY.
- DD is the two-digit day.
- YYYY is the four-digit year.

The following query uses TO_CHAR() to convert the dob column from the customers table to a string with the format MONTH DD, YYYY:

SELECT customer_id, TO_CHAR(dob, 'MONTH DD, YYYY') FROM customers;

The next query gets the current date and time from the database using the SYSDATE function, then converts the date and time to a string using TO_CHAR() with the format MONTH DD, YYYY, HH24:MI:SS. The time portion of this format indicates that the hours are in 24-hour format and that the minutes and seconds are also to be included in the string.

SELECT TO_CHAR(SYSDATE, 'MONTH DD, YYYY, HH24:MI:SS') FROM dual;

```
TO_CHAR(SYSDATE,'MONTHDD,YYY
------
NOVEMBER 05, 2007, 12:34:36
```

When you use TO CHAR() to convert a datetime to a string, the format has a number of parameters that affect the returned string. Some of these parameters are listed in Table 5-2.

Aspect	Parameter	Description	Example
Century	CC	Two-digit century.	21
	SCC	Two-digit century with a negative sign (–) for B.C.	-10
Quarter	Q	One-digit quarter of the year.	1
Year	YYYY	All four digits of the year.	2008
	IYYY	All four digits of the ISO year.	2008
	RRRR	All four digits of the rounded year (governed by the present year). See the section "How Oracle Interprets Two-Digit Years" later in this chapter for details.	2008
	SYYYY	All four digits of the year with a negative sign (–) for B.C.	-1001
	Υ,ΥΥΥ	All four digits of the year, with a comma after the first digit.	2,008
	YYY	Last three digits of the year.	008
	IYY	Last three digits of the ISO year.	008
	YY	Last two digits of the year.	08
	IY	Last two digits of the ISO year.	06
	RR	Last two digits of the rounded year, which depend on the present year. See the section "How Oracle Interprets Two-Digit Years" later in this chapter for details.	08
	Y	Last digit of the year.	8
	I	Last digit of the ISO year.	8
	YEAR	Name of the year in uppercase.	TWO THOUSAND-EIGHT
	Year	Name of the year with the first letter in uppercase.	Two Thousand-Eight
Month	MM	Two-digit month of the year.	01
	MONTH	Full name of the month in uppercase, right- padded with spaces to a total length of nine characters.	JANUARY
	Month	Full name of the month with first letter in uppercase, right-padded with spaces to a total length of nine characters.	January
	MON	First three letters of the name of the month in uppercase.	JAN
	Mon	First three letters of the name of the month with the first letter in uppercase.	Jan

 TABLE 5-2
 Datetime Formatting Parameters

Aspect	Parameter	Description	Example
	RM	Roman numeral month.	The Roman numeral month for the fourth month (April) is IV.
Week	WW	Two-digit week of the year.	02
	IW	Two-digit ISO week of the year.	02
	W	One-digit week of the month.	2
Day	DDD	Three-digit day of the year.	103
	DD	Two-digit day of the month.	31
	D	One-digit day of the week.	5
	DAY	Full name of the day in uppercase.	SATURDAY
	Day	Full name of the day with the first letter in uppercase.	Saturday
	DY	First three letters of the name of the day in uppercase.	SAT
	Dy	First three letters of the name of the day with the first letter in uppercase.	Sat
	J	Julian day—the number of days that have passed since January 1, 4713 B.C.	2439892
Hour	HH24	Two-digit hour in 24-hour format.	23
	НН	Two-digit hour in 12-hour format.	11
Minute	MI	Two-digit minute.	57
Second	SS	Two-digit second.	45
	FF[19]	Fractional seconds with an optional number of digits to the right of the decimal point. Only applies timestamps, which you'll learn about in the section "Using Timestamps" later in this chapter.	When dealing with 0.123456789 seconds, FF3 would round the seconds to 0.123.
	SSSSS	Number of seconds past 12 A.M.	46748
	MS	Millisecond (millionths of a second).	100
	CS	Centisecond (hundredths of a second).	10
Separators	-/,.;: "text"	Characters that allow you to separate the aspects of a date and time. You can supply freeform text in quotes as a separator.	For the date December 13, 1969, DD-MM-YYYY would produce 12-13-1969, and DD/MM/YYYY would produce 12/13/1969.
Suffixes	AM Or PM	AM or PM as appropriate.	AM
	A.M. or P.M.	A.M. or P.M. as appropriate.	P.M.
	AD or BC	AD or BC as appropriate.	AD

 TABLE 5-2
 Datetime Formatting Parameters (continued)

Aspect	Parameter	Description	Example
	A.D. or B.C.	A.D. or B.C. as appropriate.	B.C.
	TH	Suffix to a number. You can make the suffix uppercase by specifying the numeric format in uppercase and vice versa for lowercase.	For a day number of 28, ddTH would produce 28 th , and DDTH would produce 28TH.
	SP	Number is spelled out.	For a day number of 28, DDSP would produce TWENTY-EIGHT, and ddSP would produce twenty-eight.
	SPTH	Combination of TH and SP.	For a day number of 28, DDSPTH would produce TWENTY-EIGHTH, and ddSPTH would produce twenty-eighth.
Era	EE	Full era name for Japanese Imperial, ROC Official, and Thai Buddha calendars.	No example
	E	Abbreviated era name.	No example
Time zones	TZH	Time zone hour. You'll learn about time zones later in the section "Using Time Zones."	12
	TZM	Time zone minute.	30
	TZR	Time zone region.	PST
	TZD	Time zone with daylight savings information.	No example

 TABLE 5-2
 Datetime Formatting Parameters (continued)

The following table shows examples of strings to format the date February 5, 1968, along with the string returned from a call to TO_CHAR().

Format String	Returned String
MONTH DD, YYYY	FEBRUARY 05, 1968
MM/DD/YYYY	02/05/1968
MM-DD-YYYY	02-05-1968
DD/MM/YYYY	05/02/1968
DAY MON, YY AD	MONDAY FEB, 68 AD
DDSPTH "of" MONTH, YEAR A.D.	FIFTH of FEBRUARY, NINETEEN SIXTY-EIGHT A.D.
CC, SCC	20, 20
Q	1

Format String

```
YYYY, IYYY, RRRR, SYYYY, Y,YYY,
YYY, IYY, YY, IY, RR, Y, I,
YEAR,
Year
MM, MONTH, Month,
MON, Mon, RM
WW, IW, W
DDD, DD, DAY,
Day, DY, Dy, J
ddTH, DDTH, ddSP, DDSP, DDSPTH
```

Returned String

1968, 1968, 1968, 1968, 1,968, 968, 968, 68, 68, 68, 8, 8, NINETEEN SIXTY-EIGHT, Nineteen Sixty-Eight 02, FEBRUARY, February, FEB, Feb, II 06, 06, 1 036, 05, MONDAY, Monday, MON, Mon, 2439892 05th, 05TH, five, FIVE, FIFTH

You can see the results shown in this table by calling TO_CHAR() in a query. For example, the following query converts February 5, 1968, to a string with the format MONTH DD, YYYY:

SELECT TO_CHAR(TO_DATE('05-FEB-1968'), 'MONTH DD, YYYY')
FROM dual;

```
TO_CHAR(TO_DATE('0
------
FEBRUARY 05, 1968
```



NOTE

The TO_DATE() function converts a string to a datetime. You'll learn more about the TO_DATE() function shortly.

The following table shows examples of strings to format the time 19:32:36 (32 minutes and 36 seconds past 7 P.M.)—along with the output that would be returned from a call to ${\tt TO_CHAR}$ () with that time and format string.

Format String Returned String
HH24:MI:SS 19:32:36
HH.MI.SS AM 7.32.36 PM

Using TO_DATE() to Convert a String to a Datetime

You use TO_DATE (x [, format]) to convert the x string to a datetime. You can provide an optional format string to indicate the format of x. If you omit format, the date must be in the default database format (usually DD-MON-YYYYY or DD-MON-YYY).



NOTE

The NLS_DATE_FORMAT database parameter specifies the default date format for the database. As you'll learn later in the section "Setting the Default Date Format," you can change the setting of NLS_DATE_FORMAT.

The following query uses TO_DATE() to convert the strings 04-JUL-2007 and 04-JUL-07 to the date July 4, 2007; notice that the final date is displayed in the default format of DD-MON-YY:

SELECT TO DATE('04-JUL-2007'), TO DATE('04-JUL-07') FROM dual:

```
TO DATE (' TO DATE ('
_____
04-JUL-07 04-JUL-07
```

Specifying a Datetime Format

As mentioned earlier, you can supply an optional format for a datetime to TO DATE(). You use the same format parameters as those defined previously in Table 5-2. The following query uses TO DATE() to convert the string July 4, 2007 to a date, passing the format string MONTH DD, YYYY to TO_DATE():

```
SELECT TO_DATE('July 4, 2007', 'MONTH DD, YYYY')
    FROM dual;
```

```
TO DATE ('
_____
04-JUL-07
```

The next query passes the format string MM. DD. YY to TO DATE () and converts the string 7.4.07 to the date July 4, 2007; again, the final date is displayed in the default format DD-MON-YY:

```
SELECT TO DATE('7.4.07', 'MM.DD.YY')
    FROM dual;
    TO DATE ('
    -----
    04-JUL-07
```

Specifying Times

You can also specify a time with a datetime. If you don't supply a time with a datetime, the time part of your datetime defaults to 12:00:00 A.M. You can supply the format for a time using the various formats shown earlier in Table 5-3. One example time format is HH24:MI:SS, where

- HH24 is a two-digit hour in 24-hour format from 00 to 23.
- MI is a two-digit minute from 00 to 59.
- SS is a two-digit second from 00 to 59.

An example of a time that uses the HH24:MI:SS format is 19:32:36. A full example of a datetime that uses this time is

```
05-FEB-1968 19:32:36
```

with the format for this datetime being

```
DD-MON-YYYY HH24:MI:SS
```

The following TO DATE () call shows the use of this datetime format and value:

```
TO DATE('05-FEB-1968 19:32:36', 'DD-MON-YYYY HH24:MI:SS')
```

The datetime returned by TO_DATE() in the previous example is used in the following INSERT that adds a row to the customers table; notice that the dob column for the new row is set to the datetime returned by TO_DATE():

```
INSERT INTO customers (
   customer_id, first_name, last_name,
   dob,
   phone
) VALUES (
   6, 'Fred', 'Brown',
   TO_DATE('05-FEB-1968 19:32:36', 'DD-MON-YYYY HH24:MI:SS'),
   '800-555-1215'
);
```

You use TO_CHAR() to view the time part of a datetime. For example, the following query retrieves the rows from the customers table and uses TO_CHAR() to convert the dob column values; notice that customer #6 has the time previously set in the INSERT:

```
SELECT customer_id, TO_CHAR(dob, 'DD-MON-YYYY HH24:MI:SS')
FROM customers;
```

Notice the time for the dob column for customers #1, #2, #3, and #5 is set to 00:00:00 (12 a.m.). This is the default time substituted by the database when you don't provide a time in a datetime.

The next statement rolls back the addition of the new row:

ROLLBACK;



NOTE

If you actually ran the earlier INSERT statement, make sure you undo the change using ROLLBACK.

Combining TO_CHAR() and TO_DATE() Calls

You can combine ${\tt TO_CHAR}$ () and ${\tt TO_DATE}$ () calls; doing this allows you to use datetimes in different formats. For example, the following query combines ${\tt TO_CHAR}$ () and ${\tt TO_DATE}$ () in order to view just the time part of a datetime; notice that the output from ${\tt TO_DATE}$ () is passed to ${\tt TO_CHAR}$ ():

```
SELECT TO CHAR (TO DATE ('05-FEB-1968 19:32:36',
   'DD-MON-YYYY HH24:MI:SS'), 'HH24:MI:SS')
  FROM dual;
  TO CHAR (
  _____
  19:32:36
```

Setting the Default Date Format

The default date format is specified in the NLS DATE FORMAT database parameter. A DBA can change the setting of NLS DATE FORMAT by setting this parameter's value in the database's init.ora or spfile.ora file, both of which are read when the database is started. A DBA can also set NLS DATE FORMAT using the ALTER SYSTEM command. You can also set the NLS DATE FORMAT parameter for your own session using SQL*Plus, which you do by using the ALTER SESSION command. For example, the following ALTER SESSION statement sets NLS DATE FORMAT to MONTH-DD-YYYY:

```
ALTER SESSION SET NLS DATE FORMAT = 'MONTH-DD-YYYY';
```

Session altered



A session is started when you connect to a database and is ended when you disconnect.

You can see the use of this new date format in the results from the following query that retrieves the dob column for customer #1:

```
SELECT dob
   FROM customers
   WHERE customer id = 1;
    _____
    JANUARY -01-1965
```

You may also use the new date format when inserting a row in the database. For example, the following INSERT adds a new row to the customers table; notice the use of the format MONTH-DD-YYYY when supplying the dob column's value:

```
INSERT INTO customers (
     customer id, first name, last name, dob, phone
   ) VALUES (
     6, 'Fred', 'Brown', 'MARCH-15-1970', '800-555-1215'
   );
```

Go ahead and disconnect from the database and connect again as the store user; you'll find that the date format is back to the default. That's because any changes you make using the ALTER SESSION statement last only for that particular session—when you disconnect, you lose the change.



NOTE

If you ran the previous INSERT statement, go ahead and delete the row using DELETE FROM customers WHERE customer_id = 6.

How Oracle Interprets Two-Digit Years

The Oracle database stores all four digits of the year, but if you supply only two digits, the database will interpret the century according to whether the YY or RR format is used.



TIP

You should always specify all four digits of the year. That way, you won't get confused as to which year you mean.

Let's take a look at the YY format first, followed by the RR format.

Using the YY Format

If your date format uses YY for the year and you supply only two digits of a year, then the century for your year is assumed to be the same as the present century currently set on the database server. Therefore, the first two digits of your supplied year are set to the first two digits of the present year. For example, if your supplied year is 15 and the present year is 2007, your supplied year is set to 2015; similarly, a supplied year of 75 is set to 2075.



NOTE

If you use the YYYY format but only supply a two-digit year, then your year is interpreted using the YY format.

Let's take a look at a query that uses the YY format for interpreting the years 15 and 75. In the following query, notice that the input dates 15 and 75 are passed to TO_DATE(), whose output is passed to TO_CHAR(), which converts the dates to a string with the format DD-MON-YYYY. (The YYYY format is used here, so you can see that all four digits of the year returned by TO_DATE().)

SELECT

As expected, the years 15 and 75 are interpreted as 2015 and 2075.

Using the RR Format

If your date format is RR and you supply the last two digits of a year, the first two digits of your year are determined using the two-digit year you supply (your *supplied year*) and the last two

digits of the present date on the database server (the present year). The rules used to determine the century of your supplied year are as follows:

- **Rule 1** If your supplied year is between 00 and 49 and the present year is between 00 and 49, the century is the same as the present century. Therefore, the first two digits of your supplied year are set to the first two digits of the present year. For example, if your supplied year is 15 and the present year is 2007, your supplied year is set to 2015.
- Rule 2 If your supplied year is between 50 and 99 and the present year is between 00 and 49, the century is the present century minus 1. Therefore, the first two digits of your supplied year are set to the present year's first two digits minus 1. For example, if your supplied year is 75 and the present year is 2007, your supplied year is set to 1975.
- **Rule 3** If your supplied year is between 00 and 49 and the present year is between 50 and 99, the century is the present century plus 1. Therefore, the first two digits of your supplied year are set to the present year's first two digits plus 1. For example, if your supplied year is 15 and the present year is 2075, your supplied year is set to 2115.
- **Rule 4** If your supplied year is between 50 and 99 and the present year is between 50 and 99, the century is the same as the present century. Therefore, the first two digits of your supplied year are set to the first two digits of the present year. For example, if your supplied year is 55 and the present year is 2075, your supplied year is set to 2055.

Table 5-3 summarizes these results.



NOTE

If you use the RRRR format but supply only a two-digit year, then your year is interpreted using the RR format.

		Two-Digit Supplied Year	
		00–49	50-99
Last Two Digits of Present Year	00–49	Rule 1: First two digits of supplied year are set to first two digits of present year.	Rule 2: First two digits of supplied year are set to present year's first two digits minus 1.
	50–99	Rule 3: First two digits of supplied year are set to present year's first two digits plus 1.	Rule 4: First two digits of supplied year are set to first two digits of present year.

TABLE 5-3 How Two-Digit Years Are Interpreted

Let's take a look at a query that uses the RR format for interpreting the years 15 and 75. (In the following query, you should assume the present year is 2007.)

As expected from rules 1 and 2, the years 15 and 75 are interpreted as 2015 and 1975. In the next query, you should assume the present year is 2075.

As expected from rules 3 and 4, the years 15 and 75 are interpreted as 2115 and 2055.

Using Datetime Functions

You use the datetime functions to get or process datetimes and timestamps (you'll learn about timestamps later in this chapter). Table 5-4 shows some of the datetime functions. In this table, x represents a datetime or a timestamp.

You'll learn more about the functions shown in Table 5-4 in the following sections.

ADD_MONTHS()

ADD_MONTHS (x, y) returns the result of adding y months to x. If y is negative, then y months are subtracted from x. The following example adds 13 months to January 1, 2007:

```
SELECT ADD_MONTHS('01-JAN-2007', 13)
FROM dual;

ADD_MONTH
------
01-FEB-08
```

The next example subtracts 13 months from the January 1, 2008; notice that –13 months are "added" to this date using ADD MONTHS():

```
SELECT ADD_MONTHS('01-JAN-2008', -13)
FROM dual;

ADD_MONTH
-----
01-DEC-06
```

Function	Description
ADD_MONTHS(x, y)	Returns the result of adding y months to x . If y is negative, y months are subtracted from x .
LAST_DAY(x)	Returns the last day of the month part of x.
MONTHS_BETWEEN(x, y)	Returns the number of months between x and y . If x appears before y on the calendar, the number returned is positive; otherwise the number is negative.
NEXT_DAY(x, day)	Returns the datetime of the next day following x ; day is specified as a literal string (SATURDAY, for example).
<pre>ROUND(x [, unit])</pre>	Rounds x . By default, x is rounded to the beginning of the nearest day. You may supply an optional $unit$ string that indicates the rounding unit; for example, YYYY rounds x to the first day of the nearest year.
SYSDATE	Returns the current datetime set in the database server's operating system.
TRUNC(x [, unit])	Truncates x . By default, x is truncated to the beginning of the day. You may supply an optional $unit$ string that indicates the truncating unit; for example, MM truncates x to the first day of the month.

TABLE 5-4 Datetime Functions

You can provide a time and date to the ADD MONTHS () function. For example, the following guery adds two months to the datetime 7:15:26 P.M. on January 1, 2007:

```
SELECT ADD MONTHS (TO DATE ('01-JAN-2007 19:15:26',
   'DD-MON-YYYY HH24:MI:SS'), 2)
  FROM dual;
  ADD MONTH
  _____
  01-MAR-07
```

The next query rewrites the previous example to convert the returned datetime from ADD MONTHS() to a string using TO CHAR() with the format DD-MON-YYYY HH24:MI:SS:

```
SELECT TO CHAR (ADD MONTHS (TO DATE ('01-JAN-2007 19:15:26',
 'DD-MON-YYYY HH24:MI:SS'), 2), 'DD-MON-YYYY HH24:MI:SS')
FROM dual;
TO CHAR (ADD MONTHS (T
_____
01-MAR-2007 19:15:26
```



NOTE

You can provide a date and time to any of the functions shown earlier in Table 5-4.

LAST DAY()

LAST_DAY (x) returns the date of the last day of the month part of x. The following example displays the last date in January 2008:

```
SELECT LAST_DAY('01-JAN-2008')
FROM dual;

LAST_DAY(
------
31-JAN-08
```

MONTHS_BETWEEN()

MONTHS_BETWEEN (x, y) returns the number of months between x and y. If x occurs before y in the calendar, then the number returned by MONTHS_BETWEEN() is negative.



NOTE

The ordering of the dates in your call to the MONTHS_BETWEEN() function is important: the later date must appear first if you want the result as a positive number.

The following example displays the number of months between May 25, 2008, and January 15, 2008. Notice that since the later date (May 25, 2008) appears first, the result returned is a positive number:

SELECT MONTHS_BETWEEN('25-MAY-2008', '15-JAN-2008')
FROM dual;

The next example reverses the same dates in the call to the MONTHS_BETWEEN() function, and therefore the returned result is a negative number of months:

SELECT MONTHS_BETWEEN('15-JAN-2008', '25-MAY-2008')
FROM dual;

NEXT_DAY()

NEXT_DAY (x, day) returns the date of the next day following x; you specify day as a literal string (SATURDAY, for example).

The following example displays the date of the next Saturday after January 1, 2008:

```
SELECT NEXT_DAY('01-JAN-2008', 'SATURDAY')
FROM dual;
```

```
NEXT_DAY(
-----
05-JAN-08
```

ROUND()

ROUND (x [, unit]) rounds x, by default, to the beginning of the nearest day. If you supply an optional unit string, x is rounded to that unit; for example, YYYY rounds x to the first day of the nearest year. You can use many of the parameters shown earlier in Table 5-2 to round a datetime.

The following example uses ROUND () to round October 25, 2008, up to the first day in the nearest year, which is January 1, 2009. Notice that the date is specified as 25-OCT-2008 and is contained within a call to the TO DATE () function:

The next example rounds May 25, 2008, to the first day in the nearest month, which is June 1, 2008, because May 25 is closer to the beginning of June than it is to the beginning of May:

The next example rounds 7:45:26 P.M. on May 25, 2008, to the nearest hour, which is 8:00 P.M.:

SYSDATE

SYSDATE returns the current datetime set in the database server's operating system. The following example gets the current date:

```
SELECT SYSDATE
FROM dual;
SYSDATE
```

05-NOV-07

TRUNC()

TRUNC (x [, unit]) truncates x. By default, x is truncated to the beginning of the day. If you supply an optional unit string, x is truncated to that unit; for example, MM truncates x to the first day in the month. You can use many of the parameters shown earlier in Table 5-2 to truncate a datetime.

The following example uses TRUNC () to truncate May 25, 2008, to the first day in the year, which is January 1, 2008:

```
SELECT TRUNC (TO_DATE ('25-MAY-2008'), 'YYYY')
FROM dual;

TRUNC (TO_______
01-JAN-08
```

The next example truncates May 25, 2008, to the first day in the month, which is May 1, 2008:

```
SELECT TRUNC(TO_DATE('25-MAY-2008'), 'MM')
FROM dual;

TRUNC(TO_
-----
01-MAY-08
```

The next example truncates 7:45:26 P.M. on May 25, 2008, to the hour, which is 7:00 P.M.:

Using Time Zones

Oracle Database 9*i* introduced the ability to use different time zones. A time zone is an offset from the time in Greenwich, England. The time in Greenwich was once known as Greenwich Mean Time (GMT), but is now known as Coordinated Universal Time (UTC, which comes from the initials of the French wording).

You specify a time zone using either an offset from UTC or a geographic region (e.g., PST). When you specify an offset, you use HH:MI prefixed with a plus or minus sign:

```
+|-HH:MI
```

where

- + or indicates an increase or decrease for the offset from UTC.
- HH:MI specifies the offset in hours and minutes for the time zone.



NOTE

The time zone hour and minute use the format parameters TZH and TZR, shown earlier in Table 5-2.

The following example shows offsets of 8 hours behind UTC and 2 hours 15 minutes ahead of UTC:



You may also specify a time zone using the geographical region. For example, PST indicates Pacific Standard Time, which is 8 hours behind UTC. EST indicates Eastern Standard Time, which is 5 hours behind UTC.



The time zone region uses the format parameter TZR, shown earlier in Table 5-2.

Time Zone Functions

There are a number of functions that are related to time zones; these functions are shown in Table 5-5.

You'll learn more about these functions in the following sections.

The Database Time Zone and Session Time Zone

If you're working for a large worldwide organization, the database you access may be located in a different time zone than your local time zone. The time zone for the database is known as the database time zone, and the time zone set for your database session is known as the session time zone. You'll learn about the database and session time zones in the following sections.

Function	Description
CURRENT_DATE	Returns the current date in the local time zone set for the database session
DBTIMEZONE	Returns the time zone for the database
<pre>NEW_TIME(x, time_zone1, time_zone2)</pre>	Converts x from time_zone1 to time_zone2 and returns the new datetime
SESSIONTIMEZONE	Returns the time zone for the database session
TZ_OFFSET(time_zone)	Returns the offset for time_zone in hours and minutes

TABLE 5-5 Time Zone Functions

The Database Time Zone

The database time zone is controlled using the TIME_ZONE database parameter. Your database administrator can change the setting of the TIME_ZONE parameter in the database's init.ora or spfile.ora file, or by using ALTER DATABASE SET TIME_ZONE = offset | region (e. g., ALTER DATABASE SET TIME_ZONE = '-8:00' or ALTER DATABASE SET TIME_ZONE = 'PST').

You can get the database time zone using the DBTIMEZONE function. For example, the following query gets the time zone for my database:

SELECT DBTIMEZONE FROM dual:

DBTIME -----+00:00

As you can see, +00:00 is returned. This means my database uses the time zone set in the operating system, which is set to PST on my computer.



NOTE

The Windows operating system is typically set up to adjust the clock for daylight savings. For California, this means that in the summer the clock is only 7 hours behind UTC, rather than 8 hours. When I wrote this chapter, I set the date to November 5, 2007, which means my clock is 8 hours behind UTC (I'm located in California).

The Session Time Zone

The session time zone is the time zone for a particular session. By default, the session time zone is the same as the operating system time zone. You can change your session time zone using the ALTER SESSION statement to set the session TIME_ZONE parameter (e.g., ALTER SESSION SET TIME_ZONE = 'PST' sets the local time zone to Pacific Standard Time). You can also set the session TIME_ZONE to LOCAL, which sets the time zone to the one used by the operating system of the computer on which the ALTER SESSION statement was run. You can also set the session TIME_ZONE to DBTIMEZONE, which sets the time zone to the one used by the database.

You can get the session time zone using the SESSIONTIMEZONE function. For example, the following query gets the time zone for my session:

SELECT SESSIONTIMEZONE FROM dual;

As you can see, my session time zone is 8 hours behind UTC.

Getting the Current Date in the Session Time Zone

The SYSDATE function gets the date from the database. This gives you the date in the database time zone. You can get the date in your session time zone using the CURRENT_DATE function. For example:

SELECT CURRENT_DATE FROM dual;

```
CURRENT_D ------ 05-NOV-07
```

Obtaining Time Zone Offsets

You can get the time zone offset hours using the TZ_OFFSET() function, passing the time zone region name to TZ_OFFSET(). For example, the following query uses TZ_OFFSET() to get the time zone offset hours for PST, which is 8 hours behind UTC:

SELECT TZ_OFFSET('PST') FROM dual;

```
TZ_OFFS
------
```



NOTE

In the summer, this will be –7:00 when using Windows, which sets the clock automatically to adjust for daylight savings.

Obtaining Time Zone Names

You can obtain all the time zone names by selecting all the rows from v\$timezone_names. To query v\$timezone_names, you should first connect to the database as the system user. The following query shows the first five rows from v\$timezone_names:

SELECT *

FROM v\$timezone_names
WHERE ROWNUM <= 5
ORDER BY tzabbrev;

TZNAME	TZABBREV
Africa/Algiers	CET
Africa/Algiers	LMT
Africa/Algiers	PMT
Africa/Algiers	WEST
Africa/Algiers	WET

You may use any of the tzname or tzabbrev values for your time zone setting.



NOTE

The ROWNUM pseudo column contains the row number. For example, the first row returned by a query has a row number of 1, the second has a row number of 2, and so on. Therefore, the WHERE clause in the previous query causes the query to return the first five rows.

Converting a Datetime from One Time Zone to Another

You use the NEW_TIME() function to convert a datetime from one time zone to another. For example, the following query uses NEW_TIME() to convert 7:45 P.M. on May 13, 2008, from PST to EST:

EST is 3 hours ahead of PST: therefore, 3 hours are added to 7:45 P.M. to give 10:45 P.M. (or 22:45 in 24-hour format).

Using Timestamps

Oracle Database 9*i* introduced the ability to store timestamps. A timestamp stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second. The advantages of a timestamp over a DATE are

- A timestamp can store a fractional second.
- A timestamp can store a time zone.

Let's examine the timestamp types.

Using the Timestamp Types

There are three timestamp types, which are shown in Table 5-6.

Description

TIMESTAMP[(seconds_precision)]

Type

Stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second. You can specify an optional precision for the seconds using <code>seconds_precision</code>, which can be an integer from 0 to 9; the default is 9, which means you can store up to 9 digits to the right of the decimal point for your second. If you try to add a row with more digits in the fractional second than the <code>TIMESTAMP</code> can store, then your fractional amount is rounded.

Extends TIMESTAMP to store a time zone.

(seconds_precision)
] WITH TIME ZONE
TIMESTAMP[
 (seconds_precision)
] WITH LOCAL TIME ZONE

TIMESTAMP[

Extends TIMESTAMP to convert a supplied datetime to the local time zone set for the database. The process of conversion is known as *normalizing* the datetime.

You'll learn how to use these timestamp types in the following sections.

Using the TIMESTAMP Type

As with the other types, you can use the TIMESTAMP type to define a column in a table. The following statement creates a table named purchases with timestamp that stores customer purchases. This table contains a TIMESTAMP column named made on to record when a purchase was made; notice a precision of 4 is set for the TIMESTAMP (this means up to four digits may be stored to the right of the decimal point for the second):

```
CREATE TABLE purchases with timestamp (
      product id INTEGER REFERENCES products (product id),
      customer id INTEGER REFERENCES customers (customer id),
      made on TIMESTAMP(4)
    );
```

NOTE

The purchases with timestamp table is created and populated with rows by the store schema.sql script. You'll see other tables in the rest of this chapter that are also created by the script, so you don't need to type in the CREATE TABLE or any of the INSERT statements shown in this chapter.

To supply a TIMESTAMP literal value to the database, you use the TIMESTAMP keyword along with a datetime in the following format:

```
TIMESTAMP 'YYYY-MM-DD HH24:MI:SS.SSSSSSSS'
```

Notice there are nine S characters after the decimal point, which means you can supply up to nine digits for the fractional second in your literal string. The number of digits you can actually store in a TIMESTAMP column depends on how many digits were set for storage of fractional seconds when the column was defined. For example, up to four digits can be stored in the made on column of the purchases with timestamp table. If you try to add a row with more than four fractional second digits, your fractional amount is rounded. For example,

```
2005-05-13 07:15:31.123456789
```

would be rounded to

```
2005-05-13 07:15:31.1235
```

The following INSERT statement adds a row to the purchases with timestamp table; notice the use of the TIMESTAMP keyword to supply a datetime literal:

```
INSERT INTO purchases with timestamp (
      product id, customer id, made on
     ) VALUES (
      1, 1, TIMESTAMP '2005-05-13 07:15:31.1234'
```

The following query retrieves the row:

```
SELECT *
```

```
FROM purchases_with_timestamp;
```

Using the TIMESTAMP WITH TIME ZONE Type

The TIMESTAMP WITH TIME ZONE type extends TIMESTAMP to store a time zone. The following statement creates a table named purchases_timestamp_with_tz that stores customer purchases; this table contains a TIMESTAMP WITH TIME ZONE column named made on to record when a purchase was made:

```
CREATE TABLE purchases_timestamp_with_tz (
    product_id INTEGER REFERENCES products(product_id),
    customer_id INTEGER REFERENCES customers(customer_id),
    made_on TIMESTAMP(4) WITH TIME ZONE
);
```

To supply a timestamp literal with a time zone to the database, you simply add the time zone to your TIMESTAMP. For example, the following TIMESTAMP includes a time zone offset of -07:00:

```
TIMESTAMP '2005-05-13 07:15:31.1234 -07:00'
```

You may also supply a time zone region, as shown in the following example that specifies PST as the time zone:

```
TIMESTAMP '2005-05-13 07:15:31.1234 PST'
```

The following example adds two rows to the purchases_timestamp_with_tz table:

```
INSERT INTO purchases_timestamp_with_tz (
    product_id, customer_id, made_on
) VALUES (
    1, 1, TIMESTAMP '2005-05-13 07:15:31.1234 -07:00'
);

INSERT INTO purchases_timestamp_with_tz (
    product_id, customer_id, made_on
) VALUES (
    1, 2, TIMESTAMP '2005-05-13 07:15:31.1234 PST'
);
```

The following query retrieves the rows:

```
SELECT *
```

```
FROM purchases_timestamp_with_tz;
```

```
PRODUCT ID CUSTOMER ID MADE ON
                    1 13-MAY-05 07.15.31.1234 AM -07:00
                    2 13-MAY-05 07.15.31.1234 AM PST
```

Using the TIMESTAMP WITH LOCAL TIME ZONE Type

The TIMESTAMP WITH LOCAL TIME ZONE type extends TIMESTAMP to store a timestamp with the local time zone set for your database. When you supply a timestamp for storage in a TIMESTAMP WITH LOCAL TIME ZONE column, your timestamp is converted—or normalized—to the time zone set for the database. When you retrieve the timestamp, it is normalized to the time zone set for your session.



TIP

You should use TIMESTAMP WITH LOCAL TIME ZONE to store timestamps when your organization has a global system accessed throughout the world. This is because TIMESTAMP WITH LOCAL TIME ZONE stores a timestamp using the local time where the database is located, but users see the timestamp normalized to their own time zone.

My database time zone is PST (PST is 8 hours behind UTC), and I want to store the following timestamp in my database:

```
2005-05-13 07:15:30 EST
```

EST is 5 hours behind UTC, and the difference between PST and EST is 3 hours (8 - 5 = 3). Therefore, to normalize the previous timestamp for PST, 3 hours must be subtracted from it to give the following normalized timestamp:

```
2005-05-13 04:15:30
```

This is the timestamp that would be stored in a TIMESTAMP WITH LOCAL TIME ZONE column in my database.

The following statement creates a table named purchases with local tz that stores customer purchases; this table contains a TIMESTAMP WITH LOCAL TIME ZONE column named made on to record when a purchase was made:

```
CREATE TABLE purchases with local tz (
      product id INTEGER REFERENCES products(product_id),
      customer id INTEGER REFERENCES customers(customer_id),
      made on TIMESTAMP(4) WITH LOCAL TIME ZONE
    );
```

The following INSERT adds a row to the purchases with local tz table with the made on column set to 2005-05-13 07:15:30 EST:

```
INSERT INTO purchases with local tz (
     product id, customer id, made on
   ) VALUES (
     1, 1, TIMESTAMP '2005-05-13 07:15:30 EST'
   );
```

Although the timestamp for the made_on column is set to 2005-05-13 07:15:30 EST, the actual timestamp stored in my database is 2005-05-13 04:15:30 (the timestamp normalized for PST).

The following query retrieves the row:

SELECT *

```
FROM purchases_with_local_tz;
```

```
PRODUCT_ID CUSTOMER_ID MADE_ON

1 1 13-MAY-05 04.15.30.0000 AM
```

Because my database time zone and session time zone are both PST, the timestamp returned by the query is for PST.



CAUTION

The timestamp returned by the previous query is normalized for PST. If your database time zone or session time zone are not PST, the timestamp returned when you run the query will be different (it will be normalized for your time zone).

If I set the local time zone for my session to EST and repeat the previous query, I get the timestamp normalized for EST:

```
ALTER SESSION SET TIME_ZONE = 'EST';
```

Session altered.

SELECT *

```
FROM purchases_with_local_tz;
```

As you can see, the timestamp returned by the query is 13-MAY-05 07.15.30.0000 AM, which is the timestamp normalized for the session time zone of EST. Because EST is three hours ahead of PST, three hours must be added to 13-MAY-05 04:15:30 (the timestamp stored in the database) to give 13-MAY-05 07.15.30 AM (the timestamp returned by the query).

The following statement sets my session time zone back to PST:

```
ALTER SESSION SET TIME ZONE = 'PST';
```

Session altered.

Timestamp Functions

There are a number of functions that allow you to get and process timestamps. These functions are shown in Table 5-7.

You'll learn more about the functions shown in Table 5-7 in the following sections.

Function	Description
CURRENT_TIMESTAMP	Returns a TIMESTAMP WITH TIME ZONE containing the current date and time for the session, plus the session time zone.
EXTRACT ({ YEAR MONTH DAY HOUR MINUTE SECOND } { TIMEZONE_HOUR TIMEZONE_MINUTE } { TIMEZONE_REGION } TIMEZONE_ABBR } FROM x)	Extracts and returns the year, month, day, hour, minute, second, or time zone from x ; x may be a timestamp or a DATE.
FROM_TZ(x, time_zone)	Converts the TIMESTAMP x to the time zone specified by time_zone and returns a TIMESTAMP WITH TIMEZONE; time_zone must be specified as a string of the form + - HH:MI. The function basically merges x and time_zone into one value.
LOCALTIMESTAMP	Returns a TIMESTAMP containing the current date and time for the session.
SYSTIMESTAMP	Returns a TIMESTAMP WITH TIME ZONE containing the current date and time for the database, plus the database time zone.
SYS_EXTRACT_UTC(x)	Converts the TIMESTAMP WITH TIMEZONE x to a TIMESTAMP containing the date and time in UTC.
TO_TIMESTAMP(x, [format])	Converts the string x to a TIMESTAMP. You may also specify an optional format for x.
TO_TIMESTAMP_TZ(x, [format])	Converts the string x to a TIMESTAMP WITH TIMEZONE. You may also specify an optional format for x .

 TABLE 5-7
 Timestamp Functions

CURRENT_TIMESTAMP, LOCALTIMESTAMP, and SYSTIMESTAMP

The following query calls the CURRENT TIMESTAMP, LOCALTIMESTAMP, and SYSTIMESTAMP functions (my session time zone and database time zone are both PST, which is 8 hours behind UTC):

SELECT CURRENT TIMESTAMP, LOCALTIMESTAMP, SYSTIMESTAMP FROM dual;

```
CURRENT TIMESTAMP
LOCALTIMESTAMP
_____
SYSTIMESTAMP
05-NOV-07 12.15.32.734000 PM PST
05-NOV-07 12.15.32.734000 PM
05-NOV-07 12.15.32.734000 PM -08:00
```

If I change my session <code>TIME_ZONE</code> to EST and repeat the previous query, I get the following results:

```
ALTER SESSION SET TIME ZONE = 'EST';
```

Session altered.

SELECT CURRENT_TIMESTAMP, LOCALTIMESTAMP, SYSTIMESTAMP FROM dual;

The following statement sets my session time zone back to PST:

```
ALTER SESSION SET TIME ZONE = 'PST';
```

Session altered.

EXTRACT()

EXTRACT() extracts and returns the year, month, day, hour, minute, second, or time zone from x; x may be a timestamp or a DATE. The following query uses EXTRACT() to get the year, month, and day from a DATE returned by TO_DATE():

SELECT

```
EXTRACT (YEAR FROM TO_DATE ('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS YEAR,

EXTRACT (MONTH FROM TO_DATE ('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS MONTH,

EXTRACT (DAY FROM TO_DATE ('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS DAY

FROM dual;

YEAR MONTH DAY
```

The next query uses EXTRACT() to get the hour, minute, and second from a TIMESTAMP returned by TO TIMESTAMP():

SELECT

```
EXTRACT (HOUR FROM TO_TIMESTAMP('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS HOUR,

EXTRACT (MINUTE FROM TO_TIMESTAMP('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS MINUTE,

EXTRACT (SECOND FROM TO_TIMESTAMP('01-JAN-2008 19:15:26',
    'DD-MON-YYYY HH24:MI:SS')) AS SECOND
```

FROM dual;

SECOND	MINUTE	HOUR
26	15	19

The following query uses EXTRACT () to get the time zone hour, minute, second, region, and region abbreviation from a TIMESTAMP WITH TIMEZONE returned by TO TIMESTAMP TZ():

SELECT

```
EXTRACT (TIMEZONE HOUR FROM TO TIMESTAMP TZ (
   '01-JAN-2008 19:15:26 -7:15', 'DD-MON-YYYY HH24:MI:SS TZH:TZM'))
   AS TZH,
 EXTRACT (TIMEZONE MINUTE FROM TO TIMESTAMP TZ (
   '01-JAN-2008 19:15:26 -7:15', 'DD-MON-YYYY HH24:MI:SS TZH:TZM'))
 EXTRACT (TIMEZONE REGION FROM TO TIMESTAMP TZ (
   '01-JAN-2008 19:15:26 PST', 'DD-MON-YYYY HH24:MI:SS TZR'))
 EXTRACT (TIMEZONE ABBR FROM TO TIMESTAMP TZ (
    '01-JAN-2008 19:15:26 PST', 'DD-MON-YYYY HH24:MI:SS TZR'))
   AS TZA
FROM dual;
           ._____
       -7 -15 PST
                                       PST
```

FROM TZ()

FROM TZ (x, time zone) converts the TIMESTAMP x to the time zone specified by time zone and returns a TIMESTAMP WITH TIMEZONE; time zone must be specified as a string of the form + | - HH:MI. The function basically merges x and time zone into one value.

For example, the following query merges the timestamp 2008-05-13 07:15:31.1234 and the time zone offset of -7:00 from UTC:

```
SELECT FROM TZ(TIMESTAMP '2008-05-13 07:15:31.1234', '-7:00')
    FROM dual:
```

```
FROM TZ (TIMESTAMP'2008-05-1307:15:31.1234','-7:00')
______
13-MAY-08 07.15.31.123400000 AM -07:00
```

SYS EXTRACT UTC()

SYS EXTRACT UTC(x) converts the TIMESTAMP WITH TIMEZONE x to a TIMESTAMP containing the date and time in UTC.

The following query converts 2008–11–17 19:15:26 PST to UTC:

SELECT SYS EXTRACT UTC(TIMESTAMP '2008-11-17 19:15:26 PST') FROM dual;

```
SYS EXTRACT UTC (TIMESTAMP'2008-11-1719:15:26PST')
_____
18-NOV-08 03.15.26.000000000 AM
```

Because PST is 8 hours behind UTC in the winter, the query returns a TIMESTAMP 8 hours ahead of 2008-11-17 19:15:26 PST, which is 18-NOV-08 03.15.26 AM.

For a date in the summer, the returned TIMESTAMP is only 7 hours ahead of UTC:

SELECT SYS_EXTRACT_UTC(TIMESTAMP '2008-05-17 19:15:26 PST') FROM dual;

TO TIMESTAMP()

TO_TIMESTAMP(x, format) converts the string x (which may be a CHAR, VARCHAR2, NCHAR, or NVARCHAR2) to a TIMESTAMP. You may also specify an optional format for x.

The following query converts the string 2005-05-13 07:15:31.1234 with the format YYYY-MM-DD HH24:MI:SS.FF to a TIMESTAMP:

SELECT TO_TIMESTAMP('2008-05-13 07:15:31.1234', 'YYYY-MM-DD HH24:MI:SS.FF') FROM dual;

TO TIMESTAMP TZ()

TO_TIMESTAMP_TZ(x, [format]) converts x to a TIMESTAMP WITH TIMEZONE. You may specify an optional format for x.

The following query passes the PST time zone (identified using TZR in the format string) to TO TIMESTAMP TZ():

The next query uses a time zone offset of -7:00 from UTC (-7:00 is identified using TZR and TZM in the format string):

SELECT TO_TIMESTAMP_TZ('2008-05-13 07:15:31.1234 -7:00', 'YYYY-MM-DD HH24:MI:SS.FF TZH:TZM') FROM dual;

```
TO_TIMESTAMP_TZ('2008-05-1307:15:31.1234-7:00','YYYY-MM-DDHH24:MI:SS.FFTZH
```

Converting a String to a TIMESTAMP WITH LOCAL TIME ZONE

You can use the CAST() function to convert a string to a TIMESTAMP WITH LOCAL TIME ZONE. You were introduced to CAST () in the previous chapter. As a reminder, CAST (x AS type) converts x to a compatible database type specified by type.

The following query uses CAST () to convert the string 13-JUN-08 to a TIMESTAMP WITH LOCAL TIME ZONE:

SELECT CAST('13-JUN-08' AS TIMESTAMP WITH LOCAL TIME ZONE) FROM dual;

```
CAST ('13-JUN-08'ASTIMESTAMPWITHLOCALTIMEZONE)
13-JUN-08 12.00.00.000000 AM
```

The timestamp returned by this query contains the date June 13, 2008 and the time of 12 A.M. The next query uses CAST() to convert a more complex string to a TIMESTAMP WITH LOCAL TIME ZONE:

SELECT CAST(TO TIMESTAMP TZ('2008-05-13 07:15:31.1234 PST', 'YYYY-MM-DD HH24:MI:SS.FF TZR') AS TIMESTAMP WITH LOCAL TIME ZONE) FROM dual:

```
CAST (TO TIMESTAMP TZ ('2008-05-1307:15:31.1234PST', 'YYYY-MM-DDHH24:MI:SS.FF
13-MAY-08 06.15.31.123400 AM
```

The timestamp returned by this query contains the date May 13, 2008 and the time of 6:15:31.1234 AM PST (PST is the time zone for both my database and session).

The following query does the same thing as the previous one, except this time EST is the time zone:

SELECT CAST(TO TIMESTAMP TZ('2008-05-13 07:15:31.1234 EST', 'YYYY-MM-DD HH24:MI:SS.FF TZR') AS TIMESTAMP WITH LOCAL TIME ZONE) FROM dual:

```
CAST (TO TIMESTAMP TZ ('2008-05-1307:15:31.1234EST', 'YYYY-MM-DDHH24:MI:SS.FF
13-MAY-08 04.15.31.123400 AM
```

The timestamp returned by this query contains the date May 13, 2008 and the time of 4:15:31.1234 AM PST (because PST is 3 hours behind EST, the time returned in the timestamp is 3 hours earlier than the time in the actual query).

Using Time Intervals

Oracle Database 9i introduced data types that allow you to store time intervals. Examples of time intervals include

- 1 year 3 months
- 25 months

- -3 days 5 hours 16 minutes
- 1 day 7 hours
- -56 hours



NOTE

Do not confuse time intervals with datetimes or timestamps. A time interval records a length of time (e.g., 1 year 3 months), whereas a datetime or timestamp records a specific date and time (e.g., 7:32:16 P.M. on October 28, 2006).

In our imaginary online store, you might want to offer limited time discounts on products. For example, you could give customers a coupon that is valid for a few months, or you could run a special promotion for a few days. You'll see examples that feature coupons and promotions later in this section.

Table 5-8 shows the interval types.

You'll learn how to use the time interval types in the following sections.

Using the INTERVAL YEAR TO MONTH Type

INTERVAL YEAR TO MONTH stores a time interval measured in years and months. The following statement creates a table named coupons that stores coupon information. The coupons table contains an INTERVAL YEAR TO MONTH column named duration to record the interval of

Type

INTERVAL YEAR[(years_precision)] TO MONTH

```
INTERVAL DAY[
  (days_precision)
]
TO SECOND[
  (seconds_precision)
]
```

Description

Stores a time interval measured in years and months. You can specify an optional precision for the years using <code>years_precision</code>, which may be an integer from 0 to 9. The default precision is 2, which means you can store two digits for the years in your interval. If you try to add a row with more year digits than your <code>INTERVAL YEAR TO MONTH</code> column can store, you'll get an error. You can store a positive or negative time interval.

Stores a time interval measured in days and seconds. You can specify an optional precision for the days using <code>days_precision</code> (an integer from 0 to 9; the default is 2). In addition, you can also specify an optional precision for the fractional seconds using <code>seconds_precision</code> (an integer from 0 to 9; the default is 6). You can store a positive or negative time interval.

time for which the coupon is valid; notice that I've provided a precision of 3 for the duration column, which means that up to three digits may be stored for the year:

```
CREATE TABLE coupons (
      coupon id INTEGER CONSTRAINT coupons pk PRIMARY KEY,
      name VARCHAR2(30) NOT NULL,
      duration INTERVAL YEAR(3) TO MONTH
     );
```

To supply an INTERVAL YEAR TO MONTH literal value to the database, you use the following simplified syntax:

```
INTERVAL '[+|-][y][-m]' [YEAR[(years precision)])] [TO MONTH]
```

where

- + or is an optional indicator that specifies whether the time interval is positive or negative (the default is positive).
- y is the optional number of years for the interval.
- m is the optional number of months for the interval. If you supply years and months, you must include TO MONTH in your literal.
- years precision is the optional precision for the years (the default is 2).

The following table shows some examples of year-to-month interval literals.

Literal

INTERVAL '1' YEAR INTERVAL '11' MONTH INTERVAL '14' MONTH INTERVAL '1-3' YEAR TO MONTH INTERVAL '0-5' YEAR TO MONTH INTERVAL '123' YEAR(3) TO MONTH INTERVAL '-1-5' YEAR TO MONTH INTERVAL '1234' YEAR(3)

Description

Interval of 1 year.

Interval of 11 months.

Interval of 14 months (equivalent to 1 year 2 months).

Interval of 1 year 3 months.

Interval of 0 years 5 months.

Interval of 123 years with a precision of 3 digits.

A negative interval of 1 year 5 months.

Invalid interval: 1234 contains four digits and therefore contains one too many digits allowed by the precision of 3 (three digits maximum).

The following INSERT statements add rows to the coupons table with the duration column set to some of the intervals shown in the previous table:

```
INSERT INTO coupons (coupon id, name, duration)
    VALUES (1, '$1 off Z Files', INTERVAL '1' YEAR);
    INSERT INTO coupons (coupon id, name, duration)
    VALUES (2, '$2 off Pop 3', INTERVAL '11' MONTH);
```

```
INSERT INTO coupons (coupon id, name, duration)
VALUES (3, '$3 off Modern Science', INTERVAL '14' MONTH);
INSERT INTO coupons (coupon id, name, duration)
VALUES (4, '$2 off Tank War', INTERVAL '1-3' YEAR TO MONTH);
INSERT INTO coupons (coupon id, name, duration)
VALUES (5, '$1 off Chemistry', INTERVAL '0-5' YEAR TO MONTH);
INSERT INTO coupons (coupon id, name, duration)
VALUES (6, '$2 off Creative Yell', INTERVAL '123' YEAR(3));
```

If you try to add a row with the duration column set to the invalid interval of INTERVAL '1234' YEAR(3), you'll get an error because the precision of the duration column is 3 and is therefore too small to accommodate the number 1234. The following INSERT shows the error:

```
SQL> INSERT INTO coupons (coupon id, name, duration)
      2 VALUES (7, '$1 off Z Files', INTERVAL '1234' YEAR(3));
    VALUES (7, '$1 off Z Files', INTERVAL '1234' YEAR(3))
    ERROR at line 2:
    ORA-01873: the leading precision of the interval is too small
```

The following query retrieves the rows from the coupons table; notice the formatting of the duration values:

SELECT * FROM coupons;

COUPON_ID	NAMI	€		DURATION
1	\$1	off	Z Files	+001-00
2	\$2	off	Pop 3	+000-11
3	\$3	off	Modern Science	+001-02
4	\$2	off	Tank War	+001-03
5	\$1	off	Chemistry	+000-05
6	\$2	off	Creative Yell	+123-00

Using the INTERVAL DAY TO SECOND Type

INTERVAL DAY TO SECOND stores time intervals measured in days and seconds. The following statement creates a table named promotions that stores promotion information. The promotions table contains an INTERVAL DAY TO SECOND column named duration to record the interval of time for which the promotion is valid:

```
CREATE TABLE promotions (
      promotion id INTEGER CONSTRAINT promotions pk PRIMARY KEY,
      name VARCHAR2(30) NOT NULL,
      duration INTERVAL DAY(3) TO SECOND (4)
    );
```

Notice I've provided a precision of 3 for the day and a precision of 4 for the fractional seconds of the duration column. This means that up to three digits may be stored for the day of the interval and up to four digits to the right of the decimal point for the fractional seconds.

To supply an INTERVAL DAY TO SECOND literal value to the database, you use the following simplified syntax:

```
INTERVAL '[+|-][d] [h[:m[:s]]]' [DAY[(days precision)]])
     [TO HOUR | MINUTE | SECOND[(seconds precision)]]
```

where

- + or is an optional indicator that specifies whether the time interval is positive or negative (the default is positive).
- d is the number of days for the interval.
- h is the optional number of hours for the interval; if you supply days and hours, you must include TO HOUR in your literal.
- m is the optional number of minutes for the interval; if you supply days and minutes, you must include TO MINUTES in your literal.
- s is the optional number of seconds for the interval; if you supply days and seconds, you must include TO SECOND in your literal.
- days precision is the optional precision for the days (the default is 2).
- seconds precision is the optional precision for the fractional seconds (the default is 6).

The following table shows some examples of day-to-second interval literals.

Literal

INTERVAL '3' DAY

INTERVAL '2' HOUR INTERVAL '25' MINUTE INTERVAL '45' SECOND INTERVAL '3 2' DAY TO HOUR INTERVAL '3 2:25' DAY TO MINUTE INTERVAL '3 2:25:45' DAY TO SECOND INTERVAL '123 2:25:45.12' DAY(3) TO SECOND(2) INTERVAL '3 2:00:45' DAY TO SECOND INTERVAL '-3 2:25:45' DAY TO SECOND INTERVAL '1234 2:25:45' DAY(3) TO SECOND INTERVAL '123 2:25:45.123' DAY TO SECOND(2)

Description

Interval of 3 days. Interval of 2 hours. Interval of 25 minutes. Interval of 45 seconds. Interval of 3 days 2 hours.

Interval of 3 days 2 hours 25 minutes. Interval of 3 days 2 hours 25 minutes 45 seconds.

Interval of 123 days 2 hours 25 minutes 45.12 seconds; the precision for days is 3 digits, and the precision for the fractional seconds is 2 digits.

Interval of 3 days 2 hours 0 minutes 45 seconds.

Negative interval of 3 days 2 hours 25 minutes 45 seconds.

Invalid interval because the number of digits in the days exceeds the specified precision of 3.

Invalid interval because the number of digits in the fractional seconds exceeds the specified precision of 2.

The following INSERT statements add rows to the promotions table:

```
INSERT INTO promotions (promotion_id, name, duration)
VALUES (1, '10% off Z Files', INTERVAL '3' DAY);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (2, '20% off Pop 3', INTERVAL '2' HOUR);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (3, '30% off Modern Science', INTERVAL '25' MINUTE);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (4, '20% off Tank War', INTERVAL '45' SECOND);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (5, '10% off Chemistry', INTERVAL '3 2:25' DAY TO MINUTE);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (6, '20% off Creative Yell', INTERVAL '3 2:25:45' DAY TO SECOND);

INSERT INTO promotions (promotion_id, name, duration)
VALUES (7, '15% off My Front Line',
INTERVAL '123 2:25:45.12' DAY(3) TO SECOND(2));
```

The following query retrieves the rows from the promotions table; notice the formatting of the duration values:

SELECT * FROM promotions;

PROMOTION_ID	NAME		DURA'	TION
1	10% c	off Z Files	+003	00:00:00.0000
2	20% c	off Pop 3	+000	02:00:00.0000
3	30% c	off Modern Science	+000	00:25:00.0000
4	20% c	off Tank War	+000	00:00:45.0000
5	10% c	off Chemistry	+003	02:25:00.0000
6	20% c	off Creative Yell	+003	02:25:45.0000
7	15% c	off My Front Line	+123	02:25:45.1200

Time Interval Functions

There are a number of functions that allow you to get and process time intervals; these functions are shown in Table 5-9.

You'll learn more about the functions shown in Table 5-9 in the following sections.

NUMTODSINTERVAL()

NUMTODSINTERVAL(x, $interval_unit$) converts the number x to an INTERVAL DAY TO SECOND. The interval for x is specified in $interval_unit$, which may be DAY, HOUR, MINUTE, or SECOND.

For example, the following query converts several numbers to time intervals using NUMTODSINTERVAL():

Function	Description
<pre>NUMTODSINTERVAL(x, interval_unit)</pre>	Converts the number x to an INTERVAL DAY TO SECOND. The interval for x is specified in $interval_unit$, which may be DAY, HOUR, MINUTE, or SECOND.
<pre>NUMTOYMINTERVAL(x, interval_unit)</pre>	Converts the number x to an INTERVAL YEAR TO MONTH. The interval for x is specified in $interval_unit$, which may be YEAR or MONTH.
TO_DSINTERVAL(x)	Converts the string x to an INTERVAL DAY TO SECOND.
TO_YMINTERVAL(x)	Converts the string x to an INTERVAL YEAR TO MONTH.

TABLE 5-9 Time Interval Functions

```
SELECT
    NUMTODSINTERVAL(1.5, 'DAY'),
    NUMTODSINTERVAL (3.25, 'HOUR'),
    NUMTODSINTERVAL (5, 'MINUTE'),
    NUMTODSINTERVAL (10.123456789, 'SECOND')
   FROM dual;
   NUMTODSINTERVAL(1.5,'DAY')
   NUMTODSINTERVAL (3.25, 'HOUR')
   _____
   NUMTODSINTERVAL(5, 'MINUTE')
   _____
   NUMTODSINTERVAL (10.123456789, 'SECOND')
   _____
   +00000001 12:00:00.00000000
   +00000000 03:15:00.00000000
   +00000000 00:05:00.00000000
   +000000000 00:00:10.123456789
```

NUMTOYMINTERVAL()

NUMTOYMINTERVAL (x, interval unit) converts the number x to an INTERVAL YEAR TO MONTH. The interval for x is specified in *interval unit*, which may be YEAR or MONTH.

For example, the following query converts two numbers to time intervals using NUMTOYMINTERVAL():

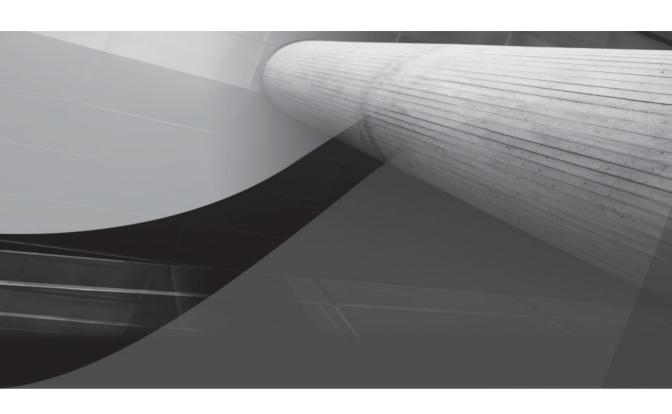
```
SELECT
    NUMTOYMINTERVAL (1.5, 'YEAR'),
    NUMTOYMINTERVAL (3.25, 'MONTH')
   FROM dual:
  NUMTOYMINTERVAL (1.5, 'YEAR')
   _____
  NUMTOYMINTERVAL (3.25, 'MONTH')
   _____
   +00000001-06
   +000000000-03
```

Summary

In this chapter, you learned the following:

- You may store a datetime using the DATE type. The DATE type stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second.
- You may use TO CHAR() and TO DATE() to convert between strings and dates and times.
- The Oracle database always stores all four digits of a year and interprets two-digit years using a set of rules. You should always specify all four digits of the year.
- There are a number of functions that process dates and times. An example is ADD_MONTHS (x, y), which returns the result of adding y months to x.
- Oracle Database 9*i* introduced the ability to use different time zones. A time zone is an offset from the time in Greenwich, England. The time in Greenwich was once known as Greenwich Mean Time (GMT), but it is now known as Coordinated Universal Time (UTC). You specify a time zone using either an offset from UTC or the name of the region (e.g., PST).
- Oracle Database 9*i* introduced the ability to store timestamps. A timestamp stores the century, all four digits of a year, the month, the day, the hour (in 24-hour format), the minute, and the second. The advantages of a timestamp over a DATE are a timestamp can store a fractional second and a time zone.
- Oracle Database 9*i* introduced the ability to handle time intervals, which allow you to store a length of time. An example time interval is 1 year 3 months.

In the next chapter, you'll learn how to nest one guery within another.



CHAPTER 6

Subqueries



Il the queries you've seen so far in this book have contained just one SELECT statement. In this chapter, you will learn the following:

- How to place an inner SELECT statement within an outer SELECT, UPDATE, or DELETE statement. The inner SELECT is known as a *subquery*.
- The features of the different types of subqueries.
- How subqueries allow you to build up very complex statements from simple components.

Types of Subqueries

There are two basic types of subqueries:

- **Single-row subqueries** return zero rows or one row to the outer SQL statement. There is a special case of a single-row subquery that contains exactly one column; this type of subquery is called a *scalar subquery*.
- Multiple-row subqueries return one or more rows to the outer SQL statement.

In addition, there are three subtypes of subqueries that may return single or multiple rows:

- Multiple-column subqueries return more than one column to the outer SQL statement.
- Correlated subqueries reference one or more columns in the outer SQL statement. These are called "correlated" subqueries because they are related to the outer SQL statement through the same columns.
- Nested subqueries are placed within another subquery. You can nest subqueries to a depth of 255.

You'll learn about each of these types of subqueries and how to add them to SELECT, UPDATE, and DELETE statements. Let's plunge in and look at how to write single-row subqueries.

Writing Single-Row Subqueries

A single-row subquery is one that returns either zero rows or one row to the outer SQL statement. As you'll see in this section, you may place a subquery in a WHERE clause, a HAVING clause, or a FROM clause of a SELECT statement. You'll also see some errors you might encounter when running subqueries.

Subqueries in a WHERE Clause

You may place a subquery in the WHERE clause of another query. The following query contains a subquery placed in its WHERE clause; notice that the subquery is placed within parentheses (...):

SELECT first_name, last_name
FROM customers
WHERE customer_id =
(SELECT customer_id
FROM customers

```
WHERE last name = 'Brown');
FIRST NAME LAST NAME
_____
John
        Brown
```

This example retrieves the first name and last name of the row from the customers table whose last name is Brown. Let's break this query down and analyze what's going on. The subquery in the WHERE clause is

```
SELECT customer id
     FROM customers
     WHERE last name = 'Brown';
```

This subguery is executed first (and only once) and returns the customer id for the row whose last name is Brown. The customer id for this row is 1, which is passed to the WHERE clause of the outer query. Therefore, the outer query returns the same result as the following query:

```
SELECT first name, last name
    FROM customers
    WHERE customer id = 1;
```

Using Other Single-Row Operators

The subquery example shown at the start of the previous section used the equality operator (=) in the WHERE clause. You may also use other comparison operators, such as <>, <, >, <=, and >=, with a single-row subquery. The following example uses > in the outer query's WHERE clause; the subquery uses the AVG () function to get the average price of the products, which is passed to the WHERE clause of the outer query. The entire query returns the product id, name, and price of products whose price is greater than that average price.

```
SELECT product id, name, price
    FROM products
    WHERE price >
      (SELECT AVG(price)
       FROM products);
```

PRODUCT_ID	NAME	PRICE
1	Modern Science	19.95
2	Chemistry	30
3	Supernova	25.99
5	Z Files	49.99

Let's break the example down to understand how it works. The following example shows the subquery run on its own:

SELECT AVG(price) FROM products;

```
AVG (PRICE)
_____
19.7308333
```



NOTE

This subquery is an example of a scalar subquery, because it returns exactly one row containing one column. The value returned by a scalar subquery is treated as a single scalar value.

The value 19.7308333 returned by the subquery is used in the WHERE clause of the outer query, which is therefore equivalent to the following:

```
SELECT product_id, name, price
FROM products
WHERE price > 19.7308333;
```

Subqueries in a HAVING Clause

As you saw in Chapter 4, you use the HAVING clause to filter groups of rows. You may place a subquery in the HAVING clause of an outer query. Doing this allows you to filter groups of rows based on the result returned by your subquery.

The following example uses a subquery in the HAVING clause of the outer query. The example retrieves the product_type_id and the average price for products whose average price is less than the maximum of the average for the groups of the same product type:

Notice the subquery uses AVG() to first compute the average price for each product type. The result returned by AVG() is then passed to MAX(), which returns the maximum of the averages.

Let's break the example down to understand how it works. The following example shows the output from the subquery when it is run on its own:

This value of 26.22 is used in the HAVING clause of the outer query to filter the group's rows to those having an average price less than 26.22. The following query shows a version of the outer query that retrieves the product type id and average price of the products grouped by product type id:

```
SELECT product type id, AVG(price)
    FROM products
    GROUP BY product type id
    ORDER BY product type id;
    PRODUCT TYPE ID AVG (PRICE)
    _____
              1
                  24.975
                  2
               3
                   13.24
                   13.99
                    13.49
```

The groups with a product type id of 1, 3, 4, and null have an average price less than 26.22. As expected, these are the same groups returned by the query at the start of this section.

Subqueries in a FROM Clause (Inline Views)

You may place a subquery in the FROM clause of an outer query. These types of subqueries are also known as inline views, because the subquery provides data in line with the FROM clause. The following simple example retrieves the products whose product id is less than 3:

```
SELECT product id
     FROM
       (SELECT product id
       FROM products
       WHERE product id < 3);
     PRODUCT ID
```

Notice that the subquery returns the rows from the products table whose product id is less than 3 to the outer query, which then retrieves and displays those product id values. As far as the FROM clause of the outer query is concerned, the output from the subquery is just another source of data.

The next example is more useful and retrieves the product id and price from the products table in the outer query, and the subquery retrieves the number of times a product has been purchased:

```
SELECT prds.product id, price, purchases data.product count
    FROM products prds,
      (SELECT product id, COUNT (product id) product count
```

```
FROM purchases
GROUP BY product_id) purchases_data
WHERE prds.product_id = purchases_data.product_id;
```

PRODUCT_ID	PRICE	PRODUCT_COUNT
1	19.95	4
2	30	4
3	25.99	1

Notice that the subquery retrieves the product_id and COUNT (product_id) from the purchases table and returns them to the outer query. As you can see, the output from subquery is just another source of data to the FROM clause of the outer query.

Errors You Might Encounter

In this section, you'll see some errors you might encounter. Specifically, you'll see that a single-row subquery may return a maximum of one row and that a subquery may not contain an ORDER BY clause.

Single-Row Subqueries May Return a Maximum of One Row

If your subquery returns more than one row, you'll get the following error:

```
ORA-01427: single-row subquery returns more than one row.
```

For example, the subquery in the following statement attempts to pass multiple rows to the equality operator (=) in the outer query:

```
SQL> SELECT product_id, name

2  FROM products

3  WHERE product_id =

4  (SELECT product_id)

5  FROM products

6  WHERE name LIKE '%e%');

(SELECT product_id

*

ERROR at line 4:

ORA-01427: single-row subquery returns more than one row
```

There are nine rows in the products table whose names contain the letter e, and the subquery attempts to pass these rows to the equality operator in the outer query. Because the equality operator can handle only a single row, the query is invalid and an error is returned.

You'll learn how to return multiple rows from a subquery later in the section "Writing Multiple-Row Subqueries."

Subqueries May Not Contain an ORDER BY Clause

A subquery may not contain an ORDER BY clause. Instead, any ordering must be done in the outer query. For example, the following outer query has an ORDER BY clause at the end that sorts the product id values in descending order:

```
SELECT product id, name, price
   FROM products
   WHERE price >
      (SELECT AVG(price)
      FROM products)
   ORDER BY product id DESC;
```

PRODUCT_ID	NAME	PRICE
5	Z Files	49.99
3	Supernova	25.99
2	Chemistry	30
1	Modern Science	19.95

Writing Multiple-Row Subqueries

You use a multiple-row subquery to return one or more rows to an outer SQL statement. To handle a subquery that returns multiple rows, your outer query may use the IN, ANY, or ALL operator. As you saw in Chapter 2, you can use these operators to see if a column value is contained in a list of values; for example:

```
SELECT product id, name
   FROM products
   WHERE product id IN (1, 2, 3);
   PRODUCT ID NAME
    _____
           1 Modern Science
           2 Chemistry
           3 Supernova
```

As you'll see in this section, the list of values can come from a subquery.



NOTE

You can also use the EXISTS operator to check if a value is in a list returned by a correlated subquery. You'll learn about this later, in the section "Writing Correlated Subqueries."

Using IN with a Multiple-Row Subquery

As you saw in Chapter 2, you use IN to check if a value is in a specified list of values. The list of values may come from the results returned by a subquery. You can also use NOT IN to perform the logical opposite of IN: to check if a value is not in a specified list of values.

The following simple example uses IN to check if a product id is in the list of values returned by the subquery; the subquery returns the product id for products whose name contains the letter e:

SELECT product id, name FROM products WHERE product id IN

```
(SELECT product id
  FROM products
  WHERE name LIKE '%e%');
PRODUCT ID NAME
_____
        1 Modern Science
        2 Chemistry
        3 Supernova
        5 Z Files
       6 2412: The Return
       7 Space Force 9
       8 From Another Planet
       11 Creative Yell
       12 My Front Line
```

The next example uses NOT IN to get the products that are not in the purchases table:

```
SELECT product id, name
    FROM products
    WHERE product id NOT IN
      (SELECT product id
      FROM purchases);
    PRODUCT ID NAME
    _____
            4 Tank War
            5 Z Files
            6 2412: The Return
            7 Space Force 9
            8 From Another Planet
            9 Classical Music
           10 Pop 3
           11 Creative Yell
           12 My Front Line
```

Using ANY with a Multiple-Row Subquery

You use the ANY operator to compare a value with any value in a list. You must place an =, <>, <, >, <=, or >= operator before ANY in your query. The following example uses ANY to get the employees whose salary is less than any of the lowest salaries in the salary grades table:

```
SELECT employee id, last name
     FROM employees
     WHERE salary < ANY
       (SELECT low salary
        FROM salary_grades);
     EMPLOYEE ID LAST NAME
     _____
              2 Johnson
              3 Hobbs
              4 Jones
```

Using ALL with a Multiple-Row Subquery

You use the ALL operator to compare a value with any value in a list. You must place an =, <>, <, >, <=, or >= operator before ALL in your query. The following example uses ALL to get the employees whose salary is greater than all of the highest salaries in the salary grades table:

```
SELECT employee id, last name
    FROM employees
    WHERE salary > ALL
       (SELECT high salary
       FROM salary grades);
     no rows selected
```

As you can see, no employee has a salary greater than the highest salary.

Writing Multiple-Column Subqueries

The subqueries you've seen so far have returned rows containing one column. You're not limited to one column: you can write subqueries that return multiple columns. The following example retrieves the products with the lowest price for each product type group:

```
SELECT product id, product type id, name, price
    FROM products
    WHERE (product type id, price) IN
       (SELECT product type id, MIN(price)
       FROM products
       GROUP BY product_type_id);
```

PRODUCT_ID	PRODUCT_TYPE_ID	NAME	PRICE
1	1	Modern Science	19.95
4	2	Tank War	13.95
8	3	From Another Planet	12.99
9	4	Classical Music	10.99

Notice that the subquery returns the product type id and the minimum price for each group of products, and these are compared in the outer query's WHERE clause with the product type id and price for each product.

Writing Correlated Subqueries

A correlated subguery references one or more columns in the outer SQL statement. These are called correlated subqueries, because they are related to the outer SQL statement through the same columns.

You typically use a correlated subguery when you need an answer to a question that depends on a value in each row contained in an outer query. For example, you might want to see whether there is a relationship between the data, but you don't care how many rows are returned by the subguery, that is, you just want to check whether any rows are returned, but you don't care how many.

A correlated subquery is run once for each row in the outer query; this is different from a noncorrelated subquery, which is run once prior to running the outer query. In addition, a correlated

subquery can resolve null values. You'll see examples in the following sections that illustrate these concepts.

A Correlated Subquery Example

The following correlated subquery retrieves the products that have a price greater than the average for their product type:

```
SELECT product_id, product_type_id, name, price
FROM products outer
WHERE price >
    (SELECT AVG(price)
    FROM products inner
    WHERE inner.product_type_id = outer.product_type_id);
```

PRODUCT_ID	PRODUCT_TYPE_ID	NAME	PRICE
2	1	Chemistry	30
5	2	Z Files	49.99
7	3	Space Force 9	13.49
10	4	Pop 3	15.99
11	4	Creative Yell	14.99

Notice that I've used the alias outer to label the outer query and the alias inner for the inner subquery. The reference to the product_type_id column in both the inner and outer parts is what makes the inner subquery correlated with the outer query. Also, the subquery returns a single row containing the average price for the product.

In a correlated subquery, each row in the outer query is passed one at a time to the subquery. The subquery reads each row in turn from the outer query and applies it to the subquery until all the rows from the outer query have been processed. The results from the entire query are then returned.

In the previous example, the outer query retrieves each row from the products table and passes them to the inner query. Each row is read by the inner query, which calculates the average price for each product where the product_type_id in the inner query is equal to the product_type_id in the outer query.

Using EXISTS and NOT EXISTS with a Correlated Subquery

You use the EXISTS operator to check for the existence of rows returned by a subquery. Although you can use EXISTS with non-correlated subqueries, you'll typically use it with correlated subqueries. The NOT EXISTS operator does the logical opposite of EXISTS: it checks if rows do not exist in the results returned by a subquery.

Using EXISTS with a Correlated Subquery

The following example uses EXISTS to retrieve employees who manage other employees; notice that I don't care how many rows are returned by the subquery; I only care whether any rows are returned at all:

```
SELECT employee_id, last_name
FROM employees outer
WHERE EXISTS
(SELECT employee_id
```

```
FROM employees inner
  WHERE inner.manager id = outer.employee id);
EMPLOYEE ID LAST NAME
_____
       1 Smith
        2 Johnson
```

Because EXISTS just checks for the existence of rows returned by the subquery, a subquery doesn't have to return a column—it can just return a literal value. This feature can improve the performance of your query. For example, the following query rewrites the previous example with the subquery returning the literal value 1:

```
SELECT employee id, last name
     FROM employees outer
     WHERE EXISTS
       (SELECT 1
        FROM employees inner
        WHERE inner.manager id = outer.employee id);
     EMPLOYEE ID LAST NAME
              1 Smith
               2 Johnson
```

As long as the subquery returns one or more rows, EXISTS returns true; if the subquery returns no rows, EXISTS returns false. In the examples, I didn't care how many rows are returned by the subguery: All I cared about was whether any rows (or no rows) are returned, so that EXISTS returns true (or false). Because the outer query requires at least one column, the literal value 1 is returned by the subquery in the previous example.

Using NOT EXISTS with a Correlated Subquery

The following example uses NOT EXISTS to retrieve products that haven't been purchased:

```
SELECT product id, name
    FROM products outer
    WHERE NOT EXISTS
      (SELECT 1
       FROM purchases inner
       WHERE inner.product id = outer.product id);
    PRODUCT ID NAME
    -----
             4 Tank War
             5 Z Files
             6 2412: The Return
             7 Space Force 9
             8 From Another Planet
            9 Classical Music
            10 Pop 3
            11 Creative Yell
            12 My Front Line
```

EXISTS and NOT EXISTS Versus IN and NOT IN

Earlier in the section "Using IN with a Multiple-Row Subquery," you saw how the IN operator is used to check if a value is contained in a list. EXISTS is different from IN: EXISTS checks just for the existence of rows, whereas IN checks for actual values.



TIP

EXISTS typically offers better performance than IN with subqueries. Therefore, you should use EXISTS rather than IN wherever possible.

You should be careful when writing queries that use NOT EXISTS or NOT IN. When a list of values contains a null value, NOT EXISTS returns true, but NOT IN returns false. Consider the following example that uses NOT EXISTS and retrieves the product types that don't have any products of that type in the products table:

```
SELECT product_type_id, name

FROM product_types outer

WHERE NOT EXISTS

(SELECT 1

FROM products inner

WHERE inner.product_type_id = outer.product_type_id);

PRODUCT_TYPE_ID NAME

5 Magazine
```

Notice one row is returned by this example. The next example rewrites the previous query to use NOT IN; notice that no rows are returned:

```
SELECT product_type_id, name
FROM product_types
WHERE product_type_id NOT IN
(SELECT product_type_id
FROM products);

no rows selected
```

No rows are returned because the subquery returns a list of product_id values, one of which is null (the product_type_id for product #12 is null). Because of this, NOT IN in the outer query returns false, and therefore no rows are returned. You can get around this by using the NVL() function to convert nulls to a value. In the following example, NVL() is used to convert null product type id values to 0:

This time the row appears.

These examples illustrate another difference between correlated and non-correlated subqueries: a correlated query can resolve null values.

Writing Nested Subqueries

You can nest subqueries inside other subqueries to a depth of 255. You should use this technique sparingly—you may find your query performs better using table joins. The following example contains a nested subquery; notice that it is contained within a subquery, which is itself contained in an outer query:

```
SELECT product type id, AVG(price)
     FROM products
     GROUP BY product type id
     HAVING AVG(price) <
       (SELECT MAX (AVG (price))
        FROM products
        WHERE product type id IN
          (SELECT product id
           FROM purchases
           WHERE quantity > 1)
        GROUP BY product type id)
     ORDER BY product type id;
     PRODUCT TYPE ID AVG (PRICE)
                  1 24.975
                  3 13.24
4 13.99
                         13.49
```

As you can see, this example is quite complex and contains three queries: a nested subquery, a subquery, and the outer query. These query parts are run in that order. Let's break the example down into the three parts and examine the results returned. The nested subquery is

```
SELECT product id
    FROM purchases
    WHERE quantity > 1
```

This subquery returns the product id for the products that have been purchased more than once. The rows returned by this subquery are

```
PRODUCT ID
```

The subquery that receives this output is

```
SELECT MAX (AVG (price))
     FROM products
     WHERE product type id IN
       (... output from the nested subguery ...)
     GROUP BY product type id
```

This subquery returns the maximum average price for the products returned by the nested subquery. The row returned is

```
MAX (AVG (PRICE))
------
26.22
```

This row is returned to the following outer query:

This query returns the product_type_id and average price of products that are less than average returned by the subquery. The rows returned are

PRODUCT_TYPE_ID	AVG(PRICE)
1	24.975
3	13.24
4	13.99
	13.49

These are the rows returned by the complete query shown at the start of this section.

Writing UPDATE and DELETE Statements Containing Subqueries

So far, you've only seen subqueries contained in a SELECT statement. As you'll see in this section, you can also put subqueries inside UPDATE and DELETE statements.

Writing an UPDATE Statement Containing a Subquery

In an UPDATE statement, you can set a column to the result returned by a single-row subquery. For example, the following UPDATE statement sets employee #4's salary to the average of the high salary grades returned by a subquery:

```
UPDATE employees

SET salary =

(SELECT AVG(high_salary)

FROM salary_grades)

WHERE employee_id = 4;
```

1 row updated.

Doing this increases employee #4's salary from \$500,000 to \$625,000 (this is the average of the high salaries from the salary grades table).



NOTE

If you execute the UPDATE statement, remember to execute a ROLLBACK to undo the change. That way, your results will match those shown later in this book.

Writing a DELETE Statement Containing a Subquery

You can use the rows returned by a subquery in the WHERE clause of a DELETE statement. For example, the following DELETE statement removes the employee whose salary is greater than the average of the high salary grades returned by a subquery:

DELETE FROM employees WHERE salary > (SELECT AVG(high salary) FROM salary grades);

1 row deleted.

This DELETE statement removes employee #1.



NOTE

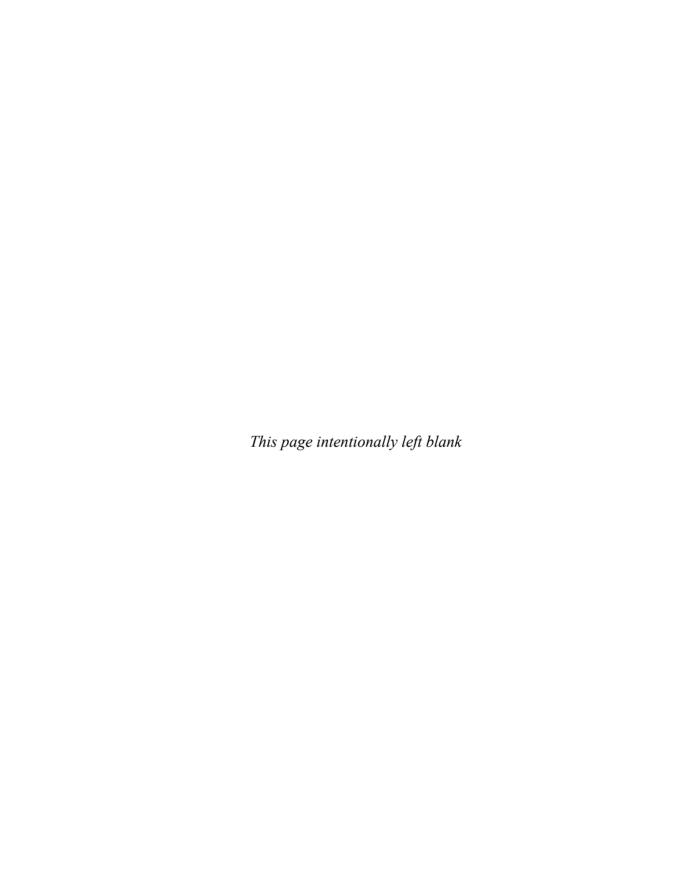
If you execute the DELETE statement, remember to execute a ROLLBACK to undo the removal of the row.

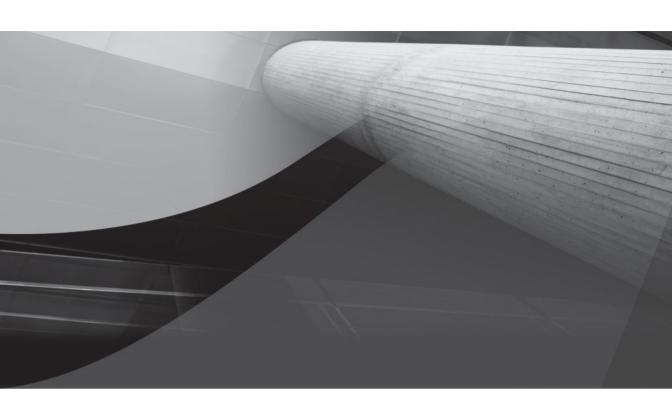
Summary

In this chapter, you learned the following:

- A subquery is a query placed within a SELECT, UPDATE, or DELETE statement.
- Single-row subqueries return zero or one row.
- Multiple-row subqueries return one or more rows.
- Multiple-column subqueries return more than one column.
- Correlated subqueries reference one or more columns in the outer SQL statement.
- Nested subqueries are subqueries placed within another subquery.

In the next chapter, you'll learn about advanced queries.





CHAPTER 7

Advanced Queries



n this chapter, you will see how to

- Use the set operators, which allow you to combine rows returned by two or more queries.
- Use the TRANSLATE() function to translate characters in one string to characters in another string.
- Use the DECODE () function to search for a certain value in a set of values.
- Use the CASE expression to perform if-then-else logic in SQL.
- Perform gueries on hierarchical data.
- Use the ROLLUP and CUBE clauses to get subtotals and totals for groups of rows.
- Take advantage of the analytic functions, which perform complex calculations, such as finding the top-selling product type for each month, the top salespersons, and so on.
- Perform inter-row calculations with the MODEL clause.
- Use the new Oracle Database 11*g* PIVOT and UNPIVOT clauses, which are useful for seeing overall trends in large amounts of data.

Let's plunge in and examine the set operators.

Using the Set Operators

The set operators allow you to combine rows returned by two or more queries. Table 7-1 shows the four set operators.

You must keep in mind the following restriction when using a set operator: The number of columns and the column types returned by the queries must match, although the column names may be different.

You'll learn how to use each of the set operators shown in Table 7-1 shortly, but first let's look at the example tables used in this section.

Operator	Description
UNION ALL	Returns all the rows retrieved by the queries, including duplicate rows.
UNION	Returns all non-duplicate rows retrieved by the queries.
INTERSECT	Returns rows that are retrieved by both queries.
MINUS	Returns the remaining rows when the rows retrieved by the second query are subtracted from the rows retrieved by the first query.

The Example Tables

The products and more products tables are created by the store schema.sql script using the following statements:

```
CREATE TABLE products (
      product id INTEGER
        CONSTRAINT products pk PRIMARY KEY,
      product type id INTEGER
        CONSTRAINT products fk product types
        REFERENCES product types (product type id),
      name VARCHAR2(30) NOT NULL,
      description VARCHAR2 (50),
      price NUMBER (5, 2)
    );
    CREATE TABLE more products (
      prd id INTEGER
        CONSTRAINT more products pk PRIMARY KEY,
      prd type id INTEGER
        CONSTRAINT more products fk product types
        REFERENCES product types (product type id),
      name VARCHAR2(30) NOT NULL,
      available CHAR(1)
    );
```

The following query retrieves the product id, product type id, and name columns from the products table:

SELECT product id, product type id, name FROM products;

```
PRODUCT_ID PRODUCT_TYPE_ID NAME
-----
                   1 Modern Science
       2
                   1 Chemistry
       3
                   2 Supernova
                   2 Tank War
       5
                    2 Z Files
                   2 2412: The Return
       6
       7
                    3 Space Force 9
       8
                    3 From Another Planet
       9
                   4 Classical Music
      10
                    4 Pop 3
      11
                   4 Creative Yell
                     My Front Line
```

The next query retrieves the prd id, prd type id, and name columns from the more products table:

```
SELECT prd id, prd type id, name
   FROM more products;
```

PRD_ID	PRD_TYPE_ID	NAME
1	1	Modern Science
2	1	Chemistry
3		Supernova
4	2	Lunar Landing
5	2	Submarine

Using the UNION ALL Operator

The UNION ALL operator returns all the rows retrieved by the queries, including duplicate rows. The following query uses UNION ALL; notice that all the rows from products and more_products are retrieved, including duplicates:

```
SELECT product_id, product_type_id, name
FROM products
UNION ALL
SELECT prd_id, prd_type_id, name
FROM more products;
```

```
PRODUCT ID PRODUCT TYPE ID NAME
                      1 Modern Science
        2
                      1 Chemistry
        3
                      2 Supernova
        4
                       2 Tank War
        5
                       2 Z Files
                       2 2412: The Return
        6
        7
                       3 Space Force 9
                       3 From Another Planet
        9
                       4 Classical Music
       10
                       4 Pop 3
       11
                       4 Creative Yell
       12
                        My Front Line
        1
                      1 Modern Science
                      1 Chemistry
        3
                        Supernova
                       2 Lunar Landing
                       2 Submarine
```

17 rows selected.

You can sort the rows using the ORDER BY clause followed by the position of the column. The following example uses ORDER BY 1 to sort the rows by the first column retrieved by the two queries (product id and prd id):

```
SELECT product_id, product_type_id, name
FROM products
UNION ALL
SELECT prd_id, prd_type_id, name
FROM more_products
ORDER BY 1;
```

```
PRODUCT_ID PRODUCT_TYPE_ID NAME
                       1 Modern Science
        1
                       1 Modern Science
        2
                       1 Chemistry
                       1 Chemistry
        2.
        3
                       2 Supernova
        3
                        Supernova
        4
                       2 Tank War
                        2 Lunar Landing
        4
        5
                        2 Z Files
        5
                        2 Submarine
                       2 2412: The Return
        7
                       3 Space Force 9
        8
                       3 From Another Planet
                       4 Classical Music
        9
       10
                        4 Pop 3
       11
                       4 Creative Yell
       12
                         My Front Line
```

17 rows selected.

Using the UNION Operator

The UNION operator returns only the non-duplicate rows retrieved by the queries. The following example uses UNION; notice the duplicate "Modern Science" and "Chemistry" rows are not retrieved, and so only 15 rows are returned:

```
SELECT product_id, product_type_id, name
    FROM products
    UNION
    SELECT prd id, prd type id, name
    FROM more products;
```

```
PRODUCT ID PRODUCT TYPE ID NAME
        1
                       1 Modern Science
                       1 Chemistry
        2
        3
                       2 Supernova
        3
                         Supernova
                       2 Lunar Landing
        4
        4
                        2 Tank War
        5
                        2 Submarine
        5
                        2 Z Files
                        2 2412: The Return
        6
        7
                       3 Space Force 9
        8
                        3 From Another Planet
        9
                       4 Classical Music
       10
                       4 Pop 3
       11
                       4 Creative Yell
                         My Front Line
       12
```

15 rows selected.

Using the INTERSECT Operator

The INTERSECT operator returns only rows that are retrieved by both queries. The following example uses INTERSECT; notice that the "Modern Science" and "Chemistry" rows are returned:

```
SELECT product_id, product_type_id, name
FROM products
INTERSECT
SELECT prd_id, prd_type_id, name
FROM more_products;

PRODUCT_ID PRODUCT_TYPE_ID NAME

1 1 Modern Science
2 1 Chemistry
```

Using the MINUS Operator

The MINUS operator returns the remaining rows when the rows retrieved by the second query are subtracted from the rows retrieved by the first query. The following example uses MINUS; notice that the rows from more_products are subtracted from products and the remaining rows are returned:

```
SELECT product_id, product_type_id, name
FROM products
MINUS
SELECT prd_id, prd_type_id, name
FROM more products;
```

PRODUCT_ID	PRODUCT_TYPE_ID	NAME
3	2	Supernova
4	2	Tank War
5	2	Z Files
6	2	2412: The Return
7	3	Space Force 9
8	3	From Another Planet
9	4	Classical Music
10	4	Pop 3
11	4	Creative Yell
12		My Front Line

10 rows selected.

Combining Set Operators

You can combine more than two queries with multiple set operators, with the returned results from one operator feeding into the next operator. By default, set operators are evaluated from top to bottom, but you should indicate the order using parentheses in case Oracle Corporation changes this default behavior in future software releases.

In the examples in this section, I'll use the following product_changes table (created by the store_schema.sql script):

```
CREATE TABLE product changes (
       product id INTEGER
         CONSTRAINT prod changes pk PRIMARY KEY,
       product_type id INTEGER
         CONSTRAINT prod changes fk product types
         REFERENCES product types (product type id),
       name VARCHAR2(30) NOT NULL,
       description VARCHAR2 (50),
       price NUMBER(5, 2)
     );
```

The following query returns the product id, product type id, and name columns from the product changes table:

SELECT product id, product type id, name FROM product changes;

```
PRODUCT ID PRODUCT TYPE ID NAME
_____
              1 Modern Science
                  1 New Chemistry
      3
                 1 Supernova
                 2 Lunar Landing
     13
     1.4
                 2 Submarine
     1.5
                  2 Airplane
```

The next query does the following:

- Uses the UNION operator to combine the results from the products and more products tables. (The UNION operator returns only the non-duplicate rows retrieved by the queries.)
- Uses the INTERSECT operator to combine the results from the previous UNION operator with the results from the product changes table. (The INTERSECT operator only returns rows that are retrieved by both queries.)
- Uses parentheses to indicate the order of evaluation, which is: (1) the UNION between the products and more products tables; (2) the INTERSECT.

```
(SELECT product id, product type id, name
    FROM products
    UNION
    SELECT prd id, prd type id, name
    FROM more products)
    INTERSECT
    SELECT product id, product type id, name
    FROM product changes;
    PRODUCT ID PRODUCT TYPE ID NAME
    _____
                         1 Modern Science
```

The following query has the parentheses set so that the INTERSECT is performed first; notice the different results returned by the query compared with the previous example:

```
SELECT product id, product type id, name
    FROM products
    UNION
    (SELECT prd id, prd type id, name
    FROM more products
    INTERSECT
    SELECT product id, product type id, name
    FROM product changes);
    PRODUCT ID PRODUCT TYPE ID NAME
                            1 Modern Science
                            1 Chemistry
             3
                            2 Supernova
             4
                             2 Tank War
             5
                            2 Z Files
             6
                             2 2412: The Return
             7
                            3 Space Force 9
             8
                            3 From Another Planet
             9
                             4 Classical Music
            10
                            4 Pop 3
            11
                            4 Creative Yell
```

This concludes the discussion of the set operators.

12

Using the TRANSLATE() Function

TRANSLATE (x, from_string, to_string) converts the occurrences of characters in from_string found in x to corresponding characters in to_string. The easiest way to understand how TRANSLATE () works is to see some examples.

The following example uses TRANSLATE () to shift each character in the string SECRET MESSAGE: MEET ME IN THE PARK by four places to the right: A becomes E, B becomes F, and so on:

My Front Line

The next example takes the output of the previous example and shifts the characters four places to the left: E becomes A, F becomes B, and so on:

```
SELECT TRANSLATE('WIGVIX QIWWEKI: QIIX QI MR XLI TEVO',
'EFGHIJKLMNOPQRSTUVWXYZABCD',
```

```
'ABCDEFGHIJKLMNOPQRSTUVWXYZ')
FROM dual:
TRANSLATE ('WIGVIXOIWWEKI: OIIXOIMRXL
_____
SECRET MESSAGE: MEET ME IN THE PARK
```

You can of course pass column values to TRANSLATE (). The following example passes the name column from the products table to TRANSLATE(), which shifts the letters in the product name four places to the right:

```
SELECT product id, TRANSLATE(name,
     'ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz',
     'EFGHIJKLMNOPQRSTUVWXYZABCDefghijklmnopgrstuvwxyzabcd')
   FROM products;
   PRODUCT ID TRANSLATE (NAME, 'ABCDEFGHIJKLMN
   _____
            1 Qshivr Wqmirqi
            2 Gligmwxvc
            3 Wytivrsze
            4 Xero Aev
            5 D Jmpiw
            6 2412: Xli Vixyvr
            7 Wtegi Jsvgi 9
            8 Jvsq Ersxliv Tperix
            9 Gpewwmgep Qywmg
           10 Tst 3
           11 Gviexmzi Cipp
           12 Oc Jvsrx Pmri
```

You can also use TRANSLATE() to convert numbers. The following example takes the number 12345 and converts 5 to 6, 4 to 7, 3 to 8, 2 to 9, and 1 to 0:

```
SELECT TRANSLATE (12345,
        54321,
        67890)
      FROM dual;
     TRANS
      ____
      09876
```

Using the DECODE() Function

DECODE (value, search value, result, default value) compares value with search value. If the values are equal, DECODE () returns result; otherwise, default value is returned. DECODE() allows you to perform if-then-else logic in SQL without having to use PL/SQL. Each of the parameters to DECODE () can be a column, a literal value, a function, or a subquery.



NOTE

DECODE () is an old Oracle proprietary function, and therefore you should use CASE expressions instead if you are using Oracle Database 9i and above (you will learn about CASE in the next section). The DECODE () function is mentioned here because you may encounter it when using older Oracle databases.

The following example illustrates the use of DECODE () with literal values; DECODE () returns 2 (1 is compared with 1, and because they are equal 2 is returned):

The next example uses DECODE () to compare 1 to 2, and because they are not equal 3 is returned:

The next example compares the available column in the more_products table; if available equals Y, the string 'Product is available' is returned; otherwise, 'Product is not available' is returned:

```
SELECT prd_id, available,

DECODE(available, 'Y', 'Product is available',
    'Product is not available')

FROM more_products;

PRD_ID A DECODE(AVAILABLE, 'Y', 'PR

1 Y Product is available
2 Y Product is available
3 N Product is not available
4 N Product is not available
5 Y Product is available
```

You can pass multiple search and result parameters to DECODE (), as shown in the following example, which returns the product type id column as the name of the product type:

```
SELECT product_id, product_type_id,

DECODE(product_type_id,

1, 'Book',

2, 'Video',

3, 'DVD',

4, 'CD',

'Magazine')
```

FROM products;

PRODUCT_ID	PRODUCT_TYPE_ID	DECODE (P
1	1	Book
2	1	Book
3	2	Video
4	2	Video
5	2	Video
6	2	Video
7	3	DVD
8	3	DVD
9	4	CD
10	4	CD
11	4	CD
12		Magazine

Notice that

- If product type id is 1, Book is returned.
- If product type id is 2, Video is returned.
- If product type id is 3, DVD is returned.
- If product type id is 4, CD is returned.
- If product type id is any other value, Magazine is returned.

Using the CASE Expression

The CASE expression performs if-then-else logic in SQL and is supported in Oracle Database 9i and above. The CASE expression works in a similar manner to DECODE (), but you should use CASE because it is ANSI-compliant and forms part of the SQL/92 standard. In addition, the CASE expression is easier to read.

There are two types of CASE expressions:

- Simple case expressions, which use expressions to determine the returned value
- Searched case expressions, which use conditions to determine the returned value

You'll learn about both of these types of CASE expressions next.

Using Simple CASE Expressions

Simple CASE expressions use embedded expressions to determine the value to return. Simple CASE expressions have the following syntax:

```
CASE search expression
      WHEN expression1 THEN result1
      WHEN expression2 THEN result2
      WHEN expressionN THEN resultN
      ELSE default result
    END
```

where

- search expression is the expression to be evaluated.
- expression1, expression2, ..., expressionN are the expressions to be evaluated against search expression.
- result1, result2, ..., resultN are the returned results (one for each possible expression). If expression1 evaluates to search_expression, result1 is returned, and similarly for the other expressions.
- default_result is returned when no matching expression is found.

The following example shows a simple CASE expression that returns the product types as names:

```
SELECT product id, product type id,
      CASE product type id
       WHEN 1 THEN 'Book'
       WHEN 2 THEN 'Video'
       WHEN 3 THEN 'DVD'
       WHEN 4 THEN 'CD'
       ELSE 'Magazine'
      END
    FROM products;
    PRODUCT_ID PRODUCT_TYPE_ID CASEPROD
    _____
            1
                         1 Book
                          1 Book
            3
                          2 Video
                          2 Video
            5
                          2 Video
```

6

7

8

9

10

11

12

Using Searched CASE Expressions

2 Video

3 DVD

3 DVD

4 CD

4 CD

4 CD

Magazine

Searched CASE expressions use conditions to determine the returned value. Searched CASE expressions have the following syntax:

```
CASE

WHEN condition1 THEN result1

WHEN condition2 THEN result2

...

WHEN conditionN THEN resultN

ELSE default_result

END
```

where

- condition1, condition2, ..., conditionN are the expressions to be evaluated.
- result1, result2, ..., resultN are the returned results (one for each possible condition). If condition1 is true, result1 is returned, and similarly for the other expressions.
- default result is returned when there is no condition that returns true.

The following example illustrates the use of a searched CASE expression:

```
SELECT product id, product type id,
     CASE
       WHEN product type id = 1 THEN 'Book'
       WHEN product type id = 2 THEN 'Video'
       WHEN product type id = 3 THEN 'DVD'
       WHEN product type id = 4 THEN 'CD'
       ELSE 'Magazine'
     END
   FROM products;
   PRODUCT_ID PRODUCT_TYPE_ID CASEPROD
   _____
           1
                         1 Book
            2
                          1 Book
            3
                          2 Video
                          2 Video
           4
           5
                          2 Video
            6
                          2 Video
           7
                          3 DVD
           8
                          3 DVD
           9
                          4 CD
           10
                          4 CD
                          4 CD
           11
                            Magazine
```

You can use operators in a searched CASE expression, as shown in the following example:

```
SELECT product id, price,
        WHEN price > 15 THEN 'Expensive'
        ELSE 'Cheap'
      END
    FROM products;
```

PRODUCT_ID	PRICE	CASEWHENP
1	19.95	Expensive
2	30	Expensive
3	25.99	Expensive
4	13.95	Cheap
5	49.99	Expensive

```
6 14.95 Cheap
7 13.49 Cheap
8 12.99 Cheap
9 10.99 Cheap
10 15.99 Expensive
11 14.99 Cheap
12 13.49 Cheap
```

You will see more advanced examples of CASE expressions later in this chapter and in Chapter 16.

Hierarchical Queries

You'll quite often encounter data that is organized in a hierarchical manner. Examples include the people who work in a company, a family tree, and the parts that make up an engine. In this section, you'll see queries that access a hierarchy of employees who work for our imaginary store.

The Example Data

You'll see the use of a table named more_employees, which is created by the store_schema.sql script as follows:

```
CREATE TABLE more_employees (
employee_id INTEGER

CONSTRAINT more_employees_pk PRIMARY KEY,
manager_id INTEGER

CONSTRAINT more_empl_fk_fk_more_empl

REFERENCES more_employees(employee_id),
first_name VARCHAR2(10) NOT NULL,
last_name VARCHAR2(10) NOT NULL,
title VARCHAR2(20),
salary NUMBER(6, 0)
);
```

The manager_id column is a self-reference back to the employee_id column of the more_employees table; manager_id indicates the manager of an employee (if any). The following query returns the rows from more employees:

SELECT * FROM more employees;

EMPLOYEE_ID	MANAGER_ID	FIRST_NAME	LAST_NAME	TITLE	SALARY
1		James	Smith	CEO	800000
2	1	Ron	Johnson	Sales Manager	600000
3	2	Fred	Hobbs	Sales Person	200000
4	1	Susan	Jones	Support Manager	500000
5	2	Rob	Green	Sales Person	40000
6	4	Jane	Brown	Support Person	45000
7	4	John	Grey	Support Manager	30000
8	7	Jean	Blue	Support Person	29000

9	6 Henry	Heyson	Support Person	30000
10	1 Kevin	Black	Ops Manager	100000
11	10 Keith	Long	Ops Person	50000
12	10 Frank	Howard	Ops Person	45000
13	10 Doreen	Penn	Ops Person	47000

As you can see, it's difficult to pick out the employee relationships from this data. Figure 7-1 shows the relationships in a graphical form.

As you can see from Figure 7-1, the elements—or nodes—form a tree. Trees of nodes have the following technical terms associated with them:

- **Root node** The root is the node at the top of the tree. In the example shown in Figure 7-1, the root node is James Smith, the CEO.
- **Parent node** A parent is a node that has one or more nodes beneath it. For example, James Smith is the parent to the following nodes: Ron Johnson, Susan Jones, and Kevin Black.
- **Child node** A child is a node that has one parent node above it. For example, Ron Johnson's parent is James Smith.
- **Leaf node** A leaf is a node that has no children. For example, Fred Hobbs and Rob Green are leaf nodes.

You use the CONNECT BY and START WITH clauses of a SELECT statement to perform hierarchical gueries, as described next.

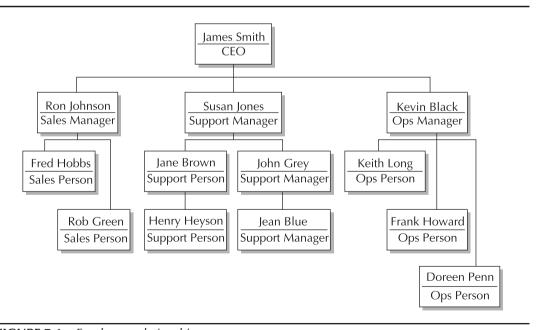


FIGURE 7-1 *Employee relationships*

Using the CONNECT BY and START WITH Clauses

The syntax for the CONNECT BY and START WITH clauses of a SELECT statement is

```
SELECT [LEVEL], column, expression, ...

FROM table
[WHERE where_clause]
[[START WITH start_condition] [CONNECT BY PRIOR prior_condition]];
```

where

- LEVEL is a pseudo column that tells you how far into a tree you are. LEVEL returns 1 for a root node, 2 for a child of the root, and so on.
- start_condition specifies where to start the hierarchical query. You must specify a START WITH clause when writing a hierarchical query. An example start_condition is employee id = 1, which specifies the query starts from employee #1.
- prior_condition specifies the relationship between the parent and child rows. You must specify a CONNECT BY PRIOR clause when writing a hierarchical query. An example prior_condition is employee_id = manager_id, which specifies the relationship is between the parent employee_id and the child manager_id—that is, the child's manager_id points to the parent's employee id.

The following query illustrates the use of the START WITH and CONNECT BY PRIOR clauses; notice that the first row contains the details of James Smith (employee #1), the second row contains the details of Ron Johnson, whose manager id is 1, and so on:

```
SELECT employee_id, manager_id, first_name, last_name
FROM more_employees
START WITH employee_id = 1
CONNECT BY PRIOR employee id = manager id;
```

EMPLOYEE_ID	MANAGER_ID	FIRST_NAME	LAST_NAME
1		James	Smith
2	1	Ron	Johnson
3	2	Fred	Hobbs
5	2	Rob	Green
4	1	Susan	Jones
6	4	Jane	Brown
9	6	Henry	Heyson
7	4	John	Grey
8	7	Jean	Blue
10	1	Kevin	Black
11	10	Keith	Long
12	10	Frank	Howard
13	10	Doreen	Penn

Using the LEVEL Pseudo Column

The next query illustrates the use of the LEVEL pseudo column to display the level in the tree:

```
SELECT LEVEL, employee id, manager id, first name, last name
    FROM more employees
    START WITH employee id = 1
    CONNECT BY PRIOR employee id = manager id
    ORDER BY LEVEL;
```

LEVEL	EMPLOYEE_ID	MANAGER_ID	FIRST_NAME	LAST_NAME
1	1		James	Smith
2	2	1	Ron	Johnson
2	4	1	Susan	Jones
2	10	1	Kevin	Black
3	3	2	Fred	Hobbs
3	7	4	John	Grey
3	12	10	Frank	Howard
3	13	10	Doreen	Penn
3	11	10	Keith	Long
3	5	2	Rob	Green
3	6	4	Jane	Brown
4	9	6	Henry	Heyson
4	8	7	Jean	Blue

The next query uses the COUNT () function and LEVEL to get the number of levels in the tree:

```
SELECT COUNT (DISTINCT LEVEL)
   FROM more employees
   START WITH employee id = 1
   CONNECT BY PRIOR employee id = manager id;
   COUNT (DISTINCTLEVEL)
   _____
```

2 Ron Johnson

Formatting the Results from a Hierarchical Query

You can format the results from a hierarchical query using LEVEL and the LPAD() function, which left-pads values with characters. The following query uses LPAD(' ', 2 * LEVEL -1) to left-pad a total of 2 * LEVEL - 1 spaces; the result indents an employee's name with spaces based on their LEVEL (that is, LEVEL 1 isn't padded, LEVEL 2 is padded by two spaces, LEVEL 3 by four spaces, and so on):

```
SET PAGESIZE 999
    COLUMN employee FORMAT A25
    SELECT LEVEL,
      LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee
    FROM more employees
    START WITH employee id = 1
    CONNECT BY PRIOR employee id = manager id;
        LEVEL EMPLOYEE
    -----
            1 James Smith
```

```
3 Fred Hobbs
3 Rob Green
2 Susan Jones
3 Jane Brown
4 Henry Heyson
3 John Grey
4 Jean Blue
2 Kevin Black
3 Keith Long
3 Frank Howard
3 Doreen Penn
```

The employee relationships are easy to pick out from these results.

Starting at a Node Other than the Root

You don't have to start at the root node when traversing a tree: you can start at any node using the START WITH clause. The following query starts with Susan Jones; notice that LEVEL returns 1 for Susan Jones, 2 for Jane Brown, and so on:

```
SELECT LEVEL,

LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee
FROM more_employees
START WITH last_name = 'Jones'
CONNECT BY PRIOR employee_id = manager_id;

LEVEL EMPLOYEE

1 Susan Jones
2 Jane Brown
3 Henry Heyson
2 John Grey
3 Jean Blue
```

If the store had more than one employee with the same name, you could simply use the employee_id in the query's START WITH clause. For example, the following query uses Susan Jones' employee_id of 4:

```
SELECT LEVEL,

LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee

FROM more_employees

START WITH employee_id = 4

CONNECT BY PRIOR employee_id = manager_id;
```

This query returns the same rows as the previous one.

Using a Subquery in a START WITH Clause

You can use a subquery in a START WITH clause. For example, the following query uses a subquery to select the <code>employee_id</code> whose name is Kevin Black; this <code>employee_id</code> is passed to the START WITH clause:

```
SELECT LEVEL,

LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee
```

```
FROM more employees
START WITH employee id = (
 SELECT employee id
 FROM more employees
 WHERE first name = 'Kevin'
 AND last name = 'Black'
CONNECT BY PRIOR employee id = manager id;
    LEVEL EMPLOYEE
        1 Kevin Black
        2 Keith Long
            Frank Howard
            Doreen Penn
```

Traversing Upward Through the Tree

You don't have to traverse a tree downward from parents to children: you can start at a child and traverse upward. You do this by switching child and parent columns in the CONNECT BY PRIOR clause. For example, CONNECT BY PRIOR manager id = employee id connects the child's manager id to the parent's employee id.

The following query starts with Jean Blue and traverses upward all the way to James Smith; notice that LEVEL returns 1 for Jean Blue, 2 for John Grey, and so on:

```
SELECT LEVEL,
     LPAD(' ', 2 * LEVEL - 1) || first name || ' ' || last name AS employee
   FROM more employees
    START WITH last name = 'Blue'
    CONNECT BY PRIOR manager id = employee id;
       LEVEL EMPLOYEE
    _____
           1 Jean Blue
           2 John Grey
               Susan Jones
           3
                  James Smith
```

Eliminating Nodes and Branches from a Hierarchical Query

You can eliminate a particular node from a query tree using a WHERE clause. The following query eliminates Ron Johnson from the results using WHERE last name != 'Johnson':

```
SELECT LEVEL,
      LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee
    FROM more employees
    WHERE last name != 'Johnson'
    START WITH employee id = 1
    CONNECT BY PRIOR employee id = manager id;
        LEVEL EMPLOYEE
    _____
            1 James Smith
            3 Fred Hobbs
```

```
3
   Rob Green
2
 Susan Jones
3
   Jane Brown
     Henry Heyson
4
   John Grev
3
4
     Jean Blue
2 Kevin Black
3 Keith Long
   Frank Howard
3
    Doreen Penn
```

You'll notice that although Ron Johnson is eliminated from the results, his employees Fred Hobbs and Rob Green are still included. To eliminate an entire branch of nodes from the results of a query, you add an AND clause to your CONNECT BY PRIOR clause. For example, the following query uses AND last_name != 'Johnson' to eliminate Ron Johnson and all his employees from the results:

```
SELECT LEVEL,
      LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee
    FROM more employees
    START WITH employee id = 1
    CONNECT BY PRIOR employee_id = manager_id
    AND last name != 'Johnson';
         LEVEL EMPLOYEE
    _____
            1 James Smith
            2 Susan Jones
            3
                 Jane Brown
            4
                   Henry Heyson
            3
                 John Grey
            4
                    Jean Blue
            2 Kevin Black
            3
                Keith Long
            3
                 Frank Howard
                Doreen Penn
```

Including Other Conditions in a Hierarchical Query

You can include other conditions in a hierarchical query using a WHERE clause. The following example uses a WHERE clause to show only employees whose salaries are less than or equal to \$50,000:

```
SELECT LEVEL,

LPAD(' ', 2 * LEVEL - 1) || first_name || ' ' || last_name AS employee,
salary

FROM more_employees
WHERE salary <= 50000
START WITH employee_id = 1
CONNECT BY PRIOR employee id = manager id;
```

LEVEL	EMPLOYEE	SALARY
3	Rob Green	40000
3	Jane Brown	45000
4	Henry Heyson	30000
3	John Grey	30000
4	Jean Blue	29000
3	Keith Long	50000
3	Frank Howard	45000
3	Doreen Penn	47000

This concludes the discussion of hierarchical queries. In the next section, you'll learn about advanced group clauses.

Using the Extended GROUP BY Clauses

In this section, you'll learn about

- ROLLUP, which extends the GROUP BY clause to return a row containing a subtotal for each group of rows, plus a row containing a grand total for all the groups.
- CUBE, which extends the GROUP BY clause to return rows containing a subtotal for all combinations of columns, plus a row containing the grand total.

First, let's look at the example tables used in this section.

The Example Tables

You'll see the use of the following tables that refine the representation of employees in our imaginary store:

- divisions, which stores the divisions within the company
- jobs, which stores the jobs within the company
- employees2, which stores the employees

These tables are created by the store schema.sql script. The divisions table is created using the following statement:

```
CREATE TABLE divisions (
      division id CHAR(3)
        CONSTRAINT divisions pk PRIMARY KEY,
      name VARCHAR2(15) NOT NULL
    );
```

The following query retrieves the rows from the divisions table:

SELECT *

FROM divisions:

```
DIV NAME
--- -----
SAL Sales
```

```
OPE Operations
SUP Support
BUS Business
```

The jobs table is created using the following statement:

```
CREATE TABLE jobs (
    job_id CHAR(3)
        CONSTRAINT jobs_pk PRIMARY KEY,
    name VARCHAR2(20) NOT NULL
);
```

The next query retrieves the rows from the jobs table:

SELECT *

FROM jobs;

The employees2 table is created using the following statement:

The following query retrieves the first five rows from the employees2 table:

SELECT *

FROM employees2 WHERE ROWNUM <= 5;</pre>

EMPLOYEE_ID	DIV	JOB	FIRST_NAME	LAST_NAME	SALARY
1	BUS	PRE	James	Smith	800000
2	SAL	MGR	Ron	Johnson	350000
3	SAL	WOR	Fred	Hobbs	140000
4	SUP	MGR	Susan	Jones	200000
5	SAL	WOR	Rob	Green	350000

Using the ROLLUP Clause

The ROLLUP clause extends GROUP BY to return a row containing a subtotal for each group of rows, plus a row containing a total for all the groups.

As you saw in Chapter 4, you use GROUP BY to group rows into blocks with a common column value. For example, the following query uses GROUP BY to group the rows from the employees2 table by department id and uses SUM() to get the sum of the salaries for each division id:

```
SELECT division id, SUM(salary)
    FROM employees2
    GROUP BY division id
    ORDER BY division id;
```

```
DIV SUM(SALARY)
___ ____
BUS
        1610000
OPE 1320000
SAL 4936000
SUP 1015000
```

Passing a Single Column to ROLLUP

The following query rewrites the previous example to use ROLLUP; notice the additional row at the end, which contains the total salaries for all the groups:

```
SELECT division_id, SUM(salary)
    FROM employees2
    GROUP BY ROLLUP (division id)
    ORDER BY division id;
```

DIV	SUM (SALARY)
BUS	1610000
OPE	1320000
SAL	4936000
SUP	1015000
	8881000



NOTE

If you need the rows in a specific order, you should use an ORDER BY clause. You need to do this just in case Oracle Corporation decides to change the default order of rows returned by ROLLUP.

Passing Multiple Columns to ROLLUP

You can pass multiple columns to ROLLUP, which then groups the rows into blocks with the same column values. The following example passes the division id and job id columns of the employees2 table to ROLLUP, which groups the rows by those columns; in the output, notice that the salaries are summed by division id and job id, and that ROLLUP returns a row

with the sum of the salaries in each division_id, plus a row at the end with the salary grand total:

```
SELECT division_id, job_id, SUM(salary)
FROM employees2
GROUP BY ROLLUP(division_id, job_id)
ORDER BY division_id, job_id;
```

Changing the Position of Columns Passed to ROLLUP

The next example switches division_id and job_id; this causes ROLLUP to calculate the sum of the salaries for each job_id:

```
SELECT job_id, division_id, SUM(salary)
FROM employees2
GROUP BY ROLLUP(job_id, division_id)
ORDER BY job id, division id;
```

```
        JOB
        DIV
        SUM(SALARY)

        ---
        ---

        ENG
        OPE
        245000

        ENG
        245000
        MGR
        530000

        MGR
        OPE
        805000

        MGR
        SAL
        4446000

        MGR
        G246000
        PRE
        800000

        PRE
        BUS
        800000

        TEC
        SUP
        115000

        WOR
        BUS
        280000

        WOR
        OPE
        270000
```

WOR	SAL	490000
WOR	SUP	435000
WOR		1475000
		8881000

Using Other Aggregate Functions with ROLLUP

You can use any of the aggregate functions with ROLLUP (for a list of the main aggregate functions, see Table 4-8 in Chapter 4). The following example uses AVG () to calculate the average salaries:

```
SELECT division id, job id, AVG(salary)
     FROM employees2
     GROUP BY ROLLUP (division_id, job_id)
     ORDER BY division id, job id;
```

```
DIV JOB AVG (SALARY)
--- --- -----
BUS MGR 176666.667
BUS PRE 800000
BUS WOR 280000
BUS 322000
OPE ENG 245000
OPE MGR 201250
OPE WOR 135000
OPE 188571.429
SAL MGR 261529.412
SAL WOR 245000
SAL 259789.474

        SUP MGR
        232500

        SUP TEC
        115000

        SUP WOR
        145000

SUP 169166.667
           240027.027
```

Using the CUBE Clause

The CUBE clause extends GROUP BY to return rows containing a subtotal for all combinations of columns, plus a row containing the grand total. The following example passes division id and job id to CUBE, which groups the rows by those columns:

```
SELECT division_id, job_id, SUM(salary)
     FROM employees2
      GROUP BY CUBE (division_id, job_id)
     ORDER BY division id, job id;
```

```
DIV JOB SUM(SALARY)
--- --- -----
BUS MGR 530000
BUS PRE 800000
BUS WOR 280000
BUS 1610000
```

208

```
OPE ENG
        245000
         805000
OPE MGR
OPE WOR
         270000
OPE 1320000
SAL MGR 4446000
SAL WOR
        490000
SAL
        4936000
SUP MGR
        465000
         115000
SUP TEC
SUP WOR
         435000
        1015000
SUP
        245000
   ENG
   MGR
        6246000
   PRE
         800000
   TEC
         115000
        1475000
   WOR
          8881000
```

Notice that the salaries are summed by division_id and job_id. CUBE returns a row with the sum of the salaries for each division_id, along with the sum of all salaries for each job_id near the end. At the very end is a row with the grand total of the salaries.

The next example switches division id and job id:

SELECT job_id, division_id, SUM(salary)
FROM employees2
GROUP BY CUBE(job_id, division_id)
ORDER BY job id, division id;

```
JOB DIV SUM(SALARY)
--- --- -----
ENG OPE 245000
ENG
             245000
MGR BUS 530000
MGR OPE 805000
MGR SAL 4446000
MGR SUP 465000
MGR 6246000
PRE BUS 800000
PRE
             800000
           115000
115000
TEC SUP
TEC
WOR BUS 280000
WOR OPE 270000
WOR SAL
            490000
WOR SUP
             435000
WOR
           1475000
    BUS
           1610000
           1320000
    OPE
    SAL
           4936000
           1015000
    SUP
             8881000
```

Using the GROUPING() Function

The GROUPING () function accepts a column and returns 0 or 1. GROUPING () returns 1 when the column value is null and returns 0 when the column value is non-null. GROUPING () is used only in queries that use ROLLUP or CUBE. GROUPING() is useful when you want to display a value when a null would otherwise be returned.

Using GROUPING() with a Single Column in a ROLLUP

As you saw earlier in the section "Passing a Single Column to ROLLUP," the last row in the example's result set contained a total of the salaries:

```
SELECT division id, SUM(salary)
    FROM employees2
    GROUP BY ROLLUP (division id)
    ORDER BY division id;
    DIV SUM (SALARY)
    --- -----
    BUS 1610000
    OPE
           1320000
    SAL 4936000
SUP 1015000
           8881000
```

The division id column for the last row is null. You can use the GROUPING() function to determine whether this column is null, as shown in the following query; notice GROUPING() returns 0 for the rows that have non-null division id values and returns 1 for the last row that has a null division id:

```
SELECT GROUPING (division id), division id, SUM(salary)
      FROM employees2
      GROUP BY ROLLUP (division id)
      ORDER BY division id;
      GROUPING (DIVISION ID) DIV SUM (SALARY)
                            0 BUS 1610000
                            0 OPE 1320000
0 SAL 4936000
0 SUP 1015000
1 8881000
                                      8881000
```

Using CASE to Convert the Returned Value from GROUPING()

You can use the CASE expression to convert the 1 in the previous example to a meaningful value. The following example uses CASE to convert 1 to the string 'All divisions':

```
SELECT
       CASE GROUPING (division id)
         WHEN 1 THEN 'All divisions'
         ELSE division id
       END AS div,
       SUM(salary)
```

```
FROM employees2
GROUP BY ROLLUP(division_id)
ORDER BY division_id;
```

DIV		SUM(SALARY)
BUS		1610000
OPE		1320000
SAL		4936000
SUP		1015000
All	divisions	8881000

Using CASE and GROUPING() to Convert Multiple Column Values

The next example extends the idea of replacing null values to a ROLLUP containing multiple columns (division_id and job_id); notice that null division_id values are replaced with the string 'All divisions' and that null job_id values are replaced with 'All jobs':

SELECT

```
CASE GROUPING (division_id)

WHEN 1 THEN 'All divisions'

ELSE division_id

END AS div,

CASE GROUPING (job_id)

WHEN 1 THEN 'All jobs'

ELSE job_id

END AS job,

SUM (salary)

FROM employees2

GROUP BY ROLLUP (division_id, job_id)

ORDER BY division_id, job_id;
```

DIV		JOB		SUM(SALARY)
BUS		MGR		530000
BUS		PRE		800000
BUS		WOR		280000
BUS		All	jobs	1610000
OPE		ENG		245000
OPE		MGR		805000
OPE		WOR		270000
OPE		All	jobs	1320000
SAL		MGR		4446000
SAL		WOR		490000
SAL		All	jobs	4936000
SUP		MGR		465000
SUP		TEC		115000
SUP		WOR		435000
SUP		All	jobs	1015000
All	divisions	All	jobs	8881000

Using GROUPING() with CUBE

You can use the GROUPING() function with CUBE, as in this example:

```
SELECT
```

```
CASE GROUPING (division id)
   WHEN 1 THEN 'All divisions'
   ELSE division id
 END AS div,
 CASE GROUPING (job id)
   WHEN 1 THEN 'All jobs'
   ELSE job id
 END AS job,
 SUM(salary)
FROM employees2
GROUP BY CUBE (division id, job id)
ORDER BY division id, job id;
```

DIV		JOB		SUM(SALARY)
BUS		MGR		530000
BUS		PRE		800000
BUS		WOR		280000
BUS		All	jobs	1610000
OPE		ENG		245000
OPE		MGR		805000
OPE		WOR		270000
OPE		All	jobs	1320000
SAL		MGR		4446000
SAL		WOR		490000
SAL		All	jobs	4936000
SUP		MGR		465000
SUP		TEC		115000
SUP		WOR		435000
SUP		All	jobs	1015000
All	divisions	ENG		245000
All	divisions	MGR		6246000
All	divisions	PRE		800000
	divisions			115000
	divisions			1475000
All	divisions	All	jobs	8881000

Using the GROUPING SETS Clause

You use the GROUPING SETS clause to get just the subtotal rows. The following example uses GROUPING SETS to get the subtotals for salaries by division id and job id:

```
SELECT division_id, job_id, SUM(salary)
     FROM employees2
     GROUP BY GROUPING SETS (division id, job id)
     ORDER BY division_id, job_id;
     DIV JOB SUM(SALARY)
     --- --- -----
     BUS 1610000
OPE 1320000
```

SAL	4936000
SUP	1015000
ENG	245000
MGR	6246000
PRE	800000
TEC	115000
WOR	1475000

Notice that only subtotals for the division_id and job_id columns are returned; the total for all salaries is not returned. You'll see how to get the total as well as the subtotals using the GROUPING ID() function in the next section.



TIP

The GROUPING SETS clause typically offers better performance than CUBE. Therefore, you should use GROUPING SETS rather than CUBE wherever possible.

Using the GROUPING_ID() Function

You can use the <code>GROUPING_ID()</code> function to filter rows using a <code>HAVING</code> clause to exclude rows that don't contain a subtotal or total. The <code>GROUPING_ID()</code> function accepts one or more columns and returns the decimal equivalent of the <code>GROUPING</code> bit vector. The <code>GROUPING</code> bit vector is computed by combining the results of a call to the <code>GROUPING()</code> function for each column in order.

Computing the GROUPING Bit Vector

Earlier in the section "Using the GROUPING() Function," you saw that GROUPING() returns 1 when the column value is null and returns 0 when the column value is non-null; for example:

- If both division_id and job_id are non-null, GROUPING() returns 0 for both columns. The result for division_id is combined with the result for job_id, giving a bit vector of 00, whose decimal equivalent is 0. GROUPING_ID() therefore returns 0 when division id and job id are non-null.
- If division_id is non-null (the GROUPING bit is 0), but job_id is null (the GROUPING bit is 1), the resulting bit vector is 01 and GROUPING ID() returns 1.
- If division_id is null (the GROUPING bit is 1), but job_id is non-null (the GROUPING bit is 0), the resulting bit vector is 10 and GROUPING_ID() returns 2.
- If both division_id and job_id are null (both GROUPING bits are 0), the bit vector is 11 and GROUPING_ID() returns 3.

The following table summarizes these results.

division_id	job_id	Bit Vector	GROUPING_ID() Return Value
non-null	non-null	00	0
non-null	null	01	1
null	non-null	10	2
null	null	11	3

An Example Query That Illustrates the Use of GROUPING_ID()

The following example passes division id and job id to GROUPING ID(); notice that the output from the GROUPING ID() function agrees with the expected returned values documented in the previous section:

SELECT

```
division id, job id,
 GROUPING (division id) AS DIV GRP,
 GROUPING (job id) AS JOB GRP,
 GROUPING ID (division id, job id) AS grp id,
 SUM(salary)
FROM employees2
GROUP BY CUBE (division id, job id)
ORDER BY division id, job id;
```

DIV	JOB	DIV_GRP	JOB_GRP	GRP_ID	SUM (SALARY)
BUS	MGR	0	0	0	530000
BUS	PRE	0	0	0	800000
BUS	WOR	0	0	0	280000
BUS		0	1	1	1610000
OPE	ENG	0	0	0	245000
OPE	MGR	0	0	0	805000
OPE	WOR	0	0	0	270000
OPE		0	1	1	1320000
SAL	MGR	0	0	0	4446000
SAL	WOR	0	0	0	490000
SAL		0	1	1	4936000
SUP	MGR	0	0	0	465000
SUP	TEC	0	0	0	115000
SUP	WOR	0	0	0	435000
SUP		0	1	1	1015000
	ENG	1	0	2	245000
	MGR	1	0	2	6246000
	PRE	1	0	2	800000
	TEC	1	0	2	115000
	WOR	1	0	2	1475000
		1	1	3	8881000

A Useful Application of GROUPING_ID()

One useful application of GROUPING ID() is to filter rows using a HAVING clause. The HAVING clause can exclude rows that don't contain a subtotal or total by simply checking if GROUPING ID() returns a value greater than 0. For example:

```
division_id, job_id,
  GROUPING ID (division id, job id) AS grp id,
  SUM(salary)
FROM employees2
GROUP BY CUBE (division_id, job_id)
```

HAVING GROUPING_ID(division_id, job_id) > 0
ORDER BY division id, job id;

DIV	JOB	GRP_ID	SUM(SALARY)
BUS		1	1610000
OPE		1	1320000
SAL		1	4936000
SUP		1	1015000
	ENG	2	245000
	MGR	2	6246000
	PRE	2	800000
	TEC	2	115000
	WOR	2	1475000
		3	8881000

Using a Column Multiple Times in a GROUP BY Clause

You can use a column many times in a GROUP BY clause. Doing this allows you to reorganize your data or report on different groupings of data. For example, the following query contains a GROUP BY clause that uses division_id twice, once to group by division_id and again in a ROLLUP:

SELECT division_id, job_id, SUM(salary)
FROM employees2
GROUP BY division id, ROLLUP(division id, job id);

DIV	JOB	SUM(SALARY)
BUS	MGR	530000
BUS	PRE	800000
BUS	WOR	280000
OPE	ENG	245000
OPE	MGR	805000
OPE	WOR	270000
SAL	MGR	4446000
SAL	WOR	490000
SUP	MGR	465000
SUP	TEC	115000
SUP	WOR	435000
BUS		1610000
OPE		1320000
SAL		4936000
SUP		1015000
BUS		1610000
OPE		1320000
SAL		4936000
SUP		1015000

Notice, however, that the last four rows are duplicates of the previous four rows. You can eliminate these duplicates using the <code>GROUP_ID()</code> function, which you'll learn about next.

Using the GROUP_ID() Function

You can use the GROUP ID() function to remove duplicate rows returned by a GROUP BY clause. GROUP ID() doesn't accept any parameters. If n duplicates exist for a particular grouping, GROUP ID returns numbers in the range 0 to n-1.

The following example rewrites the query shown in the previous section to include the output from GROUP ID(); notice that GROUP ID() returns 0 for all rows except the last four, which are duplicates of the previous four rows, and that GROUP ID() returns 1:

SELECT division id, job id, GROUP ID(), SUM(salary) FROM employees2 GROUP BY division id, ROLLUP(division id, job id);

DIV	JOB	GROUP_	_ID()	SUM(SALARY)
BUS	MGR		0	530000
BUS	PRE		0	800000
BUS	WOR		0	280000
OPE	ENG		0	245000
OPE	MGR		0	805000
OPE	WOR		0	270000
SAL	MGR		0	4446000
SAL	WOR		0	490000
SUP	MGR		0	465000
SUP	TEC		0	115000
SUP	WOR		0	435000
BUS			0	1610000
OPE			0	1320000
SAL			0	4936000
SUP			0	1015000
BUS			1	1610000
OPE			1	1320000
SAL			1	4936000
SUP			1	1015000

You can eliminate duplicate rows using a HAVING clause that allows only rows whose GROUP ID() is 0; for example:

SELECT division id, job id, GROUP ID(), SUM(salary) FROM employees2 GROUP BY division_id, ROLLUP(division_id, job_id) $HAVING GROUP_ID() = 0;$

DIV	JOB	<pre>GROUP_ID()</pre>	SUM(SALARY)
BUS	MGR	0	530000
BUS	PRE	0	800000
BUS	WOR	0	280000
OPE	ENG	0	245000
OPE	MGR	0	805000
OPE	WOR	0	270000

SAL	MGR	0	4446000
SAL	WOR	0	490000
SUP	MGR	0	465000
SUP	TEC	0	115000
SUP	WOR	0	435000
BUS		0	1610000
OPE		0	1320000
SAL		0	4936000
SUP		0	1015000

This concludes the discussion of the extended GROUP BY clauses.

Using the Analytic Functions

The database has many built-in analytic functions that enable you to perform complex calculations, such as finding the top-selling product type for each month, the top salespersons, and so on. The analytic functions are organized into the following categories:

- **Ranking functions** enable you to calculate ranks, percentiles, and *n*-tiles (tertiles, quartiles, and so on).
- Inverse percentile functions enable you to calculate the value that corresponds to a percentile.
- Window functions enable you to calculate cumulative and moving aggregates.
- **Reporting functions** enable you to calculate things like market share.
- Lag and lead functions enable you to get a value in a row where that row is a certain number of rows away from the current row.
- **First and last functions** enable you to get the first and last values in an ordered group.
- **Linear regression functions** enable you to fit an ordinary-least-squares regression line to a set of number pairs.
- **Hypothetical rank and distribution functions** enable you to calculate the rank and percentile that a new row would have if you inserted it into a table.

You'll learn about these functions shortly, but first let's examine the example table used next.

The Example Table

You'll see the use of the all_sales table in the following sections. The all_sales table stores the sum of all the sales by dollar amount for a particular year, month, product type, and employee. The all_sales table is created by the store_schema.sql script as follows:

```
CREATE TABLE all_sales (
year INTEGER NOT NULL,
month INTEGER NOT NULL,
prd_type_id INTEGER

CONSTRAINT all_sales_fk_product_types
REFERENCES product_types(product_type_id),
emp_id_INTEGER
```

```
CONSTRAINT all sales fk employees2
   REFERENCES employees2 (employee id),
  amount NUMBER(8, 2),
  CONSTRAINT all sales pk PRIMARY KEY (
    year, month, prd type id, emp id
  )
);
```

As you can see, the all sales table contains five columns, which are as follows:

- **YEAR** stores the year the sales took place.
- **MONTH** stores the month the sales took place (1 to 12).
- PRD TYPE ID stores the product type id of the product.
- **EMP** ID stores the employee id of the employee who handled the sales.
- **AMOUNT** stores the total dollar amount of the sales.

The following query retrieves the first 12 rows from the all sales table:

```
SELECT *
    FROM all sales
    WHERE ROWNUM <= 12;
```

YEAR	MONTH	PRD_TYPE_ID	EMP_ID	AMOUNT
2003	1	1	21	10034.84
2003	2	1	21	15144.65
2003	3	1	21	20137.83
2003	4	1	21	25057.45
2003	5	1	21	17214.56
2003	6	1	21	15564.64
2003	7	1	21	12654.84
2003	8	1	21	17434.82
2003	9	1	21	19854.57
2003	10	1	21	21754.19
2003	11	1	21	13029.73
2003	12	1	21	10034.84



The all sales table actually contains a lot more rows than this, but for space considerations I've omitted listing them all here.

Let's examine the ranking functions next.

Using the Ranking Functions

You use the ranking functions to calculate ranks, percentiles, and *n*-tiles. The ranking functions are shown in Table 7-2.

Let's examine the RANK () and DENSE RANK () functions first.

Function	Description
RANK()	Returns the rank of items in a group. RANK() leaves a gap in the sequence of rankings in the event of a tie.
DENSE_RANK()	Returns the rank of items in a group. DENSE_RANK() doesn't leave a gap in the sequence of rankings in the event of a tie.
CUME_DIST()	Returns the position of a specified value relative to a group of values. CUME_DIST() is short for cumulative distribution.
PERCENT_RANK()	Returns the percent rank of a value relative to a group of values.
NTILE()	Returns <i>n</i> -tiles: tertiles, quartiles, and so on.
ROW_NUMBER()	Returns a number with each row in a group.

TABLE 7-2 The Ranking Functions

Using the RANK() and DENSE_RANK() Functions

You use RANK() and DENSE_RANK() to rank items in a group. The difference between these two functions is in the way they handle items that tie: RANK() leaves a gap in the sequence when there is a tie, but DENSE_RANK() leaves no gaps. For example, if you were ranking sales by product type and two product types tie for first place, RANK() would put the two product types in first place, but the next product type would be in third place. DENSE_RANK() would also put the two product types in first place, but the next product type would be in second place.

The following query illustrates the use of RANK () and DENSE_RANK () to get the ranking of sales by product type for the year 2003; notice the use of the keyword OVER in the syntax when calling the RANK () and DENSE_RANK () functions:

```
SELECT
```

```
prd_type_id, SUM(amount),
  RANK() OVER (ORDER BY SUM(amount) DESC) AS rank,
  DENSE_RANK() OVER (ORDER BY SUM(amount) DESC) AS dense_rank
FROM all_sales
WHERE year = 2003
AND amount IS NOT NULL
GROUP BY prd_type_id
ORDER BY prd_type_id;
```

PRD_TYPE_ID	SUM (AMOUNT)	RANK	DENSE_RANK
1	905081.84	1	1
2	186381.22	4	4
3	478270.91	2	2
4	402751.16	3	3

Notice that sales for product type #1 are ranked first, sales for product type #2 are ranked fourth, and so on. Because there are no ties, RANK() and DENSE RANK() return the same ranks.

The all sales table actually contains nulls in the AMOUNT column for all rows whose PRD TYPE ID column is 5; the previous query omits these rows because of the inclusion of the line "AND amount IS NOT NULL" in the WHERE clause. The next example includes these rows by leaving out the AND line from the WHERE clause:

SELECT

```
prd type id, SUM(amount),
 RANK() OVER (ORDER BY SUM(amount) DESC) AS rank,
 DENSE RANK() OVER (ORDER BY SUM(amount) DESC) AS dense rank
FROM all sales
WHERE year = 2003
GROUP BY prd type id
ORDER BY prd type id;
```

PRD_TYPE_ID	SUM (AMOUNT)	RANK	DENSE_RANK
1	905081.84	2	2
2	186381.22	5	5
3	478270.91	3	3
4	402751.16	4	4
5		1	1

Notice that the last row contains null for the sum of the AMOUNT column and that RANK () and DENSE RANK () return 1 for this row. This is because by default RANK () and DENSE RANK () assign the highest rank of 1 to null values in descending rankings (that is, DESC is used in the OVER clause) and the lowest rank in ascending rankings (that is, ASC is used in the OVER clause).

Controlling Ranking of Null Values Using the NULLS FIRST and NULLS LAST Clauses When using an analytic function, you can explicitly control whether nulls are the highest or lowest in a group using NULLS FIRST or NULLS LAST. The following example uses NULLS LAST to specify that nulls are the lowest:

```
prd type id, SUM(amount),
 RANK() OVER (ORDER BY SUM(amount) DESC NULLS LAST) AS rank,
 DENSE RANK() OVER (ORDER BY SUM(amount) DESC NULLS LAST) AS dense rank
FROM all sales
WHERE year = 2003
GROUP BY prd_type_id
ORDER BY prd type id;
```

PRD_TYPE_ID	SUM (AMOUNT)	RANK	DENSE_RANK
1	905081.84	1	1
2	186381.22	4	4
3	478270.91	2	2
4	402751.16	3	3
5		5	5

Using the PARTITION BY Clause with Analytic Functions You use the PARTITION BY clause with the analytic functions when you need to divide the groups into subgroups. For example, if you need to subdivide the sales amount by month, you can use PARTITION BY month, as shown in the following query:

```
prd_type_id, month, SUM(amount),
  RANK() OVER (PARTITION BY month ORDER BY SUM(amount) DESC) AS rank
FROM all_sales
WHERE year = 2003
AND amount IS NOT NULL
GROUP BY prd_type_id, month
ORDER BY prd_type_id, month;
```

PRD_TYPE_ID	MONTH	SUM (AMOUNT)	RANK
1	1	38909.04	1
1	2	70567.9	1
1	3	91826.98	1
1	4	120344.7	1
1	5	97287.36	1
1	6	57387.84	1
1	7	60929.04	2
1	8	75608.92	1
1	9	85027.42	1
1	10	105305.22	1
1	11	55678.38	1
1	12	46209.04	2
2	1	14309.04	4
2	2	13367.9	4
2	3	16826.98	4
2	4	15664.7	4
2	5	18287.36	4
2	6	14587.84	4
2	7	15689.04	3
2	8	16308.92	4
2	9	19127.42	4
2	10	13525.14	4
2	11	16177.84	4
2	12	12509.04	4
3	1	24909.04	2
3	2	15467.9	3
3	3	20626.98	3
3	4	23844.7	2
3	5	18687.36	3
3	6	19887.84	3
3	7	81589.04	1
3	8	62408.92	2
3	9	46127.42	3
3	10	70325.29	3
3	11	46187.38	2
3	12	48209.04	1

4	1	17398.43	3
4	2	17267.9	2
4	3	31026.98	2
4	4	16144.7	3
4	5	20087.36	2
4	6	33087.84	2
4	7	12089.04	4
4	8	58408.92	3
4	9	49327.42	2
4	10	75325.14	2
4	11	42178.38	3
4	12	30409.05	3

Using ROLLUP, CUBE, and GROUPING SETS Operators with Analytic Functions You can use the ROLLUP, CUBE, and GROUPING SETS operators with the analytic functions. The following query uses ROLLUP and RANK() to get the sales rankings by product type ID:

SELECT

```
prd type id, SUM(amount),
 RANK() OVER (ORDER BY SUM(amount) DESC) AS rank
FROM all_sales
WHERE year = 2003
GROUP BY ROLLUP (prd type id)
ORDER BY prd type id;
```

PRD_TYPE_ID	SUM (AMOUNT)	RANK	
1	905081.84	3	
2	186381.22	6	
3	478270.91	4	
4	402751.16	5	
5		1	
	1972485.13	2	

The next query uses CUBE and RANK () to get all rankings of sales by product type ID and employee ID:

```
prd type id, emp id, SUM(amount),
 RANK() OVER (ORDER BY SUM(amount) DESC) AS rank
FROM all_sales
WHERE year = 2003
GROUP BY CUBE (prd type id, emp id)
ORDER BY prd_type_id, emp_id;
```

EMP_ID	SUM (AMOUNT)	RANK
21	197916.96	19
22	214216.96	17
23	98896.96	26
24	207216.96	18
25	93416.96	28
26	93417.04	27
	21 22 23 24 25	22 214216.96 23 98896.96 24 207216.96 25 93416.96

```
1
            905081.84
        21
            20426.96
                          40
       22
            19826.96
                         41
2
       23
            19726.96
                         42
2
      24 43866.96
                         34
2
      25 32266.96
                         38
      26 50266.42
2
                          31
2
           186381.22
                         21
3
      21 140326.96
                         22
3
      22 116826.96
                         23
3
      23 112026.96
                         2.4
3
          34829.96
      24
                          36
3
      25 29129.96
                         39
3
       26
            45130.11
                         33
3
           478270.91
                         10
      21 108326.96
                         2.5
4
          81426.96
      22
                          30
                         29
4
      23 92426.96
       24 47456.96
                         32
4
4
      25 33156.96
                         37
      26 39956.36
                         35
           402751.16
                         13
5
      21
                          1
5
       22
                          1
5
       23
                           1
5
       24
                          1
5
       25
                           1
5
                          1
       26
5
                          1
        21
          466997.84
                         11
        22
          432297.84
                         12
        23 323077.84
                         15
        24 333370.84
                         14
        25 187970.84
                         20
        26 228769.93
                         16
           1972485.13
                          8
```

The next query uses ${\tt GROUPING}$ SETS and RANK() to get just the sales amount subtotal rankings:

```
prd_type_id, emp_id, SUM(amount),
  RANK() OVER (ORDER BY SUM(amount) DESC) AS rank
FROM all_sales
WHERE year = 2003
GROUP BY GROUPING SETS(prd_type_id, emp_id)
ORDER BY prd_type_id, emp_id;
```

RANK	SUM (AMOUNT)	EMP_ID	PRD_TYPE_ID
2	905081.84		1
11	186381.22		2

3		478270.91	3
4		402751.16	6
5			1
	21	466997.84	4
	22	432297.84	5
	23	323077.84	8
	24	333370.84	7
	25	187970.84	10
	26	228769.93	9

Using the CUME_DIST() and PERCENT_RANK() Functions

You use CUME DIST() to calculate the position of a specified value relative to a group of values; CUME DIST() is short for cumulative distribution. You use PERCENT RANK() to calculate the percent rank of a value relative to a group of values.

The following query illustrates the use of CUME DIST() and PERCENT RANK() to get the cumulative distribution and percent rank of sales:

```
SELECT
```

```
prd type id, SUM(amount),
 CUME DIST() OVER (ORDER BY SUM(amount) DESC) AS cume_dist,
 PERCENT RANK() OVER (ORDER BY SUM(amount) DESC) AS percent rank
FROM all sales
WHERE year = 2003
GROUP BY prd type id
ORDER BY prd type id;
DDD TVDF ID CHM/AMOHNT) CHMF DICT DEDCENT DANK
```

PRD_LABE_ID	SUM (AMOUNT)	COME_DIST	PERCENT_RANK
1	905081.84	. 4	.25
2	186381.22	1	1
3	478270.91	.6	.5
4	402751.16	.8	.75
5		.2	0

Using the NTILE() Function

You use NTILE (buckets) to calculate n-tiles (tertiles, quartiles, and so on); buckets specifies the number of "buckets" into which groups of rows are placed. For example, NTILE (2) specifies two buckets and therefore divides the rows into two groups of rows; NTILE (4) divides the groups into four buckets and therefore divides the rows into four groups.

The following query illustrates the use of NTILE(); notice that 4 is passed to NTILE() to split the groups of rows into four buckets:

```
prd_type_id, SUM(amount),
 NTILE (4) OVER (ORDER BY SUM(amount) DESC) AS ntile
FROM all sales
WHERE year = 2003
AND amount IS NOT NULL
GROUP BY prd type id
ORDER BY prd type id;
```

PRD_TYPE_ID	SUM (AMOUNT)	NTILE
1	905081.84	1
2	186381.22	4
3	478270.91	2
4	402751.16	3

Using the ROW_NUMBER() Function

You use ROW_NUMBER() to return a number with each row in a group, starting at 1. The following query illustrates the use of ROW NUMBER():

This concludes the discussion of ranking functions.

Using the Inverse Percentile Functions

In the section "Using the CUME_DIST() and PERCENT_RANK() Functions," you saw that CUME_DIST() is used to calculate the position of a specified value relative to a group of values. You also saw that PERCENT_RANK() is used to calculate the percent rank of a value relative to a group of values.

In this section, you'll see how to use the inverse percentile functions to get the value that corresponds to a percentile. There are two inverse percentile functions: $PERCENTILE_DISC(x)$ and $PERCENTILE_CONT(x)$. They operate in a manner the reverse of $CUME_DIST()$ and $PERCENT_RANK()$. $PERCENTILE_DISC(x)$ examines the cumulative distribution values in each group until it finds one that is greater than or equal to x. $PERCENTILE_CONT(x)$ examines the percent rank values in each group until it finds one that is greater than or equal to x.

The following query illustrates the use of PERCENTILE_CONT() and PERCENTILE_DISC() to get the sum of the amount whose percentile is greater than or equal to 0.6:

```
PERCENTILE_CONT(0.6) WITHIN GROUP (ORDER BY SUM(amount) DESC)
   AS percentile_cont,
   PERCENTILE_DISC(0.6) WITHIN GROUP (ORDER BY SUM(amount) DESC)
   AS percentile_disc
FROM all_sales
WHERE year = 2003
GROUP BY prd_type_id;
```

```
PERCENTILE CONT PERCENTILE DISC
 -----
    417855.11
               402751.16
```

If you compare the sum of the amounts shown in these results with those shown in the earlier section "Using the CUME_DIST() and PERCENT_RANK() Functions," you'll see that the sums correspond to those whose cumulative distribution and percent rank are 0.6 and 0.75, respectively.

Using the Window Functions

You use the window functions to calculate things like cumulative sums and moving averages within a specified range of rows, a range of values, or an interval of time. As you know, a query returns a set of rows known as the result set. The term "window" is used to describe a subset of rows within the result set. The subset of rows "seen" through the window is then processed by the window functions, which return a value. You can define the start and end of the window.

You can use a window with the following functions: SUM(), AVG(), MAX(), MIN(), COUNT(), VARIANCE (), and STDDEV(); you saw these functions in Chapter 4. You can also use a window with FIRST VALUE() and LAST VALUE(), which return the first and last values in a window. (You'll learn more about the FIRST VALUE() and LAST VALUE() functions later in the section "Getting the First and Last Rows Using FIRST_VALUE() and LAST_VALUE().")

In the next section, you'll see how to perform a cumulative sum, a moving average, and a centered average.

Performing a Cumulative Sum

The following query performs a cumulative sum to compute the cumulative sales amount for 2003, starting with January and ending in December; notice that each monthly sales amount is added to the cumulative amount that grows after each month:

```
month, SUM (amount) AS month amount,
  SUM(SUM(amount)) OVER
    (ORDER BY month ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW)
    AS cumulative amount
FROM all sales
WHERE year = 2003
GROUP BY month
ORDER BY month;
```

MONTH	MONTH_AMOUNT	CUMULATIVE_AMOUNT
1	95525.55	95525.55
2	116671.6	212197.15
3	160307.92	372505.07
4	175998.8	548503.87
5	154349.44	702853.31
6	124951.36	827804.67
7	170296.16	998100.83
8	212735.68	1210836.51
9	199609.68	1410446.19
10	264480.79	1674926.98
11	160221.98	1835148.96
12	137336.17	1972485.13

This query uses the following expression to compute the cumulative aggregate:

```
SUM(SUM(amount)) OVER

(ORDER BY month ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW)

AS cumulative amount
```

Let's break down this expression:

- SUM(amount) computes the sum of an amount. The outer SUM() computes the cumulative amount.
- ORDER BY month orders the rows read by the query by month.
- ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW defines the start and end of the window. The start is set to UNBOUNDED PRECEDING, which means the start of the window is fixed at the first row in the result set returned by the query. The end of the window is set to CURRENT ROW; CURRENT ROW represents the current row in the result set being processed, and the end of the window slides down one row after the outer SUM() function computes and returns the current cumulative amount.

The entire query computes and returns the cumulative total of the sales amounts, starting at month 1, and then adding the sales amount for month 2, then month 3, and so on, up to and including month 12. The start of the window is fixed at month 1, but the bottom of the window moves down one row in the result set after each month's sales amounts are added to the cumulative total. This continues until the last row in the result set is processed by the window and the SUM() functions.

Don't confuse the end of the window with the end of the result set. In the previous example, the end of the window slides down one row in the result set as each row is processed (i.e., the sum of the sales amount for that month is added to the cumulative total). In the example, the end of the window starts at the first row, the sum sales amount for that month is added to the cumulative total, and then the end of the window moves down one row to the second row. At this point, the window sees two rows. The sum of the sales amount for that month is added to the cumulative total, and the end of the window moves down one row to the third row. At this point, the window sees three rows. This continues until the twelfth row is processed. At this point, the window sees twelve rows.

The following query uses a cumulative sum to compute the cumulative sales amount, starting with June of 2003 (month 6) and ending in December of 2003 (month 12):

```
SELECT
```

```
month, SUM(amount) AS month_amount,
SUM(SUM(amount)) OVER
   (ORDER BY month ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS
   cumulative_amount
FROM all_sales
WHERE year = 2003
AND month BETWEEN 6 AND 12
GROUP BY month
ORDER BY month;
```

```
MONTH MONTH AMOUNT CUMULATIVE AMOUNT

6 124951.36 124951.36
7 170296.16 295247.52
```

8	212735.68	507983.2
9	199609.68	707592.88
10	264480.79	972073.67
11	160221.98	1132295.65
12	137336.17	1269631.82

Performing a Moving Average

The following query computes the moving average of the sales amount between the current month and the previous three months:

SELECT

```
month, SUM (amount) AS month amount,
 AVG(SUM(amount)) OVER
    (ORDER BY month ROWS BETWEEN 3 PRECEDING AND CURRENT ROW)
   AS moving average
FROM all sales
WHERE year = 2003
GROUP BY month
ORDER BY month:
```

MONTH	MONTH_AMOUNT	MOVING_AVERAGE
1	95525.55	95525.55
2	116671.6	106098.575
3	160307.92	124168.357
4	175998.8	137125.968
5	154349.44	151831.94
6	124951.36	153901.88
7	170296.16	156398.94
8	212735.68	165583.16
9	199609.68	176898.22
10	264480.79	211780.578
11	160221.98	209262.033
12	137336.17	190412.155

Notice that the query uses the following expression to compute the moving average:

```
AVG(SUM(amount)) OVER
       (ORDER BY month ROWS BETWEEN 3 PRECEDING AND CURRENT ROW)
      AS moving average
```

Let's break down this expression:

- SUM (amount) computes the sum of an amount. The outer AVG () computes the average.
- ORDER BY month orders the rows read by the query by month.
- ROWS BETWEEN 3 PRECEDING AND CURRENT ROW defines the start of the window as including the three rows preceding the current row; the end of the window is the current row being processed.

So, the entire expression computes the moving average of the sales amount between the current month and the previous three months. Because for the first two months less than the full three months of data are available, the moving average is based on only the months available.

Both the start and the end of the window begin at row #1 read by the query. The end of the window moves down after each row is processed. The start of the window moves down only after row #4 has been processed, and subsequently moves down one row after each row is processed. This continues until the last row in the result set is read.

Performing a Centered Average

The following query computes the moving average of the sales amount centered between the previous and next month from the current month:

SELECT

```
month, SUM (amount) AS month amount,
 AVG(SUM(amount)) OVER
    (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
    AS moving average
FROM all sales
WHERE year = 2003
GROUP BY month
ORDER BY month;
```

MONTH MONTH AMOUNT MOVING AVERAGE -----1 95525.55 106098.575 2 116671.6 124168.357 3 160307.92 150992.773 4 175998.8 163552.053 5 154349.44 151766.533 6 124951.36 149865.653 7 170296.16 169327.733 8 212735.68 194213.84 9 199609.68 225608.717 10 264480.79 208104.15 10 264480.79 208104.15 160221.98 187346.313 137336.17 148779.075 11 160221.98

Notice that the query uses the following expression to compute the moving average:

```
AVG(SUM(amount)) OVER
       (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
       AS moving average
```

Let's break down this expression:

12

- SUM (amount) computes the sum of an amount. The outer AVG () computes the average.
- ORDER BY month orders the rows read by the query by month.
- ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING defines the start of the window as including the row preceding the current row being processed. The end of the window is the row following the current row.

So, the entire expression computes the moving average of the sales amount between the current month and the previous month. Because for the first and last month less than the full three months of data are available, the moving average is based on only the months available.

The start of the window begins at row #1 read by the query. The end of the window begins at row #2 and moves down after each row is processed. The start of the window moves down only once row #2 has been processed. Processing continues until the last row read by the query is processed.

Getting the First and Last Rows Using FIRST VALUE() and LAST VALUE()

You use the FIRST VALUE() and LAST VALUE() functions to get the first and last rows in a window. The following query uses FIRST VALUE() and LAST VALUE() to get the previous and next month's sales amount:

SELECT

```
month, SUM(amount) AS month_amount,
 FIRST VALUE (SUM (amount)) OVER
    (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
    AS previous month amount,
 LAST VALUE (SUM (amount)) OVER
    (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
    AS next month amount
FROM all sales
WHERE year = 2003
GROUP BY month
ORDER BY month;
```

MONTH	MONTH_AMOUNT	PREVIOUS_MONTH_AMOUNT	NEXT_MONTH_AMOUNT
1	95525.55	95525.55	116671.6
2	116671.6	95525.55	160307.92
3	160307.92	116671.6	175998.8
4	175998.8	160307.92	154349.44
5	154349.44	175998.8	124951.36
6	124951.36	154349.44	170296.16
7	170296.16	124951.36	212735.68
8	212735.68	170296.16	199609.68
9	199609.68	212735.68	264480.79
10	264480.79	199609.68	160221.98
11	160221.98	264480.79	137336.17
12	137336.17	160221.98	137336.17

The next query divides the current month's sales amount by the previous month's sales amount (labeled as curr div prev) and also divides the current month's sales amount by the next month's sales amount (labeled as curr div next):

```
month, SUM(amount) AS month amount,
  SUM(amount)/FIRST VALUE(SUM(amount)) OVER
    (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
    AS curr div prev,
  SUM(amount)/LAST VALUE(SUM(amount)) OVER
    (ORDER BY month ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING)
    AS curr div next
FROM all sales
WHERE year = 2003
```

```
GROUP BY month ORDER BY month;
```

MONTH	MONTH_AMOUNT	CURR_DIV_PREV	CURR_DIV_NEXT
1	95525.55	1	.818755807
2	116671.6	1.22136538	.727796855
3	160307.92	1.37400978	.910846665
4	175998.8	1.09787963	1.14026199
5	154349.44	.876991434	1.23527619
6	124951.36	.809535558	.733729756
7	170296.16	1.36289961	.800505867
8	212735.68	1.24921008	1.06575833
9	199609.68	.93829902	.754722791
10	264480.79	1.3249898	1.65071478
11	160221.98	.605798175	1.16664081
12	137336.17	.857161858	1

This concludes the discussion of window functions.

Using the Reporting Functions

You use the reporting functions to perform calculations across groups and partitions within groups. You can perform reporting with the following functions: SUM(), AVG(), MAX(), MIN(), COUNT(), VARIANCE(), and STDDEV(). You can also use the RATIO_TO_REPORT() function to compute the ratio of a value to the sum of a set of values.

In this section, you'll see how to perform a report on a sum and use the RATIO_TO_ REPORT() function.

Reporting on a Sum

For the first three months of 2003, the following query reports

- The total sum of all sales for all three months (labeled as total month amount).
- The total sum of all sales for all product types (labeled as total_product_type_amount).

```
month, prd_type_id,
SUM(SUM(amount)) OVER (PARTITION BY month)
AS total_month_amount,
SUM(SUM(amount)) OVER (PARTITION BY prd_type_id)
AS total_product_type_amount
FROM all_sales
WHERE year = 2003
AND month <= 3
GROUP BY month, prd_type_id
ORDER BY month, prd type id;</pre>
```

```
MONTH PRD_TYPE_ID TOTAL_MONTH_AMOUNT TOTAL_PRODUCT_TYPE_AMOUNT

1 1 95525.55 201303.92
1 2 95525.55 44503.92
```

1	3	95525.55	61003.92
1	4	95525.55	65693.31
1	5	95525.55	
2	1	116671.6	201303.92
2	2	116671.6	44503.92
2	3	116671.6	61003.92
2	4	116671.6	65693.31
2	5	116671.6	
3	1	160307.92	201303.92
3	2	160307.92	44503.92
3	3	160307.92	61003.92
3	4	160307.92	65693.31
3	5	160307.92	

Notice that the query uses the following expression to report the total sum of all sales for all months (labeled as total month amount):

```
SUM(SUM(amount)) OVER (PARTITION BY month)
      AS total month amount
```

Let's break down this expression:

- SUM (amount) computes the sum of an amount. The outer SUM() computes the total sum.
- OVER (PARTITION BY month) causes the outer SUM() to compute the sum for each month.

The previous query also uses the following expression to report the total sum of all sales for all product types (labeled as total product type amount):

```
SUM(SUM(amount)) OVER (PARTITION BY prd type id)
      AS total product type amount
```

Let's break down this expression:

- SUM(amount) computes the sum of an amount. The outer SUM() computes the total sum.
- OVER (PARTITION BY prd type id) causes the outer SUM() to compute the sum for each product type.

Using the RATIO TO REPORT() Function

You use the RATIO TO REPORT () function to compute the ratio of a value to the sum of a set of values.

For the first three months of 2003, the following query reports

- The sum of the sales amount by product type for each month (labeled as prd type amount).
- The ratio of the product type's sales amount to the entire month's sales (labeled as prd type ratio), which is computed using RATIO TO REPORT().

```
month, prd_type_id,
SUM(amount) AS prd_type_amount,
RATIO TO REPORT (SUM(amount)) OVER (PARTITION BY month) AS prd type ratio
```

```
FROM all_sales
WHERE year = 2003
AND month <= 3
GROUP BY month, prd_type_id
ORDER BY month, prd_type_id;
```

MONTH	PRD_TYPE_ID	PRD_TYPE_AMOUNT	PRD_TYPE_RATIO
1	1	38909.04	.40731553
1	2	14309.04	.149792804
1	3	24909.04	.260757881
1	4	17398.43	.182133785
1	5		
2	1	70567.9	.604842138
2	2	13367.9	.114577155
2	3	15467.9	.132576394
2	4	17267.9	.148004313
2	5		
3	1	91826.98	.57281624
3	2	16826.98	.104966617
3	3	20626.98	.128670998
3	4	31026.98	.193546145
3	5		

Notice that the query uses the following expression to compute the ratio (labeled as prd_type ratio):

RATIO TO REPORT(SUM(amount)) OVER (PARTITION BY month) AS prd type ratio

Let's break down this expression:

- SUM(amount) computes the sum of the sales amount.
- OVER (PARTITION BY month) causes the outer SUM() to compute the sum of the sales amount for each month.
- The ratio is computed by dividing the sum of the sales amount for each product type by the sum of the entire month's sales amount.

This concludes the discussion of reporting functions.

Using the LAG() and LEAD() Functions

You use the LAG() and LEAD() functions to get a value in a row where that row is a certain number of rows away from the current row. The following query uses LAG() and LEAD() to get the previous and next month's sales amount:

```
month, SUM(amount) AS month_amount,
  LAG(SUM(amount), 1) OVER (ORDER BY month) AS previous_month_amount,
  LEAD(SUM(amount), 1) OVER (ORDER BY month) AS next_month_amount
FROM all_sales
WHERE year = 2003
```

```
GROUP BY month
ORDER BY month;
```

MONTH	MONTH_AMOUNT	PREVIOUS_MONTH_AMOUNT	NEXT_MONTH_AMOUNT
1	95525.55		116671.6
2	116671.6	95525.55	160307.92
3	160307.92	116671.6	175998.8
4	175998.8	160307.92	154349.44
5	154349.44	175998.8	124951.36
6	124951.36	154349.44	170296.16
7	170296.16	124951.36	212735.68
8	212735.68	170296.16	199609.68
9	199609.68	212735.68	264480.79
10	264480.79	199609.68	160221.98
11	160221.98	264480.79	137336.17
12	137336.17	160221.98	

Notice that the query uses the following expressions to get the previous and next month's sales:

```
LAG(SUM(amount), 1) OVER (ORDER BY month) AS previous_month_amount,
    LEAD(SUM(amount), 1) OVER (ORDER BY month) AS next month amount
```

LAG(SUM(amount), 1) gets the previous row's sum of the amount. LEAD(SUM(amount), 1) gets the next row's sum of the amount.

Using the FIRST and LAST Functions

You use the FIRST and LAST functions to get the first and last values in an ordered group. You can use FIRST and LAST with the following functions: MIN(), MAX(), COUNT(), SUM(), AVG(), STDDEV(), and VARIANCE().

The following query uses FIRST and LAST to get the months in 2003 that had the highest and lowest sales:

SELECT

```
MIN (month) KEEP (DENSE RANK FIRST ORDER BY SUM (amount))
   AS highest sales month,
 MIN(month) KEEP (DENSE RANK LAST ORDER BY SUM(amount))
   AS lowest sales month
FROM all sales
WHERE year = 2003
GROUP BY month
ORDER BY month;
HIGHEST SALES MONTH LOWEST SALES MONTH
-----
```

Using the Linear Regression Functions

You use the linear regression functions to fit an ordinary-least-squares regression line to a set of number pairs. You can use the linear regression functions as aggregate, windowing, or reporting functions. The following table shows the linear regression functions. In the function syntax, y is interpreted by the functions as a variable that depends on x.

Function	Description
REGR_AVGX(y, x)	Returns the average of x after eliminating x and y pairs where either x or y is null
REGR_AVGY(y, x)	Returns the average of y after eliminating x and y pairs where either x or y is null
REGR_COUNT(y, x)	Returns the number of non-null number pairs that are used to fit the regression line
REGR_INTERCEPT (y, x)	Returns the intercept on the y-axis of the regression line
REGR_R2(y, x)	Returns the coefficient of determination (R-squared) of the regression line
REGR_SLOPE (y, x)	Returns the slope of the regression line
REGR_SXX(y , x)	Returns REG_COUNT (y, x) * VAR_POP(x)
REGR_SXY(y , x)	Returns REG_COUNT $(y, x) * COVAR_POP(y, x)$
$REGR_SYY(y, x)$	Returns REG_COUNT (y, x) * VAR_POP (y)

The following query shows the use of the linear regression functions:

```
SELECT
```

```
prd_type_id,
REGR_AVGX(amount, month) AS avgx,
REGR_AVGY(amount, month) AS avgy,
REGR_COUNT(amount, month) AS count,
REGR_INTERCEPT(amount, month) AS inter,
REGR_R2(amount, month) AS r2,
REGR_SLOPE(amount, month) AS slope,
REGR_SXX(amount, month) AS sxx,
REGR_SXY(amount, month) AS sxx,
REGR_SYY(amount, month) AS sxy,
REGR_SYY(amount, month) AS syy
FROM all_sales
WHERE year = 2003
GROUP BY prd_type_id;
```

PRD_TYPE_ID	AVGX	AVGY	COUNT	INTER	R2
SLOPE	SXX	SXY	SYY		
1 -115.05741	6.5 12570 858 -9871	.5811 9.26 30319		13318.4543	.003746289
2 -2.997634	6.5 2588. 858 -257	62806 1.97 1517		2608.11268	.0000508
3	6 5 6642	65153	72	2154 23119	126338815

```
690.526206 858 592471.485 3238253324
              6.5 5593.76611 72 2043.47164 .128930297
            858 468638.87 1985337488
546.199149
```

Using the Hypothetical Rank and Distribution Functions

You use the hypothetical rank and distribution functions to calculate the rank and percentile that a new row would have if you inserted it into a table. You can perform hypothetical calculations with the following functions: RANK(), DENSE RANK(), PERCENT RANK(), and CUME DIST().

An example of a hypothetical function will be given after the following query, which uses RANK() and PERCENT RANK() to get the rank and percent rank of sales by product type for 2003:

SELECT

```
prd type id, SUM(amount),
 RANK() OVER (ORDER BY SUM(amount) DESC) AS rank,
  PERCENT RANK() OVER (ORDER BY SUM(amount) DESC) AS percent rank
FROM all sales
WHERE year = 2003
AND amount IS NOT NULL
GROUP BY prd type id
ORDER BY prd type id;
```

SUM (AMOUNT)	RANK	PERCENT_RANK
905081.84	1	0
186381.22	4	1
478270.91	2	.333333333
402751.16	3	.666666667
	905081.84 186381.22 478270.91	905081.84 1 186381.22 4 478270.91 2

The next query shows the hypothetical rank and percent rank of a sales amount of \$500,000:

SELECT

```
RANK(500000) WITHIN GROUP (ORDER BY SUM(amount) DESC)
   AS rank.
 PERCENT RANK (500000) WITHIN GROUP (ORDER BY SUM (amount) DESC)
   AS percent rank
FROM all sales
WHERE year = 2003
AND amount IS NOT NULL
GROUP BY prd type id
ORDER BY prd type id;
    RANK PERCENT RANK
_____
```

As you can see, the hypothetical rank and percent rank of a sales amount of \$500,000 are 2 and .25.

This concludes the discussion of hypothetical functions.

Using the MODEL Clause

The MODEL clause was introduced with Oracle Database 10g and enables you to perform interrow calculations. The MODEL clause allows you to access a column in a row like a cell in an array. This gives you the ability to perform calculations in a similar manner to spreadsheet calculations. For example, the all_sales table contains sales information for the months in 2003. You can use the MODEL clause to calculate sales in future months based on sales in 2003.

An Example of the MODEL Clause

The easiest way to learn how to use the MODEL clause is to see an example. The following query retrieves the sales amount for each month in 2003 made by employee #21 for product types #1 and #2 and computes the predicted sales for January, February, and March of 2004 based on sales in 2003:

```
SELECT prd_type_id, year, month, sales_amount

FROM all_sales

WHERE prd_type_id BETWEEN 1 AND 2

AND emp_id = 21

MODEL

PARTITION BY (prd_type_id)

DIMENSION BY (month, year)

MEASURES (amount sales_amount) (

sales_amount[1, 2004] = sales_amount[1, 2003],

sales_amount[2, 2004] = sales_amount[2, 2003] + sales_amount[3, 2003],

sales_amount[3, 2004] = ROUND(sales_amount[3, 2003] * 1.25, 2)

)

ORDER BY prd_type_id, year, month;
```

Let's break down this query:

- PARTITION BY (prd_type_id) specifies that the results are partitioned by prd_type_id.
- DIMENSION BY (month, year) specifies that the dimensions of the array are month and year. This means that a cell in the array is accessed by specifying a month and year.
- MEASURES (amount sales_amount) specifies that each cell in the array contains an amount and that the array name is sales_amount. To access the cell in the sales_amount array for January 2003, you use sales_amount[1, 2003], which returns the sales amount for that month and year.
- After MEASURES come three lines that compute the future sales for January, February, and March of 2004:
 - sales_amount[1, 2004] = sales_amount[1, 2003] sets the sales amount for January 2004 to the amount for January 2003.
 - sales_amount[2, 2004] = sales_amount[2, 2003] + sales_ amount[3, 2003] sets the sales amount for February 2004 to the amount for February 2003 plus March 2003.

- \blacksquare sales amount[3, 2004] = ROUND(sales amount[3, 2003] * 1.25, 2) sets the sales amount for March 2004 to the rounded value of the sales amount for March 2003 multiplied by 1.25.
- ORDER BY prd type id, year, month simply orders the results returned by the entire query.

The output from the query is shown in the following listing; notice that the results contain the sales amounts for all months in 2003 for product types #1 and #2, plus the predicted sales amounts for the first three months in 2004 (which I've made bold to make them stand out):

PRD_TYPE_ID	YEAR	MONTH	SALES_AMOUNT
1	2003	1	10034.84
1	2003	2	15144.65
1	2003	3	20137.83
1	2003	4	25057.45
1	2003	5	17214.56
1	2003	6	15564.64
1	2003	7	12654.84
1	2003	8	17434.82
1	2003	9	19854.57
1	2003	10	21754.19
1	2003	11	13029.73
1	2003	12	10034.84
1	2004	1	10034.84
1	2004	2	35282.48
1	2004	3	25172.29
2	2003	1	1034.84
2	2003	2	1544.65
2	2003	3	2037.83
2	2003	4	2557.45
2	2003	5	1714.56
2	2003	6	1564.64
2	2003	7	1264.84
2	2003	8	1734.82
2	2003	9	1854.57
2	2003	10	2754.19
2	2003	11	1329.73
2	2003	12	1034.84
2	2004	1	1034.84
2	2004	2	3582.48
2	2004	3	2547.29

Using Positional and Symbolic Notation to Access Cells

In the previous example, you saw how to access a cell in an array using the following notation: sales amount [1, 2004], where 1 is the month and 2004 is the year. This is referred to as positional notation because the meaning of the dimensions is determined by their position: the first position contains the month and the second position contains the year.

You can also use symbolic notation to explicitly indicate the meaning of the dimensions, as in, for example, sales_amount[month=1, year=2004]. The following query rewrites the previous query to use symbolic notation:

```
SELECT prd_type_id, year, month, sales_amount
FROM all_sales
WHERE prd_type_id BETWEEN 1 AND 2
AND emp_id = 21
MODEL
PARTITION BY (prd_type_id)
DIMENSION BY (month, year)
MEASURES (amount sales_amount) (
    sales_amount[month=1, year=2004] = sales_amount[month=1, year=2003],
    sales_amount[month=2, year=2004] =
        sales_amount[month=3, year=2003] + sales_amount[month=3, year=2003],
    sales_amount[month=3, year=2004] =
        ROUND(sales_amount[month=3, year=2003] * 1.25, 2)
)
ORDER BY prd_type_id, year, month;
```

When using positional or symbolic notation, it is important to be aware of the different way they handle null values in the dimensions. For example, sales_amount[null, 2003] returns the amount whose month is null and year is 2003, but sales_amount[month=null, year=2004] won't access a valid cell because null=null always returns false.

Accessing a Range of Cells Using BETWEEN and AND

You can access a range of cells using the BETWEEN and AND keywords. For example, the following expression sets the sales amount for January 2004 to the rounded average of the sales between January and March of 2003:

```
sales_amount[1, 2004] =

ROUND(AVG(sales amount) [month BETWEEN 1 AND 3, 2003], 2)
```

The following query shows the use of this expression:

```
SELECT prd_type_id, year, month, sales_amount
FROM all_sales
WHERE prd_type_id BETWEEN 1 AND 2
AND emp_id = 21
MODEL
PARTITION BY (prd_type_id)
DIMENSION BY (month, year)
MEASURES (amount sales_amount) (
    sales_amount[1, 2004] =
        ROUND(AVG(sales_amount) [month BETWEEN 1 AND 3, 2003], 2)
)
ORDER BY prd_type_id, year, month;
```

Accessing All Cells Using ANY and IS ANY

You can access all cells in an array using the ANY and IS ANY predicates. You use ANY with positional notation and IS ANY with symbolic notation. For example, the following expression sets the sales amount for January 2004 to the rounded sum of the sales for all months and years:

```
sales amount[1, 2004] =
      ROUND(SUM(sales amount)[ANY, year IS ANY], 2)
```

The following query shows the use of this expression:

```
SELECT prd type id, year, month, sales amount
   FROM all sales
   WHERE prd_type id BETWEEN 1 AND 2
   AND emp id = 21
   MODEL
   PARTITION BY (prd type id)
   DIMENSION BY (month, year)
   MEASURES (amount sales amount) (
     sales amount[1, 2004] =
       ROUND (SUM (sales amount) [ANY, year IS ANY], 2)
   ORDER BY prd type id, year, month;
```

Getting the Current Value of a Dimension **Using CURRENTV()**

You can get the current value of a dimension using the CURRENTV() function. For example, the following expression sets the sales amount for the first month of 2004 to 1.25 times the sales of the same month in 2003; notice the use of CURRENTV() to get the current month, which is 1:

```
sales amount[1, 2004] =
      ROUND(sales amount[CURRENTV(), 2003] * 1.25, 2)
```

The following query shows the use of this expression:

```
SELECT prd type id, year, month, sales amount
  FROM all sales
  WHERE prd_type_id BETWEEN 1 AND 2
  AND emp id = 21
  MODEL
  PARTITION BY (prd type id)
  DIMENSION BY (month, year)
  MEASURES (amount sales amount) (
    sales amount[1, 2004] =
      ROUND(sales amount[CURRENTV(), 2003] * 1.25, 2)
  )
  ORDER BY prd type id, year, month;
```

The output from this query is as follows (I've highlighted the values for 2004 in bold):

PRD_TYPE_ID	YEAR	MONTH	SALES_AMOUNT
1	2003	1	10034.84
1	2003	2	15144.65
1	2003	3	20137.83
1	2003	4	25057.45
1	2003	5	17214.56
1	2003	6	15564.64
1	2003	7	12654.84

1	2003	8	17434.82
1	2003	9	19854.57
1	2003	10	21754.19
1	2003	11	13029.73
1	2003	12	10034.84
1	2004	1	12543.55
2	2003	1	1034.84
2	2003	2	1544.65
2	2003	3	2037.83
2	2003	4	2557.45
2	2003	5	1714.56
2	2003	6	1564.64
2	2003	7	1264.84
2	2003	8	1734.82
2	2003	9	1854.57
2	2003	10	2754.19
2	2003	11	1329.73
2	2003	12	1034.84
2	2004	1	1293.55

Accessing Cells Using a FOR Loop

You can access cells using a FOR loop. For example, the following expression sets the sales amount for the first three months of 2004 to 1.25 times the sales of the same months in 2003; notice the use of the FOR loop and the INCREMENT keyword that specifies the amount to increment month by during each iteration of the loop:

```
sales_amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
ROUND(sales amount[CURRENTV(), 2003] * 1.25, 2)
```

The following query shows the use of this expression:

```
SELECT prd_type_id, year, month, sales_amount
FROM all_sales
WHERE prd_type_id BETWEEN 1 AND 2
AND emp_id = 21
MODEL
PARTITION BY (prd_type_id)
DIMENSION BY (month, year)
MEASURES (amount sales_amount) (
    sales_amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
    ROUND(sales_amount[CURRENTV(), 2003] * 1.25, 2)
)
ORDER BY prd type id, year, month;
```

The output from this query is as follows (I've highlighted the values for 2004 in bold):

PRD_TYPE_ID	YEAR	MONTH	SALES_AMOUNT
1	2003	1	10034.84
1	2003	2	15144.65
1	2003	3	20137.83
1	2003	4	25057.45

1	2003	5	17214.56
1	2003	6	15564.64
1	2003	7	12654.84
1	2003	8	17434.82
1	2003	9	19854.57
1	2003	10	21754.19
1	2003	11	13029.73
1	2003	12	10034.84
1	2004	1	12543.55
1	2004	2	18930.81
1	2004	3	25172.29
2	2003	1	1034.84
2	2003	2	1544.65
2	2003	3	2037.83
2	2003	4	2557.45
2	2003	5	1714.56
2	2003	6	1564.64
2	2003	7	1264.84
2	2003	8	1734.82
2	2003	9	1854.57
2	2003	10	2754.19
2	2003	11	1329.73
2	2003	12	1034.84
2	2004	1	1293.55
2	2004	2	1930.81
2	2004	3	2547.29

Handling Null and Missing Values

In this section, you'll learn how to handle null and missing values using the MODEL clause.

Using IS PRESENT

IS PRESENT returns true if the row specified by the cell reference existed prior to the execution of the MODEL clause. For example:

```
sales amount[CURRENTV(), 2003] IS PRESENT
```

will return true if sales amount [CURRENTV(), 2003] exists.

The following expression sets the sales amount for the first three months of 2004 to 1.25 multiplied by the sales of the same months in 2003:

```
sales amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
       CASE WHEN sales amount[CURRENTV(), 2003] IS PRESENT THEN
         ROUND(sales amount[CURRENTV(), 2003] * 1.25, 2)
       ELSE
         0
       END
```

The following query shows the use of this expression:

```
SELECT prd type id, year, month, sales amount
    FROM all sales
    WHERE prd_type_id BETWEEN 1 AND 2
```

```
AND emp_id = 21

MODEL

PARTITION BY (prd_type_id)

DIMENSION BY (month, year)

MEASURES (amount sales_amount) (

sales_amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =

CASE WHEN sales_amount[CURRENTV(), 2003] IS PRESENT THEN

ROUND(sales_amount[CURRENTV(), 2003] * 1.25, 2)

ELSE

0

END

)

ORDER BY prd_type_id, year, month;
```

The output of this query is the same as the example in the previous section.

Using PRESENTV()

PRESENTV (cell, exprl, exprl) returns the expression exprl if the row specified by the cell reference existed prior to the execution of the MODEL clause. If the row doesn't exist, the expression exprl is returned. For example:

```
PRESENTV(sales_amount[CURRENTV(), 2003],
ROUND(sales_amount[CURRENTV(), 2003] * 1.25, 2), 0)
```

will return the rounded sales amount if sales_amount[CURRENTV(), 2003] exists; otherwise 0 will be returned.

The following query shows the use of this expression:

```
SELECT prd_type_id, year, month, sales_amount
FROM all_sales
WHERE prd_type_id BETWEEN 1 AND 2
AND emp_id = 21
MODEL
PARTITION BY (prd_type_id)
DIMENSION BY (month, year)
MEASURES (amount sales_amount) (
    sales_amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
    PRESENTV(sales_amount[CURRENTV(), 2003],
        ROUND(sales_amount[CURRENTV(), 2003] * 1.25, 2), 0)
)
ORDER BY prd_type_id, year, month;
```

Using PRESENTNNV()

PRESENTNNV (cell, exprl, exprl) returns the expression exprl if the row specified by the cell reference existed prior to the execution of the MODEL clause and the cell value is not null. If the row doesn't exist or the cell value is null, the expression exprl is returned. For example,

```
PRESENTNNV(sales_amount[CURRENTV(), 2003],
ROUND(sales_amount[CURRENTV(), 2003] * 1.25, 2), 0)
```

will return the rounded sales amount if sales amount [CURRENTV(), 2003] exists and is not null; otherwise 0 will be returned.

Using IGNORE NAV and KEEP NAV

IGNORE NAV returns

- 0 for null or missing numeric values.
- An empty string for null or missing string values.
- 01-JAN-2000 for null or missing date values.
- Null for all other database types.

KEEP NAV returns null for null or missing numeric values. Be aware that KEEP NAV is the default.

The following query shows the use of IGNORE NAV:

```
SELECT prd type id, year, month, sales amount
   FROM all sales
   WHERE prd type id BETWEEN 1 AND 2
   AND emp id = 21
   MODEL IGNORE NAV
   PARTITION BY (prd_type_id)
   DIMENSION BY (month, year)
   MEASURES (amount sales amount) (
     sales amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
       ROUND(sales amount[CURRENTV(), 2003] * 1.25, 2)
   ORDER BY prd type id, year, month;
```

Updating Existing Cells

By default, if the cell referenced on the left side of an expression exists, then it is updated. If the cell doesn't exist, then a new row in the array is created. You can change this default behavior using RULES UPDATE, which specifies that if the cell doesn't exist, a new row will not be created. The following query shows the use of RULES UPDATE:

```
SELECT prd type id, year, month, sales amount
    FROM all sales
    WHERE prd type id BETWEEN 1 AND 2
    AND emp id = 21
    MODEL
    PARTITION BY (prd type id)
    DIMENSION BY (month, year)
    MEASURES (amount sales amount)
    RULES UPDATE (
      sales amount[FOR month FROM 1 TO 3 INCREMENT 1, 2004] =
        ROUND (sales amount [CURRENTV(), 2003] * 1.25, 2)
    ORDER BY prd type id, year, month;
```

Because cells for 2004 don't exist and RULES UPDATE is used, no new rows are created in the array for 2004; therefore, the query doesn't return rows for 2004. The following listing shows the output for the query—notice there are no rows for 2004:

PRD_TYPE_ID	YEAR	MONTH	SALES_AMOUNT
1	2003	1	10034.84
1	2003	2	15144.65
1	2003	3	20137.83
1	2003	4	25057.45
1	2003	5	17214.56
1	2003	6	15564.64
1	2003	7	12654.84
1	2003	8	17434.82
1	2003	9	19854.57
1	2003	10	21754.19
1	2003	11	13029.73
1	2003	12	10034.84
2	2003	1	1034.84
2	2003	2	1544.65
2	2003	3	2037.83
2	2003	4	2557.45
2	2003	5	1714.56
2	2003	6	1564.64
2	2003	7	1264.84
2	2003	8	1734.82
2	2003	9	1854.57
2	2003	10	2754.19
2	2003	11	1329.73
2	2003	12	1034.84

Using the PIVOT and UNPIVOT Clauses

The PIVOT clause is new for Oracle Database 11g and enables you to rotate rows into columns in the output from a query, and, at the same time, to run an aggregation function on the data. Oracle Database 11g also has an UNPIVOT clause that rotates columns into rows in the output from a query.

PIVOT and UNPIVOT are useful to see overall trends in large amounts of data, such as trends in sales over a period of time. You'll see queries that show the use of PIVOT and UNPIVOT in the following sections.

A Simple Example of the PIVOT Clause

The easiest way to learn how to use the PIVOT clause is to see an example. The following query shows the total sales amount of product types #1, #2, and #3 for the first four months in 2003; notice that the cells in the query's output show the sum of the sales amounts for each product type in each month:

```
SELECT *
   FROM (
     SELECT month, prd type id, amount
     FROM all sales
     WHERE year = 2003
     AND prd_type_id IN (1, 2, 3)
   )
   PIVOT (
     SUM(amount) FOR month IN (1 AS JAN, 2 AS FEB, 3 AS MAR, 4 AS APR)
   ORDER BY prd type id;
```

APR	MAR	FEB	JAN	PRD_TYPE_ID
120344.7	91826.98	70567.9	38909.04	1
15664.7	16826.98	13367.9	14309.04	2
23844.7	20626.98	15467.9	24909.04	3

Starting with the first line of output, you can see there was

- \$38,909.04 of product type #1 sold in January.
- \$70,567.90 of product type #1 sold in February.
- ...and so on for the rest of the first line.

The second line of output shows there was

- \$14,309.04 of product type #2 sold in January.
- \$13,367.90 of product type #2 sold in February.
- ...and so on for the rest of the output.



NOTE

PIVOT is a powerful tool that allows you to see trends in sales of types of products over a period of months. Based on such trends, a real store could use the information to alter their sales tactics and formulate new marketing campaigns.

The previous SELECT statement has the following structure:

```
SELECT *
   FROM (
     inner_query
   PIVOT (
     aggregate function FOR pivot column IN (list of values)
   ORDER BY ...;
```

Let's break down the previous example into the structural elements:

- There is an inner and outer query. The inner query gets the month, product type, and amount from the all sales table and passes the results to the outer query.
- SUM(amount) FOR month IN (1 AS JAN, 2 AS FEB, 3 AS MAR, 4 AS APR) is the line in the PIVOT clause.
 - The SUM() function adds up the sales amounts for the product types in the first four months (the months are listed in the IN part). Instead of returning the months as 1, 2, 3, and 4 in the output, the AS part renames the numbers to JAN, FEB, MAR, and APR to make the months more readable in the output.
 - The month column from the all_sales table is used as the pivot column. This means that the months appear as columns in the output. In effect, the rows are rotated—or *pivoted*—to view the months as columns.
- At the very end of the example, the ORDER BY prd_type_id line simply orders the results by the product type.

Pivoting on Multiple Columns

You can pivot on multiple columns by placing those columns in the FOR part of the PIVOT. The following example pivots on both the month and prd_type_id columns, which are referenced in the FOR part; notice that the list of values in the IN part of the PIVOT contains a value for the month and prd_type_id columns:

```
SELECT *
    FROM (
      SELECT month, prd type id, amount
      FROM all sales
      WHERE year = 2003
      AND prd type id IN (1, 2, 3)
    )
    PIVOT (
      SUM(amount) FOR (month, prd type id) IN (
        (1, 2) AS JAN PRDTYPE2,
        (2, 3) AS FEB PRDTYPE3,
        (3, 1) AS MAR PRDTYPE1,
        (4, 2) AS APR PRDTYPE2
      )
    );
    JAN PRDTYPE2 FEB PRDTYPE3 MAR PRDTYPE1 APR PRDTYPE2
        14309.04 15467.9 91826.98 15664.7
```

The cells in the output show the sum of the sales amounts for each product type in the specified month (the product type and month to query are placed in the list of values in the IN part). As you can see from the query output, there were the following sales amounts:

■ \$14,309.04 of product type #2 in January

- \$15,467.90 of product type #3 in February
- \$91,826.98 of product type #1 in March
- \$15,664.70 of product type #2 in April

You can put any values in the IN part to get the values of interest to you. In the following example, the values of the product types are shuffled in the IN part to get the sales for those product types in the specified months:

```
SELECT *
    FROM (
      SELECT month, prd type id, amount
      FROM all sales
      WHERE year = 2003
      AND prd type id IN (1, 2, 3)
    PIVOT (
      SUM(amount) FOR (month, prd type id) IN (
        (1, 1) AS JAN PRDTYPE1,
        (2, 2) AS FEB PRDTYPE2,
        (3, 3) AS MAR PRDTYPE3,
        (4, 1) AS APR PRDTYPE1
      )
    );
    JAN PRDTYPE1 FEB PRDTYPE2 MAR PRDTYPE3 APR PRDTYPE1
    -----
       38909.04 13367.9 20626.98 120344.7
```

As you can see from this output, there were the following sales amounts:

- \$38,909.04 of product type #1 in January
- \$13,367.90 of product type #2 in February
- \$20,626.98 of product type #3 in March
- \$120,344.70 of product type #1 in April

Using Multiple Aggregate Functions in a Pivot

You can use multiple aggregate functions in a pivot. For example, the following query uses SUM() to get the total sales for the product types in January and February and AVG() to get the averages of the sales:

```
SELECT *
     FROM (
       SELECT month, prd type id, amount
       FROM all sales
       WHERE year = 2003
       AND prd type id IN (1, 2, 3)
     )
```

As you can see, the first line of output shows for product type #1:

- A total of \$38,909.04 and an average of \$6,484.84 sold in January
- A total of \$70,567.90 and an average of \$11,761.32 sold in February

The second line of output shows for product type #2:

- A total of \$14,309.04 and an average of \$2,384.84 sold in January
- A total of \$13,367.90 and an average of \$2,227.98 sold in February

...and so on for the rest of the output.

Using the UNPIVOT Clause

The UNPIVOT clause rotates columns into rows. The examples in this section use the following table named pivot_sales_data (created by the store_schema.sql script); pivot_sales_data is populated by a query that returns a pivoted version of the sales data:

```
CREATE TABLE pivot_sales_data AS

SELECT *

FROM (

SELECT month, prd_type_id, amount

FROM all_sales

WHERE year = 2003

AND prd_type_id IN (1, 2, 3)
)

PIVOT (

SUM(amount) FOR month IN (1 AS JAN, 2 AS FEB, 3 AS MAR, 4 AS APR)
)

ORDER BY prd_type_id;
```

The following query returns the contents of the pivot_sales_data table:

```
SELECT *
FROM pivot sales data;
```

PRD_TYPE_ID	JAN	FEB	MAR	APR
1	38909.04	70567.9	91826.98	120344.7
2	14309.04	13367.9	16826.98	15664.7
3	24909.04	15467.9	20626.98	23844.7

The next query uses UNPIVOT to get the sales data in an unpivoted form:

```
SELECT *
   FROM pivot sales data
   UNPIVOT (
     amount FOR month IN (JAN, FEB, MAR, APR)
   ORDER BY prd type id;
```

PRD_TYPE_ID	MON	AMOUNT
1	JAN	38909.04
1	FEB	70567.9
1	MAR	91826.98
1	APR	120344.7
2	JAN	14309.04
2	FEB	13367.9
2	APR	15664.7
2	MAR	16826.98
3	JAN	24909.04
3	MAR	20626.98
3	FEB	15467.9
3	APR	23844.7

Notice that the query rotates the pivoted data. For example, the monthly sales totals that appear in the horizontal rows of pivot sales data are shown in the vertical AMOUNT column.



Consider using UNPIVOT when you have a query that returns rows with many columns and you want to view those columns as rows.

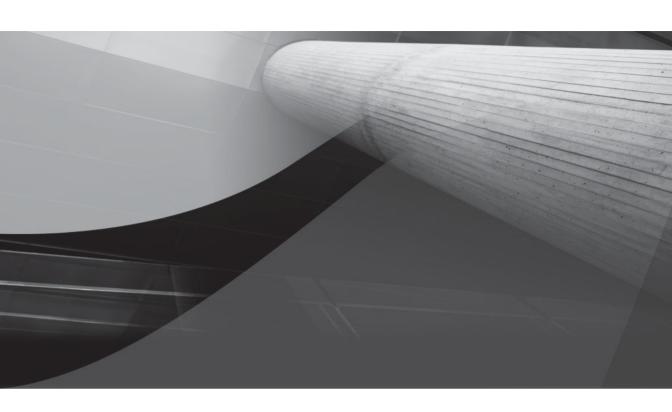
Summary

In this chapter, you learned the following:

- The set operators (UNION ALL, UNION, INTERSECT, and MINUS) allow you to combine rows returned by two or more queries.
- TRANSLATE (x, from string, to string) translates characters in one string to characters in another string.
- DECODE(value, search value, result, default value) compares value with search value. If the values are equal, DECODE() returns search value; otherwise default value is returned. DECODE () allows you to perform if-then-else logic in SQL.

- CASE is similar to DECODE (). You should use CASE because it is ANSI-compliant.
- Queries may be run against data that is organized into a hierarchy.
- ROLLUP extends the GROUP BY clause to return a row containing a subtotal for each group of rows, plus a row containing a grand total for all the groups.
- CUBE extends the GROUP BY clause to return rows containing a subtotal for all combinations of columns, plus a row containing the grand total.
- The database has many built-in analytic functions that enable you to perform complex calculations, such as finding the top-selling product type for each month, the top salespersons, and so on.
- The MODEL clause performs inter-row calculations and allows you to treat table data as an array. This gives you the ability to perform calculations in a similar manner to spreadsheet calculations.
- The Oracle Database 11*g* PIVOT and UNPIVOT clauses are useful for seeing overall trends in large amounts of data.

In the next chapter, you'll learn about changing the contents of a table.



CHAPTER 8

Changing Table Contents



n this chapter, you'll learn more about changing the contents of tables. Specifically, you'll learn the following:

- How to add, modify, and remove rows using the INSERT, UPDATE, and DELETE statements
- How database transactions may consist of multiple INSERT, UPDATE, and DELETE statements
- How to make the results of your transactions permanent using the COMMIT statement or undo their results entirely using the ROLLBACK statement
- How an Oracle database can process multiple transactions at the same time
- How to use query flashbacks to view rows as they originally were before you made changes to them

Adding Rows Using the INSERT Statement

You use the INSERT statement to add rows to a table. You can specify the following information in an INSERT statement:

- The table into which the row is to be inserted
- A list of columns for which you want to specify column values
- A list of values to store in the specified columns

When adding a row, you typically supply a value for the primary key and all other columns that are defined as NOT NULL. You don't have to specify values for NULL columns if you don't want to; by default they will be set to null.

You can find out which columns are defined as NOT NULL using the SQL*Plus DESCRIBE command. The following example describes the customers table:

DESCRIBE customers

Name	Null?		Type	
CUSTOMER ID	NOT	NIII.T.	NUMBER (38)	
FIRST NAME			VARCHAR2 (10)	
LAST NAME			VARCHAR2 (10)	
DOB			DATE	
PHONE			VARCHAR2 (12)	

As you can see, the customer_id, first_name, and last_name columns are NOT NULL, meaning that you must supply a value for these columns. The dob and phone columns don't require a value: If you omit these values when adding a row, the columns would be set to null.

The following INSERT statement adds a row to the customers table. Notice that the order of values in the VALUES clause matches the order in which the columns are specified in the column list. Also notice that the statement has three parts: the table name, the column list, and the values to be added:

```
INSERT INTO customers (
     customer id, first name, last name, dob, phone
   ) VALUES (
     6, 'Fred', 'Brown', '01-JAN-1970', '800-555-1215'
   );
   1 row created.
```

SQL*Plus responds that one row has been created. You can verify this by performing the following SELECT statement:

SELECT * FROM customers:

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	
6	Fred	Brown	01-JAN-70	800-555-1215

Notice the new row appears in the results returned by the query.

Omitting the Column List

You may omit the column list when supplying values for every column, as in this example:

```
INSERT INTO customers
    VALUES (7, 'Jane', 'Green', '01-JAN-1970', '800-555-1216');
```

When you omit the column list, the order of the values you supply must match the order of the columns as listed in the output from the DESCRIBE command.

Specifying a Null Value for a Column

You can specify a null value for a column using the NULL keyword. For example, the following INSERT specifies a null value for the dob and phone columns:

```
INSERT INTO customers
    VALUES (8, 'Sophie', 'White', NULL, NULL);
```

When you view this row using a query, you won't see a value for the dob and phone columns, because they've been set to null values:

```
SELECT *
    FROM customers
    WHERE customer id = 8;
    CUSTOMER ID FIRST NAME LAST NAME DOB PHONE
             8 Sophie White
```

Notice the dob and phone column values are blank.

Including Single and Double Quotes in a Column Value

You can include a single and double quote in a column value. For example, the following INSERT specifies a last name of O'Malley for a new customer; notice the use of two single quotes in the last name after the letter O:

```
INSERT INTO customers

VALUES (9, 'Kyle', 'O''Malley', NULL, NULL);
```

The next example specifies the name The "Great" Gatsby for a new product:

```
INSERT INTO products (
    product_id, product_type_id, name, description, price
) VALUES (
    13, 1, 'The "Great" Gatsby', NULL, 12.99
);
```

Copying Rows from One Table to Another

You can copy rows from one table to another using a query in the place of the column values in the INSERT statement. The number of columns and the column types in the source and destination must match. The following example uses a SELECT to retrieve the first_name and last name columns for customer #1 and supplies those columns to an INSERT statement:

```
INSERT INTO customers (customer_id, first_name, last_name)
SELECT 10, first_name, last_name
FROM customers
WHERE customer id = 1;
```

Notice that the customer id for the new row is set to 10.



NOTE

Oracle Database 9i introduced the MERGE statement, which allows you to merge rows from one table to another. MERGE is much more flexible than combining an INSERT and a SELECT to copy rows from one table to another. You'll learn about MERGE later in the section "Merging Rows Using MERGE."

Modifying Rows Using the UPDATE Statement

You use the UPDATE statement to modify rows in a table. When you use the UPDATE statement, you typically specify the following information:

- The table name
- A WHERE clause that specifies the rows to be changed
- A list of column names, along with their new values, specified using the SET clause

You can change one or more rows using the same UPDATE statement. If more than one row is specified, the same change will be implemented for all of those rows. For example, the following UPDATE statement sets the last_name column to Orange for the row whose customer_id is 2:

```
UPDATE customers
    SET last name = 'Orange'
    WHERE customer id = 2;
    1 row updated.
```

SQL*Plus confirms that one row was updated. If the WHERE clause were omitted, then all the rows would be updated. The following query confirms the change was made:

```
SELECT *
   FROM customers
   WHERE customer id = 2;
   CUSTOMER ID FIRST NAME LAST NAME DOB PHONE
   ______
          2 Cynthia Orange 05-FEB-68 800-555-1212
```

You can change multiple rows and multiple columns in the same UPDATE statement. For example, the following UPDATE raises the price by 20 percent for all products whose current price is greater than or equal to \$20. The UPDATE also changes those products' names to lowercase:

```
UPDATE products
    SET
      price = price * 1.20,
      name = LOWER(name)
    WHERE
      price >= 20;
    3 rows updated.
```

As you can see, three rows are updated by this statement. The following query confirms the change:

DDTCE

```
SELECT product id, name, price
    FROM products
    WHERE price \geq= (20 * 1.20);
```

TD MAME

ID	NAME	PRICE
2	chemistry	36
3	supernova	31.19
5	z-files	59.99



You can also use a subquery with an UPDATE statement. This is covered in Chapter 6 in the section "Writing an UPDATE Statement Containing a Subquery."

The RETURNING Clause

In Oracle Database 10g and above, you can use the RETURNING clause to return the value from an aggregate function such as AVG(). Aggregate functions were covered in Chapter 4.

The following tasks are performed by the next example:

- Declares a variable named average product price
- Decreases the price column of the rows in the products table and saves the average price in the average product price variable using the RETURNING clause
- Prints the value of the average product price variable

VARIABLE average_product_price NUMBER

Removing Rows Using the DELETE Statement

You use the DELETE statement to remove rows from a table. Generally, you should specify a WHERE clause that limits the rows that you wish to delete; if you don't, *all* the rows will be deleted.

The following DELETE statement removes the row from the customers table whose customer id is 10:

```
DELETE FROM customers
WHERE customer_id = 10;
```

1 row deleted.

SQL*Plus confirms that one row has been deleted.

You can also use a subquery with a DELETE statement. This is covered in Chapter 6 in the section "Writing a DELETE Statement Containing a Subquery."



NOTE

If you've been following along with the previous INSERT, UPDATE, and DELETE statements, roll them back using ROLLBACK so that your results match those shown in the rest of this chapter. Don't worry if you've already disconnected from the database: simply rerun the store_schema.sql script to re-create everything.

Database Integrity

When you execute a DML statement (an INSERT, UPDATE, or DELETE, for example), the database ensures that the rows in the tables maintain their integrity. This means that any changes made to the rows do not affect the primary key and foreign key relationships for the tables.

Enforcement of Primary Key Constraints

Let's examine some examples that show the enforcement of a primary key constraint. The customers table's primary key is the customer id column, which means that every value stored in the customer id column must be unique. If you try to insert a row with a duplicate value for a primary key, the database returns the error ORA-00001, as in this example:

```
SOL> INSERT INTO customers (
         customer id, first name, last name, dob, phone
      3 ) VALUES (
          1, 'Jason', 'Price', '01-JAN-60', '800-555-1211'
      5);
    INSERT INTO customers (
    ERROR at line 1:
    ORA-00001: unique constraint (STORE.CUSTOMERS PK) violated
```

If you attempt to update a primary key value to a value that already exists in the table, the database returns the same error:

```
SOL> UPDATE customers
     2 SET customer id = 1
     3 WHERE customer id = 2;
   UPDATE customers
   ERROR at line 1:
   ORA-00001: unique constraint (STORE.CUSTOMERS PK) violated
```

Enforcement of Foreign Key Constraints

A foreign key relationship is one in which a column from one table is referenced in another. For example, the product type id column in the products table references the product type id column in the product types table. The product types table is known as the parent table, and the products table is known as the child table, reflecting the dependence of the product type id column in the products table on the product type id column in the product types table.

If you try to insert a row into the products table with a nonexistent product type id, the database will return the error ORA-02291. This error indicates the database couldn't find a matching parent key value (the parent key is the product_type id column of the product types table). In the following example, the error is returned because there is no row in the product types table whose product type id is 6:

```
SQL> INSERT INTO products (
       2 product id, product type id, name, description, price
       3 ) VALUES (
       4 13, 6, 'Test', 'Test', NULL
       5);
     INSERT INTO products (
     ERROR at line 1:
     ORA-02291: integrity constraint (STORE.PRODUCTS FK PRODUCT TYPES)
      violated - parent key not found
```

Similarly, if you attempt to update the product_type_id of a row in the products table to a nonexistent parent key value, the database returns the same error, as in this example:

```
SQL> UPDATE products
2   SET product_type_id = 6
3   WHERE product_id = 1;
UPDATE products
*
ERROR at line 1:
ORA-02291: integrity constraint (STORE.PRODUCTS_FK_PRODUCT_TYPES)
violated - parent key not found
```

Finally, if you attempt to delete a row in the parent table that has dependent child rows, the database returns error ORA-02292. For example, if you attempt to delete the row whose product_type_id is 1 from the product_types table, the database will return this error because the products table contains rows whose product type id is 1:

```
SQL> DELETE FROM product_types
    2 WHERE product_type_id = 1;
DELETE FROM product_types
    *
ERROR at line 1:
ORA-02292: integrity constraint (STORE.PRODUCTS_FK_PRODUCT_TYPES)
    violated - child record found
```

If the database were to allow this deletion, the child rows would be invalid because they wouldn't point to valid values in the parent table.

Using Default Values

Oracle Database 9*i* introduced a feature that allows you to define a default value for a column. For example, the following statement creates a table named order_status; the status column is defaulted to 'Order placed' and the last_modified column is defaulted to the date and time returned by SYSDATE:

```
CREATE TABLE order_status (
    order_status_id INTEGER
    CONSTRAINT default_example_pk PRIMARY KEY,
    status VARCHAR2(20) DEFAULT 'Order placed' NOT NULL,
    last_modified DATE DEFAULT SYSDATE
);
```



NOTF

The order_status table is created by the store_schema.sql script. This means you don't have to type in the previous CREATE TABLE statement yourself. Also, you don't have to type in the INSERT statements shown in this section.

When you add a new row to the order_status table but don't specify the values for the status and last_modified columns, those columns are set to the default values. For example, the following INSERT statement omits values for the status and last modified columns:

```
INSERT INTO order status (order status id)
   VALUES (1);
```

The status column is set to the default value of 'Order placed', and the last modified column is set to the current date and time.

You can override the defaults by specifying a value for the columns, as shown in the following example:

```
INSERT INTO order status (order status id, status, last modified)
     VALUES (2, 'Order shipped', '10-JUN-2004');
```

The following guery retrieves the rows from order status:

SELECT *

FROM order status;

ORDER_STATUS_ID	STATUS	LAST_MODI
1	Order placed	25-JUL-07
2	Order shipped	10-JUN-04

You can set a column back to the default using the DEFAULT keyword in an UPDATE statement. For example, the following UPDATE sets the status column to the default:

```
UPDATE order status
    SET status = DEFAULT
    WHERE order status id = 2;
```

The following query shows the change made by this UPDATE statement:

SELECT *

FROM order status;

ORDER_STATUS_ID	STATUS	LAST_MODI
1	Order placed	25-JUL-07
2	Order placed	10-JUN-04

Merging Rows Using MERGE

Oracle Database 9i introduced the MERGE statement, which allows you to merge rows from one table into another. For example, you might want to merge changes to products listed in one table into the products table.

The store schema contains a table named product changes that was created using the following CREATE TABLE statement in store schema.sql:

```
CREATE TABLE product_changes (
      product id INTEGER
        CONSTRAINT prod_changes_pk PRIMARY KEY,
      product type id INTEGER
        CONSTRAINT prod changes fk product types
        REFERENCES product types (product type id),
      name VARCHAR2(30) NOT NULL,
```

```
description VARCHAR2(50),
  price NUMBER(5, 2)
);
```

The following query retrieves the product_id, product_type_id, name, and price columns from this table:

SELECT product_id, product_type_id, name, price FROM product_changes;

PRODUCT_ID	PRODUCT_TYPE_ID	NAME	PRICE
1	1	Modern Science	40
2	1	New Chemistry	35
3	1	Supernova	25.99
13	2	Lunar Landing	15.99
14	2	Submarine	15.99
15	2	Airplane	15.99

Let's say you want to merge the rows from the product_changes table into the products table as follows:

- For rows with matching product_id values in the two tables, update the existing rows in products with the column values from product_changes. For example, product #1 has a different price in product_changes from the one in products; therefore, product #1's price must be updated in the products table. Similarly, product #2 has a different name and price, so both values must be updated in products. Finally, product #3 has a different product type id, and so this value must be updated in products.
- For new rows in product_changes, insert those new rows into the products table. Products #13, #14, and #15 are new in product_changes and must therefore be inserted into products.

The easiest way to learn how to use the MERGE statement is to see an example. The following example performs the merge as defined in the previous bullet points:

```
MERGE INTO products p

USING product_changes pc ON (
    p.product_id = pc.product_id
)

WHEN MATCHED THEN

UPDATE

SET
    p.product_type_id = pc.product_type_id,
    p.name = pc.name,
    p.description = pc.description,
    p.price = pc.price

WHEN NOT MATCHED THEN

INSERT (
    p.product_id, p.product_type_id, p.name,
    p.description, p.price
) VALUES (
```

```
pc.product id, pc.product type id, pc.name,
   pc.description, pc.price
 ):
6 rows merged.
```



NOTE

You'll find a script named merge example.sql in the SQL directory. This script contains the previous MERGE statement.

Notice the following points about the MERGE statement:

- The MERGE INTO clause specifies the name of the table to merge the rows into. In the example, this is the products table, which is given an alias of p.
- The USING ... ON clause specifies a table join. In the example, the join is made on the product id columns in the products and product changes tables. The product changes table is also given an alias of pc.
- The WHEN MATCHED THEN clause specifies the action to take when the USING ... ON clause is satisfied for a row. In the example, this action is an UPDATE statement that sets the product type id, name, description, and price columns of the existing row in the products table to the column values for the matching row in the product changes table.
- The WHEN NOT MATCHED THEN clause specifies the action to take when the USING ... ON clause is *not* satisfied for a row. In the example, this action is an INSERT statement that adds a row to the products table, taking the column values from the row in the product changes table.

If you run the previous MERGE statement, you'll see that it reports six rows are merged; these are the rows with product id values of 1, 2, 3, 13, 14, and 15. The following query retrieves the six merged rows from the products table:

SELECT product id, product type id, name, price FROM products WHERE product id IN (1, 2, 3, 13, 14, 15);

PRODUCT_ID	PRODUCT_TYPE_ID	NAME	PRICE
1	1	Modern Science	40
2	1	New Chemistry	35
3	1	Supernova	25.99
13	2	Lunar Landing	15.99
14	2	Submarine	15.99
15	2	Airplane	15.99

The following changes were made to these rows:

- Product #1 has a new price.
- Product #2 has a new name and price.

- Product #3 has a new product type ID.
- Products #13, #14, and #15 are new.

Now that you've seen how to make changes to the contents of tables, let's move on to database transactions.

Database Transactions

A database *transaction* is a group of SQL statements that perform a *logical unit of work*. You can think of a transaction as an inseparable set of SQL statements whose results should be made permanent in the database as a whole (or undone as a whole).

An example of a database transaction is a transfer of money from one bank account to another. One UPDATE statement would subtract from the total amount of money from one account, and another UPDATE would add money to the other account. Both the subtraction and the addition must be permanently recorded in the database; otherwise, money will be lost. If there is a problem with the money transfer, then the subtraction and addition must both be undone. The simple example outlined in this paragraph uses only two UPDATE statements, but a transaction may consist of many INSERT, UPDATE, and DELETE statements.

Committing and Rolling Back a Transaction

To permanently record the results made by SQL statements in a transaction, you perform a *commit*, using the SQL COMMIT statement. If you need to undo the results, you perform a *rollback*, using the SQL ROLLBACK statement, which resets all the rows back to what they were originally.

The following example adds a row to the customers table and then makes the change permanent by performing a COMMIT:

```
INSERT INTO customers
VALUES (6, 'Fred', 'Green', '01-JAN-1970', '800-555-1215');

1 row created.
COMMIT;
Commit complete.
```

The following example updates customer #1 and then undoes the change by performing a ROLLBACK:

```
UPDATE customers

SET first_name = 'Edward'

WHERE customer_id = 1;

1 row updated.

ROLLBACK;

Rollback complete.
```

The following query shows the new row from the COMMIT statement:

SELECT *

FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	
6	Fred	Green	01-JAN-70	800-555-1215

Notice that customer #6 has been made permanent by the COMMIT, but the change to customer #1's first name has been undone by the ROLLBACK.

Starting and Ending a Transaction

A transaction is a logical unit of work that enables you to split up your SQL statements. A transaction has a beginning and an end; it begins when one of the following events occurs:

- You connect to the database and perform a DML statement (an INSERT, UPDATE, or DELETE).
- A previous transaction ends and you enter another DML statement.

A transaction ends when one of the following events occurs:

- You perform a COMMIT or a ROLLBACK.
- You perform a DDL statement, such as a CREATE TABLE statement, in which case a COMMIT is automatically performed.
- You perform a DCL statement, such as a GRANT statement, in which case a COMMIT is automatically performed. You'll learn about GRANT in the next chapter.
- You disconnect from the database. If you exit SQL*Plus normally, by entering the EXIT command, a COMMIT is automatically performed for you. If SQL*Plus terminates abnormally—for example, if the computer on which SQL*Plus was running were to crash—a ROLLBACK is automatically performed. This applies to any program that accesses a database. For example, if you wrote a Java program that accessed a database and your program crashed, a ROLLBACK would be automatically performed.
- You perform a DML statement that fails, in which case a ROLLBACK is automatically performed for that individual DML statement.



It is poor practice not to explicitly commit or roll back your transactions, so perform a COMMIT or ROLLBACK at the end of your transactions.

Savepoints

You can also set a *savepoint* at any point within a transaction. These allow you to roll back changes to that savepoint. Savepoints can be useful to break up very long transactions, because, if you make a mistake after you've set a savepoint, you don't have to roll back the transaction all the way to the start. However, you should use savepoints sparingly: you might be better off restructuring your transaction into smaller transactions instead.

You'll see an example of a savepoint shortly, but first let's see the current price for products #4 and #5:

```
SELECT product_id, price
FROM products
WHERE product_id IN (4, 5);
```

PRODUCT_ID	PRICE
4	13.95
5	49.99

The price for product #4 is \$13.95, and the price for product #5 is \$49.99. The following UPDATE increases the price of product #4 by 20 percent:

```
UPDATE products

SET price = price * 1.20

WHERE product_id = 4;

1 row updated.
```

The following statement sets a savepoint named save1:

```
SAVEPOINT save1;
```

Savepoint created.

Any DML statements run after this point can be rolled back to the savepoint, and the change made to product #4 will be kept.

The following UPDATE increases the price of product #5 by 30 percent:

```
UPDATE products

SET price = price * 1.30

WHERE product_id = 5;

1 row updated.
```

The following query gets the prices of the two products:

```
SELECT product_id, price

FROM products

WHERE product id IN (4, 5);
```

```
PRODUCT_ID PRICE

4 16.74
5 64.99
```

Product #4's price is 20 percent greater, and product #5's price is 30 percent greater. The following statement rolls back the transaction to the savepoint established earlier:

ROLLBACK TO SAVEPOINT save1;

```
Rollback complete.
```

This has undone the price change for product #5, but left the price change for product #4 intact. The following query shows this:

SELECT product id, price FROM products WHERE product id IN (4, 5);

PRODUCT_ID	PRICE
4	16.74
5	49.99

As expected, product #4 has kept its increased price, but product #5's price is back to the original. The following ROLLBACK undoes the entire transaction:

ROLLBACK:

```
Rollback complete.
```

This has undone the change made to product #4's price, as is shown by the following query:

```
SELECT product id, price
    FROM products
    WHERE product id IN (4, 5);
```

PRODUCT_ID	PRICE
4	13.95
5	49.99

ACID Transaction Properties

Earlier, I defined a transaction as being a logical unit of work, that is, a group of related SQL statements that are either committed or rolled back as one unit. Database theory's more rigorous definition of a transaction states that a transaction has four fundamental properties, known as ACID properties (from the first letter of each property in the following list):

- **Atomic** Transactions are atomic, meaning that the SQL statements contained in a transaction make up a single unit of work.
- Consistent Transactions ensure that the database state remains consistent, meaning that the database is in a consistent state when a transaction begins and that it ends in another consistent state when the transaction finishes.
- **Isolated** Separate transactions should not interfere with each other.
- **Durable** Once a transaction has been committed, the database changes are preserved, even if the machine on which the database software is running crashes later.

The Oracle database software handles these ACID properties and has extensive recovery facilities for restoring databases after system crashes.

Concurrent Transactions

The Oracle database software supports many users interacting with a database, and each user can run their own transactions at the same time. These transactions are known as *concurrent* transactions.

If users are running transactions that affect the same table, the effects of those transactions are separated from each other until a COMMIT is performed. The following sequence of events, based on two transactions named T1 and T2 that access the customers table, illustrates the separation of transactions:

- 1. T1 and T2 perform a SELECT that retrieves all the rows from the customers table.
- 2. T1 performs an INSERT to add a row in the customers table, but T1 doesn't perform a COMMIT.
- **3.** T2 performs another SELECT and retrieves the same rows as those in step 1. T2 doesn't "see" the new row added by T1 in step 2.
- **4.** T1 finally performs a COMMIT to permanently record the new row added in step 2.
- 5. T2 performs another SELECT and finally "sees" the new row added by T1.

To summarize: T2 doesn't see the changes made by T1 until T1 commits its changes. This is the default level of isolation between transactions, but, as you'll learn later in the section "Transaction Isolation Levels," you can change the level of isolation.

Table 8-1 shows sample SQL statements that further illustrate how concurrent transactions work. The table shows the interleaved order in which the statements are performed by two transactions named T1 and T2. T1 retrieves rows, adds a row, and updates a row in the customers table. T2 retrieves rows from the customers table. T2 doesn't see the changes made by T1 until T1 commits its changes. You can enter the statements shown in Table 8-1 and see their results by starting two separate SQL*Plus sessions and connecting as the store user for both sessions; you enter the statements in the interleaved order shown in the table into the SQL*Plus sessions.

Transaction Locking

To support concurrent transactions, the Oracle database software must ensure that the data in the tables remains valid. It does this through the use of *locks*. Consider the following example in which two transactions named T1 and T2 attempt to modify customer #1 in the customers table:

- 1. T1 performs an UPDATE to modify customer #1, but T1 doesn't perform a COMMIT. T1 is said to have "locked" the row.
- 2. T2 also attempts to perform an UPDATE to modify customer #1, but since this row is already locked by T1, T2 is prevented from getting a lock on the row. T2's UPDATE statement has to wait until T1 ends and frees the lock on the row.
- **3.** T1 ends by performing a COMMIT, thus freeing the lock on the row.
- **4.** T2 gets the lock on the row and the UPDATE is performed. T2 holds the lock on the row until T2 ends.

To summarize: A transaction cannot get a lock on a row while another transaction already holds the lock on that row.

Transaction 1 T1 (1) SELECT * FROM customers; (3) INSERT INTO customers (customer id, first name, last name) VALUES (7, 'Jason', 'Price' (4) UPDATE customers SET last name = 'Orange' WHERE customer id = 2; (5) **SELECT** * FROM customers; The returned result set contains the new row and the update. (7) **COMMIT**; This commits the new row and the update.

(2) **SELECT** * FROM customers;

Transaction 2 T2

(6) SELECT *

FROM customers:

The returned result set doesn't contain the new row or the update made by T1. Instead, the result set contains the original rows retrieved in step 2.

(8) **SELECT** * FROM customers;

The returned result set contains the new row and the update made by T1 in steps 3 and 4.

TABLE 8-1 Concurrent Transactions



The easiest way to understand default locking is as follows: readers don't block readers, writers don't block readers, and writers only block writers when they attempt to modify the same row.

Transaction Isolation Levels

The transaction isolation level is the degree to which the changes made by one transaction are separated from other transactions running concurrently. Before you see the various transaction isolation levels available, you need to understand the types of problems that may occur when current transactions attempt to access the same rows in a table.

In the following list, you'll see examples of two concurrent transactions named T1 and T2 that are accessing the same rows; listed are the three types of potential transaction processing problems:

Phantom reads T1 reads a set of rows returned by a specified WHERE clause. T2 then inserts a new row, which also happens to satisfy the WHERE clause of the query previously used by T1. T1 then reads the rows again using the same query, but now sees the additional row just inserted by T2. This new row is known as a "phantom" because to T1 this row seems to have magically appeared.

- Nonrepeatable reads T1 reads a row, and T2 updates the same row just read by T1. T1 then reads the same row again and discovers that the row it read earlier is now different. This is known as a "nonrepeatable" read, because the row originally read by T1 has been changed.
- **Dirty reads** T1 updates a row, but doesn't commit the update. T2 then reads the updated row. T1 then performs a rollback, undoing the previous update. Now the row just read by T2 is no longer valid (it's "dirty") because the update made by T1 wasn't committed when the row was read by T2.

To deal with these potential problems, databases implement various levels of transaction isolation to prevent concurrent transactions from interfering with each other. The SQL standard defines the following transaction isolation levels, shown in order of increasing isolation:

- **READ UNCOMMITTED** Phantom reads, nonrepeatable reads, and dirty reads are permitted.
- **READ COMMITTED** Phantom reads and nonrepeatable reads are permitted, but dirty reads are not.
- **REPEATABLE READ** Phantom reads are permitted, but nonrepeatable and dirty reads are not.
- SERIALIZABLE Phantom reads, nonrepeatable reads, and dirty reads are not permitted.

The Oracle database software supports the READ COMMITTED and SERIALIZABLE transaction isolation levels. It doesn't support READ UNCOMMITTED or REPEATABLE READ levels.

The default transaction isolation level defined by the SQL standard is SERIALIZABLE, but the default used by the Oracle database is READ COMMITTED, which is acceptable for nearly all applications.



CAUTION

Although you can use SERIALIZABLE with the Oracle database, it may increase the time your SQL statements take to complete. You should only use SERIALIZABLE if you absolutely have to.

You set the transaction isolation level using the SET TRANSACTION statement. For example, the following statement sets the transaction isolation level to SERIALIZABLE:

You'll see an example of a transaction that uses the isolation level of SERIALIZABLE next.

A SERIALIZABLE Transaction Example

In this section, you'll see an example that shows the effect of setting the transaction isolation level to SERIALIZABLE.

The example uses two transactions named T1 and T2. T1 has the default isolation level of READ COMMITTED; T2 has a transaction isolation level of SERIALIZABLE. T1 and T2 will read the rows in the customers table, and then T1 will insert a new row and update an existing row in the customers table. Because T2 is SERIALIZABLE, it doesn't see the inserted row or the update made to the existing row made by T1, even after T1 commits its changes. That's because reading the inserted row would be a phantom read, and reading the update would be a nonrepeatable read, which are not permitted by SERIALIZABLE transactions.

Table 8-2 shows the SQL statements that make up T1 and T2 in the interleaved order in which the statements are to be performed.

Transaction 1 T1 (READ COMMITTED)

(3) **SELECT** * FROM customers:

```
(4) INSERT INTO customers (
  customer id, first name, last name
) VALUES (
  8, 'Steve', 'Button'
);
```

(5) UPDATE customers SET last name = 'Yellow' WHERE customer id = 3;

(6) **COMMIT**; (7) **SELECT** *

FROM customers:

The returned result set contains the new row and the update.

Transaction 2 T2 (SERIALIZABLE)

- (1) SET TRANSACTION ISOLATION LEVEL SERIALIZABLE:
- (2) **SELECT** * FROM customers;

(8) **SELECT** * FROM customers;

The returned result set still doesn't contain the new row or the update made by T1. That's because T2 is SERIALIZABLE.

 TABLE 8-2
 SERIALIZABLE Transactions

Query Flashbacks

If you mistakenly commit changes and you want to view rows as they originally were, you can use a query flashback. You can then use the results of a query flashback to manually change rows back to their original values if you need to.

Query flashbacks can be based on a datetime or system change number (SCN). The database uses SCNs to track changes made to data, and you can use them to flash back to a particular SCN in the database.

Granting the Privilege for Using Flashbacks

Flashbacks use the PL/SQL DBMS_FLASHBACK package, for which you must have the EXECUTE privilege to run. The following example connects as the sys user and grants the EXECUTE privilege on DBMS FLASHBACK to the store user:

CONNECT sys/change_on_install AS sysdba GRANT EXECUTE ON SYS.DBMS FLASHBACK TO store;



NOTE

Speak with your DBA if you are unable to perform these statements. You'll learn about privileges in the next chapter, and you'll learn about PL/SQL packages in Chapter 11.

Time Query Flashbacks

The following example connects as store and retrieves the product_id, name, and price columns for the first five rows from the products table:

CONNECT store/store_password

SELECT product_id, name, price

FROM products

WHERE product id <= 5;

PRODUCT_ID	NAME	PRICE
1	Modern Science	19.95
2	Chemistry	30
3	Supernova	25.99
4	Tank War	13.95
5	Z Files	49.99



NOTE

If you see different prices for any of these products, go ahead and rerun the store schema.sql file.

The next example reduces the price of these rows, commits the change, and retrieves the rows again so you can see the new prices:

```
UPDATE products

SET price = price * 0.75

WHERE product id <= 5;
```

COMMIT:

SELECT product id, name, price FROM products WHERE product id <= 5;

PRODUCT_ID	NAME	PRICE
1	Modern Science	14.96
2	Chemistry	22.5
3	Supernova	19.49
4	Tank War	10.46
5	Z Files	37.49

The following statement executes the DBMS FLASHBACK. ENABLE AT TIME () procedure, which enables you to perform a flashback to a particular datetime; notice the DBMS FLASHBACK. ENABLE AT TIME () procedure accepts a datetime and the example passes SYSDATE - 10 / 1440 to the procedure (this expression evaluates to a datetime ten minutes in the past):

EXECUTE DBMS FLASHBACK.ENABLE AT TIME(SYSDATE - 10 / 1440);



NOTE

24 hours × 60 minutes per hour = 1440 minutes. Therefore SYSDATE - 10 / 1440 is a datetime ten minutes in the past.

Any queries you execute now will display the rows as they were ten minutes ago. Assuming you performed the earlier UPDATE less than ten minutes ago, the following query will display the prices as they were before you updated them:

SELECT product id, name, price FROM products WHERE product id <= 5;

PRODUCT_ID	NAME	PRICE
1	Modern Science	19.95
2	Chemistry	30
3	Supernova	25.99
4	Tank War	13.95
5	Z Files	49.99

To disable a flashback, you execute DBMS FLASHBACK.DISABLE(), as shown in the following example:

EXECUTE DBMS_FLASHBACK.DISABLE();



CAUTION

You must disable a flashback before you can enable it again.

Now when you perform queries, the rows as they currently exist will be retrieved, as shown here:

```
SELECT product_id, name, price
FROM products
WHERE product_id <= 5;
```

PRODUCT_ID	NAME	PRICE
1	Modern Science	14.96
2	Chemistry	22.5
3	Supernova	19.49
4	Tank War	10.46
5	Z Files	37.49

System Change Number Query Flashbacks

Flashbacks based on system change numbers (SCNs) can be more precise than those based on a time, because the database uses SCNs to track changes made to data. To get the current SCN, you can execute <code>DBMS_FLASHBACK.GET_SYSTEM_CHANGE_NUMBER()</code>, as shown in the following example:

VARIABLE current scn NUMBER

The next example adds a row to the products table, commits the change, and retrieves the new row:

The next example executes the following procedure, DBMS FLASHBACK.ENABLE AT SYSTEM CHANGE NUMBER(), which enables you to perform a flashback to an SCN; notice that this procedure accepts an SCN and that the example passes the current scn variable to the procedure:

EXECUTE DBMS FLASHBACK.ENABLE AT SYSTEM CHANGE NUMBER(:current scn);

Any queries you execute now will display the rows as they were at the SCN stored in current scn before you performed the INSERT. The following query attempts to get the row with a product id of 15; it fails because that new row was added after the SCN stored in current scn:

SELECT product id FROM products WHERE product id = 15;

no rows selected

To disable a flashback, you execute DBMS FLASHBACK.DISABLE(), as shown in the following example:

EXECUTE DBMS FLASHBACK.DISABLE();

If you perform the previous query again, you'll see the new row that was added by the INSERT.



NOTE

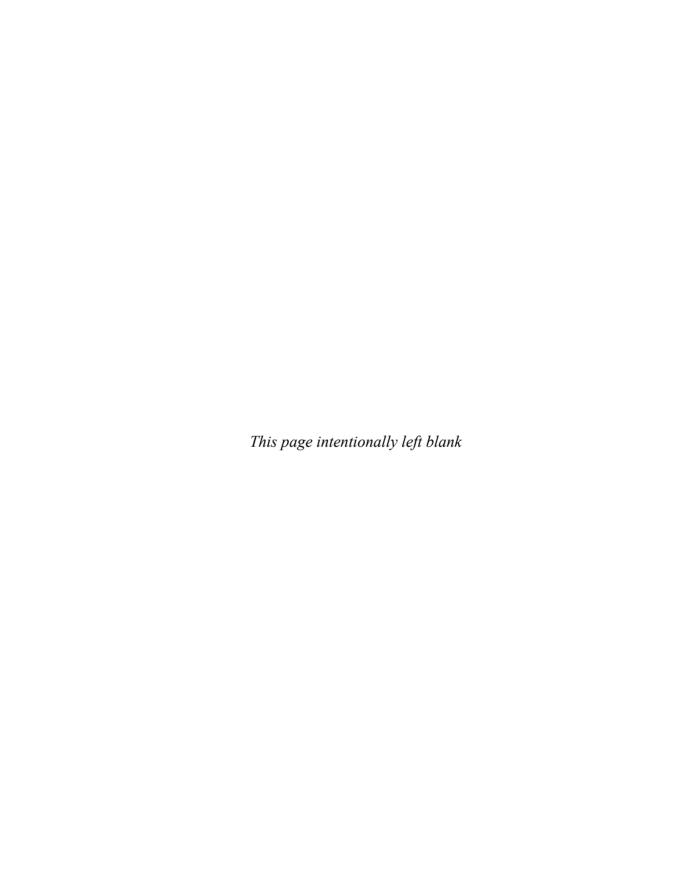
If you followed along with the examples, go ahead and rerun the store schema.sql script to recreate everything. That way, the results of your SQL statements will match mine as you progress through the rest of this book.

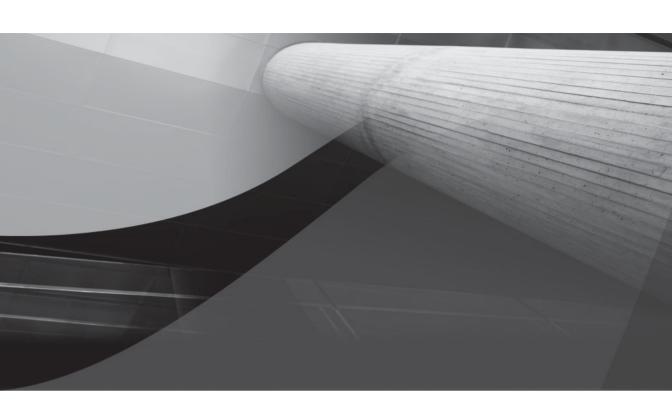
Summary

In this chapter, you have learned the following:

- How to add rows using the INSERT statement.
- How to modify rows using the UPDATE statement.
- How to remove rows using the DELETE statement.
- How the database maintains referential integrity through the enforcement of constraints.
- How to use the DEFAULT keyword to specify default values for columns.
- How to merge rows using the MERGE statement.
- A database transaction is a group of SQL statements that comprise a logical unit of work.
- The Oracle database software can handle multiple concurrent transactions.
- How to use query flashbacks to view rows as they originally were before you made changes to them.

In the next chapter, you'll learn about users, privileges, and roles.





CHAPTER 9

Users, Privileges, and Roles



n this chapter, you will do the following:

- Learn more about users
- See how privileges are used to enable users to perform tasks in the database
- Explore the two types of privileges: system privileges and object privileges
- Learn how system privileges allow you to perform actions such as executing DDL statements
- See how object privileges allow you to perform actions such as executing DML statements
- Explore how to group privileges together into roles
- Learn how to audit the execution of SQL statements



NOTE

You'll need to type in the SQL statements shown in this chapter if you want to follow the examples: The statements are not contained in any script.

Users

In this section, you'll learn how to create a user, alter a user's password, and drop a user. You will see the term "tablespace" used in this chapter. Tablespaces are used by the database to store separate objects, which can include tables, types, PL/SQL code, and so on. Typically, related objects are grouped together and stored in the same tablespace. For example, you might create an order entry application and store all the objects for that application in one tablespace, and you might create a supply chain application and store the objects for that application in a different tablespace. For more details on tablespaces, you should read the *Oracle Database Concepts* manual published by Oracle Corporation.

Creating a User

To create a user in the database, you use the CREATE USER statement. The simplified syntax for the CREATE USER statement is as follows:

CREATE USER user_name IDENTIFIED BY password [DEFAULT TABLESPACE default_tablespace] [TEMPORARY TABLESPACE temporary_tablespace];

where

- user name is the name of the database user.
- password is the password for the database user.
- default_tablespace is the default tablespace where database objects are stored. If you omit a default tablespace, the default SYSTEM tablespace, which always exists in a database, is used.

temporary tablespace is the default tablespace where temporary objects are stored. These objects include temporary tables that you'll learn about in the next chapter. If you omit a temporary tablespace, the default SYSTEM tablespace is used.

The following example connects as system and creates a user named jason with a password of price:

CONNECT system/manager CREATE USER jason IDENTIFIED BY price;



NOTE

If you want to follow along with these examples you'll need to connect to the database as a privileged user. I used the system user in the example, which has a password of manager in my database.

The next example creates a user named henry and specifies a default and temporary tablespace:

CREATE USER henry IDENTIFIED BY hooray DEFAULT TABLESPACE users TEMPORARY TABLESPACE temp:



NOTE

If your database doesn't have tablespaces named users and temp, you can skip this example. The henry user isn't used elsewhere in this book, and I included the example only so you can see how to specify tablespaces for a user. You can view all the tablespaces in a database by connecting as the system user and running the query SELECT tablespace name FROM dba tablespaces.

If you want a user to be able to do things in the database, that user must be granted the necessary permissions to do those things. For example, to connect to the database a user must be granted the permission to create a session, which is the CREATE SESSION system privilege. Permissions are granted by a privileged user (system, for example) using the GRANT statement.

The following example grants the CREATE SESSION permission to jason:

GRANT CREATE SESSION TO jason;

The jason user will now be able to connect to the database.

The following example creates other users that are used in this chapter and grants the CREATE SESSION privilege to those users:

CREATE USER steve IDENTIFIED BY button; CREATE USER gail IDENTIFIED BY seymour; GRANT CREATE SESSION TO steve, gail;

Changing a User's Password

You can change a user's password using the ALTER USER statement. For example, the following statement changes the password for jason to marcus:

ALTER USER jason IDENTIFIED BY marcus;

You can also change the password for the user you're currently logged in as using the PASSWORD command. After you enter PASSWORD, SQL*Plus prompts you to enter the old password and the new password twice for confirmation. The following example connects as jason and executes PASSWORD; notice the password itself is masked using asterisks:

CONNECT jason/marcus PASSWORD

Changing password for JASON Old password: *****
New password: *****
Retype new password: *****
Password changed

Deleting a User

You delete a user using the DROP USER statement. The following example connects as system and uses DROP USER to delete jason:

CONNECT system/manager DROP USER jason;



NOTE

You must add the keyword CASCADE after the user's name in the DROP USER statement if that user's schema contains objects such as tables. However, you should ensure no other users need access to those objects before doing this.

System Privileges

A *system privilege* allows a user to perform certain actions within the database, such as executing DDL statements. For example, CREATE TABLE allows a user to create a table in their schema. Some of the commonly used system privileges are shown in Table 9-1.



NOTE

You can get the full list of system privileges in the Oracle Database SQL Reference manual published by Oracle Corporation.

As you'll see later, privileges can be grouped together into *roles*. Two useful roles to grant to a user are CONNECT and RESOURCE; CONNECT allows a user to connect to the database; RESOURCE allows a user to create various database objects like tables, sequences, PL/SQL code, and so on.

Granting System Privileges to a User

You use GRANT to grant a system privilege to a user. The following example grants some system privileges to steve (assuming you're still connected to the database as system):

GRANT CREATE SESSION, CREATE USER, CREATE TABLE TO steve;

System Privilege	Allows You to
CREATE SESSION	Connect to a database.
CREATE SEQUENCE	Create a sequence, which is a series of numbers that are typically used to automatically populate a primary key column. You'll learn about sequences in the next chapter.
CREATE SYNONYM	Create a synonym. A synonym allows you to reference a table in another schema. You'll learn about synonyms later in this chapter.
CREATE TABLE	Create a table in the user's schema.
CREATE ANY TABLE	Create a table in any schema.
DROP TABLE	Drop a table from the user's schema.
DROP ANY TABLE	Drop a table from any schema.
CREATE PROCEDURE	Create a stored procedure.
EXECUTE ANY PROCEDURE	Execute a procedure in any schema.
CREATE USER	Create a user.
DROP USER	Drop a user.
CREATE VIEW	Create a view. A view is a stored query that allows you to access multiple tables and columns. You may then query the view as you would a table. You'll learn about views in the next chapter.

Commonly Used System Privileges **TABLE 9-1**

You can also use WITH ADMIN OPTION to allow a user to grant a privilege to another user. The following example grants the EXECUTE ANY PROCEDURE privilege with the ADMIN option to steve:

GRANT EXECUTE ANY PROCEDURE TO steve WITH ADMIN OPTION;

EXECUTE ANY PROCEDURE can then be granted to another user by steve. The following example connects as steve and grants EXECUTE ANY PROCEDURE to gail:

CONNECT steve/button GRANT EXECUTE ANY PROCEDURE TO gail;

You can grant a privilege to all users by granting to PUBLIC. The following example connects as system and grants EXECUTE ANY PROCEDURE to PUBLIC:

CONNECT system/manager GRANT EXECUTE ANY PROCEDURE TO PUBLIC;

Every user in the database now has the EXECUTE ANY PROCEDURE privilege.

Checking System Privileges Granted to a User

You can check which system privileges a user has by querying user_sys_privs. Table 9-2 describes some of the columns in user sys privs.



NOTE

user_sys_privs forms part of the Oracle database's data dictionary. The data dictionary stores information about the database itself.

The following example connects as steve and queries user sys privs:

CONNECT steve/button SELECT * FROM user_sys_privs ORDER BY privilege;

USERNAME	PRIVILEGE	ADM
STEVE	CREATE SESSION	NO
STEVE	CREATE TABLE	NO
STEVE	CREATE USER	NO
PUBLIC	EXECUTE ANY PROCEDURE	NO
STEVE	EXECUTE ANY PROCEDURE	YES

The next example connects as gail and queries user sys privs:

CONNECT gail/seymour SELECT * FROM user_sys_privs ORDER BY privilege;

USERNAME	PRIVILEGE	
GAIL	CREATE SESSION	NO
GAIL	EXECUTE ANY PROCEDURE	NO
PUBLIC	EXECUTE ANY PROCEDURE	NO

Notice gail has the EXECUTE ANY PROCEDURE privilege that was granted earlier by steve.

Column	Туре	Description
username	VARCHAR2(30)	Name of the current user
privilege	VARCHAR2 (40)	The system privilege the user has
admin_option	VARCHAR2(3)	Whether the user is able to grant the privilege to another user

Making Use of System Privileges

Once a user has been granted a system privilege, they can use it to perform the specified task. For example, steve has the CREATE USER privilege, so he is able to create a user:

CONNECT steve/button CREATE USER roy IDENTIFIED BY williams;

If steve were to attempt to use a system privilege he doesn't have, the database will return the error ORA-01031: insufficient privileges. For example, steve doesn't have the DROP USER privilege, and in the following example steve attempts to drop roy and fails:

```
SOL> DROP USER roy;
    DROP USER roy
    ERROR at line 1:
    ORA-01031: insufficient privileges
```

Revoking System Privileges from a User

You revoke system privileges from a user using REVOKE. The following example connects as system and revokes the CREATE TABLE privilege from steve:

```
CONNECT system/manager
   REVOKE CREATE TABLE FROM steve;
```

The next example revokes EXECUTE ANY PROCEDURE from steve:

```
REVOKE EXECUTE ANY PROCEDURE FROM steve;
```

When you revoke EXECUTE ANY PROCEDURE from steve—who has already passed on this privilege to gail—then gail still keeps the privilege:

```
CONNECT gail/seymour
   SELECT *
   FROM user_sys_privs
   ORDER BY privilege;
```

USERNAME	PRIVILEGE	
GAIL	CREATE SESSION	NO
GAIL	EXECUTE ANY PROCEDURE	NO
PUBLIC	EXECUTE ANY PROCEDURE	NO

Object Privileges

An object privilege allows a user to perform certain actions on database objects, such as executing DML statements on tables. For example, INSERT ON store.products allows a user to insert rows into the products table of the store schema. Some of the commonly used object privileges are shown in Table 9-3.

Object Privilege	Allows a User to
SELECT	Perform a select.
INSERT	Perform an insert.
UPDATE	Perform an update.
DELETE	Perform a delete.
EXECUTE	Execute a stored procedure.

TABLE 9-3 Commonly Used Object Privileges



NOTE

You can get the full list of system privileges in the Oracle Database SQL Reference manual published by Oracle Corporation.

Granting Object Privileges to a User

You use GRANT to grant an object privilege to a user. The following example connects as store and grants the SELECT, INSERT, and UPDATE object privileges on the products table to steve along with the SELECT privilege on the employees table:

CONNECT store/store_password

GRANT SELECT, INSERT, UPDATE ON store.products TO steve;

GRANT SELECT ON store.employees TO steve;

The next example grants the UPDATE privilege on the last_name and salary columns to steve:

GRANT UPDATE (last name, salary) ON store.employees TO steve;

You can also use the GRANT option to enable a user to grant a privilege to another user. The following example grants the SELECT privilege on the customers table with the GRANT option to steve:

GRANT SELECT ON store.customers TO steve WITH GRANT OPTION;



NOTE

You use the GRANT option to allow a user to grant an object privilege to another user, and you use the ADMIN option to allow a user to grant a system privilege to another user.

The SELECT ON store.customers privilege can then be granted to another user by steve. The following example connects as steve and grants this privilege to gail:

CONNECT steve/button
GRANT SELECT ON store.customers TO gail;

Checking Object Privileges Made

You can check which table object privileges a user has made to other users by querying user tab privs made. Table 9-4 documents the columns in user tab privs made.

The following example connects as store and queries user tab privs made. Because there are so many rows, I'll limit the retrieved rows to those where table name is PRODUCTS:

CONNECT store/store password SELECT * FROM user_tab_privs_made WHERE table name = 'PRODUCTS';

GRANTEE TABLE_NAME			
GRANTOR	PRIVILEGE	GRA	HIE
STEVE STORE	PRODUCTS INSERT	NO	NO
STEVE STORE	PRODUCTS SELECT	NO	NO
STEVE STORE	PRODUCTS UPDATE	NO	NO

You can check which column object privileges a user has made by querying user col privs made. Table 9-5 documents the columns in user col privs made.

Column	Туре	Description
grantee	VARCHAR2(30)	User to whom the privilege was granted
table_name	VARCHAR2(30)	Name of the object (such as a table) on which the privilege was granted
grantor	VARCHAR2(30)	User who granted the privilege
privilege	VARCHAR2(40)	Privilege on the object
grantable	VARCHAR2(3)	Whether the grantee can grant the privilege to another (YES or NO)
hierarchy	VARCHAR2(3)	Whether the privilege forms part of a hierarchy (YES or NO)

 TABLE 9-4
 Some Columns in user_tab_privs_made

Column	Туре	Description
grantee	VARCHAR2(30)	User to whom the privilege was granted
table_name	VARCHAR2(30)	Name of the object on which the privilege was granted
column_name	VARCHAR2(30)	Name of the object on which the privilege was granted
grantor	VARCHAR2(30)	User who granted the privilege
privilege	VARCHAR2 (40)	Privilege on the object
grantable	VARCHAR2(3)	Whether the grantee can grant the privilege to another (YES or NO)

 TABLE 9-5
 Some Columns in user col privs made

The following example queries user col privs made:

SELECT *

FROM user_col_privs_made;

GRANTEE	TABLE_NAM	Ε
COLUMN_NAME	GRANTOR	
PRIVILEGE		GRA
STEVE LAST_NAME UPDATE	EMPLOYEES STORE	NO
STEVE SALARY UPDATE	EMPLOYEES STORE	NO

Checking Object Privileges Received

You can check which object privileges on a table a user has received by querying the user_tab_privs_recd table. Table 9-6 documents the columns in user_tab_privs_recd.

The next example connects as steve and queries user tab privs recd:

Column	Туре	Description
owner	VARCHAR2(30)	User who owns the object
table_name	VARCHAR2 (30)	Name of the object on which the privilege was granted
grantor	VARCHAR2(30)	User who granted the privilege
privilege	VARCHAR2 (40)	Privilege on the object
grantable	VARCHAR2(3)	Whether the grantee can grant the privilege to another (YES or NO)
hierarchy	VARCHAR2(3)	Whether the privilege forms part of a hierarchy (YES or NO)

 TABLE 9-6
 Some Columns in user_tab_privs_recd

CONNECT steve/button SELECT * FROM user_tab_privs_recd ORDER BY privilege;

OWNER	TABLE_NAME		
GRANTOR	PRIVILEGE	GRA	HIE
STORE STORE	PRODUCTS INSERT	NO	NO
STORE STORE	CUSTOMERS SELECT	YES	NO
STORE STORE	EMPLOYEES SELECT	NO	NO
STORE	PRODUCTS SELECT	NO	NO
STORE STORE	PRODUCTS UPDATE	NO	NO

You can check which column object privileges a user has received by querying user_col_ privs_recd. Table 9-7 documents the columns in user_col_privs_recd.

Column	Туре	Description
owner	VARCHAR2(30)	User who owns the object
table_name	VARCHAR2(30)	Name of the table on which the privilege was granted
column_name	VARCHAR2(30)	Name of the column on which the privilege was granted
grantor	VARCHAR2(30)	User who granted the privilege
privilege	VARCHAR2 (40)	Privilege on the object
grantable	VARCHAR2(3)	Whether the grantee can grant the privilege to another (YES or NO)

 TABLE 9-7
 Some Columns in user_col_privs_recd

The following example queries user_col_privs_recd:

SELECT *

FROM user col privs recd;

OWNER	TABLE_NAM	Ξ
COLUMN_NAME	GRANTOR	
PRIVILEGE		GRA
STORE LAST_NAME UPDATE	EMPLOYEES STORE	NO
STORE SALARY UPDATE	EMPLOYEES STORE	NO

Making Use of Object Privileges

Once a user has been granted an object privilege, they can use it to perform the specified task. For example, steve has the SELECT privilege on store.customers:

CONNECT steve/button SELECT * FROM store.customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

If steve were to attempt to retrieve from the purchases table—for which he doesn't have any permissions—the database will return the error ORA-00942: table or view does not exist:

```
SOL> SELECT *
     2 FROM store.purchases;
    FROM store.purchases
   ERROR at line 2:
    ORA-00942: table or view does not exist
```

Synonyms

In the examples in the previous section, you saw that you can access tables in another schema by specifying the schema name followed by the table. For example, when steve retrieved rows from the customers table in the store schema, he performed a query on store.customers. You can avoid having to enter the schema name by creating a synonym for a table, which you do by using the CREATE SYNONYM statement.

Let's take a look at an example. First, connect as system and grant the CREATE SYNONYM system privilege to steve:

```
CONNECT system/manager
    GRANT CREATE SYNONYM TO steve;
```

Next, connect as steve and perform a CREATE SYNONYM statement to create a synonym for the store.customers table:

```
CONNECT steve/button
   CREATE SYNONYM customers FOR store.customers;
```

To retrieve rows from store.customers, all steve has to do is to reference the customers synonym in the FROM clause of a SELECT statement. For example:

SELECT * FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

Public Synonyms

You can also create a *public* synonym for a table. When you do this, all users see the synonym. The following tasks

- Connect as system
- Grant the CREATE PUBLIC SYNONYM system privilege to store

- Connect as store
- Create a public synonym named products for store.products

are performed by the following statements:

```
CONNECT system/manager

GRANT CREATE PUBLIC SYNONYM TO store;

CONNECT store/store_password

CREATE PUBLIC SYNONYM products FOR store.products;
```

If you connect as steve, who has the SELECT privilege on store.products, you can now retrieve rows from store.products through the products public synonym:

```
CONNECT steve/button
SELECT *
FROM products;
```

Even though a public synonym has been created for store.products, a user still needs object privileges on that table to actually access the table. For example, gail can see the products public synonym, but gail doesn't have any object privileges on store.products. Therefore, if gail attempts to retrieve rows from products, the database returns the error ORA-00942: table or view does not exist:

```
SQL> CONNECT gail/seymour
Connected.
SQL> SELECT * FROM products;
SELECT * FROM products

*
ERROR at line 1:
ORA-00942: table or view does not exist
```

If gail had the SELECT object privilege on the store.products table, the previous SELECT would succeed.

If a user has other object privileges, that user can exercise those object privileges through a synonym. For example, if gail had the INSERT object privilege on the store.products table, gail would be able to add a row to store.products through the products synonym.

Revoking Object Privileges

You revoke object privileges using REVOKE. The following example connects as store and revokes the INSERT privilege on the products table from steve:

```
CONNECT store/store_password

REVOKE INSERT ON products FROM steve;
```

The next example revokes the SELECT privilege on the customers table from steve:

```
REVOKE SELECT ON store.customers FROM steve;
```

When you revoke SELECT ON store.customers from steve—who has already passed on this privilege to gail—gail also loses the privilege.

Roles

A role is a group of privileges that you can assign to a user or to another role. The following points summarize the benefits and features of roles:

- Rather than assigning privileges one at a time directly to a user, you can create a role, assign privileges to that role, and then grant that role to multiple users and roles.
- When you add or delete a privilege from a role, all users and roles assigned that role automatically receive or lose that privilege.
- You can assign multiple roles to a user or role.
- You can assign a password to a role.

As you can see from these points, roles can help you manage multiple privileges assigned to multiple users.

Creating Roles

To create a role, you must have the CREATE ROLE system privilege. As you'll see in a later example, the store user also needs the ability to grant the CREATE USER system privilege with the ADMIN option. The following example connects as system and grants the required privileges to store:

```
CONNECT system/manager
   GRANT CREATE ROLE TO store;
   GRANT CREATE USER TO store WITH ADMIN OPTION;
```

Table 9-8 shows the roles you'll create shortly.

You create a role using the CREATE ROLE statement. The following statements connect as store and create the three roles shown in Table 9-8:

```
CONNECT store/store password
   CREATE ROLE product manager;
   CREATE ROLE hr manager;
   CREATE ROLE overall manager IDENTIFIED by manager password;
```

Notice overall manager has a password of manager password.

Role Name	Has Permissions to
product_manager	Perform SELECT, INSERT, UPDATE, and DELETE operations on the product_types and products tables.
hr_manager	Perform SELECT, INSERT, UPDATE, and DELETE operations on the salary_grades and employees tables. Also, hr_manager is able to create users.
overall_manager	Perform SELECT, INSERT, UPDATE, and DELETE operations on all the tables shown in the previous roles; overall_manager will be granted the previous roles.

Granting Privileges to Roles

You grant privileges to a role using the GRANT statement. You can grant both system and object privileges to a role as well as grant another role to a role. The following example grants the required privileges to the product_manager and hr_manager roles and grants these two roles to overall_manager:

```
GRANT SELECT, INSERT, UPDATE, DELETE ON product_types TO product_manager;
GRANT SELECT, INSERT, UPDATE, DELETE ON products TO product_manager;
GRANT SELECT, INSERT, UPDATE, DELETE ON salary_grades TO hr_manager;
GRANT SELECT, INSERT, UPDATE, DELETE ON employees TO hr_manager;
GRANT CREATE USER TO hr_manager;
GRANT product manager, hr manager TO overall manager;
```

Granting Roles to a User

You grant a role to a user using GRANT. The following example grants the overall_manager role to steve:

GRANT overall manager TO steve;

Checking Roles Granted to a User

You can check which roles have been granted to a user by querying user_role_privs. Table 9-9 defines the columns in user role privs.

The following example connects as steve and queries user_role_privs:

```
CONNECT steve/button

SELECT *

FROM user role privs;
```

USERNAME	GRANTED_ROLE	ADM	DEF	os_
STEVE	OVERALL_MANAGER	NO	YES	NO

Column	Туре	Description
username	VARCHAR2(30)	Name of the user to whom the role has been granted
granted_role	VARCHAR2(30)	Name of the role granted to the user
admin_option	VARCHAR2(3)	Whether the user is able to grant the role to another user or role (YES or NO)
default_role	VARCHAR2(3)	Whether the role is enabled by default when the user connects to the database (YES or NO)
os_granted	VARCHAR2(3)	Whether the role was granted by the operating system (YES or NO)

A user who creates a role is also granted that role by default. The following example connects as store and queries user role privs:

CONNECT store/store password SELECT * FROM user role privs;

USERNAME	GRANTED_ROLE	ADM	DEF	os_
STORE	CONNECT	NO	YES	NO
STORE	HR_MANAGER	YES	YES	NO
STORE	OVERALL_MANAGER	YES	YES	NO
STORE	PRODUCT_MANAGER	YES	YES	NO
STORE	RESOURCE	NO	YES	NO

Notice store has the roles CONNECT and RESOURCE in addition to the roles store created earlier.



NOTE

CONNECT and RESOURCE are built-in roles that were granted to store when you ran the store schema.sql script. As you'll see in the next section, the CONNECT and RESOURCE roles contain multiple privileges.

Checking System Privileges Granted to a Role

You can check which system privileges have been granted to a role by querying role sys privs. Table 9-10 defines the columns in role sys privs.

The following example retrieves the rows from role sys privs (assuming you're still connected as store):

SELECT * FROM role sys privs ORDER BY privilege;

ROLE	PRIVILE	EGE	ADM
RESOURCE	CREATE	CLUSTER	NO
RESOURCE	CREATE	INDEXTYPE	NO
RESOURCE	CREATE	OPERATOR	NO
RESOURCE	CREATE	PROCEDURE	NO
RESOURCE	CREATE	SEQUENCE	NO
CONNECT	CREATE	SESSION	NO
RESOURCE	CREATE	TABLE	NO
RESOURCE	CREATE	TRIGGER	NO
RESOURCE	CREATE	TYPE	NO
HR MANAGER	CREATE	USER	NO

Notice that the RESOURCE role has many privileges assigned to it.



NOTE

The previous query was run using Oracle Database 11g. If you are using a different version of the database software, you may get slightly different results.

Column	Туре	Description
role	VARCHAR2(30)	Name of the role
privilege	VARCHAR2(40)	System privilege granted to the role
admin_option	VARCHAR2(3)	Whether the privilege was granted with the ADMIN option (YES or NO)

 TABLE 9-10
 Some Columns in role_sys_privs

Checking Object Privileges Granted to a Role

You can check which object privileges have been granted to a role by querying role_tab_privs. Table 9-11 defines the columns in role_tab_privs.

The following example queries role tab privs where role equals HR MANAGER:

SELECT *

FROM role_tab_privs

WHERE role='HR_MANAGER'

ORDER BY table name;

ROLE	OWNER
TABLE_NAME	COLUMN_NAME
PRIVILEGE	GRA
HR_MANAGER EMPLOYEES DELETE	STORE
HR_MANAGER EMPLOYEES INSERT	STORE
HR_MANAGER EMPLOYEES SELECT	STORE
HR_MANAGER EMPLOYEES UPDATE	STORE
HR_MANAGER SALARY_GRADES DELETE	STORE
HR_MANAGER SALARY_GRADES INSERT	STORE

HR_MANAGER	STORE	
SALARY_GRADES		
SELECT		NO
HR_MANAGER	STORE	
SALARY_GRADES		
UPDATE		NO

Making Use of Privileges Granted to a Role

Once a user has been granted a privilege via a role, they can use that privilege to perform the authorized tasks. For example, steve has the overall manager role. The overall manager was granted the product manager and hr manager roles. The product manager was granted the SELECT object privilege on the products and product types tables. Therefore, steve is able to retrieve rows from these tables, as shown in the following example:

```
CONNECT steve/button
   SELECT p.name, pt.name
   FROM store.products p, store.product types pt
   WHERE p.product type id = pt.product type id;
```

NAME	NAME
Modern Science Chemistry Supernova Tank War Z Files 2412: The Return Space Force 9 From Another Planet Classical Music Pop 3	Book Book Video Video Video Video DVD DVD CD CD
Creative Yell	CD

Column	Туре	Description
role	VARCHAR2(30)	User to whom the privilege was granted
owner	VARCHAR2(30)	User who owns the object
table_name	VARCHAR2(30)	Name of the object on which the privilege was granted
column_name	VARCHAR2(30)	Name of the column (if applicable)
privilege	VARCHAR2(40)	Privilege on the object
grantable	VARCHAR2(3)	Whether the privilege was granted with the ${\tt GRANT}$ option (YES or NO)

Default Roles

By default, when a role is granted to a user, that role is enabled for that user. This means that when the user connects to the database, the role is automatically available to them. To enhance security, you can disable a role by default; when the user connects, they will have to enable the role themselves before they can use it. If the role has a password, the user must enter that password before the role is enabled. For example, the <code>overall_manager</code> role has a password of <code>manager_passsword</code>, and <code>overall_manager</code> is granted to <code>steve</code>. In the example you'll see next, you'll disable <code>overall_manager</code> so that <code>steve</code> has to enable this role and enter the password before he can use it. You do this by altering a role so that it is no longer a default role using the <code>ALTER_ROLE</code> statement. The following example connects as <code>system</code> and alters <code>steve</code> so that <code>overall_manager</code> is no longer a default role:

CONNECT system/manager

ALTER USER steve DEFAULT ROLE ALL EXCEPT overall manager;

When you connect as steve, you need to enable overall manager using SET ROLE:

CONNECT steve/button

SET ROLE overall manager IDENTIFIED BY manager password;

Once you've set the role, you can use the privileges granted to that role. You can set your role to "none" (i.e. no role) using the following statement:

SET ROLE NONE;

You can also set your role to "all roles" except overall_manager using the following statement:

```
SET ROLE ALL EXCEPT overall manager;
```

By assigning passwords to roles and setting roles to not be enabled by default for a user, you introduce an additional level of security.

Revoking a Role

You revoke a role using REVOKE. The following example connects as store and revokes the overall manager role from steve:

CONNECT store/store_password REVOKE overall manager FROM steve;

Revoking Privileges from a Role

You revoke a privilege from a role using REVOKE. The following example connects as store and revokes all privileges on the products and product_types tables from product_manager (assuming you're still connected as store):

```
REVOKE ALL ON products FROM product_manager;
REVOKE ALL ON product_types FROM product_manager;
```

Dropping a Role

You drop a role using DROP ROLE. The following example drops the overall manager, product manager, and hr manager roles (assuming you're still connected as store):

```
DROP ROLE overall manager;
    DROP ROLE product manager;
    DROP ROLE hr manager;
```

Auditing

The Oracle database software contains auditing capabilities that enable you to keep track of database operations. Some operations may be audited at a high level, such as failed attempts to log into the database, while others may be audited at a detailed level, such as when a user retrieved rows from a specific table. Typically, your database administrator will be responsible for enabling auditing and monitoring the output for security violations. In this section, you will see some simple examples of auditing, which is performed using the AUDIT statement.

Privileges Required to Perform Auditing

Before a user can issue AUDIT statements, that user must have been granted certain privileges:

- For auditing high-level operations, the user must have the AUDIT SYSTEM privilege. An example of a high-level operation is the issuance of any SELECT statement, regardless of the table involved.
- For tracking operations on specific database objects, the user must either have the AUDIT ANY privilege or the database object must be in their schema. An example of specific database object operation is the issuance of a SELECT statement for a particular table.

The following example connects to the database as the system user and grants the AUDIT SYSTEM and AUDIT ANY privileges to the store user:

```
CONNECT system/manager
    GRANT AUDIT SYSTEM TO store:
    GRANT AUDIT ANY TO store;
```

Auditing Examples

The following example connects to the database as the store user and audits the issuance of CREATE TABLE statements:

```
CONNECT store/store password
   AUDIT CREATE TABLE;
```

As a result of this AUDIT statement, any CREATE TABLE statements issued will be audited; for example, the following statement creates a simple test table:

```
CREATE TABLE test (
     id INTEGER
   );
```

You can view the audit trail of information for the user you are currently logged in as through the USER_AUDIT_TRAIL view. The following example shows the audit record generated by the previous CREATE TABLE statement:

SELECT username, extended_timestamp, audit_option FROM user_audit_trail WHERE audit option='CREATE TABLE';

You may also audit the issuance of statements by a particular user. The following example audits all SELECT statements issued by the store user:

AUDIT SELECT TABLE BY store;

The next example audits all INSERT, UPDATE, and DELETE statements made by the store and steve users:

AUDIT INSERT TABLE, UPDATE TABLE, DELETE TABLE BY store, steve;

You may also audit the issuance of statements made for a particular database object. The following example audits all SELECT statements issued for the products table:

AUDIT SELECT ON store.products;

The next example audits all statements issued for the employees table:

AUDIT ALL ON store.employees;

You may also use the WHENEVER SUCCESSFUL and WHENEVER NOT SUCCESSFUL options to indicate when auditing should be performed. WHENEVER SUCCESSFUL indicates auditing will be performed when the statement executed successfully. WHENEVER NOT SUCCESSFUL indicates auditing will be performed when the statement did not execute successfully. The default is to do both, that is, audit regardless of success. The following examples use the WHENEVER NOT SUCCESSFUL option:

AUDIT UPDATE TABLE BY steve WHENEVER NOT SUCCESSFUL; AUDIT INSERT TABLE WHENEVER NOT SUCCESSFUL;

The next example uses the WHENEVER SUCCESSFUL option to audit the creation and deletion of a user:

AUDIT CREATE USER, DROP USER WHENEVER SUCCESSFUL;

The next example uses the WHENEVER SUCCESSFUL option to audit the creation and deletion of a user by the store user:

AUDIT CREATE USER, DROP USER BY store WHENEVER SUCCESSFUL;

You may also use the BY SESSION and BY ACCESS options. The BY SESSION option causes only one audit record to be logged when the same type of statement is issued during the same user database session; a database session starts when the user logs into the database and ends when the user logs out. The BY ACCESS option causes one audit record to be logged every time the same type of statement is issued, regardless of the user session. The following examples show the use of the BY SESSION and BY ACCESS options:

```
AUDIT SELECT ON store.products BY SESSION;
    AUDIT DELETE ON store.employees BY ACCESS;
    AUDIT INSERT, UPDATE ON store.employees BY ACCESS;
```

Audit Trail Views

Earlier, you saw the use of the USER AUDIT TRAIL view. This and the other audit trail views are outlined in the following list:

- **USER AUDIT OBJECT** displays the audit records for all objects accessible to the current user.
- **USER AUDIT SESSION** displays the audit records for connections and disconnections of the current user.
- **USER AUDIT STATEMENT** displays the audit records for GRANT, REVOKE, AUDIT, NOAUDIT, and ALTER SYSTEM statements issued by the current user.
- **USER AUDIT TRAIL** displays all audit trail entries related to the current user.

You may use these views to examine the contents of the audit trail. There are a number of similarly named views that the database administrator may use to examine the audit trail; these views are named DBA AUDIT OBJECT, DBA AUDIT SESSION, DBA AUDIT STATEMENT, DBA AUDIT TRAIL, plus others. These views allow the DBA to view audit records across all users. For more details on these views, you should consult the Oracle Database Reference manual published by Oracle Corporation.

Summary

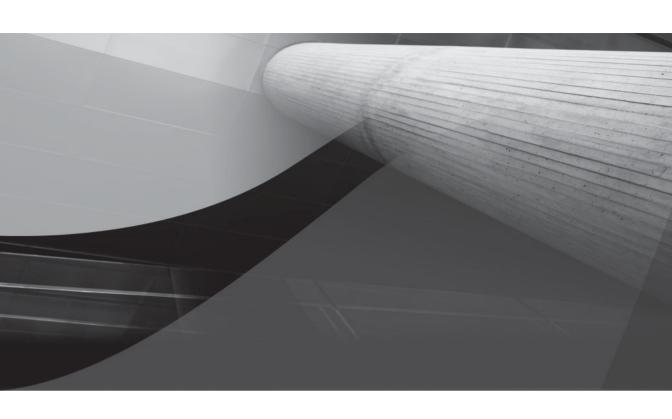
In this chapter, you've learned the following:

- A user is created using the CREATE USER statement.
- System privileges allow you to perform certain actions within the database, such as executing DDL statements.

298 Oracle Database 11*g* SQL

- Object privileges allow you to perform certain actions on database objects, such as executing DML statements on tables.
- You can avoid having to enter the schema name by creating a synonym for a table.
- A role is a group of privileges that you can assign to a user or another role.
- Auditing the execution of SQL statements can be performed using the AUDIT statement.

In the next chapter, you'll learn more about creating tables and see how to create indexes, sequences, and views.



CHAPTER 10

Creating Tables, Sequences, Indexes, and Views



n this chapter, you will do the following:

- Learn more about tables
- See how to create and use sequences, which generate a series of numbers
- Explore how to create and use indexes, which can improve the performance of queries
- Learn how to create and use views, which are predefined queries that allow you to hide complexity from users, among other benefits
- Examine flashback data archives, new for Oracle Database 11g, which store changes made to a table over a period of time

Let's plunge in and examine tables.

Tables

In this section, you'll learn more about creating a table. You'll see how to modify and drop a table as well as how to retrieve information about a table from the data dictionary. The data dictionary contains information about all the database items, such as tables, sequences, indexes, and so on.

Creating a Table

You use the CREATE TABLE statement to create a table. The simplified syntax for the CREATE TABLE statement is as follows:

where

- GLOBAL TEMPORARY means the table's rows are temporary (these tables are known as temporary tables). The rows in a temporary table are specific to a user session, and how long the rows persist is set in the ON COMMIT clause.
- table name is the name of the table.
- column_name is the name of a column.
- \blacksquare type is the type of a column.
- constraint_def is a constraint on a column.
- default exp is an expression to assign a default value to a column.

- ON COMMIT controls the duration of the rows in a temporary table. DELETE means the rows are deleted at the end of a transaction. PRESERVE means the rows are kept until the end of a user session, at which point the rows are deleted. If you omit ON COMMIT for a temporary table, then the default DELETE is used.
- tab space is the tablespace for the table. If you omit a tablespace, then the table is stored in the user's default tablespace.



NOTE

The full CREATE TABLE syntax is far richer than that shown above. For full details, see the Oracle Database SQL Reference book published by Oracle Corporation.

The following example connects as the store user and creates a table named order status2:

```
CONNECT store/store password
   CREATE TABLE order status2 (
     id INTEGER CONSTRAINT order status2 pk PRIMARY KEY,
     status VARCHAR2(10),
     last modified DATE DEFAULT SYSDATE
   );
```



If you want to follow along with the examples in this chapter, you'll need to enter and run the SQL statements using SQL*Plus.

The next example creates a temporary table named order status temp whose rows will be kept until the end of a user session (ON COMMIT PRESERVE ROWS):

```
CREATE GLOBAL TEMPORARY TABLE order status temp (
     id INTEGER,
     status VARCHAR2(10),
     last modified DATE DEFAULT SYSDATE
   ON COMMIT PRESERVE ROWS;
```

The next example performs the following:

- Adds a row to order status temp.
- Disconnects from the database to end the session, which causes the row in order status temp to be deleted.
- Reconnects as store and queries order status temp, which shows there are no rows in this table.

```
INSERT INTO order_status_temp (
   id, status
 ) VALUES (
   1, 'New'
 );
```

```
1 row created.
DISCONNECT
CONNECT store/store_password
SELECT *
FROM order_status_temp;
```

no rows selected

Getting Information on Tables

You can get information about your tables by

- Performing a DESCRIBE command on the table. You've already seen examples that use the DESCRIBE command in earlier chapters.
- Querying the user_tables view, which forms part of the data dictionary.

Table 10-1 describes some of the columns in the user tables view.



NOTE

You can retrieve information on all the tables you have access to by querying the all_tables view.

The following example retrieves some of the columns from user_tables where the table_name is order_status2 or order_status_temp:

```
SELECT table_name, tablespace_name, temporary
FROM user_tables
WHERE table_name IN ('ORDER_STATUS2', 'ORDER_STATUS_TEMP');
```

TABLE_NAME	TABLESPACE_NAME	Τ
		-
ORDER_STATUS2	USERS	N
ORDER_STATUS_TEMP		Υ

Notice the order_status_temp table is temporary, as indicated by the Y in the last column.

Column	Туре	Description
table_name	VARCHAR2(30)	Name of the table.
tablespace_name	VARCHAR2(30)	Name of the tablespace in which the table is stored. A tablespace is an area used by the database to store objects such as tables.
temporary	VARCHAR2(1)	Whether the table is temporary. This is set to Y if temporary or N if not temporary.

Getting Information on Columns in Tables

You can retrieve information about the columns in your tables from the user tab columns view. Table 10-2 describes some of the columns in user tab columns.



NOTE

You can retrieve information on all the columns in tables you have access to by querying the all tab columns view.

The following example retrieves some of the columns from user tab columns for the products table:

COLUMN column name FORMAT a15 COLUMN data_type FORMAT a10 SELECT column name, data type, data length, data precision, data scale FROM user tab columns WHERE table name = 'PRODUCTS';

COLUMN_NAME	DATA_TYPE	DATA_LENGTH	DATA_PRECISION	DATA_SCALE
PRODUCT_ID	NUMBER	22	38	0
PRODUCT_TYPE_ID	NUMBER	22	38	0
NAME	VARCHAR2	30		
DESCRIPTION	VARCHAR2	50		
PRICE	NUMBER	22	5	2

Altering a Table

You alter a table using the ALTER TABLE statement. You can use ALTER TABLE to perform tasks such as

- Adding, modifying, or dropping a column
- Adding or dropping a constraint
- Enabling or disabling a constraint

In the following sections, you'll learn how to use ALTER TABLE to perform each of these tasks.

Column	Туре	Description
table_name	VARCHAR2(30)	Name of the table
column_name	VARCHAR2(30)	Name of the column
data_type	VARCHAR2(106)	Data type of the column
data_length	NUMBER	Length of the data
data_precision	NUMBER	Precision of a numeric column if a precision was specified for the column.
data_scale	NUMBER	Scale of a numeric column

 TABLE 10-2
 Some Columns in the user_tab_columns View

Adding a Column

The following example uses ALTER TABLE to add an INTEGER column named modified_by to the order status2 table:

ALTER TABLE order_status2 ADD modified by INTEGER;

The next example adds a column named initially created to order status2:

ALTER TABLE order_status2 ADD initially created DATE DEFAULT SYSDATE NOT NULL;

You can verify the addition of the new column by executing a DESCRIBE command on order status2:

DESCRIBE order status2

Name	Nul	L?	Type
ID	NOT	NULL	NUMBER (38)
STATUS			VARCHAR2(10)
LAST_MODIFIED			DATE
MODIFIED_BY			NUMBER (38)
INITIALLY_CREATED	NOT	NULL	DATE

Adding a Virtual Column

In Oracle Database 11g, you can add a virtual column, which is a column that refers only to other columns already in the table. For example, the following ALTER TABLE statement adds a virtual column named average_salary to the salary_grades table:

```
ALTER TABLE salary_grades

ADD (average_salary AS ((low_salary + high_salary)/2));
```

Notice average_salary is set to the average of the low_salary and high_salary values. The following DESCRIBE command confirms the addition of the average_salary column to the salary grades table:

DESCRIBE salary grades

Name	Null	L?	Type
SALARY_GRADE_ID	NOT	NULL	NUMBER (38)
LOW_SALARY			NUMBER(6)
HIGH_SALARY			NUMBER(6)
AVERAGE_SALARY			NUMBER

The following query retrieves the rows from the salary_grades table:

SELECT *

FROM salary_grades;

SALARY_GRADE_ID	LOW_SALARY	HIGH_SALARY	AVERAGE_SALARY
1	1	250000	125000.5
2	250001	500000	375000.5

3	500001	750000	625000.5
4	750001	999999	875000

Modifying a Column

The following list shows some of the column aspects you can modify using ALTER TABLE:

- Change the size of a column (if the data type is one whose length may be changed, such as CHAR or VARCHAR2)
- Change the precision of a numeric column
- Change the data type of a column
- Change the default value of a column

You'll see examples of how to change these column aspects in the following sections.

Changing the Size of a Column

The following ALTER TABLE statement increases the maximum length of the order status2 .status column to 15 characters:

ALTER TABLE order status2 MODIFY status VARCHAR2(15);



CAUTION

You can only decrease the length of a column if there are no rows in the table or all the rows contain null values for that column.

Changing the Precision of a Numeric Column

The following ALTER TABLE statement changes the precision of the order status2.id column to 5:

ALTER TABLE order status2 MODIFY id NUMBER(5);



CAUTION

You can only decrease the precision of a numeric column if there are no rows in the table or the column contains null values.

Changing the Data Type of a Column

The following ALTER TABLE statement changes the data type of the order status2.status column to CHAR:

ALTER TABLE order status2 MODIFY status CHAR(15);

> If the table is empty or the column contains null values, you can change the column to any data type (including a data type that is shorter); otherwise, you can change the data type of a column only to a compatible data type. For example, you can change a VARCHAR2 to CHAR (and vice versa) as long as you don't make the column shorter; you cannot change a DATE to a NUMBER.

Changing the Default Value of a Column

The following ALTER TABLE statement changes the default value for the order_status2 .last modified column to SYSDATE - 1:

ALTER TABLE order_status2

MODIFY last modified DEFAULT SYSDATE - 1;

The default value applies only to new rows added to the table. New rows will get their last_modified column set to the current date minus one day.

Dropping a Column

The following ALTER TABLE statement drops the order_status2.initially_created column:

ALTER TABLE order_status2
DROP COLUMN initially_created;

Adding a Constraint

In earlier chapters, you've seen examples of tables with PRIMARY KEY, FOREIGN KEY, and NOT NULL constraints. These constraints, along with the other types of constraints, are summarized in Table 10-3.

You'll see how to add some of the constraints shown in Table 10-3 in the following sections.

Constraint	Constraint Type	Meaning
CHECK	С	The value for a column, or group of columns, must satisfy a certain condition.
NOT NULL	С	The column cannot store a null value. This is actually enforced as a CHECK constraint.
PRIMARY KEY	P	The primary key of a table. A primary key is made up of one or more columns that uniquely identify each row in a table.
FOREIGN KEY	R	A foreign key for a table. A foreign key references a column in another table or a column in the same table (known as a self-reference).
UNIQUE	U	The column, or group of columns, can store only unique values.
CHECK OPTION	V	Changes to the table rows made through a view must pass a check first. (You'll learn about this later in the section "Views.")
READ ONLY	0	The view may only be read from. (You'll learn about this later in the section "Views.")

Adding a CHECK Constraint

The following ALTER TABLE statement adds a CHECK constraint to the order status 2 table:

```
ALTER TABLE order status2
   ADD CONSTRAINT order status2 status ck
   CHECK (status IN ('PLACED', 'PENDING', 'SHIPPED'));
```

This constraint ensures the status column is always set to PLACED, PENDING, or SHIPPED. The following INSERT adds a row to the order status2 table (status is set to PENDING):

```
INSERT INTO order status2 (
      id, status, last modified, modified by
    ) VALUES (
      1, 'PENDING', '01-JAN-2005', 1
    );
```

If you attempt to add a row that doesn't satisfy the CHECK constraint, the database returns the error ORA-02290. For example, the following INSERT attempts to add a row whose status is not in the list:

```
INSERT INTO order status2 (
     id, status, last_modified, modified_by
   ) VALUES (
     2, 'CLEARED', '01-JAN-2005', 2
   INSERT INTO order status2 (
   ERROR at line 1:
   ORA-02290: check constraint (STORE.ORDER STATUS2 STATUS CK) violated
```

Because the CHECK constraint is violated, the database rejects the new row.

You can use other comparison operators with a CHECK constraint. The next example adds a CHECK constraint that enforces that the id value is greater than zero:

```
ALTER TABLE order status2
   ADD CONSTRAINT order status2 id ck CHECK (id > 0);
```

When adding a constraint, the existing rows in the table must satisfy the constraint. For example, if the order status2 table had rows in it, then the id column for the rows would need to be greater than zero.



NOTE

There are exceptions to the rule requiring that existing rows satisfy the constraint. You can disable a constraint when you initially add it, and you can set a constraint to apply only to new data, by specifying ENABLE NOVALIDATE. You'll learn more about this later.

Adding a NOT NULL Constraint

The following ALTER TABLE statement adds a NOT NULL constraint to the status column of the order_status2 table:

```
ALTER TABLE order_status2
  MODIFY status CONSTRAINT order status2 status nn NOT NULL;
```

Notice that you use MODIFY to add a NOT NULL constraint rather than ADD CONSTRAINT. The next example adds a NOT NULL constraint to the modified by column:

ALTER TABLE order_status2

MODIFY modified_by CONSTRAINT order_status2_modified_by_nn NOT NULL;

The following example adds a NOT NULL constraint to the last modified column:

ALTER TABLE order_status2
MODIFY last modified NOT NULL;

Notice that I didn't supply a name for this constraint. In this case, the database automatically assigns an unfriendly name to the constraint, like SYS C003381.



TIP

Always specify a meaningful name to your constraints. That way, when a constraint error occurs, you can easily identify the problem.

Adding a FOREIGN KEY Constraint

Before you see an example of adding a FOREIGN KEY constraint, the following ALTER TABLE statement drops the order status2.modified by column:

ALTER TABLE order_status2 DROP COLUMN modified by;

The next statement adds a FOREIGN KEY constraint that references the employees .employee id column:

ALTER TABLE order_status2

ADD CONSTRAINT order_status2_modified_by_fk

modified by REFERENCES employees(employee id);

You use the ON DELETE CASCADE clause with a FOREIGN KEY constraint to specify that when a row in the parent table is deleted, any matching rows in the child table are also deleted. The following example drops the modified_by column and rewrites the previous example to include the ON DELETE CASCADE clause:

ALTER TABLE order_status2 DROP COLUMN modified_by;

ALTER TABLE order_status2
ADD CONSTRAINT order_status2_modified_by_fk
modified by REFERENCES employees(employee id) ON DELETE CASCADE;

When a row is deleted from the employees table, any matching rows in order_status2 are also deleted.

You use the ON DELETE SET NULL clause with a FOREIGN KEY constraint to specify that when a row in the parent table is deleted, the foreign key column for the row (or rows) in the child table is set to null. The following example drops the modified_by column from order_status2 and rewrites the previous example to include the ON DELETE SET NULL clause:

ALTER TABLE order_status2 DROP COLUMN modified_by;

```
ALTER TABLE order status2
ADD CONSTRAINT order status2 modified by fk
modified by REFERENCES employees (employee id) ON DELETE SET NULL;
```

When a row is deleted from the employees table, the modified by column for any matching rows in order status2 is set to null.

To clean up before moving onto the next section, the following statement drops the modified by column:

```
ALTER TABLE order status2
    DROP COLUMN modified by;
```

Adding a UNIQUE Constraint

The following ALTER TABLE statement adds a UNIQUE constraint to the order status2 .status column:

```
ALTER TABLE order status2
    ADD CONSTRAINT order status2 status uq UNIQUE (status);
```

Any existing or new rows must always have a unique value in the status column.

Dropping a Constraint

You drop a constraint using the DROP CONSTRAINT clause of ALTER TABLE. The following example drops the order status2 status uq constraint:

```
ALTER TABLE order status2
    DROP CONSTRAINT order status2 status uq;
```

Disabling a Constraint

By default, a constraint is enabled when you create it. You can initially disable a constraint by adding DISABLE to the end of the CONSTRAINT clause. The following example adds a constraint to order status2, but also disables it:

```
ALTER TABLE order status2
    ADD CONSTRAINT order status2 status uq UNIQUE (status) DISABLE;
```

You can disable an existing constraint using the DISABLE CONSTRAINT clause of ALTER TABLE. The following example disables the order status2 status nn constraint:

```
ALTER TABLE order status2
    DISABLE CONSTRAINT order status2 status nn;
```

You can add CASCADE after DISABLE CONSTRAINT to disable all constraints that depend on the specified constraint. You use CASCADE when disabling a primary key or unique constraint that is part of a foreign key constraint of another table.

Enabling a Constraint

You can enable an existing constraint using the ENABLE CONSTRAINT clause of ALTER TABLE. The following example enables the order status2 status ug constraint:

```
ALTER TABLE order status2
    ENABLE CONSTRAINT order status2 status uq;
```

To enable a constraint, all the rows in the table must satisfy the constraint. For example, if the order_status2 table contained rows, then the status column would have to contain unique values.

You can apply a constraint to new data only by specifying ENABLE NOVALIDATE; for example:

ALTER TABLE order_status2

ENABLE NOVALIDATE CONSTRAINT order_status2_status_uq;



NOTE

The default is ENABLE VALIDATE, which means existing rows must pass the constraint check.

Deferred Constraints

A deferred constraint is one that is enforced when a transaction is committed; you use the DEFERRABLE clause when you initially add the constraint. Once you've added a constraint, you cannot change it to DEFERRABLE; instead, you must drop and re-create the constraint.

When you add a DEFERRABLE constraint, you can mark it as INITIALLY IMMEDIATE or INITIALLY DEFERRED. Marking as INITIALLY IMMEDIATE means that the constraint is checked whenever you add, update, or delete rows from a table (this is the same as the default behavior of a constraint). INITIALLY DEFERRED means that the constraint is only checked when a transaction is committed. Let's take a look at an example.

The following statement drops the order_status2_status_uq constraint:

ALTER TABLE order_status2

DROP CONSTRAINT order status2 status uq;

The next example adds the order_status2_status_uq constraint, setting it to DEFERRABLE INITIALLY DEFERRED:

ALTER TABLE order_status2

ADD CONSTRAINT order_status2_status_uq UNIQUE (status)

DEFERRABLE INITIALLY DEFERRED:

If you add rows to order_status2, the order_status2_status_uq constraint isn't enforced until you perform a commit.

Getting Information on Constraints

You can retrieve information on your constraints by querying the user_constraints view. Table 10-4 describes some of the columns in user constraints.



NOTE

You can retrieve information on all the constraints you have access to by querying the all constraints view.

The following example retrieves some of the columns from user_constraints for the order status2 table:

SELECT constraint_name, constraint_type, status, deferrable, deferred FROM user constraints WHERE table name = 'ORDER STATUS2';

CONSTRAINT_NAME	С	STATUS	DEF	ERRABLE	DEFERRED
	-				
ORDER_STATUS2_PK	Р	ENABLED	NOT	DEFERRABLE	IMMEDIATE
ORDER_STATUS2_STATUS_CK	С	ENABLED	NOT	DEFERRABLE	IMMEDIATE
ORDER_STATUS2_ID_CK	С	ENABLED	NOT	DEFERRABLE	IMMEDIATE
ORDER_STATUS2_STATUS_NN	С	DISABLED	NOT	DEFERRABLE	IMMEDIATE
ORDER_STATUS2_STATUS_UQ	U	ENABLED	DEF	ERRABLE	DEFERRED
SYS C004807	С	ENABLED	NOT	DEFERRABLE	IMMEDIATE

Notice that all the constraints except one have a helpful name. One constraint has the database-generated name of SYS C004807 (this name is automatically generated, and it will be different in your database). This constraint is the one for which I omitted the name when creating it earlier.



TIP

Always add a descriptive name for your constraints.

Getting Information on the Constraints for a Column

You can retrieve information on the constraints for a column by querying the user cons columns view. Table 10-5 describes some of the columns in user cons columns.



You can retrieve information on all the column constraints you have access to by querying the all cons columns view.

Column	Туре	Description
owner	VARCHAR2(30)	Owner of the constraint.
constraint_name	VARCHAR2(30)	Name of the constraint.
constraint_type	VARCHAR2(1)	Constraint type (P, R, C, U, V, or O). See Table 10-3 for the constraint type meanings.
table_name	VARCHAR2(30)	Name of the table on which the constraint is defined.
status	VARCHAR2(8)	Constraint status (ENABLED or DISABLED).
deferrable	VARCHAR2 (14)	Whether the constraint is deferrable (DEFERRABLE OR NOT DEFERRABLE).
deferred	VARCHAR2(9)	Whether the constraint is enforced immediately or deferred (IMMEDIATE or DEFERRED).

 TABLE 10-4
 Some Columns in the user_constraints View

Column	Туре	Description
owner	VARCHAR2(30)	Owner of the constraint
constraint_name	VARCHAR2(30)	Name of the constraint
table_name	VARCHAR2(30)	Name of the table on which the constraint is defined
column_name	VARCHAR2 (4000)	Name of the column on which the constraint is defined

 TABLE 10-5
 Some Columns in the user_cons_columns View

The following example retrieves the constraint_name and column_name from user_cons columns for the order_status2 table:

COLUMN column_name FORMAT a15

SELECT constraint_name, column_name
FROM user_cons_columns

WHERE table_name = 'ORDER_STATUS2'

ORDER BY constraint name;

CONSTRAINT_NAME	COLUMN_NAME
ORDER_STATUS2_ID_CK	ID
ORDER_STATUS2_PK	ID
ORDER_STATUS2_STATUS_CK	STATUS
ORDER_STATUS2_STATUS_NN	STATUS
ORDER_STATUS2_STATUS_UQ	STATUS
SYS C004807	LAST MODIFIED

The next query joins user_constraints and user_cons_columns to get the column_name, constraint_name, constraint_type, and status:

SELECT ucc.column_name, ucc.constraint_name, uc.constraint_type, uc.status
FROM user_constraints uc, user_cons_columns ucc
WHERE uc.table_name = ucc.table_name
AND uc.constraint_name = ucc.constraint_name
AND ucc.table_name = 'ORDER_STATUS2'
ORDER BY ucc.constraint_name;

COLUMN_NAME	CONSTRAINT_NAME	С	STATUS
		_	
ID	ORDER_STATUS2_ID_CK	С	ENABLED
ID	ORDER_STATUS2_PK	Ρ	ENABLED
STATUS	ORDER_STATUS2_STATUS_CK	С	ENABLED
STATUS	ORDER_STATUS2_STATUS_NN	С	DISABLED
STATUS	ORDER_STATUS2_STATUS_UQ	U	ENABLED
LAST_MODIFIED	SYS_C004807	С	ENABLED

Renaming a Table

You rename a table using the RENAME statement. The following example renames order status2 to order state:

RENAME order status2 TO order state;



NOTE

If you have used the table name in your constraint names, then you should change the names of your constraints.

The next example changes the table name back to the original:

RENAME order state TO order status2;

Adding a Comment to a Table

A comment can help you remember what the table or column is used for. You add a comment table or column using the COMMENT statement. The following example adds a comment to the order status2 table:

COMMENT ON TABLE order status2 IS 'order status2 stores the state of an order';

The next example adds a comment to the order status2.last modified column:

COMMENT ON COLUMN order status2.last modified IS 'last modified stores the date and time the order was modified last';

Retrieving Table Comments

FROM user tab comments

You can retrieve the comments on your tables from the user tab comments view, as shown here:

SELECT *

```
WHERE table name = 'ORDER STATUS2';
TABLE NAME
                              TABLE TYPE
COMMENTS
```

ORDER STATUS2 TABLE order status2 stores the state of an order

Retrieving Column Comments

You can retrieve the comments on your columns from the user col comments view; for example:

SELECT *

```
FROM user col comments
WHERE table name = 'ORDER STATUS2';
TABLE NAME
                                COLUMN NAME
                                ΙD
```

ORDER STATUS2

```
ORDER_STATUS2 STATUS

ORDER_STATUS2 LAST_MODIFIED

last modified stores the date and time the order was modified last
```

Truncating a Table

You truncate a table using the TRUNCATE statement. This removes *all* the rows from a table and resets the storage area for a table. The following example truncates order status2:

TRUNCATE TABLE order_status2;



TIP

If you need to remove all the rows from a table, you should use TRUNCATE rather than DELETE. This is because TRUNCATE resets the storage area for a table ready to receive new rows. A TRUNCATE statement doesn't require any undo space in the database, and you don't have to run a COMMIT to make the delete permanent. Undo space is an area that the database software uses to record database changes.

Dropping a Table

You drop a table using the DROP TABLE statement. The following example drops the order_status2 table:

DROP TABLE order_status2;

This concludes the discussion of tables. In the next section, you'll learn about sequences.

Sequences

A *sequence* is a database item that generates a sequence of integers. You typically use the integers generated by a sequence to populate a numeric primary key column. In this section, you'll learn how to

- Create a sequence.
- Retrieve information on a sequence from the data dictionary.
- Use a sequence.
- Modify a sequence.
- Drop a sequence.

Creating a Sequence

You create a sequence using the CREATE SEQUENCE statement, which has the following syntax:

```
CREATE SEQUENCE sequence_name
[START WITH start_num]
[INCREMENT BY increment_num]
[{ MAXVALUE maximum num | NOMAXVALUE } ]
```

```
[ { MINVALUE minimum num | NOMINVALUE } ]
[ { CYCLE | NOCYCLE } ]
[ { CACHE cache num | NOCACHE } ]
[ { ORDER | NOORDER } ];
```

where

- sequence name is the name of the sequence.
- start num is the integer to start the sequence. The default start number is 1.
- increment num is the integer to increment the sequence by. The default increment number is 1. The absolute value of increment num must be less than the difference between maximum num and minimum num.
- maximum num is the maximum integer of the sequence; maximum num must be greater than or equal to start num, and maximum num must be greater than minimum num.
- NOMAXVALUE specifies the maximum is 10^{27} for an ascending sequence or -1 for a descending sequence. NOMAXVALUE is the default.
- minimum num is the minimum integer of the sequence; minimum num must be less than or equal to start num, and minimum num must be less than maximum num.
- NOMINVALUE specifies the minimum is 1 for an ascending sequence or -10^{26} for a descending sequence. NOMINVALUE is the default.
- CYCLE means the sequence generates integers even after reaching its maximum or minimum value. When an ascending sequence reaches its maximum value, the next value generated is the minimum. When a descending sequence reaches its minimum value, the next value generated is the maximum.
- NOCYCLE means the sequence cannot generate any more integers after reaching its maximum or minimum value. NOCYCLE is the default.
- cache num is the number of integers to keep in memory. The default number of integers to cache is 20. The minimum number of integers that may be cached is 2. The maximum integers that may be cached is determined by the formula CEIL (maximum num minimum num) / ABS (increment num).
- NOCACHE means no caching. This stops the database from pre-allocating values for the sequence, which prevents numeric gaps in the sequence but reduces performance. Gaps occur because cached values are lost when the database is shut down. If you omit CACHE and NOCACHE, the database caches 20 sequence numbers by default.
- ORDER guarantees the integers are generated in the order of the request. You typically use ORDER when using Real Application Clusters, which are set up and managed by database administrators. Real Application Clusters are multiple database servers that share the same memory. Real Application Clusters can improve performance.
- NOORDER doesn't guarantee the integers are generated in the order of the request. NOORDER is the default.

The following example connects as the store user and creates a sequence named s_test (I always put s at the beginning of sequences):

CONNECT store/store_password CREATE SEQUENCE s test;

Because this CREATE SEQUENCE statement omits the optional parameters, the default values are used. This means that *start num* and *increment num* are set to the default of 1.

The next example creates a sequence named s_test2 and specifies values for the optional parameters:

CREATE SEQUENCE s_test2

START WITH 10 INCREMENT BY 5

MINVALUE 10 MAXVALUE 20

CYCLE CACHE 2 ORDER;

The final example creates a sequence named s test3 that starts at 10 and counts down to 1:

CREATE SEQUENCE s_test3

START WITH 10 INCREMENT BY -1

MINVALUE 1 MAXVALUE 10

CYCLE CACHE 5;

Retrieving Information on Sequences

You can retrieve information on your sequences from the user_sequences view. Table 10-6 describes the columns in user sequences.



NOTE

You can retrieve information on all the sequences you have access to by querying the all sequences view.

Column	Туре	Description
sequence_name	VARCHAR2(30)	Name of the sequence
min_value	NUMBER	Minimum value
max_value	NUMBER	Maximum value
increment_by	NUMBER	Number to increment or decrement sequence by
cycle_flag	VARCHAR2(1)	Whether the sequence cycles (Y or \mathbb{N})
order_flag	VARCHAR2(1)	Whether the sequence is ordered (Y or N)
cache_size	NUMBER	Number of sequence values stored in memory
last_number	NUMBER	Last number that was generated or cached by the sequence

The following example retrieves the details for the sequences from user sequences:

COLUMN sequence name FORMAT a13 SELECT * FROM user sequences ORDER BY sequence name;

SEQUENCE_NAME	MIN_VALUE	MAX_VALUE	INCREMENT_BY	C O	CACHE_SIZE	LAST_NUMBER
S_TEST	1	1.0000E+27	1	N N	20	1
S_TEST2	10	20	5	Y Y	2	10
S_TEST3	1	10	-1	Y N	5	10

Using a Sequence

A sequence generates a series of numbers. A sequence contains two pseudo columns named currval and nextval that you use to get the current value and the next value from the sequence.

Before retrieving the current value, you must first initialize the sequence by retrieving the next value. When you select s test.nextval the sequence is initialized to 1. For example, the following query retrieves s test.nextval; notice that the dual table is used in the FROM clause:

SELECT s test.nextval FROM dual:

The first value in the stest sequence is 1. Once the sequence is initialized, you can get the current value from the sequence by retrieving currval. For example:

SELECT s test.currval FROM dual;

```
CURRVAL
```

When you retrieve currval, nextval remains unchanged; nextval only changes when you retrieve nextval to get the next value. The following example retrieves s test.nextval and s test.currval; notice that these values are both 2:

SELECT s test.nextval, s test.currval FROM dual;

NEXTVAL	CURRVAL
2	2

Retrieving s test.nextval gets the next value in the sequence, which is 2; s test .currval is also 2.

The next example initializes s_test2 by retrieving s_test2.nextval; notice that the first value in the sequence is 10:

SELECT s_test2.nextval FROM dual; NEXTVAL 10

The maximum value for s test2 is 20, and the sequence was created with the CYCLE option, meaning that the sequence will cycle back to 10 once it reaches the maximum of 20:

```
SELECT s test2.nextval
   FROM dual;
      NEXTVAL
           15
   SELECT s test2.nextval
   FROM dual:
      NEXTVAL
   _____
           20
   SELECT s test2.nextval
   FROM dual;
      NEXTVAL
          10
```

The s test3 sequence starts at 10 and counts down to 1:

```
SELECT s test3.nextval
   FROM dual;
     NEXTVAL
           10
   SELECT s test3.nextval
   FROM dual;
      NEXTVAL
   SELECT s test3.nextval
   FROM dual;
```

```
NEXTVAL
```

Populating a Primary Key Using a Sequence

Sequences are useful for populating integer primary key column values. Let's take a look at an example. The following statement re-creates the order status2 table:

```
CREATE TABLE order status2 (
      id INTEGER CONSTRAINT order status2 pk PRIMARY KEY,
       status VARCHAR2(10),
       last modified DATE DEFAULT SYSDATE
     );
```

Next, the following statement creates a sequence named s order status2 (this sequence will be used to populate the order status2.id column shortly):

CREATE SEQUENCE s order status2 NOCACHE;



When using a sequence to populate a primary key column, you should typically use NOCACHE to avoid gaps in the sequence of numbers (gaps occur because cached values are lost when the database is shut down). However, using NOCACHE reduces performance. If you are absolutely sure you can live with gaps in the primary key values, then consider using CACHE.

The following INSERT statements add rows to order status2; notice that the value for the id column is set using s order status2.nextval (returns 1 for the first INSERT and 2 for the second INSERT):

```
INSERT INTO order status2 (
      id, status, last_modified
     ) VALUES (
      s order status2.nextval, 'PLACED', '01-JAN-2006'
     INSERT INTO order status2 (
      id, status, last modified
     ) VALUES (
       s_order_status2.nextval, 'PENDING', '01-FEB-2006'
```

The following query retrieves the rows from order status2; notice that the id column is set to the first two values (1 and 2) from the sorder status 2 sequence:

```
SELECT *
    FROM order_status2;
```

```
ID STATUS LAST MODI
 1 PLACED 01-JAN-06
 2 PENDING 01-FEB-06
```

Modifying a Sequence

You modify a sequence using the ALTER SEQUENCE statement. There are some limitations on what you can modify in a sequence:

- You cannot change the start value of a sequence.
- The minimum value cannot be more than the current value of the sequence.
- The maximum value cannot be less than the current value of the sequence.

The following example modifies s test to increment the sequence of numbers by 2:

```
ALTER SEQUENCE s_test
INCREMENT BY 2;
```

When this is done, the new values generated by s_test will be incremented by 2. For example, if s_test.currval is 2, then s_test.nextval is 4. This is shown in the following example:

Dropping a Sequence

You drop a sequence using DROP SEQUENCE. The following example drops s_test3:

```
DROP SEQUENCE s test3;
```

This concludes the discussion of sequences. In the next section, you'll learn about indexes.

Indexes

When looking for a particular topic in a book, you can either scan the whole book, or you can use the index to find the location. An index for a database table is similar in concept to a book index, except that database indexes are used to find specific rows in a table. The downside of indexes is that when a row is added to the table, additional time is required to update the index for the new row.

Generally, you should create an index on a column when you are retrieving a small number of rows from a table containing many rows. A good rule of thumb is

Create an index when a query retrieves <= 10 percent of the total rows in a table.

This means the column for the index should contain a wide range of values. These types of indexes are called "B-tree" indexes, a name which comes from a tree data structure used in computer science. A good candidate for B-tree indexing would be a column containing a unique value for each row (for example, a social security number). A poor candidate for B-tree indexing would be a column that contains only a small range of values (for example, N, S, E, W or 1, 2, 3, 4, 5, 6). An Oracle database automatically creates a B-tree index for the primary key of a table and for columns included in a unique constraint. For columns that contain a small range of values, you can use a "bitmap" index.

In this section, you'll learn how to

- Create a B-tree index.
- Create a function-based index.
- Retrieve information on an index from the data dictionary.
- Modify an index.
- Drop an index.
- Create a bitmap index.

Creating a B-tree Index

You create a B-tree index using CREATE INDEX, which has the following simplified syntax:

```
CREATE [UNIQUE] INDEX index name ON
    table name(column name[, column name ...])
    TABLESPACE tab space;
```

where

- UNIQUE means that the values in the indexed columns must be unique.
- index name is the name of the index.
- table name is a database table.
- column name is the indexed column. You can create an index on multiple columns (such an index is known as a *composite index*).
- tab space is the tablespace for the index. If you don't provide a tablespace, the index is stored in the user's default tablespace.



For performance reasons, you should typically store indexes in a different tablespace from tables. For simplicity, the examples in this chapter use the default tablespace. In a production database, the database administrator should create separate tablespaces for the tables and indexes.

I'll now guide you through the thought processes you should follow when creating a B-tree index for the customers.last_name column. Assume that the customers table contains a large number of rows and that you regularly retrieve rows using the following type of query:

```
SELECT customer_id, first_name, last_name
FROM customers
WHERE last_name = 'Brown';
```

Also assume that the <code>last_name</code> column contains somewhat unique values, so that any query using the <code>last_name</code> column in a <code>WHERE</code> clause will return less than 10 percent of the total number of rows in the table. This means the <code>last_name</code> column is therefore a good candidate for indexing.

The following CREATE INDEX statement creates an index named i_customers_last_name on the last_name column of the customers table (I always put i_ at the start of index names):

```
CREATE INDEX i_customers_last_name ON customers(last_name);
```

Once the index has been created, the previous query will take less time to complete. You can enforce uniqueness of column values using a unique index. For example, the following statement creates a unique index named i_customers_phone on the customers.phone column:

```
CREATE UNIQUE INDEX i customers phone ON customers(phone);
```

You can also create a composite index on multiple columns. For example, the following statement creates a composite index named i_employees_first_last_name on the first_name and last_name columns of the employees table:

```
CREATE INDEX i_employees_first_last_name ON employees(first name, last name);
```

Creating a Function-Based Index

In the previous section you saw the index i_customers_last_name. Let's say you run the following query:

```
SELECT first_name, last_name
FROM customers
WHERE last name = UPPER('BROWN');
```

Because this query uses a function—UPPER(), in this case—the i_customers_last_ name index isn't used. If you want an index to be based on the results of a function, you must create a function-based index, such as:

```
CREATE INDEX i_func_customers_last_name
ON customers(UPPER(last name));
```

In addition, the database administrator must set the initialization parameter QUERY_REWRITE_ ENABLED to true (the default is false) in order to take advantage of function-based indexes. The following example sets QUERY REWRITE ENABLED to true: CONNECT system/manager ALTER SYSTEM SET QUERY REWRITE ENABLED=TRUE;

Retrieving Information on Indexes

You can retrieve information on your indexes from the user indexes view. Table 10-7 describes some of the columns in user indexes.



NOTE

You can retrieve information on all the indexes you have access to by querying the all indexes view.

The following example connects as the store user and retrieves some of the columns from user indexes for the customers and employees tables; notice that the list of indexes includes customers pk, which is a unique index automatically created by the database for the customer id primary key column of the customers table:

CONNECT store/store password SELECT index name, table name, uniqueness, status FROM user indexes WHERE table_name IN ('CUSTOMERS', 'EMPLOYEES') ORDER BY index name;

INDEX_NAME	TABLE_NAME	UNIQUENES	STATUS
CUSTOMERS_PK	CUSTOMERS	UNIQUE	VALID
EMPLOYEES_PK	EMPLOYEES	UNIQUE	VALID
I_CUSTOMERS_LAST_NAME	CUSTOMERS	NONUNIQUE	VALID
I_CUSTOMERS_PHONE	CUSTOMERS	UNIQUE	VALID
I_EMPLOYEES_FIRST_LAST_NAME	EMPLOYEES	NONUNIQUE	VALID
I_FUNC_CUSTOMERS_LAST_NAME	CUSTOMERS	NONUNIQUE	VALID

Retrieving Information on the Indexes on a Column

You can retrieve information on the indexes on a column by querying the user ind columns view. Table 10-8 describes some of the columns in user ind columns.

Column	Туре	Description
index_name	VARCHAR2(30)	Name of the index
table_owner	VARCHAR2(30)	The user who owns the table
table_name	VARCHAR2(30)	The name of the table on which the index was created
uniqueness	VARCHAR2(9)	Indicates whether the index is unique (UNIQUE or NONUNIQUE)
status	VARCHAR2(8)	Indicates whether the index is valid (VALID or INVALID)

 TABLE 10-7
 Some Columns in the user indexes View

Column	Туре	Description
index_name	VARCHAR2(30)	Name of the index
table_name	VARCHAR2(30)	Name of the table
column_name	VARCHAR2(4000)	Name of the indexed column

TABLE 10-8 Some Columns in the user ind columns View



NOTE

You can retrieve information on all the indexes you have access to by querying the all ind columns view.

The following query retrieves some of the columns from user_ind_columns for the customers and employees tables:

COLUMN table_name FORMAT a15

COLUMN column_name FORMAT a15

SELECT index_name, table_name, column_name

FROM user_ind_columns

WHERE table_name IN ('CUSTOMERS', 'EMPLOYEES')

ORDER BY index_name;

INDEX_NAME	TABLE_NAME	COLUMN_NAME
CUSTOMERS_PK	CUSTOMERS	CUSTOMER_ID
EMPLOYEES_PK	EMPLOYEES	EMPLOYEE_ID
I_CUSTOMERS_LAST_NAME	CUSTOMERS	LAST_NAME
I_CUSTOMERS_PHONE	CUSTOMERS	PHONE
I_EMPLOYEES_FIRST_LAST_NAME	EMPLOYEES	LAST_NAME
I_EMPLOYEES_FIRST_LAST_NAME	EMPLOYEES	FIRST_NAME
I FUNC CUSTOMERS LAST NAME	CUSTOMERS	SYS NC00006\$

Modifying an Index

You modify an index using ALTER INDEX. The following example renames the i_customers_phone index to i customers phone number:

ALTER INDEX i customers phone RENAME TO i customers phone number;

Dropping an Index

You drop an index using the DROP INDEX statement. The following example drops the i_customers_phone_number index:

DROP INDEX i_customers_phone_number;

Creating a Bitmap Index

Bitmap indexes are typically used in *data warehouses*, which are databases containing very large amounts of data. The data in a data warehouse is typically read using many queries, but the data

is not modified by many concurrent transactions. Data warehouses are typically used by organizations for business intelligence analysis, like monitoring sales trends.

A candidate for a bitmap index is a column that is referenced in many queries, but that contains only a small range of values; for example:

- N, S, E, W
- 1, 2, 3, 4, 5, 6
- "Order placed", "Order shipped"

An index basically contains a pointer to a row in a table that contains a given index key value; the key value is used to get the rowid for the row in the table. (As discussed in Chapter 2, a rowid is used internally by the database to store the physical location of the row.) In a B-tree index, a list of rowids is stored for each key corresponding to the rows with that key value. In a B-tree index, the database stores a list of key values with each rowid, which enables the database to locate an actual row in a table.

In a bitmap index, however, a bitmap is used for each key value; the bitmap enables the database to locate a row. Each bit in the bitmap corresponds to a possible rowid. If the bit is set, then it means that the row with the corresponding rowid contains the key value. A mapping function converts the bit position to an actual rowid.

Bitmap indexes are typically used in tables containing large amounts of data and whose contents are not modified very often. Also, a bitmap index should only be created on columns that contain a small number of distinct values. If the number of distinct values of a column is less than 1 percent of the number of rows in the table, or if the values in a column are repeated more than 100 times, then the column is a candidate for a bitmap index. For example, if you had a table with 1 million rows, a column with 10,000 distinct values or less is a good candidate for a bitmap index; also, updates to the rows in the table should be rare, and the column would need to be frequently used in the WHERE clause of queries.

The following statement creates a bitmap index on the status column of the order status table:

CREATE BITMAP INDEX i order status ON order status(status);



NOTE

Of course, this example is not a real-world example because the order status table does not contain enough rows.

You can find more information on bitmap indexes in Oracle Database Performance Tuning Guide and Oracle Database Concepts, both books published by Oracle Corporation. These books also contain information about other exotic types of indexes you can use.

This concludes the discussion of indexes. In the next section, you'll learn about views.

Views

A view is a predefined query on one or more tables (known as base tables). Retrieving information from a view is done in the same manner as retrieving from a table: you simply include the view in the FROM clause of a query. With some views you can also perform DML operations on the base tables.



NOTE

Views don't store rows. Rows are always stored in tables.

You've already seen some examples of retrieving information from views when you selected rows from the data dictionary, which is accessed through views—for example, user_tables, user sequences, and user indexes are all views.

Views offer several benefits, such as the following:

- You can put a complex query into a view and grant users access to the view. This allows you to hide complexity from users.
- You can stop users from directly querying the base tables by granting them access only to the view.
- You can allow a view to access only certain rows in the base tables. This allows you to hide rows from an end user.

In this section, you'll learn how to

- Create and use a view.
- Get the details of a view from the data dictionary.
- Modify a view.
- Drop a view.

Creating and Using a View

You create a view using CREATE VIEW, which has the following simplified syntax:

```
CREATE [OR REPLACE] VIEW [{FORCE | NOFORCE}] VIEW view_name
[(alias_name[, alias_name ...])] AS subquery
[WITH {CHECK OPTION | READ ONLY} CONSTRAINT constraint name];
```

where

- OR REPLACE means the view replaces an existing view.
- FORCE means the view is to be created even if the base tables don't exist.
- NOFORCE means the view is not created if the base tables don't exist. NOFORCE is the default.
- view name is the name of the view.
- alias_name is the name of an alias for an expression in the subquery. There must be the same number of aliases as there are expressions in the subquery.
- subquery is the subquery that retrieves from the base tables. If you've supplied aliases, you can use those aliases in the list after the SELECT.
- WITH CHECK OPTION means that only the rows that would be retrieved by the subquery can be inserted, updated, or deleted. By default, the rows are not checked.

- constraint name is the name of the WITH CHECK OPTION or WITH READ ONLY constraint.
- WITH READ ONLY means the rows may only read from the base tables.

There are two basic types of views:

- Simple views, which contain a subquery that retrieves from one base table
- Complex views, which contain a subquery that
 - Retrieves from multiple base tables
 - Groups rows using a GROUP BY or DISTINCT clause
 - Contains a function call

You'll learn how to create and use these types of views in the following sections.

Privilege for Views

In order to create a view, the user must have the CREATE VIEW privilege. The following example connects as the system user and grants the CREATE VIEW privilege to the store user:

```
CONNECT system/manager
   GRANT CREATE VIEW TO store:
```

Creating and Using Simple Views

Simple views access one base table. The following example connects as the store user and creates a view named cheap products view whose subquery retrieves products only where the price is less than \$15:

```
CONNECT store/store password
    CREATE VIEW cheap products view AS
    SELECT *
    FROM products
    WHERE price < 15;
```

The next example creates a view named employees view whose subquery retrieves all the columns from the employees table except salary:

```
CREATE VIEW employees view AS
   SELECT employee id, manager id, first name, last name, title
   FROM employees;
```

Performing a Query on a View

Once you've created a view, you can use it to access the base table. The following query retrieves rows from cheap products view:

SELECT product id, name, price FROM cheap products view;

PRODUCT_ID	NAME	PRICE
4	Tank War	13.95
6	2412: The Return	14.95

```
7 Space Force 9 13.49
8 From Another Planet 12.99
9 Classical Music 10.99
11 Creative Yell 14.99
12 My Front Line 13.49
```

The next example retrieves rows from employees view:

SELECT *

FROM employees view;

```
EMPLOYEE_ID MANAGER_ID FIRST_NAME LAST_NAME TITLE

1 James Smith CEO
2 1 Ron Johnson Sales Manager
3 2 Fred Hobbs Salesperson
4 2 Susan Jones Salesperson
```

Performing an INSERT Using a View

You can perform DML statements using cheap_products_view. The following example performs an INSERT using cheap_products_view and then retrieves the row:

```
INSERT INTO cheap_products_view (
    product_id, product_type_id, name, price
) VALUES (
    13, 1, 'Western Front', 13.50
);

1 row created.

SELECT product_id, name, price
FROM cheap_products_view
WHERE product_id = 13;

PRODUCT_ID NAME PRICE

13 Western Front 13.5
```



You can perform DML statements only with simple views. Complex views don't support DML.

Because cheap_products_view didn't use WITH CHECK OPTION, you can insert, update, and delete rows that aren't retrievable by the view. The following example inserts a row whose price is \$16.50 (this is greater than \$15 and therefore not retrievable by the view):

```
INSERT INTO cheap_products_view (
    product_id, product_type_id, name, price
) VALUES (
    14, 1, 'Eastern Front', 16.50
);

1 row created.
```

```
SELECT *
FROM cheap products view
WHERE product id = 14;
no rows selected
```

The employees view contains a subquery that selects every column from employees except salary. When you perform an INSERT using employees view, the salary column in the employees base table will be set to null; for example:

```
INSERT INTO employees view (
      employee id, manager id, first name, last name, title
    ) VALUES (
      5, 1, 'Jeff', 'Jones', 'CTO'
    );
    1 row created.
    SELECT employee id, first name, last name, salary
    FROM employees
    WHERE employee id = 5;
    EMPLOYEE_ID FIRST_NAME LAST NAME SALARY
    ______
            5 Jeff
                       Jones
```

The salary column is null.

Creating a View with a CHECK OPTION Constraint

You can specify that DML statements on a view must satisfy the subquery using a CHECK OPTION constraint. For example, the following statement creates a view named cheap products view2 that has a CHECK OPTION constraint:

```
CREATE VIEW cheap_products_view2 AS
    SELECT *
    FROM products
    WHERE price < 15
    WITH CHECK OPTION CONSTRAINT cheap products view2 price;
```

The next example attempts to insert a row using cheap products view2 with a price of \$19.50; notice that the database returns an error because the row isn't retrievable by the view:

```
INSERT INTO cheap products view2 (
      product id, product type id, name, price
    ) VALUES (
      15, 1, 'Southern Front', 19.50
    INSERT INTO cheap products view2 (
    ERROR at line 1:
    ORA-01402: view WITH CHECK OPTION where-clause violation
```

Creating a View with a READ ONLY Constraint

You can make a view read-only by adding a READ ONLY constraint to the view. For example, the following statement creates a view named cheap_products_view3 that has a READ ONLY constraint:

```
CREATE VIEW cheap_products_view3 AS

SELECT *

FROM products

WHERE price < 15

WITH READ ONLY CONSTRAINT cheap products view3 read only;
```

The following example attempts to insert a row using cheap_products_view3; notice that the database returns an error because the view is read-only and doesn't allow DML statements:

```
INSERT INTO cheap_products_view3 (
    product_id, product_type_id, name, price
) VALUES (
    16, 1, 'Northern Front', 19.50
);
    product_id, product_type_id, name, price
    *
ERROR at line 2:
ORA-42399: cannot perform a DML operation on a read-only view
```

Getting Information on View Definitions

You can retrieve information on view definitions using the DESCRIBE command. The following example uses DESCRIBE with cheap products view3:

DESCRIBE cheap_products_view3

Name	Nul	1?	Туре
PRODUCT_ID	NOT	NULL	NUMBER (38)
PRODUCT_TYPE_ID			NUMBER (38)
NAME	NOT	NULL	VARCHAR2(30)
DESCRIPTION			VARCHAR2(50)
PRICE			NUMBER (5,2)

You can also retrieve information about your views from the user_views view. Table 10-9 describes some of the columns in user views.

Column	Туре	Description
view_name	VARCHAR2(30)	Name of the view
text_length	NUMBER	Number of characters in the view's subquery
text	LONG	Text of the view's subquery



NOTE

You can retrieve information on all the indexes you have access to by querying all views.

To see the entire view definition stored in the text column, you use the SQL*Plus command SET LONG, which sets the number of characters displayed by SQL*Plus when retrieving LONG columns. For example, the following command sets LONG to 200:

SET LONG 200

The following guery retrieves the view name, text length, and text columns from user views:

SELECT view name, text length, text FROM user views ORDER BY view name;

```
VIEW NAME
                          TEXT LENGTH
TEXT
CHEAP PRODUCTS VIEW
SELECT "PRODUCT ID", "PRODUCT TYPE ID", "NAME", "DESCRIPTION", "PRICE"
FROM products
WHERE price < 15
CHEAP PRODUCTS VIEW2
                                       116
SELECT "PRODUCT ID", "PRODUCT TYPE ID", "NAME", "DESCRIPTION", "PRICE"
FROM products
WHERE price < 15
WITH CHECK OPTION
CHEAP PRODUCTS VIEW3
                                       113
SELECT "PRODUCT ID", "PRODUCT TYPE ID", "NAME", "DESCRIPTION", "PRICE"
FROM products
WHERE price < 15
WITH READ ONLY
EMPLOYEES VIEW
SELECT employee id, manager id, first name, last name, title
FROM employees
```

Retrieving Information on View Constraints

Earlier you saw that you can add CHECK OPTION and READ ONLY constraints to a view; cheap products view2 contained a CHECK OPTION constraint to ensure the price was less than \$15; cheap products view3 contained a READ ONLY constraint to prevent modifications to the rows in the base table.

You retrieve information on view constraints from the user constraints view; for example:

```
SELECT constraint name, constraint_type, status, deferrable, deferred
     FROM user constraints
     WHERE table name IN ('CHEAP_PRODUCTS_VIEW2', 'CHEAP_PRODUCTS_VIEW3')
```

ORDER BY constraint name;

The constraint_type for CHEAP_PRODUCTS_VIEW2_PRICE is V, which, as shown earlier in Table 10-3, corresponds to a CHECK OPTION constraint. The constraint_type for CHEAP PRODUCTS VIEW3 READ ONLY is O, which corresponds to a READ ONLY constraint.

Creating and Using Complex Views

Complex views contain subqueries that

- Retrieve rows from multiple base tables.
- Group rows using a GROUP BY or DISTINCT clause.
- Contain a function call.

The following example creates a view named products_and_types_view whose subquery performs a full outer join on the products and product types tables using the SQL/92 syntax:

```
CREATE VIEW products_and_types_view AS

SELECT p.product_id, p.name product_name, pt.name product_type_name, p.price

FROM products p FULL OUTER JOIN product_types pt

USING (product_type_id)

ORDER BY p.product_id;
```

The following example queries products and types view:

SELECT * FROM products and types view;

PRODUCT_ID	PRODUCT_NAME	PRODUCT_TY	PRICE
1	Modern Science	Book	19.95
2	Chemistry	Book	30
3	Supernova	Video	25.99
4	Tank War	Video	13.95
5	Z Files	Video	49.99
6	2412: The Return	Video	14.95
7	Space Force 9	DVD	13.49
8	From Another Planet	DVD	12.99
9	Classical Music	CD	10.99
10	Pop 3	CD	15.99
11	Creative Yell	CD	14.99
12	My Front Line		13.49
13	Western Front	Book	13.5
14	Eastern Front	Book	16.5
		Magazine	

The next example creates a view named <code>employee_salary_grades_view</code> whose subquery uses an inner join to retrieve the salary grades for the employees:

CREATE VIEW employee_salary_grades_view AS SELECT e.first name, e.last name, e.title, e.salary, sq.salary grade id FROM employees e INNER JOIN salary grades sq ON e.salary BETWEEN sg.low salary AND sg.high salary ORDER BY sg.salary grade id;

The following example queries employee salary grades view:

SELECT *

FROM employee salary grades view;

FIRST_NAME	LAST_NAME	TITLE	SALARY	SALARY_GRADE_ID
Fred	Hobbs	Salesperson	150000	1
Susan	Jones	Salesperson	500000	2
Ron	Johnson	Sales Manager	600000	3
James	Smith	CEO	800000	4

The next example creates a view named product average view whose subquery uses

- A WHERE clause to filter the rows from the products table to those whose price is less than \$15.
- A GROUP BY clause to group the remaining rows by the product type id column.
- A HAVING clause to filter the row groups to those whose average price is greater than \$13.

```
CREATE VIEW product average view AS
    SELECT product_type_id, AVG(price) average_price
    FROM products
    WHERE price < 15
    GROUP BY product type id
    HAVING AVG(price) > 13
    ORDER BY product type id;
```

The following example queries product average view:

SELECT *

FROM product average view;

PRODUCT	_TYPE_ID	AVERAGE_	_PRICE
	1		13.5
	2		14.45
	3		13.24
			13.49

Modifying a View

You can completely replace a view using CREATE OR REPLACE VIEW. The following example uses CREATE OR REPLACE VIEW to replace product average view:

CREATE OR REPLACE VIEW product average view AS SELECT product type id, AVG(price) average price FROM products

```
WHERE price < 12
GROUP BY product_type_id
HAVING AVG(price) > 11
ORDER BY product type id;
```

You can alter the constraints on a view using ALTER VIEW. The following example uses ALTER VIEW to drop the cheap products view2 price constraint from cheap products view2:

```
ALTER VIEW cheap_products_view2
DROP CONSTRAINT cheap products view2 price;
```

Dropping a View

You drop a view using DROP VIEW. The following example drops cheap products view2:

```
DROP VIEW cheap_products_view2;
```

This concludes the discussion of views. In the next section, you'll learn about flashback data archives.

Flashback Data Archives

Flashback data archives, which are new for Oracle Database 11*g*, store changes made to a table over a period of time and provide you with a full audit trail. Once you've created a flashback archive and added a table to it you can do the following:

- View rows as they were at a specific timestamp
- View rows as they were between two timestamps

You create a flashback archive using the CREATE FLASHBACK ARCHIVE statement. The following example connects as the system user and creates a flashback archive named test_archive:

CONNECT system/manager

CREATE FLASHBACK ARCHIVE test_archive

TABLESPACE example

QUOTA 1 M

RETENTION 1 DAY;

Notice the following:

- The archive is created in the example tablespace; you can see the full list of tablespaces by running the query SELECT tablespace name FROM dba tablespaces.
- The test_archive has a quota of 1 megabyte, which means it can store up to 1 megabyte of data in the example tablespace.
- The data in test_archive is retained for 1 day, after which time the data is purged.

You may alter an existing table to store data in the archive; for example:

ALTER TABLE store.products FLASHBACK ARCHIVE test_archive;

Any subsequent changes made to the store.products table are now recorded in the archive. The following INSERT statement adds a row to the store.products table:

```
INSERT INTO store.products (
       product id, product type id, name, description, price
     ) VALUES (
       15, 1, 'Using Linux', 'How to Use Linux', 39.99
     );
```

The following query retrieves this row:

```
SELECT product id, name, price
   FROM store.products
   WHERE product id = 15;
```

```
PRODUCT ID NAME
______
   15 Using Linux
                       39.99
```

You can view the rows as they were 5 minutes ago using the following query:

```
SELECT product id, name, price
    FROM store.products
    AS OF TIMESTAMP
     (SYSTIMESTAMP - INTERVAL '5' MINUTE);
```

PRODUCT_ID	NAME	PRICE
1	Modern Science	19.95
2	Chemistry	30
3	Supernova	25.99
4	Tank War	13.95
5	Z Files	49.99
6	2412: The Return	14.95
7	Space Force 9	13.49
8	From Another Planet	12.99
9	Classical Music	10.99
10	Pop 3	15.99
11	Creative Yell	14.99
12	My Front Line	13.49
13	Western Front	13.5
14	Eastern Front	16.5

Notice that the new row is missing. This is because it was added to the table after the date and time specified in the query (assuming the previous INSERT was run less than 5 minutes ago).

You can also view the rows as they were at a specific timestamp using the following query (if you run this query, you need to change the timestamp to a date and time before you ran the INSERT statement earlier):

```
SELECT product id, name, price
     FROM store.products
     AS OF TIMESTAMP
     TO TIMESTAMP('2007-08-12 13:05:00', 'YYYY-MM-DD HH24:MI:SS');
```

The new row will be missing from the results again, because it was added to the table after the date and time specified in the query.

You can view the rows as they were between two timestamps using the following query (you need to change the timestamps):

```
SELECT product_id, name, price
FROM store.products VERSIONS BETWEEN TIMESTAMP
TO_TIMESTAMP('2007-08-12 12:00:00', 'YYYY-MM-DD HH24:MI:SS')
AND TO TIMESTAMP('2007-08-12 12:59:59', 'YYYY-MM-DD HH24:MI:SS');
```

You can view the rows as they were between one timestamp and the present time using the following query (you need to change the timestamp):

```
SELECT product_id, name, price
FROM store.products VERSIONS BETWEEN TIMESTAMP
TO_TIMESTAMP('2007-08-12 13:45:52', 'YYYYY-MM-DD HH24:MI:SS')
AND MAXVALUE;
```

You can stop archiving of data for a table using ALTER TABLE; for example:

```
ALTER TABLE store.products NO FLASHBACK ARCHIVE;
```

When you create a table, you can specify a flashback archive for that table; for example:

```
CREATE TABLE store.test_table (
id INTEGER,
name VARCHAR2(10)
) FLASHBACK ARCHIVE test archive;
```

You can view the details for an archive using the following views:

- user_flashback_archive and dba_flashback_archive, which display general information about the flashback archives
- user_flashback_archive_ts and dba_flashback_archive_ts, which display information about the tablespaces containing the flashback archives
- user_flashback_archive_tables and dba_flashback_archive_tables, which display information about the tables that are enabled for flashback archiving

You can alter a flashback archive; for example, the following statement changes the data retention period to 2 years:

```
ALTER FLASHBACK ARCHIVE test_archive 
MODIFY RETENTION 2 YEAR;
```

You can purge the data from a flashback archive before a given timestamp; for example, the following statement purges data older than 1 day:

```
ALTER FLASHBACK ARCHIVE test_archive
PURGE BEFORE TIMESTAMP(SYSTIMESTAMP - INTERVAL '1' DAY);
```

You can purge all the data in a flashback archive; for example:

ALTER FLASHBACK ARCHIVE test archive PURGE ALL;

You can drop a flashback archive; for example:

DROP FLASHBACK ARCHIVE test_archive;



NOTE

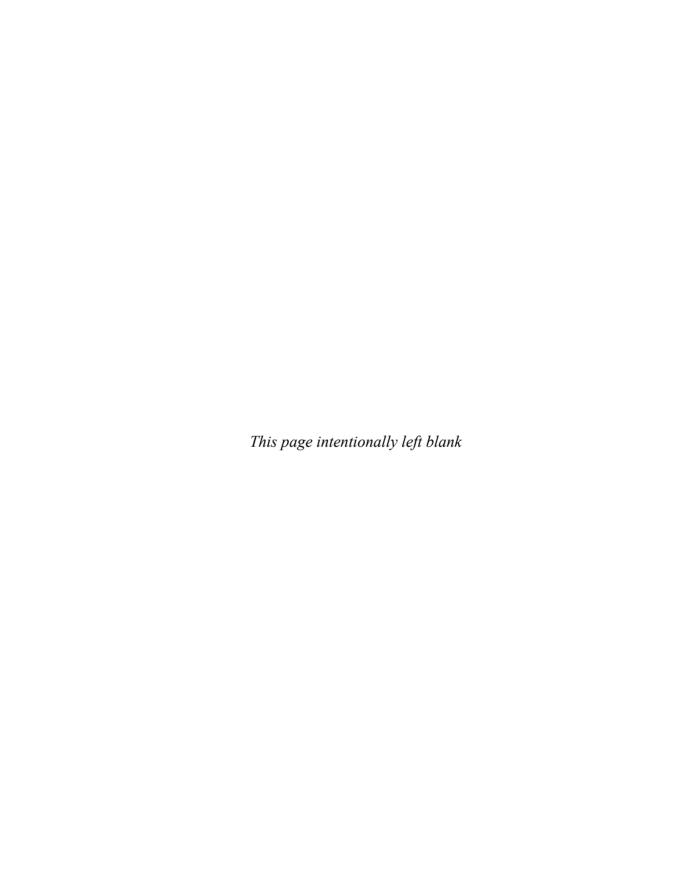
Go ahead and rerun store schema.sql to re-create the store tables so that your queries match mine in the rest of this book.

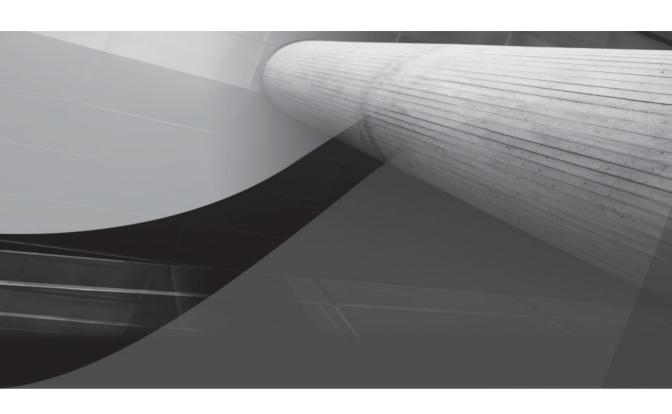
Summary

In this chapter, you have learned the following:

- A table is created using the CREATE TABLE statement.
- A sequence generates a sequence of integers.
- A database index can speed up access to rows.
- A view is a predefined query on one or more base tables.
- A flashback data archive stores changes made to a table over a period of time.

In the next chapter, you'll learn about PL/SQL programming.





CHAPTER 11

Introducing PL/SQL Programming



racle added a procedural programming language known as PL/SQL (Procedural Language/SQL) to Oracle Database 6. PL/SQL enables you to write programs that contain SQL statements. In this chapter, you'll learn about the following PL/SQL topics:

- Block structure
- Variables and types
- Conditional logic
- Loops
- Cursors, which allow PL/SQL to read the results returned by a query
- Procedures
- Functions
- Packages, which are used to group procedures and functions together in one unit
- Triggers, which are blocks of code that are run when a certain event occurs in the database
- Oracle Database 11g enhancements to PL/SQL

You can use PL/SQL to add business logic to a database application. This centralized business logic can be run by any program that can access the database, including SQL*Plus, Java programs, C# programs, and more.



NOTE

For full details on how to access a database through Java, see my book Oracle9i JDBC Programming (Oracle Press, 2002). For details on how to access a database through C#, see my book Mastering C# Database Programming (Sybex, 2003).

Block Structure

PL/SQL programs are divided up into structures known as *blocks*, with each block containing PL/SQL and SQL statements. A PL/SQL block has the following structure:

```
[DECLARE
declaration_statements
]
BEGIN
executable_statements
[EXCEPTION
exception_handling_statements
]
END;
```

where

- declaration statements declare the variables used in the rest of the PL/SQL block. DECLARE blocks are optional.
- executable statements are the actual executable statements, which may include loops, conditional logic, and so on.
- exception handling statements are statements that handle any execution errors that might occur when the block is run. EXCEPTION blocks are optional.

Every statement is terminated by a semicolon (;), and a PL/SQL block is terminated using the forward slash (/) character. Before I get into the details of PL/SQL, you'll see a simple example to get a feel for the language. The following example (contained in the area example.sql script in the SQL directory) calculates the width of a rectangle given its area and height:

SET SERVEROUTPUT ON

```
DECLARE
 v width INTEGER;
 v height INTEGER := 2;
 v area INTEGER := 6;
BEGIN
 -- set the width equal to the area divided by the height
 v width := v area / v height;
 DBMS OUTPUT.PUT LINE('v width = ' | | v width);
EXCEPTION
 WHEN ZERO DIVIDE THEN
    DBMS OUTPUT.PUT LINE('Division by zero');
END;
```

The SET SERVEROUTPUT ON command turns the server output on so you can see the lines produced by DBMS OUTPUT.PUT LINE() on the screen when you run the script in SQL*Plus. After this initial command comes the PL/SQL block itself, which is divided into the DECLARE, BEGIN, and EXCEPTION blocks.

The DECLARE block contains declarations for three INTEGER variables named v width, v height, and v area (I always put v at the start of variable names). The v height and v area variables are initialized to 2 and 6 respectively.

Next comes the BEGIN block, which contains three lines. The first line is a comment that contains the text "set the width equal to the area divided by the height." The second line sets v width to v area divided by v height; this means v width is set to 3 (= 6 / 2). The third line calls DBMS OUTPUT.PUT LINE() to display the value of v width on the screen. DBMS OUTPUT is a built-in package of code that comes with the Oracle database; among other items, DBMS OUTPUT contains procedures that allow you to output values to the screen.

Next, the EXCEPTION block handles any attempts to divide a number by zero. It does this by "catching" the ZERO DIVIDE exception; in the example, no attempt is actually made to divide by zero, but if you change v height to 0 and run the script you'll see the exception.

At the very end of the script, the forward slash character (/) marks the end of the PL/SQL block.

The following listing shows the execution of the area example.sql script in SQL*Plus:

```
SQL> @ C:\SQL\area_example.sql
v width = 3
```



NOTE

If your area_example.sql script is in a directory other than C:\SQL, use your own directory in the previous command.

Variables and Types

Variables are declared within a DECLARE block. As you saw in the previous example, a variable declaration has both a name and a type. For example, the v width variable was declared as

v width INTEGER;



NOTE

The PL/SQL types are similar to the database column types. You can see all the types in the appendix.

The following example shows more variable declarations (these variables could be used to store the column values from the products table):

```
v_product_id INTEGER;
v_product_type_id INTEGER;
v_name VARCHAR2(30);
v_description VARCHAR2(50);
v price NUMBER(5, 2);
```

You may also specify a variable's type using the %TYPE keyword, which tells PL/SQL to use the same type as a specified column in a table. The following example uses %TYPE to declare a variable of the same type as the price column of the products table, which is NUMBER(5, 2):

v product price products.price%TYPE;

Conditional Logic

You use the IF, THEN, ELSE, ELSIF, and END IF keywords to perform conditional logic:

```
IF condition1 THEN
statements1
ELSIF condition2 THEN
statements2
ELSE
statements3
END IF;
```

where

- condition1 and condition2 are Boolean expressions that evaluate to true or false.
- statements1, statements2, and statements3 are PL/SQL statements.

The conditional logic flows as follows:

- If condition1 is true, then statements1 are executed.
- If condition1 is false but condition2 is true, then statements2 are executed.
- If neither condition1 nor condition2 is true, then statements3 are executed.

You can also embed an IF statement within another IF statement, as shown in the following example:

```
IF v count > 0 THEN
    v_message := 'v_count is positive';
    IF v area > 0 THEN
      v message := 'v count and v area are positive';
    END IF
  ELSIF v count = 0 THEN
    v message := 'v count is zero';
    v message := 'v count is negative';
  END IF:
```

In this example, if v count is greater than 0, then v_message is set to 'v_count is positive'. If v count and v area are greater than 0, then v message is set to 'v count and v area are positive'. The rest of the logic is straightforward.

Loops

You use a loop to run statements zero or more times. There are three types of loops in PL/SQL:

- **Simple loops** run until you explicitly end the loop.
- **WHILE loops** run until a specified condition occurs.
- **FOR loops** run a predetermined number of times.

You'll learn about these loops in the following sections.

Simple Loops

A simple loop runs until you explicitly end the loop. The syntax for a simple loop is as follows:

```
LOOP
      statements
    END LOOP;
```

To end the loop, you use either an EXIT or an EXIT WHEN statement. The EXIT statement ends a loop immediately; the EXIT WHEN statement ends a loop when a specified condition occurs.

The following example shows a simple loop. A variable named v counter is initialized to 0 prior to the beginning of the loop. The loop adds 1 to v counter and exits when v counter is equal to 5 using an EXIT WHEN statement.

```
v counter := 0;
    LOOP
      v counter := v counter + 1;
      EXIT WHEN v counter = 5;
    END LOOP;
```



NOTE

The EXIT WHEN statement can appear anywhere in the loop code.

In Oracle Database 11g you can also end the current iteration of a loop using the CONTINUE or CONTINUE WHEN statement. The CONTINUE statement ends the current iteration of the loop unconditionally and continues with the next iteration; the CONTINUE WHEN statement ends the current iteration of the loop when a specified condition occurs and then continues with the next iteration. The following example shows the use of the CONTINUE statement:

```
v_counter := 0;
LOOP
   -- after the CONTINUE statement is executed, control returns here
   v_counter := v_counter + 1;
   IF v_counter = 3 THEN
        CONTINUE; -- end current iteration unconditionally
   END IF;
   EXIT WHEN v_counter = 5;
END LOOP;
```

The next example shows the use of the CONTINUE WHEN statement:

```
v_counter := 0;
LOOP
   -- after the CONTINUE WHEN statement is executed, control returns here
   v_counter := v_counter + 1;
   CONTINUE WHEN v_counter = 3; -- end current iteration when v_counter = 3
   EXIT WHEN v_counter = 5;
END LOOP;
```



NOTE

A CONTINUE or CONTINUE WHEN statement cannot cross a procedure, function, or method boundary.

WHILE Loops

A WHILE loop runs until a specified condition occurs. The syntax for a WHILE loop is as follows:

```
WHILE condition LOOP statements
END LOOP;
```

The following example shows a WHILE loop that executes while the v_counter variable is less than 6:

```
v_counter := 0;
WHILE v_counter < 6 LOOP
v_counter := v_counter + 1;
END LOOP;
```

FOR Loops

A FOR loop runs a predetermined number of times; you determine the number of times the loop runs by specifying the *lower* and *upper bounds* for a loop variable. The loop variable is then incremented (or decremented) each time around the loop. The syntax for a FOR loop is as follows:

```
FOR loop_variable IN [REVERSE] lower bound..upper bound LOOP
      statements
    END LOOP;
```

where

- 100p variable is the loop variable. You can use a variable that already exists as the loop variable, or you can just have the loop create a new variable for you (this occurs if the variable you specify doesn't exist). The loop variable value is increased (or decreased if you use the REVERSE keyword) by 1 each time through the loop.
- REVERSE means that the loop variable value is to be decremented each time through the loop. The loop variable is initialized to the upper boundary, and is decremented by 1 until the loop variable reaches the lower boundary. You must specify the lower boundary before the upper boundary.
- lower bound is the loop's lower boundary. The loop variable is initialized to this lower boundary provided REVERSE is not used.
- upper bound is the loop's upper boundary. If REVERSE is used, the loop variable is initialized to this upper boundary.

The following example shows a FOR loop. Notice that the variable v counter2 isn't explicitly declared—so the FOR loop automatically creates a new INTEGER variable named v counter2:

```
FOR v counter2 IN 1..5 LOOP
       DBMS OUTPUT.PUT LINE(v counter2);
     END LOOP;
```

The following example uses REVERSE:

```
FOR v counter2 IN REVERSE 1..5 LOOP
      DBMS OUTPUT.PUT LINE(v counter2);
    END LOOP;
```

In this example, v counter2 starts at 5, is decremented by 1 each time through the loop, and ends at 1.

Cursors

You use a *cursor* to fetch rows returned by a query. You retrieve the rows into the cursor using a query and then fetch the rows one at a time from the cursor. You typically use the following five steps when using a cursor:

- 1. Declare variables to store the column values for a row.
- 2. Declare the cursor, which contains a query.
- **3.** Open the cursor.
- 4. Fetch the rows from the cursor one at a time and store the column values in the variables declared in Step 1. You would then do something with those variables, such as display them on the screen, use them in a calculation, and so on.
- 5. Close the cursor.

You'll learn the details of these five steps in the following sections, and you'll see a simple example that gets the product_id, name, and price columns from the products table.

Step 1: Declare the Variables to Store the Column Values

The first step is to declare the variables that will be used to store the column values. These variables must be compatible with the column types.



TIP

Earlier you saw that %TYPE may be used to get the type of a column. If you use %TYPE when declaring your variables, your variables will automatically be of the correct type.

The following example declares three variables to store the product_id, name, and price columns from the products table; notice that %TYPE is used to automatically set the type of the variables to the same type as the columns:

DECLARE

```
v_product_id products.product_id%TYPE;
v_name products.name%TYPE;
v price products.price%TYPE;
```

Step 2: Declare the Cursor

Step 2 is to declare the cursor. A cursor declaration consists of a name that you assign to the cursor and the query you want to execute. The cursor declaration, like all other declarations, is placed in the declaration section. The syntax for declaring a cursor is as follows:

```
CURSOR cursor_name IS
SELECT_statement;
```

where

- cursor name is the name of the cursor.
- SELECT statement is the query.

The following example declares a cursor named v_product_cursor whose query retrieves the product id, name, and price columns from the products table:

```
CURSOR v_product_cursor IS

SELECT product_id, name, price
FROM products

ORDER BY product id;
```

The query isn't executed until you open the cursor.

Step 3: Open the Cursor

Step 3 is to open the cursor. You open a cursor using the OPEN statement, which must be placed in the executable section of the block.

The following example opens v_product_cursor, which executes the query:

```
OPEN v product cursor;
```

Step 4: Fetch the Rows from the Cursor

Step 4 is to fetch the rows from the cursor, which you do using the FETCH statement. The FETCH statement reads the column values into the variables declared in Step 1. FETCH uses the following svntax:

```
FETCH cursor name
    INTO variable[, variable ...];
    where
```

- cursor name is the name of the cursor.
- variable is the variable into which a column value from the cursor is stored. You need to provide matching variables for each column value.

The following FETCH example retrieves a row from v product cursor and stores the column values in the v product id, v name, and v price variables created earlier in Step 1:

```
FETCH v product cursor
    INTO v product id, v name, v price;
```

Because a cursor may contain many rows, you need a loop to read them. To figure out when to end the loop, you can use the Boolean variable v product cursor%NOTFOUND. This variable is true when there are no more rows to read in v product cursor. The following example shows a loop:

```
LOOP
       -- fetch the rows from the cursor
       FETCH v product cursor
       INTO v product id, v name, v price;
       -- exit the loop when there are no more rows, as indicated by
       -- the Boolean variable v product cursor%NOTFOUND (= true when
       -- there are no more rows)
       EXIT WHEN v product cursor%NOTFOUND;
       -- use DBMS OUTPUT.PUT LINE() to display the variables
       DBMS OUTPUT.PUT LINE (
         'v product id = ' || v product id || ', v name = ' || v name ||
         ', v_price = ' || v_price
       );
     END LOOP;
```

Notice that I've used DBMS OUTPUT.PUT LINE() to display the v product id, v name, and v price variables that were read for each row. In a real application, you might use v price in a complex calculation.

Step 5: Close the Cursor

Step 5 is to close the cursor using the CLOSE statement. Closing a cursor frees up system resources. The following example closes v product cursor:

```
CLOSE v product cursor;
```

The following section shows a complete script that contains all five steps.

Complete Example: product cursor.sql

The following product cursor.sql script is contained in the SQL directory:

```
-- product cursor.sql displays the product id, name,
     -- and price columns from the products table using a cursor
     SET SERVEROUTPUT ON
     DECLARE
       -- step 1: declare the variables
       v product id products.product id%TYPE;
       v_name products.name%TYPE;
v_price products.price%TYPE;
       -- step 2: declare the cursor
       CURSOR v product cursor IS
         SELECT product id, name, price
         FROM products
         ORDER BY product id;
     BEGIN
       -- step 3: open the cursor
       OPEN v product cursor;
       LOOP
         -- step 4: fetch the rows from the cursor
         FETCH v product cursor
         INTO v product id, v name, v price;
         -- exit the loop when there are no more rows, as indicated by
         -- the Boolean variable v product cursor%NOTFOUND (= true when
         -- there are no more rows)
         EXIT WHEN v product cursor%NOTFOUND;
         -- use DBMS OUTPUT.PUT LINE() to display the variables
         DBMS OUTPUT.PUT LINE (
           'v_product_id = ' || v_product_id || ', v_name = ' || v name ||
           ', v price = ' || v price
         );
       END LOOP;
       -- step 5: close the cursor
       CLOSE v product cursor;
     END;
     /
```

To run this script, follow these steps:

- 1. Connect to the database as store with the password store password.
- 2. Run the product cursor.sql script using SQL*Plus:

```
SQL> @ C:\SQL\product cursor.sql
```



NOTE

If your product cursor.sql script is in a different directory from C:\SOL, use your own directory in the previous command.

The output from product cursor.sql is as follows:

```
v product id = 1, v name = Modern Science, v price = 19.95
    v product id = 2, v name = Chemistry, v price = 30
    v product id = 3, v name = Supernova, v price = 25.99
    v product id = 4, v name = Tank War, v price = 13.95
    v product id = 5, v name = Z Files, v price = 49.99
    v product id = 6, v name = 2412: The Return, v price = 14.95
    v product id = 7, v name = Space Force 9, v price = 13.49
    v product id = 8, v name = From Another Planet, v price = 12.99
    v product id = 9, v name = Classical Music, v price = 10.99
    v product id = 10, v name = Pop 3, v price = 15.99
    v product id = 11, v name = Creative Yell, v price = 14.99
    v product id = 12, v name = My Front Line, v price = 13.49
```

Cursors and FOR Loops

You can use a FOR loop to access the rows in a cursor. When you do this, you don't have to explicitly open and close the cursor—the FOR loop does this automatically for you. The following product cursor2.sql script uses a FOR loop to access the rows in v product cursor; notice that this script contains less code than product cursor.sql:

```
-- product cursor2.sql displays the product id, name,
      -- and price columns from the products table using a cursor
      -- and a FOR loop
      SET SERVEROUTPUT ON
      DECLARE
        CURSOR v product cursor IS
          SELECT product id, name, price
          FROM products
          ORDER BY product id;
      BEGIN
        FOR v product IN v product cursor LOOP
          DBMS OUTPUT.PUT LINE (
            'product id = ' || v product.product id ||
            ', name = ' || v product.name ||
            ', price = ' || v product.price
          );
        END LOOP;
      END;
```

To run the product cursor2.sql script, you issue a command similar to the following:

The output from this script is as follows:

```
product id = 1, name = Modern Science, price = 19.95
     product id = 2, name = Chemistry, price = 30
     product id = 3, name = Supernova, price = 25.99
     product id = 4, name = Tank War, price = 13.95
     product id = 5, name = Z Files, price = 49.99
     product id = 6, name = 2412: The Return, price = 14.95
     product id = 7, name = Space Force 9, price = 13.49
     product id = 8, name = From Another Planet, price = 12.99
     product id = 9, name = Classical Music, price = 10.99
     product id = 10, name = Pop 3, price = 15.99
     product id = 11, name = Creative Yell, price = 14.99
     product id = 12, name = My Front Line, price = 13.49
```

OPEN-FOR Statement

You may also use the OPEN-FOR statement with a cursor, which adds even more flexibility when processing cursors because you can assign the cursor to a different query. This is shown in the following product cursor3.sql script:

```
-- product_cursor3.sql displays the product_id, name,
     -- and price columns from the products table using a cursor
     -- variable and the OPEN-FOR statement
     SET SERVEROUTPUT ON
     DECLARE
       -- declare a REF CURSOR type named t product cursor
       TYPE t product cursor IS
       REF CURSOR RETURN products%ROWTYPE;
       -- declare a t product cursor object named v product cursor
       v product cursor t product cursor;
       -- declare an object to store columns from the products table
       -- named v product (of type products%ROWTYPE)
       v product products%ROWTYPE;
       -- assign a query to v product cursor and open it using OPEN-FOR
       OPEN v product cursor FOR
       SELECT * FROM products WHERE product id < 5;
       -- use a loop to fetch the rows from v product cursor into v product
       LOOP
         FETCH v product cursor INTO v product;
         EXIT WHEN v product cursor%NOTFOUND;
         DBMS OUTPUT.PUT LINE (
           'product id = ' || v product.product id ||
            ', name = ' || v product.name ||
            ', price = ' || v product.price
         );
```

```
END LOOP;
 -- close v product cursor
 CLOSE v product cursor;
END:
```

In the DECLARE block, the following statement declares a REF CURSOR type named t product cursor (I always put t at the start of type names):

```
TYPE t product cursor IS
    REF CURSOR RETURN products%ROWTYPE;
```

A REF CURSOR is a pointer to a cursor, and is similar to a pointer in the C++ programming language. The previous statement declares a user-defined type named t product cursor, and returns a row containing the various columns of the products table (this is indicated using %ROWTYPE). This user-defined type may be used to declare an actual object, as shown in the following statement, which declares an object named v product cursor:

```
v product cursor t product cursor;
```

The following statement declares an object to store columns from the products table named v product (of type products%ROWTYPE):

```
v product products%ROWTYPE;
```

In the BEGIN block, v product cursor is assigned a query and opened by the following OPEN-FOR statement:

```
OPEN v product cursor FOR
    SELECT * FROM products WHERE product id < 5;
```

After this statement is executed, v product cursor will be loaded with the first four rows in the products table. The query assigned to v product cursor can be any valid SELECT statement; this means you can re-use the cursor and assign another query to the cursor later in the PL/SQL code.

Next, the following loop fetches the rows from v product cursor into v product and displays the row details:

```
LOOP
```

```
FETCH v product cursor INTO v product;
 EXIT WHEN v product cursor%NOTFOUND;
 DBMS OUTPUT.PUT LINE (
    'product_id = ' || v_product.product id ||
    ', name = ' || v product.name ||
    ', price = ' || v product.price
 );
END LOOP;
```

After the loop, v product cursor is closed using the following statement:

```
CLOSE v product cursor;
```

The output from this script is the same as the output from product cursor2.sql.

Unconstrained Cursors

The cursors in the previous section all have a specific return type: these cursors are known as constrained cursors. The return type for a constrained cursor must match the columns in the query that is run by the cursor. An unconstrained cursor has no return type, and can therefore run any query.

The use of an unconstrained cursor is shown in the following unconstrained cursor .sgl script; notice v cursor in the code is used to run two different gueries:

```
-- This script shows the use of unconstrained cursors
    SET SERVEROUTPUT ON
    DECLARE
      -- declare a REF CURSOR type named t cursor (this has no return
      -- type and can therefore run any query)
      TYPE t cursor IS REF CURSOR;
      -- declare a t cursor object named v cursor
      v cursor t cursor;
      -- declare an object to store columns from the products table
      -- named v product (of type products%ROWTYPE)
      v product products%ROWTYPE;
      -- declare an object to store columns from the customers table
      -- named v customer (of type customers%ROWTYPE)
      v customer customers%ROWTYPE;
    BEGIN
      -- assign a query to v cursor and open it using OPEN-FOR
      OPEN v cursor FOR
      SELECT * FROM products WHERE product id < 5;
      -- use a loop to fetch the rows from v cursor into v product
        FETCH v cursor INTO v product;
        EXIT WHEN v cursor%NOTFOUND;
        DBMS OUTPUT.PUT LINE (
          'product id = ' || v product.product id ||
          ', name = ' || v_product.name ||
          ', price = ' || v product.price
        );
      END LOOP;
      -- assign a new query to v cursor and open it using OPEN-FOR
      OPEN v cursor FOR
      SELECT * FROM customers WHERE customer id < 3;
      -- use a loop to fetch the rows from v cursor into v product
        FETCH v cursor INTO v customer;
```

EXIT WHEN v cursor%NOTFOUND;

```
DBMS OUTPUT.PUT LINE (
      'customer id = ' || v customer.customer id ||
      ', first name = ' || v customer.first name ||
      ', last name = ' || v customer.last name
   );
 END LOOP:
 -- close v cursor
 CLOSE v cursor;
END:
```

To run the unconstrained cursor.sql script, you issue a command similar to the following:

```
SQL> @ "C:\SQL\unconstrained cursor.sql"
```

The output from this script is as follows:

```
product id = 1, name = Modern Science, price = 19.95
    product id = 2, name = Chemistry, price = 30
    product id = 3, name = Supernova, price = 25.99
    product id = 4, name = Tank War, price = 13.95
    customer id = 1, first name = John, last name = Brown
    customer id = 2, first name = Cynthia, last name = Green
```

You'll learn more about REF CURSOR variables later in this chapter and more about userdefined types in the next chapter.

Exceptions

Exceptions are used to handle run-time errors in your PL/SQL code. Earlier, you saw the following PL/SQL example that contains an EXCEPTION block:

```
DECLARE
      v width INTEGER;
       v height INTEGER := 2;
       v area INTEGER := 6;
     BEGIN
       -- set the width equal to the area divided by the height
       v width := v area / v height;
       DBMS OUTPUT.PUT LINE('v width = ' || v width);
     EXCEPTION
       WHEN ZERO DIVIDE THEN
         DBMS OUTPUT.PUT LINE('Division by zero');
     END;
     /
```

The EXCEPTION block in this example handles an attempt to divide a number by zero. In PL/ SQL terminology, the EXCEPTION block catches a ZERO DIVIDE exception that is raised in the BEGIN block (although in the example code, ZERO DIVIDE is never actually raised). The ZERO DIVIDE exception and the other common exceptions are shown in Table 11-1.

Exception	Error	Description
ACCESS_INTO_NULL	ORA-06530	An attempt was made to assign values to the attributes of an uninitialized object. (You'll learn about objects in Chapter 12.)
CASE_NOT_FOUND	ORA-06592	None of the WHEN clauses of a CASE statement was selected, and there is no default ELSE clause.
COLLECTION_IS_NULL	ORA-06531	An attempt was made to call a collection method (other than EXISTS) on an uninitialized nested table or varray, or an attempt was made to assign values to the elements of an uninitialized nested table or varray. (You'll learn about collections in Chapter 13.)
CURSOR_ALREADY_OPEN	ORA-06511	An attempt was made to open an already open cursor. The cursor must be closed before it can be reopened.
DUP_VAL_ON_INDEX	ORA-00001	An attempt was made to store duplicate values in a column that is constrained by a unique index.
INVALID_CURSOR	ORA-01001	An attempt was made to perform an illegal cursor operation, such as closing an unopened cursor.
INVALID_NUMBER	ORA-01722	An attempt to convert a character string into a number failed because the string does not represent a valid number. Note: In PL/SQL statements, VALUE_ERROR is raised instead of INVALID_NUMBER.
LOGIN_DENIED	ORA-01017	An attempt was made to connect to a database using an invalid user name or password.
NO_DATA_FOUND	ORA-01403	A SELECT INTO statement returned no rows, or an attempt was made to access a deleted element in a nested table or an uninitialized element in an "index by" table.
NOT_LOGGED_ON	ORA-01012	An attempt was made to access a database item without being connected to the database.
PROGRAM_ERROR	ORA-06501	PL/SQL had an internal problem.
ROWTYPE_MISMATCH	ORA-06504	The host cursor variable and the PL/SQL cursor variable involved in an assignment have incompatible return types. For example, when an open host cursor variable is passed to a stored procedure or function, the return types of the actual and formal parameters must be compatible.

 TABLE 11-1
 Predefined Exceptions

Exception	Error	Description
SELF_IS_NULL	ORA-30625	An attempt was made to call a MEMBER method on a null object. That is, the built-in parameter SELF (which is always the first parameter passed to a MEMBER method) is null.
STORAGE_ERROR	ORA-06500	The PL/SQL module ran out of memory or the memory has been corrupted.
SUBSCRIPT_BEYOND_COUNT	ORA-06533	An attempt was made to reference a nested table or varray element using an index number larger than the number of elements in the collection.
SUBSCRIPT_OUTSIDE_LIMIT	ORA-06532	An attempt was made to reference a nested table or varray element using an index number that is outside the legal range (-1 for example).
SYS_INVALID_ROWID	ORA-01410	The conversion of a character string to a universal rowid failed because the character string does not represent a valid rowid.
TIMEOUT_ON_RESOURCE	ORA-00051	A timeout occurred while the database was waiting for a resource.
TOO_MANY_ROWS	ORA-01422	A SELECT INTO statement returned more than one row.
VALUE_ERROR	ORA-06502	An arithmetic, conversion, truncation, or size-constraint error occurred. For example, when selecting a column value into a character variable, if the value is longer than the declared length of the variable, PL/SQL aborts the assignment and raises VALUE_ERROR. Note: In PL/SQL statements, VALUE_ERROR is raised if the conversion of a character string into a number fails. In SQL statements, INVALID_NUMBER is raised instead of VALUE_ERROR.
ZERO_DIVIDE	ORA-01476	An attempt was made to divide a number by zero.

 TABLE 11-1
 Predefined Exceptions (continued)

The following sections show examples that raise some of the exceptions shown in Table 11-1.

ZERO_DIVIDE Exception

The ZERO DIVIDE exception is raised when an attempt is made to divide a number by zero. The following example attempts to divide 1 by 0 in the BEGIN block and therefore raises the ZERO DIVIDE exception:

BEGIN

```
DBMS_OUTPUT.PUT_LINE(1 / 0);
EXCEPTION
 WHEN ZERO_DIVIDE THEN
```

```
DBMS_OUTPUT.PUT_LINE('Division by zero');
END;
/
Division by zero
```

When an exception is raised, program control passes to the EXCEPTION block and the WHEN clause is examined for a matching exception; the code inside the matching clause is then executed. In the previous example, the ZERO_DIVIDE exception is raised in the BEGIN block, and program control then passes to the EXCEPTION block; a matching exception is found in the WHEN clause, and the code inside the clause is executed.

If no matching exception is found, the exception is propagated to the enclosing block. For example, if the EXCEPTION block was omitted from the previous code, the exception is propagated up to SQL*Plus:

BEGIN

```
DBMS_OUTPUT.PUT_LINE(1 / 0);
END;
BEGIN
*
ERROR at line 1:
ORA-01476: divisor is equal to zero
ORA-06512: at line 2
```

As you can see, SQL*Plus displays a default error that shows the line numbers, the Oracle error codes, and a simple description.

DUP_VAL_ON_INDEX Exception

The DUP_VAL_ON_INDEX exception is raised when an attempt is made to store duplicate values in a column that is constrained by a unique index. The following example attempts to insert a row in the customers table with a customer_id of 1; this causes DUP_VAL_ON_INDEX to be raised, because the customers table already contains a row with a customer_id of 1:

BEGIN

```
INSERT INTO customers (
    customer_id, first_name, last_name
) VALUES (
    1, 'Greg', 'Green'
);
EXCEPTION
    WHEN DUP_VAL_ON_INDEX THEN
        DBMS_OUTPUT.PUT_LINE('Duplicate value on an index');
END;
/
```

INVALID_NUMBER Exception

Duplicate value on an index

The INVALID_NUMBER exception is raised when an attempt is made to convert an invalid character string into a number. The following example attempts to convert the string 123X to

a number that is used in an INSERT, which causes INVALID NUMBER to be raised because 123X is not a valid number:

BEGIN

```
INSERT INTO customers (
    customer id, first name, last name
  ) VALUES (
    '123X', 'Greg', 'Green'
 );
EXCEPTION
 WHEN INVALID NUMBER THEN
    DBMS OUTPUT.PUT LINE('Conversion of string to number failed');
END;
```

Conversion of string to number failed

OTHERS Exception

You can use the OTHERS exception to handle all exceptions, as shown here:

BEGIN DBMS OUTPUT.PUT LINE(1 / 0); EXCEPTION WHEN OTHERS THEN DBMS OUTPUT.PUT LINE('An exception occurred'); END: /

An exception occurred

Because OTHERS matches all exceptions, you must list it after any specific exceptions in your EXCEPTION block. If you attempt to list OTHERS elsewhere, the database returns the error PLS-00370; for example:

```
SQL> BEGIN
      2
        DBMS OUTPUT.PUT LINE(1 / 0);
      3 EXCEPTION
           WHEN OTHERS THEN
      5
             DBMS OUTPUT.PUT_LINE('An exception occurred');
      6
           WHEN ZERO DIVIDE THEN
      7
             DBMS OUTPUT.PUT LINE('Division by zero');
      8 END;
      9 /
      WHEN OTHERS THEN
    ERROR at line 4:
    ORA-06550: line 4, column 3:
    PLS-00370: OTHERS handler must be last among the exception
     handlers of a block
    ORA-06550: line 0, column 0:
    PL/SQL: Compilation unit analysis terminated
```

Procedures

A procedure contains a group of SQL and PL/SQL statements. Procedures allow you to centralize your business logic in the database and may be used by any program that accesses the database. In this section, you'll learn how to

- Create a procedure.
- Call a procedure.
- Get information on procedures.
- Drop a procedure.
- View errors in a procedure.

Creating a Procedure

You create a procedure using the CREATE PROCEDURE statement. The simplified syntax for the CREATE PROCEDURE statement is as follows:

```
CREATE [OR REPLACE] PROCEDURE procedure_name
[(parameter_name [IN | OUT | IN OUT] type [, ...])]
{IS | AS}
BEGIN
procedure_body
END procedure_name;
```

where

- OR REPLACE means the procedure is to replace an existing procedure.
- procedure name is the name of the procedure.
- parameter_name is the name of a parameter that is passed to the procedure. You may pass multiple parameters to a procedure.
- IN | OUT | IN OUT is the *mode* of the parameter. You may pick one of the following modes for each parameter:
 - IN, which is the default mode for a parameter. An IN parameter must be set to a value when the procedure is run. The value of an IN parameter cannot be changed in the procedure body.
 - OUT, which means the parameter is set to a value in the procedure body.
 - IN OUT, which means the parameter can have a value when the procedure is run, and the value can be changed in the body.
- \blacksquare type is the type of the parameter.
- procedure body contains the actual code for the procedure.

The following example creates a procedure named update_product_price()—this procedure, and the other PL/SQL code shown in the rest of this chapter, was created when you

ran the store schema.sql script. The update product price() procedure multiplies the price of a product by a factor; the product ID and the factor are passed as parameters to the procedure. If the product exists, the procedure multiplies the product price by the factor and commits the change.

```
CREATE PROCEDURE update product price (
      p product id IN products.product id%TYPE,
      p factor
                  IN NUMBER
    ) AS
      v product count INTEGER;
    BEGIN
      -- count the number of products with the
      -- supplied product id (will be 1 if the product exists)
      SELECT COUNT(*)
      INTO v product count
      FROM products
      WHERE product id = p product id;
      -- if the product exists (v product count = 1) then
      -- update that product's price
      IF v product count = 1 THEN
        UPDATE products
        SET price = price * p factor
        WHERE product id = p product id;
      END IF;
    EXCEPTION
      WHEN OTHERS THEN
        ROLLBACK;
    END update product price;
```

The procedure accepts two parameters named p product id and p factor (I always put p at the start of parameter names). Both of these parameters use the IN mode, which means that their values must be set when the procedure is run and that the parameter values cannot be changed in the procedure body.

The declaration section contains an INTEGER variable named v product count:

```
v product count INTEGER;
```

The body of the procedure starts after BEGIN. The SELECT statement in the body gets the number of rows from the products table whose product id is equal to p product id:

```
SELECT COUNT (*)
    INTO v product count
     FROM products
    WHERE product id = p product id;
```



COUNT (*) returns number of rows found.

If the product is found, v product count will be set to 1; otherwise, v product count will be set to 0. If v product count is 1, the price column is multiplied by p factor using the UPDATE statement, and the change is committed:

```
IF v product count = 1 THEN
      UPDATE products
      SET price = price * p factor
      WHERE product id = p product id;
      COMMIT:
    END IF;
```

The EXCEPTION block performs a ROLLBACK if an exception is raised:

```
EXCEPTION
      WHEN OTHERS THEN
       ROLLBACK;
```

Finally, the END keyword is used to mark the end of the procedure:

```
END update product price;
```



NOTE

The repetition of the procedure name after the END keyword is not required, but it is good programming practice to put it in.

Calling a Procedure

You run (or call) a procedure using the CALL statement. The example you'll see in this section will multiply the price of product #1 by 1.5 using the procedure shown in the previous section. First, the following query retrieves the price of product #1 so you can compare it with the modified price later:

```
SELECT price
   FROM products
   WHERE product id = 1;
        PRICE
    _____
        19.95
```

The following statement calls update product price(), passing the parameter values 1 (the product id) and 1.5 (the factor by which the product price is multiplied):

```
CALL update product price(1, 1.5);
```

```
Call completed.
```

This statement shows the use of *positional notation* to indicate the values to be passed to the procedure or function. In positional notation, the position of parameters is used to assign the values passed to the procedure. In the example, the first value in the call is 1, and this is passed to the first parameter in the procedure (p product id); the second value in the call is 1.5, and this is passed to the second parameter (p factor). In Oracle Database 11g, you can also use named and mixed notation in addition to positional notation, and you'll learn about these types of notation shortly.

The next query retrieves the details for product #1 again; notice the price has been multiplied by 1.5:

```
SELECT price
   FROM products
   WHERE product id = 1;
        PRICE
        29.93
```

In Oracle Database 11g you can pass parameters using named and mixed notation. In named notation, you include the name of the parameter when calling a procedure. For example, the following statement calls update product price() using named notation; notice that the values for the p factor and p product id parameters are indicated using =>:

```
CALL update product price(p factor => 1.3, p product id => 2);
```



TIP

Named notation makes your code easier to read and maintain because the parameters are explicitly shown.

In mixed notation, you use both positional and named notation; you use positional notation for the first set of parameters and named notation for the last set of parameters. Mixed notation is useful when you have procedures and functions that have both required and optional parameters; you use positional notation for the required parameters, and named notation for the optional parameters. The following example uses mixed notation; notice that positional notation comes before named notation when specifying the parameter values:

CALL update_product_price(3, p_factor => 1.7);

Getting Information on Procedures

You can get information on your procedures from the user procedures view. Table 11-2 describes some of the columns in user procedures.

Column	Туре	Description
OBJECT_NAME	VARCHAR2(30)	The object name, which may be a procedure, function, or package name
PROCEDURE_NAME	VARCHAR2(30)	The procedure name
AGGREGATE	VARCHAR2(3)	Whether the procedure is an aggregate function (YES or NO)
IMPLTYPEOWNER	VARCHAR2(30)	The owner of the type (if any)
IMPLTYPENAME	VARCHAR2(30)	The name of the type (if any)
PARALLEL	VARCHAR2(3)	Whether the procedure is enabled for parallel queries (YES or NO)

 TABLE 11-2
 Some Columns in the user procedures View



NOTE

You can get information on all the procedures you have access to using all procedures.

The following example retrieves the object_name, aggregate, and parallel columns from user procedures for update product price():

Dropping a Procedure

You drop a procedure using DROP PROCEDURE. For example, the following statement drops update product price():

DROP PROCEDURE update product price;

Viewing Errors in a Procedure

If the database reports an error when you create a procedure, you can view the errors using the SHOW ERRORS command. For example, the following CREATE PROCEDURE statement attempts to create a procedure that has a syntax error at line 6 (the parameter should be p_dob, not p_dobs):

```
SQL> CREATE PROCEDURE update_customer_dob (

p_customer_id INTEGER, p_dob DATE

3 ) AS

4 BEGIN

5 UPDATE customers

6 SET dob = p_dobs

7 WHERE customer_id = p_customer_id;

8 END update_customer_dob;

9 /
```

Warning: Procedure created with compilation errors.

As you can see, there is a compilation error. To view the errors, you use SHOW ERRORS:

SQL> SHOW ERRORS

```
Errors for PROCEDURE UPDATE_CUSTOMER_DOB:

LINE/COL ERROR

5/3 PL/SQL: SQL Statement ignored
6/13 PL/SQL: ORA-00904: invalid column name
```

Line 5 was ignored because an invalid column name was referenced in line 6. You can fix the error by issuing an EDIT command to edit the CREATE PROCEDURE statement, changing p_dobs to p_dob, and rerunning the statement by entering /.

Functions

A function is similar to a procedure, except that a function must return a value. Together, stored procedures and functions are sometimes referred to as stored subprograms because they are small programs.

In this section, you'll learn how to

- Create a function.
- Call a function.
- Get information on functions.
- Drop a function.

Creating a Function

You create a function using the CREATE FUNCTION statement. The simplified syntax for the CREATE FUNCTION statement is as follows:

```
CREATE [OR REPLACE] FUNCTION function name
    [(parameter name [IN | OUT | IN OUT] type [, ...])]
    RETURN type
    {IS | AS}
    BEGIN
      function body
    END function name;
```

where

- OR REPLACE means the procedure is to replace an existing function.
- function name is the name of the function.
- parameter name is the name of a parameter that is passed to the function. You may pass multiple parameters to a function.
- IN |OUT | IN OUT is the mode of the parameter.
- type is the type of the parameter.
- function body contains actual code for the function. Unlike a procedure, the body of a function must return a value of the type specified in the RETURN clause.

The following example creates a function named circle area(), which returns the area of a circle. The radius of the circle is passed as a parameter named p radius to circle area(); notice that circle area() returns a NUMBER:

```
CREATE FUNCTION circle area (
      p radius IN NUMBER
    ) RETURN NUMBER AS
      v pi NUMBER := 3.1415926;
      v area NUMBER;
    BEGIN
      -- circle area is pi multiplied by the radius squared
```

```
v_area := v_pi * POWER(p_radius, 2);
RETURN v_area;
END circle_area;
/
```

The next example creates a function named average_product_price(), which returns the average price of products whose product type id equals the parameter value:

```
CREATE FUNCTION average_product_price (
    p_product_type_id IN INTEGER
) RETURN NUMBER AS
    v_average_product_price NUMBER;

BEGIN
    SELECT AVG(price)
    INTO v_average_product_price
    FROM products
    WHERE product_type_id = p_product_type_id;
    RETURN v_average_product_price;

END average_product_price;
//
```

Calling a Function

You call your own functions as you would call any of the built-in database functions; you saw how to call built-in functions in Chapter 4. (Just to refresh your memory, you can call a function using a SELECT statement that uses the dual table in the FROM clause.) The following example calls circle_area(), passing a radius of 2 meters to the function using positional notation:

```
SELECT circle_area(2)
FROM dual;

CIRCLE_AREA(2)
```

12.5663704

In Oracle Database 11*g*, you can also use named and mixed notation when calling functions. For example, the following query uses named notation when calling circle area():

```
SELECT circle_area(p_radius => 4)
FROM dual;

CIRCLE_AREA(P_RADIUS=>4)
```

50.2654816

The next example calls average_product_price(), passing the parameter value 1 to the function to get the average price of products whose product type id is 1:

Getting Information on Functions

You can get information on your functions from the user procedures view; this view was covered earlier in the section "Getting Information on Procedures." The following example retrieves the object name, aggregate, and parallel columns from user procedures for circle area() and average product price():

```
SELECT object name, aggregate, parallel
     FROM user procedures
     WHERE object name IN ('CIRCLE AREA', 'AVERAGE PRODUCT PRICE');
```

OBJECT_NAME	AGG	PAR
AVERAGE_PRODUCT_PRICE	NO	NO
CIRCLE AREA	NO	NO

Dropping a Function

You drop a function using DROP FUNCTION. For example, the following statement drops circle area():

```
DROP FUNCTION circle area;
```

Packages

In this section, you'll learn how to group procedures and functions together into packages. Packages allow you to encapsulate related functionality into one self-contained unit. By modularizing your PL/SQL code through the use of packages, you build up your own libraries of code that other programmers can reuse. In fact, the Oracle database comes with a library of packages, which allow you to access external files, manage the database, generate HTML, and much more; to see all the packages, you should consult the Oracle Database PL/SQL Packages and Types Reference manual from Oracle Corporation.

Packages are typically made up of two components: a specification and a body. The package specification lists the available procedures, functions, types, and objects. You can make the items listed in the specification available to all database users, and I refer to these items as being public (although only users you have granted privileges to access your package can use it). The specification doesn't contain the code that makes up the procedures and functions; the code is contained in the package body.

Any items in the body that are not listed in the specification are *private* to the package. Private items can be used only inside the package body. By using a combination of public and private items, you can build up a package whose complexity is hidden from the outside world. This is one of the primary goals of all programming: hide complexity from your users.

Creating a Package Specification

You create a package specification using the CREATE PACKAGE statement. The simplified syntax for the CREATE PACKAGE statement is as follows:

```
CREATE [OR REPLACE] PACKAGE package name
     {IS | AS}
       package specification
     END package name;
```

where

- package name is the name of the package.
- package_specification lists the public procedures, functions, types, and objects available to your package's users.

The following example creates a specification for a package named product package:

The t_ref_cursor type is a PL/SQL REF CURSOR type. A REF CURSOR is similar to a pointer in the C++ programming language, and it points to a cursor; as you saw earlier, a cursor allows you to read the rows returned by a query. The get_products_ref_cursor() function returns a t_ref_cursor, and, as you'll see in the next section, it points to a cursor that contains the rows retrieved from the products table.

The update_product_price() procedure multiplies the price of a product and commits the change.

Creating a Package Body

You create a package body using the CREATE PACKAGE BODY statement. The simplified syntax for the CREATE PACKAGE BODY statement is as follows:

```
CREATE [OR REPLACE] PACKAGE BODY package_name
{IS | AS}

package_body

END package name;
```

where

- package_name is the name of the package, which must match the package name in the specification.
- package body contains the code for the procedures and functions.

The following example creates the package body for product package:

```
CREATE PACKAGE BODY product_package AS
FUNCTION get_products_ref_cursor
RETURN t_ref_cursor IS
    v_products_ref_cursor t_ref_cursor;
BEGIN
    -- get the REF CURSOR
    OPEN v_products_ref_cursor FOR
    SELECT product_id, name, price
    FROM products;
```

```
-- return the REF CURSOR
   RETURN v products ref cursor;
 END get products ref cursor;
 PROCEDURE update product price (
   p product id IN products.product id%TYPE,
   p_factor IN NUMBER
 ) AS
   v product count INTEGER;
 BEGIN
   -- count the number of products with the
   -- supplied product id (will be 1 if the product exists)
   SELECT COUNT(*)
   INTO v product count
   FROM products
   WHERE product id = p product id;
   -- if the product exists (v product count = 1) then
   -- update that product's price
   IF v product count = 1 THEN
     UPDATE products
     SET price = price * p factor
     WHERE product id = p product id;
     COMMIT;
   END IF;
 EXCEPTION
   WHEN OTHERS THEN
     ROLLBACK;
 END update product price;
END product package;
```

The get products ref cursor() function opens the cursor and retrieves the product id, name, and price columns from the products table The reference to this cursor (the REF CURSOR) is stored in v products ref cursor and returned by the function.

The update product price () procedure multiplies the price of a product and commits the change. This procedure is identical to the one shown earlier in the section "Creating a Procedure," so I won't go into the details on how it works again.

Calling Functions and Procedures in a Package

When calling functions and procedures in a package, you must include the package name in the call. The following example calls product package.get products ref cursor(), which returns a reference to a cursor containing the product id, name and price for the products:

SELECT product package.get products ref cursor FROM dual;

```
GET PRODUCTS REF CUR
_____
CURSOR STATEMENT: 1
CURSOR STATEMENT: 1
```

PRODUCT_ID	NAME	PRICE
1	Modern Science	19.95
2	Chemistry	30
3	Supernova	25.99
4	Tank War	13.95
5	Z Files	49.99
6	2412: The Return	14.95
7	Space Force 9	13.49
8	From Another Planet	12.99
9	Classical Music	10.99
10	Pop 3	15.99
11	Creative Yell	14.99
12	My Front Line	13.49

The next example calls product_package.update_product_price() to multiply product #3's price by 1.25:

```
CALL product package.update product price(3, 1.25);
```

The next query retrieves the details for product #3; notice that the price has increased:

```
FROM products
WHERE product_id = 3;

PRICE
------
32.49
```

Getting Information on Functions and Procedures in a Package

You can get information on your functions and procedures in a package from the user_procedures view; this view was covered earlier in the section "Getting Information on Procedures." The following example retrieves the object_name and procedure_name columns from user_procedures for product_package:

Dropping a Package

SELECT object name, procedure name

You drop a package using DROP PACKAGE. For example, the following statement drops product_package:

```
DROP PACKAGE product package;
```

Triggers

A trigger is a procedure that is run (or fired) automatically by the database when a specified DML statement (INSERT, UPDATE, or DELETE) is run against a certain database table. Triggers are useful for doing things like advanced auditing of changes made to column values in a table.

When a Trigger Fires

A trigger may fire before or after a DML statement runs. Also, because a DML statement can affect more than one row, the code for the trigger may be run once for every row affected (a row-level trigger), or just once for all the rows (a statement-level trigger). For example, if you create a rowlevel trigger that fires for an UPDATE on a table, and you run an UPDATE statement that modified ten rows of that table, then that trigger would run ten times. If, however, your trigger was a statement-level trigger, the trigger would fire once for the whole UPDATE statement, regardless of the number of rows affected.

There is another difference between a row-level trigger and a statement-level trigger: A rowlevel trigger has access to the old and new column values when the trigger fires as a result of an UPDATE statement on that column. The firing of a row-level trigger may also be limited using a trigger condition; for example, you could set a condition that limits the trigger to fire only when a column value is less than a specified value.

Set Up for the Example Trigger

As mentioned, triggers are useful for doing advanced auditing of changes made to column values. In the next section, you'll see a trigger that records when a product's price is lowered by more than 25 percent; when this occurs, the trigger will add a row to the product price audit table. The product price audit table is created by the following statement in the store schema.sql script:

```
CREATE TABLE product price audit (
      product id INTEGER
        CONSTRAINT price audit fk products
        REFERENCES products (product id),
      old price NUMBER(5, 2),
      new price NUMBER(5, 2)
    );
```

As you can see, the product id column of the product price audit table is a foreign key to the product id column of the products table. The old price column will be used to store the old price of a product prior to the change, and the new price column will be used to store the new price after the change.

Creating a Trigger

You create a trigger using the CREATE TRIGGER statement. The simplified syntax for the CREATE TRIGGER statement is as follows:

```
CREATE [OR REPLACE] TRIGGER trigger name
     {BEFORE | AFTER | INSTEAD OF | FOR} trigger event
    ON table name
     [FOR EACH ROW]
     [{FORWARD | REVERSE} CROSSEDITION]
```

```
[{FOLLOWS | PRECEDES} schema.other_trigger}
[{ENABLE | DISABLE}]
[WHEN trigger_condition]]
BEGIN
   trigger_body
END trigger_name;
```

where

- OR REPLACE means the trigger is to replace an existing trigger, if present.
- trigger_name is the name of the trigger.
- BEFORE means the trigger fires before the triggering event is performed. AFTER means the trigger fires after the triggering event is performed. INSTEAD OF means the trigger fires instead of performing the triggering event. FOR, which is new for Oracle Database 11g, allows you to create a compound trigger consisting of up to four sections in the trigger body.
- trigger_event is the event that causes the trigger to fire.
- table name is the table that the trigger references.
- FOR EACH ROW means the trigger is a row-level trigger, that is, the code contained within trigger_body is run for each row when the trigger fires. If you omit FOR EACH ROW, the trigger is a statement-level trigger, which means the code within trigger_body is run once when the trigger fires.
- {FORWARD | REVERSE} CROSSEDITION is new for Oracle Database 11g and will typically be used by database administrators or application administrators. A FORWARD cross edition trigger is intended to fire when a DML statement makes a change in the database while an online application currently accessing the database *is being patched or upgraded* (FORWARD is the default); the code in the trigger body must be designed to handle the DML changes when the application patching or upgrade is complete. A REVERSE cross edition trigger is similar, except it is intended to fire and handle DML changes made *after the online application has been patched or upgraded*.
- {FOLLOWS | PRECEDES} schema.other_trigger is new for Oracle Database 11g and specifies whether the firing of the trigger follows or precedes the firing of another trigger specified in schema.other_trigger. You can create a series of triggers that fire in a specific order.
- {ENABLE | DISABLE} is new for Oracle Database 11g and indicates whether the trigger is initially enabled or disabled when it is created (the default is ENABLE). You enable a disabled trigger by using the ALTER TRIGGER trigger_name ENABLE statement or by enabling all triggers for a table using ALTER TABLE table_name ENABLE ALL TRIGGERS.
- trigger_condition is a Boolean condition that limits when a trigger actually runs its code.
- *trigger body* contains the code for the trigger.

The example trigger you'll see in this section fires before an update of the price column in the products table; therefore, I'll name the trigger before product price update. Also, because I want to use the price column values before and after an UPDATE statement modifies the price column's value, I must use a row-level trigger. Finally, I want to audit a price change when the new price is lowered by more than 25 percent of the old price; therefore, I'll need to specify a trigger condition to compare the new price with the old price. The following statement creates the before product price update trigger:

```
CREATE TRIGGER before product price update
    BEFORE UPDATE OF price
    ON products
    FOR EACH ROW WHEN (new.price < old.price * 0.75)
      dbms output.put line('product id = ' || :old.product id);
      dbms output.put line('Old price = ' || :old.price);
      dbms output.put line('New price = ' || :new.price);
      dbms output.put line('The price reduction is more than 25%');
      -- insert row into the product price audit table
      INSERT INTO product price audit (
        product id, old price, new price
      ) VALUES (
        :old.product id, :old.price, :new.price
    END before product price update;
```

There are five things you should notice about this statement:

- BEFORE UPDATE OF price means the trigger fires before an update of the price column.
- FOR EACH ROW means this as a row-level trigger, that is, the trigger code contained within the BEGIN and END keywords runs once for each row modified by the update.
- The trigger condition is (new.price < old.price * 0.75), which means the trigger fires only when the new price is less than 75 percent of the old price (that is, when the price is reduced by more than 25 percent).
- The new and old column values are accessed using the :old and :new aliases in the trigger.
- The trigger code displays the product id, the old and new prices, and a message stating that the price reduction is more than 25 percent. The code then adds a row to the product price audit table containing the product id and the old and new prices.

Firing a Trigger

To see the output from the trigger, you need to run the SET SERVEROUTPUT ON command:

SET SERVEROUTPUT ON

To fire the before product price update trigger, you must reduce a product's price by more than 25 percent. Go ahead and perform the following UPDATE statement to reduce the price

of products #5 and #10 by 30 percent (this is achieved by multiplying the price column by .7). The following UPDATE statement causes the before product price update trigger to fire:

```
UPDATE products

SET price = price * .7

WHERE product_id IN (5, 10);

product_id = 10
Old price = 15.99
New price = 11.19
The price reduction is more than 25%
product_id = 5
Old price = 49.99
New price = 34.99
The price reduction is more than 25%
2 rows updated.
```

As you can see, the trigger fired for products #10 and #5. You can see that the trigger did indeed add the two required rows containing the product_ids, along with the old and new prices, to the product price audit table using the following query:

```
SELECT *

FROM product_price_audit

ORDER BY product_id;
```

PRODUCT_ID	OLD_PRICE	NEW_PRICE
5	49.99	34.99
10	15.99	11.19

Getting Information on Triggers

You can get information on your triggers from the user_triggers view. Table 11-3 describes some of the columns in user triggers.

Column	Туре	Description
TRIGGER_NAME	VARCHAR2(30)	Name of the trigger.
TRIGGER_TYPE	VARCHAR2(16)	Type of the trigger.
TRIGGERING_EVENT	VARCHAR2 (227)	Event that causes the trigger to fire.
TABLE_OWNER	VARCHAR2(30)	User who owns the table that the trigger references.
BASE_OBJECT_TYPE	VARCHAR2(16)	Type of the object referenced by the trigger.
TABLE_NAME	VARCHAR2(30)	Name of the table referenced by the trigger.
COLUMN_NAME	VARCHAR2 (4000)	Name of the column referenced by the trigger.

 TABLE 11-3
 Some Columns in the user triggers View

Column	Туре	Description
REFERENCING_NAMES	VARCHAR2(128)	Name of the old and new aliases.
WHEN_CLAUSE	VARCHAR2(4000)	Trigger condition that limits when the trigger runs its code.
STATUS	VARCHAR2(8)	Whether the trigger is enabled or disabled (ENABLED or DISABLED).
DESCRIPTION	VARCHAR2(4000)	Description of the trigger.
ACTION_TYPE	VARCHAR2(11)	Action type of the trigger (CALL or PL/SQL).
TRIGGER_BODY	LONG	Code contained in the trigger body. (The LONG type allows storage of large amounts of text. You'll learn about the LONG type in Chapter 14.)

 TABLE 11-3
 Some Columns in the user triggers View (continued)



NOTE

You can get information on all the triggers you have access to using all triggers.

The following example retrieves the details of the before product price update trigger from user triggers (the output is printed pretty for clarity):

SELECT trigger name, trigger type, triggering event, table owner base object type, table name, referencing names, when clause, status, description, action_type, trigger_body FROM user_triggers WHERE trigger name = 'BEFORE PRODUCT PRICE UPDATE';

TRIGGER NAME TRIGGER TYPE BEFORE PRODUCT PRICE UPDATE BEFORE EACH ROW

TRIGGERING EVENT

UPDATE

BASE_OBJECT_TYPE TABLE_NAME STORE TABLE PRODUCTS

REFERENCING NAMES

REFERENCING NEW AS NEW OLD AS OLD

WHEN CLAUSE

new.price < old.price * 0.75

```
STATUS
-----
ENABLED

DESCRIPTION
------
before_product_price_update
BEFORE UPDATE OF
  price
ON
  products
FOR EACH ROW

ACTION_TYPE
------
PL/SQL

TRIGGER_BODY
-----
BEGIN
  dbms_output.put_line('product_id = ' || :old.product_id);
  dbms_output...
```



NOTE

You can see all the code for the trigger using the SQL*Plus SET LONG command, for example, SET LONG 1000.

Disabling and Enabling a Trigger

You can stop a trigger from firing by disabling it by using the ALTER TRIGGER statement. The following example disables the before product price update trigger:

ALTER TRIGGER before_product_price_update DISABLE;

The next example enables the before product price update trigger:

ALTER TRIGGER before_product_price_update ENABLE;

Dropping a Trigger

You drop a trigger using DROP TRIGGER. The following example drops the before_product_price update trigger:

DROP TRIGGER before product price update;

New Oracle Database 11g PL/SQL Features

In this section, you'll see some of the new PL/SQL features introduced in Oracle Database 11*g*. Specifically, the following will be discussed:

- The SIMPLE_INTEGER type
- Support for sequences in PL/SQL
- PL/SQL native machine code generation

SIMPLE INTEGER Type

The SIMPLE INTEGER type is a subtype of BINARY INTEGER; the SIMPLE INTEGER can store the same range as BINARY INTEGER, except SIMPLE INTEGER cannot store a NULL value. The range of values SIMPLE INTEGER can store is -2^{31} (-2,147,483,648) to 2^{31} (2,147,483,648).

Arithmetic overflow is truncated when using SIMPLE INTEGER values; therefore, calculations don't raise an error when overflow occurs. Because overflow errors are ignored, the values stored in a SIMPLE INTEGER can wrap from positive to negative and from negative to positive, as, for example:

```
2^{30} + 2^{30} = 0x40000000 + 0x40000000 = 0x80000000 = -2^{31}
-2^{31} + -2^{31} = 0 \times 800000000 + 0 \times 800000000 = 0 \times 0000000000 = 0
```

In the first example, two positive values are added, and a negative total is produced. In the second example, two negative values are added, and zero is produced.

Because overflow is ignored and truncated when using SIMPLE INTEGER values in calculations, SIMPLE INTEGER offers much better performance than BINARY INTEGER when the DBA configures the database to compile PL/SQL to native machine code. Because of this benefit, you should use SIMPLE INTEGER in your PL/SQL code when you don't need to store a NULL and you don't care about overflow truncation occurring in your calculations; otherwise, you should use BINARY INTEGER.

The following get area() procedure shows the use of the SIMPLE INTEGER type; get area() calculates and displays the area of a rectangle:

```
CREATE PROCEDURE get area
     AS
       v width SIMPLE INTEGER := 10;
       v height SIMPLE INTEGER := 2;
       v area SIMPLE INTEGER := v width * v height;
       DBMS OUTPUT.PUT LINE('v area = ' || v area);
     END get area;
```



You'll find this example, and the other examples in this section, in a script named plsql 11g examples.sql in the SQL directory. You may run this script if you are using Oracle Database 11g.

The following example shows the execution of get area():

```
SET SERVEROUTPUT ON
  CALL get area();
   v area = 20
```

As expected, the calculated area is 20.

Sequences in PL/SQL

In the previous chapter you saw how to create and use sequences of numbers in SQL. In Oracle Database 11g, you can also use sequences in PL/SQL code.

As a reminder, a sequence generates a series of numbers. When you create a sequence in SQL, you can specify its initial value and an increment for the series of subsequent numbers.

You use the currval pseudo column to get the current value in the sequence and nextval to generate the next number. Before you access currval, you must first use nextval to generate an initial number.

The following statement creates a table named new products; this table will be used shortly:

```
CREATE TABLE new products (
      product id INTEGER CONSTRAINT new products pk PRIMARY KEY,
      name VARCHAR2(30) NOT NULL,
      price NUMBER(5, 2)
    );
```

The next statement creates a sequence named s product id:

```
CREATE SEQUENCE s_product_id;
```

The following statement creates a procedure named add new products, which uses s product id to set the product id column in a row added to the new products table; notice the use of the nextval and currval pseudo columns in the PL/SQL code (this is new for Oracle Database 11g):

```
CREATE PROCEDURE add new products
       v product id BINARY INTEGER;
     BEGIN
       -- use nextval to generate the initial sequence number
       v product id := s product id.nextval;
       DBMS OUTPUT.PUT LINE('v product id = ' || v product id);
       -- add a row to new products
       INSERT INTO new products
       VALUES (v product id, 'Plasma Physics book', 49.95);
       DBMS OUTPUT.PUT LINE('s product id.currval = ' || s product id.currval);
       -- use nextval to generate the next sequence number
       v product id := s product id.nextval;
       DBMS OUTPUT.PUT LINE('v product id = ' || v product id);
       -- add another row to new products
       INSERT INTO new products
       VALUES (v product id, 'Quantum Physics book', 69.95);
       DBMS OUTPUT.PUT LINE('s product id.currval = ' || s product id.currval);
     END add new products;
```

The following example runs add new products () and shows the contents of the new products table:

```
SET SERVEROUTPUT ON
   CALL add_new_products();
   v product id = 1
```

```
s product id.currval = 1
v product id = 2
s product id.currval = 2
```

SELECT * FROM new products;

PRODUCT_ID	NAME	PRICE
1	Plasma Physics book	49.95
2	Quantum Physics book	69.95

As expected, two rows were added to the table.

PL/SQL Native Machine Code Generation

By default, each PL/SQL program unit is compiled into intermediate form, machine-readable code. This machine-readable code is stored in the database and interpreted every time the code is run. With PL/SQL native compilation, the PL/SQL is turned into native code and stored in shared libraries. Native code runs much faster than intermediate code because native code doesn't have to be interpreted before it runs.

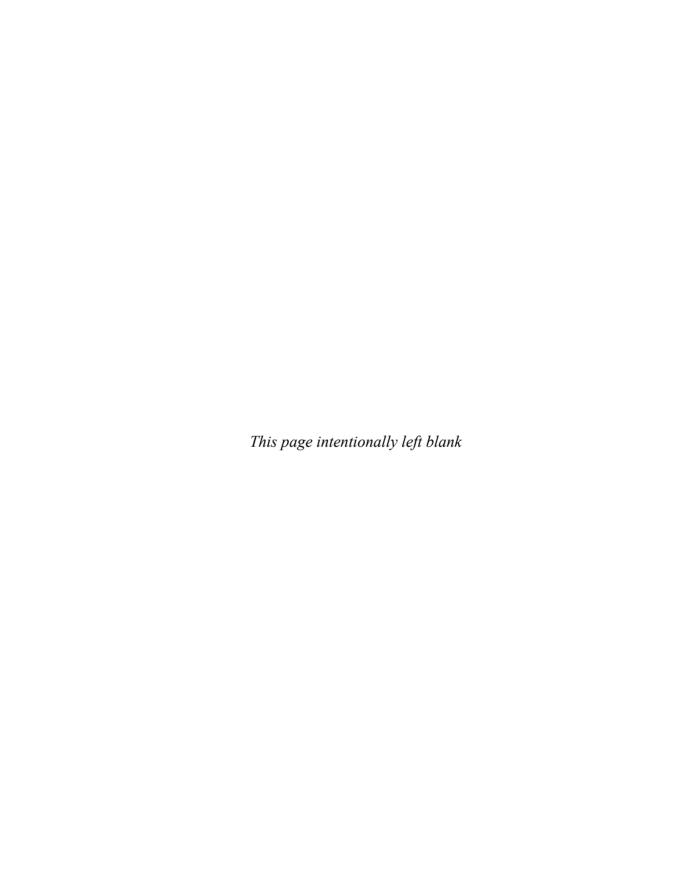
In certain versions of the database prior to Oracle Database 11g, you can compile PL/SQL code to C code, and then compile the C code into machine code; this is a very laborious and problematic process. In Oracle Database 11g, the PL/SQL complier can generate native machine code directly. Setting up the database to generate native machine code should be done only by an experienced DBA (as such, its coverage is beyond the scope of this book). You can read all about PL/SQL native machine code generation in the PL/SQL User's Guide and Reference manual from Oracle Corporation.

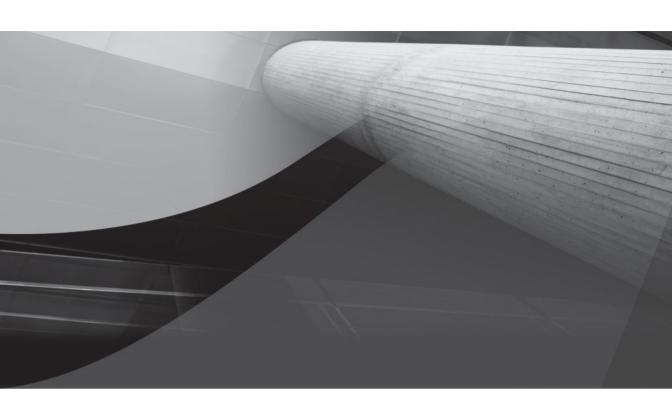
Summary

In this chapter, you learned the following:

- PL/SQL programs are divided up into blocks containing PL/SQL and SQL statements.
- A loop, such as a WHILE or FOR loop, runs statements multiple times.
- A cursor allows PL/SQL to read the rows returned by a query.
- Exceptions are used to handle run-time errors that occur in your PL/SQL code.
- A procedure contains a group of statements. Procedures allow you to centralize your business logic in the database and may be run by any program that accesses the database.
- A function is similar to a procedure except that a function must return a value.
- You can group procedures and functions together into packages, which encapsulate related functionality into one self-contained unit.
- A trigger is a procedure that is run automatically by the database when a specific INSERT, UPDATE, or DELETE statement is run. Triggers are useful for doing things like advanced auditing of changes made to column values in a table.

In the next chapter, you'll learn about database objects.





CHAPTER 12

Database Objects



n this chapter, you will do the following:

- Learn about objects in the database
- Learn how to create object types containing attributes and methods
- Use object types to define column objects and object tables
- Create and manipulate objects in SQL and PL/SQL
- Learn how a type may inherit from another type and create hierarchies of types
- Define your own constructors to set the attributes of an object
- See how to override a method in one type with a method from another type

Introducing Objects

Object-oriented programming languages such as Java, C++, and C# allow you to define classes, and these classes act as templates from which you can create objects. Classes define attributes and methods; attributes are used to store an object's state, and methods are used to model an object's behaviors.

With the release of Oracle Database 8, objects became available within the database, and object features have been improved upon in subsequent product releases. The availability of objects in the database was a major breakthrough because they enable you to define your own classes, known as *object types*, in the database. Like classes in Java and C#, database object types can contain attributes and methods. Object types are also sometimes known as user-defined types.

A simple example of an object type would be a type that represents a product. This object type could contain attributes for the product's name, description, price, and, in the case of a product that is perishable, the number of days the product can sit on the shelf before it must be thrown away. This product object type could also contain a method that returns the sell-by date of the product, based on the shelf life of the product and the current date. Another example of an object type is one that represents a person; this object type could store attributes for the person's first name, last name, date of birth, and address; the person's address could itself be represented by an object type, and it could store things like the street, city, state, and zip code. In this chapter you'll see examples of object types that represent a product, person, and address. You'll also see how to create tables from those object types, populate those tables with actual objects, and manipulate those objects in SQL and PL/SQL.

I've provided an SQL*Plus script named object_schema.sql in the SQL directory, which creates a user named object_user with a password of object_password. This script also creates the types and tables, performs the various INSERT statements, and creates the PL/SQL code shown in the first part of this chapter. You must run this script while logged in as a user with the required privileges to create a new user with the CONNECT, RESOURCE, and CREATE PUBLIC SYNONYM privileges; I log in as the system user on my database to run the scripts. After the script completes, you will be logged in as object user.

Creating Object Types

You create an object type using the CREATE TYPE statement. The following example uses the CREATE TYPE statement to create an object type named t address. This object type is used to represent an address and contains four attributes named street, city, state, and zip:

```
CREATE TYPE t address AS OBJECT (
     street VARCHAR2(15),
     city VARCHAR2 (15),
     state CHAR(2),
     zip VARCHAR2(5)
   );
```

The example shows that each attribute is defined using a database type. For example, street is defined as VARCHAR2 (15). As you'll see shortly, the type of an attribute can itself be an object type.

The next example creates an object type named t person; notice that t person has an attribute named address, which is of type t address:

```
CREATE TYPE t person AS OBJECT (
      id
         INTEGER,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      dob DATE,
      phone VARCHAR2(12),
      address t address
    );
```

The following example creates an object type named to product that will be used to represent products; notice that this type declares a function named get sell by date() using the MEMBER FUNCTION clause:

```
CREATE TYPE t product AS OBJECT (
     id INTEGER,
     name VARCHAR2(15),
     description VARCHAR2 (22),
     price NUMBER(5, 2),
     days valid INTEGER,
     -- get sell by date() returns the date by which the
     -- product must be sold
     MEMBER FUNCTION get sell by date RETURN DATE
    );
```

Because t product contains a method declaration, a body for t product must also be created. The body contains the actual code for the method, and the body is created using the CREATE TYPE BODY statement. The following example creates the body for t product; notice the body contains the code for the get sell by date() function.

```
CREATE TYPE BODY t_product AS

-- get_sell_by_date() returns the date by which the

-- product must be sold

MEMBER FUNCTION get_sell_by_date RETURN DATE IS

v_sell_by_date DATE;

BEGIN

-- calculate the sell by date by adding the days_valid attribute

-- to the current date (SYSDATE)

SELECT days_valid + SYSDATE

INTO v_sell_by_date

FROM dual;

-- return the sell by date

RETURN v_sell_by_date;

END;

END;
```

As you can see, get_sell_by_date() calculates and returns the date by which the product must be sold; it does this by adding the days_valid attribute to the current date returned by the built-in database SYSDATE() function.

You can also create a public synonym for a type, which enables all users to see the type and use it to define columns in their own tables. The following example creates a public synonym named t pub product for t product:

CREATE PUBLIC SYNONYM t_pub_product FOR t_product;

Using DESCRIBE to Get Information on Object Types

You can use the DESCRIBE command to get information on an object type. The following examples show the t_address, t_person, and t_product types:

DESCRIBE t address

1	Name	Null?	Type
(STREET CITY STATE ZIP		VARCHAR2 (15) VARCHAR2 (15) CHAR (2) VARCHAR2 (5)
	ESCRIBE t_person	Null2	Type

Name	Null?	Type
ID		NUMBER (38)
FIRST_NAME		VARCHAR2(10)
LAST_NAME		VARCHAR2(10)
DOB		DATE
PHONE		VARCHAR2(12)
ADDRESS		T_ADDRESS

DESCRIBE t product

```
Null? Type
Name
 ΙD
                                                      NUMBER (38)
                                                      VARCHAR2 (10)
 NAME
                                                      VARCHAR2 (22)
 DESCRIPTION
 PRICE
                                                      NUMBER (5,2)
 DAYS VALID
                                                      INTEGER
METHOD
 MEMBER FUNCTION GET SELL BY DATE RETURNS DATE
```

You can set the depth to which DESCRIBE will show information for embedded types using SET DESCRIBE DEPTH. The following example sets the depth to 2 and then describes t person again; notice that the attributes of address are displayed, which is an embedded object of type t address:

SET DESCRIBE DEPTH 2 DESCRIBE t person

Name	Null?	Type
ID		NUMBER (38)
FIRST_NAME		VARCHAR2 (10)
LAST_NAME		VARCHAR2 (10)
DOB		DATE
PHONE		VARCHAR2 (12)
ADDRESS		T_ADDRESS
STREET		VARCHAR2 (15)
CITY		VARCHAR2(15)
STATE		CHAR(2)
ZIP		VARCHAR2(5)

Using Object Types in Database Tables

Now that you've seen how to create object types, let's look at how you use these types in database tables. You can use an object type to define an individual column in a table, and the objects subsequently stored in that column are known as *column objects*. You can also use an object type to define an entire row in a table; the table is then known as an object table. Finally, you can use an object reference to access an individual row in an object table; an object reference is similar to a pointer in C++. You'll see examples of column objects, object tables, and object references in this section.

Column Objects

The following example creates a table named products that contains a column named product of type t product; the table also contains a column named quantity in stock, which is used to store the number of those products currently in stock:

```
CREATE TABLE products (
     product t product,
     quantity in stock INTEGER
    );
```

When adding a row to this table, you must use a *constructor* to supply the attribute values for the new t_product object; as a reminder, the t_product type was created using the following statement:

A constructor is a built-in method for the object type, and it has the same name as the object type; the constructor accepts parameters that are used to set the attributes of the new object. The constructor for the t_product type is named t_product and accepts five parameters, one to set each of the attributes; for example, t_product(1, pasta, 20 oz bag of pasta, 3.95, 10) creates a new t_product object and sets its id to 1, name to pasta, description to 20 oz bag of pasta, price to 3.95, and days valid to 10.

The following INSERT statements add two rows to the products table; notice the use of the t_product constructor to supply the attribute values for the product column objects:

```
INSERT INTO products (
    product,
    quantity_in_stock
) VALUES (
    t_product(1, 'pasta', '20 oz bag of pasta', 3.95, 10),
    50
);

INSERT INTO products (
    product,
    quantity_in_stock
) VALUES (
    t_product(2, 'sardines', '12 oz box of sardines', 2.99, 5),
    25
);
```

The following query retrieves these rows from the products table; notice that the product column objects' attributes are displayed within a constructor for t product:

SELECT * FROM products;

```
PRODUCT(ID, NAME, DESCRIPTION, PRICE, DAYS VALID)
QUANTITY IN STOCK
_____
T PRODUCT(1, 'pasta', '20 oz bag of pasta', 3.95, 10)
T PRODUCT(2, 'sardines', '12 oz box of sardines', 2.99, 5)
```

You can also retrieve an individual column object from a table; to do this, you must supply a table alias through which you select the object. The following query retrieves product #1 from the products table; notice the use of the table alias p for the products table, through which the product object's id attribute is specified in the WHERE clause:

```
SELECT p.product
    FROM products p
    WHERE p.product.id = 1;
    PRODUCT(ID, NAME, DESCRIPTION, PRICE, DAYS VALID)
    _____
    T PRODUCT(1, 'pasta', '20 oz bag of pasta', 3.95, 10)
```

The next query explicitly includes the product object's id, name, price, and days valid attributes in the SELECT statement, plus the quantity in stock:

```
SELECT p.product.id, p.product.name,
    p.product.price, p.product.days valid, p.quantity in stock
    FROM products p
    WHERE p.product.id = 1;
    PRODUCT.ID PRODUCT.NA PRODUCT.PRICE PRODUCT.DAYS VALID QUANTITY IN STOCK
    3.95
```

The t product object type contains a function named get sell by date(), which calculates and returns the date by which the product must be sold. The function does this by adding the days valid attribute to the current date, which is obtained from the database using the SYSDATE() function. You can call the get sell by date() function using a table alias, as shown in the following query that uses the table alias p for the products table:

```
SELECT p.product.get sell by date()
   FROM products p;
```

```
P.PRODUCT
_____
19-JUN-07
13-JUN-07
```

Of course, if you run this guery your dates will be different, because they are calculated using SYSDATE (), which returns the current date and time.

The following UPDATE statement modifies the description of product #1; notice that the table alias p is used again:

```
UPDATE products p
SET p.product.description = '30 oz bag of pasta'
WHERE p.product.id = 1;

1 row updated.
```

The following DELETE statement removes product #2:

```
DELETE FROM products p
WHERE p.product.id = 2;

1 row deleted.
```

ROLLBACK;



NOTE

If you run these UPDATE and DELETE statements, make sure you execute the ROLLBACK so that your example data matches that shown in the rest of this chapter.

Object Tables

You can use an object type to define an entire table, and such a table is known as an object table. The following example creates an object table named <code>object_products</code>, which stores objects of type <code>t_product</code>; notice the use of the <code>OF</code> keyword to identify the table as an object table of type <code>t_product</code>:

```
CREATE TABLE object products OF t product;
```

When inserting a row into an object table, you can choose whether to use a constructor to supply attribute values or to supply the values in the same way that you would supply column values in a relational table. The following INSERT statement adds a row to the <code>object_products</code> table using the constructor for <code>t_product</code>:

```
INSERT INTO object_products VALUES (
    t_product(1, 'pasta', '20 oz bag of pasta', 3.95, 10)
);
```

The next INSERT statement omits the constructor for t_product; notice that the attribute values for t_product are supplied in the same way that columns would be in a relational table:

```
INSERT INTO object_products (
   id, name, description, price, days_valid
) VALUES (
   2, 'sardines', '12 oz box of sardines', 2.99, 5
);
```

The following query retrieves these rows from the object products table:

```
SELECT *
FROM object products;
```

ID	NAME	DESCRIPTION	PRICE	DAYS_VALID
1	pasta	20 oz bag of pasta	3.95	10
2	sardines	12 oz box of sardines	2.99	5

You can also specify individual object attributes in a query; for example, by doing this:

SELECT id, name, price FROM object products op WHERE id = 1:

ID	NAME	PRICE
1	pasta	3.95

or this:

SELECT op.id, op.name, op.price FROM object products op WHERE op.id = 1;

ID	NAME	PRICE
1	pasta	3.95

You can use the built-in Oracle database VALUE () function to select a row from an object table. VALUE () treats the row as an actual object and returns the attributes for the object within a constructor for the object type. VALUE () accepts a parameter containing a table alias, as shown in the following query:

SELECT VALUE (op) FROM object products op;

```
VALUE (OP) (ID, NAME, DESCRIPTION, PRICE, DAYS VALID)
_____
T PRODUCT(1, 'pasta', '20 oz bag of pasta', 3.95, 10)
T PRODUCT(2, 'sardines', '12 oz box of sardines', 2.99, 5)
```

You can also add an object attribute after VALUE():

SELECT VALUE(op).id, VALUE(op).name, VALUE(op).price FROM object products op;

```
VALUE (OP) . ID VALUE (OP) . VALUE (OP) . PRICE
_____
                         3.95
       1 pasta
       2 sardines
                         2.99
```

The following UPDATE statement modifies the description of product #1:

UPDATE object products SET description = '25 oz bag of pasta' WHERE id = 1;

1 row updated.

The following DELETE statement removes product #2:

```
DELETE FROM object_products
WHERE id = 2;

1 row deleted.
```

ROLLBACK:

Let's take a look at a more complex object table. The following CREATE TABLE statement creates an object table named object customers, which stores objects of type t person:

```
CREATE TABLE object customers OF t person;
```

The t_person type contains an embedded t_address object; t_person was created using the following statement:

```
CREATE TYPE t_person AS OBJECT (
id INTEGER,
first_name VARCHAR2(10),
last_name VARCHAR2(10),
dob DATE,
phone VARCHAR2(12),
address t_address
);
```

The following INSERT statements add two rows into object_customers. The first INSERT uses constructors for t_person and t_address, while the second INSERT omits the t_person constructor:

```
INSERT INTO object_customers VALUES (
   t_person(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
        t_address('2 State Street', 'Beantown', 'MA', '12345')
   )
);

INSERT INTO object_customers (
   id, first_name, last_name, dob, phone,
   address
) VALUES (
   2, 'Cynthia', 'Green', '05-FEB-1968', '800-555-1212',
   t_address('3 Free Street', 'Middle Town', 'CA', '12345')
);
```

The following query retrieves these rows from the <code>object_customers</code> table; notice that the attributes for the embedded <code>address</code> column object are displayed within the <code>t_address</code> constructor:

```
SELECT *
FROM object customers;
```

```
ID FIRST NAME LAST NAME DOB
______ _____
ADDRESS (STREET, CITY, STATE, ZIP)
_____
      1 John Brown 01-FEB-55 800-555-1211
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345')
      2 Cvnthia Green 05-FEB-68 800-555-1212
T ADDRESS('3 Free Street', 'Middle Town', 'CA', '12345')
```

The next query retrieves customer #1 from object customers; notice the use of the table alias oc through which the id attribute is specified in the WHERE clause:

```
SELECT *
   FROM object customers oc
   WHERE oc.id = 1;
```

```
ID FIRST NAME LAST NAME DOB
______ _____
ADDRESS (STREET, CITY, STATE, ZIP)
_____
           Brown 01-FEB-55 800-555-1211
     1 John
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345')
```

In the following query, a customer is retrieved based on the state attribute of the address column object:

```
SELECT *
```

```
FROM object_customers oc
WHERE oc.address.state = 'MA';
```

```
ID FIRST NAME LAST NAME DOB PHONE
______
ADDRESS (STREET, CITY, STATE, ZIP)
_____
            Brown
                   01-FEB-55 800-555-1211
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345')
```

In the next query, the id, first name, and last name attributes of customer #1 are explicitly included in the SELECT statement, along with the attributes of the embedded address column object:

```
SELECT oc.id, oc.first name, oc.last name,
     oc.address.street, oc.address.city, oc.address.state, oc.address.zip
    FROM object customers oc
    WHERE oc.id = 1;
```

```
ID FIRST NAME LAST_NAME ADDRESS.STREET ADDRESS.CITY AD ADDRE
1 John Brown 2 State Street Beantown
                                 MA 12345
```

Object Identifiers and Object References

Each object in an object table has a unique *object identifier* (OID), and you can retrieve the OID for an object using the REF() function. For example, the following query retrieves the OID for customer #1 in the object customers table:

```
SELECT REF(oc)
FROM object_customers oc
WHERE oc.id = 1;

REF(OC)
-----
0000280209D66AB93F991647649D78D08B267EE44858C7B9989D9D40689FB4DA92820
AFFE2010003280000
```

The long string of numbers and letters are the OID, which identifies the location of the object in the database. You can store an OID in an object reference and later access the object it refers to. An object reference, which is similar to a pointer in C++, points to an object stored in an object table using the OID. You may use object references to model relationships between object tables, and, as you'll see later, you can use object references in PL/SQL to access objects.

You use the REF type to define an object reference; the following statement creates a table named purchases that contains two object reference columns named customer_ref and product ref:

```
CREATE TABLE purchases (

id INTEGER PRIMARY KEY,

customer_ref REF t_person SCOPE IS object_customers,

product_ref REF t_product SCOPE IS object_products
);
```

The SCOPE IS clause restricts an object reference to point to objects in a specific table. For example, the customer_ref column is restricted to point to objects in the object_customers table only; similarly, the product_ref column is restricted to point to objects in the object_products table only.

As I mentioned earlier, each object in an object table has a unique object identifier (OID) that you can store in an object reference; you can retrieve an OID using the REF() function and store it in an object reference. For example, the following INSERT statement adds a row to the purchases table; notice that the REF() function is used in the queries to get the object identifiers for customer #1 and product #1 from the object customers and object products tables:

```
INSERT INTO purchases (
    id,
    customer_ref,
    product_ref
) VALUES (
    1,
    (SELECT REF(oc) FROM object_customers oc WHERE oc.id = 1),
    (SELECT REF(op) FROM object_products op WHERE op.id = 1)
);
```

This example records that customer #1 purchased product #1.

The following query selects the row from the purchases table; notice that the customer ref and product ref columns contain references to the objects in the object customers and object products tables:

SELECT * FROM purchases;

```
ΙD
_____
CUSTOMER REF
PRODUCT REF
______
0000220208D66AB93F991647649D78D08B267EE44858C7B9989D9D40689FB4DA92820
0000220208662E2AB4256711D6A1B50010A4E7AE8A662E2AB2256711D6A1B50010A4E
7AE8A
```

You can retrieve the actual objects stored in an object reference using the DEREF() function, which accepts an object reference as a parameter and returns the actual object. For example, the following query uses DEREF() to retrieve customer #1 and product #1 through the customer ref and product ref columns of the purchases table:

SELECT DEREF(customer ref), DEREF(product ref) FROM purchases;

```
DEREF (CUSTOMER REF) (ID, FIRST NAME, LAST NAME, DOB, PHONE,
ADDRESS (STREET, CITY,
_____
DEREF (PRODUCT REF) (ID, NAME, DESCRIPTION, PRICE, DAYS VALID)
_____
T_PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'))
T PRODUCT(1, 'pasta', '20 oz bag of pasta', 3.95, 10)
```

The next query retrieves the customer's first name and address.street attributes, plus the product's name attribute:

SELECT DEREF(customer ref).first name, DEREF(customer ref).address.street, DEREF(product ref).name FROM purchases;

```
DEREF (CUST DEREF (CUSTOMER_ DEREF (PROD
_____
       2 State Street pasta
```

The following UPDATE statement modifies the product ref column to point to product #2:

```
UPDATE purchases SET product ref = (
     SELECT REF(op) FROM object products op WHERE op.id = 2
   ) WHERE id = 1:
   1 row updated.
```

The following query verifies this change:

```
SELECT DEREF(customer_ref), DEREF(product_ref)
FROM purchases;
```

```
DEREF(CUSTOMER_REF)(ID, FIRST_NAME, LAST_NAME, DOB, PHONE, ADDRESS(STREET, CITY,

DEREF(PRODUCT_REF)(ID, NAME, DESCRIPTION, PRICE, DAYS_VALID)

T_PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
 T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'))
T PRODUCT(2, 'sardines', '12 oz box of sardines', 2.99, 5)
```

Comparing Object Values

You can compare the value of two objects in a WHERE clause of a query using the equality operator (=). For example, the following query retrieves customer #1 from the object_customers table:

The next query retrieves product #1 from the object_products table:

```
SELECT op.id, op.name, op.price, op.days_valid
FROM object_products op
WHERE VALUE(op) = t_product(1, 'pasta', '20 oz bag of pasta', 3.95, 10);
```

You can also use the <> and IN operators in the WHERE clause:

```
SELECT op.id, op.name, op.price, op.days valid
FROM object products op
WHERE VALUE(op) IN t product(1, 'pasta', '20 oz bag of pasta', 3.95, 10);
     ID NAME
              PRICE DAYS VALID
______
                     3.95 10
      1 pasta
```

If you want to use an operator like <, >, <=, >=, LIKE, or BETWEEN, you need to provide a map function for the type. A map function must return a single value of one of the built-in types that the database can then use to compare two objects. The value returned by the map function will be different for every object type, and you need to figure out what the best attribute, or concatenation of attributes, represents an object's value. For example, with the t product type, I'd return the price attribute; with the t_person type, I'd return a concatenation of the last name and first name attributes.

The following statements create a type named t person2 that contains a map function named get string(); notice that get string() returns a VARCHAR2 string containing a concatenation of the last name and first name attributes:

```
CREATE TYPE t person2 AS OBJECT (
      id INTEGER,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      dob DATE,
      phone VARCHAR2(12),
      address t address,
      -- declare the get string() map function,
      -- which returns a VARCHAR2 string
      MAP MEMBER FUNCTION get string RETURN VARCHAR2
    );
    CREATE TYPE BODY t person2 AS
      -- define the get string() map function
      MAP MEMBER FUNCTION get string RETURN VARCHAR2 IS
      BEGIN
        -- return a concatenated string containing the
        -- last name and first name attributes
        RETURN last name || ' ' || first_name;
      END get string;
    END;
```

As you'll see shortly, the database will automatically call get string() when comparing t person2 objects.

The following statements create a table named object customers2 and add rows to it:

```
CREATE TABLE object customers2 OF t person2;
     INSERT INTO object customers2 VALUES (
```

```
t_person2(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
    t_address('2 State Street', 'Beantown', 'MA', '12345')
);

INSERT INTO object_customers2 VALUES (
    t_person2(2, 'Cynthia', 'Green', '05-FEB-1968', '800-555-1212',
        t_address('3 Free Street', 'Middle Town', 'CA', '12345')
);
```

The following query uses > in the WHERE clause:

```
SELECT oc2.id, oc2.first_name, oc2.last_name, oc2.dob

FROM object_customers2 oc2

WHERE VALUE(oc2) >

t_person2(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',

t_address('2 State Street', 'Beantown', 'MA', '12345')

);

ID FIRST_NAME LAST_NAME DOB

2 Cynthia Green 05-FEB-68
```

When the query is executed, the database automatically calls get_string() to compare the objects in the object_customers2 table to the object after the > in the WHERE clause. The get_string() function returns a concatenation of the last_name and first_name attributes of the objects, and because Green Cynthia is greater than Brown John, she is returned by the query.

Using Objects in PL/SQL

You can create and manipulate objects in PL/SQL. In this section, you'll see the use of a package named product_package, which is created when you run the object_schema.sql script; product_package contains the following methods:

- A function named get_products () that returns a REF CURSOR that points to the objects in the object products table
- A procedure named display_product() that displays the attributes of a single object in the object products table
- A procedure named insert_product() that adds an object to the object_products table
- A procedure named update_product_price() that updates the price attribute of an object in the object products table
- A function named get_product() that returns a single object from the object_products table
- A procedure named update_product() that updates all the attributes of an object in the object products table

- A function named get product ref() that returns a reference to a single object from the object products table
- A procedure named delete product () that deletes a single object from the object products table

The object schema.sql script contains the following package specification:

```
CREATE PACKAGE product package AS
     TYPE t ref cursor IS REF CURSOR;
     FUNCTION get products RETURN t ref cursor;
     PROCEDURE display product (
       p id IN object products.id%TYPE
     PROCEDURE insert product (
       p id
               IN object products.id%TYPE,
       p name IN object products.name%TYPE,
       p description IN object products.description%TYPE,
       p price IN object products.price%TYPE,
       p days valid IN object products.days valid%TYPE
     PROCEDURE update product price (
       p id IN object products.id%TYPE,
       p factor IN NUMBER
     FUNCTION get product(
       p id IN object products.id%TYPE
     ) RETURN t product;
     PROCEDURE update product(
       p product t product
     );
     FUNCTION get product ref(
       p id IN object products.id%TYPE
     ) RETURN REF t product;
     PROCEDURE delete product (
       p id IN object products.id%TYPE
   END product package;
```

You'll see the methods in the body of product package in the following sections.

The get_products() Function

The get products () function returns a REF CURSOR that points to the objects in the object products table; get products () is defined as follows in the body of product package:

```
FUNCTION get products
   RETURN t ref cursor IS
     -- declare a t ref cursor object
     v products ref cursor t ref cursor;
   BEGIN
```

```
-- get the REF CURSOR

OPEN v_products_ref_cursor FOR

SELECT VALUE(op)

FROM object_products op

ORDER BY op.id;

-- return the REF CURSOR

RETURN v_products_ref_cursor;

END get_products;
```

The following query calls product_package.get_products() to retrieve the products from object products:

SELECT product_package.get_products FROM dual;

```
GET_PRODUCTS

CURSOR STATEMENT: 1

CURSOR STATEMENT: 1

VALUE(OP)(ID, NAME, DESCRIPTION, PRICE, DAYS_VALID)

T_PRODUCT(1, 'pasta', '20 oz bag of pasta', 3.95, 10)

T PRODUCT(2, 'sardines', '12 oz box of sardines', 2.99, 5)
```

The display_product() Procedure

The display_product() procedure displays the attributes of a single object in the object_products table; display product() is defined as follows in the body of product package:

```
PROCEDURE display product (
     p id IN object products.id%TYPE
    ) AS
     -- declare a t product object named v product
      v product t product;
    BEGIN
      -- attempt to get the product and store it in v product
      SELECT VALUE (op)
      INTO v product
      FROM object products op
      WHERE id = p id;
      -- display the attributes of v product
      DBMS OUTPUT.PUT LINE('v product.id=' ||
        v product.id);
      DBMS OUTPUT.PUT LINE('v product.name=' ||
        v product.name);
      DBMS OUTPUT.PUT LINE('v product.description=' ||
        v product.description);
      DBMS OUTPUT.PUT LINE('v product.price=' ||
```

```
v product.price);
  DBMS OUTPUT.PUT LINE('v product.days valid=' ||
    v product.days valid);
  -- call v product.get sell by date() and display the date
  DBMS OUTPUT.PUT LINE('Sell by date=' ||
    v product.get sell by date());
END display product;
```

The following example calls product package.display product (1) to retrieve product #1 from the object products table:

SET SERVEROUTPUT ON

```
CALL product package.display product(1);
v product.id=1
v product.name=pasta
v product.description=20 oz bag of pasta
v product.price=3.95
v product.days valid=10
Sell by date=25-JUN-07
```

The insert_product() Procedure

The insert product() procedure adds an object to the object products table; insert product () is defined as follows in the body of product package:

```
PROCEDURE insert product(
                    IN object products.id%TYPE,
      p id
      p name
                   IN object products.name%TYPE,
      p description IN object products.description%TYPE,
      p price IN object products.price%TYPE,
      p days valid IN object products.days valid%TYPE
       -- create a t product object named v product
      v product t product :=
        t product(
          p id, p name, p description, p price, p days valid
        );
    BEGIN
      -- add v product to the object products table
      INSERT INTO object products VALUES (v product);
      COMMIT;
    EXCEPTION
      WHEN OTHERS THEN
        ROLLBACK;
    END insert product;
```

The following example calls product package.insert product() to add a new object to the object products table:

```
CALL product package.insert product(3, 'salsa',
     '15 oz jar of salsa', 1.50, 20);
```

The update_product_price() Procedure

The update product price() procedure updates the price attribute of an object in the object products table; update product price() is defined as follows in the body of product package:

```
PROCEDURE update product price (
           IN object products.id%TYPE,
     p factor IN NUMBER
   ) AS
     -- declare a t product object named v product
     v product t product;
   REGIN
     -- attempt to select the product for update and
     -- store the product in v product
     SELECT VALUE (op)
     INTO v product
     FROM object products op
     WHERE id = p id
     FOR UPDATE;
     -- display the current price of v product
     DBMS OUTPUT.PUT LINE('v product.price=' ||
       v product.price);
     -- multiply v product.price by p factor
     v product.price := v product.price * p factor;
     DBMS OUTPUT.PUT LINE('New v product.price=' ||
       v product.price);
     -- update the product in the object products table
     UPDATE object products op
     SET op = v product
     WHERE id = p id;
     COMMIT;
   EXCEPTION
     WHEN OTHERS THEN
      ROLLBACK;
   END update product price;
```

The following example calls product package.update product price() to update the price of product #3 in the object products table:

```
CALL product package.update product price(3, 2.4);
   v product.price=1.5
   New v product.price=3.6
```

The get_product() Function

The get product() function returns a single object from the object products table; get product () is defined as follows in the body of product package:

```
FUNCTION get product(
     p id IN object products.id%TYPE
   RETURN t product IS
     -- declare a t product object named v product
     v product t product;
     -- get the product and store it in v product
     SELECT VALUE (op)
     INTO v product
     FROM object products op
     WHERE op.id = p id;
     -- return v product
     RETURN v product;
    END get product;
```

The following query calls product package.get product() to get product #3 from the object products table:

```
SELECT product package.get product(3)
   FROM dual;
```

```
PRODUCT PACKAGE.GET PRODUCT(3)(ID, NAME, DESCRIPTION
_____
T PRODUCT(3, 'salsa', '15 oz jar of salsa', 3.6, 20)
```

The update product() Procedure

The update product () procedure updates all the attributes of an object in the object products table; update product() is defined as follows in the body of product package:

```
PROCEDURE update product(
      p product IN t product
    ) AS
    BEGIN
      -- update the product in the object products table
      UPDATE object products op
      SET op = p product
      WHERE id = p product.id;
      COMMIT;
    EXCEPTION
      WHEN OTHERS THEN
        ROLLBACK;
    END update product;
```

The following example calls product package.update product() to update product #3 in the object products table:

```
CALL product_package.update_product(t_product(3, 'salsa',
     '25 oz jar of salsa', 2.70, 15));
```

The get_product_ref() Function

The get_product_ref() function returns a reference to a single object from the object_products table; get_product_ref() is defined as follows in the body of product_package:

```
FUNCTION get_product_ref(
    p_id IN object_products.id%TYPE
)

RETURN REF t_product IS
    -- declare a reference to a t_product
    v_product_ref REF t_product;

BEGIN
    -- get the REF for the product and
    -- store it in v_product_ref

SELECT REF(op)
    INTO v_product_ref
    FROM object_products op
    WHERE op.id = p_id;

    -- return v_product_ref
    RETURN v_product_ref;
END get product ref;
```

The following query calls product_package.get_product_ref() to get the reference to product #3 from the object products table:

```
SELECT product_package.get_product_ref(3)
FROM dual;
```

The next example calls product_package.get_product_ref() again, this time using DEREF() to get to the actual product:

SELECT DEREF(product_package.get_product_ref(3)) FROM dual;

The delete_product() Procedure

The delete_product() procedure deletes a single object from the object_products table; delete product() is defined as follows in the body of product package:

```
PROCEDURE delete_product(
    p_id IN object_products.id%TYPE
) AS
BEGIN
```

```
-- delete the product
 DELETE FROM object products op
 WHERE op.id = p id;
 COMMIT:
EXCEPTION
 WHEN OTHERS THEN
   ROLLBACK;
END delete product;
```

The following example calls product package.delete product() to delete product #3 from the object products table:

CALL product package.delete product(3);

Now that you've seen all the methods in product package, it's time for you to see two procedures named product lifecycle() and product lifecycle2() that call the various methods in the package. Both procedures are created when you run the object schema.sql script.

The product lifecycle() Procedure

The product lifecycle() procedure is defined as follows:

```
CREATE PROCEDURE product lifecycle AS
      -- declare object
      v product t product;
    BEGIN
      -- insert a new product
      product package.insert product(4, 'beef',
       '25 lb pack of beef', 32, 10);
      -- display the product
      product package.display product(4);
      -- get the new product and store it in v product
      SELECT product package.get product(4)
      INTO v product
      FROM dual;
      -- change some attributes of v product
      v product.description := '20 lb pack of beef';
      v product.price := 36;
      v product.days valid := 8;
      -- update the product
      product package.update product(v product);
      -- display the product
      product package.display product(4);
      -- delete the product
      product package.delete product(4);
    END product lifecycle;
```

The following example calls product lifecycle():

```
CALL product_lifecycle();
v_product.id=4
v_product.name=beef
v_product.description=25 lb pack of beef
v_product.price=32
v_product.days_valid=10
Sell by date=27-JUN-07
v_product.id=4
v_product.name=beef
v_product.description=20 lb pack of beef
v_product.price=36
v_product.days_valid=8
Sell by date=25-JUN-07
```

The product_lifecycle2() Procedure

The product_lifecycle2() procedure uses an object reference to access a product; product_lifecycle2() is defined as follows:

```
CREATE PROCEDURE product lifecycle2 AS
      -- declare object
      v product t product;
      -- declare object reference
      v product ref REF t product;
    BEGIN
      -- insert a new product
      product package.insert product(4, 'beef',
       '25 lb pack of beef', 32, 10);
      -- display the product
      product package.display product(4);
      -- get the new product reference and store it in v product ref
      SELECT product package.get product ref(4)
      INTO v_product ref
      FROM dual;
      -- dereference v product ref using the following query
      SELECT DEREF(v product ref)
      INTO v product
      FROM dual;
      -- change some attributes of v product
      v product.description := '20 lb pack of beef';
      v product.price := 36;
      v product.days valid := 8;
      -- update the product
      product package.update product(v product);
```

```
-- display the product
 product package.display product(4);
 -- delete the product
 product package.delete product(4);
END product lifecycle2;
```

One point to note in this procedure is that, in order to dereference v product ref, you have to use the following query:

```
SELECT DEREF(v product ref)
    INTO v product
    FROM dual;
```

The reason you have to use this query is that you cannot use DEREF () directly in PL/SQL code. For example, the following statement won't compile in PL/SQL:

```
v product := DEREF(v product ref);
```

The following example calls product lifecycle2():

```
CALL product lifecycle2();
```

```
v product.id=4
v product.name=beef
v product.description=25 lb pack of beef
v product.price=32
v product.days valid=10
Sell by date=27-JUN-07
v product.id=4
v product.name=beef
v product.description=20 lb pack of beef
v product.price=36
v product.days valid=8
Sell by date=25-JUN-07
```

Type Inheritance

Oracle Database 9i introduced object type inheritance, which allows you to define hierarchies of object types. For example, you might want to define a business person object type and have that type inherit the existing attributes from t person. The business person type could extend t person with attributes to store the person's job title and the name of the company they work for. For t person to be inherited from, the t person definition must include the NOT FINAL clause:

```
CREATE TYPE t person AS OBJECT (
      id INTEGER,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      dob DATE,
     phone
              VARCHAR2 (12),
      address t address,
     MEMBER FUNCTION display details RETURN VARCHAR2
    ) NOT FINAL;
```

The NOT FINAL clause indicates that t_person can be inherited from when defining another type. (The default when defining types is FINAL, meaning that the object type cannot be inherited from.)

The following statement creates the body for t_person; notice that the display_details() function returns a VARCHAR2 containing the id and name of the person:

```
CREATE TYPE BODY t_person AS

MEMBER FUNCTION display_details RETURN VARCHAR2 IS

BEGIN

RETURN 'id=' || id || ', name=' || first_name || ' ' || last_name;

END;

END;
```

NOTE

I've provided an SQL*Plus script named object_schema2.sql, which creates all the items shown in this and the following sections. You can run the script if you are using Oracle Database 9i or above. After the script completes, you will be logged in as object_user2.

To have a new type inherit attributes and methods from an existing type, you use the UNDER keyword when defining your new type. Our business person type, which I'll name t_business person, uses the UNDER keyword to inherit the attributes from t_person:

```
CREATE TYPE t_business_person UNDER t_person (
    title VARCHAR2(20),
    company VARCHAR2(20)
);
/
```

In this example, t_person is known as the *supertype*, and t_business_person is known as the *subtype*. You can then use t_business_person when defining column objects or object tables. For example, the following statement creates an object table named object_business_customers:

```
CREATE TABLE object_business_customers OF t_business_person;
```

The following INSERT statement adds an object to object_business_customers; notice that the two additional title and company attributes are supplied at the end of the t_business_person constructor:

```
INSERT INTO object_business_customers VALUES (
    t_business_person(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
        t_address('2 State Street', 'Beantown', 'MA', '12345'),
        'Manager', 'XYZ Corp'
    )
);
```

The following query retrieves this object:

```
FROM object_business_customers
WHERE id = 1;
```

```
ID FIRST NAME LAST NAME DOB PHONE
ADDRESS(STREET, CITY, STATE, ZIP)
_____
TITLE
             COMPANY
_____
      1 John Brown 01-FEB-55 800-555-1211
T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345')
             XYZ Corp
Manager
```

The following query calls the display details () function for this object:

```
SELECT o.display details()
   FROM object business customers o
   WHERE id = 1:
   O.DISPLAY DETAILS()
   ______
   id=1, name=John Brown
```

When you call a method, the database searches for that method in the subtype first; if the method isn't found, the supertype is searched. If you have a hierarchy of types, the database will search for the method up the hierarchy; if the method cannot be found, the database will report an error.

Using a Subtype Object in Place of a Supertype Object

In this section you'll see how you can use a subtype object in place of a supertype object; doing this gives you great flexibility when storing and manipulating related types. In the examples, you'll see how you use at business person object (a subtype object) in place of at person object (a supertype object).

SQL Examples

The following statement creates a table named object customers of type t person:

```
CREATE TABLE object customers OF t person;
```

The following INSERT statement adds a t person object to this table (the name is Jason Bond):

```
INSERT INTO object customers VALUES (
      t person(1, 'Jason', 'Bond', '03-APR-1965', '800-555-1212',
       t address('21 New Street', 'Anytown', 'CA', '12345')
   );
```

There's nothing unusual about the previous statement: The INSERT simply adds at person object to the object customers table. Now, because the object customers table stores objects of type t person, and t person is a supertype of t business person, you can

store a t_business_person object in object_customers; the following INSERT shows this, adding a customer named Steve Edwards:

```
INSERT INTO object_customers VALUES (
    t_business_person(2, 'Steve', 'Edwards', '03-MAR-1955', '800-555-1212',
        t_address('1 Market Street', 'Anytown', 'VA', '12345'),
        'Manager', 'XYZ Corp'
    )
);
```

The object_customers table now contains two objects: the t_person object added earlier (Jason Bond) and the new t_business_person object (Steve Edwards). The following query retrieves these two objects; notice that the title and company attributes for Steve Edwards are missing from the output:

SELECT *

FROM object customers o;

```
ID FIRST_NAME LAST_NAME DOB PHONE

ADDRESS(STREET, CITY, STATE, ZIP)

1 Jason Bond 03-APR-65 800-555-1212

T_ADDRESS('21 New Street', 'Anytown', 'CA', '12345')

2 Steve Edwards 03-MAR-55 800-555-1212

T_ADDRESS('1 Market Street', 'Anytown', 'VA', '12345')
```

You can get the full set of attributes for Steve Edwards by using VALUE() in the query, as shown in the following example; notice the different types of the objects for Jason Bond (a t_person object) and Steve Edwards (a t_business_person object) and that the title and company attributes for Steve Edwards now appear in the output:

SELECT VALUE(o) FROM object customers o;

```
VALUE(O)(ID, FIRST_NAME, LAST_NAME, DOB, PHONE,
ADDRESS(STREET, CITY, STATE, ZIP

T_PERSON(1, 'Jason', 'Bond', '03-APR-65', '800-555-1212',
T_ADDRESS('21 New Street', 'Anytown', 'CA', '12345'))

T_BUSINESS_PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212',
T_ADDRESS('1 Market Street', 'Anytown', 'VA', '12345'),
'Manager', 'XYZ Corp')
```

PL/SQL Examples

You can also manipulate subtype and supertype objects in PL/SQL. For example, the following procedure named $subtypes_and_supertypes()$ manipulates $t_business_person$ and t_person objects:

```
CREATE PROCEDURE subtypes_and_supertypes AS -- create objects
```

```
v business person t business person :=
    t business person(
     1, 'John', 'Brown',
     '01-FEB-1955', '800-555-1211',
      t address('2 State Street', 'Beantown', 'MA', '12345'),
      'Manager', 'XYZ Corp'
    );
  v person t person :=
    t person(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
      t address('2 State Street', 'Beantown', 'MA', '12345'));
  v business person2 t business person;
  v person2 t person;
BEGIN
  -- assign v business person to v person2
  v person2 := v business person;
  DBMS OUTPUT.PUT LINE('v person2.id = ' || v person2.id);
  DBMS OUTPUT.PUT LINE('v person2.first name = ' ||
    v person2.first name);
  DBMS OUTPUT.PUT LINE('v person2.last name = ' ||
   v person2.last name);
  -- the following lines will not compile because v person2
 -- is of type t person, and t person does not know about the
 -- additional title and company attributes
 -- DBMS OUTPUT.PUT LINE('v person2.title = ' ||
 -- v person2.title);
  -- DBMS OUTPUT.PUT LINE('v person2.company = ' ||
 -- v person2.company);
 -- the following line will not compile because you cannot
 -- directly assign a t person object to a t business person
 -- object
 -- v business person2 := v person;
END subtypes and supertypes;
```

The following example shows the result of calling subtypes and supertypes ():

```
SET SERVEROUTPUT ON
   CALL subtypes and supertypes();
    v person2.id = 1
    v person2.first name = John
    v person2.last name = Brown
```

NOT SUBSTITUTABLE Objects

If you want to prevent the use of a subtype object in place of a supertype object, you can mark an object table or object column as "not substitutable"; for example, the following statement creates a table named object customers2:

```
CREATE TABLE object customers not subs OF t person
   NOT SUBSTITUTABLE AT ALL LEVELS;
```

The NOT SUBSTITUTABLE AT ALL LEVELS clause indicates that no objects of a type other than t_person can be inserted into the table. If an attempt is made to add an object of type t business_person to this table, an error is returned:

```
SQL> INSERT INTO object_customers_not_subs VALUES (

2     t_business_person(1, 'Steve', 'Edwards', '03-MAR-1955', '800-555-1212',

3     t_address('1 Market Street', 'Anytown', 'VA', '12345'),

4     'Manager', 'XYZ Corp'

5  )

6 );

t_business_person(1, 'Steve', 'Edwards', '03-MAR-1955', '800-555-1212',

*

ERROR at line 2:

ORA-00932: inconsistent datatypes: expected OBJECT_USER2.T_PERSON got
OBJECT_USER2.T_BUSINESS_PERSON
```

You can also mark an object column as not substitutable; for example, the following statement creates a table with an object column named product that can store only objects of type t product:

```
CREATE TABLE products (

product t_product,

quantity_in_stock INTEGER
)

COLUMN product NOT SUBSTITUTABLE AT ALL LEVELS;
```

Any attempts to add an object not of type t_product to the product column will result in an error.

Other Useful Object Functions

In the earlier sections of this chapter you saw the use of the REF(), DEREF(), and VALUE() functions. In this section, you'll see the following additional functions that may be used with objects:

- IS OF() checks if an object is of a particular type or subtype.
- TREAT() does a run-time check to see if an object's type may be treated as a supertype.
- SYS_TYPEID() returns the ID of an object's type.

IS OF()

You use IS OF() to check whether an object is of a particular type or subtype. For example, the following query uses IS OF() to check whether the objects in the object_business_customers table are of type t_business_person—because they are, a row is returned by the query:

```
SELECT VALUE(o)

FROM object_business_customers o

WHERE VALUE(o) IS OF (t_business_person);

VALUE(O)(ID, FIRST_NAME, LAST_NAME, DOB, PHONE,

ADDRESS(STREET, CITY, STATE, ZIP
```

```
T BUSINESS PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
'Manager', 'XYZ Corp')
```

You can also use IS OF() to check whether an object is of a subtype of the specified type. For example, the objects in the object business customers table are of type t business person, which is a subtype of t person; therefore, the following query returns the same result as that shown in the previous example:

SELECT VALUE (o) FROM object business customers o WHERE VALUE(o) IS OF (t person); VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE, ADDRESS (STREET, CITY, STATE, ZIP T BUSINESS PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211', T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'), 'Manager', 'XYZ Corp')

You can include more than one type in IS OF(); for example:

```
SELECT VALUE (o)
   FROM object business customers o
   WHERE VALUE(o) IS OF (t business person, t person);
  VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE,
   ADDRESS (STREET, CITY, STATE, ZIP
   _____
   T BUSINESS PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
   T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
   'Manager', 'XYZ Corp')
```

In the earlier section entitled "Using a Subtype Object in Place of a Supertype Object," you saw the addition of a t person object (Jason Bond) and t business person object (Steve Edwards) to the object customers table. As a reminder, the following query shows these objects:

```
SELECT VALUE (o)
    FROM object customers o;
```

```
VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE,
ADDRESS (STREET, CITY, STATE, ZIP
______
T PERSON(1, 'Jason', 'Bond', '03-APR-65', '800-555-1212',
T ADDRESS('21 New Street', 'Anytown', 'CA', '12345'))
T BUSINESS PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212',
 T ADDRESS ('1 Market Street', 'Anytown', 'VA', '12345'),
 'Manager', 'XYZ Corp')
```

Because t business person type is a subtype of t person, IS OF (t person) returns true when a t business person object or a t person object is checked; this is illustrated in the following query that retrieves both Jason Bond and Steve Edwards using IS OF (t person):

```
SELECT VALUE (o)
    FROM object customers o
    WHERE VALUE(o) IS OF (t person);
    VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE,
     ADDRESS (STREET, CITY, STATE, ZIP
    ______
    T PERSON(1, 'Jason', 'Bond', '03-APR-65', '800-555-1212',
     T ADDRESS('21 New Street', 'Anytown', 'CA', '12345'))
    T BUSINESS PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212',
     T ADDRESS('1 Market Street', 'Anytown', 'VA', '12345'),
     'Manager', 'XYZ Corp')
```

You can also use the ONLY keyword in conjunction with IS OF () to check for objects of a specific type only: IS OF () returns false for objects of another type in the hierarchy. For example, IS OF (ONLY t person) returns true for objects of type t person only and returns false for objects of type t business person. In this way, you can use IS OF (ONLY t person) to restrict the object returned by a query against the object customers table to Jason Bond, as shown in the following example:

```
SELECT VALUE (o)
    FROM object customers o
    WHERE VALUE(o) IS OF (ONLY t person);
    VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE,
    ADDRESS (STREET, CITY, STATE, ZIP
    _____
    T_PERSON(1, 'Jason', 'Bond', '03-APR-65', '800-555-1212',
     T ADDRESS('21 New Street', 'Anytown', 'CA', '12345'))
```

Similarly, IS OF (ONLY t business person) returns true for objects of type t business person only, and returns false for objects of type t person. For example, the following query retrieves the t business person object only and therefore Steve Edwards is returned:

```
SELECT VALUE (o)
     FROM object customers o
     WHERE VALUE(o) IS OF (ONLY t business person);
     VALUE (O) (ID, FIRST NAME, LAST NAME, DOB, PHONE,
     ADDRESS (STREET, CITY, STATE, ZIP
     T BUSINESS PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212',
      T ADDRESS('1 Market Street', 'Anytown', 'VA', '12345'),
      'Manager', 'XYZ Corp')
```

You can include multiple types after ONLY. For example, IS OF (ONLY t person, t business person) returns true for t person and t business person objects only; the following query shows this by returning, as expected, both Jason Bond and Steve Edwards:

SELECT VALUE (o) FROM object customers o WHERE VALUE(o) IS OF (ONLY t person, t business person); VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE, ADDRESS (STREET, CITY, STATE, ZIP T PERSON(1, 'Jason', 'Bond', '03-APR-65', '800-555-1212', T ADDRESS('21 New Street', 'Anytown', 'CA', '12345')) T BUSINESS PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212', T ADDRESS('1 Market Street', 'Anytown', 'VA', '12345'),

'Manager', 'XYZ Corp')

You can also use IS OF () in PL/SQL. For example, the following procedure named check types() creates t business person and t person objects, and it uses IS OF() to check their types:

```
CREATE PROCEDURE check types AS
      -- create objects
      v business person t business person :=
        t business person(
          1, 'John', 'Brown',
          '01-FEB-1955', '800-555-1211',
          t_address('2 State Street', 'Beantown', 'MA', '12345'),
          'Manager', 'XYZ Corp'
      v person t person :=
        t person(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
          t address('2 State Street', 'Beantown', 'MA', '12345'));
    BEGIN
      -- check the types of the objects
      IF v business person IS OF (t business person) THEN
        DBMS OUTPUT.PUT LINE('v business person is of type ' ||
          't business person');
      END IF;
      IF v person IS OF (t person) THEN
        DBMS OUTPUT.PUT LINE('v person is of type t person');
      IF v business person IS OF (t person) THEN
        DBMS OUTPUT.PUT LINE('v business person is of type t person');
      IF v business person IS OF (t business person, t person) THEN
        DBMS OUTPUT.PUT LINE('v business person is of ' ||
          'type t business person or t person');
      END IF;
      IF v business person IS OF (ONLY t business person) THEN
        DBMS OUTPUT.PUT LINE('v business person is of only ' ||
          'type t business person');
      END IF;
      IF v business person IS OF (ONLY t person) THEN
        DBMS OUTPUT.PUT LINE('v business person is of only ' ||
```

```
'type t_person');
ELSE
   DBMS_OUTPUT_LINE('v_business_person is not of only ' ||
    'type t_person');
END IF;
END check_types;
/
```

The following example shows the result of calling check types ():

```
SET SERVEROUTPUT ON CALL check_types();
```

```
v_business_person is of type t_business_person
v_person is of type t_person
v_business_person is of type t_person
v_business_person is of type t_business_person or t_person
v_business_person is of only type t_business_person
v_business_person is not of only type t_person
```

TREAT()

You use TREAT() to do a run-time check to see whether an object of a subtype may be treated as an object of a supertype; if this is so, TREAT() returns an object, and if not so, TREAT() returns null. For example, because t_business_person is a subtype of t_person, a t_business_person object can be treated as a t_person object; you saw this earlier in the section entitled "Using a Subtype Object in Place of a Supertype Object," where a t_business_person object (Steve Edwards) was inserted into the object_customers table, which normally holds t_person objects. The following query uses TREAT() to check that Steve Edwards can be treated as a t_person object:

```
SELECT NVL2(TREAT(VALUE(o) AS t_person), 'yes', 'no')
FROM object_customers o
WHERE first_name = 'Steve' AND last_name = 'Edwards';

NVL
---
yes
```

NVL2() returns yes because TREAT (VALUE(o) AS t_person) returns an object (that is, not a null value). This means that Steve Edwards can be treated as a t person object.

The next query checks whether Jason Bond (a t_person object) can be treated as a t_business_person object—he cannot, and, therefore, TREAT() returns null, and NVL2() returns no:

```
SELECT NVL2(TREAT(VALUE(o) AS t_business_person), 'yes', 'no')
FROM object_customers o
WHERE first_name = 'Jason' AND last_name = 'Bond';

NVL
---
no
```

Because TREAT() returns null for the whole object, all the individual attributes for the object are also null. For example, the following query attempts to access the first name attribute through Jason Bond—null is returned (as expected):

SELECT

```
NVL2(TREAT(VALUE(o) AS t business person).first name, 'not null', 'null')
FROM object customers o
WHERE first name = 'Jason' AND last name = 'Bond';
NVL2
____
null
```

The next query uses TREAT() to check whether Jason Bond can be treated as a t person object—he is a t person object and therefore yes is returned:

```
SELECT NVL2(TREAT(VALUE(o) AS t person).first name, 'yes', 'no')
    FROM object customers o
    WHERE first name = 'Jason' AND last_name = 'Bond';
    NVL
    ___
    yes
```

You can also retrieve an object through the use of TREAT(); for example, the following query retrieves Steve Edwards:

```
SELECT TREAT(VALUE(o) AS t business person)
    FROM object customers o
   WHERE first name = 'Steve' AND last name = 'Edwards';
   TREAT (VALUE (O) AST BUSINESS PERSON) (ID, FIRST NAME, LAST NAME, DOB, PHONE,
    ADDRESS
    T BUSINESS PERSON(2, 'Steve', 'Edwards', '03-MAR-55', '800-555-1212',
    T ADDRESS ('1 Market Street', 'Anytown', 'VA', '12345'),
    'Manager', 'XYZ Corp')
```

If you try this query with Jason Bond, null is returned, as expected; therefore, nothing appears in the output of the following query:

```
SELECT TREAT(VALUE(o) AS t business person)
   FROM object customers o
   WHERE first name = 'Jason' AND last name = 'Bond';
   TREAT (VALUE (O) AST BUSINESS PERSON) (ID, FIRST NAME, LAST NAME, DOB, PHONE,
   ______
```

Let's take look at using TREAT() with the object business customers table, which contains the t business person object John Brown:

```
SELECT VALUE (o)
   FROM object business customers o;
```

```
VALUE(O)(ID, FIRST NAME, LAST NAME, DOB, PHONE,
ADDRESS (STREET, CITY, STATE, ZIP
______
T BUSINESS PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
'Manager', 'XYZ Corp')
```

The following guery uses TREAT() to check whether John Brown can be treated as a t person object—he can, because t business person is a subtype of t person; therefore, yes is returned by the guery:

```
SELECT NVL2(TREAT(VALUE(o) AS t person), 'yes', 'no')
    FROM object business customers o
    WHERE first name = 'John' AND last name = 'Brown';
    NVL
    ___
    yes
```

The following example shows the object returned by TREAT () when querying the object business customers table; notice that you still get the title and company attributes for John Brown:

```
SELECT TREAT(VALUE(o) AS t person)
    FROM object business customers o;
```

```
TREAT (VALUE (O) AST PERSON) (ID, FIRST NAME, LAST NAME, DOB, PHONE,
ADDRESS (STREET,
T BUSINESS PERSON(1, 'John', 'Brown', '01-FEB-55', '800-555-1211',
 T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
 'Manager', 'XYZ Corp')
```

You can also use TREAT () in PL/SQL. For example, the following procedure named treat example () illustrates the use of TREAT () (you should study the comments in the code to understand how TREAT () works in PL/SQL):

```
CREATE PROCEDURE treat example AS
        -- create objects
        v business person t business person :=
          t business person(
            1, 'John', 'Brown',
            '01-FEB-1955', '800-555-1211',
            t address('2 State Street', 'Beantown', 'MA', '12345'),
            'Manager', 'XYZ Corp'
        v person t person :=
          t person(1, 'John', 'Brown', '01-FEB-1955', '800-555-1211',
            t address('2 State Street', 'Beantown', 'MA', '12345'));
        v business person2 t business person;
        v person2 t person;
```

```
BEGIN
 -- assign v business person to v person2
 v person2 := v business person;
  DBMS OUTPUT.PUT LINE('v_person2.id = ' || v_person2.id);
  DBMS OUTPUT.PUT LINE('v person2.first name = ' | |
   v person2.first name);
  DBMS OUTPUT.PUT LINE('v person2.last name = ' ||
   v person2.last name);
 -- the following lines will not compile because v person2
  -- is of type t person, and t person does not know about the
  -- additional title and company attributes
  -- DBMS OUTPUT.PUT LINE('v person2.title = ' ||
  -- v person2.title);
  -- DBMS OUTPUT.PUT LINE('v person2.company = ' ||
  -- v person2.company);
  -- use TREAT when assigning v business person to v person2
  DBMS OUTPUT.PUT LINE ('Using TREAT');
  v person2 := TREAT(v business person AS t_person);
  DBMS OUTPUT.PUT LINE('v person2.id = ' || v person2.id);
  DBMS OUTPUT.PUT LINE('v person2.first name = ' ||
   v person2.first name);
  DBMS_OUTPUT.PUT_LINE('v person2.last name = ' ||
   v person2.last name);
  -- the following lines will still not compile because v person2
  -- is of type t person, and t person does not know about the
  -- additional title and company attributes
  -- DBMS OUTPUT.PUT LINE('v person2.title = ' ||
  -- v person2.title);
  -- DBMS OUTPUT.PUT LINE('v person2.company = ' ||
  -- v person2.company);
  -- the following lines do compile because TREAT is used
  DBMS OUTPUT.PUT LINE('v person2.title = ' ||
   TREAT(v person2 AS t business person).title);
  DBMS OUTPUT.PUT LINE('v person2.company = ' ||
   TREAT (v person2 AS t business person).company);
 -- the following line will not compile because you cannot
  -- directly assign a t person object to a t business person
  -- object
 -- v business person2 := v person;
 -- the following line throws a runtime error because you cannot
 -- assign a supertype object (v person) to a subtype object
 -- (v business person2)
 -- v business person2 := TREAT(v person AS t business person);
END treat example;
```

The following example shows the result of calling treat example ():

```
SET SERVEROUTPUT ON

CALL treat_example();

v_person2.id = 1

v_person2.first_name = John

v_person2.last_name = Brown

Using TREAT

v_person2.id = 1

v_person2.first_name = John

v_person2.last_name = Brown

v_person2.title = Manager

v_person2.company = XYZ Corp
```

SYS_TYPEID()

You use SYS_TYPEID() to get the ID of an object's type. For example, the following query uses SYS_TYPEID() to get the ID of the object type in the object_business_customers table:

```
SELECT first_name, last_name, SYS_TYPEID(VALUE(o))
FROM object_business_customers o;
```

You can get details on the types defined by the user through the user_types view. The following query retrieves the details of the type with a typeid of '02' (the ID returned by SYS_TYPEID() earlier) and the type name of T BUSINESS PERSON:

From the output of this query you can see that the supertype of t_business_person is t_person. Also, t_business_person has eight attributes and one method.

NOT INSTANTIABLE Object Types

You can mark an object type as NOT INSTANTIABLE, which prevents objects of that type from being created. You might want to mark an object type as NOT INSTANTIABLE when you use the type as an abstract supertype only and never create any objects of that type. For example, you could create a t_vehicle abstract type and use it as a supertype for a t_car subtype and a t_motorcycle subtype; you would then create actual t_car and t_motorcycle objects, but never t_vehicle objects.

The following statement creates a type named t vehicle, which is marked as NOT INSTANTIABLE:

```
CREATE TYPE t vehicle AS OBJECT (
         INTEGER,
      make VARCHAR2(15),
      model VARCHAR2(15)
    ) NOT FINAL NOT INSTANTIABLE;
```

NOTE

The t vehicle type is also marked as NOT FINAL, because a NOT INSTANTIABLE type cannot be FINAL. If it were FINAL, you wouldn't be able to use it as a supertype, which is the whole point of creating it in the first place.

The next example creates a subtype named t car under the t vehicle supertype; notice that t car has an additional attribute named convertible, which will be used to record whether the car has a convertible roof (Y for yes, N for no):

```
CREATE TYPE t car UNDER t vehicle (
     convertible CHAR(1)
    );
```

The following example creates a subtype named t motorcycle under the t vehicle supertype; notice that t motorcycle has an additional attribute named sidecar, which will be used to record whether the motorcycle has a sidecar (Y for yes, N for no):

```
CREATE TYPE t motorcycle UNDER t vehicle (
      sidecar CHAR(1)
    );
```

The next example creates tables named vehicles, cars, and motorcycles, which are object tables of the types t vehicle, t car, and t motorcycle, respectively:

```
CREATE TABLE vehicles OF t vehicle;
   CREATE TABLE cars OF t car;
   CREATE TABLE motorcycles OF t motorcycle;
```

Because t vehicle is NOT INSTANTIABLE, you cannot add an object to the vehicles table. If you attempt to do so, the database returns an error:

```
SOL> INSERT INTO vehicles VALUES (
          t vehicle(1, 'Toyota', 'MR2', '01-FEB-1955')
      3 );
      t vehicle(1, 'Toyota', 'MR2', '01-FEB-1955')
    ERROR at line 2:
    ORA-22826: cannot construct an instance of a non instantiable type
```

The following examples add objects to the cars and motorcycles tables:

```
INSERT INTO cars VALUES (
    t_car(1, 'Toyota', 'MR2', 'Y')
);

INSERT INTO motorcycles VALUES (
    t_motorcycle(1, 'Harley-Davidson', 'V-Rod', 'N')
);
```

The following queries retrieve the objects from the cars and motorcycles tables:

SELECT *

FROM cars;

SELECT *

FROM motorcycles;

ID	MAKE	MODEL	S
			_
1	Harlev-Davidson	V-Rod	Ν

User-Defined Constructors

As in other object-oriented languages like Java and C#, you can define your own constructors in PL/SQL to initialize a new object. You can define your own constructor to do such things as programmatically setting the attributes of a new object to default values.

The following example creates a type named t_person2 that declares two constructor methods with differing numbers of parameters:

```
CREATE TYPE t person2 AS OBJECT (
      id INTEGER,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      dob DATE,
      phone VARCHAR2(12),
      CONSTRUCTOR FUNCTION t person2 (
        p id INTEGER,
       p first name VARCHAR2,
       p last name VARCHAR2
      ) RETURN SELF AS RESULT,
      CONSTRUCTOR FUNCTION t person2(
        p id INTEGER,
        p first name VARCHAR2,
        p_last_name VARCHAR2,
        p_dob DATE
      ) RETURN SELF AS RESULT
     );
```

Notice the following about the constructor declarations:

- The CONSTRUCTOR FUNCTION keywords are used to identify the constructors.
- The RETURN SELF AS RESULT keywords indicate the current object being processed is returned by each constructor; SELF represents the current object being processed. What this means is that the constructor returns the new object it creates.
- The first constructor accepts three parameters (p id, p first name, and p last name), and the second constructor accepts four parameters (p id, p first name, p last name, and p dob).

The constructor declarations don't contain the actual code definitions for the constructors; the definitions are contained in the type body, which is created by the following statement:

```
CREATE TYPE BODY t person2 AS
       CONSTRUCTOR FUNCTION t person2 (
         p id INTEGER,
         p first name VARCHAR2,
         p last name VARCHAR2
       ) RETURN SELF AS RESULT IS
       BEGIN
         SELF.id := p id;
         SELF.first name := p first name;
         SELF.last name := p last name;
         SELF.dob := SYSDATE;
         SELF.phone := '555-1212';
         RETURN:
       CONSTRUCTOR FUNCTION t person2 (
         p id INTEGER,
         p first name VARCHAR2,
         p last name VARCHAR2,
         p dob DATE
       ) RETURN SELF AS RESULT IS
       BEGIN
         SELF.id := p id;
         SELF.first name := p first name;
         SELF.last name := p last name;
         SELF.dob := p dob;
         SELF.phone := '555-1213';
         RETURN;
       END;
     END;
```

Notice the following:

The constructors use SELF to reference the new object being created. For example, SELF.id := p id sets the id attribute of the new object to the value of the p id parameter passed into the constructor.

- The first constructor sets the id, first_name, and last_name attributes to the p_id, p_first_name, and p_last_name parameter values passed into the constructor; the dob attribute is set to the current datetime returned by SYSDATE(), and the phone attribute is set to 555-1212.
- The second constructor sets the id, first_name, last_name, and dob attributes to the p_id, p_first_name, p_last_name, and p_dob parameter values passed into the constructor; the remaining phone attribute is set to 555–1213.

Although not shown, the database automatically provides a default constructor that accepts five parameters and sets each attribute to the appropriate parameter value passed into the constructor. You'll see an example of this shortly.



NOTE

The constructors show an example of method overloading, whereby methods of the same name but different parameters are defined in the same type. A method may be overloaded by providing different numbers of parameters, types of parameters, or ordering of parameters.

The following example describes t person2; notice the constructor definitions in the output:

DESCRIBE	t_person2
----------	-----------

Name	Null?	Type
ID FIRST_NAME LAST_NAME DOB PHONE		NUMBER (38) VARCHAR2 (10) VARCHAR2 (10) DATE VARCHAR2 (12)
METHOD		
FINAL CONSTRUCTOR FUNCTION T_PE		In/Out Default?
P_ID	NUMBER VARCHAR2 VARCHAR2	IN IN IN
METHOD		
FINAL CONSTRUCTOR FUNCTION T_PE Argument Name	Type	AS RESULT In/Out Default?
	NUMBER VARCHAR2 VARCHAR2 DATE	IN IN IN

The following statement creates a table of type t_person2:

CREATE TABLE object customers2 OF t person2;

The following INSERT statement adds an object to the table; notice that three parameters are passed to the t person2 constructor:

```
INSERT INTO object customers2 VALUES (
      t person2(1, 'Jeff', 'Jones')
    );
```

Because three parameters are passed to t person2, this INSERT statement exercises the first constructor. This constructor sets the id, first name, and last name attributes of the new object to 1, Jeff, and Jones; the remaining dob and phone attributes are set to the result returned by SYSDATE () and the literal 555-1212. The following query retrieves the new object:

SELECT * FROM object customers2 WHERE id = 1:

```
ID FIRST_NAME LAST_NAME DOB PHONE
______
           Jones 17-JUN-07 555-1212
    1 Jeff
```

The next INSERT statement adds another object to the table; notice that four parameters are passed to the t person2 constructor:

```
INSERT INTO object customers2 VALUES (
       t person2(2, 'Gregory', 'Smith', '03-APR-1965')
     );
```

Because four parameters are passed to t person2, this INSERT statement exercises the second constructor. This constructor sets the id, first name, last name, and dob attributes of the object to 2, Gregory, Smith, and 03-APR-1965, respectively; the remaining phone attribute is set to 555-1213. The following query retrieves the new object:

SELECT *

FROM object customers2 WHERE id = 2;

```
ID FIRST NAME LAST NAME DOB PHONE
2 Gregory Smith 03-APR-65 555-1213
```

The next INSERT statement adds another object to the table; notice that five parameters are passed to the t person2 constructor:

```
INSERT INTO object customers2 VALUES (
       t person2(3, 'Jeremy', 'Hill', '05-JUN-1975', '555-1214')
    );
```

Because five parameters are passed to t person2, this INSERT statement exercises the default constructor. This constructor sets the id, first_name, last_name, dob, and phone attributes to 3, Jeremy, Hill, 05-JUN-1975, and 555-1214, respectively. The following query retrieves the new object:

```
SELECT *
  FROM object customers2
  WHERE id = 3:
       ID FIRST_NAME LAST NAME DOB PHONE
   3 Jeremy Hill 05-JUN-75 555-1214
```

Overriding Methods

When you create a subtype under a supertype, you can override a method in the supertype with a method in the subtype. This gives you a very flexible way of defining methods in a hierarchy of types.

The following statements create a supertype named t person3; notice that the display details () function returns a VARCHAR2 containing the attribute values of the object:

```
CREATE TYPE t person3 AS OBJECT (
             INTEGER,
       first name VARCHAR2(10),
       last name VARCHAR2(10),
       MEMBER FUNCTION display details RETURN VARCHAR2
      ) NOT FINAL;
     CREATE TYPE BODY t person3 AS
       MEMBER FUNCTION display details RETURN VARCHAR2 IS
         RETURN 'id=' || id ||
           ', name=' || first name || ' ' || last name;
       END;
     END;
```

The next set of statements creates a subtype named t business person3 under t person3; notice that the display details() function is overridden using the OVERRIDING keyword and that the function returns a VARCHAR2 containing the original and extended attribute values of the object:

```
CREATE TYPE t business person3 UNDER t person3 (
      title VARCHAR2(20),
      company VARCHAR2 (20),
      OVERRIDING MEMBER FUNCTION display details RETURN VARCHAR2
    );
    CREATE TYPE BODY t business person3 AS
      OVERRIDING MEMBER FUNCTION display details RETURN VARCHAR2 IS
      BEGIN
        RETURN 'id=' || id ||
          ', name=' || first name || ' ' || last name ||
          ', title=' || title || ', company=' || company;
```

```
END;
END:
```

The use of the OVERRIDING keyword indicates that display details () in t business person3 overrides display details() in t person3; therefore, when display details() in t business person3 is called, it calls display details() in t business person3, not display details() in t person3.



NOTE

In the next section of this chapter, you'll see how you can directly call a method in a supertype from a subtype. This saves you from having to recreate code in the subtype that is already in the supertype. You do this direct calling by using a new feature called generalized invocation in Oracle Database 11g.

The following statements create a table named object business customers3 and add an object to this table:

```
CREATE TABLE object business customers3 OF t business person3;
   INSERT INTO object business customers3 VALUES (
     t business person3(1, 'John', 'Brown', 'Manager', 'XYZ Corp')
   );
```

The following example calls display details() using object business customers3:

```
SELECT o.display details()
   FROM object business customers3 o
   WHERE id = 1;
   O.DISPLAY DETAILS()
   id=1, name=John Brown, title=Manager, company=XYZ Corp
```

Because the display details () function as defined in t business person3 is called, the VARCHAR2 returned by the function contains the id, first name, and last name attributes, along with the title and company attributes.

Generalized Invocation

As you saw in the previous section, you can override a method in the supertype with a method in the subtype. Generalized invocation is a new feature in Oracle Database 11g and allows you to call a method in a supertype from a subtype. As you'll see, generalized invocation saves you from having to recreate code in the subtype that is already in the supertype.



NOTE

I've provided an SQL*Plus script named object schema3.sql, which creates all the items shown in the rest of this chapter. You can run the object schema3.sql script only if you are using Oracle Database 11g. After the script completes, you will be logged in as object user3.

The following statements create a supertype named t_person; notice that the display_details() function returns a VARCHAR2 containing the attribute values:

The next set of statements creates a subtype named t_business_person under t_person; notice that the display details() function is overridden using the OVERRIDING keyword:

As you can see, display_details() in t_business_person overrides display_details() in t_person. The following line in display_details() uses generalized invocation to call a method in a supertype from a subtype:

```
RETURN (SELF AS t_person).display_details ||
', title=' || title || ', company=' || company;
```

What (SELF AS t_person).display_details does is to treat an object of the current type (which is t_business_person) as an object of type t_person and then to call display_details() in t_person. So, when display_details() in t_business_person is called, it first calls display_details() in t_person (which displays the id, first_name, and

last name attribute values), then displays the title and company attribute values. This meant I didn't have to re-create the code already in t person.display details() in t business person.display details(), thereby saving some work. If you have more complex methods in your types, this feature can save a lot of work and make your code easier to maintain.

The following statements create a table named object business customers and add an object to this table:

```
CREATE TABLE object business customers OF t business person;
   INSERT INTO object business customers VALUES (
     t business person(1, 'John', 'Brown', 'Manager', 'XYZ Corp')
   );
```

The following query calls display details() using object business customers:

```
SELECT o.display details()
   FROM object business customers o;
```

```
O.DISPLAY DETAILS()
______
id=1, name=John Brown, dob=01-FEB-55, title=Manager, company=XYZ Corp
```

As you can see, the id, name, and date of birth (dob) are displayed (which come from display details () in t person), followed by the title and company (which come from display details() in t business person).

Summary

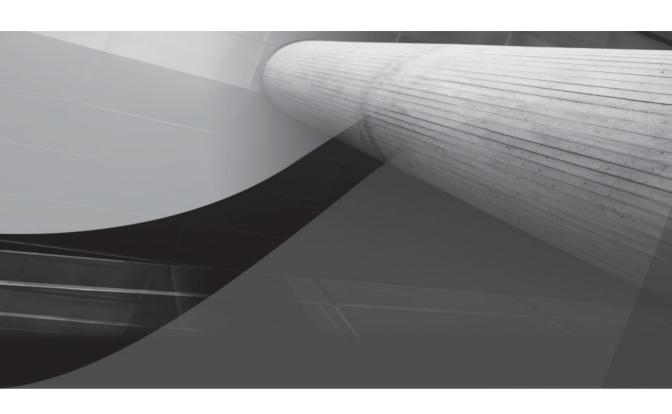
In this chapter, you learned the following:

- The Oracle database allows you to create object types. An object type is like a class in Java, C++, and C#. An object type may contain attributes and methods; you create an object type using the CREATE TYPE statement.
- You can use an object type to define a column object or an object table.
- You can use an object reference to access an individual row in an object table. An object reference is similar to a pointer in C++.
- You can create and manipulate objects in SQL and PL/SQL.
- With the release of Oracle Database 9*i*, you can use object type inheritance. This allows you to define hierarchies of database types.
- You can use a subtype object in place of a supertype object, which gives you great flexibility when storing and manipulating related types. If you want to prevent the use of a subtype object in place of supertype object, you can mark an object table or object column as NOT SUBSTITUTABLE.
- You can use a number of useful functions with objects, such as REF(), DEREF(), VALUE(), IS OF(), SYS TYPEID(), and TREAT().

426 Oracle Database 11*g* SQL

- You can mark an object type as NOT INSTANTIABLE, which prevents objects of that type from being created. You'll want to mark an object type as NOT INSTANTIABLE when you use that type as an abstract supertype and never actually create objects of that type.
- You can define your own constructors to do things like programmatically setting a default for attributes of an object.
- You can override a method in a supertype with a method in a subtype, giving you a very flexible way of defining methods in a hierarchy of types.
- You can use the new Oracle Database 11g generalized invocation feature to call methods in supertype from a subtype. Doing this can save you a lot of work and make your code easier to maintain.

In the next chapter, you'll learn about collections.



CHAPTER 13

Collections



n this chapter, you will do the following:

- Learn about collections
- Learn how to create collection types
- Use collection types to define columns in tables
- Create and manipulate collection data in SQL and PL/SQL
- Learn how a collection may itself contain embedded collections (a "multilevel" collection)
- Examine the enhancements to collections that were introduced in Oracle Database 10g

Introducing Collections

Oracle Database 8 introduced two new database types, known as *collections*, that allow you to store sets of elements. Oracle Database 9*i* extended these features to include multilevel collections, which allow you to embed a collection within another collection. Oracle Database 10*g* further extended collections to include associative arrays and much more.

There are three types of collections:

- Varrays A varray is similar to an array in Java, C++, and C#. A varray stores an ordered set of elements, and each element has an index that records its position in the array. Elements in a varray can be modified only as a whole, not individually; this means that even if you only want to modify one element, you must supply all the elements for the varray. A varray has a maximum size that you set when creating it, but you can change the size later.
- Nested tables A nested table is a table that is embedded within another table. You can insert, update, and delete individual elements in a nested table; this makes them more flexible than a varray, whose elements can be modified only as a whole. A nested table doesn't have a maximum size, and you can store an arbitrary number of elements in a nested table.
- Associative arrays (formerly known as index-by tables) An associative array is similar to a hash table in Java. Introduced in Oracle Database 10g, an associative array is a set of key and value pairs. You can get the value from the array using the key (which may be a string) or an integer that specifies the position of the value in the array. An associative array can be used only in PL/SQL and cannot be stored in the database.

You might be asking yourself why you would want to use collections in the first place. After all, using two tables with a foreign key already allows you to model relationships between data. The answer is that collections follow the object-oriented style of modern programming; in addition, the data stored in the collection may be accessed more rapidly by the database than if you were to use two relational tables to store the same data.

I've provided an SQL*Plus script named collection schema.sql in the SQL directory. The script creates a user named collection user with a password of collection password, and creates the collection types, tables, and PL/SQL code used in the first part of this chapter. You must run this script while logged in as a user with the required privileges to create a new user with the CONNECT and RESOURCE privileges; I log in as the system user on my database to run the scripts. After the script completes, you will be logged in as collection user.

Creating Collection Types

In this section, you'll see how to create a varray type and a nested table type.

Creating a Varray Type

A varray stores an ordered set of elements, all of the same type, and the type can be a built-in database type or a user-defined object type. Each element has an index that corresponds to its position in the array, and you can modify elements in the varray only as a whole.

You create a varray type using the CREATE TYPE statement, in which you specify the maximum size and the type of elements stored in the varray. The following example creates a type named t varray address that can store up to three VARCHAR2 strings:

```
CREATE TYPE t varray address AS VARRAY(3) OF VARCHAR2(50);
```

Each VARCHAR2 will be used to represent a different address for a customer of our example store. In Oracle Database 10g and higher, you can change the maximum number of elements of a varray using the ALTER TYPE statement. For example, the following statement alters the maximum number of elements to ten:

```
ALTER TYPE t varray address MODIFY LIMIT 10 CASCADE;
```

The CASCADE option propagates the change to any dependent objects in the database.

Creating a Nested Table Type

A nested table stores an unordered set of any number of elements. You can insert, update, and delete individual elements in a nested table. A nested table doesn't have a maximum size, and you can store an arbitrary number of elements in a nested table.

In this section, you'll see a nested table type that stores t address object types. You saw the use of t address in the previous chapter; it is used to represent an address and is defined as follows:

```
CREATE TYPE t address AS OBJECT (
      street VARCHAR2 (15),
      city VARCHAR2(15),
      state CHAR(2),
      zip VARCHAR2(5)
    );
```

You create a nested table type using the CREATE TYPE statement, and the following example creates a type named t nested table address that stores t address objects:

```
CREATE TYPE t nested table address AS TABLE OF t address;
```

Notice that you don't specify the maximum size of a nested table. That's because a nested table can store any number of elements.

Using a Collection Type to Define a Column in a Table

Once you've created a collection type, you can use it to define a column in a table. You'll see how to use the varray type and nested table type created in the previous section to define a column in a table.

Using a Varray Type to Define a Column in a Table

The following statement creates a table named customers_with_varray, which uses t varray address to define a column named addresses:

```
CREATE TABLE customers_with_varray (
id INTEGER PRIMARY KEY,
first_name VARCHAR2(10),
last_name VARCHAR2(10),
addresses t_varray_address
);
```

The elements in a varray are stored directly inside the table when the size of the varray is 4KB or less; otherwise, the varray is stored outside of the table. When a varray is stored with the table, accessing its elements is faster than accessing elements in a nested table.

Using a Nested Table Type to Define a Column in a Table

The following statement creates a table named customers_with_nested_table, which uses t nested table address to define a column named addresses:

The NESTED TABLE clause identifies the name of the nested table column (addresses in the example), and the STORE AS clause specifies the name of the nested table (nested_addresses in the example) where the actual elements are stored. You cannot access the nested table independently of the table in which it is embedded.

Getting Information on Collections

As you'll see in this section, you can use the DESCRIBE command and a couple of user views to get information on your collections.

Getting Information on a Varray

The following example describes t varray address:

DESCRIBE t varray address

```
t varray address VARRAY(3) OF VARCHAR2(50)
```

The next example describes the customers with varray table, whose addresses column is of the t varray address type:

DESCRIBE customers with varray

Name Null		L?	Type	
ID	NOT	NULL	NUMBER (38)	
FIRST_NAME			VARCHAR2(10)	
LAST_NAME			VARCHAR2(10)	
ADDRESSES			T_VARRAY_ADDRESS	

You can also get information on your varrays from the user varrays view. Table 13-1 describes the columns in user varrays.

Column	Туре	Description
	, ·	•
parent_table_name	VARCHAR2(30)	Name of the table that contains the varray.
parent_table_column	VARCHAR2 (4000)	Name of the column in the parent table containing the varray.
type_owner	VARCHAR2(30)	User who owns the varray type.
type_name	VARCHAR2(30)	Name of the varray type.
lob_name	VARCHAR2(30)	Name of the large object (LOB) when the varray is stored in an LOB. You'll learn about LOBs in the next chapter.
storage_spec	VARCHAR2(30)	Storage specification for the varray.
return_type	VARCHAR2(20)	Return type of the column.
element_ substitutable	VARCHAR2 (25)	Whether or not (Y/N) the varray element is substitutable for a subtype.

 TABLE 13-1
 Columns in the user_varrays View



NOTE

ADDRESSES

T VARRAY ADDRESS

You can get information on all the varrays you have access to using the all_varrays view.

The following example retrieves the details for t varray address from user varrays:

Getting Information on a Nested Table

You can also use DESCRIBE with a nested table, as shown in the following example that describes t nested table address:

DESCRIBE t nested table address

The next example describes the customers_with_nested_table table, whose addresses column is of type t nested table address:

DESCRIBE customers with nested table

Name	Null?	Type
ID	NOT NULL	NUMBER (38)
FIRST_NAME		VARCHAR2(10)
LAST_NAME		VARCHAR2(10)
ADDRESSES		T NESTED TABLE ADDRESS

If you set the depth to 2 and describe customers_with_nested_table, you can see the attributes that make up t nested table address:

SET DESCRIBE DEPTH 2

DESCRIBE customers_with_nested_table

Name	Null?	Type
ID	NOT NULL	NUMBER (38)
FIRST NAME	VARCHAR2	(10)

```
LAST NAME
                                      VARCHAR2 (10)
ADDRESSES
                                      T NESTED TABLE ADDRESS
  STREET
                                      VARCHAR2 (15)
  CITY
                                       VARCHAR2 (15)
  STATE
                                       CHAR(2)
  ZIP
                                       VARCHAR2 (5)
```

You can also get information on your nested tables from the user nested tables view. Table 13-2 describes the columns in user nested tables.



NOTE

You can get information on all the nested tables you have access to using the all nested tables view.

The following example retrieves the details for the nested addresses table from user nested tables:

SELECT table name, table type name, parent table name, parent table column FROM user nested tables WHERE table name = 'NESTED ADDRESSES';

```
TABLE NAME
                   TABLE TYPE NAME
PARENT TABLE NAME
_____
PARENT TABLE COLUMN
_____
NESTED ADDRESSES
                   T NESTED TABLE ADDRESS
CUSTOMERS WITH NESTED TABLE
ADDRESSES
```

Column	Туре	Description
table_name	VARCHAR2(30)	Name of the nested table
table_type_owner	VARCHAR2(30)	User who owns the nested table type
table_type_name	VARCHAR2(30)	Name of the nested table type
parent_table_name	VARCHAR2(30)	Name of the parent table that contains the nested table
parent_table_column	VARCHAR2 (4000)	Name of the column in the parent table containing the nested table
storage_spec	VARCHAR2(30)	Storage specification for the nested table
return_type	VARCHAR2(20)	Return type of the column
element_substitutable	VARCHAR2(25)	Whether or not (Y/N) the nested table element is substitutable for a subtype.

Populating a Collection with Elements

In this section, you'll see how to populate a varray and a nested table with elements using INSERT statements. You don't have to run the INSERT statements shown in this section: they are executed when you run the collection schema.sql script.

Populating a Varray with Elements

The following INSERT statements add rows to the customers_with_varray table; notice the use of the t varray address constructor to specify the strings for the elements of the varray:

```
INSERT INTO customers_with_varray VALUES (
   1, 'Steve', 'Brown',
   t_varray_address(
       '2 State Street, Beantown, MA, 12345',
       '4 Hill Street, Lost Town, CA, 54321'
   )
);

INSERT INTO customers_with_varray VALUES (
   2, 'John', 'Smith',
   t_varray_address(
       '1 High Street, Newtown, CA, 12347',
       '3 New Street, Anytown, MI, 54323',
       '7 Market Street, Main Town, MA, 54323'
)
);
```

As you can see, the first row has two addresses and the second has three. Any number of addresses up to the maximum limit for the varray can be stored.

Populating a Nested Table with Elements

The following INSERT statements add rows to customers_with_nested_table; notice the use of the t_nested_table_address and t_address constructors to specify the elements of the nested table:

```
INSERT INTO customers_with_nested_table VALUES (
    1, 'Steve', 'Brown',
    t_nested_table_address(
        t_address('2 State Street', 'Beantown', 'MA', '12345'),
        t_address('4 Hill Street', 'Lost Town', 'CA', '54321')
    )
);

INSERT INTO customers_with_nested_table VALUES (
    2, 'John', 'Smith',
    t_nested_table_address(
```

```
t address('1 High Street', 'Newtown', 'CA', '12347'),
    t address('3 New Street', 'Anytown', 'MI', '54323'),
    t address('7 Market Street', 'Main Town', 'MA', '54323')
);
```

As you can see, the first row has two addresses and the second has three. Any number of addresses can be stored in a nested table.

Retrieving Elements from Collections

In this section, you'll see how to retrieve elements from a varray and a nested table using queries. The output from the queries has been formatted slightly to make the results more readable.

Retrieving Elements from a Varray

The following query retrieves customer #1 from the customers with varray table; one row is returned, and it contains the two addresses stored in the varray:

```
SELECT *
    FROM customers_with_varray
   WHERE id = 1:
         ID FIRST NAME LAST NAME
    _____
    ADDRESSES
           1 Steve Brown
    T VARRAY ADDRESS ('2 State Street, Beantown, MA, 12345',
    '4 Hill Street, Lost Town, CA, 54321')
```

The next query specifies the actual column names:

```
SELECT id, first name, last name, addresses
    FROM customers with varray
    WHERE id = 1:
          ID FIRST NAME LAST NAME
    ADDRESSES
             1 Steve Brown
    T VARRAY ADDRESS('2 State Street, Beantown, MA, 12345',
     '4 Hill Street, Lost Town, CA, 54321')
```

These examples all return the addresses in the varray as a single row. Later, in the section "Using TABLE() to Treat a Collection as a Series of Rows," you'll see how you can treat the data stored in a collection as a series of rows.

Retrieving Elements from a Nested Table

The following query retrieves customer #1 from customers_with_nested_table; one row is returned, and it contains the two addresses stored in the nested table:

```
FROM customers_with_nested_table
WHERE id = 1;

ID FIRST_NAME LAST_NAME

ADDRESSES(STREET, CITY, STATE, ZIP)

1 Steve Brown

T_NESTED_TABLE_ADDRESS(
T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
T_ADDRESS('4 Hill Street', 'Lost Town', 'CA', '54321'))

The post grown appoints the actual solumn names:
```

The next query specifies the actual column names:

```
FROM customers_with_nested_table
WHERE id = 1;

ID FIRST_NAME LAST_NAME

ADDRESSES(STREET, CITY, STATE, ZIP)

1 Steve Brown

T_NESTED_TABLE_ADDRESS(
T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
T ADDRESS('4 Hill Street', 'Lost Town', 'CA', '54321'))
```

The next query gets just the addresses nested table; as in the previous examples, one row is returned, and it contains the two addresses stored in the nested table:

```
SELECT addresses

FROM customers_with_nested_table

WHERE id = 1;

ADDRESSES(STREET, CITY, STATE, ZIP)

_______

T_NESTED_TABLE_ADDRESS(

T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),

T ADDRESS('4 Hill Street', 'Lost Town', 'CA', '54321'))
```

Using TABLE() to Treat a Collection as a Series of Rows

The previous queries you've seen in this chapter return the contents of a collection as a single row. Sometimes, you may wish to treat the data stored in a collection as a series of rows; for

example, you might be working with a legacy application that can only use rows. To treat a collection as a series of rows, you use the TABLE () function. In this section, you'll see how to use TABLE () with a varray and a nested table.

Using TABLE() with a Varray

The following guery uses TABLE () to retrieve customer #1's two addresses from the customers with varray table; two separate rows are returned:

```
SELECT a.*
    FROM customers with varray c, TABLE(c.addresses) a
    WHERE id = 1;
    COLUMN VALUE
     2 State Street, Beantown, MA, 12345
     4 Hill Street, Lost Town, CA, 54321
```

Notice how the Oracle database software automatically adds the column name of COLUMN VALUE to the rows returned by the query. COLUMN VALUE is a pseudo column alias, and it is automatically added when a collection contains data of one of the built-in data types, like VARCHAR2, CHAR, NUMBER, or DATE. Because the example varray contains VARCHAR2 data, the COLUMN VALUE alias is added. If the varray had contained data of a user-defined object type, then TABLE() would return objects of that type and COLUMN VALUE would not appear; you'll see an example of this in the next section.

You can also embed an entire SELECT statement inside TABLE (). For example, the following query rewrites the previous example, placing a SELECT inside TABLE ():

```
SELECT *
     FROM TABLE (
       -- get the addresses for customer #1
       SELECT addresses
       FROM customers with varray
       WHERE id = 1
     );
     COLUMN VALUE
     2 State Street, Beantown, MA, 12345
     4 Hill Street, Lost Town, CA, 54321
```

The following query shows another example that uses TABLE() to get the addresses:

```
SELECT c.id, c.first name, c.last name, a.*
  FROM customers with varray c, TABLE(c.addresses) a
  WHERE id = 1;
        ID FIRST NAME LAST NAME
   -----
  COLUMN VALUE
         1 Steve Brown
```

```
2 State Street, Beantown, MA, 12345

1 Steve Brown
4 Hill Street, Lost Town, CA, 54321
```

Using TABLE() with a Nested Table

The following query uses TABLE() to retrieve customer #1's two addresses from customers_with nested table; notice that two separate rows are returned:

SELECT a.*

```
FROM customers_with_nested_table c, TABLE(c.addresses) a
WHERE id = 1;
```

S.	FREET	CITY	ST	ZIP
2	State Street	Beantown	MA	12345
4	Hill Street	Lost Town	CA	54321

The next query gets the street and state attributes of the addresses:

The following guery shows another example that uses TABLE() to get the addresses:

```
SELECT c.id, c.first_name, c.last_name, a.*

FROM customers_with_nested_table c, TABLE(c.addresses) a

WHERE c.id = 1:
```

ID	FIRST_NAME	LAST_NAME	STREET	CITY	ST	ZIP
1	Steve	Brown	2 State Street	Beantown	MA	12345
1	Steve	Brown	4 Hill Street	Lost Town	CA	54321

You'll see an important use of TABLE () later in the section "Modifying Elements of a Nested Table."

Modifying Elements of Collections

In this section, you'll see how to modify the elements in a varray and a nested table. You should feel free to run the UPDATE, INSERT, and DELETE statements shown in this section.

Modifying Elements of a Varray

The elements in a varray can be modified only as a whole, which means that even if you only want to modify one element, you must supply all the elements for the varray. The following UPDATE statement modifies the addresses of customer #2 in the customers with varray table:

```
UPDATE customers with varray
   SET addresses = t varray address(
      '6 Any Street, Lost Town, GA, 33347',
     '3 New Street, Anytown, MI, 54323',
      '7 Market Street, Main Town, MA, 54323'
   WHERE id = 2:
   1 row updated.
```

Modifying Elements of a Nested Table

Unlike in a varray, elements in a nested table can be modified individually. You can insert, update, and delete individual elements in a nested table; you'll see how to do all three of these modifications in this section.

The following INSERT statement adds an address to customer #2 in customer with nested table; notice that TABLE() is used to get the addresses as a series of rows:

```
INSERT INTO TABLE (
     -- get the addresses for customer #2
     SELECT addresses
     FROM customers with nested table
     WHERE id = 2
   ) VALUES (
      t address('5 Main Street', 'Uptown', 'NY', '55512')
   );
   1 row created.
```

The following UPDATE statement changes the '1 High Street' address of customer #2 to '9 Any Street'; notice the use of the alias addr in the VALUE clauses when specifying the addresses:

```
UPDATE TABLE (
     -- get the addresses for customer #2
     SELECT addresses
     FROM customers with nested table
     WHERE id = 2
   ) addr
   SET VALUE (addr) =
     t_address('9 Any Street', 'Lost Town', 'VA', '74321')
   WHERE VALUE(addr) =
     t address('1 High Street', 'Newtown', 'CA', '12347');
   1 row updated.
```

The following DELETE statement removes the '3 New Street...' address from customer #2:

```
DELETE FROM TABLE (
      -- get the addresses for customer #2
      SELECT addresses
      FROM customers_with_nested_table
```

```
WHERE id = 2
) addr
WHERE VALUE(addr) =
  t_address('3 New Street', 'Anytown', 'MI', '54323');
1 row_deleted.
```

Using a Map Method to Compare the Contents of Nested Tables

You can compare the contents of one nested table with the contents of another. Two nested tables are equal only if

- They are of the same type.
- They have the same number of rows.
- All their elements contain the same values.

If the elements of the nested table are of a built-in database type, like NUMBER, VARCHAR2, and so on, then the database will automatically compare the contents of the nested tables for you. If, however, the elements are of a user-defined object type, then you will need to provide a map function that contains code to compare the objects (map functions were shown in the section "Comparing Object Values" of the previous chapter).

The following statements create a type named t_address2 that contains a map function named get_string(); notice that get_string() returns a VARCHAR2 containing the values for the zip, state, city, and street attributes:

```
CREATE TYPE t address2 AS OBJECT (
      street VARCHAR2(15),
      city VARCHAR2(15),
      state CHAR(2),
      zip VARCHAR2(5),
      -- declare the get string() map function,
      -- which returns a VARCHAR2 string
      MAP MEMBER FUNCTION get string RETURN VARCHAR2
    );
    CREATE TYPE BODY t address2 AS
      -- define the get string() map function
      MAP MEMBER FUNCTION get string RETURN VARCHAR2 IS
        -- return a concatenated string containing the
        -- zip, state, city, and street attributes
        RETURN zip || ' ' || state || ' ' || city || ' ' || street;
      END get string;
    END;
```

As you'll see shortly, the database will automatically call get string () when comparing t address2 objects.

The following statements create a nested table type and a table, and add a row to the table:

```
CREATE TYPE t nested table address2 AS TABLE OF t address2;
    CREATE TABLE customers with nested table2 (
      id INTEGER PRIMARY KEY,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      addresses t nested table address2
    )
    NESTED TABLE
      addresses
    STORE AS
      nested addresses2;
     INSERT INTO customers with nested table2 VALUES (
      1, 'Steve', 'Brown',
      t nested table address2(
        t address2('2 State Street', 'Beantown', 'MA', '12345'),
        t_address2('4 Hill Street', 'Lost Town', 'CA', '54321')
     );
```

The following query includes a nested table in the WHERE clause; notice that the addresses after the = in the WHERE clause are the same as those in the previous INSERT statement:

```
SELECT cn.id, cn.first name, cn.last name
    FROM customers with nested table2 cn
    WHERE cn.addresses =
      t nested table address2(
        t address2('2 State Street', 'Beantown', 'MA', '12345'),
        t address2('4 Hill Street', 'Lost Town', 'CA', '54321')
           ID FIRST NAME LAST NAME
    -----
            1 Steve Brown
```

When the query is executed, the database automatically calls get string() to compare the t address2 objects in cn.addresses to the t address2 objects after the = in the WHERE clause. The get string () function returns a VARCHAR2 string containing the zip, state, city, and street attributes of the objects, and when the strings are equal for every object, the nested tables are also equal.

The next query returns no rows because the single address after the = in the WHERE clause matches only one of the addresses in cn.addresses (remember: two nested tables are equal only if they are of the same type, have the same number of rows, and their elements contain the same values):

```
SELECT cn.id, cn.first_name, cn.last_name
    FROM customers with nested table2 cn
```

```
WHERE cn.addresses =
  t_nested_table_address2(
    t_address2('4 Hill Street', 'Lost Town', 'CA', '54321')
);
no rows selected
```

In Oracle Database 10g and higher, you can use the SUBMULTISET operator to check whether the contents of one nested table are a subset of another nested table. The following query rewrites the previous example and returns a row:

```
SELECT cn.id, cn.first_name, cn.last_name
FROM customers_with_nested_table2 cn
WHERE

t_nested_table_address2(
t_address2('4 Hill Street', 'Lost Town', 'CA', '54321')
)
SUBMULTISET OF cn.addresses;

ID FIRST_NAME LAST_NAME

1 Steve Brown
```

Because the address in the first part of the WHERE clause is a subset of the addresses in cn.addresses, a match is found and a row is returned.

The following query shows another example; this time the addresses in cn.addresses are a subset of the addresses after OF in the WHERE clause:

```
SELECT cn.id, cn.first_name, cn.last_name
FROM customers_with_nested_table2 cn
WHERE
cn.addresses SUBMULTISET OF
t_nested_table_address2(
    t_address2('2 State Street', 'Beantown', 'MA', '12345'),
    t_address2('4 Hill Street', 'Lost Town', 'CA', '54321'),
    t_address2('6 State Street', 'Beantown', 'MA', '12345')
);

ID FIRST_NAME_LAST_NAME

1 Steve Brown
```

You'll learn more about the SUBMULTISET operator later in this chapter in the section "SUBMULTISET Operator." Also, in the section "Equal and Not-Equal Operators," you'll see how to use the ANSI operators implemented in Oracle Database 10g to compare nested tables.



NOTE

There is no direct mechanism for comparing the contents of varrays.

Using CAST() to Convert Collections from One Type to Another

You may use CAST () to convert a collection of one type to another collection type. In this section, you'll see how to use CAST() to convert a varray to a nested table and vice versa.

Using CAST() to Convert a Varray to a Nested Table

The following statements create and populate a table named customers with varray2 that contains an addresses column of type t varray address2:

```
CREATE TYPE t varray address2 AS VARRAY(3) OF t address;
      CREATE TABLE customers with varray2 (
       id INTEGER PRIMARY KEY,
       first name VARCHAR2(10),
       last name VARCHAR2(10),
       addresses t varray address2
     );
      INSERT INTO customers with varray2 VALUES (
       1, 'Jason', 'Bond',
       t varray address2(
         t address('9 Newton Drive', 'Sometown', 'WY', '22123'),
         t address('6 Spring Street', 'New City', 'CA', '77712')
       )
      );
```

The following query uses CAST () to return the varray addresses for customer #1 as a nested table; notice that the addresses appear in a constructor for the T NESTED TABLE ADDRESS type, indicating the conversion of the elements to this type:

```
SELECT CAST(cv.addresses AS t nested table address)
   FROM customers with varray2 cv
   WHERE cv.id = 1;
   CAST(CV.ADDRESSESAST NESTED TABLE ADDRESS)(STREET, CITY, STATE, ZIP)
   T NESTED TABLE ADDRESS (
    T_ADDRESS('9 Newton Drive', 'Sometown', 'WY', '22123'),
T_ADDRESS('6 Spring Street', 'New City', 'CA', '77712'))
```

Using CAST() to Convert a Nested Table to a Varray

The following query uses CAST() to return the addresses for customer #1 in customers with nested table as a varray; notice that the addresses appear in a constructor for T VARRAY ADDRESS2:

```
SELECT CAST(cn.addresses AS t_varray_address2)
   FROM customers with nested table cn
```

Using Collections in PL/SQL

You can use collections in PL/SQL. In this section, you'll see how to perform the following tasks in PL/SQL:

- Manipulate a varray
- Manipulate a nested table
- Use the PL/SQL collection methods to access and manipulate collections

All the packages you'll see in this section are created when you run the collection_schema.sql script. If you performed any of the INSERT, UPDATE, or DELETE statements shown in the earlier sections of this chapter, go ahead and rerun the collection_schema.sql script so that your output matches mine in this section.

Manipulating a Varray

In this section, you'll see a package named varray_package; this package contains the following items:

- A REF CURSOR type named t_ref_cursor
- A function named get_customers(), which returns a t_ref_cursor object that points to the rows in the customers with varray table
- A procedure named insert_customer(), which adds a row to the customers_with_varray table

The collection_schema.sql script contains the following package specification and body for varray_package:

```
CREATE PACKAGE BODY varray package AS
      -- get customers() function returns a REF CURSOR
      -- that points to the rows in customers with varray
      FUNCTION get customers
      RETURN t ref cursor IS
       --declare the REF CURSOR object
       v customers ref cursor t ref cursor;
      BEGIN
        -- get the REF CURSOR
        OPEN v customers ref cursor FOR
         SELECT *
         FROM customers with varray;
        -- return the REF CURSOR
       RETURN customers ref cursor;
      END get customers;
      -- insert customer() procedure adds a row to
      -- customers with varray
      PROCEDURE insert customer (
                 IN customers with varray.id%TYPE,
       p first name IN customers with varray.first name%TYPE,
       p last name IN customers with varray.last name%TYPE,
       p addresses IN customers with varray.addresses%TYPE
      ) IS
      BEGIN
        INSERT INTO customers_with_varray
       VALUES (p_id, p_first name, p last name, p addresses);
       COMMIT;
      EXCEPTION
       WHEN OTHERS THEN
         ROLLBACK;
      END insert customer;
    END varray package;
       The following example calls insert customer() to add a new row to the customers
    with varray table:
CALL varray package.insert customer(
      3, 'James', 'Red',
      t varray address(
        '10 Main Street, Green Town, CA, 22212',
        '20 State Street, Blue Town, FL, 22213'
    );
   Call completed.
       The next example calls get products () to retrieve the rows from customers with
```

varray:

SELECT varray_package.get_customers FROM dual;

```
GET CUSTOMERS
CURSOR STATEMENT: 1
CURSOR STATEMENT: 1
     ID FIRST NAME LAST NAME
_____
ADDRESSES
______
       1 Steve Brown
T VARRAY ADDRESS('2 State Street, Beantown, MA, 12345',
'4 Hill Street, Lost Town, CA, 54321')
       2 John Smith
T VARRAY ADDRESS('1 High Street, Newtown, CA, 12347',
 '3 New Street, Anytown, MI, 54323',
 '7 Market Street, Main Town, MA, 54323')
       3 James
                 Red
T VARRAY ADDRESS ('10 Main Street, Green Town, CA, 22212',
 '20 State Street, Blue Town, FL, 22213')
```

Manipulating a Nested Table

In this section, you'll see a package named nested_table_package; this package contains the following items:

- A REF CURSOR type named t ref cursor
- A function named get_customers(), which returns a t_ref_cursor object that points to the rows in customers_with_nested_table
- A procedure named insert_customer(), which adds a row to customers_with_ nested table

The collection_schema.sql script contains the following package specification and body for nested table package:

```
END nested table package;
   CREATE PACKAGE BODY nested table package AS
     -- get customers() function returns a REF CURSOR
     -- that points to the rows in customers with nested table
     FUNCTION get customers
     RETURN t ref cursor IS
       -- declare the REF CURSOR object
       v customers ref cursor t ref cursor;
     BEGIN
       -- get the REF CURSOR
       OPEN v customers ref cursor FOR
         SELECT *
         FROM customers with nested table;
       -- return the REF CURSOR
       RETURN customers ref cursor;
     END get customers;
     -- insert customer() procedure adds a row to
     -- customers with nested table
     PROCEDURE insert customer (
                    IN customers with nested table.id%TYPE,
       p first name IN customers with nested table.first name%TYPE,
       p last name IN customers with nested table.last name%TYPE,
       p_addresses IN customers with nested table.addresses%TYPE
     ) IS
     BEGIN
       INSERT INTO customers with nested table
       VALUES (p id, p first name, p last name, p addresses);
       COMMIT;
     EXCEPTION
       WHEN OTHERS THEN
        ROLLBACK;
     END insert customer;
   END nested table package;
      The following example calls insert customer() to add a new row to customers
   with nested table:
CALL nested_table_package.insert_customer(
     3, 'James', 'Red',
     t nested table address(
       t address('10 Main Street', 'Green Town', 'CA', '22212'),
       t address('20 State Street', 'Blue Town', 'FL', '22213')
     )
   );
   Call completed.
```

The next example calls get_customers() to retrieve the rows from customers_with_nested table:

SELECT nested_table_package.get_customers FROM dual;

```
GET CUSTOMERS
_____
CURSOR STATEMENT: 1
CURSOR STATEMENT: 1
      ID FIRST NAME LAST NAME
_____
ADDRESSES (STREET, CITY, STATE, ZIP)
______
               Brown
       1 Steve
T NESTED TABLE ADDRESS (
T ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),
T ADDRESS('4 Hill Street', 'Lost Town', 'CA', '54321'))
       2 John
                   Smith
T NESTED TABLE ADDRESS (
T ADDRESS ('1 High Street', 'Newtown', 'CA', '12347'),
T_ADDRESS('3 New Street', 'Anytown', 'MI', '54323'),
T ADDRESS('7 Market Street', 'Main Town', 'MA', '54323'))
        3 James
                   Red
T NESTED TABLE ADDRESS (
T ADDRESS('10 Main Street', 'Green Town', 'CA', '22212'),
T ADDRESS('20 State Street', 'Blue Town', 'FL', '22213'))
```

PL/SQL Collection Methods

In this section, you'll see the PL/SQL methods you can use with collections. Table 13-3 summarizes the collection methods. These methods can be used only in PL/SQL.

The following sections use a package named <code>collection_method_examples</code>; the examples illustrate the use of the methods shown in the previous table. The package is created by the <code>collection_schema.sql</code> script, and you'll see the individual methods defined in this package in the following sections.

COUNT()

COUNT returns the number of elements in a collection. Because a nested table can have individual elements that are empty, COUNT returns the number of non-empty elements in a nested table. For example, let's say you have a nested table named <code>v_nested_table</code> that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Empty
2	Not empty
3	Empty
4	Not empty

Method	Description
COUNT	Returns the number of elements in a collection. Because a nested table can have individual elements that are empty, COUNT returns the number of non-empty elements in a nested table.
DELETE	Removes elements from a collection. There are three forms of DELETE:
DELETE(n) DELETE(n, m)	■ DELETE removes all elements.
	■ DELETE (n) removes element n.
	■ DELETE (n, m) removes elements n through m .
	Because varrays always have consecutive subscripts, you cannot delete individual elements from a varray (except from the end by using TRIM).
EXISTS(n)	Returns true if element n in a collection exists: EXISTS returns true for non-empty elements and false for empty elements of nested tables or elements beyond the range of a collection.
EXTEND EXTEND(n)	Adds elements to the end of a collection. There are three forms of EXTEND:
EXTEND (n, m)	■ EXTEND adds one element, which is set to null.
	■ EXTEND (n) adds n elements, which are set to null.
	■ EXTEND (n, m) adds n elements, which are set to a copy of the m element.
FIRST	Returns the index of the first element in a collection. If the collection is completely empty, FIRST returns null. Because a nested table can have individual elements that are empty, FIRST returns the lowest index of a non-empty element in a nested table.
LAST	Returns the index of the last element in a collection. If the collection is completely empty, LAST returns null. Because a nested table can have individual elements that are empty, LAST returns the highest index of a non-empty element in a nested table.
LIMIT	For nested tables, which have no declared size, LIMIT returns null. For varrays, LIMIT returns the maximum number of elements that the varray can contain. You specify the limit in the type definition. The limit is changed when using TRIM and EXTEND, or when you use ALTER TYPE to change the limit.
NEXT (n)	Returns the index of the element after n . Because a nested table can have individual elements that are empty, NEXT returns the index of a non-empty element after n . If there are no elements after n , NEXT returns null.
PRIOR(n)	Returns the index of the element before n . Because a nested table can have individual elements that are empty, PRIOR returns the index of a non-empty element before n . If there are no elements before n , PRIOR returns null.
TRIM TRIM(n)	Removes elements from the end of a collection. There are two forms of TRIM:
	■ TRIM removes one element from the end.
	■ TRIM(n) removes n elements from the end.

Given this configuration, v_nested_table.COUNT returns 2, the number of non-empty elements.

COUNT is used in the get_addresses() and display_addresses() methods of the collection_method_examples package. The get_addresses() function returns the specified customer's addresses from customers_with_nested_table, whose id is passed to the function:

```
FUNCTION get addresses (
      p id customers with nested table.id%TYPE
    ) RETURN t nested table address IS
      -- declare object named v addresses to store the
      -- nested table of addresses
      v addresses t nested table address;
    BEGIN
      -- retrieve the nested table of addresses into v addresses
      SELECT addresses
      INTO v addresses
      FROM customers with nested table
      WHERE id = p id;
      -- display the number of addresses using v addresses.COUNT
      DBMS OUTPUT.PUT LINE (
        'Number of addresses = '|| v_addresses.COUNT
      );
      -- return v addresses
      RETURN v addresses;
    END get addresses;
```

The following example sets the server output on and calls get addresses () for customer #1:

SET SERVEROUTPUT ON SELECT collection_method_examples.get_addresses(1) addresses FROM dual;

```
ADDRESSES(STREET, CITY, STATE, ZIP)

T_NESTED_TABLE_ADDRESS(

T_ADDRESS('2 State Street', 'Beantown', 'MA', '12345'),

T_ADDRESS('4 Hill Street', 'Lost Town', 'CA', '54321'))

Number of addresses = 2
```

The following display_addresses () procedure accepts a parameter named $p_addresses$, which contains a nested table of addresses; the procedure displays the number of addresses in $p_addresses$ using COUNT, and then displays those addresses using a loop:

```
PROCEDURE display_addresses(
    p_addresses t_nested_table_address
) IS
    v count INTEGER;
```

```
BEGIN
 -- display the number of addresses in p addresses
 DBMS OUTPUT.PUT LINE (
   'Current number of addresses = '|| p addresses.COUNT
  -- display the addresses in p addresses using a loop
  FOR v count IN 1..p addresses.COUNT LOOP
    DBMS_OUTPUT.PUT_LINE('Address #' || v_count || ':');
    DBMS OUTPUT.PUT(p addresses(v count).street | | ', ');
    DBMS OUTPUT.PUT(p addresses(v count).city | ', ');
    DBMS OUTPUT.PUT(p addresses(v count).state || ', ');
    DBMS_OUTPUT.PUT_LINE(p_addresses(v_count).zip);
 END LOOP;
END display addresses;
```

You'll see the use of display addresses () shortly.

DELETE()

DELETE removes elements from a collection. There are three forms of DELETE:

- DELETE removes all elements.
- DELETE (n) removes element n.
- DELETE (n, m) removes elements n through m.

For example, let's say you have a nested table named v nested table that has seven elements, then v nested table.DELETE(2, 5) removes elements 2 through 5.

The following delete address() procedure gets the addresses for customer #1 and then uses DELETE to remove the address whose index is specified by the p address num parameter:

```
PROCEDURE delete address(
     p address num INTEGER
     v addresses t nested table address;
   BEGIN
     v addresses := get addresses(1);
     display addresses (v addresses);
     DBMS OUTPUT.PUT LINE('Deleting address #' || p address num);
     -- delete the address specified by p address num
     v_addresses.DELETE(p_address_num);
     display addresses (v addresses);
   END delete address;
```

The following example calls delete address (2) to remove address #2 from customer #1:

```
CALL collection method examples.delete address(2);
    Number of addresses = 2
    Current number of addresses = 2
```

```
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Deleting address #2
Current number of addresses = 1
Address #1:
2 State Street, Beantown, MA, 12345
```

EXISTS()

EXISTS (n) returns true if element n in a collection exists: EXISTS returns true for non-empty elements, and it returns false for empty elements of nested tables or elements beyond the range of a collection. For example, let's say you have a nested table named v_nested_table that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Empty
2	Not empty
3	Empty
4	Not empty

Given this configuration, v_nested_table.EXISTS(2) returns true (because element #2 is not empty), and v_nested_table.EXISTS(3) returns false (because element #3 is empty). The following exist_addresses() procedure gets the addresses for customer #1, uses DELETE to remove address #1, and then uses EXISTS to check whether addresses #1 and #2 exist (#1 does not exist because it has been deleted, #2 does exist):

```
PROCEDURE exist_addresses IS

v_addresses t_nested_table_address;

BEGIN

v_addresses := get_addresses(1);

DBMS_OUTPUT.PUT_LINE('Deleting address #1');

v_addresses.DELETE(1);

-- use EXISTS to check if the addresses exist

IF v_addresses.EXISTS(1) THEN

DBMS_OUTPUT.PUT_LINE('Address #1 does exist');

ELSE

DBMS_OUTPUT.PUT_LINE('Address #1 does not exist');

END IF;

IF v_addresses.EXISTS(2) THEN

DBMS_OUTPUT.PUT_LINE('Address #2 does exist');

END IF;

END exist_addresses;
```

CALL collection_method_examples.exist_addresses();
Number of addresses = 2

The following example calls exist addresses ():

```
Deleting address #1
Address #1 does not exist
Address #2 does exist
```

EXTEND()

EXTEND adds elements to the end of a collection. There are three forms of EXTEND:

- EXTEND adds one element, which is set to null.
- EXTEND (n) adds n elements, which are set to null.
- EXTEND (n, m) adds n elements, which are set to a copy of the m element.

For example, let's say you have a collection named v nested table that has seven elements, then v nested table.EXTEND(2, 5) adds element #5 twice to the end of the collection.

The following extend addresses () procedure gets the addresses for customer #1 into v addresses, then uses EXTEND to copy address #1 twice to the end of v addresses:

```
PROCEDURE extend addresses IS
      v addresses t nested table address;
    BEGIN
      v addresses := get addresses(1);
      display addresses (v addresses);
      DBMS OUTPUT.PUT LINE('Extending addresses');
      -- copy address #1 twice to the end of v addresses
      v addresses.EXTEND(2, 1);
      display addresses (v addresses);
    END extend addresses;
```

The following example calls extend addresses ():

CALL collection method examples.extend addresses();

```
Number of addresses = 2
Current number of addresses = 2
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Extending addresses
Current number of addresses = 4
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Address #3:
2 State Street, Beantown, MA, 12345
Address #4:
2 State Street, Beantown, MA, 12345
```

FIRST()

You use FIRST to get the index of the first element in a collection. If the collection is completely empty, FIRST returns null. Because a nested table can have individual elements that are empty, FIRST returns the lowest index of a non-empty element in a nested table. For example, let's say you have a nested table named v_nested_table that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Empty
2	Not empty
3	Empty
4	Not empty

Given this configuration, $v_nested_table.FIRST$ returns 2, the lowest index containing a non-empty element.

The following first_address() procedure gets the addresses for customer #1 into v_addresses and then uses FIRST to display the index of the first address in v_addresses; the procedure then deletes address #1 using DELETE and displays the new index returned by FIRST:

```
PROCEDURE first_address IS

v_addresses t_nested_table_address;

BEGIN

v_addresses := get_addresses(1);

-- display the FIRST address

DBMS_OUTPUT.PUT_LINE('First address = ' || v_addresses.FIRST);

DBMS_OUTPUT.PUT_LINE('Deleting address #1');

v_addresses.DELETE(1);

-- display the FIRST address again

DBMS_OUTPUT.PUT_LINE('First address = ' || v_addresses.FIRST);

END first_address;
```

The following example calls first address():

```
CALL collection_method_examples.first_address();
Number of addresses = 2
First address = 1
Deleting address #1
```

LAST()

First address = 2

LAST returns the index of the last element in a collection. If the collection is completely empty, LAST returns null. Because a nested table can have individual elements that are empty, LAST returns the highest index of a non-empty element in a nested table. For example, let's say you have a nested table named v nested table that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Not empty
2	Empty
3	Empty
4	Not empty

Given this configuration, v nested table. LAST returns 4, the highest index containing a non-empty element.

The following last address() procedure gets the addresses for customer #1 into v addresses and then uses LAST to display the index of the last address in v addresses; the procedure then deletes address #2 using DELETE and displays the new index returned by LAST:

```
PROCEDURE last address IS
      v addresses t nested table address;
    BEGIN
      v addresses := get addresses(1);
      -- display the LAST address
      DBMS OUTPUT.PUT LINE('Last address = ' || v_addresses.LAST);
      DBMS OUTPUT.PUT LINE('Deleting address #2');
      v addresses.DELETE(2);
      -- display the LAST address again
      DBMS OUTPUT.PUT LINE('Last address = ' || v addresses.LAST);
    END last address;
```

The following example calls last address():

CALL collection method examples.last address();

```
Number of addresses = 2
Last address = 2
Deleting address #2
Last address = 1
```

NEXT()

NEXT (n) returns the index of the element after n. Because a nested table can have individual elements that are empty, NEXT returns the index of a non-empty element after n. If there are no elements after n, NEXT returns null. For example, let's say you have a nested table named v nested table that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Not empty
2	Empty
3	Empty
4	Not empty

Given this configuration, v nested table.NEXT (1) returns 4, the index containing the next non-empty element; v nested table.NEXT(4) returns null.

The following next address() procedure gets the addresses for customer #1 into v addresses and then uses NEXT (1) to get the index of the address after address #1 in v addresses; the procedure then uses NEXT (2) to attempt to get the index of the address after address #2 (there isn't one, because customer #1 only has two addresses, so null is returned):

```
PROCEDURE next address IS
      v addresses t nested table address;
    BEGIN
      v addresses := get addresses(1);
      -- use NEXT(1) to get the index of the address
      -- after address #1
      DBMS OUTPUT.PUT LINE (
         'v addresses.NEXT(1) = ' || v addresses.NEXT(1)
      );
      -- use NEXT(2) to attempt to get the index of
      -- the address after address #2 (there isn't one,
      -- so null is returned)
      DBMS OUTPUT.PUT LINE (
        'v addresses.NEXT(2) = ' || v_addresses.NEXT(2)
    END next address;
```

The following example calls next address(); v addresses.NEXT(2) is null, and so no output is shown after the = for that element:

CALL collection_method_examples.next_address();

```
Number of addresses = 2
v = 2
v addresses.NEXT(2) =
```

PRIOR()

PRIOR (n) returns the index of the element before n. Because a nested table can have individual elements that are empty, PRIOR returns the index of a non-empty element before n. If there are no elements before n, PRIOR returns null. For example, let's say you have a nested table named v nested table that has its elements set as shown in the following table.

Element Index	Empty/Not Empty
1	Not empty
2	Empty
3	Empty
4	Not empty

Given this configuration, v nested table.PRIOR(4) returns 1, the index containing the prior non-empty element; v nested table.PRIOR(1) returns null.

The following prior address () procedure gets the addresses for customer #1 into v addresses and then uses PRIOR (2) to get the index of the address before address #2 in v addresses; the procedure then uses PRIOR (1) to attempt to get the index of the address before address #1 (there isn't one, so null is returned):

```
PROCEDURE prior address IS
      v addresses t nested table address;
    BEGIN
      v addresses := get addresses(1);
      -- use PRIOR(2) to get the index of the address
      -- before address #2
      DBMS OUTPUT.PUT LINE (
        'v addresses.PRIOR(2) = ' || v addresses.PRIOR(2)
      -- use PRIOR(1) to attempt to get the index of
      -- the address before address #1 (there isn't one,
      -- so null is returned)
      DBMS OUTPUT.PUT LINE (
        'v addresses.PRIOR(1) = ' || v_addresses.PRIOR(1)
    END prior address;
```

The following example calls prior address(); v addresses.PRIOR(1) is null, and so no output is shown after the = for that element:

CALL collection method examples.prior address();

```
Number of addresses = 2
v = 1 = 1
v addresses.PRIOR(1) =
```

TRIM()

TRIM removes elements from the end of a collection. There are two forms of TRIM:

- TRIM removes one element from the end.
- TRIM(n) removes n elements from the end.

For example, let's say you have a nested table named v nested table, then v nested table.TRIM(2) removes two elements from the end.

The following trim addresses () procedure gets the addresses of customer #1, copies address #1 to the end of v addresses three times using EXTEND(3, 1), and then removes two addresses from the end of v addresses using TRIM(2):

```
PROCEDURE trim addresses IS
     v addresses t nested table address;
   BEGIN
     v addresses := get addresses(1);
     display addresses (v addresses);
     DBMS OUTPUT.PUT LINE('Extending addresses');
```

```
v addresses.EXTEND(3, 1);
 display addresses (v addresses);
 DBMS OUTPUT.PUT LINE('Trimming 2 addresses from end');
 -- remove 2 addresses from the end of v addresses
 -- using TRIM(2)
 v addresses.TRIM(2);
 display addresses (v addresses);
END trim addresses;
```

The following example calls trim addresses():

CALL collection method examples.trim addresses();

```
Number of addresses = 2
Current number of addresses = 2
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Extending addresses
Current number of addresses = 5
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Address #3:
2 State Street, Beantown, MA, 12345
Address #4:
2 State Street, Beantown, MA, 12345
Address #5:
2 State Street, Beantown, MA, 12345
Trimming 2 addresses from end
Current number of addresses = 3
Address #1:
2 State Street, Beantown, MA, 12345
Address #2:
4 Hill Street, Lost Town, CA, 54321
Address #3:
2 State Street, Beantown, MA, 12345
```

Multilevel Collections

With the release of Oracle Database 9i, you can create a collection in the database whose elements are also a collection. These "collections of collections" are known as multilevel collections. The following list shows the valid multilevel collections:

- A nested table of nested tables
- A nested table of varrays
- A varray of varrays
- A varray of nested tables

I've provided an SQL*Plus script named collection schema2.sql in the SQL directory. This script creates a user named collection user2, with a password of collection password, along with the types and the table shown in this section. You can run this script if you are using Oracle Database 9i or higher. After the script completes, you will be logged in as collection user2.

Let's say you wanted to store a set of phone numbers for each address of a customer. The following example creates a varray type of three VARCHAR2 strings named t varray phone to represent phone numbers:

```
CREATE TYPE t varray phone AS VARRAY(3) OF VARCHAR2(14);
```

Next, the following example creates an object type named t address that contains an attribute named phone numbers; this attribute is of type t varray phone:

```
CREATE TYPE t address AS OBJECT (
      street VARCHAR2(15),
                 VARCHAR2(15),
      city
      state
                 CHAR (2),
                 VARCHAR2(5),
      phone numbers t varray phone
    );
```

The next example creates a nested table type of t address objects:

```
CREATE TYPE t nested table address AS TABLE OF t address;
```

The following example creates a table named customers with nested table, which contains a column named addresses of type t nested table address:

```
CREATE TABLE customers with nested table (
               INTEGER PRIMARY KEY,
       first name VARCHAR2(10),
       last name VARCHAR2(10),
       addresses t nested table address
     NESTED TABLE
       addresses
     STORE AS
       nested addresses;
```

So, customers with nested table contains a nested table whose elements contain an address with a varray of phone numbers.

The following INSERT statement adds a row to customers with nested table; notice the structure and content of the INSERT statement, which contains elements for the nested table of addresses, each of which has an embedded varray of phone numbers:

```
INSERT INTO customers with nested table VALUES (
      1, 'Steve', 'Brown',
      t nested table address(
        t address('2 State Street', 'Beantown', 'MA', '12345',
```

```
t_varray_phone(
    '(800)-555-1211',
    '(800)-555-1212',
    '(800)-555-1213'
)
),
t_address('4 Hill Street', 'Lost Town', 'CA', '54321',
    t_varray_phone(
    '(800)-555-1211',
    '(800)-555-1212'
)
)
));
```

You can see that the first address has three phone numbers, while the second address has two. The following query retrieves the row from customers_with_nested_table:

```
SELECT *
```

```
FROM customers_with_nested_table;
```

You can use TABLE () to treat the data stored in the collections as a series of rows, as shown in the following query:

```
SELECT cn.first_name, cn.last_name, a.street, a.city, a.state, p.*
FROM customers_with_nested_table cn,
TABLE(cn.addresses) a, TABLE(a.phone numbers) p;
```

FIRST_NAME	LAST_NAME	STREET		CITY	ST	COLUMN_VALUE
Steve	Brown	2 5	State Street	Beantown	MA	(800) -555-1211
Steve	Brown	2 5	State Street	Beantown	MA	(800) -555-1212
Steve	Brown	2 5	State Street	Beantown	MA	(800) -555-1213
Steve	Brown	4 I	Hill Street	Lost Town	CA	(800) -555-1211
Steve	Brown	4 I	Hill Street	Lost Town	CA	(800) -555-1212

The following UPDATE statement shows how to update the phone numbers for the 2 State Street address; notice that TABLE () is used to get the addresses as a series of rows and that a varray containing the new phone numbers is supplied in the SET clause:

```
UPDATE TABLE (
     -- get the addresses for customer #1
     SELECT cn.addresses
     FROM customers with nested table cn
     WHERE cn.id = 1
   ) addrs
   SET addrs.phone numbers =
     t varray phone (
        '(800)-555-1214',
        '(800)-555-1215'
   WHERE addrs.street = '2 State Street';
   1 row updated.
```

The following query verifies the change:

```
SELECT cn.first name, cn.last name, a.street, a.city, a.state, p.*
   FROM customers with nested table cn,
    TABLE(cn.addresses) a, TABLE(a.phone numbers) p;
```

FIRST_NAME	LAST_NAME	STREET		CITY	ST	COLUMN_VALUE
Steve	Brown	2 Stat	e Street	Beantown	MA	(800)-555-1214
Steve	Brown	2 Stat	e Street	Beantown	MA	(800)-555-1215
Steve	Brown	4 Hill	Street	Lost Town	CA	(800)-555-1211
Steve	Brown	4 Hill	Street	Lost Town	CA	(800) -555-1212

Support for multilevel collection types is a very powerful extension to the Oracle database software, and you might want to consider using them in any database designs you contribute to.

Oracle Database 10g Enhancements to Collections

In this section, you'll learn about the following enhancements made to collections in Oracle Database 10g:

- Support for associative arrays
- Ability to change the size or precision of an element type
- Ability to increase the number of elements in a varray
- Ability to use varray columns in temporary tables
- Ability to use a different tablespace for a nested table's storage table
- ANSI support for nested tables

The various statements that create the items shown in this section are contained in the collection schema3.sql script. This script creates a user named collection user3 with a password of collection_password and creates the collection types, tables, and PL/SQL code. You can run this script if you are using Oracle Database 10g or higher. After the script completes, you will be logged in as collection user3.

Associative Arrays

An associative array is a set of key and value pairs. You can get the value from the array using the key (which may be a string) or an integer that specifies the position of the value in the array. The following example procedure named customers_associative_array() illustrates the use of associative arrays:

```
CREATE PROCEDURE customers associative array AS
      -- define an associative array type named t assoc array;
      -- the value stored in each array element is a NUMBER,
      -- and the index key to access each element is a VARCHAR2
      TYPE t assoc array IS TABLE OF NUMBER INDEX BY VARCHAR2(15);
      -- declare an object named v customer array of type t assoc array;
      -- v customer array will be used to store the ages of customers
      v customer array t assoc array;
    BEGIN
      -- assign the values to v customer array; the VARCHAR2 key is the
      -- customer name and the NUMBER value is the age of the customer
      v customer array('Jason') := 32;
      v customer array('Steve') := 28;
      v customer array('Fred') := 43;
      v customer array('Cynthia') := 27;
      -- display the values stored in v customer array
      DBMS OUTPUT.PUT LINE (
        'v customer array[''Jason''] = ' || v customer array('Jason')
      DBMS OUTPUT.PUT LINE (
        'v customer array[''Steve''] = ' || v customer array('Steve')
      DBMS OUTPUT.PUT LINE (
        'v customer array[''Fred''] = ' || v customer array('Fred')
      );
      DBMS OUTPUT.PUT LINE (
        'v customer array[''Cynthia''] = ' || v customer array('Cynthia')
    END customers associative array;
```

The following example sets the server output on and calls customers_associative_array():

SET SERVEROUTPUT ON

```
CALL customers_associative_array();
v_customer_array['Jason'] = 32
v_customer_array['Steve'] = 28
v_customer_array['Fred'] = 43
v customer array['Cynthia'] = 27
```

Changing the Size of an Element Type

You can change the size of an element type in a collection when the element type is one of the character, numeric, or raw types (raw is used to store binary data—you'll learn about this in the next chapter). Earlier in this chapter, you saw the following statement that creates a varray type named t varray address:

```
CREATE TYPE t varray address AS VARRAY(2) OF VARCHAR2(50);
```

The following example changes the size of the VARCHAR2 elements in t varray address

```
ALTER TYPE t varray address
   MODIFY ELEMENT TYPE VARCHAR2 (60) CASCADE;
```

```
Type altered.
```

The CASCADE option propagates the change to any dependent objects in the database, which, in the example, is the customers with varray table that contains a column named addresses of type t varray address. You can also use the INVALIDATE option to invalidate any dependent objects and immediately recompile the PL/SQL code for the type.

Increasing the Number of Elements in a Varray

You can increase the number of elements in a varray. The following example increases the number of elements in t varray address to 5:

```
ALTER TYPE t varray address
   MODIFY LIMIT 5 CASCADE;
```

Type altered.

Using Varrays in Temporary Tables

You can use varrays in temporary tables, which are tables whose rows are temporary and are specific to a user session (temporary tables were covered in the section "Creating a Table" in Chapter 10). The following example creates a temporary table named cust with varray temp table that contains a varray named addresses of type t varray address:

```
CREATE GLOBAL TEMPORARY TABLE cust with varray temp table (
           INTEGER PRIMARY KEY,
      first name VARCHAR2(10),
      last name VARCHAR2(10),
      addresses t varray address
```

Using a Different Tablespace for a Nested Table's **Storage Table**

By default, a nested table's storage table is created in the same tablespace as the parent table (a tablespace is an area used by the database to store objects such as tables—see the section "Creating a Table" in Chapter 10 for details).

In Oracle Database 10g and higher, you can specify a different tablespace for a nested table's storage table. The following example creates a table named <code>cust_with_nested_table</code> that contains a nested table named <code>addresses</code> of type <code>t_nested_table_address</code>; notice that the tablespace for the <code>nested_addresses2</code> storage table is the <code>users</code> tablespace:

You must have a tablespace named users in order for this example to work, and for this reason I've commented out the example in the collection_schema3.sql script. You can see all the tablespaces you have access to by performing the following query:

SELECT tablespace_name FROM user_tablespaces;

If you want to run the previous CREATE TABLE statement, you can edit the example in the collection_schema3.sql script to reference one of your tablespaces and then copy the statement into SQL*Plus and run it.

ANSI Support for Nested Tables

The American National Standards Institute (ANSI) specification includes a number of operators that may be used with nested tables. You'll learn about these operators in the following sections.

Equal and Not-Equal Operators

The equal (=) and not-equal (<>) operators compare two nested tables, which are considered equal when they satisfy all the following conditions:

- The tables are the same type.
- The tables are the same cardinality, that is, they contain the same number of elements.
- All the elements of the table have the same value.

The following equal_example() procedure illustrates the use of the equal and not-equal operators:

```
CREATE PROCEDURE equal example AS
       -- declare a type named t nested table
       TYPE t nested table IS TABLE OF VARCHAR2(10);
       -- create t nested table objects named v customer nested table1,
       -- v customer nested table2, and v customer nested table3;
       -- these objects are used to store the names of customers
       v customer nested table1 t nested table :=
         t nested table('Fred', 'George', 'Susan');
       v customer nested table2 t nested table :=
         t_nested_table('Fred', 'George', 'Susan');
       v customer nested table3 t nested table :=
         t nested table('John', 'George', 'Susan');
       v result BOOLEAN;
     BEGIN
       -- use = operator to compare v customer nested table1 with
       -- v customer nested table2 (they contain the same names, so
       -- v result is set to true)
       v result := v customer nested table1 = v customer nested table2;
       IF v result THEN
         DBMS OUTPUT.PUT LINE (
           'v customer nested table1 equal to v customer nested table2'
         );
       END IF;
       -- use <> operator to compare v customer nested table1 with
       -- v customer nested table3 (they are not equal because the first
       -- names, 'Fred' and 'John', are different and v result is set
       -- to true)
       v result := v customer nested_table1 <> v_customer_nested_table3;
       IF v result THEN
         DBMS OUTPUT.PUT LINE (
           'v customer nested table1 not equal to v customer nested table3'
         );
       END IF;
     END equal example;
```

The following example calls equal example ():

CALL equal example();

```
v customer nested table1 equal to v customer nested table2
v customer nested table1 not equal to v customer nested table3
```

IN and NOT IN Operators

The IN operator checks if the elements of one nested table appear in another nested table. Similarly, NOT IN checks if the elements of one nested table do not appear in another nested table. The following in example () procedure illustrates the use of IN and NOT IN:

```
CREATE PROCEDURE in example AS
       TYPE t nested table IS TABLE OF VARCHAR2(10);
```

```
v customer nested table1 t nested table :=
   t nested table('Fred', 'George', 'Susan');
 v customer nested table2 t nested table :=
   t nested table('John', 'George', 'Susan');
 v customer nested table3 t nested table :=
   t nested table ('Fred', 'George', 'Susan');
 v result BOOLEAN;
BEGIN
 -- use IN operator to check if elements of v customer nested table3
 -- are in v customer nested table1 (they are, so v result is
 -- set to true)
 v result := v_customer_nested_table3 IN
    (v customer nested table1);
 IF v result THEN
   DBMS OUTPUT.PUT LINE (
      'v customer nested table3 in v customer nested table1'
 END IF;
 -- use NOT IN operator to check if the elements of
 -- v customer nested table3 are not in v customer nested table2
 -- (they are not, so v result is set to true)
 v result := v customer nested table3 NOT IN
    (v customer nested table2);
 IF v result THEN
   DBMS OUTPUT.PUT LINE (
      'v customer nested table3 not in v customer nested table2'
 END IF;
END in example;
```

The following example calls in example ():

CALL in example();

```
v customer nested table3 in v customer nested table1
v customer nested table3 not in v customer nested table2
```

SUBMULTISET Operator

The SUBMULTISET operator checks whether the elements of one nested table are a subset of another nested table. The following submultiset example () procedure illustrates the use of SUBMULTISET:

```
CREATE PROCEDURE submultiset example AS
      TYPE t nested table IS TABLE OF VARCHAR2(10);
      v customer nested table1 t nested table :=
        t nested table ('Fred', 'George', 'Susan');
      v customer nested table2 t nested table :=
        t nested table('George', 'Fred', 'Susan', 'John', 'Steve');
      v result BOOLEAN;
    BEGIN
```

```
-- use SUBMULTISET operator to check if elements of
 -- v customer nested table1 are a subset of v customer nested table2
 -- (they are, so v result is set to true)
 v result :=
   v customer nested table1 SUBMULTISET OF v customer nested table2;
 IF v result THEN
   DBMS OUTPUT.PUT LINE (
      'v customer nested table1 subset of v customer nested table2'
 END IF;
END submultiset example;
```

The following example calls submultiset example():

```
CALL submultiset example();
   customer nested table1 subset of customer nested table2
```

MULTISET Operator

The MULTISET operator returns a nested table whose elements are set to certain combinations of elements from two supplied nested tables. There are three MULTISET operators:

- **MULTISET UNION** returns a nested table whose elements are set to the sum of the elements from two supplied nested tables.
- **MULTISET INTERSECT** returns a nested table whose elements are set to the elements that are common to two supplied nested tables.
- **MULTISET EXCEPT** returns a nested table whose elements are in the first supplied nested table but not in the second.

You may also use one of the following options with MULTISET:

- **ALL** indicates that all the applicable elements are in the returned nested table. ALL is the default. For example, MULTISET UNION ALL returns a nested table whose elements are set to the sum of elements from two supplied nested tables, and all elements, including duplicates, are in the returned nested table.
- **DISTINCT** indicates that only the non-duplicate (that is, distinct) elements are in the returned nested table. For example, MULTISET UNION DISTINCT returns a nested table whose elements are set to the sum of elements from two supplied nested tables, but duplicates are removed from the returned nested table.

The following multiset example () procedure illustrates the use of MULTISET:

```
CREATE PROCEDURE multiset example AS
       TYPE t nested table IS TABLE OF VARCHAR2(10);
       v customer nested table1 t nested table :=
         t nested table ('Fred', 'George', 'Susan');
       v customer nested table2 t nested table :=
         t nested table('George', 'Steve', 'Rob');
       v customer nested table3 t nested table;
```

```
v count INTEGER;
BEGIN
 -- use MULTISET UNION (returns a nested table whose elements
  -- are set to the sum of the two supplied nested tables)
 v customer nested table3 :=
    v customer nested table1 MULTISET UNION
      v customer nested table2;
  DBMS OUTPUT.PUT('UNION: ');
  FOR v count IN 1..v customer nested table3.COUNT LOOP
    DBMS OUTPUT.PUT(v customer nested table3(v count) || ' ');
  END LOOP:
  DBMS OUTPUT.PUT LINE(' ');
  -- use MULTISET UNION DISTINCT (DISTINCT indicates that only
  -- the non-duplicate elements of the two supplied nested tables
  -- are set in the returned nested table)
  v customer nested table3 :=
   v customer nested table1 MULTISET UNION DISTINCT
      v customer nested table2;
  DBMS OUTPUT.PUT('UNION DISTINCT: ');
  FOR v count IN 1..v customer nested table3.COUNT LOOP
    DBMS OUTPUT.PUT(v customer nested table3(v count) || ' ');
  END LOOP;
  DBMS OUTPUT.PUT LINE(' ');
  -- use MULTISET INTERSECT (returns a nested table whose elements
  -- are set to the elements that are common to the two supplied
  -- nested tables)
  v customer nested table3 :=
    v customer nested table1 MULTISET INTERSECT
      v customer nested table2;
  DBMS OUTPUT.PUT('INTERSECT: ');
  FOR v count IN 1..v customer nested table3.COUNT LOOP
    DBMS OUTPUT.PUT(v customer nested table3(v count) || ' ');
  END LOOP;
  DBMS OUTPUT.PUT LINE(' ');
  -- use MULTISET EXCEPT (returns a nested table whose
 -- elements are in the first nested table but not in
  -- the second)
  v customer nested table3 :=
    v customer nested table1 MULTISET EXCEPT
      v customer nested table2;
  DBMS OUTPUT.PUT LINE('EXCEPT: ');
  FOR v count IN 1..v customer nested table3.COUNT LOOP
    DBMS OUTPUT.PUT(v customer nested table3(v count) || ' ');
  END LOOP;
END multiset example;
```

The following example calls multiset example():

CALL multiset example(); UNION: Fred George Susan George Steve Rob UNION DISTINCT: Fred George Susan Steve Rob INTERSECT: George EXCEPT:

CARDINALITY() Function

The CARDINALITY () function returns the number of elements in a collection. The following cardinality example() procedure illustrates the use of CARDINALITY():

```
CREATE PROCEDURE cardinality example AS
      TYPE t nested table IS TABLE OF VARCHAR2(10);
      v customer nested table1 t nested table :=
        t nested table ('Fred', 'George', 'Susan');
      v cardinality INTEGER;
    BEGIN
      -- call CARDINALITY() to get the number of elements in
      -- v customer nested table1
      v cardinality := CARDINALITY(v customer nested table1);
      DBMS OUTPUT.PUT LINE('v cardinality = ' || v cardinality);
    END cardinality example;
```

The following example calls cardinality example ():

```
CALL cardinality example();
   v cardinality = 3
```

MEMBER OF Operator

The MEMBER OF operator checks whether an element is in a nested table. The following member of example () procedure illustrates the use of MEMBER OF:

```
CREATE PROCEDURE member of example AS
    TYPE t nested table IS TABLE OF VARCHAR2(10);
    v customer nested table1 t nested table :=
     t nested table('Fred', 'George', 'Susan');
    v result BOOLEAN;
    -- use MEMBER OF to check if 'George' is in
    -- v customer nested table1 (he is, so v result is set
    -- to true)
    v result := 'George' MEMBER OF v_customer_nested_table1;
    IF v result THEN
      DBMS OUTPUT.PUT LINE('''George'' is a member');
    END IF;
  END member of example;
```

The following example calls member of example ():

```
CALL member_of_example();
'George' is a member
```

SET() Function

The SET() function first converts a nested table into a set, then removes duplicate elements from the set, and finally returns the set as a nested table. The following set_example() procedure illustrates the use of SET():

```
CREATE PROCEDURE set example AS
      TYPE t nested table IS TABLE OF VARCHAR2(10);
      v customer nested table1 t nested table :=
        t nested table('Fred', 'George', 'Susan', 'George');
      v customer nested table2 t nested table;
      v count INTEGER;
    BEGIN
      -- call SET() to convert a nested table into a set,
      -- remove duplicate elements from the set, and get the set
      -- as a nested table
      v customer nested table2 := SET(v customer nested table1);
      DBMS OUTPUT.PUT('v customer nested table2: ');
      FOR v count IN 1..v customer nested table2.COUNT LOOP
        DBMS OUTPUT.PUT(v customer nested table2(v count) || ' ');
      END LOOP;
      DBMS OUTPUT.PUT LINE(' ');
    END set example;
```

The following example calls set example():

```
CALL set_example();
v customer nested table2: Fred George Susan
```

IS A SET Operator

The IS A SET operator checks if the elements in a nested table are distinct. The following is_a_set_example() procedure illustrates the use of IS A SET:

```
CREATE PROCEDURE is_a_set_example AS

TYPE t_nested_table IS TABLE OF VARCHAR2(10);

v_customer_nested_table1 t_nested_table :=
    t_nested_table('Fred', 'George', 'Susan', 'George');

v_result BOOLEAN;

BEGIN

-- use IS A SET operator to check if the elements in
    -- v_customer_nested_table1 are distinct (they are not, so
    -- v_result is set to false)

v_result := v_customer_nested_table1 IS A SET;

IF v_result THEN
    DBMS_OUTPUT.PUT_LINE('Elements are all unique');

ELSE
    DBMS_OUTPUT.PUT_LINE('Elements contain duplicates');

END IF;
```

```
END is a set example;
```

The following example calls is a set example():

```
CALL is a set example();
    Elements contain duplicates
```

IS EMPTY Operator

The IS EMPTY operator checks if a nested table doesn't contain elements. The following is empty example() procedure illustrates the use of IS EMPTY:

```
CREATE PROCEDURE is empty example AS
     TYPE t nested table IS TABLE OF VARCHAR2(10);
     v customer nested table1 t nested table :=
       t nested table('Fred', 'George', 'Susan');
     v result BOOLEAN;
   BEGIN
     -- use IS EMPTY operator to check if
     -- v customer nested table1 is empty (it is not, so
     -- v result is set to false)
     v result := v customer nested table1 IS EMPTY;
     IF v result THEN
       DBMS OUTPUT.PUT LINE('Nested table is empty');
       DBMS OUTPUT.PUT LINE('Nested table contains elements');
     END IF;
   END is empty example;
```

The following example calls is empty example():

CALL is empty example();

Nested table contains elements

COLLECT() Function

The COLLECT () function returns a nested table from a set of elements. The following query illustrates the use of COLLECT():

```
SELECT COLLECT(first name)
   FROM customers with varray;
   COLLECT (FIRST NAME)
   ______
   SYSTPfrFhAg+WRJGwW7ma9zy1KA==('Steve', 'John')
```

You can use CAST() to convert the elements returned by COLLECT() to a specific type, as shown in the following query:

SELECT CAST(COLLECT(first name) AS t table) FROM customers with varray;

```
CAST (COLLECT (FIRST_NAME) AST_TABLE)
-----
T TABLE ('Steve', 'John')
```

For your reference, the t_table type used in the previous example is created by the following statement in the collection schema3.sql script:

```
CREATE TYPE t_table AS TABLE OF VARCHAR2(10);
```

POWERMULTISET() Function

The POWERMULTISET() function returns all combinations of elements in a given nested table, as shown in the following query:

```
SELECT *
     FROM TABLE (
       POWERMULTISET(t table('This', 'is', 'a', 'test'))
     );
     COLUMN VALUE
     -----
     T TABLE ('This')
     T TABLE ('is')
     T TABLE ('This', 'is')
     T TABLE ('a')
     T TABLE ('This', 'a')
     T TABLE('is', 'a')
     T TABLE ('This', 'is', 'a')
     T TABLE ('test')
     T TABLE ('This', 'test')
     T TABLE('is', 'test')
     T TABLE('This', 'is', 'test')
     T TABLE('a', 'test')
     T_TABLE('This', 'a', 'test')
     T TABLE('is', 'a', 'test')
     T TABLE('This', 'is', 'a', 'test')
```

POWERMULTISET BY CARDINALITY() Function

The POWERMULTISET_BY_CARDINALITY() function returns the combinations of elements in a given nested table that have a specified number of elements (or "cardinality"). The following query illustrates the use of POWERMULTISET BY CARDINALITY(), specifying a cardinality of 3:

```
FROM TABLE(

POWERMULTISET_BY_CARDINALITY(

t_table('This', 'is', 'a', 'test'), 3

);

COLUMN_VALUE

T_TABLE('This', 'is', 'a')

T_TABLE('This', 'is', 'test')

T_TABLE('This', 'a', 'test')

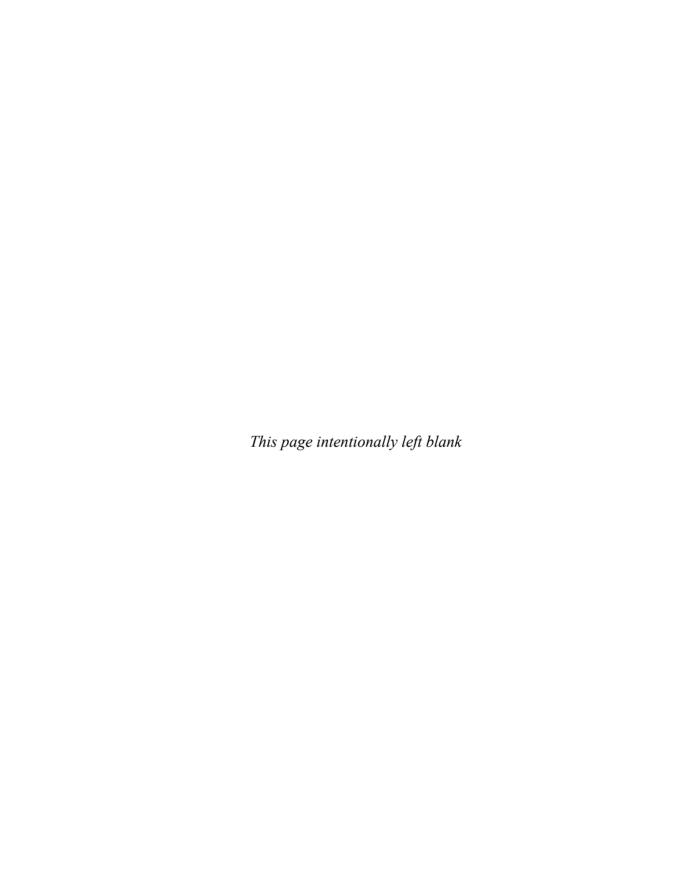
T_TABLE('Is', 'a', 'test')

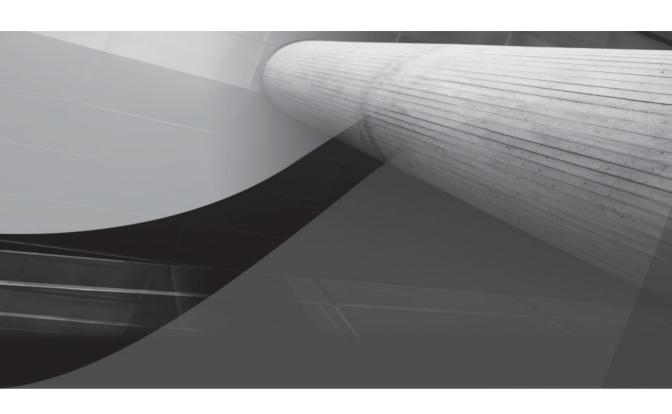
T_TABLE('is', 'a', 'test')
```

In this chapter, you have learned the following:

- Collections allow you to store sets of elements.
- There are three types of collections: varrays, nested tables, and associative arrays.
- A varray is similar to an array in Java; you can use a varray to store an ordered set of elements with each element having an index associated with it. The elements in a varray are of the same type, and a varray has one dimension. A varray has a maximum size that you set when creating it, but you can change the size later.
- A nested table is a table that is embedded within another table, and you can insert, update, and delete individual elements in a nested table. Because you can modify individual elements in a nested table, they are more flexible than a varray—a varray can be modified only as a whole. A nested table doesn't have a maximum size, and you can store an arbitrary number of elements in a nested table.
- An associative array is a set of key and value pairs. You can get the value from the array using the key (which may be a string) or an integer that specifies the position of the value in the array. An associative array is similar to a hash table in programming languages such as Java.
- A collection may itself contain embedded collections. Such a collection is known as a multilevel collection.

In the next chapter, you'll learn about large objects.





CHAPTER 14

Large Objects



n this chapter, you will do the following:

- Learn about large objects (LOBs)
- See files whose content will be used to populate example LOBs
- Examine the differences between the different types of LOBs
- Create tables containing LOBs
- Use LOBs in SQL and PL/SQL
- Examine the LONG and LONG RAW types
- See the Oracle Database 10g and 11g enhancements to LOBs

Introducing Large Objects (LOBs)

Today's websites demand more than just the storage and retrieval of text and numbers: they also require multimedia. Consequently, databases are now being called upon to store items like music and video. Prior to the release of Oracle Database 8, you had to store large blocks of character data using the LONG database type, and large blocks of binary data had to be stored using either the LONG RAW type or the shorter RAW type.

With the release Oracle Database 8, a new class of database types known as *large objects*, or LOBs for short, was introduced. LOBs may be used to store binary data, character data, and references to files. The binary data can contain images, music, video, documents, executables, and so on. LOBs can store up to 128 terabytes of data, depending on the database configuration.

The Example Files

You'll see the use of the following two files in this chapter:

- textContent.txt A text file
- binaryContent.doc A Microsoft Word file



NOTE

These files are contained in the sample_files directory, which is created when you extract the Zip file for this book. If you want to follow along with the examples, you should copy the sample_files directory to the C partition on your database server. If you're using Linux or Unix, you can copy the directory to one of your partitions.

The file textContent.txt contains an extract from Shakespeare's play Macbeth. The following text is the speech made by Macbeth shortly before he is killed:

To-morrow, and to-morrow, and to-morrow, Creeps in this petty pace from day to day, To the last syllable of recorded time; And all our yesterdays have lighted fools The way to a dusty death. Out, out, brief candle! Life's but a walking shadow; a poor player, That struts and frets his hour upon the stage, And then is heard no more: it is a tale Told by an idiot, full of sound and fury, Signifying nothing.

The binaryContent.doc file is a Word document that contains the same text as textContent.txt. (A Word document is a binary file.) Although a Word document is used in the examples, you can use any binary file, for example, MP3, DivX, JPEG, MPEG, PDF, or EXE. I have tested the examples with all these types of files.

Large Object Types

There are four LOB types:

- **CLOB** The character LOB type, which is used to store character data.
- **NCLOB** The National Character Set LOB type, which is used to store multiple byte character data (typically used for non-English characters). You can learn all about non-English character sets in the Oracle Database Globalization Support Guide published by Oracle Corporation.
- **BLOB** The binary LOB type, which is used to store binary data.
- **BFILE** The binary FILE type, which is used to store a pointer to a file. The file can be on a hard disk, CD, DVD, Blu-Ray disk, HD-DVD, or any other device that is accessible through the database server's file system. The file itself is never stored in the database, only a pointer to the file.

Prior to Oracle Database 8 your only choice for storing large amounts of data was to use the LONG and LONG RAW types (you could also use the RAW type for storing binary data of less than 4 kilobytes in size). The LOB types have three advantages over these older types:

- A LOB can store up to 128 terabytes of data. This is far more data than you can store in a LONG and LONG RAW column, which may only store up to 2 gigabytes of data.
- A table can have multiple LOB columns, but a table can only have one LONG or LONG RAW column.
- LOB data can be accessed in random order; LONG and LONG RAW data can be accessed only in sequential order.

A LOB consists of two parts:

- The LOB locator A pointer that specifies the location of the LOB data
- The LOB data The actual character or byte data stored in the LOB

Depending on the amount of data stored in a CLOB, NCLOB or BLOB column, the data will be stored either inside or outside of the table. If the data is less than 4 kilobytes, the data is stored in the same table; otherwise, the data is stored outside the table. With a BFILE column, only the locator is stored in the table—and the locator points to an external file stored in the file system.

Creating Tables Containing Large Objects

You'll see the use of the following three tables in this section:

- The clob_content table, which contains a CLOB column named clob_column
- The blob_content table, which contains a BLOB column named blob_column
- The bfile_content table, which contains a BFILE column named bfile_column

I've provided an SQL*Plus script named lob_schema.sql in the SQL directory. This script may be run using Oracle Database 8 and higher. The script creates a user named lob_user with a password of lob_password, and it creates the tables and PL/SQL code used in the first part of this chapter. After the script completes, you will be logged in as lob user.

The three tables are created using the following statements in the script:

Using Large Objects in SQL

In this section, you'll learn how to use SQL to manipulate large objects. You'll start by examining CLOB and BLOB objects and then move on to BFILE objects.

Using CLOBs and BLOBs

The following sections show how to populate CLOB and BLOB objects with data, retrieve the data, and then modify the data.

Populating CLOBs and BLOBs with Data

The following INSERT statements add two rows to the clob content table; notice the use of the TO CLOB() function to convert the text to a CLOB:

```
INSERT INTO clob content (
     id, clob column
   ) VALUES (
     1, TO CLOB('Creeps in this petty pace')
   );
   INSERT INTO clob content (
     id, clob column
   ) VALUES (
     2, TO CLOB(' from day to day')
   );
```

The following INSERT statements add two rows to the blob content table; notice the use of the TO BLOB () function to convert the numbers to a BLOB (the first statement contains a binary number, and the second contains a hexadecimal number):

```
INSERT INTO blob content (
      id, blob column
    ) VALUES (
      1, TO BLOB('100111010101011111')
    INSERT INTO blob content (
      id, blob column
    ) VALUES (
      2, TO BLOB('A0FFB71CF90DE')
     );
```

Retrieving Data from CLOBs

The following query retrieves the rows from the clob content table:

SELECT *

```
FROM clob_content;
```

```
TD
CLOB COLUMN
        1
Creeps in this petty pace
 from day to day
```

The next query attempts to retrieve the row from the blob content table and fails:

SELECT *

```
FROM blob_content;
```

SP2-0678: Column or attribute type can not be displayed by SQL*Plus

This example fails because SQL*Plus cannot display the binary data in a BLOB. You'll learn how to retrieve the data from a BLOB later in the section "Using Large Objects in PL/SQL."

You can, however, get the non-BLOB columns from the table:

```
SELECT id
FROM blob_content;

ID
------
1
```

Modifying the Data in CLOBs and BLOBs

You should feel free to run the UPDATE and INSERT statements shown in this section. The following UPDATE statements show how you modify the contents of a CLOB and a BLOB:

```
UPDATE clob_content
SET clob_column = TO_CLOB('What light through yonder window breaks')
WHERE id = 1;

UPDATE blob_content
SET blob_column = TO_BLOB('11100110101011111')
WHERE id = 1;
```

You can also initialize the LOB locator, but not store actual data in the LOB. You do this using the EMPTY_CLOB() function to store an empty CLOB, and EMPTY_BLOB() to store an empty BLOB:

```
INSERT INTO clob_content(
   id, clob_column
) VALUES (
   3, EMPTY_CLOB()
);

INSERT INTO blob_content(
   id, blob_column
) VALUES (
   3, EMPTY_BLOB()
);
```

These statements initialize the LOB locator, but set the LOB data to empty.

You can also use EMPTY_CLOB() and EMPTY_BLOB() in UPDATE statements when you want to empty out the LOB data. For example:

```
UPDATE clob_content
SET clob_column = EMPTY_CLOB()
WHERE id = 1;

UPDATE blob_content
SET blob_column = EMPTY_BLOB()
WHERE id = 1;
```

If you ran any of the INSERT and UPDATE statements shown in this section, go ahead and roll back the changes so that your output matches mine in the rest of this chapter:

ROLLBACK:

Using BFILEs

A BFILE stores a pointer to a file that is accessible through the database server's file system. The important point to remember is that these files are stored outside of the database. A BFILE can point to files located on any media: a hard disk, CD, DVD, Blu-Ray, HD-DVD, and so on.



NOTE

A BFILE contains a pointer to an external file. The actual file itself is never stored in the database, only a pointer to that file is stored. The file must be accessible through the database server's file system.

Creating a Directory Object

Before you can store a pointer to a file in a BFILE, you must first create a directory object in the database. The directory object stores the directory in the file system where the files are located. You create a directory object using the CREATE DIRECTORY statement, and to run this statement you must have the CREATE ANY DIRECTORY database privilege.

The following example (contained in lob schema.sql) creates a directory object named SAMPLE FILES DIR for the file system directory C:\sample files:

CREATE DIRECTORY SAMPLE FILES DIR AS 'C:\sample files';



Windows uses the backslash character (\) in directories, while Linux and Unix use the forward slash character (/). Also, if your sample files directory is not stored in the C partition, then you need to specify the appropriate path in the previous example.

When you create a directory object you must ensure that

- The actual directory exists in the file system.
- The user account in the operating system that was used to install the Oracle software has read permission on the directory and on any files that are to be pointed to by a BFILE in the database.

If you're using Windows, you shouldn't need to worry about the second point. The Oracle database software should have been installed using a user account that has administrator privileges, and such a user account has read permission on everything in the file system. If you're using Linux or Unix, you'll need to grant read access to the appropriate Oracle user account that owns the database (you do this using the chmod command).

Populating a BFILE Column with a Pointer to a File

Because a BFILE is just a pointer to an external file, populating a BFILE column is very simple. All you have to do is to use the Oracle database's BFILENAME() function to populate a BFILE with a pointer to your external file. The BFILENAME() function accepts two parameters: the directory object's name and the name of the file.

For example, the following INSERT adds a row to the bfile_content table; notice that the BFILENAME() function is used to populate bfile_column with a pointer to the textContent.txt file:

```
INSERT INTO bfile_content (
    id, bfile_column
) VALUES (
    1, BFILENAME('SAMPLE_FILES_DIR', 'textContent.txt')
);
```

The next INSERT adds a row to the bfile_content table; notice that the BFILENAME() function is used to populate bfile column with a pointer to the binaryContent.doc file:

```
INSERT INTO bfile_content (
   id, bfile_column
) VALUES (
   2, BFILENAME('SAMPLE_FILES_DIR', 'binaryContent.doc')
);
```

The following query attempts to retrieve the rows from bfile_content and fails because SQL*Plus cannot display the content in a BFILE:

SELECT *

```
FROM bfile_content;
SP2-0678: Column or attribute type can not be displayed by SQL*Plus
```

You may use PL/SQL to access the content in a BFILE or a BLOB, and you'll learn how to do that next.

Using Large Objects in PL/SQL

In this section, you'll learn how to use LOBs in PL/SQL. You'll start off by examining the methods in the DBMS_LOB package, which comes with the database. Later, you'll see plenty of PL/SQL programs that show how to use the DBMS_LOB methods to read data in a LOB, copy data from one LOB to another, search data in a LOB, copy data from a file to a LOB, copy data from a LOB to a file, and much more.

Table 14-1 summarizes the most commonly used methods in the DBMS_LOB package. In the following sections, you'll see the details of some of the methods shown in the previous table. You can see all the DBMS_LOB methods in the *Oracle Database PL/SQL Packages and Types Reference* manual published by Oracle Corporation.

Method

APPEND(dest lob, src lob) CLOSE (lob)

COMPARE (lob1, lob2, amount, offset1, offset2)

CONVERTTOBLOB (dest blob, src clob, amount, dest offset, src offset, blob csid, lang context, warning)

CONVERTTOCLOB(dest clob, src blob, amount, dest offset, src offset, blob csid, lang context, warning)

COPY (dest lob, src lob, amount, dest offset, src offset)

CREATETEMPORARY (lob, cache, duration)

ERASE(lob, amount, offset)

FILECLOSE (bfile)

FILECLOSEALL()

FILEEXISTS (bfile)

FILEGETNAME (bfile, directory, filename)

FILEISOPEN (bfile)

FILEOPEN (bfile, open mode)

Description

Adds the data read from src lob to the end of dest lob.

Closes a previously opened LOB.

Compares the data stored in 1ob1 and 1ob2, starting at offset1 in 10b1 and offset2 in 10b2. Offsets always start at 1, which is the position of the first character or byte in the data.

The data in the LOBs are compared over a maximum number of characters or bytes (the maximum is specified in amount).

Converts the character data read from src clob into binary data written to dest blob.

The read begins at src offset in src clob, and the write begins at dest offset in dest blob.

blob csid is the desired character set for the converted data written to dest blob. You should typically use DBMS LOB. DEFAULT CSID, which is the default character set for the database. lang context is the language context to use when converting the

characters read from src clob. You should typically use DBMS LOB. DEFAULT LANG CTX, which is the default language context for the database.

warning is set to DBMS LOB.WARN INCONVERTIBLE_CHAR if there was a character that could not be converted.

Converts the binary data read from src blob into character data written to dest clob.

blob csid is the character set for the data read from dest blob. You should typically use DBMS LOB. DEFAULT CSID.

lang context is the language context to use when writing the converted characters to dest clob. You should typically use DBMS LOB. DEFAULT LANG CTX.

warning is set to DBMS LOB. WARN INCONVERTIBLE CHAR if there was a character that could not be converted.

Copies data from src lob to dest lob, starting at the offsets for a total amount of characters or bytes.

Creates a temporary LOB in the user's default temporary tablespace.

Erases data from a LOB, starting at the offset for a total amount of characters or bytes.

Closes bfile. You should use the newer CLOSE () method instead of FILECLOSE().

Closes all previously opened BFILES.

Checks if the external file pointed to by bfile actually exists.

Returns the directory and filename of the external file pointed to by bfile.

Checks if bfile is currently open. You should use the newer ISOPEN() method instead of FILEISOPEN().

Opens bfile in the indicated mode, which can be set only to DBMS LOB.FILE READONLY, which indicates the file may only be read from (and never written to). You should use the newer OPEN() method instead of FILEOPEN().

Method	Description
FREETEMPORARY (10b)	Frees a temporary LOB.
GETCHUNKSIZE (10b)	Returns the chunk size used when reading and writing the data stored in the LOB. A chunk is a unit of data.
<pre>GET_STORAGE_LIMIT()</pre>	Returns the maximum allowable size for a LOB.
GETLENGTH (10b)	Gets the length of the data stored in the LOB.
<pre>INSTR(lob, pattern, offset, n)</pre>	Returns the starting position of characters or bytes that match the <i>nth</i> occurrence of a pattern in the LOB data. The data is read from the LOB starting at the offset.
ISOPEN(10b)	Checks if the LOB was already opened.
ISTEMPORARY (10b)	Checks if the LOB is a temporary LOB.
LOADFROMFILE (dest_lob, src_bfile, amount, dest_offset, src_offset)	Loads the data retrieved via <code>src_bfile</code> to <code>dest_lob</code> , starting at the offsets for a total amount of characters or bytes; <code>src_bfile</code> is a <code>BFILE</code> that points to an external file. <code>LOADFROMFILE()</code> is old, and you should use the higher-performance <code>LOADBLOBFROMFILE()</code> or <code>LOADCLOBFROMFILE()</code> methods.
LOADBLOBFROMFILE (dest_blob, src_bfile, amount, dest_offset, src_offset)	Loads the data retrieved via <code>src_bfile</code> to <code>dest_blob</code> , starting at the offsets for a total amount of bytes; <code>src_bfile</code> is a BFILE that points to an external file. LOADBLOBFROMFILE() offers improved performance over LOADFROMFILE() when using a BLOB.
LOADCLOBFROMFILE (dest_clob, src_bfile, amount, dest_offset, src_offset, src_csid, lang_context, warning)	Loads the data retrieved via <code>src_bfile</code> to <code>dest_clob</code> , starting at the offsets for a total amount of characters; <code>src_bfile</code> is a <code>BFILE</code> that points to an external file. LOADCLOBFROMFILE () offers improved performance over <code>LOADFROMFILE</code> () when using a <code>CLOB/NCLOB</code> .
LOBMAXSIZE	Returns the maximum size for a LOB in bytes (currently 2 ⁶⁴).
OPEN(lob, open_mode)	Opens the LOB in the indicated mode, which may be set to ■ DBMS_LOB.FILE_READONLY, which indicates the LOB may only be read from ■ DBMS_LOB.FILE_READWRITE, which indicates the LOB may read from and written to
DEAD(lob amount offeet buffer)	read from and written to Reads the data from the LOB and stores them in the buffer variable,
READ(lob, amount, offset, buffer)	starting at the offset in the LOB for a total amount of characters or bytes.
SUBSTR(lob, amount, offset)	Returns part of the LOB data, starting at the offset in the LOB for a total amount of characters or bytes.
TRIM(lob, newlen)	Trims the LOB data to the specified shorter length.
WRITE(lob, amount, offset, buffer)	Writes the data from the <code>buffer</code> variable to the LOB, starting at the offset in the LOB for a total amount of characters or bytes.

Writes the data from the buffer variable to the end of the LOB,

starting at the offset in the LOB for a total amount of characters or bytes.

WRITEAPPEND(lob, amount, buffer)

APPEND()

APPEND () adds the data in a source LOB to the end of a destination LOB. There are two versions of APPEND():

```
DBMS LOB.APPEND (
     dest lob IN OUT NOCOPY BLOB,
     src lob IN
    );
    DBMS LOB.APPEND(
     dest lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
     src lob IN CLOB/NCLOB CHARACTER SET dest lob%CHARSET
    );
```

where

- dest_lob is the destination LOB to which the data is appended.
- src lob is the source LOB from which the data is read.
- CHARACTER SET ANY CS means the data in dest lob can be any character set.
- CHARACTER SET dest lob%CHARSET is the character set of dest lob.

The following table shows the exception thrown by APPEND().

Exception

Thrown When

VALUE ERROR

Either dest lob or src lob is null.

CLOSE()

CLOSE() closes a previously opened LOB. There are three versions of CLOSE():

```
DBMS LOB.CLOSE(
      lob IN OUT NOCOPY BLOB
    );
    DBMS LOB.CLOSE (
      lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS
    );
    DBMS LOB.CLOSE (
      lob IN OUT NOCOPY BFILE
    );
```

where lob is the LOB to be closed.

COMPARE()

COMPARE () compares the data stored in two LOBs, starting at the offsets over a total amount of characters or bytes. There are three versions of COMPARE ():

```
DBMS LOB.COMPARE(
       10b1 IN BLOB,
       10b2 IN BLOB,
       amount IN INTEGER := 4294967295,
       offset1 IN INTEGER := 1,
       offset2 IN INTEGER := 1
     ) RETURN INTEGER;
      DBMS LOB.COMPARE (
       10b1 IN CLOB/NCLOB CHARACTER SET ANY CS,
       10b2 IN CLOB/NCLOB CHARACTER SET lob 1%CHARSET,
       amount IN INTEGER := 4294967295,
       offset1 IN INTEGER := 1,
       offset2 IN INTEGER := 1
     ) RETURN INTEGER;
     DBMS LOB.COMPARE (
       10b1 IN BFILE,
       1ob2 IN BFILE,
       amount IN INTEGER,
       offset1 IN INTEGER := 1,
       offset2 IN INTEGER := 1
     ) RETURN INTEGER;
```

where

- 10b1 and 10b2 are the LOBs to compare.
- amount is the maximum number of characters to read from a CLOB/NCLOB, or the maximum number of bytes to read from a BLOB/BFILE.
- offset1 and offset2 are the offsets in characters or bytes in 10b1 and 10b2 to start the comparison (the offsets start at 1).

COMPARE() returns

- 0 if the LOBs are identical.
- 1 if the LOBs aren't identical.
- Null if
 - \blacksquare amount < 1
 - amount > LOBMAXSIZE (Note: LOBMAXSIZE is the maximum size of the LOB)
 - \blacksquare offset1 or offset2 < 1
 - offset1 or offset2 > LOBMAXSIZE

The following table shows the exceptions thrown by COMPARE ().

Exception	Thrown When
UNOPENED_FILE	The file hasn't been opened yet.
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	You don't have privileges to access the directory.
INVALID_DIRECTORY	The directory is invalid.
INVALID OPERATION	The file exists, but you don't have privileges to access the file.

COPY()

COPY () copies data from a source LOB to a destination LOB, starting at the offsets for a total amount of characters or bytes. There are two versions of COPY():

```
DBMS LOB.COPY(
          dest lob IN OUT NOCOPY BLOB,
         dest_lob IN OUT NOOT _____,
src_lob IN BLOB,
amount IN INTEGER,
dest_offset IN INTEGER := 1,
src_offset IN INTEGER := 1
       );
       DBMS LOB.COPY(
         dest lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
         src_lob IN CLOB/NCLOB CHARACTER SET dest_lob%CHARSET,
amount IN INTEGER,
dest_offset IN INTEGER := 1,
src_offset IN INTEGER := 1
       );
```

where

- dest lob and src lob are the LOBs to write to and read from, respectively.
- amount is the maximum number of characters to read from a CLOB/NCLOB, or the maximum number of bytes to read from a BLOB/BFILE.
- dest offset and src offset are the offsets in characters or bytes in dest lob and src lob to start the copy (the offsets start at 1).

The following table shows the exceptions thrown by COPY ().

Exception

Thrown When

```
VALUE ERROR
                      Any of the parameters are null.
INVALID ARGVAL
                       Either:
                      ■ src offset < 1
                       \blacksquare dest offset < 1
                       ■ src offset > LOBMAXSIZE
                       ■ dest offset > LOBMAXSIZE
                       \blacksquare amount < 1
                       ■ amount > LOBMAXSIZE
```

CREATETEMPORARY()

CREATETEMPORARY () creates a temporary LOB in the user's default temporary tablespace. There are two versions of CREATETEMPORARY ():

where

- 10b is the temporary LOB to create.
- cache indicates whether the LOB should be read into the buffer cache (true for yes, false for no).
- duration is a hint (can be set to SESSION, TRANSACTION, or CALL) as to whether the temporary LOB is removed at the end of the session, transaction, or call (the default is SESSION).

The following table shows the exception thrown by CREATETEMPORARY ().

Exception

Thrown When

VALUE ERROR

The 10b parameter is null.

ERASE()

ERASE() removes data from a LOB, starting at the offset for a total amount of characters or bytes. There are two versions of ERASE():

```
DBMS_LOB.ERASE(

lob IN OUT NOCOPY BLOB,

amount IN OUT NOCOPY INTEGER,

offset IN INTEGER := 1
);

DBMS_LOB.ERASE(

lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY_CS,

amount IN OUT NOCOPY INTEGER,

offset IN INTEGER := 1
);
```

where

■ 10b is the LOB to erase.

- amount is the maximum number of characters to read from a CLOB/NCLOB, or the number of bytes to read from a BLOB.
- offset is the offset in characters or bytes in 10b to start the erasure (the offset starts at 1).

The following table shows the exceptions thrown by ERASE().

Exception	Thrown When
VALUE_ERROR	Any of the parameters are null.
INVALID_ARGVAL	Either:
	\blacksquare amount < 1
	\blacksquare amount > LOBMAXSIZE
	\blacksquare offset < 1
	■ offset > LOBMAXSIZE

FILECLOSE()

FILECLOSE() closes a BFILE. You should use the newer CLOSE() procedure, as Oracle Corporation does not plan to extend the older FILECLOSE () procedure. I'm only including coverage of FILECLOSE () here so you can understand older programs.

```
DBMS LOB.FILECLOSE(
      bfile IN OUT NOCOPY BFILE
```

where bfile is the BFILE to close.

The following table shows the exceptions thrown by FILECLOSE ().

Exception	Thrown When
VALUE_ERROR	The bfile parameter is null.
UNOPENED_FILE	The file hasn't been opened yet.
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	You don't have privileges to access the directory.
INVALID_DIRECTORY	The directory is invalid.
INVALID_OPERATION	The file exists, but you don't have privileges to access the file.

FILECLOSEALL()

FILECLOSEALL() closes all BFILE objects.

```
DBMS LOB.FILECLOSEALL;
```

The following table shows the exception thrown by FILECLOSEALL().

Exception	Thrown When
UNOPENED FILE	No files have been opened in the session.

FILEEXISTS()

FILEEXISTS () checks if a file exists.

```
DBMS_LOB.FILEEXISTS(

bfile IN BFILE

) RETURN INTEGER;
```

where bfile is a BFILE that points to an external file.

FILEEXISTS() returns

- 0 if the file doesn't exist.
- 1 if the file exists.

The following table shows the exceptions thrown by FILEEXISTS().

Exception	Thrown When
VALUE_ERROR	The bfile parameter is null.
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	You don't have privileges to access the directory.
INVALID DIRECTORY	The directory is invalid.

FILEGETNAME()

FILEGETNAME () returns the directory and filename from a BFILE.

```
DBMS_LOB.FILEGETNAME(
bfile IN BFILE,
directory OUT VARCHAR2,
filename OUT VARCHAR2
);
```

where

- *bfile* is the pointer to the file.
- *directory* is the directory where the file is stored.
- filename is the name of the file.

The following table shows the exceptions thrown by FILEGETNAME ().

Exception	Thrown When
VALUE_ERROR	Any of the input parameters are null or invalid.
INVALID ARGVAL	The directory or filename parameters are null.

FILEISOPEN()

FILEISOPEN() checks if a file is open. You should use the newer ISOPEN() procedure to check if a file is open in your own programs, as Oracle Corporation does not plan to extend the older FILEISOPEN() method. I'm including coverage of FILEISOPEN() here only so you can understand older programs.

```
DBMS LOB.FILEISOPEN (
      bfile IN BFILE
    ) RETURN INTEGER;
```

where bfile is the pointer to the file.

FILEISOPEN() returns

- 0 if the file isn't open.
- 1 if the file is open.

The following table shows the exceptions thrown by FILEISOPEN().

Exception	Thrown When
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	You don't have privileges to access the directory.
INVALID_DIRECTORY	The directory is invalid.
INVALID OPERATION	The file doesn't exist or you don't have privileges to access the file.

FILEOPEN()

FILEOPEN() opens a file. You should use the newer OPEN() procedure to open a file in your own programs, as Oracle Corporation does not plan to extend the older FILEOPEN() procedure. I'm including coverage of FILEOPEN() here only so you can understand older programs.

```
DBMS LOB.FILEOPEN (
    bfile IN OUT NOCOPY BFILE,
     open mode IN BINARY INTEGER := DBMS LOB.FILE READONLY
   );
```

where

- bfile is the pointer to the file.
- open mode indicates the open mode; the only open mode is DBMS LOB.FILE READONLY, which indicates the file may be read from.

The following table shows the exceptions thrown by FILEOPEN().

Exception	Thrown When
VALUE_ERROR	Any of the input parameters are null or invalid.
INVALID_ARGVAL	The open_mode is not set to DBMS_LOB.FILE_READONLY.
OPEN_TOOMANY	An attempt was made to open more than SESSION_MAX_OPEN_FILES files, where SESSION_MAX_OPEN_FILES is a database initialization parameter set by a database administrator.
NOEXIST_DIRECTORY	The directory doesn't exist.
INVALID_DIRECTORY	The directory is invalid.
INVALID_OPERATION	The file exists, but you don't have privileges to access the file.

FREETEMPORARY()

FREETEMPORARY () frees a temporary LOB from the default temporary tablespace of the user. There are two versions of FREETEMPORARY ():

```
DBMS_LOB.FREETEMPORARY (
    lob IN OUT NOCOPY BLOB
);

DBMS_LOB.FREETEMPORARY (
    lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY_CS
);
```

where 10b is the lob to be freed.

The following table shows the exception thrown by FREETEMPORARY().

Exception

Thrown When

VALUE ERROR

Any of the input parameters are null or invalid.

GETCHUNKSIZE()

GETCHUNKSIZE () returns the chunk size when reading and writing LOB data (a chunk is a unit of data). There are two versions of GETCHUNKSIZE ():

```
DBMS_LOB.GETCHUNKSIZE(
    lob IN BLOB
) RETURN INTEGER;

DBMS_LOB.GETCHUNKSIZE(
    lob IN CLOB/NCLOB CHARACTER SET ANY_CS
) RETURN INTEGER;
```

where lob is the LOB to get the chunk size for.

GETCHUNKSIZE() returns

- The chunk size in bytes for a BLOB
- The chunk size in characters for a CLOB/NCLOB

The following table shows the exception thrown by GETCHUNKSIZE().

Exception

Thrown When

VALUE ERROR

The 10b parameter is null.

GET_STORAGE_LIMIT()

GET STORAGE LIMIT() returns the maximum allowable size for a LOB.

```
DBMS_LOB.GET_STORAGE_LIMIT()
RETURN INTEGER;
```

GETLENGTH()

GETLENGTH() returns the length of the LOB data. There are three versions of GETLENGTH():

```
DBMS LOB.GETLENGTH (
       lob IN BLOB
     ) RETURN INTEGER;
     DBMS LOB.GETLENGTH (
       lob IN CLOB/NCLOB CHARACTER SET ANY CS
     ) RETURN INTEGER;
     DBMS LOB.GETLENGTH (
       bfile IN BFILE
     ) RETURN INTEGER;
```

where

- 10b is the BLOB, CLOB, or NCLOB data to get the length of.
- bfile is the BFILE data to get the length of.

GETLENGTH() returns

- The length in bytes for a BLOB or BFILE
- The length in characters for a CLOB or NCLOB

The following table shows the exception thrown by GETLENGTH().

Exception

Thrown When

VALUE ERROR

The lob or bfile parameter is null.

INSTR()

INSTR() returns the starting position of characters that match the *nth* occurrence of a pattern in the LOB data, starting at an offset. There are three versions of INSTR():

```
DBMS LOB.INSTR(
       lob IN BLOB,
       pattern IN RAW,
       offset IN INTEGER := 1,
             IN INTEGER := 1
     ) RETURN INTEGER;
     DBMS LOB.INSTR(
       10b IN CLOB/NCLOB CHARACTER SET ANY CS,
       pattern IN VARCHAR2 CHARACTER SET lob%CHARSET,
       offset IN INTEGER := 1,
       n IN INTEGER := 1
     ) RETURN INTEGER;
```

494 Oracle Database 11*g* SQL

where

- 10b is the BLOB, CLOB, or NCLOB to read from.
- bfile is the BFILE to read from.
- pattern is the pattern to search for in the LOB data; the pattern is a group of RAW bytes for a BLOB or BFILE, and a VARCHAR2 character string for a CLOB; the maximum size of the pattern is 16,383 bytes.
- offset is the offset to start reading data from the LOB (the offset starts at 1).
- \blacksquare *n* is the occurrence of the pattern to search the data for.

INSTR() returns

- The offset of the start of the pattern (if found)
- Zero if the pattern isn't found
- Null if
 - Any of the IN parameters are null or invalid
 - offset < 1 or offset > LOBMAXSIZE
 - n < 1 or n > LOBMAXSIZE

The following table shows the exceptions thrown by INSTR().

Exception	Thrown When
VALUE_ERROR	Any of the input parameters are null or invalid.
UNOPENED_FILE	The BFILE isn't open.
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	The directory exists, but you don't have privileges to access the directory.
INVALID_DIRECTORY	The directory is invalid.
INVALID_OPERATION	The file exists, but you don't have privileges to access the file.

ISOPEN()

ISOPEN() checks if the LOB was already opened. There are three versions of ISOPEN():

```
DBMS LOB.ISOPEN(
      lob IN BLOB
    ) RETURN INTEGER;
    DBMS LOB.ISOPEN (
      lob IN CLOB/NCLOB CHARACTER SET ANY CS
    ) RETURN INTEGER;
    DBMS LOB. ISOPEN (
      bfile IN BFILE
    ) RETURN INTEGER;
```

- 10b is the BLOB, CLOB, or NCLOB to check.
- bfile is the BFILE to check.

ISOPEN() returns

- 0 if the LOB isn't open.
- 1 if the LOB is open.

The following table shows the exception thrown by ISOPEN().

Exception

Thrown When

VALUE ERROR

The lob or bfile parameter is null or invalid.

ISTEMPORARY()

ISTEMPORARY () checks if the LOB is a temporary LOB. There are two versions of ISTEMPORARY():

```
DBMS LOB.ISTEMPORARY (
      lob IN BLOB
    ) RETURN INTEGER;
    DBMS LOB.ISTEMPORARY (
      lob IN CLOB/NCLOB CHARACTER SET ANY CS
    ) RETURN INTEGER;
```

where

10b is the LOB to check.

ISTEMPORARY() returns

- 0 if the LOB isn't temporary.
- 1 if the LOB is temporary.

The following table shows the exception thrown by ISTEMPORARY().

Exception

Thrown When

VALUE ERROR

The 10b parameter is null or invalid.

LOADFROMFILE()

LOADFROMFILE () loads data retrieved via a BFILE into a CLOB, NCLOB, or BLOB, starting at the offsets for a total amount of characters or bytes. You should use the higher-performance LOADCLOBFROMFILE() or LOADBLOBFROMFILE() procedures in your own programs, and I'm including coverage of LOADFROMFILE () here only so you can understand older programs.

There are two versions of LOADFROMFILE():

```
DBMS LOB.LOADFROMFILE(
       dest lob IN OUT NOCOPY BLOB,
      src_bfile IN BFILE,
amount IN INTEGER,
dest_offset IN INTEGER := 1,
src_offset IN INTEGER := 1
    DBMS LOB.LOADFROMFILE (
      dest lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
      src_bfile IN BFILE,
      amount IN INTEGER,
dest_offset IN INTEGER := 1,
src_offset IN INTEGER := 1
    );
```

where

- dest lob is the LOB into which the data is to be written.
- src bfile is the pointer to the file from which the data is to be read.
- amount is the maximum number of bytes or characters to read from src bfile.
- dest offset is the offset in bytes or characters in dest 10b to start writing data (the offset starts at 1).
- src offset is the offset in bytes in src bfile to start reading data (the offset starts at 1).

The following table shows the exceptions thrown by LOADFROMFILE ().

Exception

Thrown When

INVALID ARGVAL

VALUE ERROR

Any of the input parameters are null or invalid.

```
\blacksquare src offset < 1
■ dest offset < 1
■ src offset > LOBMAXSIZE
■ dest offset > LOBMAXSIZE
```

 \blacksquare amount < 1

■ amount > LOBMAXSIZE

LOADBLOBFROMFILE()

LOADBLOBFROMFILE() loads data retrieved via a BFILE into a BLOB. LOADBLOBFROMFILE() offers improved performance over the LOADFROMFILE () method when using a BLOB.

```
DBMS LOB.LOADBLOBFROMFILE (
      dest blob IN OUT NOCOPY BLOB,
      src bfile IN BFILE,
      amount IN INTEGER,

dest_offset IN OUT INTEGER := 1,

src offset IN OUT INTEGER := 1
```

where

- dest blob is the BLOB into which the data is to be written.
- src bfile is the pointer to the file from which the data is to be read.
- amount is the maximum number of bytes to read from src bfile.
- dest offset is the offset in bytes in dest lob to start writing data (the offset starts at 1).
- src offset is the offset in bytes in src bfile to start reading data (the offset starts

The following table shows the exceptions thrown by LOADBLOBFROMFILE().

Exception

Thrown When

VALUE ERROR

Any of the input parameters are null or invalid.

INVALID ARGVAL

Either:

```
\blacksquare src offset < 1
\blacksquare dest offset < 1
```

■ src offset > LOBMAXSIZE ■ dest offset > LOBMAXSIZE

 \blacksquare amount < 1

■ amount > LOBMAXSIZE

LOADCLOBFROMFILE()

LOADCLOBFROMFILE() loads data retrieved via a BFILE into a CLOB/NCLOB. LOADCLOBFROMFILE() offers improved performance over the LOADFROMFILE() method when using a CLOB/NCLOB. LOADCLOBFROMFILE () also automatically converts binary data to character data.

```
DBMS LOB.LOADCLOBFROMFILE (
    src_bfile IN BFILE,
amount IN INTEGE
                         INTEGER,
    dest offset IN OUT
                        INTEGER,
    src_offset IN OUT src_csid IN
                       INTEGER,
                         NUMBER,
```

```
lang_context IN OUT INTEGER,
  warning OUT INTEGER
);
```

- dest blob is the CLOB/NCLOB into which the data is to be written.
- *src bfile* is the pointer to the file from which the data is to be read.
- amount is the maximum number of characters to read from src bfile.
- dest_offset is the offset in characters in dest_lob to start writing data (the offset starts at 1).
- src_offset is the offset in characters in src_bfile to start reading data (the offset starts at 1).
- src_csid is the character set of src_bfile (you should typically use DBMS_LOB.DEFAULT CSID, which is the default character set for the database).
- lang_context is the language context to use for the load (you should typically use DBMS_LOB.DEFAULT_LANG_CTX, which is the default language context for the database).
- warning is a warning message that contains information if there was a problem with the load; a common problem is that a character in src_bfile cannot be converted to a character in dest_lob (in which case, warning is set to DBMS_LOB.WARN_ INCONVERTIBLE CHAR).



NOTE

You can learn all about character sets, contexts, and how to convert characters from one language to another in the Oracle Database Globalization Support Guide published by Oracle Corporation.

The following table shows the exceptions thrown by LOADCLOBFROMFILE().

Exception

Thrown When

VALUE_ERROR
INVALID ARGVAL

Any of the input parameters are null or invalid.

Either:

- \blacksquare src offset < 1
- \blacksquare dest_offset < 1
- src offset > LOBMAXSIZE
- dest offset > LOBMAXSIZE
- \blacksquare amount < 1
- amount > LOBMAXSIZE

OPEN()

OPEN () opens a LOB. There are three versions of OPEN ():

```
DBMS LOB.OPEN(
   lob IN OUT NOCOPY BLOB,
    open_mode IN BINARY_INTEGER
   );
   DBMS LOB.OPEN (
    10b IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
    open_mode IN BINARY_INTEGER
   );
   DBMS LOB.OPEN (
    bfile IN OUT NOCOPY BFILE,
    );
```

- 10b is the LOB to open.
- *bfile* is the pointer to the file to open.
- open mode indicates the open mode; the default is DBMS LOB.FILE READONLY, which indicates the LOB may only be read from; DBMS LOB.FILE READWRITE indicates the LOB may be read from and written to.

The following table shows the exception thrown by OPEN().

Exception

Thrown When

VALUE ERROR

Any of the input parameters are null or invalid.

READ()

READ() reads data into a buffer from a LOB. There are three versions of READ():

```
DBMS_LOB.READ(

lob IN BLOB,
      amount IN OUT NOCOPY BINARY INTEGER,
     offset IN INTEGER, buffer OUT RAW
   );
   amount IN OUT NOCOPY BINARY INTEGER,
     offset IN INTEGER,
buffer OUT VARCHAR2 CHARACTER SET lob%CHARSET
   );
   DBMS LOB.READ(
     bfile IN BFILE,
      amount IN OUT NOCOPY BINARY INTEGER,
```

- 10b is the CLOB, NCLOB, or BLOB to read from.
- bfile is the BFILE to read from.
- amount is the maximum number of characters to read from a CLOB/NCLOB, or the maximum number of bytes to read from a BLOB/BFILE.
- offset is the offset to start reading (the offset starts at 1).
- buffer is the variable where the data read from the LOB is to be stored.

The following table shows the exceptions thrown by READ().

Exception	Ihrown When
VALUE_ERROR	Any of the input parameters are null.
INVALID_ARGVAL	Either:
	\blacksquare amount < 1
	■ amount > MAXBUFSIZE
	amount > capacity of buffer in bytes or characters
	■ offset < 1
	■ offset > LOBMAXSIZE
NO_DATA_FOUND	The end of the LOB was reached and there are no more bytes or characters to read from the LOB.

SUBSTR()

SUBSTR() returns part of the LOB data, starting at the offset for a total amount of characters or bytes. There are three versions of SUBSTR():

```
DBMS_LOB.SUBSTR(

lob IN BLOB,

amount IN INTEGER := 32767,

offset IN INTEGER := 1

) RETURN RAW;

DBMS_LOB.SUBSTR (

lob IN CLOB/NCLOB CHARACTER SET ANY_CS,

amount IN INTEGER := 32767,

offset IN INTEGER := 1

) RETURN VARCHAR2 CHARACTER SET lob%CHARSET;

DBMS_LOB.SUBSTR (

bfile IN BFILE,

amount IN INTEGER := 32767,

offset IN INTEGER := 32767,

offset IN INTEGER := 1

) RETURN RAW;
```

- 10b is the BLOB, CLOB, or NCLOB to read from.
- bfile is the pointer to the file to read from.
- amount is the maximum number of characters read from a CLOB/NCLOB, or the maximum number of bytes to read from a BLOB/BFILE.
- offset is the offset to start reading data from the LOB (the offset starts at 1).

SUBSTR() returns

- RAW data when reading from a BLOB/BFILE.
- VARCHAR2 data when reading from a CLOB/NCLOB.
- Null if
 - \blacksquare amount < 1
 - amount > 32767
 - \blacksquare offset < 1
 - offset > LOBMAXSIZE

The following table shows the exceptions thrown by SUBSTR().

Exception	Thrown When
VALUE_ERROR	Any of the input parameters are null or invalid.
UNOPENED_FILE	The BFILE isn't open.
NOEXIST_DIRECTORY	The directory doesn't exist.
NOPRIV_DIRECTORY	You don't have privileges on the directory.
INVALID_DIRECTORY	The directory is invalid.
INVALID_OPERATION	The file exists, but you don't have privileges to access the file.

TRIM()

TRIM() trims the LOB data to the specified shorter length. There are two versions of TRIM():

```
DBMS LOB.TRIM(
     10b IN OUT NOCOPY BLOB,
     newlen IN INTEGER
   );
   DBMS LOB.TRIM(
     10b IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
     newlen IN INTEGER
   );
```

where

- *lob* is the BLOB, CLOB, or NCLOB to trim.
- newlen is the new length (in bytes for a BLOB, or characters for a CLOB/NCLOB).

The following table shows the exceptions thrown by TRIM().

Exception

Thrown When

```
VALUE_ERROR The lob parameter is null.

INVALID_ARGVAL Either:

■ newlen < 0
■ newlen > LOBMAXSIZE
```

WRITE()

WRITE() writes data from a buffer to a LOB. There are two versions of WRITE():

```
DBMS_LOB.WRITE(

lob IN OUT NOCOPY BLOB,

amount IN BINARY_INTEGER,

offset IN INTEGER,

buffer IN RAW

);

DBMS_LOB.WRITE(

lob IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY_CS,

amount IN BINARY_INTEGER,

offset IN INTEGER,

buffer IN VARCHAR2 CHARACTER SET lob%CHARSET

);
```

where

- 10b is the LOB to write to.
- amount is the maximum number of characters to write to a CLOB/NCLOB, or the maximum number of bytes to write to a BLOB.
- offset is the offset to start writing data to the LOB (the offset starts at 1).
- buffer is the variable that contains the data to be written to the LOB.

The following table shows the exceptions thrown by WRITE().

Exception

Thrown When

VALUE_ERROR Any of the input parameters are null or invalid.

INVALID_ARGVAL

Either:

■ amount < 1
■ amount > MAXBUFSIZE

 \blacksquare offset < 1

■ offset > LOBMAXSIZE

WRITEAPPEND()

WRITEAPPEND() writes data from the buffer to the end of a LOB, starting at the offset for a total amount of characters or bytes. There are two versions of WRITEAPPEND():

```
DBMS LOB.WRITEAPPEND (
     lob IN OUT NOCOPY BLOB,
     amount IN BINARY_INTEGER,
     buffer IN RAW
   );
    DBMS LOB.WRITEAPPEND (
     10b IN OUT NOCOPY CLOB/NCLOB CHARACTER SET ANY CS,
     amount IN BINARY_INTEGER,
buffer IN VARCHAR2 CHARACTER SET lob%CHARSET
   );
```

where

- 10b is the BLOB, CLOB, or NCLOB to write to.
- amount is the maximum number of characters to write to a CLOB/NCLOB, or the maximum number of bytes to write to a BLOB.
- buffer is the variable that contains the data to be written to the LOB.

The following table shows the exceptions thrown by WRITEAPPEND ().

Exception

Thrown When

VALUE ERROR Any of the input parameters are null or invalid.

INVALID ARGVAL

 \blacksquare amount < 1

■ amount > MAXBUFSIZE

Example PL/SQL Procedures

In this section, you'll see example PL/SQL procedures that use the various methods described in the previous sections. The example procedures are created when you run the lob schema.sql script.

Retrieving a LOB Locator

The following get clob locator() procedure gets a LOB locator from the clob content table; get clob locator() performs the following tasks:

- Accepts an IN OUT parameter named p clob of type CLOB; p clob is set to a LOB locator inside the procedure. Because p_clob is IN OUT, the value is passed out of the procedure.
- Accepts an IN parameter named p id of type INTEGER, which specifies the id of a row to retrieve from the clob content table.
- Selects clob column from the clob content table into p clob; this stores the LOB locator of clob column in p clob.

```
CREATE PROCEDURE get clob locator(
      p clob IN OUT CLOB,
      p id IN INTEGER
    ) AS
    BEGIN
      -- get the LOB locator and store it in p clob
      SELECT clob column
      INTO p clob
      FROM clob content
      WHERE id = p id;
    END get clob locator;
```

The following get blob locator () procedure does the same thing as the previous procedure, except it gets the locator for a BLOB from the blob content table:

```
CREATE PROCEDURE get blob locator(
    p blob IN OUT BLOB,
    p id IN INTEGER
  ) AS
  BEGIN
    -- get the LOB locator and store it in p blob
    SELECT blob column
    INTO p blob
    FROM blob content
    WHERE id = p_id;
  END get blob locator;
```

These two procedures are used in the code shown in the following sections.

Reading Data from CLOBs and BLOBs

The following read clob example () procedure reads the data from a CLOB and displays the data on the screen; read clob example () performs the following tasks:

- Calls get clob locator () to get a locator and stores it in v clob
- Uses READ() to read the contents of v clob into a VARCHAR2 variable named v char buffer
- Displays the contents of v char buffer on the screen

```
CREATE PROCEDURE read clob example(
      p id IN INTEGER
    ) AS
     v clob CLOB;
      v offset INTEGER := 1;
      v amount INTEGER := 50;
      v char buffer VARCHAR2 (50);
    BEGIN
      -- get the LOB locator and store it in v clob
      get clob locator(v clob, p id);
```

```
-- read the contents of v clob into v char buffer, starting at
 -- the v offset position and read a total of v amount characters
 DBMS LOB.READ(v clob, v amount, v offset, v char buffer);
 -- display the contents of v char buffer
 DBMS OUTPUT.PUT LINE('v char buffer = ' || v char buffer);
  DBMS OUTPUT.PUT LINE('v amount = ' || v amount);
END read clob example;
```

The following example turns the server output on and calls read clob example():

SET SERVEROUTPUT ON

```
CALL read clob example(1);
v char buffer = Creeps in this petty pace
v = 25
```

The following read blob example () procedure reads the data from a BLOB; read blob example () performs the following tasks:

- Calls get blob locator () to get the locator and stores it in v blob
- Calls READ() to read the contents of v blob into a RAW variable named v binary buffer
- Displays the contents of v binary buffer on the screen

```
CREATE PROCEDURE read blob example (
     p id IN INTEGER
   ) AS
     v blob BLOB;
     v offset INTEGER := 1;
     v amount INTEGER := 25;
     v binary buffer RAW(25);
   BEGIN
     -- get the LOB locator and store it in v blob
     get blob locator (v blob, p id);
     -- read the contents of v blob into v binary buffer, starting at
     -- the v offset position and read a total of v amount bytes
     DBMS LOB.READ(v blob, v amount, v offset, v binary buffer);
     -- display the contents of v binary buffer
     DBMS OUTPUT.PUT_LINE('v_binary_buffer = ' || v_binary_buffer);
     DBMS OUTPUT.PUT LINE('v amount = ' || v amount);
   END read blob example;
```

The following example calls read blob example ():

```
CALL read_blob_example(1);
   v binary buffer = 100111010101011111
   v = 9
```

Writing to a CLOB

The following write_example() procedure writes a string in v_char_buffer to v_clob using WRITE(); notice that the SELECT statement in the procedure uses the FOR UPDATE clause, which is used because the CLOB is written to using WRITE():

```
CREATE PROCEDURE write example (
       p id IN INTEGER
     ) AS
       v clob CLOB;
       v offset INTEGER := 7;
       v amount INTEGER := 6;
       v char buffer VARCHAR2(10) := 'pretty';
     BEGIN
       -- get the LOB locator into v clob for update (for update
       -- because the LOB is written to using WRITE() later)
       SELECT clob column
       INTO v clob
       FROM clob content
       WHERE id = p id
       FOR UPDATE;
       -- read and display the contents of the CLOB
       read clob example(p id);
       -- write the characters in v char buffer to v clob, starting
       -- at the v offset position and write a total of v amount characters
       DBMS LOB.WRITE(v clob, v amount, v offset, v char buffer);
       -- read and display the contents of the CLOB
       -- and then rollback the write
       read clob example (p id);
       ROLLBACK;
     END write example;
```

The following example calls write example ():

CALL write_example(1);

```
v_char_buffer = Creeps in this petty pace
v_amount = 25
v_char_buffer = Creepsprettyis petty pace
v amount = 25
```

Appending Data to a CLOB

The following append_example() procedure uses APPEND() to copy the data from v_src_clob to the end of v_dest_clob:

```
CREATE PROCEDURE append_example AS

v_src_clob CLOB;

v_dest_clob CLOB;

BEGIN

-- get the LOB locator for the CLOB in row #2 of
```

```
-- the clob content table into v src clob
 get clob locator(v src clob, 2);
 -- get the LOB locator for the CLOB in row #1 of
 -- the clob content table into v dest clob for update
 -- (for update because the CLOB will be added to using
 -- APPEND() later)
 SELECT clob column
 INTO v dest clob
 FROM clob content
 WHERE id = 1
 FOR UPDATE;
 -- read and display the contents of CLOB #1
 read clob example(1);
 -- use APPEND() to copy the contents of v src clob to v dest clob
 DBMS LOB.APPEND(v dest clob, v_src_clob);
 -- read and display the contents of CLOB #1
 -- and then rollback the change
 read clob example(1);
 ROLLBACK;
END append example;
```

The following example calls append example ():

CALL append example();

```
v char buffer = Creeps in this petty pace
v amount = 25
v char buffer = Creeps in this petty pace from day to day
v amount = 41
```

Comparing the Data in Two CLOBs

The following compare example () procedure compares the data in v clob1 and v clob2 using COMPARE():

```
CREATE PROCEDURE compare example AS
     v clob1 CLOB;
     v clob2 CLOB;
     v return INTEGER;
   BEGIN
     -- get the LOB locators
     get clob_locator(v_clob1, 1);
     get clob locator(v clob2, 2);
     -- compare v clob1 with v clob2 (COMPARE() returns 1
     -- because the contents of v clob1 and v clob2 are different)
     DBMS OUTPUT.PUT LINE('Comparing v clob1 with v clob2');
     v return := DBMS LOB.COMPARE(v clob1, v clob2);
     DBMS OUTPUT.PUT_LINE('v_return = ' || v_return);
```

```
-- compare v clob1 with v clob1 (COMPARE() returns 0
 -- because the contents are the same)
 DBMS OUTPUT.PUT LINE('Comparing v clob1 with v clob1');
 v return := DBMS LOB.COMPARE(v clob1, v clob1);
 DBMS OUTPUT.PUT LINE('v return = ' || v return);
END compare example;
```

The following example calls compare example ():

CALL compare example();

```
Comparing v clob1 with v clob2
v return = 1
Comparing v clob1 with v clob1
v return = 0
```

Notice that v return is 1 when comparing v clob1 with v clob2, which indicates the LOB data is different; v return is 0 when comparing v clob1 with v clob1, which indicates the LOB data is the same.

Copying Data from One CLOB to Another

The following copy example () procedure copies some characters from v src clob to v dest_clob using COPY():

```
CREATE PROCEDURE copy example AS
      v src clob CLOB;
      v dest clob CLOB;
      v src offset INTEGER := 1;
      v dest offset INTEGER := 7;
      v amount INTEGER := 5;
    BEGIN
      -- get the LOB locator for the CLOB in row #2 of
      -- the clob content table into v dest clob
      get clob locator(v src clob, 2);
      -- get the LOB locator for the CLOB in row #1 of
      -- the clob content table into v dest clob for update
      -- (for update because the CLOB will be added to using
      -- COPY() later)
      SELECT clob column
      INTO v dest clob
      FROM clob content
      WHERE id = 1
      FOR UPDATE;
      -- read and display the contents of CLOB #1
      read clob example(1);
      -- copy characters to v_dest_clob from v_src_clob using COPY(),
      -- starting at the offsets specified by v dest offset and
      -- v src offset for a total of v amount characters
```

```
DBMS LOB.COPY(
   v dest clob, v src clob,
    v amount, v dest offset, v src offset
 -- read and display the contents of CLOB #1
 -- and then rollback the change
 read clob example(1);
 ROLLBACK;
END copy example;
```

The following example calls copy example ():

CALL copy example();

```
v char buffer = Creeps in this petty pace
v = 25
v char buffer = Creeps fromhis petty pace
v amount = 25
```

Using Temporary CLOBs

The following temporary lob example () procedure illustrates the use of a temporary CLOB:

```
CREATE PROCEDURE temporary lob example AS
     v clob CLOB;
     v amount INTEGER;
     v offset INTEGER := 1;
     v char buffer VARCHAR2(17) := 'Juliet is the sun';
     -- use CREATETEMPORARY() to create a temporary CLOB named v clob
     DBMS LOB.CREATETEMPORARY (v clob, TRUE);
     -- use WRITE() to write the contents of v char buffer to v clob
     v amount := LENGTH(v char buffer);
     DBMS LOB.WRITE(v clob, v amount, v offset, v char buffer);
     -- use ISTEMPORARY() to check if v clob is temporary
     IF (DBMS LOB.ISTEMPORARY(v clob) = 1) THEN
       DBMS OUTPUT.PUT LINE('v clob is temporary');
     END IF;
     -- use READ() to read the contents of v clob into v char buffer
     DBMS LOB.READ (
       v clob, v amount, v offset, v char buffer
     );
     DBMS OUTPUT.PUT LINE('v char buffer = ' || v char buffer);
     -- use FREETEMPORARY() to free v clob
     DBMS LOB.FREETEMPORARY (v clob);
   END temporary lob example;
```

The following example calls temporary lob example():

```
CALL temporary_lob_example();
v_clob is temporary
v char buffer = Juliet is the sun
```

Erasing Data from a CLOB

The following erase example () procedure erases part of a CLOB using ERASE ():

```
CREATE PROCEDURE erase example IS
       v clob CLOB;
       v offset INTEGER := 2;
      v amount INTEGER := 5;
     BEGIN
       -- get the LOB locator for the CLOB in row #1 of
       -- the clob content table into v dest clob for update
       -- (for update because the CLOB will be erased using
       -- ERASE() later)
       SELECT clob column
       INTO v clob
       FROM clob content
       WHERE id = 1
       FOR UPDATE;
       -- read and display the contents of CLOB #1
       read clob example(1);
       -- use ERASE() to erase a total of v amount characters
       -- from v clob, starting at v offset
       DBMS LOB.ERASE(v clob, v amount, v offset);
       -- read and display the contents of CLOB #1
       -- and then rollback the change
       read clob example(1);
       ROLLBACK;
     END erase example;
```

The following example calls erase example():

CALL erase_example();

```
v_char_buffer = Creeps in this petty pace
v_amount = 25
v_char_buffer = C in this petty pace
v_amount = 25
```

Searching the Data in a CLOB

The following instr_example() procedure uses INSTR() to search the character data stored in a CLOB:

```
CREATE PROCEDURE instr example AS
       v clob CLOB;
        v char buffer VARCHAR2(50) := 'It is the east and Juliet is the sun';
       v pattern VARCHAR2(5);
       v offset INTEGER := 1;
       v amount INTEGER;
        v occurrence INTEGER;
       v return INTEGER;
      BEGIN
       -- use CREATETEMPORARY() to create a temporary CLOB named v clob
        DBMS LOB.CREATETEMPORARY (v clob, TRUE);
        -- use WRITE() to write the contents of v char buffer to v clob
        v amount := LENGTH(v char buffer);
        DBMS LOB.WRITE(v clob, v amount, v offset, v char buffer);
        -- use READ() to read the contents of v clob into v char buffer
        DBMS LOB.READ(v clob, v amount, v offset, v char buffer);
        DBMS OUTPUT.PUT LINE('v char buffer = ' || v char buffer);
        -- use INSTR() to search v clob for the second occurrence of is,
        -- and INSTR() returns 27
        DBMS OUTPUT.PUT LINE('Searching for second ''is''');
        v pattern := 'is';
        v occurrence := 2;
        v_return := DBMS_LOB.INSTR(v clob, v pattern, v offset, v occurrence);
        DBMS OUTPUT.PUT LINE('v return = ' || v return);
        -- use INSTR() to search v clob for the first occurrence of Moon,
        -- and INSTR() returns 0 because Moon doesn't appear in v clob
        DBMS OUTPUT.PUT LINE('Searching for ''Moon''');
        v pattern := 'Moon';
        v occurrence := 1;
        v return := DBMS LOB.INSTR(v clob, v pattern, v offset, v occurrence);
        DBMS OUTPUT.PUT LINE('v return = ' || v return);
        -- use FREETEMPORARY() to free v_clob
        DBMS LOB.FREETEMPORARY(v clob);
      END instr example;
         The following example calls instr example ():
  CALL instr example();
      v char buffer = It is the east and Juliet is the sun
      Searching for second 'is'
      v return = 27
      Searching for 'Moon'
      v return = 0
```

Copying Data from a File into a CLOB and a BLOB

The following copy file data to clob() procedure shows how to read text from a file and store it in a CLOB:

```
CREATE PROCEDURE copy file data to clob(
      p clob id INTEGER,
      p directory VARCHAR2,
      p file name VARCHAR2
      v file UTL FILE.FILE TYPE;
      v chars read INTEGER;
      v dest clob CLOB;
      v amount INTEGER := 32767;
      v char buffer VARCHAR2 (32767);
    BEGIN
      -- insert an empty CLOB
      INSERT INTO clob content(
        id, clob column
      ) VALUES (
        p clob id, EMPTY CLOB()
      );
      -- get the LOB locator of the CLOB
      SELECT clob column
      INTO v dest clob
      FROM clob content
      WHERE id = p clob id
      FOR UPDATE:
      -- open the file for reading of text (up to v amount characters per line)
      v file := UTL FILE.FOPEN(p directory, p file name, 'r', v amount);
      -- copy the data from the file into v dest clob one line at a time
      LOOP
        BEGIN
          -- read a line from the file into v char buffer;
          -- GET LINE() does not copy the newline character into
          -- v char buffer
          UTL FILE.GET LINE(v file, v char buffer);
          v chars read := LENGTH(v char buffer);
          -- append the line to v dest clob
          DBMS LOB.WRITEAPPEND(v dest clob, v chars read, v char buffer);
          -- append a newline to v dest clob because v char buffer;
          -- the ASCII value for newline is 10, so CHR(10) returns newline
          DBMS LOB.WRITEAPPEND(v dest clob, 1, CHR(10));
        EXCEPTION
          -- when there is no more data in the file then exit
```

```
WHEN NO DATA FOUND THEN
       EXIT;
   END;
 END LOOP;
 -- close the file
 UTL FILE.FCLOSE(v file);
 DBMS OUTPUT.PUT LINE('Copy successfully completed.');
END copy file data to clob;
```

There are a number of things to note about this procedure:

- UTL FILE is a package included with the database and contains methods and types that enable you to read and write files. For example, UTL FILE.FILE TYPE is an object type used to represent a file.
- The v amount variable is set to 32767, which is the maximum number of characters that can be read from a file during each read operation.
- The v char buffer variable is used to store the results read from the file before they are appended to v dest clob. The maximum length of v char buffer is set to 32767; this length is large enough to store the maximum number of characters read from the file during each read operation.
- UTL FILE.FOPEN(directory, file name, open mode, amount) opens a file; open mode can be set to one of the following modes:
 - r to read text
 - w to write text
 - a to append text
 - rb to read bytes
 - wb to write bytes
 - ab to append bytes
- UTL FILE.GET LINE(v file, v char buffer) gets a line of text from v file into v char buffer. GET LINE() does not add the newline to v char buffer; because I want the newline, I add it using DBMS LOB.WRITEAPPEND(v dest clob, 1, CHR(10)).

The following example calls copy file data to clob() to copy the contents of the file textContent.txt to a new CLOB with an id of 3:

```
CALL copy file data to clob(3, 'SAMPLE FILES DIR', 'textContent.txt');
   Copy successfully completed.
```

The following <code>copy_file_data_to_blob()</code> procedure shows how to read binary data from a file and store it in a <code>BLOB</code>; notice that a <code>RAW</code> array is used to store the binary data read from the file:

```
CREATE PROCEDURE copy file data to blob(
       p blob id INTEGER,
       p directory VARCHAR2,
      p file name VARCHAR2
      v file UTL FILE.FILE TYPE;
       v bytes read INTEGER;
       v dest blob BLOB;
       v amount INTEGER := 32767;
       v binary buffer RAW(32767);
       -- insert an empty BLOB
       INSERT INTO blob content (
        id, blob column
       ) VALUES (
        p blob id, EMPTY BLOB()
       );
       -- get the LOB locator of the BLOB
       SELECT blob column
       INTO v dest blob
       FROM blob content
       WHERE id = p blob id
       FOR UPDATE;
       -- open the file for reading of bytes (up to v amount bytes at a time)
       v file := UTL FILE.FOPEN(p directory, p file name, 'rb', v amount);
       -- copy the data from the file into v dest blob
       LOOP
         BEGIN
          -- read binary data from the file into v binary buffer
           UTL FILE.GET RAW(v file, v binary buffer, v amount);
           v bytes read := LENGTH(v binary buffer);
           -- append v binary buffer to v dest blob
           DBMS LOB.WRITEAPPEND(v dest blob, v bytes read/2,
             v binary buffer);
         EXCEPTION
           -- when there is no more data in the file then exit
           WHEN NO DATA FOUND THEN
             EXIT;
         END;
       END LOOP;
       -- close the file
       UTL FILE.FCLOSE (v file);
```

```
DBMS OUTPUT.PUT LINE('Copy successfully completed.');
END copy file data to blob;
```

The following example calls copy file data to blob() to copy the contents of the file binaryContent.doc to a new BLOB with an id of 3:

```
CALL copy file data to blob(3, 'SAMPLE FILES DIR', 'binaryContent.doc');
    Copy successfully completed.
```

Of course, copy file data to blob() can be used to write any binary data contained in a file to a BLOB. The binary data can contain music, video, images, executables, and so on. Go ahead and try this using your own files.



TIP

You can also bulk-load data into a LOB using the Oracle SQL*Loader and Data Pump utilities; see the Oracle Database Large Objects Developer's Guide published by Oracle Corporation for details.

Copying Data from a CLOB and a BLOB to a File

The following copy clob data to file() procedure shows how to read text from a CLOB and save it to a file:

```
CREATE PROCEDURE copy clob data to file(
     p clob id INTEGER,
     p directory VARCHAR2,
     p file name VARCHAR2
    ) AS
     v src clob CLOB;
     v file UTL FILE.FILE TYPE;
     v offset INTEGER := 1;
     v amount INTEGER := 32767;
     v char buffer VARCHAR2 (32767);
   BEGIN
     -- get the LOB locator of the CLOB
     SELECT clob column
     INTO v src clob
     FROM clob content
     WHERE id = p clob id;
     -- open the file for writing of text (up to v amount characters at a time)
     v file := UTL FILE.FOPEN(p directory, p file name, 'w', v amount);
     -- copy the data from v src clob to the file
     LOOP
       BEGIN
         -- read characters from v src clob into v char buffer
         DBMS_LOB.READ(v_src_clob, v_amount, v_offset, v_char_buffer);
          -- copy the characters from v char buffer to the file
          UTL FILE.PUT(v file, v char buffer);
```

```
-- add v amount to v offset
     v offset := v offset + v amount;
   EXCEPTION
      -- when there is no more data in the file then exit
     WHEN NO DATA FOUND THEN
        EXIT:
   END:
 END LOOP;
 -- flush any remaining data to the file
 UTL FILE.FFLUSH(v file);
 -- close the file
 UTL FILE.FCLOSE(v file);
 DBMS OUTPUT.PUT LINE('Copy successfully completed.');
END copy clob data to file;
```

The following example calls copy clob data to file () to copy the contents of CLOB #3 to a new file named textContent2.txt:

```
CALL copy_clob_data_to_file(3, 'SAMPLE_FILES_DIR', 'textContent2.txt');
    Copy successfully completed.
```

If you look in the C:\sample files directory, you will find the new textContent2.txt file. This file contains identical text to textContent.txt.

The following copy blob data to file() procedure shows how to read binary data from a BLOB and save it to a file:

```
CREATE PROCEDURE copy blob data to file(
     p blob id INTEGER,
     p directory VARCHAR2,
     p file name VARCHAR2
    ) AS
     v src blob BLOB;
     v file UTL FILE.FILE TYPE;
     v offset INTEGER := 1;
     v amount INTEGER := 32767;
     v binary buffer RAW(32767);
    BEGIN
      -- get the LOB locator of the BLOB
      SELECT blob column
      INTO v src blob
      FROM blob content
     WHERE id = p blob id;
      -- open the file for writing of bytes (up to v amount bytes at a time)
     v file := UTL FILE.FOPEN(p directory, p file name, 'wb', v amount);
      -- copy the data from v src blob to the file
```

```
LOOP
    BEGIN
      -- read characters from v src blob into v binary buffer
      DBMS LOB.READ(v src blob, v amount, v offset, v binary buffer);
      -- copy the binary data from v binary buffer to the file
      UTL FILE.PUT RAW(v file, v binary buffer);
      -- add v amount to v offset
      v offset := v offset + v amount;
    EXCEPTION
      -- when there is no more data in the file then exit
      WHEN NO DATA FOUND THEN
       EXIT:
    END;
 END LOOP:
  -- flush any remaining data to the file
 UTL FILE.FFLUSH(v file);
 -- close the file
 UTL FILE.FCLOSE(v file);
 DBMS OUTPUT.PUT LINE('Copy successfully completed.');
END copy blob data to file;
```

The following example calls copy blob data to file () to copy the contents of BLOB #3 to a new file named binaryContent2.doc:

CALL copy blob data to file(3, 'SAMPLE FILES DIR', 'binaryContent2.doc'); Copy successfully completed.

If you look in the C:\sample files directory, you will find the new binaryContent2 . doc file. This file contains identical text to binaryContent.doc.

Of course, copy blob data to file () can be used to write any binary data contained in a BLOB to a file. The binary data can contain music, video, images, executables, and so on.

Copying Data from a BFILE to a CLOB and a BLOB

The following copy bfile data to clob() procedure shows how to read text from a BFILE and save it to a CLOB:

```
CREATE PROCEDURE copy bfile data to clob(
      p bfile id INTEGER,
      p clob id INTEGER
    ) AS
      v src bfile BFILE;
      v directory VARCHAR2 (200);
      v filename VARCHAR2(200);
      v length INTEGER;
      v dest clob CLOB;
```

```
v amount INTEGER := DBMS LOB.LOBMAXSIZE;
 v dest offset INTEGER := 1;
 v src offset INTEGER := 1;
 v src csid INTEGER := DBMS LOB.DEFAULT CSID;
 v lang context INTEGER := DBMS LOB.DEFAULT LANG CTX;
 v warning INTEGER;
BEGIN
 -- get the locator of the BFILE
 SELECT bfile column
 INTO v src bfile
  FROM bfile content
 WHERE id = p bfile id;
  -- use FILEEXISTS() to check if the file exists
  -- (FILEEXISTS() returns 1 if the file exists)
  IF (DBMS LOB.FILEEXISTS(v src bfile) = 1) THEN
   -- use OPEN() to open the file
    DBMS LOB.OPEN(v src bfile);
    -- use FILEGETNAME() to get the name of the file and the directory
    DBMS LOB.FILEGETNAME(v src bfile, v directory, v filename);
    DBMS_OUTPUT.PUT_LINE('Directory = ' || v directory);
    DBMS OUTPUT.PUT LINE('Filename = ' | | v filename);
    -- insert an empty CLOB
    INSERT INTO clob content (
     id, clob column
    ) VALUES (
     p clob id, EMPTY CLOB()
    -- get the LOB locator of the CLOB (for update)
    SELECT clob column
    INTO v dest clob
    FROM clob content
    WHERE id = p \ clob \ id
    FOR UPDATE;
    -- use LOADCLOBFROMFILE() to get up to v_amount characters
    -- from v src bfile and store them in v dest clob, starting
    -- at offset 1 in v src bfile and v dest clob
    DBMS LOB.LOADCLOBFROMFILE (
     v dest clob, v src bfile,
     v amount, v dest offset, v src offset,
     v src csid, v lang context, v warning
    );
    -- check v warning for an inconvertible character
    IF (v warning = DBMS LOB.WARN INCONVERTIBLE CHAR) THEN
     DBMS OUTPUT.PUT LINE('Warning! Inconvertible character.');
    END IF:
```

```
-- use CLOSE() to close v src bfile
   DBMS LOB.CLOSE(v src bfile);
   DBMS OUTPUT.PUT LINE('Copy successfully completed.');
   DBMS OUTPUT.PUT LINE('File does not exist');
 END IF:
END copy bfile data to clob;
```

The following example calls copy bfile data to clob() to copy the contents of BFILE #1 to a new CLOB with an id of 4:

```
CALL copy bfile data to clob(1, 4);
    Copy successfully completed.
```

The next example calls copy clob data to file () to copy the contents of CLOB #4 to a new file named textContent3.txt:

```
CALL copy clob data to file(4, 'SAMPLE FILES DIR', 'textContent3.txt');
    Copy successfully completed.
```

If you look in the C:\sample files directory, you will find the new textContent3.txt file. This file contains identical text to textContent.txt.

The following copy bfile data to blob() procedure shows how to read binary data from a BFILE and save it to a BLOB:

```
CREATE PROCEDURE copy bfile data to blob(
       p bfile id INTEGER,
      p blob id INTEGER
     ) AS
      v src bfile BFILE;
       v directory VARCHAR2 (200);
       v filename VARCHAR2(200);
       v length INTEGER;
       v dest blob BLOB;
       v amount INTEGER := DBMS LOB.LOBMAXSIZE;
       v dest offset INTEGER := 1;
       v src offset INTEGER := 1;
     BEGIN
       -- get the locator of the BFILE
       SELECT bfile column
       INTO v src bfile
       FROM bfile content
       WHERE id = p bfile id;
       -- use FILEEXISTS() to check if the file exists
       -- (FILEEXISTS() returns 1 if the file exists)
       IF (DBMS LOB.FILEEXISTS(v src_bfile) = 1) THEN
        -- use OPEN() to open the file
        DBMS LOB.OPEN(v src bfile);
         -- use FILEGETNAME() to get the name of the file and
```

```
-- the directory
   DBMS LOB.FILEGETNAME(v src bfile, v directory, v filename);
   DBMS OUTPUT.PUT LINE('Directory = ' || v directory);
   DBMS OUTPUT.PUT LINE('Filename = ' || v filename);
   -- insert an empty BLOB
   INSERT INTO blob content (
     id, blob column
   ) VALUES (
     p blob id, EMPTY BLOB()
   );
   -- get the LOB locator of the BLOB (for update)
   SELECT blob column
   INTO v dest_blob
   FROM blob content
   WHERE id = p blob id
   FOR UPDATE;
   -- use LOADBLOBFROMFILE() to get up to v amount bytes
   -- from v_src_bfile and store them in v dest blob, starting
   -- at offset 1 in v src bfile and v dest blob
   DBMS LOB.LOADBLOBFROMFILE (
     v dest blob, v src bfile,
     v amount, v dest offset, v src offset
   );
   -- use CLOSE() to close v src bfile
   DBMS LOB.CLOSE(v src bfile);
   DBMS OUTPUT.PUT LINE('Copy successfully completed.');
 ELSE
   DBMS OUTPUT.PUT LINE('File does not exist');
 END IF;
END copy bfile data to blob;
```

The following example calls <code>copy_bfile_data_to_blob()</code> to copy the contents of <code>BFILE #2</code> to a new <code>BLOB</code> with an id of 4:

```
CALL copy_bfile_data_to_blob(2, 4);
Copy successfully completed.
```

The next example calls copy_blob_data_to_file() to copy the contents of BLOB #4 to a new file named binaryContent3.doc:

```
CALL copy_blob_data_to_file(4, 'SAMPLE_FILES_DIR', 'binaryContent3.doc');
Copy successfully completed.
```

If you look in the C:\sample_files directory, you will find the new binaryContent3.doc file. This file contains identical text to binaryContent.doc.

This is the end of the coverage on large objects. In the next section, you'll learn about the LONG and LONG RAW types.

LONG and LONG RAW Types

I mentioned at the start of this chapter that LOBs are the preferred storage type for large blocks of data, but you may encounter databases that still use the following types:

- **LONG** Used to store up to 2 gigabytes of character data
- **LONG RAW** Used to store up to 2 gigabytes of binary data
- **RAW** Used to store up to 4 kilobytes of binary data

In this section, you'll learn how to use LONG and LONG RAW types. RAW is used in the same way as a LONG RAW, so I've omitted coverage of RAW.

The Example Tables

In this section, you'll see the use of the following two tables:

- long content Contains a LONG column named long column
- long raw content Contains a LONG RAW column named long raw column

These two tables are created by the lob schema.sql script using the following statements:

```
CREATE TABLE long content (
                INTEGER PRIMARY KEY,
     long column LONG NOT NULL
   );
   CREATE TABLE long raw content (
            INTEGER PRIMARY KEY,
     long raw column LONG RAW NOT NULL
   );
```

Adding Data to LONG and LONG RAW Columns

The following INSERT statements add rows to the long content table:

```
INSERT INTO long content (
     id, long column
    ) VALUES (
     1, 'Creeps in this petty pace'
   INSERT INTO long content (
     id, long column
   ) VALUES (
     2, ' from day to day'
   );
```

The following INSERT statements add rows to the long raw content table (the first INSERT contains a binary number, the second a hexadecimal number):

```
INSERT INTO long raw content (
      id, long raw column
```

```
) VALUES (
    1, '100111010101011111'
);

INSERT INTO long_raw_content (
    id, long_raw_column
) VALUES (
    2, 'A0FFB71CF90DE'
);
```

In the next section, you'll see how to convert LONG and LONG RAW columns to LOBs.

Converting LONG and LONG RAW Columns to LOBs

You can convert a LONG to a CLOB using the TO_LOB() function. For example, the following statement converts long_column to a CLOB using TO_LOB() and stores the results in the clob content table:

```
INSERT INTO clob_content

SELECT 10 + id, TO_LOB(long_column)

FROM long_content;
```

You can convert a LONG RAW to a BLOB using the TO_LOB() function. For example, the following statement converts long_raw_column to a BLOB using TO_LOB() and stores the results in the blob content table:

```
INSERT INTO blob_content
SELECT 10 + id, TO_LOB(long_raw_column)
FROM long_raw_content;
2 rows created.
```

You can also use the ALTER TABLE statement to convert LONG and LONG RAW columns directly. For example, the following statement converts long column to a CLOB:

```
ALTER TABLE long_content MODIFY (long_column CLOB);
```

The next example converts long raw column to a BLOB:

```
ALTER TABLE long_raw_content MODIFY (long_raw_column BLOB);
```



CAUTION

2 rows created.

You should not modify tables that are currently used in a production application.

Once a LONG or LONG RAW column is converted to a LOB, you can use the rich PL/SQL methods described earlier to access the LOB.

Oracle Database 10*g* **Enhancements** to Large Objects

In this section, you'll learn about the following enhancements made to large objects in Oracle Database 10g:

- Implicit conversion between CLOB and NCLOB objects
- Use of the : new attribute when using LOBs in a trigger

I've provided an SQL*Plus script named lob schema2.sql in the SQL directory. This script can be run using Oracle Database 10g and higher. The script creates a user named 10b user2 with a password of lob password and creates the tables and PL/SQL code used in this section. After the script completes, you will be logged in as lob user2.

Implicit Conversion Between CLOB and NCLOB Objects

In today's global business environment, you might have to deal with conversions between Unicode and a national language character set. Unicode is a universal character set that enables you to store text that can be converted into any language; it does this by providing a unique code for every character, regardless of the language. A national character set stores text in a specific language.

In versions of the database below Oracle Database 10g, you have to explicitly convert between Unicode text and the national character set text using the TO CLOB() and TO NCLOB() functions. TO CLOB() allows you to convert text stored in a VARCHAR2, NVARCHAR2, or NCLOB to a CLOB. Similarly, TO NCLOB() allows you to convert text stored in a VARCHAR2, NVARCHAR2, or CLOB to an NCLOB.

Oracle Database 10g and higher implicitly converts Unicode text and national character set text in CLOB and NCLOB objects, which saves you from using TO CLOB() and TO NCLOB(). You can use this implicit conversion for IN and OUT variables in gueries and DML statements as well as for PL/SQL method parameters and variable assignments.

Let's take a look at an example. The following statement creates a table named nclob content that contains an NCLOB column named nclob column:

```
CREATE TABLE nclob content (
      id INTEGER PRIMARY KEY,
      nclob column NCLOB
    );
```

The following nclob example () procedure shows the implicit conversion of a CLOB to an NCLOB, and vice versa:

```
CREATE PROCEDURE nclob example
       v clob CLOB := 'It is the east and Juliet is the sun';
       v nclob NCLOB;
     BEGIN
```

```
-- insert v clob into nclob column; this implicitly
 -- converts the CLOB v clob to an NCLOB, storing
 -- the contents of v clob in the nclob content table
 INSERT INTO nclob content (
   id, nclob column
 ) VALUES (
   1, v clob
 );
 -- select nclob column into v clob; this implicitly
 -- converts the NCLOB stored in nclob column to a
 -- CLOB, retrieving the contents of nclob column
 -- into v clob
 SELECT nclob column
 INTO v clob
 FROM nclob content
 WHERE id = 1;
 -- display the contents of v clob
 DBMS OUTPUT.PUT LINE('v clob = ' || v clob);
END nclob example;
```

The following example turns the server output on and calls nclob example():

```
SET SERVEROUTPUT ON
   CALL nclob example();
   v clob = It is the east and Juliet is the sun
```

Use of the :new Attribute When Using LOBs in a Trigger

In Oracle Database 10g and higher, you can use the : new attribute when referencing LOBs in a BEFORE UPDATE or BEFORE INSERT row level trigger. The following example creates a trigger named before clob content update; the trigger fires when the clob content table is updated and displays the length of the new data in clob column; notice that :new is used to access the new data in clob column:

```
CREATE TRIGGER before clob content update
    BEFORE UPDATE
    ON clob content
    FOR EACH ROW
    BEGIN
      DBMS OUTPUT.PUT LINE('clob content changed');
      DBMS OUTPUT.PUT LINE (
        'Length = ' | DBMS LOB.GETLENGTH(:new.clob column)
    END before clob content update;
```

The following example updates the clob content table, causing the trigger to be fired:

```
UPDATE clob content
   SET clob column = 'Creeps in this petty pace'
   WHERE id = 1:
   clob content changed
   Length = 25
```

Oracle Database 11g Enhancements to Large Objects

In this section, you'll learn about the following enhancements made to large objects in Oracle Database 11g:

- Encryption of BLOB, CLOB, and NCLOB data, which prevents unauthorized viewing and modification of the data
- Compression to squeeze BLOB, CLOB, and NCLOB data
- De-duplication of BLOB, CLOB, and NCLOB data to automatically detect and remove repeated data

Encrypting LOB Data

You can disguise your data using encryption so that unauthorized users cannot view or modify it. You should encrypt sensitive data such as credit card numbers, social security numbers, and so on.

Before you can encrypt data, either you or a database administrator needs to set up a "wallet" to store security details. The data in a wallet includes a private key for encrypting and decrypting data. In this section, you'll see how to create a wallet, encrypt LOB data, and encrypt regular column data.

Creating a Wallet

To create a wallet, you must first create a directory called wallet in the directory \$ORACLE BASE\admin\\$ORACLE SID, where ORACLE BASE is the base directory where the Oracle database software is installed, and ORACLE SID is the system identifier for the database in which the wallet is to be created. For example, on my computer running Windows XP and Oracle Database 11g, I created my wallet directory in C:\oracle 11g\admin\orcl.

Once the wallet directory is created, you need to run SQL*Plus, connect to the database using a privileged user account (for example, system), and run an ALTER SYSTEM command to set the password for the wallet encryption key, as shown here:

```
SQL> CONNECT system/manager
    SOL> ALTER SYSTEM SET ENCRYPTION KEY IDENTIFIED BY "testpassword123";
    System altered.
```

Once this is done, a file called ewallet.p12 appears in the wallet directory, and the database automatically opens the wallet. The encryption key password is stored in the wallet, and is used to encrypt and decrypt data behind the scenes.

I've provided an SQL*Plus script named lob_schema3.sql in the SQL directory. This script may be run using Oracle Database 11g. The script creates a user named lob_user3 with a password of lob_password, and it also creates the tables used later in this section. After the script completes, you will be logged in as lob user3.

Encrypting LOB Data

You can encrypt the data stored in a BLOB, CLOB, or NCLOB to prevent unauthorized access to that data; you cannot encrypt a BFILE, because the file itself is stored outside the database. You can use the following algorithms to encrypt data:

- 3DES168 The Triple-DES (Data Encryption Standard) algorithm with a key length of 168 bits.
- **AES128** The Advanced Encryption Standard algorithm with a key length of 128 bits. The AES algorithms were developed to replace the older algorithms based on DES.
- AES192 The Advanced Encryption Standard algorithm with a key length of 192 bits.
- **AES256** The Advanced Encryption Standard algorithm with a key length of 256 bits. This is the most secure encryption algorithm supported by the Oracle database.

The following statement creates a table with a CLOB whose contents are to be encrypted using the AES128 algorithm; notice the use of the ENCRYPT and SECUREFILE keywords, which are required when encrypting data:

```
CREATE TABLE clob_content (
   id INTEGER PRIMARY KEY,
   clob_column CLOB ENCRYPT USING 'AES128'
) LOB(clob_column) STORE AS SECUREFILE (
   CACHE
);
```

As you can see, the contents of clob_column will be encrypted using the AES128 algorithm. If you omit the USING keyword and the algorithm, then the default AES192 algorithm is used.

The CACHE keyword in the CREATE TABLE statement indicates that the database places data from the LOB into the buffer cache for faster access. The options you can use for buffer caching are as follows:

- **CACHE READS** Use when the LOB data will be frequently read, but written only once or occasionally.
- **CACHE** Use when the LOB data will be frequently read and frequently written.
- **NOCACHE** Use when the LOB data will be read once or occasionally and written once or occasionally. This is the default option.

The following INSERT statements add two rows to the clob content table:

```
INSERT INTO clob_content (
   id, clob_column
) VALUES (
   1, TO_CLOB('Creeps in this petty pace')
```

```
);
INSERT INTO clob content (
 id, clob column
) VALUES (
 2, TO CLOB(' from day to day')
);
```

The data supplied to clob column in these statements are automatically encrypted behind the scenes by the database.

The following query retrieves the rows from the clob content table:

SELECT *

FROM clob content;

```
ΙD
-----
CLOB COLUMN
       1
Creeps in this petty pace
        2
from day to day
```

When the data is retrieved, it is automatically decrypted by the database and then returned to SOL*Plus.

As long as the wallet is open, you can store and retrieve encrypted data; when the wallet is closed, you cannot. Let's see what happens when the wallet is closed; the following statements connect as the system user and close the wallet:

CONNECT system/manager ALTER SYSTEM SET WALLET CLOSE;

If you now attempt to connect as lob user3 and retrieve clob column from the clob content table, you get the error ORA-28365: wallet is not open:

```
CONNECT lob user3/lob password
     SELECT clob column
     FROM clob content;
     ORA-28365: wallet is not open
```

You can still retrieve and modify the contents of unencrypted columns; for example, the following query retrieves the id column from the clob content table:

SELECT id FROM clob content;

1 2

```
ΙD
_____
```

The following statements connect as the system user and re-open the wallet:

CONNECT system/manager ALTER SYSTEM SET WALLET OPEN IDENTIFIED BY "testpassword123";

Once this is done, you can retrieve and modify the contents of clob_column from the clob_content table.

Encrypting Column Data

You can also encrypt regular column data. This feature was introduced in Oracle Database 10g Release 2. For example, the following statement creates a table named <code>credit_cards</code> with an encrypted column named <code>card number:</code>

You can use the same algorithms to encrypt a column as for a LOB: 3DES168, AES128, AES192 (the default), and AES256. Because I didn't specify an algorithm after the ENCRYPT keyword for the card number column, the default AES192 algorithm is used.

The following INSERT statements add two rows to the credit cards table:

```
INSERT INTO credit_cards (
    card_number, first_name, last_name, expiration
) VALUES (
    1234, 'Jason', 'Bond', '03-FEB-2008'
);

INSERT INTO credit_cards (
    card_number, first_name, last_name, expiration
) VALUES (
    5768, 'Steve', 'Edwards', '07-MAR-2009'
);
```

As long as the wallet is open, you can retrieve and modify the contents of the <code>card_number</code> column. If the wallet is closed, you get the error <code>ORA-28365</code>: wallet is not open. You saw examples that illustrate these concepts in the previous section, so I won't repeat similar examples here.

Accessing data in an encrypted column introduces additional overhead. The overhead for encrypting or decrypting a column is estimated by Oracle Corporation to be about 5 percent; this means a SELECT or an INSERT takes about 5 percent more time to complete. The total overhead depends on the number of encrypted columns and their frequency of access; therefore, you should only encrypt columns that contain sensitive data.



NOTE

If you are interested in learning more about wallets and database security generally, you should read the Advanced Security Administrator's Guide published by Oracle Corporation.

Compressing LOB Data

You can compress the data stored in a BLOB, CLOB, or NCLOB to reduce storage space. For example, the following statement creates a table with a CLOB whose contents are to be compressed; notice the use of the COMPRESS keyword:

```
CREATE TABLE clob content3 (
                INTEGER PRIMARY KEY,
     clob column CLOB
    ) LOB(clob column) STORE AS SECUREFILE (
     COMPRESS
      CACHE
    );
```



NOTE

Even though the table does not contain encrypted data, the SECUREFILE clause must be used.

When you add data to the LOB, it will be automatically compressed by the database; similarly, when you read data from a LOB, it will be automatically decompressed. You can use COMPRESS HIGH for maximum data compression; the default is COMPRESS MEDIUM, and the MEDIUM keyword is optional. The higher the compression, the higher the overhead when reading and writing LOB data.

Removing Duplicate LOB Data

You can configure a BLOB, CLOB, or NCLOB so that any duplicate data supplied to it is automatically removed; this process is known as de-duplicating data and can save storage space. For example, the following statement creates a table with a CLOB whose contents are to be de-duplicated; notice the use of the DEDUPLICATE LOB keywords:

```
CREATE TABLE clob content2 (
                  INTEGER PRIMARY KEY,
       clob column CLOB
     ) LOB(clob column) STORE AS SECUREFILE (
       DEDUPLICATE LOB
       CACHE
     );
```

Any duplicate data added to the LOB will be automatically removed by the database. The database uses the SHA1 secure hash algorithm to detect duplicate data.

You can learn even more about large objects in the Oracle Database Large Objects Developer's Guide published by Oracle Corporation.

Summary

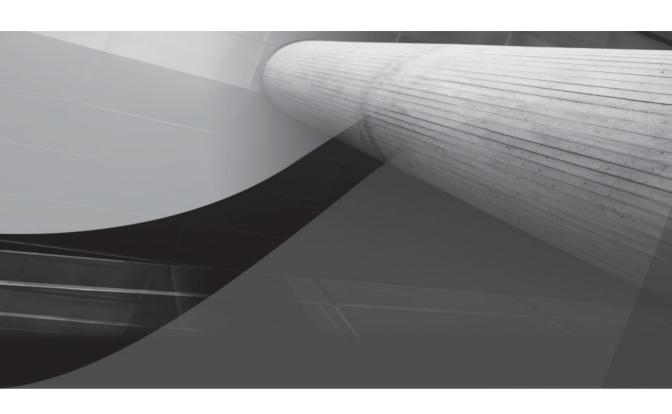
In this chapter, you have learned the following:

LOBs may be used to store binary data, character data, and references to external files. LOBs can store up to 128 terabytes of data.

530 Oracle Database 11*g* SQL

- There are four LOB types: CLOB, NCLOB, BLOB, and BFILE.
- A CLOB stores character data.
- An NCLOB stores multiple byte character data.
- A BLOB stores binary data.
- A BFILE stores a pointer to a file located in the file system.
- A LOB consists of two parts: a locator, which specifies the location of the LOB data, and the data itself.
- The DBMS LOB PL/SQL package contains methods for accessing LOBs.

In the next chapter, you'll learn how to run SQL statements from a Java program.



CHAPTER 15

Running SQL Using Java



n this chapter, you will do the following:

- Learn how to run SQL from Java programs using the Java Database Connectivity (JDBC) Application Programming Interface (API)
- Examine the various Oracle JDBC drivers that may be used to connect to an Oracle database
- Perform queries and SQL DML statements to access database tables
- Use the various Java types to get and set column values in the database
- Examine how to perform transaction control statements and SQL DDL statements
- Handle database exceptions that may occur when a Java program runs
- Examine the Oracle database software extensions to IDBC
- See complete Java programs that illustrate the use of JDBC



NOTE

This chapter gives an introduction to JDBC. For full details on using JDBC with an Oracle database, you should read my book Oracle9i JDBC Programming (McGraw-Hill/Osborne, 2002).

Getting Started

Prior to running the examples in this chapter, you'll need to install a version of Sun's Java Software Development Kit (SDK). You can download the SDK and view full installation instructions from Sun's Java website at java.sun.com.



NOTE

I used Java 1.6.0 when writing this chapter, which is installed with Java EE 5 SDK Update 2.

The directory where you installed the Oracle software on your machine is called the ORACLE_HOME directory. On my Windows computer, this directory is E:\oracle_11g\ product\11.1.0\db1. Inside ORACLE_HOME are many subdirectories, one of which is the jdbc directory. The jdbc directory contains the following:

- A text file named Readme.txt. You should open and read this file, as it contains important items such as release information and the latest installation instructions.
- A directory named lib, which contains a number of Java Archive (JAR) files.

Configuring Your Computer

Once you've downloaded and installed the required software, your next step is to configure your computer to develop and run Java programs containing JDBC statements. You must set four environment variables on your machine:

- ORACLE HOME
- JAVA HOME
- PATH
- CLASSPATH

If you're using Unix or Linux, you'll also need to set the additional LD LIBRARY PATH environment variable. You'll learn how to set these environment variables in the following sections.



CAUTION

The information in this section was correct at time of writing. You need to read the Readme.txt file in the ORACLE HOME\idbc directory to check the latest release notes and installation instructions.

Setting the ORACLE_HOME Environment Variable

The ORACLE_HOME subdirectory is located in the directory where you installed the Oracle software. You'll need to set an environment variable named ORACLE HOME on your machine that specifies this directory.

Setting an Environment Variable in Windows XP

To set an environment variable in Windows XP, you perform the following steps:

- 1. Open the Control Panel.
- 2. Double-click System. This displays the System Properties dialog box.
- 3. Select the Advanced tab.
- 4. Click the Environment Variables button. This displays the Environment Variables dialog box.
- 5. Click the New button in the System Variables area (the lower pane of the dialog box).
- 6. Set the variable name to ORACLE HOME and set the value to your ORACLE_HOME directory. (On my Windows XP machine, I have ORACLE HOME set to E:\oracle_11g\ product\11.1.0\db1.)

Setting an Environment Variable with Unix or Linux

To set an environment variable in Unix or Linux, you need to add lines to a special file; the file you need to modify depends on which shell you're using. If you're using the Bourne, Korn, or Bash shell, then you add lines similar to the following ones to your .profile (when using Bourne or Korn shell) or your .bash_profile (Bash shell):



NOTE

You'll need to replace the directory shown in the previous example with the correct ORACLE_HOME for your setup.

If you're using the C shell, you add the following line to your .login file:

setenv ORACLE_HOME /u01/app/oracle/product/11.1.0/db_1

Setting the JAVA_HOME Environment Variable

The JAVA_HOME environment variable specifies the directory where you installed the Java SDK. For example, if you installed the Java SDK in the E:\java\jdk directory, you create a JAVA_HOME system variable and set it to E:\java\jdk. To do this, you can use similar steps to those shown in the previous section.

Setting the PATH Environment Variable

The PATH environment variable contains a list of directories. When you enter a command using the operating system command line, the computer searches the directories in the PATH for the executable you are trying to run. You need to add the following two directories to your existing PATH:

- The bin subdirectory where you installed the Java SDK
- The BIN subdirectory of ORACLE_HOME

For example, if you installed the Java SDK in the E:\java\jdk directory, and your ORACLE_HOME is E:\oracle_11g\product\11.1.0\db1, then you add E:\java\jdk\bin; E:\oracle_11g\product\11.1.0\db1 to your PATH (notice a semicolon separates the two directories). To add the directories to the PATH in Windows XP, you can use steps similar to those shown earlier.

To add to an existing PATH in Unix or Linux, you need to modify the appropriate file for your shell. For example, if you're using the Bash shell with Linux, then you add lines to the .bash_ profile file that are similar to the following:

PATH=\$PATH:\$JAVA_HOME/bin:\$ORACLE_HOME/BIN export PATH

Notice that a colon (:) separates the directories.

Setting the CLASSPATH Environment Variable

The CLASSPATH environment variable contains a list of locations where Java class packages are found. A location can be a directory name or the name of a Zip file or JAR file containing classes. The ORACLE_HOME\jdbc\lib directory contains a number of JAR files; which ones you add to your CLASSPATH depends on what Java SDK you're using.

At time of writing, the following was correct for setting a CLASSPATH:

- If you're using JDK 1.6 (or higher), add ORACLE_HOME\jdbc\lib\ojdbc6.jar to your CLASSPATH.
- If you're using JDK 1.5, add ORACLE_HOME\jdbc\lib\ojdbc5.jar to your CLASSPATH.
- If you need National Language support, add ORACLE_HOME\jlib\orai18n.jar to your CLASSPATH.

- If you need the JTA and JNDI features, add ORACLE HOME\ilib\ita.jar and ORACLE HOME\jlib\jndi.jar to your CLASSPATH. JNDI is the Java Naming and Directory Interface. ITA is the Java Transaction API.
- You also need to add the current directory to your CLASSPATH. You do this by adding a period (.) to your CLASSPATH. That way, the classes in your current directory will be found by Java when you run your programs.

When Java 1.6 is used and the ORACLE HOME is E:\oracle 11g\product\11.1.0\db1, an example CLASSPATH for Windows XP is as follows:

```
.;E:\oracle 11g\product\11.1.0\db1\jdbc\lib\ojdbc6.jar;
    E:\oracle 11g\product\11.1.0\db1\jlib\orai18n.jar
```

If you're using Windows XP, you use the steps described earlier to create a system environment variable called CLASSPATH. If you're using Linux and Java 1.6, you should add the following lines to your .bash_profile:

```
CLASSPATH=$CLASSPATH:.:$ORACLE HOME/jdbc/lib/ojdbc6.jar:
   $ORACLE HOME/jlib/orai18n.jar
   export CLASSPATH
```

Setting the LD LIBRARY PATH Environment Variable

If you're using Unix or Linux, you'll also need to set the LD LIBRARY PATH environment variable to \$ORACLE_HOME/jdbc/lib. This directory contains shared libraries that are used by the JDBC OCI driver. You add LD LIBRARY PATH to the appropriate file; for example:

```
LD LIBRARY PATH=$ORACLE HOME/jdbc/lib
     export CLASSPATH
```

That concludes configuring your computer. You'll learn about the Oracle JDBC drivers next.

The Oracle JDBC Drivers

In this section, you'll learn about the various Oracle JDBC drivers. These drivers enable the JDBC statements in a Java program to access an Oracle database. There are four Oracle JDBC drivers:

- Thin driver
- OCI driver
- Server-side internal driver
- Server-side Thin driver

The following sections describe each of these drivers.

The Thin Driver

The Thin driver has the smallest footprint of all the drivers, meaning that it requires the least amount of system resources to run. The Thin driver is written entirely in Java. If you are writing a Java applet, you should use the Thin driver. The Thin driver may also be used in stand-alone Java

applications and may be used to access all versions of the Oracle database. The Thin driver works only with TCP/IP and requires that Oracle Net be up and running. For details on Oracle Net, you can read the *Oracle Database Net Services Administrator's Guide* published by Oracle Corporation.



NOTE

You don't have to install anything on the client computer to use the Thin driver, and therefore you can use it for applets.

The OCI Driver

The OCI driver requires more resources than the Thin driver, but it generally has better performance. The OCI driver is suitable for programs deployed on the middle tier—a web server, for example.



NOTE

The OCI driver requires that you install it on the client computer and is therefore not suitable for applets.

The OCI driver has a number of performance enhancing features, including the ability to pool database connections and prefetch rows from the database. The OCI driver works with all versions of the database and all of the supported Oracle Net protocols.

The Server-Side Internal Driver

The server-side internal driver provides direct access to the database, and it is used by the Oracle JVM to communicate with that database The Oracle JVM is a Java Virtual Machine that is integrated with the database. You can load a Java class into the database, then publish and run methods contained in that class using the Oracle JVM; the Java code runs on the database server and can access data from a single Oracle session.

The Server-Side Thin Driver

The server-side Thin driver is also used by the Oracle JVM and provides access to remote databases. Like the Thin driver, this driver is also written entirely in Java. Java code that uses the server-side Thin driver can access another session on the same database server or a remote server.

Importing the JDBC Packages

In order for your programs to use JDBC, you must import the required JDBC packages into your Java programs. There are two sets of JDBC packages:

- Standard JDBC packages from Sun Microsystems
- Extension packages from Oracle Corporation

The standard JDBC packages enable your Java programs to access the basic features of most databases, including the Oracle database, SQL Server, DB2, and MySQL. The Oracle extensions to JDBC enable your programs to access all of the Oracle-specific features as well as the Oracle-specific performance extensions. You'll learn about some of the Oracle-specific features later in this chapter.

To use JDBC in your programs you should import the standard java.sql.* packages, as shown in the following import statement:

```
import java.sql.*;
```

Of course, importing java.sql.* imports all of the standard JDBC packages. As you become proficient in JDBC, you'll find that you don't always need to import all the classes: you can just import those packages that your program actually uses.

Registering the Oracle JDBC Drivers

Before you can open a database connection, you must first register the Oracle JDBC drivers with your Java program. As mentioned earlier, the JDBC drivers enable your JDBC statements to access the database.

There are two ways to register the Oracle JDBC drivers:

- Use the forName () method of the class java.lang.Class
- Use the registerDriver() method of the JDBC DriverManager class

The following example illustrates the use of the forName () method:

```
Class.forName("oracle.jdbc.OracleDriver");
```

The second way to register the Oracle JDBC drivers is to use the registerDriver() method of the java.sql.DriverManager class, as shown in the following example:

```
DriverManager.registerDriver(new oracle.jdbc.OracleDriver());
```

Once you have registered the Oracle JDBC drivers, you can open a connection to a database.

Opening a Database Connection

Before you can issue SQL statements in your Java programs, you must open a database connection. There are two main ways to open a database connection:

- Use the getConnection() method of the DriverManager class
- Use an Oracle data source object, which must first be created and then connected to. This method uses a standardized way of setting database connection details, and an Oracle data source object may be used with the Java Naming and Directory Interface (JNDI).

I'll describe both of these ways to open a database connection in the following sections, starting with the getConnection() method of the DriverManager class.

Connecting to the Database Using getConnection()

The getConnection () method returns a JDBC Connection object, which should be stored in your program so it can be referenced later. The syntax of a call to the getConnection() method is as follows:

```
DriverManager.getConnection(URL, username, password);
```

where

- URL is the database that your program connects to, along with the JDBC driver you want to use. (See the following section, "The Database URL," for details on the URL.)
- username is the name of the database user that your program connects as.
- password is the password for the username.

The following example shows the getConnection() method being used to connect to a database:

```
Connection myConnection = DriverManager.getConnection(
   "jdbc:oracle:thin:@localhost:1521:ORCL",
   "store",
   "store_password"
);
```

In this example, the connection is made to a database running on the machine identified as localhost with an Oracle System Identifier (SID) of ORCL; the Oracle JDBC Thin driver is used. The connection is made with the username store and the password store_password. The Connection object returned by the call to getConnection() is stored in myConnection. The connection to a database is made through Oracle Net, which should be up and running when the program line is run.

The Database URL

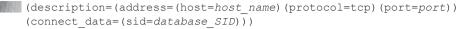
The database URL specifies the location of the database. The structure of the database URL is dependent on the vendor who provides the JDBC drivers. In the case of Oracle's JDBC drivers, the database URL structure is as follows:

driver name:@driver information

where

- driver_name is the name of the Oracle JDBC driver that your program uses. This may be set to one of the following:
 - idbc:oracle:thin The Oracle IDBC Thin driver
 - jdbc:oracle:oci The Oracle JDBC OCI driver
- driver_information The driver-specific information required to connect to the database. This is dependent on the driver being used. In the case of the Oracle JDBC Thin driver, the driver-specific information may be specified in the following format:
 - host_name:port:database_SID Where host_name is the name of the computer, port is the port to access the database, and database_SID is the database SID

For all the Oracle JDBC drivers, including the Thin driver and the various OCI drivers, the driver-specific information may also be specified using Oracle Net keyword-value pairs, which may be specified in the following format:



where

- host name is the name of the computer on which the database is running.
- port is the port number on which the Oracle Net database listener waits for requests; 1521 is the default port number. Your DBA can provide the port number.
- database SID is the Oracle SID of the database instance to which you want to connect. Your DBA can provide the database SID.

The following example shows the getConnection () method being used to connect to a database using the Oracle OCI driver, with the driver-specific information specified using Oracle Net keyword-value pairs:

```
Connection myConnection = DriverManager.getConnection(
      "jdbc:oracle:oci:@(description=(address=(host=localhost)" +
        "(protocol=tcp) (port=1521)) (connect data=(sid=ORCL)))",
      "store",
      "store password"
    );
```

As you can see, in this example a connection is made to a database running on the machine identified as localhost, with an Oracle SID of ORCL, using the Oracle OCI driver. The connection to the database is made with the username store and with a password of store password. The Connection object returned by the call to getConnection () is stored in myConnection.



NOTE

For the Oracle OCI driver, you may also use an Oracle Net TNSNAMES string. For more information on this, speak with your DBA or consult the Oracle Database Net Services Administrator's Guide published by Oracle Corporation.

Connecting to the Database Using an Oracle Data Source

You can also use an Oracle data source to connect to a database. An Oracle data source uses a more standardized way of supplying the various parameters to connect to a database than the previous method, which used DriverManager.getConnection(). In addition, an Oracle data source may also be registered with JNDI. Using JNDI with JDBC is very useful, because it allows you to register, or bind, data sources, and then look up those data sources in your program without having to provide the exact database connection details. Thus, if the database connection details change, only the JNDI object must be changed.



NOTE

You can learn about JNDI in my book Oracle9i JDBC Programming (McGraw-Hill/Osborne, 2002).

There are three steps that must be performed to use an Oracle data source:

- Create an Oracle data source object of the oracle.jdbc.pool.OracleDataSource class.
- 2. Set the Oracle data source object attributes using the set methods, which are defined in the class
- Connect to the database via the Oracle data source object using the getConnection() method.

The following sections describe these three steps.

Step 1: Create an Oracle Data Source Object

The first step is to create an Oracle data source object of the oracle.jdbc.pool .OracleDataSource class. The following example creates an OracleDataSource object named myDataSource (you may assume that the oracle.jdbc.pool.OracleDataSource class has been imported):

OracleDataSource myDataSource = new OracleDataSource();

Once you have your OracleDataSource object, the second step is to set that object's attributes using the set methods.

Step 2: Set the Oracle Data Source Object Attributes

Before you can use your OracleDataSource object to connect to a database, you must set a number of attributes in that object to indicate the connection details, using various set methods defined in the class. These details include items like the database name, the JDBC driver to use, and so on; each of these details has a corresponding attribute in an OracleDataSource object.

The oracle.jdbc.pool.OracleDataSource class actually implements the javax .sql.DataSource interface provided with JDBC. The javax.sql.DataSource interface defines a number of attributes, which are listed in Table 15-1. This table shows the name, description, and type of each attribute.

The oracle.jdbc.pool.OracleDataSource class provides an additional set of attributes, which are listed in Table 15-2.

You may use a number of methods to read from and write to each of the attributes listed in Tables 15-1 and 15-2. The methods that read from the attributes are known as *get* methods, and the methods that write to the attributes are known as *set* methods.

The set and get method names are easy to remember: you take the attribute name, convert the first letter to uppercase, and add the word "set" or "get" to the beginning. For example, to set the database name (stored in the databaseName attribute), you use the setDatabaseName() method; to get the name of the database currently set, you use the getDatabaseName() method. There is one exception to this: there is no getPassword() method (this is for security reasons—you don't want someone to be able to get your password programmatically).

Most of the attributes are Java String objects, so most of the set methods accept a single String parameter, and most of the get methods return a String. The exception to this is the portNumber attribute, which is an int. Therefore, its set method setPortNumber() accepts an int, and its get method getPortNumber() returns an int.

Attribute Name	Attribute Description	Attribute Type
databaseName	The database name (Oracle SID).	String
dataSourceName	The name of the underlying data source class.	String
description	Description of the data source.	String
networkProtocol	The network protocol to use to communicate with the database. This applies only to the Oracle JDBC OCI drivers, and defaults to tcp. For further details, read the <i>Oracle Database Net Services Administrator's Guide</i> published by Oracle Corporation.	String
password	The password for the supplied username.	String
portNumber	The port on which the Oracle Net listener waits for database connection requests. The default is 1521.	int
serverName	The database server machine name (TCP/IP address or DNS alias).	String
user	The database username.	String

TABLE 15-1 DataSource Attributes

Attribute Name	Attribute Description	Attribute Type
driverType	The JDBC driver to use. If you are using the server-side internal driver, this is set to kprb, and the other settings for the attributes are ignored.	String
url	May be used to specify an Oracle database URL, which can be used as an alternative to setting the database location. See the section earlier on database URLs for details.	String
tnsEntryName	May be used to specify an Oracle Net TNSNAMES string, which can also be used to specify the database location when using the OCI drivers.	String

 TABLE 15-2
 OracleDataSource Attributes

The following examples illustrate the use of the set methods to write to the attributes of the OracleDataSource object myDataSource that was created earlier in Step 1:

```
myDataSource.setServerName("localhost");
myDataSource.setDatabaseName("ORCL");
myDataSource.setDriverType("oci");
myDataSource.setNetworkProtocol("tcp");
myDataSource.setPortNumber(1521);
myDataSource.setUser("scott");
myDataSource.setPassword("tiger");
```

The next examples illustrate the use of some of the get methods to read the attributes previously set in myDataSource:

```
String serverName = myDataSource.getServerName();
String databaseName = myDataSource.getDatabaseName();
String driverType = myDataSource.getDriverType();
String networkProtocol = myDataSource.getNetworkProtocol();
int portNumber = myDataSource.getPortNumber();
```

Once you've set your OracleDataSource object attributes, you can use it to connect to the database.

Step 3: Connect to the Database via the Oracle Data Source Object

The third step is to connect to the database via the OracleDataSource object. You do this by calling the getConnection() method of your OracleDataSource object. The getConnection() method returns a JDBC Connection object, which must be stored.

The following example shows how to call the getConnection() method using the myDataSource object populated in the previous step:

```
Connection myConnection = myDataSource.getConnection();
```

The Connection object returned by getConnection() is stored in myConnection. You can also pass a username and password as parameters to the getConnection() method, as shown in the following example:

```
Connection myConnection = myDataSource.getConnection(
    "store", "store_password"
);
```

In this example, the username and password will override the username and password previously set in myDataSource. Therefore, the connection to the database will be made using the username of store with a password of store_password, rather than scott and tiger, which were set in myDataSource in the previous section.

Once you have your Connection object, you can use it to create a JDBC Statement object.

Creating a JDBC Statement Object

Next, you need to create a JDBC Statement object of the class java.sql.Statement. A Statement object is used to represent an SQL statement, such as a query, a DML statement (INSERT, UPDATE, or DELETE), or a DDL statement (such as CREATE TABLE). You'll learn how to issue queries, DML, and DDL statements later in this chapter.

To create a Statement object, you use the createStatement() method of a Connection object. In the following example, a Statement object named myStatement is created using the createStatement() method of the myConnection object created in the previous section:

```
Statement myStatement = myConnection.createStatement();
```

The method used in the Statement class to run the SQL statement will depend on the SQL statement you want to perform. If you want to perform a guery, you use the executeOuery() method. If you want to perform an INSERT, UPDATE, or DELETE statement, you use the executeUpdate() method. If you don't know ahead of time which type of SQL statement is to be performed, you can use the execute () method, which may be used to perform any SQL statement.

There is another IDBC class that may be used to represent an SQL statement: the PreparedStatement class. This offers more advanced functionality than the Statement class; I will discuss the PreparedStatement class after I have explained the use of the Statement class.

Once you have a Statement object, you're ready to issue SQL statements using JDBC.

Retrieving Rows from the Database

To perform a query using JDBC, you use the executeQuery () method of the Statement object, which accepts a Java String containing the text for the query.

Because a query may return more than one row, the executeQuery () method returns an object that stores the row(s) returned by your query. This object is known as a JDBC result set and is of the java.sgl.ResultSet class. When using a ResultSet object to read rows from the database, there are three steps you follow:

- 1. Create a ResultSet object and populate it with the results returned by a query.
- 2. Read the column values from the ResultSet object using get methods.
- 3. Close the ResultSet object.

I will now walk you through an example that uses a ResultSet object to retrieve the rows from the customers table.

Step 1: Create and Populate a ResultSet Object

You first create a ResultSet object and populate it with the results returned by a query. The following example creates a ResultSet object named customerResultSet and populates it with the customer id, first name, last name, dob, and phone columns from the customers table:

```
ResultSet customerResultSet = myStatement.executeQuery(
     "SELECT customer id, first name, last name, dob, phone " +
     "FROM customers"
   );
```

After this statement is executed, the ResultSet object will contain the column values for the rows retrieved by the query. The ResultSet object may then be used to access the column values for the retrieved rows. In the example, customerResultSet will contain the five rows retrieved from the customers table.

Because the execute() method accepts a Java String, you can build up your SQL statements when your program actually runs. This means that you can do some fairly powerful things in JDBC. For example, you could have the user of your program type in a string containing a WHERE clause for a query when they run the program—or even enter the whole query. The following example shows a WHERE clause string:

```
String whereClause = "WHERE customer_id = 1";
ResultSet customerResultSet2 = myStatement.executeQuery(
    "SELECT customer_id, first_name, last_name, dob, phone " +
    "FROM customers " +
    whereClause
);
```

You're not limited to queries in using this dynamic build-up method: you can build up other SQL statements in a similar manner.

Step 2: Read the Column Values from the ResultSet Object

To read the column values for the rows stored in a ResultSet object, the ResultSet class provides a series of get methods. Before I get into the details of these get methods, you need to understand how the data types used to represent values in Oracle may be mapped to compatible Java data types.

Oracle and Java Types

A Java program uses a different set of types from the Oracle database types to represent values. Fortunately, the types used by Oracle are compatible with certain Java types. This allows Java and Oracle to interchange data stored in their respective types. Table 15-3 shows one set of compatible type mappings.

From this table, you can see that an Oracle INTEGER is compatible with a Java int. (I'll talk about the other numeric types later in this chapter in the section "Handling Numbers.") This means that the customer id column of the customers table (which is defined as an Oracle INTEGER)

Oracle Type	Java Type
CHAR	String
VARCHAR2	String
DATE	<pre>java.sql.Date java.sql.Time java.sql.Timestamp</pre>
INTEGER	short int long
NUMBER	<pre>float double java.math.BigDecimal</pre>

TABLE 15-3 Compatible Type Mappings

may be stored in a Java int variable. Similarly, the first name, last name, and phone column values (VARCHAR2s) may be stored in Java String variables.

The Oracle DATE type stores a year, month, day, hour, minute, and second. You may use a java.sql.Date object to store the date part of the dob column value and a java.sql.Time variable to store the time part. You may also use a java.sql.Timestamp object to store both the date and time parts of the dob column. Later in this chapter, I'll discuss the oracle.sql . DATE type, which is an Oracle extension to the IDBC standard and provides a superior way of storing dates and times.

Getting back to the example, the customer id, first name, last name, dob, and phone columns are retrieved by the query in the previous section, and the following examples declare Java variables and objects that are compatible with those columns:

```
int customerId = 0;
    String firstName = null;
    String lastName = null;
    java.sql.Date dob = null;
    String phone = null;
```

The int and String types are part of the core Java language, while java.sql.Date is part of JDBC and is an extension of the core Java language. JDBC provides a number of such types that allow Java and a relational database to exchange data. However, JDBC doesn't have types to handle all of the types used by Oracle; one example is the ROWID type—you must use the oracle.sgl.ROWID type to store an Oracle ROWID.

To handle all of the Oracle types, Oracle provides a number of additional types, which are defined in the oracle.sql package. Also, Oracle has a number of types that may be used as an alternative to the core Java and JDBC types, and in some cases these alternatives offer more functionality and better performance than the core Java and JDBC types. I'll talk more about the Oracle types defined in the oracle.sql package later in this chapter.

Now that you understand a little bit about compatible Java and Oracle types, let's continue with the example and see how to use the get methods to read column values.

Using the Get Methods to Read Column Values

The get methods are used to read values stored in a ResultSet object. The name of each get method is simple to understand: take the name of the Java type you want the column value to be returned as and add the word "get" to the beginning. For example, use getInt() to read a column value as a Java int, and use getString() to read a column value as a Java String. To read the value as a java.sql.Date, you would use getDate(). Each get method accepts one parameter: either an int representing the position of the column in the original query or a String containing the name of the column. Let's examine some examples based on the earlier example that retrieved the columns from the customers table in the customerResultSet object.

To get the value of the customer id column, which was the first column specified in the query, you use getInt(1). You can also use the name of the column in the get method, so you could also use getInt("customer id") to get the same value.



Using the column name rather than the column position number in a get method makes your code easier to read.

To get the value of the first_name column, which was the second column specified in the query, you use getString(2) or getString("first_name"). You use similar method calls to get the last_name and phone column values because those columns are also text strings. To get the value of the dob column, you could use getDate(4) or getDate("dob"). To actually read the values stored in a ResultSet object, you must call the get methods using that ResultSet object.

Because a ResultSet object may contain more than one row, JDBC provides a method named next() that allows you to step through each row stored in a ResultSet object. You must call the next() method to access the first row in the ResultSet object, and each successive call to next() steps to the next row. When there are no more rows in the ResultSet object to read, the next() method returns the Boolean false value.

Okay, let's get back to our example: we have a ResultSet object named customerResultSet that has five rows containing the column values retrieved from the customer_id, first_name, last_name, dob, and phone columns in the customers table. The following example shows a while loop that reads the column values from customerResultSet into the customerId, firstName, lastName, dob, and phone objects created earlier:

```
while (customerResultSet.next()) {
   customerId = customerResultSet.getInt("customer_id");
   firstName = customerResultSet.getString("first_name");
   lastName = customerResultSet.getString("last_name");
   dob = customerResultSet.getDate("dob");
   phone = customerResultSet.getString("phone");

   System.out.println("customerId = " + customerId);
   System.out.println("firstName = " + firstName);
   System.out.println("lastName = " + lastName);
   System.out.println("dob = " + dob);
   System.out.println("phone = " + phone);
} // end of while loop
```

When there are no more rows to read from customerResultSet, the next() method returns false and the loop terminates. You'll notice that the example passes the name of the column to be read, rather than numeric positions, to the get methods. Also, I've copied the column values into Java variables and objects; for example, the value returned from customerResultSet.getInt("customer_id") is copied to customerId. You don't have to do that copy: you could simply use the get method call whenever you need the value. However, it is generally better if you copy it to a Java variable or object, because it will save time if you use that value more than once (the time is saved because you don't have to call the get method again).

Step 3: Close the ResultSet Object

Once you've finished with your ResultSet object, you must close that ResultSet object using the close () method. The following example closes customerResultSet:

```
customerResultSet.close();
```



NOTE

It is important that you remember to close your ResultSet object once you've finished with it. Doing so ensures that the object is scheduled for garbage collection.

Now that you've seen how to retrieve rows, I'll show you how to add rows to a database table using JDBC.

Adding Rows to the Database

You use the $\bar{S}QL$ INSERT statement to add rows to a table. There are two main ways you can perform an INSERT statement using JDBC:

- Use the executeUpdate () method defined in the Statement class.
- Use the execute () method defined in the PreparedStatement class. (I will discuss this class later in this chapter.)

The examples in this section illustrate how to add a row to the customers table. The customer id, first name, last name, dob, and phone columns for this new row will be set to 6; Jason; Price; February 22, 1969; and 800-555-1216, respectively.

To add this new row, I'll use the same Statement object declared earlier (myStatement), along with the same variables and objects that were used to retrieve the rows from the customers table in the previous section. First off, I'll set those variables and objects to the values that I want to set the database columns to in the customers table:

```
customerId = 6;
    firstName = "Jason";
    lastName = "Red";
    dob = java.sql.Date.valueOf("1969-02-22");
    phone = "800-555-1216";
```



NOTE

The java.sql.Date class stores dates using the format YYYY-MM-DD, where YYYY is the year, MM is the month number, and DD is the day number. You can also use the java.sql. Time and java.sql . Timestamp classes to represent times and dates containing times, respectively.

When you attempt to specify a date in an SQL statement, you first convert it to a format that the database can understand by using the TO DATE() built-in database function. TO DATE() accepts a string containing a date, along with the format for that date. You'll see the use of the TO DATE () function shortly in the INSERT statement example. Later in this chapter, I'll discuss the Oracle JDBC extensions, and you'll see a superior way of representing Oracle-specific dates using the oracle.sql.DATE type.

We're ready to perform an INSERT to add the new row to the customers table. The myStatement object is used to perform the INSERT statement, setting the customer id, first_name, last_name, dob, and phone column values equal to the values previously set in the customerId, firstName, lastName, dob, and phone variables.

```
myStatement.executeUpdate(
   "INSERT INTO customers " +
   "(customer_id, first_name, last_name, dob, phone) VALUES (" +
      customerId + ", '" + firstName + "', '" + lastName + "', " +
      "TO_DATE('" + dob + "', 'YYYY, MM, DD'), '" + phone + "')"
);
```

Notice the use of the TO_DATE() function to convert the contents of the dob object to an acceptable Oracle database date. Once this statement has completed, the customers table will contain the new row.

Modifying Rows in the Database

You use the SQL UPDATE statement to modify existing rows in a table. Just as with performing an INSERT statement with JDBC, you can use the executeUpdate() method defined in the Statement class or the execute() method defined in the PreparedStatement class. Use of the PreparedStatement class is covered later in this chapter.

The following example illustrates how to modify the row where the customer_id column is equal to 1:

```
first_name = "Jean";
myStatement.executeUpdate(
   "UPDATE customers " +
   "SET first_name = '" + firstName + "' " +
   "WHERE customer_id = 1"
);
```

After this statement runs, customer #1's first name will be set to "Jean".

Deleting Rows from the Database

You use the SQL DELETE statement to delete existing rows from a table. You can use the executeUpdate() method defined in the Statement class or the execute() method defined in the PreparedStatement class.

The following example illustrates how to delete customer #5 from the customers table:

```
myStatement.executeUpdate(
   "DELETE FROM customers " +
   "WHERE customer_id = 5"
);
```

After this statement runs, the row for customer #5 will have been removed from the customers table.

Handling Numbers

This section describes the issues associated with storing numbers in your lava programs. An Oracle database is capable of storing numbers with a precision of up to 38 digits. In the context of number representation, precision refers to the accuracy with which a floating-point number may be represented in a digital computer's memory. The 38 digits level of precision offered by the database allows you to store very large numbers.

That precision capability is fine when working with numbers in the database, but Java uses its own set of types to represent numbers. This means you must be careful when selecting the Java type that will be used to represent numbers in your programs, especially if those numbers are going to be stored in a database.

To store integers in your Java program, you can use the short, int, long, or java.math .BigInteger types, depending on how big the integer you want to store is. Table 15-4 shows the number of bits used to store short, int, and long types, along with the low and high values supported by each type.

To store floating-point numbers in your Java programs, you can use the float, double, or java.math.BigDecimal types. Table 15-5 shows the same columns as Table 15-4 for the float and double types, along with the precision supported by each of these types.

As you can see, a float may be used to store floating-point numbers with a precision of up to 6 digits, and a double may be used for floating-point numbers with a precision of up to 15 digits. If you have a floating-point number that requires more than 15 digits of precision for storage in your Java program, you can use the java.math.BigDecimal type, which can store an arbitrarily long floating-point number.

In addition to these types, there is one of the Oracle JDBC extension types you can use to store your integers or floating-point numbers. This type is oracle.sql.NUMBER, and it allows you to store numbers with up to 38 digits of precision. You'll learn more about the oracle.sql . NUMBER type later in this chapter. In Oracle Database 10g and above, you can use the oracle .sql.BINARY FLOAT and oracle.sql.BINARY DOUBLE types. These types allow you to store the BINARY FLOAT and BINARY DOUBLE numbers.

Let's take a look at some examples of using these integer and floating-point types to store the product id and price column values for a row retrieved from the products table. Assume that a ResultSet object named productResultSet has been populated with the product id and price columns for a row from the products table. The product id column is defined

Туре	Bits	Low Value	High Value
short	16	-32768	32767
int	32	-2147483648	2147483647
long	64	-9223372036854775808	9223372036854775807

 TABLE 15-4
 The short, int, and long Types

Туре	Bits	Low Value	High Value	Precision
float	32	-3.4E+38	3.4E+38	6 digits
double	64	-1.7E+308	1.7E+308	15 digits

TABLE 15-5 The float and double Types

as a database INTEGER, and the price column is defined as a database NUMBER. The following example creates variables of the various integer and floating-point types and retrieves the product id and price column values into those variables:

```
short productIdShort = productResultSet.getShort("product_id");
int productIdInt = productResultSet.getInt("product_id");
long productIdLong = productResultSet.getLong("product_id");
float priceFloat = productResultSet.getFloat("price");
double priceDouble = productResultSet.getDouble("price");
java.math.BigDecimal priceBigDec = productResultSet.getBigDecimal("price");
```

Notice the use of the different get methods to retrieve the column values as the different types, the output of which is then stored in a Java variable of the appropriate type.

Handling Database Null Values

A column in a database table may be defined as being NULL or NOT NULL. NULL indicates that the column may store a NULL value; NOT NULL indicates that the column may not contain a NULL value. A NULL value means that the value is unknown. When a table is created in the database and you don't specify that a column is NULL or NOT NULL, the database assumes you mean NULL.

The Java object types, such as String, may be used to store database NULL values. When a query is used to retrieve a column that contains a NULL value into a Java String, that String will contain a Java null value. For example, the phone column (a VARCHAR2) for customer #5 is NULL, and the following statement uses the getString() method to read that value into a String named phone:

```
phone = customerResultSet.getString("phone");
```

Once the statement is run, the phone Java String will contain the Java null value.

That method's fine for NULL values being stored in Java objects, but what about the Java numeric, logical, and bit type types? If you retrieve a NULL value into a Java numeric, logical, or bit variable—int, float, boolean, or byte, for example—that variable will contain the value zero. To the database, zero and NULL are different values: zero is a definite value; NULL means the value is unknown. This causes a problem if you want to differentiate between zero and NULL in your Java program.

There are two ways to get around this problem:

■ You can use the wasNull() method in the ResultSet. The wasNull() method returns true if the value retrieved from the database was NULL; otherwise, the method returns false.

You can use a Java wrapper class. A wrapper class is a Java class that allows you to define a wrapper object, which can then be used to store the column value returned from the database. A wrapper object stores database NULL values as Java null values, and non-NULL values are stored as regular values.

Let's take a look at an example that illustrates the use of first technique, using product #12 from the products table. This row has a NULL value in the product type id column, and this column is defined as a database INTEGER. Also, assume that a ResultSet object named productResultSet has been populated with the product id and product type id columns for product #12 from the products table. The following example uses the wasNull() method to check if the value read for the product type id column was NULL:

```
System.out.println("product type id = " +
       productResultSet.getInt("product type id"));
     if (productResultSet.wasNull()) {
       System.out.println("Last value read was NULL");
```

Because the product type id column contains a NULL value, wasNull() will return true, and so the string "Last value read was NULL" would be displayed.

Before you see an example of the second method that uses the Java wrapper classes, I need to explain what these wrapper classes actually are. The wrapper classes are defined in the java .lang package, with the following seven wrapper classes being defined in that package:

- java.lang.Short
- java.lang.Integer
- java.lang.Long
- java.lang.Float
- java.lang.Double
- java.lang.Boolean
- java.lang.Byte

Objects declared using these wrapper classes can be used to represent database NULL values for the various types of numbers as well as for the Boolean type. When a database NULL is retrieved into such an object, it will contain the Java null value. The following example declares a java.lang.Integer named productTypeId:

```
java.lang.Integer productTypeId;
```

A database NULL may then be stored in productTypeId using a call to the getObject() method, as shown in the following example:

```
productTypeId =
       (java.lang.Integer) productResultSet.getObject("product type id");
```

The getObject() method returns an instance of the java.lang.Object class and must be cast into an appropriate type, in this case, to a java.lang.Integer. Assuming this example reads the same row from productResultSet as the previous example, getObject() will return a Java null value, and this value will be copied into productTypeId. Of course, if the value retrieved from the database had a value other than NULL, productTypeId would contain that value. For example, if the value retrieved from the database was 1, productTypeId would contain the value 1.

You can also use a wrapper object in a JDBC statement that performs an INSERT or UPDATE to set a column to a regular value or a NULL value. If you want to set a column value to NULL using a wrapper object, you would set that wrapper object to null and use it in an INSERT or UPDATE statement to set the database column to NULL. The following example sets the price column for product #12 to NULL using a java.lang.Double object that is set to null:

```
java.lang.Double price = null;
myStatement.executeUpdate(
   "UPDATE products " +
   "SET price = " + price + " " +
   "WHERE product_id = 12"
);
```

Controlling Database Transactions

In Chapter 8 you learned about database transactions and how to use the SQL COMMIT statement to permanently record changes you make to the contents of tables. You also saw how to use the ROLLBACK statement to undo changes in a database transaction. The same concepts apply to SQL statements executed using JDBC statements within your Java programs.

By default, the results of your INSERT, UPDATE, and DELETE statements executed using JDBC are immediately committed. This is known as *auto-commit* mode. Generally, using auto-commit mode is *not* the preferred way of committing changes, because it is counter to the idea of considering transactions as logical units of work. With auto-commit mode, all statements are considered as individual transactions, and this assumption is usually incorrect. Also, auto-commit mode may cause your SQL statements to take longer to complete, due to the fact that each statement is always committed.

Fortunately, you can enable or disable auto-commit mode using the setAutoCommit() method of the Connection class, passing it a Boolean true or false value. The following example disables auto-commit mode for the Connection object named myConnection:

myConnection.setAutoCommit(false);



TIP

You should disable auto-commit mode. Doing this will usually make your programs run faster.

Once auto-commit has been disabled, you can commit your transaction changes using the commit() method of the Connection class, or you can roll back your changes using the rollback() method. In the following example, the commit() method is used to commit changes made to the database using the myConnection object:

```
myConnection.commit();
```

In the next example, the rollback() method is used to roll back changes made to the database:

```
myConnection.rollback();
```

If auto-commit has been disabled and you close your Connection object, an implicit commit is performed. Therefore, any DML statements you have performed up to that point and haven't already committed will be committed automatically.

Performing Data Definition Language Statements

The SQL Data Definition Language (DDL) statements are used to create database users, tables, and many other types of database structures that make up a database. DDL consists of statements such as CREATE, ALTER, DROP, TRUNCATE, and RENAME. DDL statements may be performed in JDBC using the execute() method of the Statement class. In the following example, the CREATE TABLE statement is used to create a table named addresses, which may be used to store customer addresses:

```
myStatement.execute(
      "CREATE TABLE addresses (" +
         address id INTEGER CONSTRAINT addresses pk PRIMARY KEY, " +
         customer id INTEGER CONSTRAINT addresses fk customers " +
          REFERENCES customers (customer id)," +
        street VARCHAR2(20) NOT NULL," +
      " city VARCHAR2(20) NOT NULL," +
         state CHAR(2) NOT NULL" +
    );
```



NOTE

Performing a DDL statement results in an implicit commit being issued. Therefore, if you've performed uncommitted DML statements prior to issuing a DDL statement, those DML statements will also be committed.

Handling Exceptions

When an error occurs in either the database or the JDBC driver, a java.sql.SQLException will be raised. The java.sql.SQLException class is a subclass of the java.lang.Exception class. For this reason, you must place all your JDBC statements within a try/catch statement in order for your code not to throw a java.sql.SQLException. When such an exception occurs, Java attempts to locate the appropriate handler to process the exception.

If you include a handler for a java.sql.SQLException in a catch clause, when an error occurs in either the database or the JDBC driver, Java will move to that handler and run the appropriate code that you've included in that catch clause. In the handler code, you can do things like display the error code and error message, which will help you determine what happened.

The following try/catch statement contains a handler for exceptions of type java.sql . SQLException that may occur in the try statement:

```
try {
    } catch (SQLException e) {
    }
```



NOTE

I'm assuming java.sql.* has been imported, so I can simply use SQLException in the catch, rather than having to reference java.sql.SQLException.

The try statement will contain your JDBC statements that may cause an SQLException to be thrown, and the catch clause will contain your error handling code.

The SQLException class defines four methods that are useful for finding out what caused the exception to occur:

- **getErrorCode ()** In the case of errors that occur in the database or the JDBC driver, this method returns the Oracle error code, which is a five-digit number.
- **getMessage ()** In the case of errors that occur in the database, this method returns the error message along with the five-digit Oracle error code. In the case of errors that occur in the JDBC driver, this method returns just the error message.
- **getSQLState()** In the case of errors that occur in the database, this method returns a five-digit code containing the SQL state. In the case of errors that occur in the JDBC driver, this method doesn't return anything of interest.
- **printStackTrace()** This method displays the contents of the stack when the exception occurred. This information may further assist you in finding out what went wrong.

The following try/catch statement illustrates the use of these four methods:

```
try {
    ...
} catch (SQLException e) {
    System.out.println("Error code = " + e.getErrorCode());
    System.out.println("Error message = " + e.getMessage());
    System.out.println("SQL state = " + e.getSQLState());
    e.printStackTrace();
}
```

If your code throws an SQLException rather than handling it locally as just shown, Java will search for an appropriate handler in the calling procedure or function until one is found. If none is found, the exception will be handled by the default exception handler, which displays the Oracle error code, the error message, and the stack trace.

Closing Your JDBC Objects

In the examples shown in this chapter, I've created a number of JDBC objects: a Connection object named myConnection, a Statement object named myStatement, and two ResultSet objects named customerResultSet and productResultSet. ResultSet objects should be closed when they are no longer needed using the close () method. Similarly, you should also close the Statement and Connection objects when those objects are no longer needed.

In the following example, the myStatement and myConnection objects are closed using the close() method:

```
myStatement.close();
    myConnection.close();
```

You should typically close your Statement and Connection objects in a finally clause. Any code contained in a finally clause is guaranteed to be run, no matter how control leaves the try statement. If you want to add a finally clause to close your Statement and Connection objects, those objects should be declared before the first try/catch statement used to trap exceptions. The following example shows how to structure the main () method so that the Statement and Connection objects may be closed in a finally clause:

```
public static void main (String args []) {
       // declare Connection and Statement objects
       Connection myConnection = null;
       Statement myStatement = null;
       try {
         // register the Oracle JDBC drivers
         DriverManager.registerDriver(
           new oracle.jdbc.driver.OracleDriver()
         );
         // connect to the database as store
         // using the Oracle JDBC Thin driver
         myConnection = DriverManager.getConnection(
           "jdbc:oracle:thin:@localhost:1521:ORCL",
           "store",
           "store password"
         );
         // create a Statement object
         myStatement = myConnection.createStatement();
         // more of your code goes here
       } catch (SQLException e) {
         e.printStackTrace();
       } finally {
         try {
           // close the Statement object using the close() method
           if (myStatement != null) {
             myStatement.close();
           // close the Connection object using the close() method
           if (myConnection != null) {
             myConnection.close();
         } catch (SQLException e) {
           e.printStackTrace();
     } // end of main()
```

Notice that the code in the finally clause checks to see if the Statement and Connection objects are not equal to null before closing them using the close () method. If they are equal to null, there is no need to close them. Because the code in the finally clause is the last thing to be run and is guaranteed to be run, the Statement and Connection objects are always closed, regardless of what else happens in your program. For the sake of brevity, only the first program featured in this chapter uses a finally clause to close the Statement and Connection objects.

You have now seen how to write JDBC statements that connect to a database, run DML and DDL statements, control transactions, handle exceptions, and close JDBC objects. The following section contains a complete program that illustrates the use of JDBC.

Example Program: BasicExample1.java

The BasicExample1.java program illustrates the concepts covered in this chapter so far. This program and the other programs featured in this chapter may be found in the Java folder where you extracted this book's Zip file. All the programs contain detailed comments that you should study.

```
/*
       BasicExample1.java shows how to:
       - import the JDBC packages
       - load the Oracle JDBC drivers
       - connect to a database
       - perform DML statements
       - control transactions
       - use ResultSet objects to retrieve rows
       - use the get methods
       - perform DDL statements
     // import the JDBC packages
     import java.sql.*;
     public class BasicExample1 {
       public static void main (String args []) {
         // declare Connection and Statement objects
         Connection myConnection = null;
         Statement myStatement = null;
         try {
           // register the Oracle JDBC drivers
           DriverManager.registerDriver(
             new oracle.jdbc.OracleDriver()
           );
           // EDIT AS NECESSARY TO CONNECT TO YOUR DATABASE
           // create a Connection object, and connect to the database
           // as the store user using the Oracle JDBC Thin driver
           myConnection = DriverManager.getConnection(
             "jdbc:oracle:thin:@localhost:1521:ORCL",
             "store",
```

```
"store password"
);
// disable auto-commit mode
myConnection.setAutoCommit(false);
// create a Statement object
myStatement = myConnection.createStatement();
// create variables and objects used to represent
// column values
int customerId = 6;
String firstName = "Jason";
String lastName = "Red";
java.sql.Date dob = java.sql.Date.valueOf("1969-02-22");
java.sql.Time dobTime;
java.sql.Timestamp dobTimestamp;
String phone = "800-555-1216";
// perform SQL INSERT statement to add a new row to the
// customers table using the values set in the previous
// step - the executeUpdate() method of the Statement
// object is used to perform the INSERT
myStatement.executeUpdate(
  "INSERT INTO customers " +
  "(customer id, first name, last name, dob, phone) VALUES (" +
    customerId + ", '" + firstName + "', '" + lastName + "', " +
  "TO DATE('" + dob + "', 'YYYY, MM, DD'), '" + phone + "')"
);
System.out.println("Added row to customers table");
// perform SQL UPDATE statement to modify the first name
// column of customer #1
firstName = "Jean";
myStatement.executeUpdate(
  "UPDATE customers " +
  "SET first name = '" + firstName + "' " +
  "WHERE customer id = 1"
);
System.out.println("Updated row in customers table");
// perform SQL DELETE statement to remove customer #5
myStatement.executeUpdate(
 "DELETE FROM customers " +
  "WHERE customer id = 5"
);
System.out.println("Deleted row from customers table");
// create a ResultSet object, and populate it with the
// result of a SELECT statement that retrieves the
// customer id, first name, last name, dob, and phone columns
```

```
// for all the rows from the customers table - the
// executeQuery() method of the Statement object is used
// to perform the SELECT
ResultSet customerResultSet = myStatement.executeOuery(
 "SELECT customer id, first name, last name, dob, phone " +
  "FROM customers"
System.out.println("Retrieved rows from customers table");
// loop through the rows in the ResultSet object using the
// next() method, and use the get methods to read the values
// retrieved from the database columns
while (customerResultSet.next()) {
 customerId = customerResultSet.getInt("customer id");
 firstName = customerResultSet.getString("first name");
 lastName = customerResultSet.getString("last name");
 dob = customerResultSet.getDate("dob");
 dobTime = customerResultSet.getTime("dob");
 dobTimestamp = customerResultSet.getTimestamp("dob");
 phone = customerResultSet.getString("phone");
 System.out.println("customerId = " + customerId);
 System.out.println("firstName = " + firstName);
 System.out.println("lastName = " + lastName);
 System.out.println("dob = " + dob);
 System.out.println("dobTime = " + dobTime);
 System.out.println("dobTimestamp = " + dobTimestamp);
 System.out.println("phone = " + phone);
} // end of while loop
// close the ResultSet object using the close() method
customerResultSet.close();
// roll back the changes made to the database
myConnection.rollback();
// create numeric variables to store the product id and price columns
short productIdShort;
int productIdInt;
long productIdLong;
float priceFloat;
double priceDouble;
java.math.BigDecimal priceBigDec;
// create another ResultSet object and retrieve the
// product id, product type id, and price columns for product #12
// (this row has a NULL value in the product type id column)
ResultSet productResultSet = myStatement.executeQuery(
 "SELECT product id, product type id, price " +
 "FROM products " +
 "WHERE product id = 12"
);
```

```
System.out.println("Retrieved row from products table");
while (productResultSet.next()) {
  System.out.println("product id = " +
    productResultSet.getInt("product id"));
 System.out.println("product type id = " +
    productResultSet.getInt("product type id"));
 // check if the value just read by the get method was NULL
 if (productResultSet.wasNull()) {
    System.out.println("Last value read was NULL");
 // use the getObject() method to read the value, and convert it
 // to a wrapper object - this converts a database NULL value to a
 // Java null value
  java.lang.Integer productTypeId =
    (java.lang.Integer) productResultSet.getObject("product type id");
 System.out.println("productTypeId = " + productTypeId);
 // retrieve the product id and price column values into
 // the various numeric variables created earlier
 productIdShort = productResultSet.getShort("product id");
 productIdInt = productResultSet.getInt("product id");
 productIdLong = productResultSet.getLong("product id");
 priceFloat = productResultSet.getFloat("price");
 priceDouble = productResultSet.getDouble("price");
 priceBigDec = productResultSet.getBigDecimal("price");
 System.out.println("productIdShort = " + productIdShort);
 System.out.println("productIdInt = " + productIdInt);
 System.out.println("productIdLong = " + productIdLong);
 System.out.println("priceFloat = " + priceFloat);
 System.out.println("priceDouble = " + priceDouble);
 System.out.println("priceBigDec = " + priceBigDec);
} // end of while loop
// close the ResultSet object
productResultSet.close();
// perform SQL DDL CREATE TABLE statement to create a new table
// that may be used to store customer addresses
myStatement.execute(
 "CREATE TABLE addresses (" +
 " address id INTEGER CONSTRAINT addresses pk PRIMARY KEY," +
 " customer id INTEGER CONSTRAINT addresses fk customers " +
    REFERENCES customers (customer id)," +
 " street VARCHAR2(20) NOT NULL," +
  " city VARCHAR2(20) NOT NULL," +
 " state CHAR(2) NOT NULL" +
  11) 11
);
System.out.println("Created addresses table");
```

```
// drop this table using the SQL DDL DROP TABLE statement
     myStatement.execute("DROP TABLE addresses");
      System.out.println("Dropped addresses table");
    } catch (SOLException e) {
      System.out.println("Error code = " + e.getErrorCode());
      System.out.println("Error message = " + e.getMessage());
      System.out.println("SQL state = " + e.getSQLState());
      e.printStackTrace();
    } finally {
      try {
       // close the Statement object using the close() method
       if (myStatement != null) {
         myStatement.close();
       // close the Connection object using the close() method
       if (myConnection != null) {
         myConnection.close();
      } catch (SQLException e) {
       e.printStackTrace();
  } // end of main()
}
```



NOTE

You may need to edit the line labeled with the text EDIT AS NECESSARY... with the correct settings to access your database.

Compile BasicExample1

To compile BasicExample1.java, you type the following command using your operating system command prompt:

javac BasicExample1.java

If you haven't set the CLASSPATH environment variable properly, you'll get the following error message when trying to compile the FirstExample.java program:

```
FirstExample.java:22: cannot resolve symbol symbol : class OracleDriver location: package jdbc new oracle.jdbc.OracleDriver()
```

You should check the setting for your CLASSPATH environment variable—it's likely your CLASSPATH is missing the Oracle JDBC classes file (ojdbc6.jar, for example). Refer to the earlier section "Setting the CLASSPATH Environment Variable."



HP

You can enter javac -help to get help on the Java compiler.

Run BasicExample1

Once BasicExample1. java is compiled, you can run the resulting executable class file (named BasicExample1.class) by entering the following command:

java BasicExample1



CAUTION

Java is case-sensitive, so make sure you enter BasicExample1 with uppercase B and E characters.

If the program fails with the following error code and message, it means the store user with a password of store password doesn't exist in your database:

```
Error code = 1017
  Error message = ORA-01017: invalid username/password; logon denied
```

If you get this error, check that the store user is in the database.

The program may also be unable to find your database, in which case you'll get the following

```
Error code = 17002
   Error message = Io exception: The Network Adapter could not establish
    the connection
```

Typically, there are two reasons why you might get this error:

- There is no database running on your localhost machine with the Oracle SID of ORCL.
- Oracle Net is not running or is not listening for connections on port 1521.

You should ensure that you have the right connection string in the program and also that the database and Oracle Net are running.

Assuming the program runs, you should get the following output:

```
Added row to customers table
    Updated row in customers table
    Deleted row from customers table
    Retrieved rows from customers table
    customerId = 1
    firstName = Jean
    lastName = Brown
    dob = 1965-01-01
    dobTime = 00:00:00
    dobTimestamp = 1965-01-01 00:00:00.0
   phone = 800-555-1211
    customerId = 2
    firstName = Cynthia
    lastName = Green
    dob = 1968-02-05
    dobTime = 00:00:00
    dobTimestamp = 1968-02-05 00:00:00.0
    phone = 800-555-1212
```

```
customerId = 3
firstName = Steve
lastName = White
dob = 1971 - 03 - 16
dobTime = 00:00:00
dobTimestamp = 1971-03-16 00:00:00.0
phone = 800-555-1213
customerId = 4
firstName = Gail
lastName = Black
dob = null
dobTime = null
dobTimestamp = null
phone = 800-555-1214
customerId = 6
firstName = Jason
lastName = Red
dob = 1969-02-22
dobTime = 00:00:00
dobTimestamp = 1969-02-22 00:00:00.0
phone = 800-555-1216
Retrieved row from products table
product id = 12
product type id = 0
Last value read was NULL
productTypeId = null
productIdShort = 12
productIdInt = 12
productIdLong = 12
priceFloat = 13.49
priceDouble = 13.49
priceBigDec = 13.49
Created addresses table
Dropped addresses table
```

Prepared SQL Statements

When you send an SQL statement to the database, the database software reads the SQL statement and verifies that it is correct. This is known as *parsing* the SQL statement. The database software then builds a plan, known as the *execution plan*, to actually run the statement. So far, all the SQL statements sent to the database through JDBC have required a new execution plan to be built. This is because each SQL statement sent to the database has been different.

Suppose you had a Java application that was performing the same INSERT statement repeatedly; an example is loading many new products to our example store, a process that would require adding lots of rows to the products table using INSERT statements. Let's see Java statements that would actually do this. Assume that a class named Product has been defined as follows:

```
class Product {
   int productId;
   int productTypeId;
   String name;
```

```
String description;
double price;
```

The following code creates an array of five Product objects. Because the products table already contains rows with product id values from 1 to 12, the productId attributes for the new Product objects start at 13:

```
Product [] productArray = new Product[5];
     for (int counter = 0; counter < productArray.length; counter ++) {</pre>
       productArray[counter] = new Product();
       productArray[counter].productId = counter + 13; //start at 13
       productArray[counter].productTypeId = 1;
       productArray[counter].name = "Test product";
       productArray[counter].description = "Test product";
       productArray[counter].price = 19.95;
     } // end of for loop
```

To add the rows to the products table, I'll use a for loop that contains a IDBC statement to perform an INSERT statement, and the column values will come from productArray:

```
Statement myStatement = myConnection.createStatement();
     for (int counter = 0; counter < productArray.length; counter ++) {</pre>
       myStatement.executeUpdate(
         "INSERT INTO products " +
         "(product id, product type id, name, description, price) VALUES (" +
         productArray[counter]. productId + ", " +
         productArray[counter]. productTypeId + ", '" +
         productArray[counter].name + "', '" +
         productArray[counter].description + "', " +
        productArray[counter].price + ")"
       );
     } // end of for loop
```

Each iteration through the loop results in an INSERT statement being sent to the database. Because the string representing each INSERT statement contains different values, the actual INSERT sent to the database is slightly different each time. This means that the database creates a different execution plan for every INSERT statement—very inefficient.

You'll be glad to know that JDBC provides a better way to run such SQL statements. Instead of using a JDBC Statement object to run your SQL statements, you can use a JDBC PreparedStatement object. A PreparedStatement object allows you to perform the same SQL statement but supply different values for actual execution of that statement. This is more efficient because the same execution plan is used by the database when the SQL statement is run. The following example creates a PreparedStatement object containing an INSERT statement similar to the one used in the previous loop:

```
PreparedStatement myPrepStatement = myConnection.prepareStatement(
      "INSERT INTO products " +
      "(product id, product type id, name, description, price) VALUES (" +
      "?, ?, ?, ?, ?"
      ")"
    );
```

There are two things you should notice about this example:

- The prepareStatement() method is used to specify the SQL statement.
- Question mark characters (?) are used to indicate the positions where you will later provide variables to be used when the SQL statement is actually run.

The positions of the question marks are important: they are referenced according to their position, with the first question mark being referenced using number 1, the second as number 2, and so on.

The process of supplying Java variables to a prepared statement is known as *binding* the variables to the statement, and the variables themselves are known as *bind variables*. To actually supply variables to the prepared SQL statement, you must use set methods. These methods are similar to the get methods discussed earlier, except that set methods are used to supply variable values, rather than read them.

For example, to bind a Java int variable named intVar to the product_id column in the PreparedStatement object, you use setInt(1, intVar). The first parameter indicates the numeric position of the question mark (?) in the string previously specified in the prepareStatement() method call. For this example, the value 1 corresponds to the first question mark, which supplies a value to the product_id column in the INSERT statement. Similarly, to bind a Java String variable named stringVar to the name column, you use setString(3, stringVar), because the third question mark corresponds to the name column. Other methods you can call in a PreparedStatement object include setFloat() and setDouble(), which are used for setting single-precision floating-point and double-precision floating-point numbers.

The following example features a loop that shows the use of set methods to bind the attributes of the Product objects in productArray to the PreparedStatement object; notice the execute () method is used to actually run the SQL statement:

```
for (int counter = 0; counter < productArray.length; counter ++) {
    myPrepStatement.setInt(1, productArray[counter]. productId);
    myPrepStatement.setInt(2, productArray[counter]. productTypeId);
    myPrepStatement.setString(3, productArray[counter].name);
    myPrepStatement.setString(4, productArray[counter].description);
    myPrepStatement.setDouble(5, productArray[counter].price);
    myPrepStatement.execute();
} // end of for loop</pre>
```

After this code is executed, the products table will contain five new rows.

To set a database column to NULL using a PreparedStatement object, you may use the setNull() method. For example, the following statement sets the description column to NULL:

```
myPrepStatement.setNull(4, java.sql.Types.VARCHAR);
```

The first parameter in the call to setNull() is the numeric position of the column you want to set to NULL. The second parameter is an int that corresponds to the database type of the column that is to be set to NULL. This second parameter should be specified using one of the constants defined in the java.sql.Types class. For a VARCHAR2 column (the description column is a VARCHAR2), you should use java.sql.Types.VARCHAR.

Example Program: BasicExample2.java

The following BasicExample2.java program contains the statements shown in the previous sections.

```
/*
       BasicExample2.java shows how to use prepared SQL statements
     // import the JDBC packages
     import java.sql.*;
     class Product {
       int productId;
       int productTypeId;
       String name;
       String description;
       double price;
     public class BasicExample2 {
       public static void main (String args []) {
         try {
           // register the Oracle JDBC drivers
           DriverManager.registerDriver(
             new oracle.jdbc.OracleDriver()
           );
           // EDIT AS NECESSARY TO CONNECT TO YOUR DATABASE
           // create a Connection object, and connect to the database
           // as the store user using the Oracle JDBC Thin driver
           Connection myConnection = DriverManager.getConnection(
             "jdbc:oracle:thin:@localhost:1521:ORCL",
             "store",
             "store password"
           );
           // disable auto-commit mode
           myConnection.setAutoCommit(false);
           Product [] productArray = new Product[5];
           for (int counter = 0; counter < productArray.length; counter ++) {</pre>
             productArray[counter] = new Product();
             productArray[counter].productId = counter + 13;
             productArray[counter].productTypeId = 1;
             productArray[counter].name = "Test product";
             productArray[counter].description = "Test product";
             productArray[counter].price = 19.95;
            } // end of for loop
           // create a PreparedStatement object
```

```
PreparedStatement myPrepStatement = myConnection.prepareStatement(
  "INSERT INTO products " +
  "(product id, product type id, name, description, price) VALUES (" +
 "?, ?, ?, ?" +
  ")"
);
// initialize the values for the new rows using the
// appropriate set methods
for (int counter = 0; counter < productArray.length; counter ++) {</pre>
 myPrepStatement.setInt(1, productArray[counter].productId);
 myPrepStatement.setInt(2, productArray[counter].productTypeId);
 myPrepStatement.setString(3, productArray[counter].name);
 myPrepStatement.setString(4, productArray[counter].description);
 myPrepStatement.setDouble(5, productArray[counter].price);
 myPrepStatement.execute();
} // end of for loop
// close the PreparedStatement object
myPrepStatement.close();
// retrieve the product id, product type id, name, description, and
// price columns for these new rows using a ResultSet object
Statement myStatement = myConnection.createStatement();
ResultSet productResultSet = myStatement.executeQuery(
  "SELECT product id, product type id, " +
  " name, description, price " +
 "FROM products " +
 "WHERE product id > 12"
);
// display the column values
while (productResultSet.next()) {
  System.out.println("product id = " +
    productResultSet.getInt("product id"));
  System.out.println("product type id = " +
    productResultSet.getInt("product type id"));
  System.out.println("name = " +
    productResultSet.getString("name"));
 System.out.println("description = " +
    productResultSet.getString("description"));
 System.out.println("price = " +
   productResultSet.getDouble("price"));
} // end of while loop
// close the ResultSet object using the close() method
productResultSet.close();
// roll back the changes made to the database
myConnection.rollback();
// close the other JDBC objects
```

```
myStatement.close();
   myConnection.close();
  } catch (SOLException e) {
    System.out.println("Error code = " + e.getErrorCode());
    System.out.println("Error message = " + e.getMessage());
   System.out.println("SQL state = " + e.getSQLState());
   e.printStackTrace();
} // end of main()
```

The output from this program is as follows:

```
product id = 13
     product type id = 1
     name = Test product
     description = Test product
     price = 19.95
     product id = 14
     product type id = 1
     name = Test product
     description = Test product
     price = 19.95
     product id = 15
     product_type id = 1
     name = Test product
     description = Test product
     price = 19.95
     product id = 16
     product_type id = 1
     name = Test product
     description = Test product
     price = 19.95
     product id = 17
     product type id = 1
     name = Test product
     description = Test product
     price = 19.95
```

The Oracle JDBC Extensions

The Oracle extensions to JDBC enable you to access all of the data types provided by Oracle, along with Oracle-specific performance extensions. You'll learn about handling of strings, numbers, dates, and row identifiers in this section. You may read my book Oracle9i JDBC Programming for all of the Oracle types and performance enhancements.

There are two JDBC extension packages supplied by Oracle Corporation:

- **oracle.sql** contains the classes that support all the Oracle database types.
- oracle.jdbc contains the interfaces that support access to an Oracle database.

To import the Oracle JDBC packages into your Java programs, you add the following import statements to your program:

```
import oracle.sql.*;
import oracle.jdbc.*;
```

Of course, you don't have to import all the packages—you could just import the classes and interfaces you actually use in your program. In the following sections, you'll learn the key features of the oracle.sql and oracle.jdbc packages.

The oracle.sql Package

The oracle.sql package contains the classes that support all of the Oracle database types. Using objects of the classes defined in this package to store database values is more efficient than using regular Java objects. This is because the database values don't need to be converted to an appropriate base Java type first. Also, using a Java float or double to represent a floating-point number may result in a loss of precision for that number. If you use an oracle.sql.NUMBER object, your numbers never lose precision.



TIP

If you are writing a program that moves a lot of data around in the database, you should use the oracle.sql.* classes.

The oracle.sql classes extend the oracle.sql.Datum class, which contains the functionality that is common to all the classes. Table 15-6 shows a subset of the oracle.sql classes, along with the mapping to the compatible Oracle database types.

From Table 15-6, you can see that an oracle.sql.NUMBER object is compatible with a database INTEGER or NUMBER type. An oracle.sql.CHAR object is compatible with a database CHAR, VARCHAR2, NCHAR, and NVARCHAR2 (NCHAR and NVARCHAR2 are typically used to store non-English characters).

Class	Compatible Database Type
oracle.sql.NUMBER	INTEGER NUMBER
oracle.sql.CHAR	CHAR VARCHAR2 NCHAR NVARCHAR2
oracle.sql.DATE	DATE
oracle.sql.BINARY_FLOAT	BINARY_FLOAT
oracle.sql.BINARY_DOUBLE	BINARY_DOUBLE
oracle.sql.ROWID	ROWID

TABLE 15-6 Classes and Compatible Oracle Database Types

Objects declared using the oracle.sql classes store data as byte arrays—also known as SOL format—and don't reformat the data retrieved from the database. This means that no information is lost when converting between an SQL format object and a value stored in the database.

Each oracle.sql object has a getBytes () method that returns the binary data stored in the object. Each object also has a toldbc() method that returns the binary data as a compatible lava type (an exception to this is the oracle.sql.ROWID, where toJdbc() always returns an oracle.sql.ROWID).

Each oracle.sql class also provides methods to convert their SQL format binary data to a core Java type. For example, stringValue() returns the value as Java String, intValue() returns a Java int, floatValue() returns a float, doubleValue() returns a double, bigDecimalValue() returns a java.math.BigDecimal, dateValue() returns a java . sql. Date, and so on. You use these methods when you want to store the SQL format data in a core Java type or to display the SQL data on the screen. Each oracle.sql class also contains a constructor that accepts a Java variable, object, or byte array as input.

As you will see later, the oracle.jdbc.OraclePreparedStatement class contains a number of set methods that you use to specify values for oracle.sql objects. The OracleResultSet class defines a number of get methods that you use to read values from oracle.sql objects.

The following sections describe the main oracle.sql classes.

The oracle.sql.NUMBER Class

The oracle.sql.NUMBER class is compatible with the database INTEGER and NUMBER types. The oracle.sql.NUMBER class may be used to represent a number with up to 38 digits of precision. The following example creates an oracle.sql.NUMBER object named customerid, which is set to the value 6 using the constructor:

```
oracle.sql.NUMBER customerId = new oracle.sql.NUMBER(6);
```

You can read the value stored in customerId using the intValue() method, which returns the value as a Java int. For example:

```
int customerIdInt = customerId.intValue();
```

You can also set an oracle.sql.NUMBER object to a floating-point number. The next example passes the value 19.95 to the constructor of an object named price:

```
oracle.sql.NUMBER price = new oracle.sql.NUMBER(19.95);
```

You can read the floating-point number stored in price using the floatValue(), doubleValue(), and bigDecimalValue() methods, which return a Java float, double, and bigDecimal respectively. You can also get a floating-point number truncated to an int using intValue() (for example, 19.95 would be returned as 19). The following examples show the use of these methods:

```
float priceFloat = price.floatValue();
    double priceDouble = price.doubleValue();
     java.math.BigDecimal priceBigDec = price.bigDecimalValue();
     int priceInt = price.intValue();
```

The stringValue () method returns the value as a Java String:

String priceString = price.stringValue();

The oracle.sql.CHAR Class

The oracle.sql.CHAR class is compatible with the database CHAR, VARCHAR2, NCHAR, and NVARCHAR2 types. Both the Oracle database and the oracle.sql.CHAR class contain globalization support for many different languages. For full details of the various languages supported by Oracle, see the *Oracle Database Globalization Support Guide* published by Oracle Corporation.

When you retrieve character data from the database into an oracle.sql.CHAR object, the Oracle JDBC driver returns that object using either the database character set (the default), the WE8ISO8859P1 character set (ISO 8859-1 West European), or the UTF8 character set (Unicode 3.0 UTF-8 Universal).

When passing an <code>oracle.sql.CHAR</code> object to the database, there are restrictions on the character set for the object. The character set depends on the database column type that the object will be stored in. If you are storing the <code>oracle.sql.CHAR</code> object in a <code>CHAR</code> or <code>VARCHAR2</code> column, you must use US7ASCII (ASCII 7-bit American), WE8ISO8859P1 (ISO 8859-1 West European), or UTF8 (Unicode 3.0 UTF-8 Universal). If you are storing the <code>oracle.sql.CHAR</code> object in an <code>NCHAR</code> or <code>NVARCHAR2</code> column, you must use the character set used by the database.

When creating your own oracle.sql.CHAR object, there are two steps you must follow:

- 1. Create an oracle.sql.CharacterSet object with the character set you wish to use.
- 2. Create an oracle.sql.CHAR object through the oracle.sql.CharacterSet object.

The following sections cover these steps.

Step I: Create an oracle.sql.CharacterSet Object The following example creates an oracle .sql.CharacterSet object named myCharSet:

```
oracle.sql.CharacterSet myCharSet =
CharacterSet.make(CharacterSet.US7ASCII CHARSET);
```

The make () method accepts an int that specifies the character set. In the example, the constant US7ASCII_CHARSET (defined in the oracle.sql.CharacterSet class) specifies the US7ASCII character set. Other int values include UTF8_CHARSET (for UTF8) and DEFAULT_CHARSET (for the character set used by the database).

Step 2: Create an oracle.sql.CHAR Object The next example creates an oracle.sql.CHAR object named firstName, using the myCharSet object created in the previous step:

```
oracle.sql.CHAR firstName = new oracle.sql.CHAR("Jason", myCharSet);
```

The firstName object is populated with the string Jason. You can read this string using the stringValue() method; for example:

```
String firstNameString = firstName.stringValue();
System.out.println("firstNameString = " + firstNameString);
```

This will display firstNameString = Jason.

Similarly, the following example creates another oracle.sql.CHAR object named lastName:

```
oracle.sql.CHAR lastName = new oracle.sql.CHAR("Price", myCharSet);
```

You can also display the value in an oracle.sql.CHAR object directly, as shown in the following example:

```
System.out.println("lastName = " + lastName);
```

This statement displays the following:

```
lastName = Price
```

The oracle.sql.DATE Class

The oracle.sql.DATE class is compatible with the database DATE type. The following example creates an oracle.sql.DATE object named dob:

```
oracle.sql.DATE dob = new oracle.sql.DATE("1969-02-22 13:54:12");
```

Notice that the constructor may accept a string in the format YYYY-MM-DD HH:MI:SS, where YYYY is the year, MM is the month, DD is the day, HH is the hour, MI is the minute, and SS is the second. You can read the value stored in dob as a Java String using the stringValue() method, as shown in the following example:

```
String dobString = dob.stringValue();
```

In this example, dobString will contain 2/22/1969 13:54:12 (the format changes to MM/DD/YYYY HH:MI:SS when using a Java String).

You can also pass a java.sql.Date object to the oracle.sql.DATE constructor, as shown in the following example:

```
oracle.sql.DATE anotherDob =
      new oracle.sql.DATE(java.sql.Date.valueOf("1969-02-22"));
```

In this example, anotherDob will contain the oracle.sql.DATE 1969-02-22 00:00:00.

The oracle.sql.ROWID Class

The oracle.sql.ROWID class is compatible with the database ROWID type. The ROWID contains the physical address of a row in the database. The following example creates an oracle.sql.ROWID object named rowid:

```
oracle.sql.ROWID rowid;
```

The oracle.jdbc Package

The classes and interfaces of the oracle.jdbc package allow you to read and write column values in the database via oracle.sql objects. The oracle.jdbc package also contains a number of performance enhancements. In this section, you'll learn about the contents of the oracle.jdbc package and see how to create a row in the customers table. Then you'll learn how to read that row using oracle.sql objects.

The Classes and Interfaces of the oracle.jdbc Package

Table 15-7 shows the classes and interfaces of the oracle.jdbc package.

Using an OraclePreparedStatement Object

The OraclePreparedStatement interface implements java.sql.PreparedStatement. This interface supports the various set methods for binding oracle.sql objects.

In the previous section, you saw the following oracle.sql objects:

- An oracle.sql.NUMBER object named customerId, which was set to 6
- An oracle.sql.CHAR object named firstName, which was set to Jason
- Another oracle.sql.CHAR object named lastName, which was set to Price
- An oracle.sql.DATE object named dob, which was set to 1969-02-22 13:54:12

To use these objects in an SQL DML statement, you must use an OraclePreparedStatement object, which contains set methods to handle oracle.sql objects. The following example creates an OraclePreparedStatement named myPrepStatement, which will be used later to add a row to the customers table:

```
OraclePreparedStatement myPrepStatement =
   (OraclePreparedStatement) myConnection.prepareStatement(
     "INSERT INTO customers " +
     "(customer_id, first_name, last_name, dob, phone) VALUES (" +
     "?, ?, ?, ?, ?" +
     ")"
);
```

Notice that the PreparedStatement object returned by the prepareStatement() method is cast to an OraclePreparedStatement object and is stored in myPrepStatement.

The next step is to bind the oracle.sql objects to myPrepStatement using the set methods. This involves assigning values to the placeholders marked by question mark (?) characters in myPrepStatement. Just as you use set methods like setInt(), setFloat(), setString(), and setDate() to bind Java variables to a PreparedStatement object, you also use set methods to bind oracle.sql objects to an OraclePreparedStatement object (these set methods include setNUMBER(), setCHAR(), setDATE(), and so on).

The following example illustrate how to bind the customerId, firstName, lastName, and dob objects to myPrepStatement using the appropriate set methods:

```
myPrepStatement.setNUMBER(1, customerId);
myPrepStatement.setCHAR(2, firstName);
myPrepStatement.setCHAR(3, lastName);
myPrepStatement.setDATE(4, dob);
```

The next example sets the fifth question mark (?) in myPrepStatement to NULL (the fifth question mark (?) corresponds to the phone column in the customers table):

```
myPrepStatement.setNull(5, OracleTypes.CHAR);
```

Name	Class or Interface	Description
OracleDriver	Class	Implements java.sql.Driver. You input an object of this class when registering the Oracle JDBC drivers using the registerDriver() method of the java.sql.DriverManager class.
OracleConnection	Interface	Implements java.sql.Connection. This interface extends the standard JDBC connection functionality to use OracleStatement objects. It also improves performance over the standard JDBC functions.
OracleStatement	Interface	Implements java.sql.Statement and is the superclass of the OraclePreparedStatement and OracleCallableStatement classes.
OraclePreparedStatement	Interface	Implements java.sql .PreparedStatement, and is the superclass of OracleCallableStatement.This interface supports the various set methods for binding oracle.sql objects.
OracleCallableStatement	Interface	Implements java.sql .CallableStatement.This interface contains various get and set methods for binding oracle.sql objects.
OracleResultSet	Interface	Implements java.sql.ResultSet. This interface contains various get methods for binding oracle.sql objects.
OracleResultSetMetaData	Interface	Implements java.sql .ResultSetMetaData. This interface contains methods for retrieving meta data about Oracle result sets (such as column names and their types).
OracleDatabaseMetaData	Class	Implements java.sql.DatabaseMetaData. This class contains methods for retrieving meta data about the Oracle database (such as the database software version).
OracleTypes	Class	Defines integer constants for the database types. This class duplicates the standard java.sql.Types class and contains additional integer constants for all of the Oracle types.

 TABLE 15-7
 Classes and Interfaces of the oracle.jdbc Package

The int constant OracleTypes.CHAR specifies that the database type is compatible with the oracle.sql.CHAR type; OracleTypes.CHAR is used because the phone column is a VARCHAR2.

The only thing left to do is to run the INSERT statement using the execute() method:

```
myPrepStatement.execute();
```

Doing this adds the row to the customers table.

Using an OracleResultSet Object

The OracleResultSet interface implements java.sql.ResultSet and contains get methods to handle oracle.sql objects. In this section, you'll see how to use an OracleResultSet object to retrieve the row previously added to the customers table.

The first thing needed is a JDBC Statement object through which an SQL statement may be run:

```
Statement myStatement = myConnection.createStatement();
```

Next, the following example creates an OracleResultSet object named customerResultSet, which is populated with the ROWID, customer_id, first_name, last_dob, and phone columns retrieved from customer #6:

```
OracleResultSet customerResultSet =
   (OracleResultSet) myStatement.executeQuery(
     "SELECT ROWID, customer_id, first_name, last_name, dob, phone " +
     "FROM customers " +
     "WHERE customer_id = 6"
);
```

I defined the following oracle.sql objects earlier: rowid, customerId, firstName, lastName, and dob. These may be used to hold the first five column values. In order to store the value for the phone column, an oracle.sql.CHAR object is needed:

```
oracle.sql.CHAR phone = new oracle.sql.CHAR("", myCharSet);
```

An OracleResultSet object contains get methods that return oracle.sql objects. You use getCHAR() to get an oracle.sql.CHAR, getNUMBER() to get an oracle.sql.NUMBER, getDATE() to get an oracle.sql.DATE, and so on.

The following while loop contains calls to the appropriate get methods to copy the values from customerResultSet to rowid, customerId, firstName, lastName, dob, and phone:

```
while (customerResultSet.next()) {
   rowid = customerResultSet.getROWID("ROWID");
   customerId = customerResultSet.getNUMBER("customer_id");
   firstName = customerResultSet.getCHAR("first_name");
   lastName = customerResultSet.getCHAR("last_name");
   dob = customerResultSet.getDATE("dob");
   phone = customerResultSet.getCHAR("phone");

   System.out.println("rowid = " + rowid.stringValue());
   System.out.println("customerId = " + customerId.stringValue());
   System.out.println("firstName = " + firstName);
   System.out.println("lastName = " + lastName);
   System.out.println("dob = " + dob.stringValue());
```

```
System.out.println("phone = " + phone);
} // end of while loop
```

To display the values, the example uses calls to the stringValue() method to convert the rowid, customerId, and dob objects to lava String values. For the firstName, lastName, and phone objects, the example simply uses these objects directly in the System.out .println() calls.

The following section shows a complete program containing the statements shown in the previous sections.

Example Program: BasicExample3.java

The following BasicExample3. java program contains the statements shown in the previous sections:

```
/*
      BasicExample3.java shows how to use the Oracle JDBC extensions
      to add a row to the customers table, and then retrieve that row
    // import the JDBC packages
    import java.sql.*;
    // import the Oracle JDBC extension packages
    import oracle.sql.*;
    import oracle.jdbc.*;
    public class BasicExample3 {
      public static void main (String args []) {
        trv {
          // register the Oracle JDBC drivers
          DriverManager.registerDriver(
            new oracle.jdbc.OracleDriver()
          );
          // EDIT AS NECESSARY TO CONNECT TO YOUR DATABASE
          // create a Connection object, and connect to the database
          // as the store user using the Oracle JDBC Thin driver
          Connection myConnection = DriverManager.getConnection(
            "jdbc:oracle:thin:@localhost:1521:ORCL",
            "store",
            "store password"
          );
          // disable auto-commit mode
          myConnection.setAutoCommit(false);
          // create an oracle.sql.NUMBER object
          oracle.sql.NUMBER customerId = new oracle.sql.NUMBER(6);
          int customerIdInt = customerId.intValue();
          System.out.println("customerIdInt = " + customerIdInt);
          // create two oracle.sql.CHAR objects
```

```
oracle.sql.CharacterSet myCharSet =
  CharacterSet.make(CharacterSet.US7ASCII CHARSET);
oracle.sql.CHAR firstName = new oracle.sql.CHAR("Jason", myCharSet);
String firstNameString = firstName.stringValue();
System.out.println("firstNameString = " + firstNameString);
oracle.sql.CHAR lastName = new oracle.sql.CHAR("Price", myCharSet);
System.out.println("lastName = " + lastName);
// create an oracle.sql.DATE object
oracle.sql.DATE dob = new oracle.sql.DATE("1969-02-22 13:54:12");
String dobString = dob.stringValue();
System.out.println("dobString = " + dobString);
// create an OraclePreparedStatement object
OraclePreparedStatement myPrepStatement =
  (OraclePreparedStatement) myConnection.prepareStatement(
    "INSERT INTO customers " +
    "(customer id, first name, last name, dob, phone) VALUES (" +
    "?, ?, ?, ?" +
    ")"
 );
// bind the objects to the OraclePreparedStatement using the
// appropriate set methods
myPrepStatement.setNUMBER(1, customerId);
myPrepStatement.setCHAR(2, firstName);
myPrepStatement.setCHAR(3, lastName);
myPrepStatement.setDATE(4, dob);
// set the phone column to NULL
myPrepStatement.setNull(5, OracleTypes.CHAR);
// run the PreparedStatement
myPrepStatement.execute();
System.out.println("Added row to customers table");
// retrieve the ROWID, customer id, first name, last name, dob, and
// phone columns for this new row using an OracleResultSet
// object
Statement myStatement = myConnection.createStatement();
OracleResultSet customerResultSet =
  (OracleResultSet) myStatement.executeQuery(
    "SELECT ROWID, customer id, first name, last name, dob, phone " +
    "FROM customers " +
    "WHERE customer id = 6"
 );
System.out.println("Retrieved row from customers table");
// declare an oracle.sql.ROWID object to store the ROWID, and
// an oracle.sql.CHAR object to store the phone column
oracle.sql.ROWID rowid;
oracle.sql.CHAR phone = new oracle.sql.CHAR("", myCharSet);
```

```
// display the column values for row using the
    // get methods to read the values
    while (customerResultSet.next()) {
     rowid = customerResultSet.getROWID("ROWID");
     customerId = customerResultSet.getNUMBER("customer id");
     firstName = customerResultSet.getCHAR("first name");
     lastName = customerResultSet.getCHAR("last name");
     dob = customerResultSet.getDATE("dob");
     phone = customerResultSet.getCHAR("phone");
     System.out.println("rowid = " + rowid.stringValue());
     System.out.println("customerId = " + customerId.stringValue());
     System.out.println("firstName = " + firstName);
     System.out.println("lastName = " + lastName);
     System.out.println("dob = " + dob.stringValue());
      System.out.println("phone = " + phone);
    } // end of while loop
    // close the OracleResultSet object using the close() method
    customerResultSet.close();
    // roll back the changes made to the database
    myConnection.rollback();
    // close the other JDBC objects
    myPrepStatement.close();
   myConnection.close();
 } catch (SOLException e) {
   System.out.println("Error code = " + e.getErrorCode());
   System.out.println("Error message = " + e.getMessage());
   System.out.println("SQL state = " + e.getSQLState());
   e.printStackTrace();
} // end of main()
```

The output from this program is as follows:

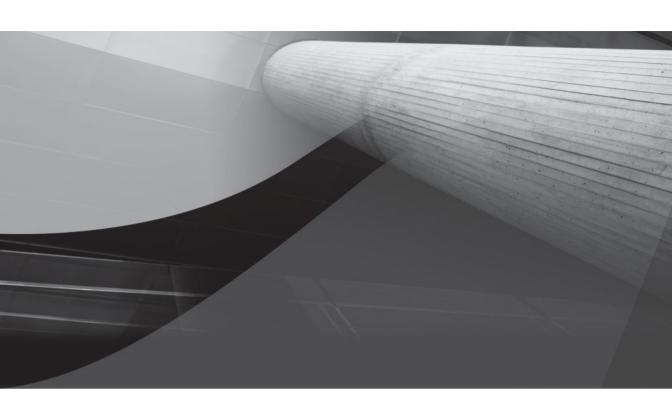
```
customerIdInt = 6
    firstNameString = Jason
    lastName = Price
    dobString = 2/22/1969 13:54:12
    Added row to customers table
    Retrieved row from customers table
    rowid = AAARk5AAEAAAAGPAAF
    customerId = 6
    firstName = Jason
    lastName = Price
    dob = 2/22/1969 13:54:12
    phone = null
    dobString2 = 2/22/1969 0:0:0
```

Summary

In this chapter, you have learned the following:

- The JDBC API enables a Java program to access a database.
- The Oracle JDBC drivers are used to connect to an Oracle database.
- SQL statements may be executed using JDBC.
- Oracle has developed a number of extensions to standard JDBC that allow you to gain access to all of the Oracle database types.

In the next chapter, you'll learn how to tune your SQL statements for maximum performance.



CHAPTER 16

SQL Tuning



n this chapter, you will do the following:

- Learn about SQL tuning
- See SQL tuning tips that you can use to shorten the length of time your queries take to execute
- Learn about the Oracle optimizer
- See how to compare the cost of performing queries
- Examine optimizer hints
- Learn about some additional tuning tools

Introducing SQL Tuning

One of the main strengths of SQL is that you don't have to tell the database exactly how to obtain the data requested. You simply run a query specifying the information you want, and the database software figures out the best way to get it. Sometimes, you can improve the performance of your SQL statements by "tuning" them. In the following sections, you'll see tuning tips that can make your queries run faster; later, you'll see more advanced tuning techniques.

Use a WHERE Clause to Filter Rows

Many novices retrieve all the rows from a table when they only want one row (or a few rows). This is very wasteful. A better approach is to add a WHERE clause to a query. That way, you restrict the rows retrieved to just those actually needed.

For example, say you want the details for customer #1 and #2. The following query retrieves all the rows from the customers table in the store schema (wasteful):

```
-- BAD (retrieves all rows from the customers table) SELECT *
```

FROM customers;

CUSTOMER_ID	FIRST_NAME	LAST_NAME	DOB	PHONE
1	John	Brown	01-JAN-65	800-555-1211
2	Cynthia	Green	05-FEB-68	800-555-1212
3	Steve	White	16-MAR-71	800-555-1213
4	Gail	Black		800-555-1214
5	Doreen	Blue	20-MAY-70	

The next query adds a WHERE clause to the previous example to just get customer #1 and #2:

```
-- GOOD (uses a WHERE clause to limit the rows retrieved)

SELECT *

FROM customers

WHERE customer id IN (1, 2);
```

```
CUSTOMER ID FIRST NAME LAST NAME DOB PHONE
-----
       1 John Brown 01-JAN-65 800-555-1211 2 Cynthia Green 05-FEB-68 800-555-1212
```

You should avoid using functions in the WHERE clause, as that increases execution time.

Use Table Joins Rather than Multiple Queries

If you need information from multiple related tables, you should use join conditions rather than multiple queries. In the following bad example, two queries are used to get the product name and the product type name for product #1 (using two queries is wasteful). The first query gets the name and product type id column values from the products table for product #1. The second query then uses that product type id to get the name column from the product types table.

```
-- BAD (two separate queries when one would work)
    SELECT name, product type id
    FROM products
    WHERE product id = 1;
                        PRODUCT TYPE ID
    ______
    Modern Science
    SELECT name
    FROM product types
    WHERE product type id = 1;
    NAME
    Book
```

Instead of using the two queries, you should write one query that uses a join between the products and product types tables. The following good query shows this:

```
-- GOOD (one query with a join)
     SELECT p.name, pt.name
     FROM products p, product types pt
     WHERE p.product type id = pt.product type id
    AND p.product id = 1;
    NAME
                              NAME
     ______
     Modern Science
                              Book
```

This query results in the same product name and product type name being retrieved as in the first example, but the results are obtained using one query. One query is generally more efficient than two.

You should choose the join order in your query so that you join fewer rows to tables later in the join order. For example, say you were joining three related tables named tab1, tab2, and tab3. Assume tab1 contains 1,000 rows, tab2 100 rows, and tab3 10 rows. You should join tab1 with tab2 first, followed by tab2 and tab3.

Also, avoid joining complex views in your queries, because doing so causes the queries for the views to be run first, followed by your actual query. Instead, write your query using the tables rather than the views.

Use Fully Qualified Column References When Performing Joins

Always include table aliases in your queries and use the alias for each column in your query (this is known as "fully qualifying" your column references). That way, the database doesn't have to search for each column in the tables used in your query.

The following bad example uses the aliases p and pt for the products and product_types tables, respectively, but the query doesn't fully qualify the description and price columns:

```
-- BAD (description and price columns not fully qualified)

SELECT p.name, pt.name, description, price

FROM products p, product_types pt

WHERE p.product_type_id = pt.product_type_id

AND p.product_id = 1;
```

MAPIE	IVAPID	
DESCRIPTION		PRICE
Modern Science	Book	
A description of modern science	е	19.95

This example works, but the database has to search both the products and product_types tables for the description and price columns; that's because there's no alias that tells the database which table those columns are in. The extra time spent by the database having to do the search is wasted time.

The following good example includes the table alias p to fully qualify the description and price columns:

```
-- GOOD (all columns are fully qualified)

SELECT p.name, pt.name, p.description, p.price

FROM products p, product_types pt

WHERE p.product_type_id = pt.product_type_id

AND p.product_id = 1;
```

NAME	NAME	
DESCRIPTION		PRICE
Modern Science A description of modern science	Book ce	19.95

Because all references to columns include a table alias, the database doesn't have to waste time searching the tables for the columns, and execution time is reduced.

Use CASE Expressions Rather than Multiple Queries

Use CASE expressions rather than multiple queries when you need to perform many calculations on the same rows in a table. The following bad example uses multiple queries to count the number of products within various price ranges:

```
-- BAD (three separate gueries when one CASE statement would work)
     SELECT COUNT(*)
     FROM products
     WHERE price < 13;
       COUNT (*)
     _____
     SELECT COUNT (*)
     FROM products
     WHERE price BETWEEN 13 AND 15;
      COUNT(*)
     _____
     SELECT COUNT(*)
     FROM products
     WHERE price > 15;
      COUNT(*)
```

Rather than using three queries, you should write one query that uses CASE expressions. This is shown in the following good example:

```
-- GOOD (one query with a CASE expression)
    COUNT(CASE WHEN price < 13 THEN 1 ELSE null END) low,
    COUNT (CASE WHEN price BETWEEN 13 AND 15 THEN 1 ELSE null END) med,
    COUNT (CASE WHEN price > 15 THEN 1 ELSE null END) high
    FROM products;
         LOW MED HIGH
    -----
```

Notice that the counts of the products with prices less than \$13 are labeled as low, products between \$13 and \$15 are labeled med, and products greater than \$15 are labeled high.



NOTE

You can, of course, use overlapping ranges and different functions in your CASE expressions.

Add Indexes to Tables

When looking for a particular topic in a book, you can either scan the whole book or use the index to find the location. An index for a database table is similar in concept to a book index, except that database indexes are used to find specific rows in a table. The downside of indexes is that when a row is added to the table, additional time is required to update the index for the new row.

Generally, you should create an index on a column when you are retrieving a small number of rows from a table containing many rows. A good rule of thumb is

Create an index when a guery retrieves <= 10 percent of the total rows in a table.

This means the column for the index should contain a wide range of values. A good candidate for indexing would be a column containing a unique value for each row (for example, a social security number). A poor candidate for indexing would be a column that contains only a small range of values (for example, N, S, E, W or 1, 2, 3, 4, 5, 6). An Oracle database automatically creates an index for the primary key of a table and for columns included in a unique constraint.

In addition, if your database is accessed using a lot of hierarchical queries (that is, a query containing a CONNECT BY), you should add indexes to the columns referenced in the START WITH and CONNECT BY clauses (see Chapter 7 for details on hierarchical queries).

Finally, for a column that contains a small range of values and is frequently used in the WHERE clause of queries, you should consider adding a bitmap index to that column. Bitmap indexes are typically used in data warehouses, which are databases containing very large amounts of data. The data in a data warehouse is typically read using many queries, but the data is not modified by many concurrent transactions.

Normally, a database administrator is responsible for creating indexes. However, as an application developer, you'll be able to provide the DBA with feedback on which columns are good candidates for indexing, because you may know more about the application than the DBA. Chapter 10 covers indexes in depth, and you should re-read the section on indexes if necessary.

Use WHERE Rather than HAVING

You use the WHERE clause to filter rows; you use the HAVING clause to filter groups of rows. Because the HAVING clause filters groups of rows *after* they have been grouped together (which takes some time to do), you should first filter rows using a WHERE clause whenever possible. That way, you avoid the time taken to group the filtered rows together in the first place.

The following bad query retrieves the product_type_id and average price for products whose product type id is 1 or 2. To do this, the query performs the following:

- It uses the GROUP BY clause to group rows into blocks with the same product type id.
- It uses the HAVING clause to filter the returned results to those groups that have a product_type_id in 1 or 2 (this is bad, because a WHERE clause would work).

```
-- BAD (uses HAVING rather than WHERE)

SELECT product_type_id, AVG(price)

FROM products

GROUP BY product_type_id

HAVING product type id IN (1, 2);
```

```
PRODUCT TYPE ID AVG(PRICE)
           1 24.975
                 26.22
```

The following good query rewrites the previous example to use WHERE rather than HAVING to first filter the rows to those whose product type id is 1 or 2:

```
-- GOOD (uses WHERE rather than HAVING)
    SELECT product type id, AVG(price)
     FROM products
    WHERE product type id IN (1, 2)
     GROUP BY product type id;
     PRODUCT TYPE ID AVG(PRICE)
     _____
                1 24.975
2 26.22
```

Use UNION ALL Rather than UNION

You use UNION ALL to get all the rows retrieved by two queries, including duplicate rows; you use UNION to get all non-duplicate rows retrieved by the queries. Because UNION removes duplicate rows (which takes some time to do), you should use UNION ALL whenever possible.

The following bad query uses UNION (bad because UNION ALL would work) to get the rows from the products and more products tables; notice that all non-duplicate rows from products and more products are retrieved:

```
-- BAD (uses UNION rather than UNION ALL)
    SELECT product_id, product_type_id, name
    FROM products
    UNION
    SELECT prd id, prd type id, name
    FROM more products;
```

```
PRODUCT_ID PRODUCT_TYPE_ID NAME
-----
       1
                     1 Modern Science
                    1 Chemistry
       3
                     2 Supernova
       3
                      Supernova
                     2 Lunar Landing
       4
                     2 Tank War
       4
       5
                     2 Submarine
       5
                     2 Z Files
                     2 2412: The Return
       6
       7
                     3 Space Force 9
       8
                     3 From Another Planet
       9
                     4 Classical Music
      10
                     4 Pop 3
                     4 Creative Yell
      11
      12
                      My Front Line
```

The following good query rewrites the previous example to use UNION ALL; notice that all the rows from products and more products are retrieved, including duplicates:

```
-- GOOD (uses UNION ALL rather than UNION)
      SELECT product id, product type id, name
      FROM products
      UNION ALL
      SELECT prd id, prd type id, name
      FROM more products;
      PRODUCT ID PRODUCT TYPE ID NAME
                              1 Modern Science
              2
                             1 Chemistry
              3
                              2 Supernova
              4
                              2 Tank War
              5
                              2 Z Files
                              2 2412: The Return
               6
              7
                              3 Space Force 9
              8
                              3 From Another Planet
              9
                              4 Classical Music
             10
                              4 Pop 3
             11
                              4 Creative Yell
             12
                               Mv Front Line
                             1 Modern Science
              1
              2
                              1 Chemistry
                               Supernova
                              2 Lunar Landing
              4
                              2 Submarine
```

Use EXISTS Rather than IN

You use IN to check if a value is contained in a list. You use EXISTS to check for the existence of rows returned by a subquery. EXISTS is different from IN: EXISTS just checks for the existence of rows, whereas IN checks actual values. EXISTS typically offers better performance than IN with subqueries. Therefore, you should use EXISTS rather than IN whenever possible.

You should refer back to the section entitled "Using EXISTS and NOT EXISTS with a Correlated Subquery" in Chapter 6 for full details on when you should use EXISTS with a correlated subquery (an important point to remember is that correlated subqueries can resolve null values).

The following bad query uses IN (bad because EXISTS would work) to retrieve products that have been purchased:

```
2 Chemistry
3 Supernova
```

The following good query rewrites the previous example to use EXISTS:

```
-- GOOD (uses EXISTS rather than IN)
    SELECT product id, name
    FROM products outer
    WHERE EXISTS
      (SELECT 1
      FROM purchases inner
      WHERE inner.product id = outer.product id);
    PRODUCT ID NAME
    _____
            1 Modern Science
            2 Chemistry
            3 Supernova
```

Use EXISTS Rather than DISTINCT

You can suppress the display of duplicate rows using DISTINCT. You use EXISTS to check for the existence of rows returned by a subquery. Whenever possible, you should use EXISTS rather than DISTINCT, because DISTINCT sorts the retrieved rows before suppressing the duplicate rows.

The following bad query uses DISTINCT (bad because EXISTS would work) to retrieve products that have been purchased:

```
-- BAD (uses DISTINCT when EXISTS would work)
    SELECT DISTINCT pr.product id, pr.name
    FROM products pr, purchases pu
    WHERE pr.product id = pu.product id;
    PRODUCT ID NAME
    _____
            1 Modern Science
            2 Chemistry
            3 Supernova
```

The following good query rewrites the previous example to use EXISTS rather than DISTINCT:

```
-- GOOD (uses EXISTS rather than DISTINCT)
    SELECT product id, name
    FROM products outer
    WHERE EXISTS
      (SELECT 1
      FROM purchases inner
      WHERE inner.product id = outer.product id);
    PRODUCT ID NAME
    -----
           1 Modern Science
            2 Chemistry
           3 Supernova
```

Use GROUPING SETS Rather than CUBE

The GROUPING SETS clause typically offers better performance than CUBE. Therefore, you should use GROUPING SETS rather than CUBE wherever possible. This is fully covered in the section entitled "Using the GROUPING SETS Clause" in Chapter 7.

Use Bind Variables

The Oracle database software caches SQL statements; a cached SQL statement is reused if an identical statement is submitted to the database. When an SQL statement is reused, the execution time is reduced. However, the SQL statement must be *absolutely identical* in order for it to be reused. This means that

- All characters in the SQL statement must be the same.
- All letters in the SQL statement must be in the same case.
- All spaces in the SQL statement must be the same.

If you need to supply different column values in a statement, you can use bind variables instead of literal column values. You'll see examples that clarify these ideas next.

Non-Identical SQL Statements

In this section, you'll see some non-identical SQL statements. The following non-identical queries retrieve products #1 and #2:

```
SELECT * FROM products WHERE product_id = 1;
SELECT * FROM products WHERE product id = 2;
```

These queries are not identical, because the value 1 is used in the first statement, but the value 2 is used in the second.

The following non-identical queries have spaces in different positions:

```
SELECT * FROM products WHERE product_id = 1;
SELECT * FROM products WHERE product id = 1;
```

The following non-identical queries use a different case for some of the characters:

```
select * from products where product_id = 1;
SELECT * FROM products WHERE product id = 1;
```

Now that you've seen some non-identical statements, let's take a look at identical SQL statements that use bind variables.

Identical SQL Statements That Use Bind Variables

You can ensure that a statement is identical by using bind variables to represent column values. You create a bind variable using the SQL*Plus VARIABLE command. For example, the following command creates a variable named v_product_id of type NUMBER:

VARIABLE v_product_id NUMBER



NOTE

You can use the types shown in Table A-1 of the appendix to define the type of a bind variable.

You reference a bind variable in an SQL or PL/SQL statement using a colon followed by the variable name (such as :v product id). For example, the following PL/SQL block sets v product id to 1:

BEGIN :v product id := 1; END; /

The following query uses v product id to set the product id column value in the WHERE clause; because v product id was set to 1 in the previous PL/SQL block, the query retrieves the details of product #1:

SELECT * FROM products WHERE product id = :v product id;

```
PRODUCT_ID PRODUCT_TYPE_ID NAME
DESCRIPTION
                                 PRICE
1 1 Modern Science
A description of modern science
                                 19.95
```

The next example sets v product id to 2 and repeats the query:

```
BEGIN
```

```
:v product id := 2;
END:
SELECT * FROM products WHERE product_id = :v_product_id;
PRODUCT ID PRODUCT TYPE ID NAME
_____
1 Chemistry
Introduction to Chemistry
                                   30
```

Because the query used in this example is identical to the previous query, the cached query is reused and there's an improvement in performance.



You should typically use bind variables if you're performing the same query many times. Also, in the example, the bind variables are session specific and need to be reset if the session is lost.

Listing and Printing Bind Variables

You list bind variables in SQL*Plus using the VARIABLE command. For example:

VARIABLE

```
variable v_product_id
datatype NUMBER
```

You display the value of a bind variable in SQL*Plus using the PRINT command. For example:

Using a Bind Variable to Store a Value Returned by a PL/SQL Function

You can also use a bind variable to store returned values from a PL/SQL function. The following example creates a bind variable named v_average_product_price and stores the result returned by the function average_product_price() (this function was described in Chapter 11 and calculates the average product price for the supplied product_type_id):

Using a Bind Variable to Store Rows from a REFCURSOR

You can also use a bind variable to store returned values from a REFCURSOR (a REFCURSOR is a pointer to a list of rows). The following example creates a bind variable named v_products_refcursor and stores the result returned by the function product_package.get_products_ref_cursor() (this function was introduced in Chapter 11; it returns a pointer to the rows in the products table):

3	Supernova	25.99
4	Tank War	13.95
5	Z Files	49.99
6	2412: The Return	14.95
7	Space Force 9	13.49
8	From Another Planet	12.99
9	Classical Music	10.99
10	Pop 3	15.99
11	Creative Yell	14.99
PRODUCT_ID	NAME	PRICE
12	My Front Line	13.49

Comparing the Cost of Performing Queries

The Oracle database software uses a subsystem known as the *optimizer* to generate the most efficient path to access the data stored in the tables. The path generated by the optimizer is known as an execution plan. Oracle Database 10g and above automatically gathers statistics about the data in your tables and indexes in order to generate the best execution plan (this is known as *cost-based* optimization).

Comparing the execution plans generated by the optimizer allows you to judge the relative cost of one SQL statement versus another. You can use the results to improve your SQL statements. In this section, you'll learn how to view and interpret a couple of example execution plans.



NOTE

Database versions prior to Oracle Database 10g don't automatically gather statistics, and the optimizer automatically defaults to rule-based optimization. Rule-based optimization uses syntactic rules to generate the execution plan. Cost-based optimization is typically better than rule-based optimization because the former uses actual information gathered from the data in the tables and indexes. If you're using Oracle Database 9i or below, you can gather statistics yourself (you'll learn how to do that later in the section "Gathering Table Statistics").

Examining Execution Plans

The optimizer generates an execution plan for an SQL statement. You can examine the execution plan using the SQL*Plus EXPLAIN PLAN command. The EXPLAIN PLAN command populates a table named plan table with the SQL statement's execution plan (plan table is often referred to as the "plan table"). You may then examine that execution plan by querying the plan table. The first thing you must do is check if the plan table currently exists in the database.

Checking if the Plan Table Currently Exists in the Database

To check if the plan table currently exists in the database, you should connect to the database as the store user and run the following DESCRIBE command:

SQL> DESCRIBE plan table

Name	Null?	Type
STATEMENT_ID		VARCHAR2(30)

PLAN ID	NUMBER
TIMESTAMP	DATE
REMARKS	VARCHAR2 (4000)
OPERATION	VARCHAR2 (30)
OPTIONS	VARCHAR2 (255)
OBJECT NODE	VARCHAR2 (128)
OBJECT OWNER	VARCHAR2 (30)
OBJECT NAME	VARCHAR2 (30)
OBJECT ALIAS	VARCHAR2 (65)
OBJECT INSTANCE	NUMBER (38)
OBJECT TYPE	VARCHAR2 (30)
OPTIMIZER	VARCHAR2 (255)
SEARCH COLUMNS	NUMBER
ID	NUMBER (38)
PARENT_ID	NUMBER (38)
DEPTH	NUMBER (38)
POSITION	NUMBER (38)
COST	NUMBER (38)
CARDINALITY	NUMBER (38)
BYTES	NUMBER (38)
OTHER TAG	VARCHAR2 (255)
PARTITION START	VARCHAR2 (255)
PARTITION STOP	VARCHAR2 (255)
PARTITION ID	NUMBER (38)
OTHER	LONG
OTHER XML	CLOB
DISTRIBUTION	VARCHAR2(30)
CPU COST	NUMBER (38)
IO COST	NUMBER (38)
TEMP SPACE	NUMBER (38)
ACCESS PREDICATES	VARCHAR2 (4000)
FILTER PREDICATES	VARCHAR2 (4000)
PROJECTION	VARCHAR2 (4000)
TIME	NUMBER (38)
QBLOCK_NAME	VARCHAR2(30)

If you get a table description similar to these results, you have the plan table already. If you get an error, then you need to create the plan table.

Creating the Plan Table

If you don't have the plan table, you must create it. To do this, you run the SQL*Plus script utlxplan.sql (on my Windows computer, the script is located in the directory E:\oracle_11g\product\11.1.0\db_1\RDBMS\ADMIN). The following example shows the command to run the utlxplan.sql script:

SQL> @ E:\oracle_11g\product\11.1.0\db_1\RDBMS\ADMIN\utlxplan.sql



NOTE

You'll need to replace the directory path with the path for your environment

The most important columns in the plan table are shown in Table 16-1.

Creating a Central Plan Table

If necessary, a database administrator can create one central plan table. That way, individual users don't have to create their own plan tables. To do this, a database administrator performs the following steps:

- 1. Creates the plan table in a schema of their choice by running the utlxplan.sql script
- 2. Creates a public synonym for the plan table
- 3. Grants access on the plan table to the public role

Here is an example of these steps:

@ E:\oracle 11g\product\11.1.0\db 1\RDBMS\ADMIN\utlxplan.sql CREATE PUBLIC SYNONYM plan table FOR plan table; GRANT SELECT, INSERT, UPDATE, DELETE ON plan table TO PUBLIC;

Column	Description
statement_id	Name you assign to the execution plan.
operation	Database operation performed, which can be Scanning a table Scanning an index Accessing rows from a table by using an index Joining two tables together Sorting a row set
	For example, the operation for accessing a table is TABLE ACCESS.
options	Name of the option used in the operation. For example, the option for a complete scan is FULL.
object_name	Name of the database object referenced in the operation.
object_type	Attribute of object. For example, a unique index has the attribute of UNIQUE.
id	Number assigned to this operation in the execution plan.
parent_id	Parent number for the current step in the execution plan. The parent_id value relates to an id value from a parent step.
position	Processing order for steps that have the same parent_id.
cost	Estimate of units of work for operation. Cost-based optimization uses disk I/O, CPU usage, and memory usage as units of work. Therefore, the cost is an estimate of the number of disk I/Os and the amount of CPU and memory used in performing an operation.

Generating an Execution Plan

Once you have a plan table, you can use the EXPLAIN PLAN command to generate an execution plan for an SQL statement. The syntax for the EXPLAIN PLAN command is as follows:

EXPLAIN PLAN SET STATEMENT_ID = statement_id FOR sql_statement;

where

- statement_id is the name you want to call the execution plan. This can be any alphanumeric text.
- sql statement is the SQL statement you want to generate an execution plan for.

The following example generates the execution plan for a query that retrieves all rows from the customers table (notice that the statement id is set to 'CUSTOMERS'):

```
EXPLAIN PLAN SET STATEMENT_ID = 'CUSTOMERS' FOR
SELECT customer_id, first_name, last_name FROM customers;
Explained
```

After the command completes, you may examine the execution plan stored in the plan table. You'll see how to do that next.



NOTE

The query in the EXPLAIN PLAN statement doesn't return rows from the customers table. The EXPLAIN PLAN statement simply generates the execution plan that would be used if the query was run.

Querying the Plan Table

For querying the plan table, I have provided an SQL*Plus script named $explain_plan.sql$ in the SQL directory. The script prompts you for the $statement_id$ and then displays the execution plan for that statement.

The explain plan.sql script is as follows:

-- Displays the execution plan for the specified statement id

```
UNDEFINE v_statement_id;

SELECT
  id ||
  DECODE(id, 0, '', LPAD(' ', 2*(level - 1))) || ' ' ||
  operation || ' ' ||
  options || ' ' ||
  object_name || ' ' ||
  object_type || ' ' ||
  DECODE(cost, NULL, '', 'Cost = ' || position)
AS execution_plan
  FROM plan_table
  CONNECT BY PRIOR id = parent_id
AND statement_id = '&&v_statement_id';
START WITH id = 0
AND statement_id = '&v_statement_id';
```

An execution plan is organized into a hierarchy of database operations similar to a tree; the details of these operations are stored in the plan table. The operation with an id of 0 is the root of the hierarchy, and all the other operations in the plan stem from this root. The guery in the script retrieves the details of the operations, starting with the root operation and then navigating the tree from the root.

The following example shows how to run the explain plan.sql script to retrieve the 'CUSTOMERS' plan created earlier:

```
SQL> @ c:\sql book\sql\explain plan.sql
    Enter value for v statement_id: CUSTOMERS
     old 12: statement id = '&&v statement id'
     new 12: statement id = 'CUSTOMERS'
     old 14: statement id = '&v statement id'
     new 14: statement id = 'CUSTOMERS'
    EXECUTION PLAN
     0 SELECT STATEMENT Cost = 3
      TABLE ACCESS FULL CUSTOMERS TABLE Cost = 1
```

The operations shown in the EXECUTION PLAN column are executed in the following order:

- The rightmost indented operation is executed first, followed by any parent operations above it.
- For operations with the same indentation, the topmost operation is executed first, followed by any parent operations above it.

Each operation feeds its results back up the chain to its immediate parent operation, and the parent operation is then executed. In the EXECUTION PLAN column, the operation ID is shown on the far left. In the example execution plan, operation 1 is run first, with the results of that operation being passed to operation 0. The following example illustrates the ordering for a more complex example:

```
0 SELECT STATEMENT Cost = 6
      MERGE JOIN Cost = 1
         TABLE ACCESS BY INDEX ROWID PRODUCT TYPES TABLE Cost = 1
            INDEX FULL SCAN PRODUCT TYPES PK INDEX (UNIQUE) Cost = 1
     4
          SORT JOIN Cost = 2
            TABLE ACCESS FULL PRODUCTS TABLE Cost = 1
```

The order in which the operations are executed in this example is 3, 2, 5, 4, 1, and 0. Now that you've seen the order in which operations are executed, it's time to move onto what the operations actually do. The execution plan for the 'CUSTOMERS' query was

```
0 SELECT STATEMENT
                      Cost = 3
       TABLE ACCESS FULL CUSTOMERS TABLE Cost = 1
```

Operation 1 is run first, with the results of that operation being passed to operation 0. Operation 1 involves a full table scan—indicated by the string TABLE ACCESS FULL—on the customers table. Here's the original command used to generate the 'CUSTOMERS' query:

```
EXPLAIN PLAN SET STATEMENT ID = 'CUSTOMERS' FOR
    SELECT customer id, first name, last name FROM customers;
```

A full table scan is performed because the SELECT statement specifies that all the rows from the customers table are to be retrieved.

The total cost of the query is three work units, as indicated in the cost part shown to the right of operation 0 in the execution plan (0 SELECT STATEMENT Cost = 3). A work unit is the amount of processing the software has to do to perform a given operation. The higher the cost, the more work the database software has to do to complete the SQL statement.



NOTE

If you're using a version of the database prior to Oracle Database 10g, then the output for the overall statement cost may be blank. That's because earlier database versions don't automatically collect table statistics. In order to gather statistics, you have to use the ANALYZE command. You'll learn how to do that later in the section "Gathering Table Statistics."

Execution Plans Involving Table Joins

Execution plans for queries with table joins are more complex. The following example generates the execution plan for a query that joins the products and product_types tables:

```
EXPLAIN PLAN SET STATEMENT_ID = 'PRODUCTS' FOR SELECT p.name, pt.name

FROM products p, product_types pt

WHERE p.product type id = pt.product type id;
```

The execution plan for this query is shown in the following example:

```
@ c:\sql_book\sql\explain_plan.sql
Enter value for v_statement_id: PRODUCTS
```

SORT JOIN Cost = 2

```
0 SELECT STATEMENT Cost = 6

1 MERGE JOIN Cost = 1

2 TABLE ACCESS BY INDEX ROWID PRODUCT_TYPES TABLE Cost = 1

3 INDEX FULL SCAN PRODUCT TYPES PK INDEX (UNIQUE) Cost = 1
```



NOTE

EXECUTION PLAN

If you run the example, you may get a slightly different execution plan depending on the version of the database you are using and on the settings of the parameters in the database's init.ora configuration file.

TABLE ACCESS FULL PRODUCTS TABLE Cost = 1

The previous execution plan is more complex, and you can see the hierarchical relationships between the various operations. The execution order of the operations is 3, 2, 5, 4, 1, and 0. Table 16-2 describes each operation in the order they are performed.

Operation ID	Description
3	Full scan of the index product_types_pk (which is a unique index) to obtain the addresses of the rows in the product_types table. The addresses are in the form of ROWID values, which are passed to operation 2.
2	Access the rows in the product_types table using the list of ROWID values passed from operation 3. The rows are passed to operation 1.
5	Access the rows in the products table. The rows are passed to operation 4.
4	Sort the rows passed from operation 5. The sorted rows are passed to operation 1.
1	Merge the rows passed from operations 2 and 5. The merged rows are passed to operation 0.
0	Return the rows from operation 1 to the user. The total cost of the query is 6 work units.

 TABLE 16-2
 Execution Plan Operations

Gathering Table Statistics

If you're using a version of the database prior to Oracle Database 10g (such as 9i), then you'll have to gather table statistics yourself using the ANALYZE command. By default, if no statistics are available then rule-based optimization is used. Rule-based optimization isn't usually as good as cost-based optimization.

The following examples use the ANALYZE command to gather statistics for the products and product types tables:

```
ANALYZE TABLE products COMPUTE STATISTICS;
ANALYZE TABLE product types COMPUTE STATISTICS;
```

Once the statistics have been gathered, cost-based optimization will be used rather than rulebased optimization.

Comparing Execution Plans

By comparing the total cost shown in the execution plan for different SQL statements, you can determine the value of tuning your SQL. In this section, you'll see how to compare two execution plans and see the benefit of using EXISTS rather than DISTINCT (a tip I gave earlier). The following example generates an execution plan for a query that uses EXISTS:

```
EXPLAIN PLAN SET STATEMENT ID = 'EXISTS QUERY' FOR
   SELECT product id, name
   FROM products outer
   WHERE EXISTS
     (SELECT 1
     FROM purchases inner
     WHERE inner.product id = outer.product id);
```

The execution plan for this query is shown in the following example:

```
@ c:\sql_book\sql\explain_plan.sql
Enter value for v_statement_id: EXISTS_QUERY

EXECUTION_PLAN

O SELECT STATEMENT Cost = 4

MERGE JOIN SEMI Cost = 1

TABLE ACCESS BY INDEX ROWID PRODUCTS TABLE Cost = 1

INDEX FULL SCAN PRODUCTS_PK INDEX (UNIQUE) Cost = 1

SORT UNIQUE Cost = 2

INDEX FULL SCAN PURCHASES PK INDEX (UNIQUE) Cost = 1
```

As you can see, the total cost of the query is 4 work units. The next example generates an execution plan for a query that uses DISTINCT:

```
EXPLAIN PLAN SET STATEMENT_ID = 'DISTINCT_QUERY' FOR SELECT DISTINCT pr.product_id, pr.name FROM products pr, purchases pu WHERE pr.product id = pu.product id;
```

The execution plan for this query is shown in the following example:

```
@ c:\sql_book\sql\explain_plan.sql
Enter value for v_statement_id: DISTINCT_QUERY

EXECUTION_PLAN

O SELECT STATEMENT Cost = 5

HASH UNIQUE Cost = 1

MERGE JOIN Cost = 1

TABLE ACCESS BY INDEX ROWID PRODUCTS TABLE Cost = 1

INDEX FULL SCAN PRODUCTS_PK INDEX (UNIQUE) Cost = 1

SORT JOIN Cost = 2

INDEX FULL SCAN PURCHASES PK INDEX (UNIQUE) Cost = 1
```

The cost for the query is 5 work units. This query is more costly than the earlier query that used EXISTS (that query had a cost of only 4 work units). These results prove it is better to use EXISTS than DISTINCT.

Passing Hints to the Optimizer

You can pass hints to the optimizer. A hint is an optimizer directive that influences the optimizer's choice of execution plan. The correct hint may improve the performance of an SQL statement. You can check the effectiveness of a hint by comparing the cost in the execution plan of an SQL statement with and without the hint.

In this section, you'll see an example query that uses one of the more useful hints: the $FIRST_ROWS(n)$ hint. The $FIRST_ROWS(n)$ hint tells the optimizer to generate an execution plan that will minimize the time taken to return the first n rows in a query. This hint can be useful when you don't want to wait around too long before getting *some* rows back from your query, but you still want to see all the rows.

The following example generates an execution plan for a query that uses FIRST ROWS (2); notice that the hint is placed within the strings /*+ and */:

```
EXPLAIN PLAN SET STATEMENT ID = 'HINT' FOR
     SELECT /*+ FIRST ROWS(2) */ p.name, pt.name
     FROM products p, product types pt
     WHERE p.product type id = pt. product type id;
```



CAUTION

EXECUTION PLAN

Your hint must use the exact syntax shown—otherwise, the hint will be ignored. The syntax is: / *+ followed by one space, the hint, followed by one space, and */.

The execution plan for this query is shown in the following example; notice that the cost is 4 work units:

```
@ c:\sql book\sql\explain plan.sql
```

```
Enter value for v statement id: HINT
```

```
0 SELECT STATEMENT Cost = 4
1 NESTED LOOPS
2 NESTED LOOPS Cost = 1
     TABLE ACCESS FULL PRODUCTS TABLE Cost = 1
```

INDEX UNIQUE SCAN PRODUCT TYPES PK INDEX (UNIQUE) Cost = 2 TABLE ACCESS BY INDEX ROWID PRODUCT TYPES TABLE Cost = 2

The next example generates an execution plan for the same query without the hint:

```
EXPLAIN PLAN SET STATEMENT ID = 'NO HINT' FOR
    SELECT p.name, pt.name
    FROM products p, product types pt
    WHERE p.product type id = pt. product type id;
```

The execution plan for the query is shown in the following example; notice the cost is 6 work units (higher than the query with the hint):

@ c:\sql book\sql\explain plan.sql

```
Enter value for v statement id: NO HINT
```

```
EXECUTION PLAN
______
0 SELECT STATEMENT Cost = 6
1 MERGE JOIN Cost = 1
   TABLE ACCESS BY INDEX ROWID PRODUCT TYPES TABLE Cost = 1
    INDEX FULL SCAN PRODUCT TYPES PK INDEX (UNIQUE) Cost = 1
   SORT JOIN Cost = 2
      TABLE ACCESS FULL PRODUCTS TABLE Cost = 1
```

These results show that the inclusion of the hint reduces the cost of running the query by 2 work units.

There are many hints that you can use, and this section has merely given you a taste of the subject.

Additional Tuning Tools

In this final section, I'll mention some other tuning tools. Full coverage of these tools is beyond the scope of this book. You can read the *Oracle Database Performance Tuning Guide*, published by Oracle Corporation, for full details of the tools mentioned in this section and for a comprehensive list of hints.

Oracle Enterprise Manager Diagnostics Pack

The Oracle Enterprise Manager Diagnostics Pack captures operating system, middle tier, and application performance data, as well as database performance data. The Diagnostics Pack analyzes this performance data and displays the results graphically. A database administrator can also configure the Diagnostics Pack to alert them immediately of performance problems via e-mail or page. Oracle Enterprise Manager also includes software guides to help resolve performance problems.

Automatic Database Diagnostic Monitor

The Automatic Database Diagnostic Monitor (ADDM) is a self-diagnostic module built into the Oracle database software. ADDM enables a database administrator to monitor the database for performance problems by analyzing system performance over a long period of time. The database administrator can view the performance information generated by ADDM in Oracle Enterprise Manager. When ADDM finds performance problems, it will suggest solutions for corrective action. Some example ADDM suggestions include

- Hardware changes—for example, adding CPUs to the database server
- Database configuration—for example, changing the database initialization parameter settings
- Application changes—for example, using the cache option for sequences or using bind variables
- Use other advisors—for example, running the SQL Tuning Advisor and SQL Access Advisor on SQL statements that are consuming the most database resources to execute

You'll learn about the SQL Tuning Advisor and SQL Access Advisor next.

SQL Tuning Advisor

The SQL Tuning Advisor allows a developer or database administrator to tune an SQL statement using the following items:

- The text of the SQL statement
- The SQL identifier of the statement (obtained from the V\$SQL_PLAN view, which is one of the views available to a database administrator)
- The range of snapshot identifiers
- The SQL Tuning Set name

An SQL Tuning Set is a set of SQL statements with their associated execution plan and execution statistics. SQL Tuning Sets are analyzed to generate SQL Profiles that help the optimizer to choose the optimal execution plan. SQL Profiles contain collections of information that enable optimization of the execution plan.

SQL Access Advisor

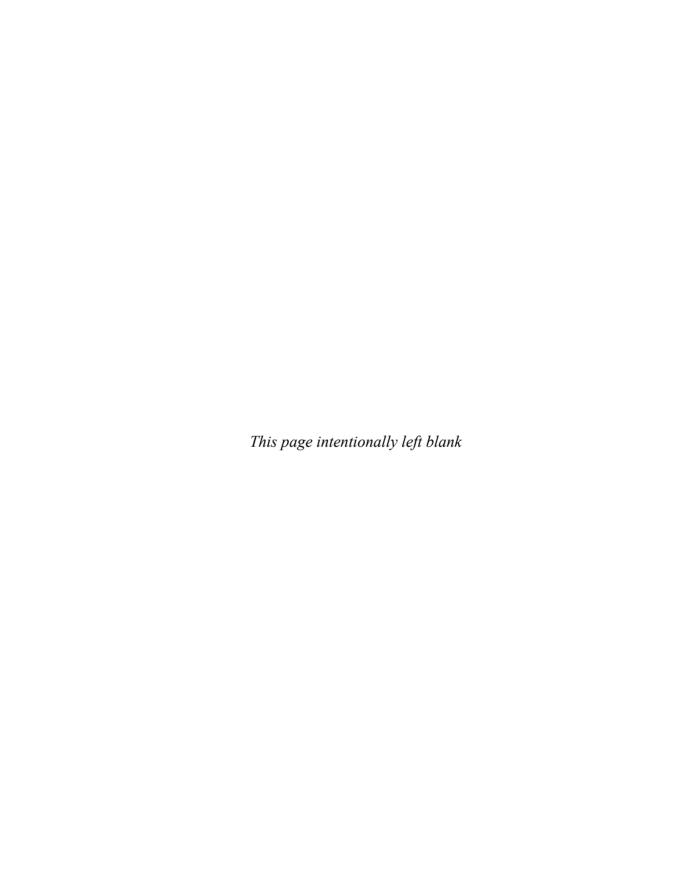
The SQL Access Advisor provides a developer or database administrator with performance advice on materialized views, indexes, and materialized view logs. The SQL Access Advisor examines space usage and query performance and recommends the most cost-effective configuration of new and existing materialized views and indexes.

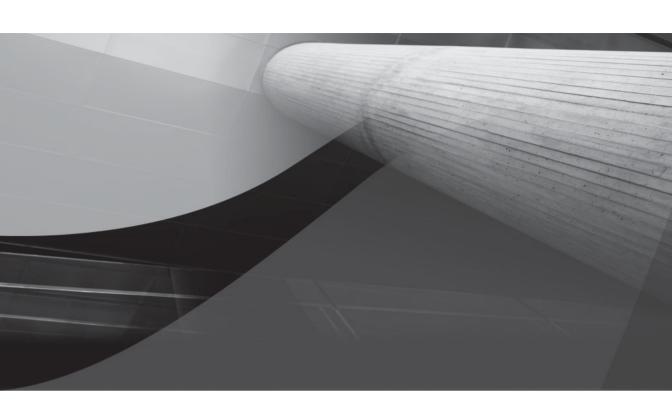
Summary

In this chapter, you have learned the following:

- Tuning is the process of making your SQL statements run faster.
- The optimizer is a subsystem of the Oracle database software that generates an execution plan, which is a set of operations used to perform a particular SQL statement.
- Hints may be passed to the optimizer to influence the generated execution plan for an SOL statement.
- There are a number of additional software tools a database administrator can use to tune the database.

In the next chapter, you'll learn about XML.





CHAPTER 17

XML and the Oracle Database



n this chapter, you will do the following:

- Become introduced to XML
- See how to generate XML from relational data
- Examine how to save XML in the database

Introducing XML

The Extensible Markup Language (XML) is a general-purpose markup language. XML enables you to share structured data across the Internet, and can be used to encode data and other documents. Some advantages of XML include the following:

- XML can be read by humans and computers, and is stored as plain text.
- XML is platform independent.
- XML supports Unicode, which means it can store information written in many human languages.
- XML uses a self-documenting format that contains the document structure, element names, and element values.

Because of these advantages, XML is widely used for document storage and processing, and it is used by many organizations to send data between their computer systems. For example, many suppliers allow their customers to send purchase orders as XML files over the Internet.

Oracle Database 9*i* introduced the ability to store XML in the database, along with extensive functionality for manipulating and processing XML. Oracle Database 10*g* Release 2 added additional XML-generating functions, and Oracle Database 11*g* adds capabilities like Java and C processing of binary XML (binary XML provides more efficient storage and manipulation of XML in the database). This chapter focuses on a useful subset of the XML capabilities in the Oracle database.

If you are new to XML, you will find a wealth of information at the following websites:

- http://www.w3.org/XML
- http://www.wikipedia.org/wiki/XML

Generating XML from Relational Data

The Oracle database contains a number of SQL functions you can use for generating XML, and in this section you'll see how to generate XML from relational data using some of these functions.

XMLELEMENT()

You use the XMLELEMENT () function to generate XML elements from relational data. You supply a name for the element, plus the column you wish to retrieve to XMLELEMENT(), and it returns the elements as XMLType objects. The XMLType is a built-in Oracle database type that is used to represent XML data. By default, an XMLType object stores the XML data as text in a CLOB (character large object).

The following example connects as the store user and gets the customer id column values as XMLType objects:

```
CONNECT store/store password
    SELECT XMLELEMENT ("customer id", customer id)
    AS xml customers
    FROM customers:
    XML CUSTOMERS
    <customer id>1</customer id>
    <customer id>2</customer id>
    <customer id>3</customer id>
    <customer id>4</customer id>
    <customer id>5</customer id>
```

As you can see from these results, XMLELEMENT ("customer id", customer id) returns the customer id values within a customer id tag. You can use whatever tag name you want, as shown in the following example which uses the tag "cust id":

```
SELECT XMLELEMENT ("cust id", customer id)
    AS xml customers
    FROM customers;
    XML CUSTOMERS
    <cust id>1</cust id>
    <cust id>2</cust id>
    <cust id>3</cust id>
    <cust id>4</cust id>
    <cust id>5</cust id>
```

The next example gets the first name and dob values for customer #2:

```
SELECT XMLELEMENT("first name", first name) || XMLELEMENT("dob", dob)
   AS xml customer
   FROM customers
   WHERE customer id = 2;
   XML CUSTOMER
   <first name>Cynthia</first name><dob>1968-02-05</dob>
```

The following example uses the TO_CHAR() function to change the date format for the dob value:

The next example embeds two calls to XMLELEMENT() within an outer call to XMLELEMENT(); notice that the returned customer_id and name elements are contained within an outer customer element:

```
SELECT XMLELEMENT (
     "customer",
     XMLELEMENT ("customer id", customer id),
     XMLELEMENT("name", first name || ' ' || last name)
   AS xml customers
   FROM customers
   WHERE customer id IN (1, 2);
   XML CUSTOMERS
   _____
   <customer>
     <customer id>1</customer id>
     <name>John Brown</name>
   </customer>
   <customer>
     <customer id>2</customer id>
     <name>Cynthia Green</name>
   </customer>
```

NOTE

I've added some line breaks and spaces in the XML returned by this query to make the XML easier to read. I've done the same thing in some of the other examples in this chapter.

You can retrieve regular relational data as well as XML, as shown in the following example, which retrieves the <code>customer_id</code> column as a regular relational result and the <code>first_name</code> and <code>last_name</code> columns concatenated together as XML elements:

```
SELECT customer_id,

XMLELEMENT("customer", first_name || ' ' || last_name) AS xml_customer

FROM customers;
```

```
CUSTOMER ID XML CUSTOMER
          1 <customer>John Brown</customer>
          2 <customer>Cynthia Green</customer>
          3 <customer>Steve White</customer>
          4 <customer>Gail Black</customer>
          5 <customer>Doreen Blue</customer>
```

You can generate XML for database objects, as shown in the next example which connects as object user and gets the id and address columns for customer #1 in the object customers table (the address column stores an object of type t address):

```
CONNECT object user/object password
    SELECT XMLELEMENT("id", id) || XMLELEMENT("address", address)
    AS xml object customer
     FROM object customers
    WHERE id = \frac{1}{1};
    XML OBJECT CUSTOMER
     _____
     <id>1</id>
     <address>
      <T ADDRESS>
        <STREET>2 State Street</STREET>
        <CITY>Beantown</CITY>
        <STATE>MA</STATE>
        <ZIP>12345</ZIP>
      </T ADDRESS>
     </address>
```

You can generate XML for collections, as shown in the next example, which connects as collection user and gets the id and addresses columns for customer #1 stored in customers with nested table (the addresses column stores an object of type t nested table address, which is a nested table of t address objects):

```
CONNECT collection user/collection password
     SELECT XMLELEMENT("id", id) || XMLELEMENT("addresses", addresses)
     AS xml customer
     FROM customers_with_nested_table
     WHERE id = 1;
     XML CUSTOMER
     <id>1</id>
     <addresses>
       <T NESTED TABLE ADDRESS>
         <T ADDRESS>
           <STREET>2 State Street</STREET><CITY>Beantown</CITY>
           <STATE>MA</STATE><ZIP>12345</ZIP>
         </T ADDRESS>
```

XMLATTRIBUTES()

You use XMLATTRIBUTES() in conjunction with XMLELEMENT() to specify the attributes for the XML elements retrieved by XMLELEMENT(). The following example connects as the store user and uses XMLATTRIBUTES() to set attribute names for the customer_id, first_name, last_name, and dob elements:

```
CONNECT store/store password
    SELECT XMLELEMENT (
     "customer",
     XMLATTRIBUTES (
       customer id AS "id",
       first name || ' ' || last name AS "name",
       TO CHAR (dob, 'MM/DD/YYYY') AS "dob"
     )
    )
   AS xml customers
    FROM customers
   WHERE customer id IN (1, 2);
    XML CUSTOMERS
    ______
    <customer id="1" name="John Brown" dob="01/01/1965"></customer>
    <customer id="2" name="Cynthia Green" dob="02/05/1968"></customer>
```

Notice that the id, name, and dob attributes are returned inside customer.

XMLFOREST()

You use XMLFOREST() to generate a "forest" of XML elements. XMLFOREST() concatenates XML elements together without you having to use the concatenation operator || with multiple calls to XMLELEMENT(). The following example uses XMLFOREST() to get the customer_id, phone, and dob for customers #1 and #2:

```
SELECT XMLELEMENT(

"customer",

XMLFOREST(

customer_id AS "id",

phone AS "phone",

TO_CHAR(dob, 'MM/DD/YYYY') AS "dob"

)

)

AS xml customers
```

```
FROM customers
WHERE customer id IN (1, 2);
XML CUSTOMERS
-----
<customer>
 <id>1</id>
 <phone>800-555-1211</phone>
  <dob>01/01/1965</dob>
</customer>
<customer>
 <id>2</id>
  <phone>800-555-1212</phone>
  <dob>02/05/1968</dob>
</customer>
```

The following command sets the SQL*Plus LONG parameter to 500, so you can see all the XML returned by the subsequent queries (LONG controls the maximum length of text data displayed by SQL*Plus):

SET LONG 500

The following query places the customer name inside the customer element tag using XMLATTRIBUTES():

```
SELECT XMLELEMENT (
     "customer",
     XMLATTRIBUTES(first_name || ' ' || last_name AS "name"),
     XMLFOREST (phone AS "phone", TO CHAR (dob, 'MM/DD/YYYY') AS "dob")
   AS xml customers
   FROM customers
   WHERE customer id IN (1, 2);
   XML CUSTOMERS
   _____
   <customer name="John Brown">
     <phone>800-555-1211</phone>
     <dob>01/01/1965</dob>
   </customer>
   <customer name="Cynthia Green">
     <phone>800-555-1212</phone>
     <dob>02/05/1968</dob>
   </customer>
```

XMLAGG()

You use XMLAGG() to generate a forest of XML elements from a collection of XML elements. XMLAGG() is typically used for grouping XML together into a common list of items underneath one parent or for retrieving data from collections. You can use the GROUP BY clause of a query to group the retuned set of rows into multiple groups, and you can use an ORDER BY clause of XMLAGG() to sort the rows.

By default, ORDER BY sorts the results in ascending order, but you can add DESC after the list of columns to sort the rows in descending order. You can add ASC to explicitly indicate an ascending sort. You can also add NULLS LAST to put any null values at the end of the results.

The following example retrieves the customer first_name and last_name values and returns them in a list named customer_list; notice that ORDER BY is used with XMLAGG() to sort the results by the first_name column. I've added ASC to explicitly indicate an ascending sort:

```
SELECT XMLELEMENT (
     "customer list",
     XMLAGG (
       XMLELEMENT("customer", first name || ' ' || last name)
       ORDER BY first name ASC
     )
    )
   AS xml customers
    FROM customers
   WHERE customer id IN (1, 2);
   XML CUSTOMERS
    _____
    <customer list>
     <customer>Cynthia Green
     <customer>John Brown</customer>
    </customer list>
```

The next example retrieves the product_type_id and average price for each group of products; notice that the products are grouped by product_type_id using the GROUP BY clause of the query, and NULLS LAST is used in the ORDER BY clause of XMLAGG() to place the row with the null product type id at the end of the returned results:

```
SELECT XMLELEMENT (
      "product list",
     XMLAGG (
       XMLELEMENT (
         "product type and avg", product type id || ' ' || AVG(price)
       ORDER BY product type id NULLS LAST
      )
    )
    AS xml products
    FROM products
    GROUP BY product type id;
    XML PRODUCTS
    ______
    cproduct list>
     cproduct type and avg>1 24.975/product type and avg>
      cproduct type and avg>2 26.22/product type and avg>
      cproduct type and avg>3 13.24/product_type_and_avg>
```

```
cproduct type and avg>4 13.99/product type and avg>
 cproduct type and avg> 13.49/product type and avg>
</product list>
```



NOTE

You can also place the null row first by specifying NULLS FIRST in the ORDER BY clause of XMLAGG().

The next example retrieves the product type id and name for the products with product type id values of 1 and 2, and the products are grouped by product type id:

```
SELECT XMLELEMENT (
     "products in group",
     XMLATTRIBUTES (product type id AS "prd type id"),
     XMLAGG (
       XMLELEMENT ("name", name)
     )
   )
   AS xml products
   FROM products
   WHERE product type id IN (1, 2)
   GROUP BY product type id;
   XML PRODUCTS
   _____
   cproducts in group prd type id="1">
     <name>Modern Science</name>
     <name>Chemistry</name>
   </products in group>
   cproducts in group prd type id="2">
     <name>Supernova</name>
     <name>2412: The Return</name>
   </products_in_group>
```

The next example connects as collection user and retrieves the addresses for customer #1 from customers with nested table:

```
CONNECT collection user/collection password
   SELECT XMLELEMENT ("customer",
       XMLELEMENT ("addresses", addresses)
     )
   )
   AS xml_customer
   FROM customers_with_nested_table
   WHERE id = 1;
   XML CUSTOMER
   _____
   <customer>
     <addresses>
```

```
<T NESTED TABLE_ADDRESS>
      <T ADDRESS>
        <STREET>2 State Street</STREET>
        <CITY>Beantown</CITY>
        <STATE>MA</STATE>
        <ZIP>21345</ZIP>
      </T ADDRESS>
      <T ADDRESS>
        <STREET>4 Hill Street</STREET>
        <CITY>Lost Town</CITY>
        <STATE>CA</STATE>
        <ZIP>54321</ZIP>
      </T ADDRESS>
   </T NESTED TABLE ADDRESS>
 </addresses>
</customer>
```

XMLCOLATTVAL()

You use XMLCOLATTVAL() to create an XML fragment and then expand the resulting XML. Each XML fragment has the name column with the attribute name. You can use the AS clause to change the attribute name.

The following example connects as the store user retrieves the customer_id, dob, and phone values for customers #1 and #2:

```
CONNECT store/store password
     SELECT XMLELEMENT (
       "customer",
       XMLCOLATTVAL (
         customer id AS "id",
         dob AS "dob",
         phone AS "phone"
       )
     AS xml customers
     FROM customers
     WHERE customer id IN (1, 2);
     XML CUSTOMERS
     <customer>
       <column name = "id">1</column>
       <column name = "dob">1965-01-01</column>
       <column name = "phone">800-555-1211</column>
     </customer>
     <customer>
       <column name = "id">2</column>
       <column name = "dob">1968-02-05</column>
       <column name = "phone">800-555-1212</column>
     </customer>
```

XMLCONCAT()

You use XMLCONCAT () to concatenate a series of elements for each row. The following example concatenates the XML elements for first name, last name, and phone values for customers #1 and #2:

```
SELECT XMLCONCAT(
     XMLELEMENT ("first name", first name),
     XMLELEMENT ("last name", last name),
     XMLELEMENT ("phone", phone)
    )
   AS xml customers
    FROM customers
   WHERE customer id IN (1, 2);
   XML CUSTOMERS
    <first name>John</first name>
    <last name>Brown</last name>
    <phone>800-555-1211</phone>
    <first name>Cynthia</first name>
    <last name>Green
    <phone>800-555-1212</phone>
```

XMLPARSE()

You use XMLPARSE() to parse and generate XML from the evaluated result of an expression. The expression must resolve to a string; if the expression resolves to null, then XMLPARSE () returns null. You must specify one of the following items before the expression:

- CONTENT, which means the expression must resolve to a valid XML value
- DOCUMENT, which means the expression must resolve to a singly rooted XML document

You can also add WELLFORMED after the expression, which means you are guaranteeing that your expression resolves to a well-formed XML document. This also means that the database will not perform validity checks on your expression.

The following example parses an expression containing the details for a customer:

```
SELECT XMLPARSE (
      CONTENT
       '<customer><customer id>1</customer id><name>John Brown</name></customer>'
      WELLFORMED
    AS xml customer
    FROM dual:
    XML CUSTOMER
    <customer>
      <customer id>1</customer id>
       <name>John Brown</name>
     </customer>
```



NOTE

You can read more about well-formed XML documents and values at http://www.w3.org/TR/REC-xml.

XMLPI()

You use XMLPI() to generate an XML processing instruction. You typically use a processing instruction to provide an application with information that is associated with XML data; the application can then use the processing instruction to determine how to process the XML data. The following example generates a processing instruction for an order status:

The next example generates a processing instruction to display an XML document using a cascading stylesheet file named example.css:

XMLCOMMENT()

You use XMLCOMMENT () to generate an XML comment, which is a text string placed within <!-- and -->. For example:

XMLSEQUENCE()

You use XMLSEQUENCE () to generate an XMLSequenceType object, which is a varray of XMLType objects. Because XMLSEQUENCE () returns a varray, you can use it in the FROM clause of a guery. For example:

```
SELECT VALUE(list of values).GETSTRINGVAL() order values
     FROM TABLE (
       XMLSEOUENCE (
         EXTRACT (
           XMLType('<A><B>PLACED</B><B>PENDING</B><B>SHIPPED</B></A>'),
           '/A/B'
         )
       ١
     ) list_of_values;
     ORDER VALUES
     _____
     <B>PLACED</B>
     <B>PENDING</B>
     <B>SHIPPED</B>
```

Let's break down this example. The call to XMLType () is

```
XMLType('<A><B>PLACED</B><B>PENDING</B><B>SHIPPED</B></A>')
```

This creates an XMLType object containing the XML

```
<A><B>PLACED</B><B>PENDING</B><B>SHIPPED</B></A>.
```

The call to the EXTRACT () function is

```
EXTRACT (
      XMLType('<A><B>PLACED</B><B>PENDING</B><B>SHIPPED</B></A>'),
      '/A/R'
```

EXTRACT() extracts the XML data from the XMLType object returned by the call to XMLType(). The second parameter to EXTRACT() is an XPath string. XPath is a language that allows you access specific elements in XML data. For example, in the previous call to EXTRACT(), '/A/B' returns all the B elements that are children of the A elements; therefore, the EXTRACT () function returns the following:

```
<B>PLACED</B>
    <B>PENDING</B>
    <B>SHIPPED</B>
```

The call to XMLSEQUENCE () in the example simply returns a varray containing the elements returned by EXTRACT(). TABLE() converts the varray into a table of rows and applies the alias list of values to the table. The SELECT statement retrieves the string value of the rows in the table using GETSTRINGVAL().

You'll see more examples of EXTRACT () and XPath later in this chapter.

XMLSERIALIZE()

You use XMLSERIALIZE () to generate a string or LOB (large object) representation of XML data from the evaluated result of an expression. You must specify one of the following items before the expression:

- CONTENT, which means the expression must resolve to a valid XML value
- DOCUMENT, which means the expression must resolve to a singly rooted XML document

The following example uses XMLSERIALIZE () with CONTENT to generate an XML value:

The next example uses XMLSERIALIZE() with DOCUMENT to generate an XML document, with the document returned in a CLOB (character large object):

A PL/SQL Example That Writes XML Data to a File

In this section, you'll see a complete PL/SQL example that writes customer names to an XML file. First, you need to connect as a privileged user (for example, the system user) and grant the CREATE ANY DIRECTORY privilege to the store user:

```
CONNECT system/manager;
GRANT CREATE ANY DIRECTORY TO store;
```

Next, you need to connect as the store user and create a directory object:

```
CONNECT store/store_password;

CREATE DIRECTORY TEMP_FILES_DIR AS 'C:\temp_files';
```

You'll also need to create a directory named temp_files in the C partition. (If you're using Linux or Unix, you can create the directory on one of your partitions and use an appropriate CREATE DIRECTORY command with the correct path. Also, make sure you grant write permissions on the directory to the Oracle user account you used to install the database software.)

Next, you need to run the xml examples.sql script located in the SQL directory, as shown here.

@ "E:\Oracle SQL book\sql book\SQL\xml examples.sql"



CAUTION

Run only the xml examples.sql script at this point. You may notice there is a script named xml schema.sql in the SQL directory. Do not run xml schema.sql yet.

The xml examples.sql script creates two procedures; the one you'll see in this section is named write xml data to file(), which retrieves the customer names and writes them to an XML file. The write xml data to file () procedure is defined as follows:

```
CREATE PROCEDURE write xml data to file(
      p directory VARCHAR2,
      p file name VARCHAR2
    ) AS
      v file UTL FILE.FILE TYPE;
      v amount INTEGER := 32767;
      v xml data XMLType;
      v char buffer VARCHAR2(32767);
    BEGIN
      -- open the file for writing of text (up to v amount
      -- characters at a time)
      v file := UTL FILE.FOPEN(p directory, p file name, 'w', v amount);
       -- write the starting line to v file
      UTL FILE.PUT LINE(v file, '<?xml version="1.0"?>');
       -- retrieve the customers and store them in v xml data
       SELECT
        EXTRACT (
          XMLELEMENT (
             "customer list",
            XMLAGG (
               XMLELEMENT("customer", first name || ' ' || last name)
               ORDER BY last name
             )
          ),
           '/customer list'
         )
      AS xml customers
       INTO v xml data
       FROM customers;
       -- get the string value from v xml data and store it in v char buffer
      v char buffer := v xml data.GETSTRINGVAL();
       -- copy the characters from v char buffer to the file
      UTL FILE.PUT(v file, v char buffer);
```

```
-- flush any remaining data to the file
UTL_FILE.FFLUSH(v_file);

-- close the file
UTL_FILE.FCLOSE(v_file);
END write_xml_data_to_file;
/
The following statement calls write_xml_data_to_file():
```

```
CALL write xml data to file('TEMP FILES DIR', 'customers.xml');
```

After you run this statement, you'll find a file named customers.xml in C:\temp_files, or whichever directory you used when using the CREATE DIRECTORY command earlier. The contents of the customers.xml file is as follows:

```
<?xml version="1.0"?>
    <customer_list><customer>Gail Black</customer><customer>Doreen Blue
    </customer><customer>John Brown</customer><customer>Cynthia Green
    </customer><customer>Steve White</customer></customer list>
```

You can modify the write_xml_data_to_file() procedure to retrieve any relational data from the database and write it out to an XML file.

XMLQUERY()

You use XMLQUERY() to construct XML or query XML. You pass an XQuery expression to XMLQUERY(). XQuery is a query language that allows you to construct and query XML. XMLQUERY() returns the result of evaluating the XQuery expression.

The following simple example illustrates the use of XMLQUERY():

Here are some notes for the example:

- The string passed to XMLQUERY() is the XQuery expression, that is, the XQuery expression is (1, 2 + 5, "d", 155 to 161, <A>text). 1 is an integer literal, 2 + 5 is an arithmetic expression, d is a string literal, 155 to 161 is a sequence of integers, and <A>text is an XML element.
- Each of the items in the XQuery is evaluated in turn. For example, 2 + 5 is evaluated, and 7 is returned. Similarly, 155 to 161 is evaluated and 155 156 157 158 159 160 161 is returned.

RETURNING CONTENT means an XML fragment is returned. This XML fragment is a single XML element with any number of "children," which can themselves be of any XML element type, including text elements. The XML fragment also conforms to the extended Infoset data model. Infoset is a specification describing an abstract data model of an XML document. You can learn more about Infoset at http://www.w3.org/TR/xml-infoset.

Let's explore a more complex example. The following statement (contained in the xml examples.sql script) creates a procedure named create xml resources(); this procedure creates XML strings for products and product types. It uses methods in the PL/SQL DBMS XDB package to delete and create XML resource files in the Oracle XML DB Repository (the XML DB Repository is a storage area for XML data within the database):

```
CREATE PROCEDURE create xml resources AS
       v result BOOLEAN;
       -- create string containing XML for products
       v products VARCHAR2(300):=
         '<?xml version="1.0"?>' ||
         ''oducts>' ||
           'product product id="1" product type id="1" name="Modern Science"'
            || ' price="19.95"/>' ||
           'roduct product id="2" product type id="1" name="Chemistry"' | |
           ' price="30"/>' ||
           'roduct product id="3" product type id="2" name="Supernova"' ||
           ' price="25.99"/>' ||
         '</products>';
       -- create string containing XML for product types
       v product types VARCHAR2(300):=
         '<?xml version="1.0"?>' ||
         ''oduct types>' ||
           'oduct type product type id="1" name="Book"/>' ||
           'oduct type product type id="2" name="Video"/>' ||
         '</product types>';
     BEGIN
       -- delete any existing resource for products
       DBMS XDB.DELETERESOURCE('/public/products.xml',
         DBMS XDB.DELETE RECURSIVE FORCE);
       -- create resource for products
       v result := DBMS XDB.CREATERESOURCE('/public/products.xml',
         v products);
       -- delete any existing resource for product types
       DBMS XDB.DELETERESOURCE('/public/product types.xml',
         DBMS XDB.DELETE RECURSIVE FORCE);
       -- create resource for product types
       v result := DBMS XDB.CREATERESOURCE('/public/product types.xml',
         v product types);
     END create xml resources;
```

Here are some notes for create xml resources ():

- The DBMS_XDB.DELETERESOURCE() procedure deletes an XML resource from the database. This procedure is called by create_xml_resources() so that you don't have to manually remove the resources if you run create_xml_resources() more than once.
- The DBMS_XDB.DELETE_RECURSIVE_FORCE constant forces the deletion of the resource, including any child objects.
- The DBMS_XDB.CREATERESOURCE() function creates an XML resource in the database and returns a Boolean true/false value indicating whether the operation was successful. The two calls to this function create resources for the products and product types in / public, which is the absolute path to store the resources.

The following statement calls create xml resources ():

CALL create_xml_resources();

The following query uses XMLQUERY() to retrieve the products from the /public/products.xml resource:

The XQuery expression inside XMLQUERY () in the previous example is

for \$product in doc("/public/products.xml")/products/product return cproduct name="{\$product/@name}"/>

Let's break down this XQuery expression:

- The for loop iterates over the products in /public/products.xml.
- \$product is a binding variable that is bound to the sequence of products returned by doc("/public/products.xml")/products/product; doc("/public/products.xml") returns the products.xml document stored in /public. With each iteration of the loop, \$product is set to each product in products.xml, one after another.
- The return part of the expression returns the product name in \$product.

The next query retrieves the product types from the /public/product types.xml resource:

```
SELECT XMLQUERY(
     'for $product type in
      doc("/public/product types.xml")/product types/product type
      return cproduct type name="{$product type/@name}"/>'
     RETURNING CONTENT
   AS xml product types
   FROM DUAL:
   XML PRODUCT TYPES
    _____
    cproduct type name="Book">
    cproduct type name="Video">
```

The following query retrieves the products whose price is greater than 20, along with their product type:

```
SELECT XMLQUERY(
      'for $product in doc("/public/products.xml")/products/product
       let $product type :=
         doc("/public/product types.xml")//product type[@product type id =
           $product/@product type id]/@name
       where $product/@price > 20
       order by $product/@product id
       return cproduct name="{$product/@name}"
         product type="{$product type}"/>'
      RETURNING CONTENT
    )
    AS xml query results
    FROM DUAL;
    XML QUERY RESULTS
    cproduct name="Chemistry" product type="Book"></product>
    oduct name="Supernova" product type="Video">
```

Let's break down the XQuery expression in this example:

- Two binding variables are used: \$product and \$product type. These variables are used to store the products and product types.
- The let part of the expression sets \$product type to the product type retrieved from sproduct. The expression on the right-hand side of the := performs a join using the product type id value stored in \$product type and \$product. The // means retrieve all elements.
- The where part retrieves only products whose price is greater than 20.
- The order by part orders the results by the product ID (in ascending order by default).

The next example shows the use of the following XQuery functions:

- count (), which counts the number of objects passed to it.
- avg (), which calculates the average of the numbers passed to it.
- integer(), which truncates a number and returns the integer. The integer() function is in the xs namespace. (The count() and avg() functions are in the fn namespace, which is automatically referenced by the database, thereby allowing you to omit this namespace when calling these functions.)

The following example returns the product type name, the number of products in each product type, and the average price of the products in each product type (truncated to an integer):

```
SELECT XMLQUERY(
      'for $product type in
       doc("/public/product types.xml")/product types/product type
       let $product :=
        doc("/public/products.xml")//product[@product type id =
          $product type/@product type id]
       return
        cproduct type name="{$product type/@name}"
         num products="{count($product)}"
         average price="{xs:integer(avg($product/@price))}"
        />'
     RETURNING CONTENT
    AS xml_query_results
    FROM DUAL;
    XML QUERY RESULTS
    ______
    duct type name="Book" num products="2" average price="24">
    </product type>
    cproduct type name="Video" num products="1" average price="25">
    </product type>
```

As you can see from the results, there are two books and one video.



NOTE

You can read more about functions at http://www.w3.org/TR/xquery-operators. You can find more information on XMLQUERY () at http://www.sqlx.org.

Saving XML in the Database

In this section, you'll see how to store an XML document in the database and retrieve information from the stored XML.

The Example XML File

You'll see the use of a file named purchase order.xml, which is an XML file that contains a purchase order. This file is contained in the xml files directory, which is created when you extracted the Zip file for this book. If you want to follow along with the examples, you should copy the xml files directory to the C partition on your database server (if you're using Linux or Unix, you can copy the directory to one of your partitions).



NOTE

If you copy the xml files directory to a location different from C, then you'll need to edit the xml schema.sql script (you'll see this script shortly).

The contents of the purchase order.xml file is as follows:

```
<?xml version="1.0"?>
    <purchase order>
      <customer order id>176</customer order id>
      <order date>2007-05-17</order date>
      <customer name>Best Products 456 Inc.</customer name>
      <street>10 Any Street</street>
      <city>Any City</city>
      <state>CA</state>
      <zip>94440</zip>
      <phone number>555-121-1234</phone number>
      cproducts>
        oduct>
          cproduct id>1/product id>
          <name>Supernova video</name>
          <quantity>5</quantity>
        </product>
        cproduct>
          cproduct id>2
          <name>Oracle SQL book</name>
          <quantity>4</quantity>
        </product>
      </products>
    </purchase order>
```

In the following sections, you'll see how to store this XML file in the database.

In a real-world example, the purchase order could be sent via the Internet to an online store, which would then dispatch the requested items to the customer.

Creating the Example XML Schema

I've provided an SQL*Plus script named xml schema.sql in the SQL directory. The script creates a user named xml user with a password of xml password, and it creates the items used in the rest of this chapter. Don't run this script yet.

The script contains the following statements that create an object type named t_product (used to represent products), a nested table type named t_nested_table_product (used to represent a nested table of products), and a table named purchase order:

```
CREATE TYPE t product AS OBJECT (
     product id INTEGER,
     name VARCHAR2(15),
     quantity INTEGER
    );
    CREATE TYPE t nested table product AS TABLE OF t product;
   CREATE TABLE purchase order (
     purchase order id INTEGER CONSTRAINT purchase order pk PRIMARY KEY,
     customer order id INTEGER,
     order date DATE,
     customer name VARCHAR2(25),
     street VARCHAR2(15),
     city VARCHAR2 (15),
      state VARCHAR2(2),
     zip VARCHAR2(5),
     phone number VARCHAR2(12),
     products t nested table product,
     xml purchase order XMLType
   NESTED TABLE products
    STORE AS nested products;
```

Notice that the xml_purchase_order column is of type XMLType, which is a built-in Oracle database type that allows you to store XML data. By default, an XMLType column stores the XML data as text in a CLOB (character large object).

The xml_schema.sql script also contains the following statement that creates a directory object named XML FILES DIR:

```
CREATE OR REPLACE DIRECTORY XML_FILES_DIR AS 'C:\xml_files';
```

You'll need to modify this line if you copied the xml_files directory to a location different from C. If you need to modify this line, go ahead and do it now and then save the script.

The following INSERT statement (also contained in the script) adds a row to the purchase_order table:

```
INSERT INTO purchase_order (
   purchase_order_id,
   xml_purchase_order
) VALUES (
   1,
   XMLType(
      BFILENAME('XML_FILES_DIR', 'purchase_order.xml'),
      NLS_CHARSET_ID('AL32UTF8')
   )
);
```

As you can see, the XMLType () constructor accepts two parameters. The first parameter is a BFILE, which is a pointer to an external file. The second parameter is the character set for the XML text in the external file. In the previous INSERT, the BFILE points to the purchase order.xml file, and the character set is AL32UTF8, which is standard UTF-8 encoding. When the INSERT is run, the XML from the purchase order.xml file is read and then stored in the database as CLOB text in the xml purchase order column.



NOTE

When you are working with XML files written in English, you should typically use the AL32UTF8 character set. You can find more information about different character sets in the Oracle Database Globalization Support Guide published by Oracle Corporation.

You may have noticed the customer order id, order date, customer name, street, city, state, zip, phone number, and products columns in the purchase order table are empty. The data for these columns can be extracted from the XML stored in the xml purchase order column. Later in this chapter, you'll see a PL/SQL procedure that reads the XML and sets the other columns accordingly.

Go ahead and run the xml schema.sql script as a privileged user (such as the system user):

CONNECT system/manager

@ "E:\Oracle SQL book\sql book\SQL\xml schema.sql"

After the script completes, you will be logged in as xml user.

Retrieving Information from the Example XML Schema

In this section, you'll see how to retrieve information from the xml user schema. The following example retrieves the row from the purchase order table:

```
SET LONG 1000
     SET PAGESIZE 500
     SELECT purchase order id, xml purchase order
     FROM purchase_order;
     PURCHASE ORDER ID
     XML PURCHASE ORDER
                    1
     <?xml version="1.0"?>
     <purchase order>
       <customer order id>176</customer order id>
       <order date>2007-05-17</order date>
       <customer name>Best Products 456 Inc.</customer name>
       <street>10 Any Street</street>
       <city>Any City</city>
       <state>CA</state>
       <zip>94440</zip>
       <phone number>555-121-1234</phone_number>
       cproducts>
```

The next query extracts the customer_order_id, order_date, customer_name, and phone_number from the XML stored in the xml_purchase_order column using the EXTRACT() function:

SELECT

```
EXTRACT (xml purchase order,
   '/purchase order/customer order id') cust order id,
 EXTRACT(xml purchase order, '/purchase order/order date') order date,
 EXTRACT(xml purchase order, '/purchase order/customer name') cust name,
 EXTRACT(xml purchase order, '/purchase order/phone number') phone number
FROM purchase order
WHERE purchase order id = 1;
CUST ORDER ID
_____
ORDER DATE
CUST NAME
______
PHONE NUMBER
<customer order id>176</customer order id>
<order date>2007-05-17</order date>
<customer name>Best Products 456 Inc.</customer name>
<phone number>555-121-1234</phone number>
```

The EXTRACT () function returns the values as XMLType objects.

You can use the EXTRACTVALUE () function to get the values as strings. For example, the following query extracts the same values as strings using the EXTRACTVALUE () function:

SELECT

```
EXTRACTVALUE(xml_purchase_order,
  '/purchase_order/customer_order_id') cust_order_id,
EXTRACTVALUE(xml_purchase_order,
  '/purchase_order/order_date') order_date,
EXTRACTVALUE(xml_purchase_order,
  '/purchase_order/customer_name') cust_name,
EXTRACTVALUE(xml_purchase_order,
```

```
'/purchase order/phone number') phone number
FROM purchase order
WHERE purchase order id = 1;
CUST ORDER ID
_____
ORDER DATE
_____
CUST NAME
_____
PHONE NUMBER
176
2007-05-17
Best Products 456 Inc.
555-121-1234
```

The next query extracts and converts order date to a DATE using the TO DATE () function; notice that the format for the date as stored in the XML is supplied using the second parameter to TO_DATE() and that TO DATE() returns the date in the default date format used by the database (DD-MON-YY):

SELECT

```
EXTRACTVALUE(xml purchase order, '/purchase order/order date'),
    'YYYY-MM-DD'
  ) AS ord date
FROM purchase order
WHERE purchase order id = 1;
ORD DATE
_____
17-MAY-07
```

The following query retrieves all the products from xml purchase order as XML using EXTRACT(); notice the use of // to get all the products:

SELECT

```
EXTRACT (xml purchase order, '/purchase order//products') xml products
FROM purchase order
WHERE purchase order id = 1;
XML PRODUCTS
-----
cproducts>
 oduct>
   cproduct id>1
   <name>Supernova video</name>
   <quantity>5</quantity>
 </product>
 oduct>
```

The next query retrieves product #2 from xml_purchase_order; notice product [2] returns product #2:

The following query retrieves the "Supernova video" product from xml_purchase_order; notice that the name of the product to retrieve is placed inside square brackets:

You use the EXISTSNODE () function to check if an XML element exists. EXISTSNODE () returns 1 if the element exists; otherwise, it returns 0. For example, the following query returns the string 'Exists' because product #1 exists:

```
SELECT 'Exists' AS "EXISTS"

FROM purchase_order

WHERE purchase_order_id = 1

AND EXISTSNODE(
```

```
xml purchase order,
  '/purchase order/products/product[product id=1]'
) = 1;
EXISTS
_____
Exists
```

The next query returns no rows because product #3 does not exist:

```
SELECT 'Exists'
   FROM purchase order
   WHERE purchase order id = 1
   AND EXISTSNODE (
     xml purchase order,
     '/purchase order/products/product[product id=3]'
   no rows selected
```

The following query retrieves the products as a varray of XMLType objects using the XMLSEQUENCE () function; notice the use of product .* to retrieve all the products and their XML elements:

```
SELECT product.*
    FROM TABLE (
      SELECT
        XMLSEQUENCE(EXTRACT(xml_purchase_order, '/purchase_order//product'))
      FROM purchase order
      WHERE purchase order id = 1
    ) product;
    COLUMN VALUE
    coduct>
      cproduct id>1/product id>
      <name>Supernova video</name>
      <quantity>5</quantity>
    </product>
    oduct>
      cproduct id>2
      <name>Oracle SQL book</name>
      <quantity>4</quantity>
    </product>
```

The next query retrieves the product id, name, and quantity for the products as strings using the EXTRACTVALUE () function:

SELECT

```
EXTRACTVALUE (product.COLUMN_VALUE, '/product/product_id') AS product_id,
EXTRACTVALUE (product.COLUMN VALUE, '/product/name') AS name,
```

Updating Information in the Example XML Schema

The customer_order_id, order_date, customer_name, street, city, state, zip, phone_number, and products columns in the purchase_order table are empty. The data for these columns can be extracted from the XML stored in the xml_purchase_order column. In this section, you'll a PL/SQL procedure named update_purchase_order() that reads the XML and sets the other columns accordingly.

The most complex aspect of update_purchase_order() is the process of reading the products from the XML and storing them in the products nested table column of the purchase_order table. In this procedure, a cursor is used to read the products from the XML, then the XML is converted to strings using EXTRACTVALUE(), and the strings are stored in a nested table.

The following statement (contained in the xml_schema.sql script) creates the update_purchase order() procedure:

```
CREATE PROCEDURE update_purchase_order(
    p_purchase_order_id IN purchase_order.purchase_order_id%TYPE)

AS
    v_count INTEGER := 1;

-- declare a nested table to store products
    v_nested_table_products t_nested_table_product :=
        t_nested_table_product();

-- declare a type to represent a product record

TYPE t_product_record IS RECORD (
    product_id INTEGER,
    name VARCHAR2(15),
    quantity INTEGER
);
```

```
-- declare a REF CURSOR type to point to product records
  TYPE t product cursor IS REF CURSOR RETURN t product record;
  -- declare a cursor
 v product cursor t product cursor;
  -- declare a variable to store a product record
 v product t product record;
BEGIN
 -- open v product cursor to read the product id, name, and quantity for
 -- each product stored in the XML of the xml purchase order column
 -- in the purchase order table
  OPEN v product cursor FOR
 SELECT
   EXTRACTVALUE (product.COLUMN VALUE, '/product/product id')
     AS product id,
   EXTRACTVALUE (product.COLUMN VALUE, '/product/name') AS name,
    EXTRACTVALUE (product.COLUMN VALUE, '/product/quantity') AS quantity
  FROM TABLE (
    SELECT
     XMLSEQUENCE (EXTRACT (xml purchase order, '/purchase order//product'))
    FROM purchase order
   WHERE purchase order id = p purchase order id
  ) product;
  -- loop over the contents of v product cursor
    -- fetch the product records from v product cursor and exit when there
    -- are no more records found
    FETCH v product cursor INTO v product;
    EXIT WHEN v product cursor%NOTFOUND;
    -- extend v nested table products so that a product can be stored in it
    v nested table products.EXTEND;
    -- create a new product and store it in v nested table products
    v nested table products(v count) :=
      t product(v product.product id, v product.name, v product.quantity);
    -- display the new product stored in v nested table products
    DBMS OUTPUT.PUT LINE('product id = ' | |
      v nested table products(v_count).product_id);
    DBMS OUTPUT.PUT LINE('name = ' | |
      v nested table products(v count).name);
    DBMS OUTPUT.PUT LINE('quantity = ' ||
      v nested table products (v count).quantity);
    -- increment v count ready for the next iteration of the loop
    v count := v count + 1;
  END LOOP;
  -- close v product cursor
```

```
CLOSE v product cursor;
 -- update the purchase order table using the values extracted from the
 -- XML stored in the xml purchase order column (the products nested
 -- table is set to v nested table products already populated by the
 -- previous loop)
 UPDATE purchase order
   customer order id =
     EXTRACTVALUE (xml purchase order,
        '/purchase order/customer order id'),
   order date =
      TO DATE (EXTRACTVALUE (xml purchase order,
        '/purchase order/order date'), 'YYYY-MM-DD'),
   customer name =
     EXTRACTVALUE(xml purchase order, '/purchase order/customer name'),
     EXTRACTVALUE (xml purchase order, '/purchase order/street'),
     EXTRACTVALUE (xml purchase order, '/purchase order/city'),
     EXTRACTVALUE (xml purchase order, '/purchase order/state'),
     EXTRACTVALUE(xml purchase order, '/purchase order/zip'),
   phone number =
     EXTRACTVALUE(xml purchase order, '/purchase order/phone number'),
   products = v nested table_products
 WHERE purchase order id = p purchase order id;
 -- commit the transaction
 COMMIT;
END update purchase order;
```

The following example sets the server output on and calls update_purchase_order() to update purchase order #1:

```
SET SERVEROUTPUT ON

CALL update_purchase_order(1);

product_id = 1

name = Supernova video
```

```
name = Supernova video
quantity = 5
product_id = 2
name = Oracle SQL book
quantity = 4
```

The following query retrieves the columns from purchase order #1:

```
SELECT purchase_order_id, customer_order_id, order_date, customer_name, street, city, state, zip, phone_number, products
FROM purchase_order
```

WHERE purchase order id = 1;

```
PURCHASE ORDER ID CUSTOMER ORDER ID ORDER DAT CUSTOMER NAME
STREET
     CITY ST ZIP PHONE NUMBER
______ _____
PRODUCTS (PRODUCT ID, NAME, QUANTITY)
_____
                    176 17-MAY-07 Best Products 456 Inc.
10 Any Street Any City CA 94440 555-121-1234
T NESTED TABLE PRODUCT (
T PRODUCT(1, 'Supernova video', 5),
T PRODUCT(2, 'Oracle SQL book', 4)
```

The products nested table contains the same data as stored in the XML product elements in the xml purchase order column. I've added some line breaks to separate the products in the example's results to make them easier to see.

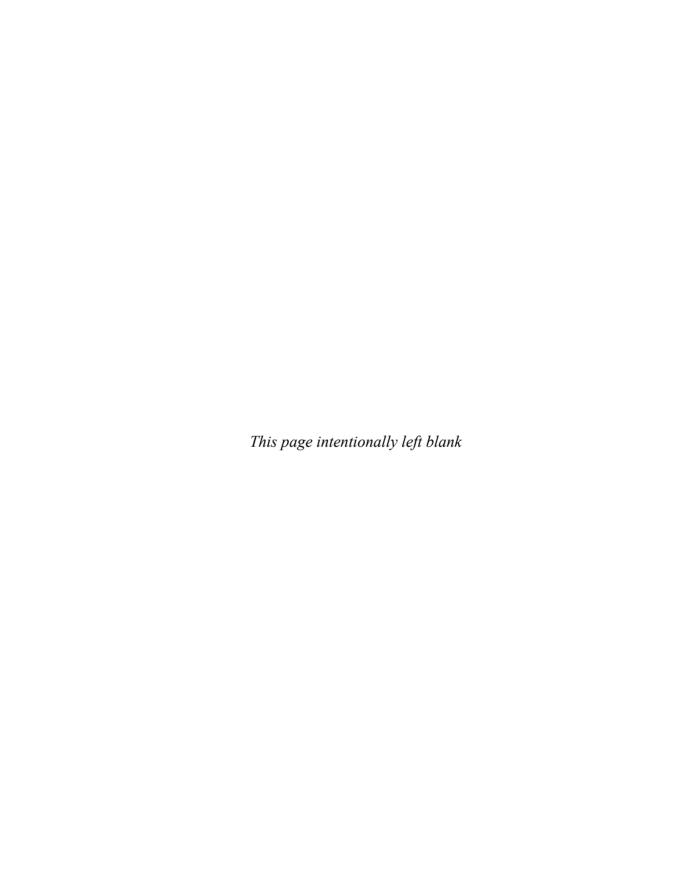
Summary

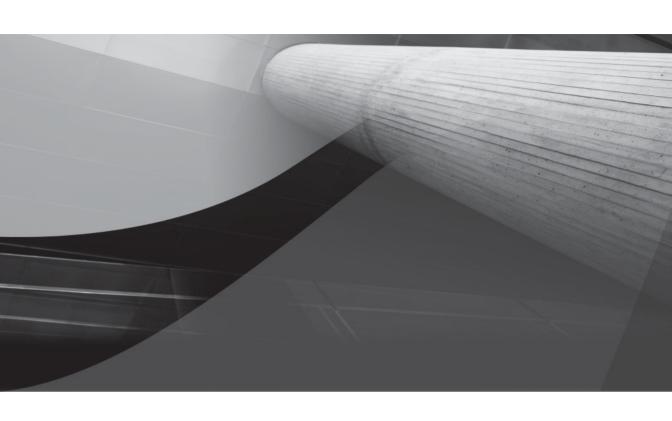
In this chapter, you have learned

- How to generate XML from relational data.
- How to save XML in the database and subsequently read that XML to update relational columns.

This short chapter has barely scratched the surface of the rich XML functionality available in the Oracle database. You can find more information in the Oracle XML Developer's Kit and the Oracle XML DB Developer's Guide, both published by Oracle Corporation.

Apart from the final appendix, this is the end of this book. I hope you've found the book informative and useful, and that I've held your interest!





APPENDIX

Oracle Data Types



his appendix contains two tables documenting the data types that are available in Oracle SQL and that may be used to define columns in a table, along with the additional types supported by Oracle PL/SQL.

Oracle SQL Types

Table A-1 shows the Oracle SQL types.

Туре	Description
CHAR[(length [BYTE CHAR])] ¹	Fixed-length character data of <code>length</code> bytes or characters and padded with trailing spaces. Maximum length is 2,000 bytes.
VARCHAR2(length [BYTE CHAR]) ¹	Variable-length character data of up to $length$ bytes or characters. Maximum length is 4,000 bytes.
NCHAR[(length)]	Fixed-length Unicode character data of <code>length</code> characters. Number of bytes stored is 2 multiplied by <code>length</code> for AL16UTF16 encoding and 3 multiplied by <code>length</code> for UTF8 encoding. Maximum length is 2,000 bytes.
NVARCHAR2 (length)	Variable-length Unicode character data of <code>length</code> characters. Number of bytes stored is 2 multiplied by <code>length</code> for AL16UTF16 encoding and 3 multiplied by <code>length</code> for UTF8 encoding. Maximum length is 4,000 bytes.
BINARY_FLOAT	Introduced in Oracle Database 10 <i>g</i> , stores a single-precision 32-bit floating-point number. Operations involving BINARY_FLOAT are typically performed faster than operations using NUMBER values. BINARY_FLOAT requires 5 bytes of storage space.
BINARY_DOUBLE	Introduced in Oracle Database 10 <i>g</i> , stores a double-precision 64-bit floating-point number. Operations involving BINARY_DOUBLE are typically performed faster than operations using NUMBER values. BINARY_DOUBLE requires 9 bytes of storage space.
NUMBER(precision, scale) and NUMERIC(precision, scale)	Variable-length number; <code>precision</code> is the maximum number of digits (left and right of a decimal point, if used) that may be used for the number. The maximum precision supported is 38; <code>scale</code> is the maximum number of digits to the right of a decimal point (if used). If neither <code>precision</code> nor <code>scale</code> is specified, then a number with up to a precision and scale of 38 digits may be supplied (meaning you can supply a number with up to 38 digits, and any of those 38 digits may be right or left of the decimal point).
DEC and DECIMAL	Subtype of NUMBER. A fixed-point decimal number with up to 38 digits of decimal precision.
DOUBLE PRECISION and FLOAT	Subtype of NUMBER. A floating-point number with up to 38 digits of precision.
REAL	Subtype of NUMBER. A floating-point number with up to 18 digits of precision.
INT, INTEGER, and SMALLINT	Subtype of NUMBER. An integer with up to 38 digits of decimal precision.
DATE	Date and time with the century; all four digits of year, month, day, hour (in 24-hour format), minute, and second. May be used to store a date and time between January 1, 4712 B.C. and December 31, 4712 A.D. Default format is specified by the NLS_DATE_FORMAT database parameter (for example: DD-MON-RR).
<pre>INTERVAL YEAR[(years_precision)] TO MONTH</pre>	Time interval measured in years and months; <code>years_precision</code> specifies the precision for the years, which may be an integer from 0 to 9 (default is 2). Can be used to represent a positive or negative time interval.

INTERVAL DAY ((days_precision)] TO SECOND[(seconds_precision)] Time interval measured in days and seconds; days_precision becifies the precision for the days, which is an integer from 0 to 9 (default is 2); seconds_precision specifies the precision for the fractional part of the seconds, which is an integer from 0 to 9 (default is 6). Can be used to represent a positive or negative time interval. TIMESTAMP [(seconds_precision)] Date and time with the century; all four digits of year, month, day, hour (in 24-hour format), minute, and second; seconds_precision specifies the number of digits for the fractional part of the seconds, which can be an integer from 0 to 9 (default is 6). Default format is specified by the NLS_TIMESTAMP_FORMAT database parameter. TIMESTAMP [(seconds_precision)] WITH LOCAL TIME ZONE	Туре	Description
format), minute, and seconds, precision specifies the number of digits for the fractional part of the seconds, which can be an integer from 0 to 9 (default is 6). Default format is specified by the NLS_TIMESTAMP_FORMAT database parameter. TIMESTAMP[(seconds_precision)] WITH Extends TIMESTAMP to store a time zone. The time zone can be an offset from UTC, such as -8:0, or a region name, such as US/Pacific or FST. Default format is specified by the NLS_TIMESTAMP_TZ_FORMAT database parameter. TIMESTAMP[(seconds_precision)] WITH LOCAL TIME ZONE TIMESTAMP[(seconds_precision)] WITH LOCAL TIME ZONE Extends TIMESTAMP to convert a supplied datetime to the local time zone set for the database. The process of conversion is known as normalizing the datetime. Default format is specified by the NLS_TIMESTAMP_FORMAT database parameter. CLOB Variable-length single-byte character data of up to 128 terabytes. Variable-length Unicode national character set data of up to 128 terabytes. LONG Variable-length Unicode national character set data of up to 128 terabytes. Pointer to an external file. The external file is not stored in the database. LONG Variable-length binary data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type, but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. Hexadecimal string used to represent a row address. Hexadecimal string used to represent a row address. Hexadecimal string used to represent a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. NESTED TABLE Nested defined object type		precision for the days, which is an integer from 0 to 9 (default is 2); seconds precision specifies the precision for the fractional part of the seconds, which is an integer from 0 to 9 (default is 6). Can be used to represent a positive or negative
TIME ZONE UTC, such as -8:0, or a region name, such as US/Pacific or PST. Default format is specified by the NLS_TIMESTAMP_TZ_FORMAT database parameter. Extends TIMESTAMP to convert a supplied datetime to the local time zone set for the database. The process of conversion is known as normalizing the datetime. Default format is specified by the NLS_TIMESTAMP_FORMAT database parameter. CLOB Variable-length single-byte character data of up to 128 terabytes. NCLOB Variable-length Unicode national character set data of up to 128 terabytes. BLOB Variable-length binary data of up to 128 terabytes. Pointer to an external file. The external file is not stored in the database. LONG Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to 1ength bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string used to represent a row address. Hexadecimal string representing the logical address of a row of an index-organized table; length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	TIMESTAMP[(seconds_precision)]	format), minute, and second; <pre>seconds_precision</pre> specifies the number of digits for the fractional part of the seconds, which can be an integer from 0 to 9 (default is 6). Default format is specified by the NLS_TIMESTAMP_FORMAT database
the database. The process of conversion is known as normalizing the datetime. Default format is specified by the NLS_TIMESTAMP_FORMAT database parameter. CLOB Variable-length single-byte character data of up to 128 terabytes. NCLOB Variable-length Unicode national character set data of up to 128 terabytes. BLOB Variable-length binary data of up to 128 terabytes. BFILE Pointer to an external file. The external file is not stored in the database. LONG Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to 1ength bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string used to represent a row address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type	<u>—</u> -	UTC, such as -8:0, or a region name, such as US/Pacific or PST. Default
NCLOB Variable-length Unicode national character set data of up to 128 terabytes. BLOB Variable-length binary data of up to 128 terabytes. BFILE Pointer to an external file. The external file is not stored in the database. LONG Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to 1 length bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type		the database. The process of conversion is known as <i>normalizing</i> the datetime.
BLOB Variable-length binary data of up to 128 terabytes. BFILE Pointer to an external file. The external file is not stored in the database. LONG Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to length bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type	CLOB	Variable-length single-byte character data of up to 128 terabytes.
Pointer to an external file. The external file is not stored in the database. LONG Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to length bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an unordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType You can define your own object type and create objects of that type. See Chapter	NCLOB	Variable-length Unicode national character set data of up to 128 terabytes.
Variable-length character data of up to 2 gigabytes. Superseded by the CLOB and NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to length bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	BLOB	Variable-length binary data of up to 128 terabytes.
NCLOB types, but supported for backwards compatibility. RAW (length) Variable-length binary data of up to length bytes. Maximum length is 2,000 bytes. Superseded by the BLOB type, but supported for backwards compatibility. LONG RAW Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	BFILE	Pointer to an external file. The external file is not stored in the database.
bytes. Superseded by the BLOB type, but supported for backwards compatibility. Variable-length binary data of up to 2 gigabytes. Superseded by the BLOB type but supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	LONG	
supported for backwards compatibility. ROWID Hexadecimal string used to represent a row address. UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	RAW(length)	
UROWID[(length)] Hexadecimal string representing the logical address of a row of an index-organized table; length specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	LONG RAW	
table; <code>length</code> specifies the number of bytes. Maximum length is 4,000 bytes (also the default length if none is specified). REF object_type Reference to an object type. Similar to a pointer in the C++ programming language. VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	ROWID	Hexadecimal string used to represent a row address.
VARRAY Variable-length array. This is a composite type and stores an ordered set of elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	UROWID[(length)]	table; <code>length</code> specifies the number of bytes. Maximum length is 4,000 bytes (also
elements. NESTED TABLE Nested table. This is a composite type and stores an unordered set of elements. XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	REF object_type	Reference to an object type. Similar to a pointer in the C++ programming language.
XMLType Stores XML data. User defined object type You can define your own object type and create objects of that type. See Chapter	VARRAY	
User defined object type You can define your own object type and create objects of that type. See Chapter	NESTED TABLE	Nested table. This is a composite type and stores an unordered set of elements.
	XMLType	Stores XML data.
	User defined object type	, , , , , , , , , , , , , , , , , , , ,

 $^{^{1}}$ The BYTE and CHAR keywords work only with Oracle Database 9i and above. If neither BYTE nor CHAR is specified, the default is BYTE.

TABLE A-1 Oracle SQL Types (continued)

Oracle PL/SQL Types
Oracle PL/SQL supports all the types previously shown in Table A-1, plus the following additional Oracle PL/SQL specific types shown in Table A-2.

Туре	Description
BOOLEAN	Boolean value (TRUE, FALSE, or NULL).
BINARY_INTEGER	Integer between -2 ³¹ (-2,147,483,648) and 2 ³¹ (2,147,483,648).
NATURAL	Subtype of BINARY_INTEGER. A non-negative integer.
NATURALN	Subtype of BINARY_INTEGER. A non-negative integer (cannot be NULL).
POSITIVE	Subtype of BINARY_INTEGER. A positive integer.
POSITIVEN	Subtype of BINARY_INTEGER. A positive integer (cannot be NULL).
SIGNTYPE	Subtype of BINARY_INTEGER. An integer set to -1, 0, or 1.
PLS_INTEGER	Integer between -2^{31} (-2,147,483,648) and 2^{31} (2,147,483,648). Identical to BINARY_INTEGER.
SIMPLE_INTEGER	New for Oracle Database 11g, SIMPLE_INTEGER is a subtype of BINARY_INTEGER. SIMPLE_INTEGER can store the same range of values as BINARY_INTEGER, except for NULL values, which cannot be stored in a SIMPLE_INTEGER. Also, arithmetic overflow does not cause an error when using SIMPLE_INTEGER values; instead, the result is simply truncated.
STRING	Same as VARCHAR2.
RECORD	Composite of a group of other types. Similar to a structure in the C++ programming language.
REF CURSOR	Pointer to a set of rows.

 TABLE A-2
 Oracle PL/SQL Types



Symbols

- (minus), 31

(pound), 107

% (percent), 40-41

%TYPE, 342, 346

() (parentheses), 34, 168

* (asterisk)

in arithmetic, 31

column retrieval, 29

customers table, 15

metacharacters, 110

/(forward slash)

in arithmetic, 31

editing SQL statements, 65-66

PL/SQL statements, 341

SQL directory paths, 10

\ (backward slash), 10

:new attribute, 524-525

; (semicolon), 6, 341

? (question mark), 564

!= (not equal), 38–39

@ filename, 67

(underscore), 40–41

|| (concatenation operator), 34

+ (plus)

in arithmetic, 31

outer joins, 51-55

< (less than)

comparing object values, 393

non-equijoins, 50

value comparison, 38-39

<= (less than or equal)

comparing object values, 393

non-equijoins, 50

value comparison, 38-39

<> (not equal)

ANSI nested table support,

464-465

comparing object values,

392-394

value comparison, 38–39

= (equal)

ANSI nested table support,

464-465

comparing object values,

392-393

equijoins, 50

value comparison, 38-39

> (greater than)

comparing object values, 393-394

non-equijoins, 50

value comparison, 38–39

>= (greater than or equal)

comparing object values, 393 non-equijoins, 50

value comparison, 38-39

" (quotes), 34, 254

Α

ABS(), 98, 100 ACCEPT, 78-79, 80-81 ACCESS_INTO_NULL, 354 ACID (atomic, consistent, isolated and durable) transaction properties, 265-266 ACOS(), 98

ADDM (Automatic Database

Diagnostic Monitor), 600-601

ADD_MONTHS(), 142-143 advanced queries. See queries.

advanced

aggregate functions

defined, 90

grouping rows, 120-127

multiple in PIVOT, 247-248

RETURNING, 255-256

with ROLLUP, 207

types of, 117–120

algorithms, LOB data encryption, 526

aliases

columns, 34-35

table, 47-48

ALL, 38-39, 175

ALTER

defined, 3

INDEX, 324

SEQUENCE, 320

TABLE, 302-306

American National Standards Institute (ANSI), 464-472

analytic functions

defined, 216-217

FIRST and LAST functions, 233

hypothetical rank and

distribution functions, 235

inverse percentile functions, 224-225

LAG() and LEAD(), 232-233

linear regression functions,

233-235

ranking functions, 217-224 reporting functions, 230-232

window functions, 225-230

AND, 43, 238

BasicExample3.java, 575-577

BETWEEN

ANSI (American National Standards	cell range access using, 238	
Institute), 47, 464–472	comparing object values, 393	CACLIE in commence 210
ANY	defined, 39	CACHE in sequences, 319
accessing all cells using,	non-equijoins, 50–51	caching LOB data, 526
238–239	overview, 42–43	calculations, numeric, 98–102
multiple-row subqueries, 174	BFILEs (binary FILE types)	calculations with MODEL. See
value comparison, 38–39	copying data from into CLOBs	MODEL
A[PPEND], 65–66	1,7	calling PL/SQL functions, 364,
APPEND()	and BLOBs, 517–520	367–368
	defined, 477–478	calling PL/SQL procedures, 360–
data to CLOBs, 506–507	FILECLOSE() and	361, 367–368
defined, 483, 485	FILECLOSEALL(), 489	CARDINALITY(), 469
arithmetic, 31–34	FILEEXISTS(), 490	Cartesian products, 48-49
arrays, associative, 428	FILEGETNAME(), 490	CASE expressions
associative, 462	FILEISOPEN(), 490–491	advanced queries, 193–196
arrays, varray. See varrays	FILEOPEN(), 491	vs. multiple queries, 583
ASCII(), 91–92	LOADBLOBFROMFILE(), 497	using with $GROUPING()$,
ASIN(), 98	LOADCLOBFROMFILE(),	209–210
associative arrays, 428, 462	497–498	CASE_NOT_FOUND, 354
asterisk (*). See * (asterisk)	LOADFROMFILE(), 496	CAST(), 108-109, 443-444
ATAN(), 98	using in SQL, 481–482	CEIL(), 99, 100
ATAN2(), 99	binary LOBs (BLOBs). See BLOBs	cells
atomic transactions, 265	(binary LOBs)	accessing all using ANY and IS
attributes	BINARY_DOUBLE, 13, 23–25	ANY, 238–239
:new, 524–525	BINARY_FLOAT, 13, 23–25	FOR loops, 240-241
OracleDataSource,	bind variables, 588-591	positional and symbolic
540–542	binds, 539	notation, 237–238
XMLATTRIBUTES(), 608	bit vectors, 212	range access using BETWEEN
auditing	BITAND(), 99	and AND, 238
overview, 295–297	bitmap indexes, 321, 324-325	updating existing, 243-244
using triggers, 369–374	BLOBs (binary LOBs)	centered average, 228-229
auto-commit mode, 552-553	compression, 529	central plan tables, 593
Automatic Database Diagnostic	copying data from BFILE into,	century formatting, 133
Monitor (ADDM), 600-601	517–520	C[HANGE], 65–66
automatically generated SQL	copying data from file into,	changing table contents. See table
statements, 86	512–515	contents, changing
averages	copying data from into files,	CHAR(), 13
centered, 228–229	515–517	character functions
moving, 227–228	defined, 477–478	defined, 90
AVG(), 117–118	encryption, 526–528	types of, 90–98
	LOADBLOBFROMFILE(), 497	character LOBs (CLOBs). See CLOBs
	reading data from, 504–506	(character LOBs)
В	using in SQL, 478–481	characters
	block structure, 340-342	variable definition, 76
B-tree indexes, 321-322	body, object type, 381-382	XML files written in English, 625
backward slash (\), 10	body, package, 365-367	CHECK, 306, 307
base tables, 325	branch elimination, 201–202	CHECK OPTION, 306, 329
BasicExample1.java, 556-562	BREAK ON, 83–85	child nodes, 197
BasicExample2.java, 565-567	BTITLE, 82–83	child tables, 257–258

buckets, 223

CHR(), 91-92

chunks of data, 492 information retrieval on indexes, defining column with type, 430 DELETE(), 451-452 323-324 classes compatible Oracle database element type sizing, 463 joins with fully qualified types, 568 EXISTS(), 452-453 references, 582 EXTEND(), 453 LONG and LONG RAW types, object types as, 380 oracle.jdbc package, 573 FIRST(), 454 521-522 CLASSPATH environment variable, generating XML for, 607-608 multiple-column increasing elements in 534-535 subaueries, 175 CLE[AR], 70-71 varrav, 463 objects, 383-386 information retrieval, 431-433 pivoting on multiple, 246-247 CLEAR, 73 CLOBs (character LOBs) introduction, 428-429 plan table, 593 appending data to, 506-507 LAST(), 454-455 reading values from comparing data, 507-508 modifying elements, 438-440 ResultSet object, 544-546 compression, 529 multilevel, 458-461 retrieving all, 29 in ROLLUP clause, 205-207 copying data from, 508-509 nested table's storage table tablespace, 463-464 copying data from BFILE into, SQL*Plus, clearing formats, 73 517-520 NEXT(), 455-456 SQL*Plus, formatting, 70-72 copying data from file into, Oracle database 10g substituting names using variables, 76-77 512-515 enhancements, 461-462 UNPIVOT, 248-249 copying data from into files, in PL/SQL, 444-448 515-517 populating with elements, using multiple in group, 122 defined, 477-478 434-435 combining functions, 98 encryption, 526-528 PRIOR(), 456-457 command line SQL*Plus startup, 6 erasing data from, 510 retrieving elements from, commands implicit conversion, 523-524 435-436 column formatting, 70-72 LOADCLOBFROMFILE(), defined variables, 77-79 summary, 473 497-498 TABLE(), 436-438 editing SQL statements, 65-66 TRIM(), 457-458 HELP, 85-86 reading data from, 504-506 searching data, 510-511 types, 429-430 saving/retrieving/running files, temporary, 509-510 varrays in temporary tables, 463 66-70 XMLAGG(), 609-612 using in SQL, 478-481 SET LINESIZE, 73 writing to, 506 SET PAGESIZE, 72-73 columns CLOSE(), 483, 485 simple reporting, 79-85 accessing using MODEL clause, closed operations, 23 236-237 comments closing adding/modifying/removing, in tables, 313-314 cursors, 347 303-306 XML, 614 JDBC objects, 554-556 aliases, 34-35 **COMMIT** ResultSet object, 546-547 in arithmetic, 32-34 auto-commit mode, 552-553 changing table contents. See Codd, E. F., 2, 3 defined, 4 table contents, changing transactions, 262-263 COLLECT(), 471–472 combining output, 35 COLLECTION_IS_NULL, 354 Communications of the ACM, 2 collections, 427-473 comments, 313-314 COMPARE() ANSI nested table support, constraint information retrieval, data in two CLOBs, 507-508 464-472 311-312 defined, 483, 486-487 associative arrays, 462 declaring variables to store comparing values comparing nested table contents values, 346 defined, 37-39 with map method, 440-442 defined, 2 table objects, 392-394 converting types with CAST(), defining with collection type, 430 complex views, 332-333 443-444 encrypting LOB data, 528 compression, LOB, 529 COUNT(), 448-451 GROUPING(), 209-211 COMPUTE, 83-85

oracle Database 118 5Q	L	
computer configuration for running SQL using Java, 533–535 computing subtotals, 83–85 concatenation combining column output, 35 CONCAT(), 91–92 XMLCONCAT(), 613 concatenation operator (), 34 concurrent transactions, 266 conditional logic, 342–343 conditions in hierarchical queries, 202–203 conditions, trigger, 369 configuring computer for running SQL using Java, 533–535 CONNECT BY, 198 connectivity, database, 537–542 consistent transactions, 265 constrained cursors, 345–351 constraints creating views with, 329–330 enforcing, 257–258 getting information on view, 331–332 table, 306–312 constructors column objects, 384 object tables, 386 user-defined, 418–422 CONTENT parsing XML expressions, 613 XMLSERIALIZE(), 616 contents, changing table. See table contents, changing table. Handle Converting datetime, 131–139, 150 CONVERTTOBLOB(), 483 CONVERTTOCLOB(), 483 CONVERTTOCLOB(), 483 CONVERTTOCLOB(), 483 CONVERTTOCLOB(), 483 CONVERTTOCLOB(), 483 COPY() copying data from CLOBs, 508–509 defined, 483 LOB methods, 487 copying LOB data	from file into CLOBs and BLOBs, 512–515 copying rows, 254 correlated subqueries, 168, 175–179 COS(), 99 COSH(), 99 cost-based optimization, 591–598 COUNT(), 117–119, 448–451 CREATE, 3, 11–12 CREATE PROCEDURE, 25–26 CREATE SYNONNYM, 287–288 CREATE USER, 276–277 CREATETEMPORARY(), 483, 488 cross joins, 61 CUBE defined, 207–208 vs. GROUPING SETS, 212, 588 GROUPING() with, 210–211 using with analytic functions, 221–223 CUME_DIST(), 218, 223 cumulative sums, 225–227 CURRENT_DATE(), 149 CURRENT_TIMESTAMP, 155–156 CURRENTV(), 239–240 currval, 317, 376 CURSOR_ALREADY_OPEN, 354 cursors FOR loops and, 349–350 OPEN-FOR statement, 350–351 overview, 345–346 product_cursor.sql example, 348–349 programming steps, 346–348 unconstrained, 352–353 customer_id, 14, 17–18 customers table, 14–15 D Data Control Language (DCL) statements	comparing in CLOBs, 507–508 compression, 529 copying. See copying LOB data defined, 478 encryption, 525–528 erasing from CLOB, 510 reading from CLOBs and BLOBs, 504–506 removing duplicate, 529 searching in CLOB, 510–511 Data Manipulation Language (DML) statements. See DML (Data Manipulation Language) statements data sources, 539–542 data types. See types database management systems, 2–3 database objects, 379–426 column, 383–386 comparing values, 392–394 creating types, 381–382 delete_product(), 400–401 display_product(), 396–397 generalized invocation, 423–425 generating XML for, 607 get_products(), 398–399 get_product_ref(), 400 get_products(), 395–396 insert_product(), 397 introduction, 380 IS OF(), 408–412 JDBC Statement, 542–543 large. See LOBs (large objects) NOT INSTANTIABLE types, 416–418 OIDs and references, 390–392 OracleDataSource, 540–542 overriding methods, 422–423 in PL/SQL, 394–395 product_lifecycle(), 401–402 product_lifecycle2(), 402–403 Resultset, 543–547
appending data to CLOBs, 506–507	defined, 3, 11–20 running SQL using Java, 553	subtype in place of supertype, 405–408
from BFILE into CLOBs and	data, LOB	summary, 425–426
BLOBs, 517-520	appending to CLOB, 506-507	SYS_TYPEID(), 416
from CLOBs and BLOBs into	CLOBs and BLOBs in SQL,	tables, 386–389
files, 515–517	478–481	TREAT(), 412–416

type inheritance, 403–405	datetime, 130–166	collection information retrieval,
update_product(),399	converting, 131–139	431–433
update_product_	default format, 139–140	object type information,
price(),398	functions, 142–146	382–383
user-defined constructors,	storing/retrieving, 130–131	row identifiers, 30
418–422	summary, 166	table information retrieval, 302
using DESCRIBE to get type	time intervals, 159–165	viewing table structure, 64–65
information, 382–383	time zones, 146–150	description, 17
wrappers, 551–552	timestamp functions, 154–159	dimensions, current values, 239–240
database table information retrieval,	timestamps, 150–154	directory objects
27–61	two-digit years, 140-142	using BFILEs, 481
all columns, 29	days	wallets, 525–526
arithmetic, 31–34	datetime functions, 143-145	directory paths, 10
Cartesian products, 48–49	default format, 130	dirty reads, 268
column aliases, 34–35	format parameters, 134	disabling constraints, 309
concatenation, 35	INTERVAL DAY TO SECOND,	disabling triggers, 374
distinct row display, 37	162–164	displaying distinct rows, 37
join conditions and types, 50	DCL (Data Control Language)	display_product(),394,
joins using SQL/92 syntax,	statements, 4	396–39 7
56–61	DD, 130	DISTINCT
non-equijoins, 50–51	DDL (Data Definition Language)	defined, 37
null values, 35-37	statements	vs. EXISTS, 587
operator precedence, 44	defined, 3, 11–20	distribution, 235
outer joins, 51–55	running SQL using Java, 553	DML (Data Manipulation Language)
row identifiers, 30	declarations, cursor, 346	statements
row numbers, 30–31	DECODE(), 191–193	changing table contents. See
SELECT statements using	default datetime format, 139-140	table contents, changing
multiple tables, 49–50	default editors, 69	defined, 3
SELECT statements using single	default locking, 266–267	performing with views, 328–329
table, 28–29	default roles, 294	triggers, 369–374
SELECT statements using two	default values, 258–259, 306	dob (date of birth), 15
tables, 45–47	deferred constraints, 310	DOCUMENT, 613, 616
self joins, 55–56	DEFINE, 77–78	drivers, JDBC
sorting rows, 44–45	defined variables	registering, 537
SQL operators, 39–43	defined, 74, 77-79	types of, 535–536
summary, 61	in scripts, 81	DROP, 3
table aliases, 47–48	definition characters, 76	DROP TABLE, 86
tables and columns in tables,	definitions, view, 330–331	DROP USER, 278
302–303	DEL, 65–66	dropping
value comparison, 37–39	DELETE(), 449, 451–452	columns, 306
WHERE, 29	DELETE 137 131 132	constraints, 309
databases	defined, 3	indexes, 324
common types, 12–13	JDBC, 548	packages, 368
creating users, 11–12	removing rows, 22–23, 256	PL/SQL functions, 365
integrity, 256–258	statements containing	procedures, 362
opening connection, 537–542	subqueries, 180–181	roles, 295
tables. <i>See</i> tables	•	sequences, 320
	delete_product(), 395,	•
time zones, 147–149	400–401 DENSE RANK() 217–223	tables, 314
transactions, 262–269, 552–553	DENSE RANK(), 217–223	triggers, 374
DATE, 13	DEREF(), 391–392	views, 334
date of high (data) 15	DESCRIBE	dual table, 32
date of birth (dob), 15	adding rows, 20	duplicate LOB data, 529

exceptions. See exceptions

duplicate rows, 215–216 duplicate values, 356 DUP_VAL_ON_INDEX, 354, 356 durable transactions, 265	single-row subqueries, 172–173 viewing in procedures, 362 exceptions DUP_VAL_ON_INDEX, 356 INVALID_NUMBER, 356–357	filtering rows with HAVING, 125–126 with WHERE, 580–581 WHERE vs. HAVING, 584–585
E	OTHERS, 357 PL/SQL, 353–355	finding triggers, 371–372 FIRST, 216, 233
FDUTL 67	running SQL using Java,	FIRST(), 449, 454
ED[IT], 67	553–554	first_name, 14
editing SQL statements, 7, 65–66	thrown by methods. See	FIRST_VALUE(), 229–230
elements	methods	flashback data archives, 334–337
adding to collection, 453	ZERO_DIVIDE, 355–356	flashbacks, query, 270–273
counting, 448–451	execution plans, 562, 591–598	FLOOR(), 99, 100–101
deleting, 451–452	EXISTS(), 449, 452–453	footers, 82–83
existing, 452–453	EXISTS	FOR loops
getting first, 454	vs. IN, 586–587	cell access using, 240–241
getting last, 454–455	correlated subqueries, 176–179	cursors and, 349–350
increasing in varray, 463	vs. DISTINCT, 587	defined, 343, 344–345
modifying in collections,	EXIT, 25, 86	foreign keys
438–440	EXP(), 99	adding, 308–309
populating collections,	EXPLAIN PLAN, 594	constraints enforcement,
434–435	expressions	257–258
removing from collection,	arithmetic, 31–34	defined, 306
457–458	CASE. See CASE expressions	forests
retrieving from collections,	regular expression functions,	XMLAGG(), 609–612
435–436	109–116	XMLFOREST(), 608–609
returning next, 455–456	using with functions, 98	FOR[MAT] format, 70-71
returning prior, 456–457	XMLPARSE(), 613	formats, datetime
type sizing, 463	EXTEND(), 449, 453	conversion, 131–139
XML, 605–608	Extensible Markup Language (XML).	date, 130
ELSE, 342–343	See XML (Extensible Markup	default, 139–140
ELSIF, 342–343	Language)	two-digit year interpretation,
employees table, 18-19	EXTRACT(), 155–157	140–142
enabling constraints, 309–310 enabling triggers, 374		formatting hierarchical query results, 199–200
encryption, LOB, 525–528	F	
END IF, 342–343	<u> </u>	forward slash (/). See / (forward slash)
ending transactions, 263	FETCH, 347	fragments, XML, 612
environment variables, 533–535	FILECLOSE(), 483, 489	FREETEMPORARY(), 484, 492
equal (=). See = (equal)	FILECLOSEALL(), 483, 489	FROM, 171–172
equijoins	FILEEXISTS(), 483, 490	FROM_TZ(), 155, 157
defined, 46	FILEGETNAME(), 483, 490	full outer joins, 60
SELECT statements using	FILEISOPEN(), 483, 490–491	fully qualified column references, 582
multiple tables, 49–50	FILEOPEN(), 483, 491	function-based indexes, 322–323
SQL/92 syntax, 57–58	files	functions, 89–127
era formatting parameters, 135	BFILEs. See BFILEs (binary	aggregate, 117–120
ERASE()	FILE types)	analytic. See analytic functions
data from CLOB, 510	copying data from CLOBs and	ANSI nested table support,
defined, 483	BLOBs into, 515–517	469–472
LOB methods, 488–489	copying data from into CLOBs	calling in packages, 367–368
errors	and BLOBs, 512–515	with CAST(), 443–444

saving/retrieving/running, 66-70

character, 90-98

conversion, 102-109	get_product(), 394, 398-399	Н
CURRENTV(), 239–240	get_product_ref(), 395, 400	
datetime, 142-146	get_products(),394,	
datetime conversion,	395–396	
131–139	GET_STORAGE_LIMIT(),	h
DECODE(), 191–193	484, 492	Н
DEREF(), 391–392	GMT (Greenwich Mean Time),	h
get product(),398-399	146–147	h
get product ref(),400	GRANT, 4	
get products(), 395-396	granting privileges	
GROUP_ID(), 215–216	object, 282	
GROUPING(), 209–211	roles, 290	
grouping rows, 120–127	to users, 278–279	
GROUPING_ID(), 212–214	using with roles, 293	
IS OF(), 408–412	granting roles, 290–291	
numeric, 98–102	greater than (>). See > (greater than)	
PL/SQL, 363–365	greater than or equal (>=). See >=	
REF(), 390–392	(greater than or equal)	
regular expression, 109–116	Greenwich Mean Time (GMT),	
single-row, 90	146–147	
storing returned values with	GROUP BY, 120–127	h
bind variables, 590	GROUP BY, extended clauses	
summary, 127	CUBE, 207–208	h
SYS_TYPEID(), 416	defined, 203–204	
TABLE(), 436–438	GROUP BY column use, 214	h
timestamp, 154–159	GROUP_ID(), 215-216	h
TRANSLATE(), 190–191	GROUPING(), 209-211	
TREAT(), 412–416	GROUPING SETS, 211–212	
using objects in PL/SQL,	GROUPING_ID(), 212–214	
394–395	multiple column use, 214	
XMLAGG(), 609-612	ROLLUP, 205–207	IE
XMLATTRIBUTES(), 608	GROUP_ID(), 215-216	
XMLCOLATTVAL(), 612	GROUPING(), 209–211	
XMLCOMMENT(), 614	grouping functions	
XMLCONCAT(), 613	grouping rows, 120–127	ic
XMLELEMENT(), 605–608	types of, 117–120	ic
XMLFOREST(), 608–609	GROUPING SETS	IF
XMLPARSE(), 613	vs. CUBE, 588	IC
XMLPI(), 614	defined, 211–212	ir
XMLQUERY(), 618–622	using with analytic functions,	ir
XMLSEQUENCE(), 615	221–223	١N
XMLSERIALIZE(), 616	GROUPING_ID(), 212–214	

G

generalized invocation, 423-425 GET filename, 67 get methods, 540-542, 545-546 GETCHUNKSIZE(), 484, 492 getConnection(), 537-538 GETLENGTH(), 484, 493

Н

handling database null values, 550-552 handling exceptions. See exceptions handling null and missing values, 241–243 handling numbers, 549-550

HAVING grouping rows, 125-127 subqueries in, 170-171 vs. WHERE, 584-585 neaders, 82-83 HEA[DING] heading, 70-71 nelp, SOL*Plus, 85–86 hierarchical gueries defined, 196-199 eliminating nodes and branches, 201-202 formatting results, 199-200 including other conditions, 202-203 START WITH subquery, 200-201 starting at node other than root, 200 traversing upward through tree, 201 nierarchies of object types, 403-405 nints, passing to optimizer, 598-600 nour formatting parameters, 134 hypothetical functions, 216, 235

ID GROUP_ID(), 215-216 GROUPING_ID(), 212-214 SYS_TYPEID(), 416 dentical SQL statements, 588-589 dentifiers, row, 30 IF, 342-343 **GNORE NAV, 243** mplicit conversion, 523-524 mporting JDBC packages, 536-537 ANSI nested table support, 465-466 comparing object values, 392-394 defined, 39 vs. EXISTS, 586-587 vs. EXISTS in correlated subqueries, 178 multiple-row subqueries, 173-174

outer join limitations, 54

IN (continued) procedure parameters, 358 using, 41–42 IN OUT, 358 inclusive range, 42 indexes adding to tables, 584 DUP_VAL_ON_INDEX exceptions, 356 overview, 320–325 summary, 337 INF (infinity), 24 information retrieval from CLOBs, 479–480 collections, 431–433 column comments, 313–314	interrow calculations with MODEL. See MODEL INTERSECT, 184, 188–189 INTERVAL DAY TO SECOND, 162–164 INTERVAL YEAR TO MONTH, 159–162 intervals, time, 159–165 INVALID_CURSOR, 354 INVALID_NUMBER, 354, 356–357 inverse percentile functions, 216, 224–225 invocation, generalized, 423–425 IS A SET, 470–471 IS ANY, 238–239 IS EMPTY, 471	execution plans involving table, 596–597 MERGE, 261 non-equijoins, 50–51 outer joins, 51–55 qualified column references, 582 SELECT statements using multiple tables, 49–50 SELECT statements using two tables, 45–47 self joins, 55–56 table joins vs. multiple queries, 581–582 using SQL/92 syntax, 56–61 JUS[TIFY], 70–71
constraints, 310–312 indexes, 323–324 packages, 368 PL/SQL functions, 365 procedures, 361–362 sequences, 316–317 table comments, 313 tables. See database table information retrieval triggers, 372–374	IS INFINITE, 39 IS NAN, 39 IS NULL, 39 IS OF(), 408–412 IS PRESENT, 241–242 isolated transactions, 265 isolation levels, transaction, 267–269 ISOPEN(), 484, 494–495 ISTEMPORARY(), 484, 495–496	K KEEP NAV, 243 keys associative arrays, 462 primary and foreign, 257–258
using DESCRIBE to get object type information, 382–383 views, 330–332 XML schema, 625–630	J	LAG(), 216, 232–233 large objects (LOBs). <i>See</i> LOBs (larg objects)

LIMIT(), 449 LOADCLOBFROMFILE(), M line size setting, 73 497-498 linear regression functions, 216, LOADFROMFILE(), 496 Macbeth (Shakespeare), 477 233-235 LONG and LONG RAW types, map method, 440-442 Linux, ORACLE HOME 521-522 MAX(), 117, 119 environment variable, 533-534 OPEN(), 498-499 MEMBER OF, 469-470 Oracle database 10g L[IST], 65-66 MERGE, 259-262 listing bind variables, 590 enhancements, 523-525 metacharacters LN(), 99 Oracle database 11g Perl-influenced, 112 LOADBLOBFROMFILE(), enhancements, 525-529 regular expressions, 109-112 in PL/SQL, 482-484 484, 497 method overloading, 420 LOADCLOBFROMFILE(), 484, READ(), 499-500 method overriding, 422-423 reading data from CLOBs and 497-498 methods BLOBs, 504-506 LOADFROMFILE(), 484, 496 APPEND(), 485 retrieving locator, LOBMAXSIZE, 484 CLOSE(), 485 LOBs (large objects) 503-504 COMPARE(), 486-487 APPEND(), 485 searching CLOB data, COPY(), 487 appending data to CLOB, 510-511 COUNT(), 448-451 506-507 in SQL, 478-482 CREATETEMPORARY(), 488 CLOSE(), 485 SUBSTR(), 500-501 DELETE(), 451-452 COMPARE(), 486-487 summary, 529-530 ERASE(), 488-489 comparing data in CLOBs, tables containing, 478 EXISTS(), 452-453 507-508 temporary CLOBs, 509-510 EXTEND(), 453 COPY(), 487 TRIM(), 501-502 FILECLOSE() and copying data from BFILE into types, 477-478 FILECLOSEALL(), 489 CLOBs and BLOBs, 517-520 WRITE(), 502 FILEEXISTS(), 490 copying data from CLOBs, WRITEAPPEND(), 503 FILEGETNAME(), 490 508-509 writing to CLOB, 506 FILEISOPEN(), 490-491 copying data from CLOBs and XMLSERIALIZE(), 616 FILEOPEN(), 491 BLOBs into files, 515-517 local time zones, 153-154, 159 FIRST(), 454 copying data from file into LOCALTIMESTAMP, 155-156 FREETEMPORARY(), 492 CLOBs and BLOBs, 512-515 locators, LOB, 478, 503-504 generalized invocation, CREATETEMPORARY(), 488 locks, transaction, 266-267 423-425 ERASE(), 488-489 LOG(), 99 GETCHUNKSIZE(), 492 erasing CLOB data, 510 logic, conditional, 342-343 getConnection(),537-538 example files, 476-477 logical operators, 43 GETLENGTH(), 493 FILECLOSE() and logical unit of work, 262 GET STORAGE LIMIT(), 492 FILECLOSEALL(), 489 LOGIN_DENIED, 354 INSTR(), 493-494 FILEEXISTS(), 490 LONG RAW types, 521-522 ISOPEN(), 494-495 LONG types, 521–522 FILEGETNAME(), 490 ISTEMPORARY(), 495–496 looking up, 539 FILEISOPEN(), 490-491 LAST(), 454-455 FILEOPEN(), 491 loops LOADBLOBFROMFILE(), 497 FREETEMPORARY(), 492 cell access using FOR, LOADCLOBFROMFILE(), GETCHUNKSIZE(), 492 240-241 497-498 GETLENGTH(), 493 cursors and, 349-350 LOADFROMFILE(), 496 GET_STORAGE_LIMIT(), 492 PL/SQL, 343-345 NEXT(), 455-456 LOWER(), 91, 94-95 INSTR(), 493-494 OPEN(), 498-499 introduction, 476 lower bounds, FOR loops, PRIOR(), 456-457 ISOPEN(), 494-495 344-345 READ(), 499-500 ISTEMPORARY(), 495–496 LPAD(), 91, 95 SUBSTR(), 500-501 LOADBLOBFROMFILE(), 497 LTRIM(), 91, 95 TRIM(), 457-458, 501-502

named notation, 361

names, time zone, 149

methods (continued)	NAN (not a number), 24	NOT INSTANTIABLE types,
WRITE(), 502	native machine code generation, 377	416–418
WRITEAPPEND(), 503	NCLOBs (National Character	NOT NULL, 306, 307–308
MIN(), 117, 119	Set LOBs)	notation
minus (–), 31	compression, 529	calling procedures, 360–361
MINUS, 184, 188	defined, 477–478	positional and symbolic,
minute formatting parameters, 134	encryption, 526–528	237–238
missing values, handling, 241–243	implicit conversion, 523–524	NOT_LOGGED_ON, 354
mixed notation, 361	LOADCLOBFROMFILE(),	NTILE(), 218, 223–224
MOD(), 99, 101 MODEL	497–498	null values controlling ranking, 219
all cells access using ANY and	nested subqueries, 168, 179–180 nested tables	database table information
IS ANY, 238–239	ANSI support, 464–472	retrieval, 35–37
cell access using FOR loop,	comparing contents with map	handling, 241–243,
240–241	method, 440–442	550–552
cell range access using	converting with CAST(),	NOT IN and, 42
BETWEEN and AND, 238	443–444	specifying for columns, 253
current dimension values using	defined, 428	NULLS FIRST, 219
CURRENTV(), 239–240	information retrieval, 432-433	NULLS LAST, 219
defined, 236–237	manipulating, 446–448	NUMBER(), 13
null and missing values	modifying elements, 439–440	NUMBER, 23
handling, 241–243	multilevel collections, 458-461	numbers
positional and symbolic	populating with elements,	handling, 549–550
notation, 237–238	434–435	rows, 30–31
updating existing cells, 243–244	retrieving elements from, 436	numeric columns, 305
modes, 358	storage table tablespace,	numeric functions, 90, 98–102
modifying elements, 438–440	463–464	NUMTODSINTERVAL(), 164–165
MON, 130	types, 429–430	NUMTOYMINTERVAL(), 165
months	using TABLE() with, 438	NVL(), 36, 91, 96
datetime functions, 142–144	NEXT(), 449, 455–456	NVL2(), 91, 96
default format, 130	NEXT_DAY(), 143–145	O
formatting parameters, 133–134	nextval, 317–319, 376	
INTERVAL YEAR TO MONTH, 159–162	NLS_DATE_FORMAT, 139–140	object identifiers (OIDs), 390–392
	NOCACHE, 319	object privileges
MONTHS_BETWEEN(), 143–144 moving average, 227–228	NO_DATA_FOUND, 354 nodes	defined, 281–287
multilevel collections, 458–461	defined, 197	granting to roles, 292–293
multiple aggregate functions,	eliminating, 201–202	revoking, 288
247–248	other than root, 200	synonyms, 287–288
multiple-column subqueries, 168, 175	non-equijoins, 50–51	objects, database. See database
multiple queries	non-identical SQL statements, 588	objects
vs. CASE expressions, 583	nonrepeatable reads, 268	OCI (Oracle Call Interface)
vs. table joins, 581–582	normalized time zones, 153–154	drivers, 536
multiple-row subqueries, 168,	NOT	offsets, time zone, 149
173–175	defined, 40, 43	OIDs (object identifiers),
multiple set operators, 188-189	vs. NOT EXISTS in correlated	390–392
MULTISET, 467–469	subqueries, 178	ON, 56–57
	not a number (NAN), 24	ONLY, 410–411
N	not equal (!=), 38-39	OPEN(), 484, 498–499
15 17	not equal (<>). See <> (not equal)	OPEN-FOR, 350–351
name, 15, 17	NOT EXISTS, 176–179	opening cursors, 346

NOT IN, 465-466

operands, 31

operators Oracle Enterprise Manager passwords, 11-12, 277-278 ANSI nested table support, Diagnostics Pack, 600 PATH environment variable, 534 464-471 Oracle introduction, 1-26 percent (%), 40-41 arithmetic, 31-34 BINARY FLOAT and BINARY PERCENT RANK(), 218, 223 DOUBLE types, 23-25 performance speed, 23 combining set operators, 188-190 DDL statements used to create Perl-influenced metacharacters, 112 comparing object values, store schema, 11-20 phantom reads, 267-268 392-394 PL/SQL, 25-26 phone, 15 correlated subqueries, 176 quitting SQL*Plus, 25 **PIVOT** relational databases, 2-3 defined, 244 joins. See joins logical, 43 rows, adding/modifying/ example, 244-246 multiple-row subqueries, 173 removing, 20-23 multiple aggregate functions in, precedence, 44 SOL, 3-4 247-248 set operators, 184–188 SOL Developer, 7-10 on multiple columns, 246–247 single-row subqueries, 169-170 SQL*Plus, 4-7 PL/SQL (Procedural Language/SQL), SOL, 39-43 store schema, 10-11 339-377 value comparison, 37-39 block structure, 340-342 summary, 26 optimizers ORACLE_HOME environment in collections, 444-448 hints, 598-600 variable, 533-534 conditional logic, 342-343 query costs comparison, oracle.jdbc package, 571-575 COUNT(), 448-451 591-598 oracle.sql package, 568-571 cursor steps, 346-348 OR, 43, 55, 65-66 ORDER BY, 44-45, 172-173 cursors, 345-346 Oracle OTHERS exceptions, 357 cursors and FOR loops, connecting to database with OUT, 358 349-350 data source, 539-542 outer joins database objects in, 394-395 IDBC. See IDBC defined, 50, 51-55 DELETE(), 451-452 two-digit year interpretation, SQL/92 syntax, 59-60 delete product(), 400-401 140-142 output, combining column, 35 display product(), types, 544-545, 635-638 output lines, 75–76 396-397 XML and. See XML (Extensible overloading methods, 420 DUP_VAL_ON_INDEX overriding methods, 422-423 Markup Language) exceptions, 356 Oracle Call Interface (OCI) exceptions, 353-355 P drivers, 536 EXISTS(), 452-453 Oracle database 10g EXTEND(), 453 ANSI nested table support, packages FIRST(), 454 464-472 defined, 365-368 functions, 363-365 associative arrays, 462 get product(), 398-399 importing JDBC, 536-537 collections enhancements. page size, setting, 72–73 get product ref(),400 461-462 parameters get products(), 395-396 cost-based optimization, 591 converting datetime using TO_ insert product(), 397 element type sizing, 463 CHAR(), 133-135 introduction, 25-26 increasing elements in procedures, 358-359 INVALID_NUMBER exceptions, varray, 463 TO_CHAR(), 105-106 356-357 LOB enhancements, 523-525 user-defined constructors, LAST(), 454-455 nested table's storage table 418-422 LOBs in, 482-484. See also parent nodes, 197 tablespace, 463-464 LOBs (large objects) varrays in temporary tables, 463 parent tables, 257-258 loops, 343-345 Oracle database 11g parentheses (), 34, 168 new Oracle Database 11g generalized invocation feature, features, 374-377 parsing 423-425 defined, 562 NEXT(), 455-456 LOB enhancements, 525-529 XMLPARSE(), 613 OPEN-FOR statement, 350-351 new PL/SQL features, 374-377 Oracle data types, 638 PARTITION BY, 220-221

PL/SQL (continued) OTHERS exceptions, 357 packages, 365–368 PRIOR(), 456–457 procedures, 358–362 product_cursor.sql example, 348–349 product_lifecycle(), 401–402 product_lifecycle2(), 402–403 storing returned values	precision, 13 predefined exceptions, 354–355 prepared SQL statements, 562–564 PRESENTNNV(), 242–243 PRESENTV(), 242 price, 17 primary keys constraints enforcement, 257 defined, 306 populating with sequences, 319 printing bind variables, 590 PRIOR(), 449, 456–457	searching CLOB data, 510–511 temporary CLOBs, 509–510 triggers, 369–374 update_product(), 399 update_product_ price(), 398 using objects in PL/SQL, 394–395 writing to CLOB, 506 processing datetime. See datetime processing instructions, 614 product_cursor.sql example,
from function with bind variable, 590 subtype in place of supertype,	private items, 365 privileges auditing, 295	348-349 product_id, 16, 17-18 product_lifecycle(),
406–407 summary, 377 TREAT() in, 414–416	checking role, 291–292 granting object to roles, 292–293	401-402 product_lifecycle2(), 402-403
triggers, 369–374 TRIM(), 457–458 unconstrained cursors,	object, 281–287 query flashbacks, 270 revoking object, 288	products table, 15-16 product_type_id, 15-17 product_types table, 15-16
352-353 update_product(),399 update_product_ price(),398	summary, 297–298 synonyms, 287–288 user, 278–281 views, 327	PROGRAM_ERROR, 354 programming languages PL/SQL. See PL/SQL (Procedural Language/SQL)
user-defined constructors, 418–422 variables and types, 342	Procedural Language/SQL (PL/SQL). See PL/SQL (Procedural Language/SQL)	SQL*Plus. See SQL*Plus XML. See XML (Extensible Markup Language)
writing XML data to file, 616–618 ZERO_DIVIDE exceptions,	procedures appending data to CLOB, 506–507	pseudo columns defined, 30 LEVEL, 198–199
355–356 plan tables, 592–593 plans, execution, 562, 591–598	calling in packages, 367–368 comparing data in CLOBs, 507–508	sequences, 317–319 public synonyms, 287–288 purchases table, 17–18
plus (+). See + (plus) pointers, 482 populating collections, 434–435	copying data from CLOBs, 508–509 copying LOB data. <i>See</i> copying	Q
populating LOBs BFILEs, 482 CLOBs and BLOBs, 479	LOB data defined, 358–362 delete_product(), 400–401	quantity, 18 quarter formatting parameters, 133 queries
populating primary keys, 319 populating ResultSet object, 543–544	display_product(), 396-397 erasing CLOB data, 510	costs comparison, 591–598 defined, 28 flashbacks, 270–273
positional notation advanced queries, 237–238 calling procedures, 360–361 pound (#), 107	<pre>insert_product(), 397 product_lifecycle(), 401-402 product lifecycle2(),</pre>	subqueries. <i>See</i> subqueries tuning SQL. <i>See</i> SQL (Structured Query Language) tuning
POWER(), 99, 101 POWERMULTISET(), 472 POWERMULTISET_BY_	402–403 reading data from CLOBs and BLOBs, 504–506	using JDBC, 543–547 XMLQUERY(), 618–622 queries, advanced, 183–250 all cells access using ANY and
CARDINALITY(), 472	retrieving LOB locators, 503-504	IS ANY, 238–239

retrieving LOB locators, 503–504

analytic functions, 216-217 summary, 249-250 reports, 79-85 result set CASE expressions, 193-196 TRANSLATE(), 190-191 cell access using FOR loop, defined, 28 traversing upward through 240-241 tree, 201 retrieving rows using JDBC, cell range access using UNPIVOT, 248-249 543-547 BETWEEN and AND, 238 updating existing cells, results, hierarchical query, 199-200 CUBE, 207-208 243-244 retrieving information. See current dimension values using window functions, 225-230 information retrieval CURRENTV(), 239-240 query statements, SQL, 3 RETURNING, 255-256 DECODE(), 191-193 question mark (?), 564 REVOKE, 4 eliminating nodes and branches, quitting SQL *Plus, 25 revoking privileges 201-202 quotes ("), 34, 254 object, 288 FIRST and LAST functions, 233 from users, 281 revoking roles, 294 formatting hierarchical results, R 199-200 right outer joins, 52-54, 59-60 GROUP BY, 203-204 roles range, 23 GROUP BY column use, 214 checking privileges, rank, hypothetical functions, 235 GROUP_ID(), 215-216 291-292 GROUPING(), 209-211 ranking functions, 216, 217–224 creating, 289 GROUPING SETS, 211-212 RATIO_TO_REPORT(), 231–232 default, 294 GROUPING_ID(), 212-214 RAW types, 521 dropping, 295 hierarchical, 196-199 READ(), 484, 499-500, granted privilege use, 293 504-506 hierarchical, including other granted to users, 290-291 conditions, 202-203 READ COMMITTED, 268 granting object privileges, READ ONLY, 306, 330 hypothetical rank and 292-293 distribution functions, 235 **READ UNCOMMITTED, 268** granting privileges, 290 REF(), 390-392 inverse percentile functions, revoking, 294 REF CURSOR, 351, 590-591 224-225 summary, 297-298 references, object, 383, **ROLLBACK** LAG() and LEAD(), 232–233 linear regression functions, 390-392 defined, 4 233 - 235REGEXP_COUNT(), 114, 116 row removal, 22-23 MODEL, 236-237 REGEXP_INSTR(), 113, 115 transactions, 262-263 REGEXP_LIKE(), 113, 114-115 multiple aggregate functions in **ROLLUP** pivot, 247-248 REGEXP_REPLACE(), 114, 116 defined, 205-207 REGEXP_SUBSTR(), 114, 116 null and missing values GROUPING(), 209-210 regular expression functions, 90, handling, 241-243 using with analytic functions, PIVOT and UNPIVOT, 244 109-116 221-223 PIVOT example, 244–246 regular expressions Romeo and Juliet functions, 113-116 pivoting on multiple columns, (Shakespeare), 110 metacharacters, 109-112 root nodes 246-247 positional and symbolic relational data, generating XML. See defined, 197 XML (Extensible Markup notation, 237-238 starting at node other than, 200 ranking functions, 217-224 Language) rotating. See PIVOT relational databases, 2-3 reporting functions, 230-232 ROUND() ROLLUP, 205-207 removing rows, 256 datetime functions, set operators, 184-188 RENAME, 3 143, 145 set operators, combining, renaming tables, 313 defined, 99, 101 188-190 REPEATABLE READ, 268 row-level triggers, 369 START WITH subquery, 200-201 REPLACE(), 91, 96-97 ROWID, 30 starting at node other than reporting functions, 216, ROWNUM, 30-31 230-232 root, 200 ROW_NUMBER(), 218, 224

rows	example file, 623	server-side internal driver, 536
adding/modifying/removing,	example schema, 623–625	server-side Thin driver, 536
20–23	example schema information	sessions
adding using JDBC, 547–548	retrieval, 625–630	defined, 139
changing table contents. See	example schema information	time zones, 147–149
table contents, changing	updating, 630-633	SET(), 470
defined, 2	scalar subqueries, 170	SET LINESIZE, 73
deleting using JDBC, 548	scale, 13	set methods, 540–542
displaying distinct, 37	schemas, 2	set operators, 184–190
filtering with WHERE, 580–581	schemas, store	SET PAGESIZE, 72–73
getting first and last, 229–230	DDL statements, 11–20	SETS, GROUPING, 211–212
grouping, 120–127	introduction, 10–11	Shakespeare, William, 110, 477
identifiers, 30	schemas, XML	SIGN(), 99, 102
modifying using JDBC, 548	creating, 623–625	simple CASE expressions,
multiple-row subqueries,	information retrieval, 625–630	193–194
173–175	information updating, 630–633	simple loops, 343–344
numbers, 30–31	SCNs (system change numbers),	simple views, 327
pivoting. See PIVOT	272–273	SIMPLE_INTEGER type, 375
removing duplicate with	scripts, 79–85	SIN(), 99
GROUP_ID(), 215–216	searched CASE expressions,	single-row functions
retrieving specific using	194–196	character, 90–98
WHERE, 29	searching CLOB data, 510–511	conversion, 102–109
retrieving using JDBC, 543–547	seconds	numeric, 98–102
retrieving with cursors. See	formatting parameters, 134	regular expression, 109–116
cursors	INTERVAL DAY TO SECOND,	types of, 90
single-row subqueries, 168–173	162–164	single-row subqueries, 168–173
sorting, 44–45	SELECT	SINH(), 99
storing with bind variable,	adding rows, 21	sizing columns, 305
590–591	performing using SQL*Plus, 6–7	sizing element types, 463
using TABLE() to treat collection	single table information	SOME, 38–39
as series of, 436–438	retrieval, 28–29	sorting rows, 44–45
ROWTYPE_MISMATCH, 354	using multiple tables, 49–50	SOUNDEX(), 91, 97
RPAD(), 91, 95	using two tables, 45–47	special values, 24–25
RR format, 140–142	SELF, 419	specifications, package,
RTRIM(), 91, 95	self joins, 50, 55–56, 60	365–366, 395
rule-based optimization, 591	SELF_IS_NULL, 355	SPO[OL], 67
RULES UPDATE, 243–244	semicolon (;), 6, 341	SQL/92
R[UN], 65	separators, 134	CASE expressions, 193
run-time checks, 412–416	sequences, 314–320	defined, 47
run-time errors handling. See	creating, 314–316	joins using, 56–61
exceptions	dropping, 320	SQL Access Advisor, 601
	modifying, 320	SQL Developer, 7–10
S	new Oracle Database 11 <i>g</i>	SQL (Structured Query Language)
	features, 375–377	introduction, 3–4
salary_grades table, 19–20	populating primary keys, 319	LOBs in, 478–482
SAVEPOINT, 4	retrieving information, 316–317	operator precedence, 44
savepoints, 264–265	summary, 337	operators, 39–43
saving SQL*Plus files, 66–70	using, 317–319	Oracle data types, 636–637
saving XML	XMLSEQUENCE(), 615 SERIALIZABLE transactions,	Procedural Language. <i>See</i> PL/SQL (Procedural
in database, 622	268–269	
•	200-203	Language/SQL)

subtype in place of supertype, GROUPING SETS vs. statistics in cost-based 405-406 CUBE, 588 optimization, 597 SQL (Structured Query Language), introduction, 580 STDDEV(), 117, 120 running using Java, 531-578 joins with qualified column storage adding rows, 547-548 references, 582 datetime. See datetime BasicExample1.java, passing hints to optimizer, defined, 23 556-562 598-600 numbers in Java, query costs comparison, BasicExample2.java, 549-550 565-567 591-598 returned values with bind BasicExample3.java, summary, 601 variable, 590 575-577 table joins vs. multiple queries, rows with bind variable. closing JDBC objects, 554-556 581-582 590-591 configuring computer, 533-535 UNION ALL vs. UNION. tables, 463-464 database connection, 537-542 STORAGE ERROR, 355 585-586 database transaction control. WHERE vs. HAVING, store schemas 552-553 584-585 DDL statements, 11-20 DDL statements, 553 SQL Tuning Advisor, 600-601 introduction, 10-11 deleting rows, 548 SQL*Plus, 63-87 stored subprograms, 363 getting started, 532 automatically generated strings handling database null values, statements, 86 converting datetime using 550-552 BINARY_FLOAT and BINARY_ TO_CHAR(), 132-136 handling exceptions, 553-554 DOUBLE types, 23-25 converting datetime using handling numbers, 549-550 column format clearing, 73 TO_DATE(), 136-139 importing JDBC packages, column formatting, 70-72 converting to TIMESTAMP 536-537 DDL statements used to create WITH LOCAL TIME store schema, 11-20 JDBC Statement object, **ZONE, 159** 542-543 disconnecting/exiting, 86 XMLSERIALIZE(), 616 modifying rows, 548 editing SQL statements, structure, block, 340-342 Oracle JDBC drivers, 535-536 65-66 Structured Query Language (SQL). Oracle JDBC extensions, files, saving/retrieving/running, See SQL (Structured Query 567-568 66-70 Language) SUBMULTISET, 442, 466-467 getting help, 85-86 oracle.jdbc package, 571-575 information retrieval. See subqueries, 167-181 database table information correlated, 175-179 oracle.sql package, retrieval multiple-column, 175 568-571 prepared statements, 562-564 introduction, 4-7 multiple-row, 173-175 registering Oracle JDBC line size, setting, 73 nested, 179-180 drivers, 537 page size, setting, 72-73 single-row, 168-173 retrieving rows, 543-547 quitting, 25 START WITH, 200-201 summary, 578 rows, adding/modifying/ summary, 181 SQL (Structured Query Language) removing, 20-23 types of, 168 tuning, 579-601 simple reporting, 79-85 UPDATE, 255 adding indexes to tables, 584 store schema, 10-11 **UPDATE** and **DELETE** statements additional tools, 600-601 summary, 87 containing, 180-181 bind variables, 588-591 table structure, 64-65 SUBSCRIPT_BEYOND_ CASE expressions vs. multiple variables, 74-79 COUNT, 355 gueries, 583 SUBSCRIPT_OUTSIDE_LIMIT, 355 SQRT(), 99, 102 substitution variables, 74-79 EXISTS vs. DISTINCT, 587 START WITH, 198, 200-201 EXISTS vs. IN, 586-587 starting transactions, 263 SUBSTR(), 91, 97-98, 484, filtering rows with WHERE, statement-level trigger, 369 500-501 580-581 Statement object, 542-543 subtotal computing, 83-85

subtypes	table joins. See also joins	temporary tables, 463		
defined, 404	defined, 46	temporary variables, 74-77, 80		
generalized invocation,	execution plans involving,	THEN, 342–343		
423–425	596–597	Thin driver, 535–536		
IS OF(), 408–412	vs. multiple queries,	time. See datetime		
overriding methods, 422–423	581–582	time query flashbacks, 270-272		
using in place of supertype,	tables, 300–314	time zones, 135, 146–150,		
405–408	adding comments, 313–314	152–154		
SUM(), 117, 120	adding indexes to, 584	TIMEOUT_ON_RESOURCE, 355		
sums, cumulative, 225–227	altering, 302–306	timestamps		
sums, reporting on, 230–231	BINARY_FLOAT and BINARY_	flashback data archives,		
	DOUBLE, 24–25	334–337		
supertypes defined, 404				
	column objects, 383–386	functions, 154–159		
generalized invocation,	comparing object values,	using, 150–154		
423–425	392–394	TO_CHAR()		
NOT INSTANTIABLE object	constraints, 306–312	combining with TO_DATE(),		
types, 416–418	containing LOBs, 478	138–139		
overriding methods, 422–423	creating, 300–302	converting datetime, 131–136		
subtype, using in place of,	defining column with collection	defined, 104–107		
405–408	type, 430	TO_DATE(), 32, 131–132, 136–139		
symbolic notation, 237–238	dropping, 314	TO_NUMBER(), 107–108		
synonyms, privileges, 287–288	examining, 12–20	tools, SQL tuning. See SQL		
SYSDATE, 6–7, 143, 145	getting information on,	(Structured Query Language)		
SYS_EXTRACT_UTC(), 155,	302–303	tuning		
157–158	information retrieval. See	TOO_MANY_ROWS, 355		
SYS_INVALID_ROWID, 355	database table information	TO_TIMESTAMP(), 155, 158		
system change numbers (SCNs),	retrieval	TO_TIMESTAMP_TZ(), 155, 158		
272–273	nested. See nested tables	trail views, auditing, 296		
system privileges. See privileges	object tables, 386–389	Transaction Control (TC)		
SYSTIMESTAMP, 155–156	OIDs and object references,	statements, 4		
SYS_TYPEID(), 408, 416	390–392	transaction isolation level,		
	plan table, 592–593	267–269		
Т	renaming, 313	transaction locking, 266-267		
	SQL*Plus structure, 64–65	transactions, database, 262-269,		
TABLE() 426 420	storage table tablespace,	552–553		
TABLE(), 436–438	463–464	TRANSLATE(), 190–191		
table aliases	substituting names using	transparent rounding, 23		
defined, 47–48	variables, 76–77	TREAT(), 408, 412–416		
self joins, 55	summary, 337	tree queries. See also hierarchical		
table contents, changing, 251–273	truncating, 314	queries		
adding rows, 252–254	varrays in temporary, 463	defined, 197		
database integrity, 256–258	tablespaces, 276, 463-464	eliminating nodes and branches,		
database transactions,	TAN(), 99	201–202		
262–269	TANH(), 99	traversing upward through, 201		
default values, 258–259	TC (Transaction Control)	triggers		
merging rows, 259–262	statements, 4	: new attribute, 524-525		
modifying rows, 254–255	temporary LOBs	defined, 369-374		
query flashbacks, 270–273	CREATETEMPORARY(), 488	TRIM()		
removing rows, 256	FREETEMPORARY(), 492	defined, 91, 95, 449,		
RETURNING, 255–256	using temporary CLOBs,	457–458, 484		
summary, 273	509–510	LOB methods, 501-502		
		•		

TRUNC() datetime functions, 143, 146 defined, 99, 102 TRUNCATE, 3, 314 try/catch statements, 553–554	SYS_TYPEID(), 416 using DESCRIBE to get information, 382–383	privileges, 278–281 summary, 297–298 USING, 57–58, 261 UTC (Coordinated Universal Time), 146–147		
defined, 99, 102		UTC (Coordinated Universal Time),		
IS OF(), 408–405 IS OF(), 408–412 NOT INSTANTIABLE, 416–418 overriding methods, 422–423 subtype in place of supertype, 405–408	varrays, 431–432 users auditing, 295–296 creating, 11–12 granted roles to, 290–291 granting object privileges, 282 overview, 276–278	defined, 428 increasing elements in, 463 information retrieval, 431–432 manipulating, 444–446 modifying elements, 438–439 multilevel collections, 458–461 populating with elements, 434		

656

Oracle Database 11g SQL

varrays (continued) retrieving elements from, 435 in temporary tables, 463 types, 429-430 using TABLE() with, 437–438 views, 325-334 creating/using, 326-333 defined, 325-326 dropping, 334 errors in procedures, 362 modifying, 333-334 nested table information retrieval, 433 summary, 337 table structure, 64-65 varray information retrieval, 431-432 virtual columns, 304-305

W

wallets, 525–526 week formatting parameters, 134 WHEN NOT MATCHED, 261 WHERE

filtering rows with, 580–581 grouping rows, 126–127 vs. HAVING, 584–585 hierarchical queries, 202–203 retrieving specific rows, 29

row modification, 22 subqueries in, 168–169 using SQL operators, 39–43 WHILE loops, 343, 344 window functions, 216, 225–230 Windows XP, ORACLE_HOME environment variable, 533 WOR[D_WRAPPED], 70–71 WRA[PPED], 70–71 wrapper classes, 551–552 WRITE(), 484, 502, 506 WRITEAPPEND(), 484, 503

X

x, 65–66
XML (Extensible Markup Language), 603–633
example file, 623
example schema, 623–625
example schema information
retrieval, 625–630
example schema information
updating, 630–633
generating from relational
data, 604
introduction, 604
PL/SQL example, 616–618
saving in database, 622
summary, 633

XMLAGG(), 609–612 XMLATTRIBUTES(), 608 XMLCOLATTVAL(), 612 XMLCOMMENT(), 614 XMLCONCAT(), 613 XMLELEMENT(), 605–608 XMLFOREST(), 608–609 XMLPARSE(), 613 XMLPI(), 614 XMLQUERY(), 618–622 XMLSEQUENCE(), 615 XMLSERIALIZE(), 616 XOuery expression, 618–622

Y

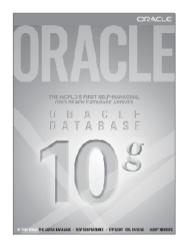
years
datetime default format, 130
datetime formatting
parameters, 133
INTERVAL YEAR TO MONTH,
159–162
two-digit year interpretation,
140–142
YY format, 140
YYYY, 130

Z

ZERO_DIVIDE exceptions, 355-356

FREE SUBSCRIPTION TO ORACLE MAGAZINE

Oracle Magazine is essential gear for today's information technology professionals. Stay informed and increase your productivity with every issue of Oracle Magazine. Inside each free bimonthly issue you'll get:



IF THERE ARE OTHER ORACLE USERS AT YOUR LOCATION WHO WOULD LIKE TO RECEIVE THEIR OWN SUBSCRIPTION TO ORACLE MAGAZINE, PLEASE PHOTOCOPY THIS FORM AND PASS IT ALONG.



- Up-to-date information on Oracle Database, Oracle Application Server,
 Web development, enterprise grid computing, database technology,
 and business trends
- · Third-party vendor news and announcements
- Technical articles on Oracle and partner products, technologies, and operating environments
- · Development and administration tips
- Real-world customer stories

Three easy ways to subscribe:

① Web

Visit our Web site at otn.oracle.com/oraclemagazine. You'll find a subscription form there, plus much more!

② Fax

Complete the questionnaire on the back of this card and fax the questionnaire side only to +1.847.763.9638.

③ Maîl

Complete the questionnaire on the back of this card and mail it to P.O. Box 1263, Skokie, IL 60076-8263



FREE SUBSCRIPTION

0	Yes, please send me a FREE To receive a free subscriptio (incomplete cards cannot be Or subscribe at our Web site	n to <i>Ora</i> process	<i>acle Magazine</i> , you sed or acknowledg	must fill out the entire (jed). You can also fax yo			
0	From time to time, Oracle Publishing		signature (requ	signature (required) date			
our partners exclusive access to our e-mail addresses for special promotions and announcements. To be included in this pro- gram, please check this circle.			X				
_			n a m e t i		title		
Oracle Publishing allows sharing of our mailing list with selected third parties. If you		s. If you					
	prefer your mailing address not included in this program, please che		company			e-mail address	
	If at any time you would like to be refrom this mailing list, please	emoved	street/p.o. box				
	Customer Service at +1.847.647.9630		city/state/zip or postal code		t e l e p h o n e		
	an e-mail to oracle@halldata.com.		country		fax	ах	
	YOU MUST ANSWER ALL TEN Q	UESTION	IS BELOW.				
	ACTIVITY OF Y OUR FIRM AT THIS LOCATION? (check one only) 101 Aerospace and Defense Manufacturing 102 Application Service Provider 103 Automotive Manufacturing 104 Chemicals, Oil and Gas 105 Communications and Media 106 Construction/Engineering 107 Consumer Sector/Consumer Packaged Goods 108 Education 109 Financial Services/Insurance 100 Government (civil) 111 Government (inilitary) 121 Healthcare 131 High Technology Manufacturing, OEM 141 Integrated Software Vendor 15 Life Sciences (Biotech, Pharmaceuticals) 16 Mining 17 Retail/Wholesale/Distribution 18 Systems Integrator, VAR/VAD 19 Telecommunications 20 Travel and Transportation 21 Utilities (electric, gas, sanitation, water) 98 Other Business and Services WHICH OF THE FOLLOWING BEST DESCRIBES YOUR PRIMARY JOB FUNCTION? (check one only) COFO, CPO, Owner, Partner, Principal) 107 Finance/Administrative Management (VP/Director/Manager/Controller, Purchasing, Administrative Management (VP/Director/Manager/Controller, Programming Management (VP/Director/Manager MIS, Operations) 1S/IT Staff 15 Systems Development/ Programming Staff 17 Consulting 18 DBA/Systems Administrator 19 B DBA/Systems Administrator 19 B DBA/Systems Administrator 19 B DBA/Systems Administrator 10 B DBA/Systems Administrator 10 B DBA/Systems Administrator 10 B DBA/Systems Administrator	RECOMMEND PURCHASE O (check all that :	INIX ash IS IFE IR IK Computing MS WIX IK DYNIX/ptx Iaris/SunOS are WS WS IT INIX IF the above LUATE, SPECIFY,	Hardware 20 Macintosh 21 Mainframe 22 Massively Parallel Processing 23 Minicomputer 24 PC 25 Metwork Computer 26 Symmetric Multiprocessing 27 Workstation Peripherals 28 Bridges/Routers/Hubs/Gateways 29 CD-ROM Drives 30 Disk Drives/Subsystems 31 Modems 32 Tape Drives/Subsystems 33 Video Boards/Multimedia Services 34 Application Service Provider 35 Consulting 36 Education/Training 37 Maintenance 38 Online Database Services 39 Support 40 Technology-Based Training 99 None of the above 60 WHAT ORACLE PRODUCTS ARE IN AT YOUR SITE? (check all that apply) Oracle E-Business Suite 01 Oracle Marketing 02 Oracle Oracle Trifllment 04 Oracle Supply Chain Management 05 Oracle Procurement 06 Oracle Manufacturing 07 Oracle Maintenance Management 08 Oracle Services 09 Oracle Contracts 10 Oracle Manufacturing 07 Oracle Maintenance Management 08 Oracle Services 19 Oracle Services 19 Oracle Interaction Center 14 Oracle Interaction Center 14 Oracle Interaction Center 14 Oracle Oracle Oracle Interaction Center 15 Oracle Pholic Sector/University (in Cracle Financial Services (module: Server/Software 17 Oracle9) Lite	USE (8	Tools	
3	WHAT IS YOUR CURRENT PRIMARY OPERATING PLATFORM? (select all that apply) OT Digital Equipment UNIX O 20 Digital Equipment VAX VMS	□ 17 Spread	: Management fic and Engineering Isheets Is Management	☐ 19 Oracle8 / ☐ 20 Other Oracle database ☐ 21 Oracle9 / Application Server ☐ 22 Oracle9 / Application Server Wirele ☐ 23 Oracle Small Business Suite	ss	WHAT IS YOUR COMPANY'S YEARLY SALES REVENUE? (please choose one) 10 1 \$500, 000, 000 and above 10 2 \$100, 000, 000 to \$500, 000, 000 10 3 \$50, 000, 000 to \$50, 000, 000 10 4 \$5, 000, 000 to \$50, 000, 000 10 10 \$50, 000, 000	



Professional development and industry recognition are not the only benefits you gain from Oracle certifications. They also facilitate career growth, improve productivity, and enhance credibility. Hiring managers who want to distinguish among candidates for critical IT positions know that the Oracle Certification Program is one of the most highly valued benchmarks in the marketplace. Hundreds of thousands of Oracle certified technologists testify to the importance of this industry-recognized credential as the best way to get ahead—and stay there.

For details about the Oracle Certification Program, go to oracle.com/education/certification.

Oracle University — Learn technology from the source





Learn Oracle Database 11g from Oracle University

Ready to learn from the experts?
Only Oracle University offers a complete curriculum to take you from beginner through Certified Master.

Our training offerings for Oracle Database 11g include courses on identity and access management, infrastructure grid, change management, security, Real Application Clusters, SQL and PL/SQL, building JavaEE applications, Bl and data warehousing, high availability, and much more.

Oracle Database 11*g*: making information more manageable, more available, and more secure. Learn all about it from Oracle University: oracle.com/education.

Oracle University — Learn technology from the source



